

## Lecture 11:

# Data Fitting and Regression



Markus Hohle

University California, Berkeley

**Numerical Methods for  
Computational Science**

## Course Map

Week 1:	Introduction to Scientific Computing and Python Libraries
Week 2:	Linear Algebra Fundamentals
Week 3:	Vector Calculus
Week 4:	Numerical Differentiation and Integration
Week 5:	Solving Nonlinear Equations
Week 6:	Probability Theory Basics
Week 7:	Random Variables and Distributions
Week 8:	Statistics for Data Science
Week 9:	Eigenvalues and Eigenvectors
Week 10:	Simulation and Monte Carlo Method
<b>Week 11:</b>	<b>Data Fitting and Regression</b>
Week 12:	Optimization Techniques
Week 13:	Machine Learning Fundamentals



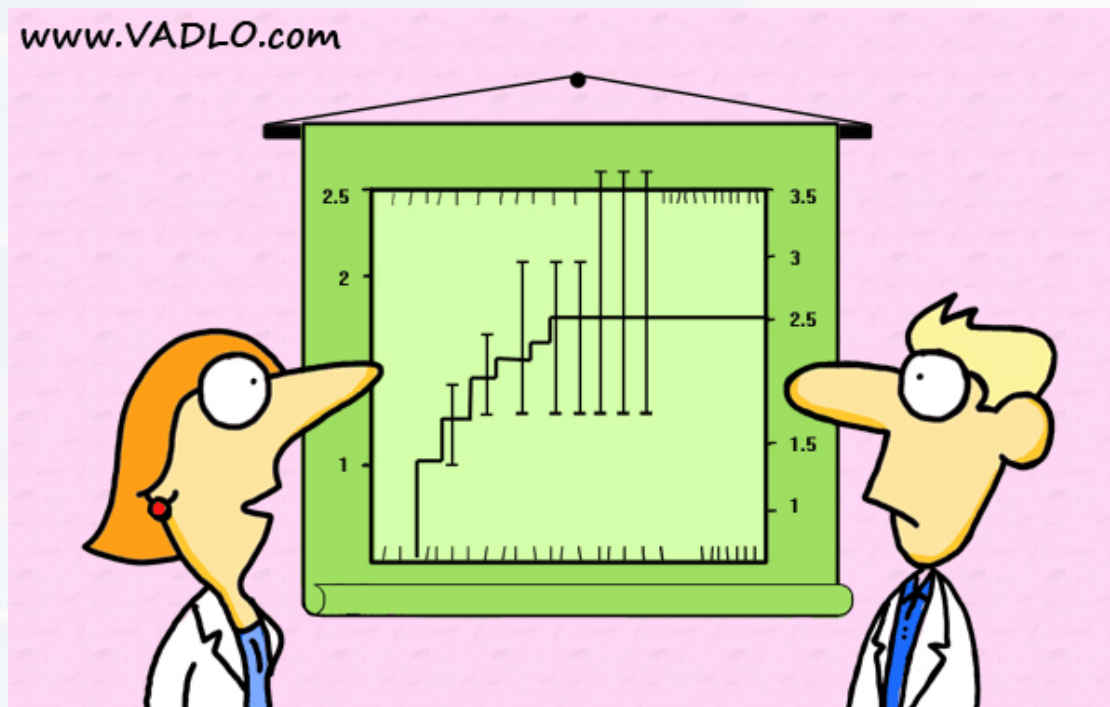
## Outline

### Linear Regression

- The Math
- What is Linear?
- Stats
- a Python Example

### Logistic Regression

### Curve Fitting



“Did you really have to show the error bars?”



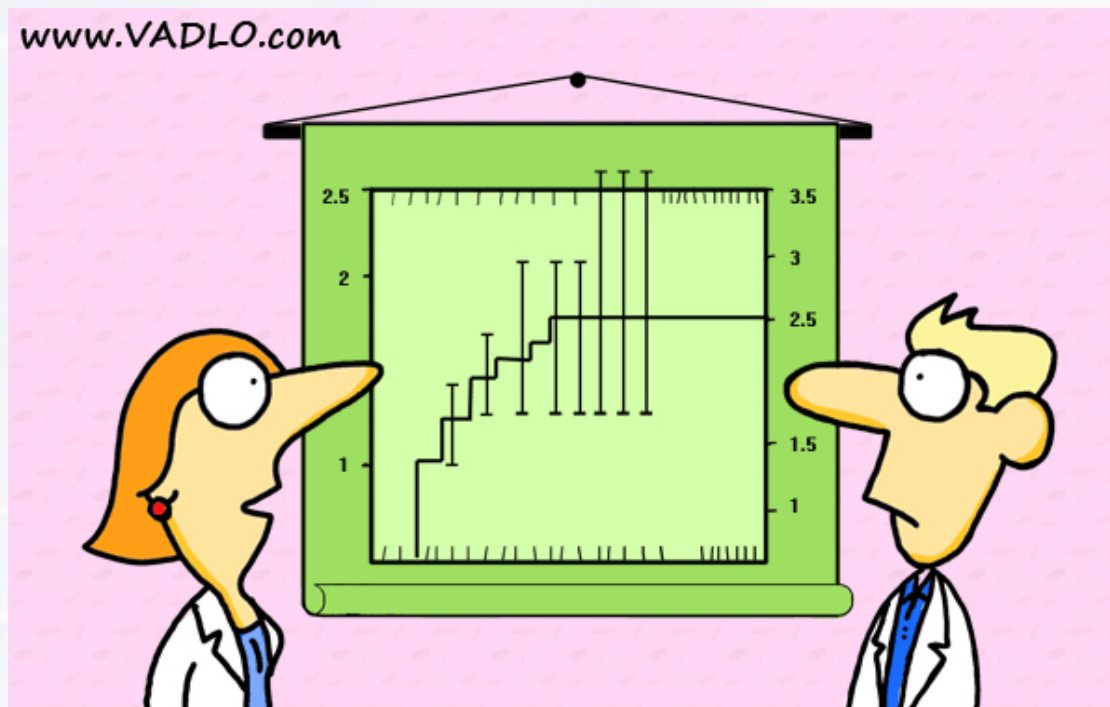
## Outline

### Linear Regression

- The Math
- What is Linear?
- Stats
- a Python Example

### Logistic Regression

### Curve Fitting



“Did you really have to show the error bars?”





**problem:** a parameter  $y$  depends on many **regressors**  $x_n$

**goal:** finding a model that tells us **how**  $y$  depends on each  $x_n$   
deriving a model to **predict**  $y$  from new data points

### The Math

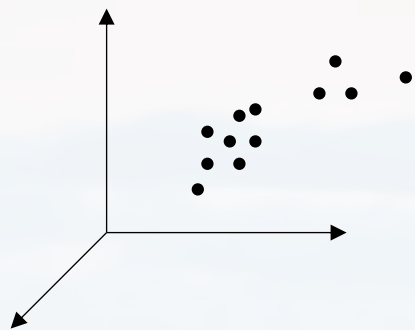
What is Linear?

Stats

a Python Example

$N$

		$x_1$	$x_2$	$x_3$	$x_4$	$x_5$	$y_k$
	Index	molecular_weight	electronegativity	bond_lengths	num_hydrogen_bonds	logP	toxicity_score
$k^{th}$ data point	0	341.704	2.65585	3.09407	2	9.11147	80.9281
	1	335.951	3.22262	2.89039	7	8.92848	83.4911
	2	235.203	2.44115	2.48203	1	6.49731	61.8406
	3	246.505	2.76656	2.71547	7	7.45089	57.0538
	4	437.939	3.4801	3.59569	3	10.9156	131.326



idea: data point  $y_k$  in  $N$  dimensional space

$$\rightarrow y_k = f(x_1, \dots, x_n, \dots, x_N) + \epsilon \quad \text{for each data point } k$$

### The Math

What is Linear?

Stats

a Python Example

ansatz:

$$y_k = \beta_0 + \sum_{n=1}^N \beta_n x_n + \epsilon$$

**linear** combination

$y$ : response

$x$ : regressors (assumed to be independent)

$\beta$ : factors (how a regressor contributes to the response)

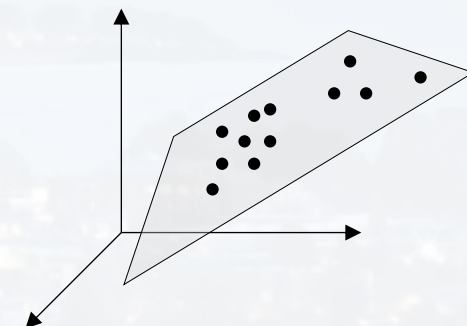
$\beta_0$ : intercept

$\epsilon$ : error (stochasticity of the data, assumed to be normally dist.)

$N$

$x_1 \quad x_2 \quad x_3 \quad x_4 \quad x_5 \quad y_k$

	Index	molecular_weight	electronegativity	bond_lengths	num_hydrogen_bonds	logP	toxicity_score
$k^{\text{th}}$ data point	0	341.704	2.65585	3.09407	2	9.11147	80.9281
	1	335.951	3.22262	2.89039	7	8.92848	83.4911
	2	235.203	2.44115	2.48203	1	6.49731	61.8406
	3	246.505	2.76656	2.71547	7	7.45089	57.0538
	4	437.939	3.4801	3.59569	3	10.9156	131.326





linear  $\neq$  not curved

$$y_k = \beta_0 + \sum_{n=1}^N \beta_n x_n^n + \epsilon$$

...is still linear

just define:  $\bar{x}_n := x_n^n$

$$y_k = \beta_1 x_n^{\beta_2}$$

...is still linear

$$\begin{aligned} \text{just use log: } \bar{y}_k &= \log(y_k) = \log(\beta_1) + \beta_2 \log(x_n) \\ &= \bar{\beta}_1 + \beta_2 \bar{x}_n \end{aligned}$$

As long as we can recover the linear structure by any transformation  $\rightarrow$  it is linear

in part. log scaling is quite common examples:

- log fold change (DESeq/RNASeq)
- log odds ratio (comparing models, HMM)
- sound  $\rightarrow$  dB is a log unit
- log incidence rates (medical studies)
- percentiles (medical studies)
- .....

The Math

**What is Linear?**

Stats

a Python Example

y:	response
x:	regressors
$\beta$ :	factors
$\beta_0$ :	intercept
$\epsilon$ :	error



...what is **not** linear?

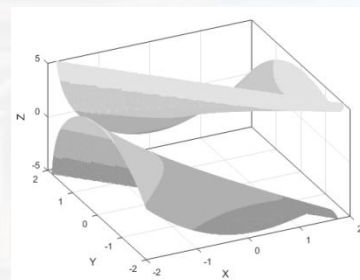
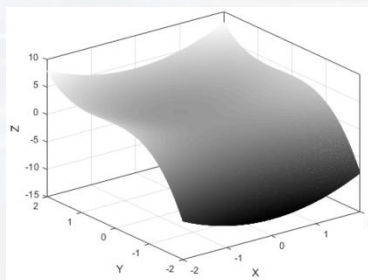
$$y_k = \beta_0 + \beta_1 x_n^{\beta_2} \quad \text{log trick does not work here}$$

general: linear refers to the **factors**

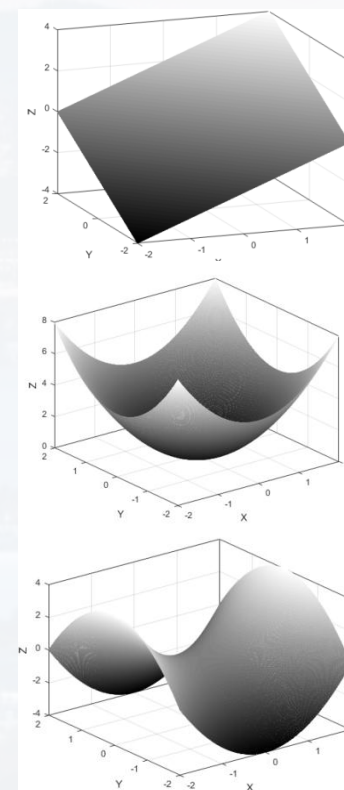
$$y_k = \beta_0 + \beta_1 x_1 + \beta_2 x_2 \quad 2D \text{ plane in } 3D \text{ space}$$

$$y_k = \beta_0 + \beta_1 x_1^2 + \beta_2 x_2^2 \quad 2D \text{ parabolic}$$

$$y_k = \beta_0 + \beta_1 x_1^2 - \beta_2 x_2^2 \quad 2D \text{ hyperbolic}$$



...and many more...



The Math

**What is Linear?**

Stats

a Python Example

y:	response
x:	regressors
$\beta$ :	factors
$\beta_0$ :	intercept
$\epsilon$ :	error

all linear

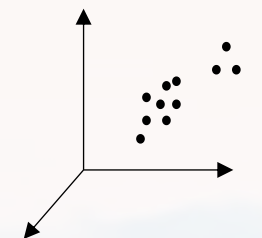
$$y_k = \beta_0 + \sum_{n=1}^N \beta_n x_n + \epsilon$$





for  $K$  data points in  $N$  dimensional space

$$y_k = \beta_0 + \sum_{n=1}^N \beta_n x_n + \epsilon$$



The Math  
What is Linear?  
**Stats**  
a Python Example

$$\underbrace{\begin{pmatrix} y_1 \\ \vdots \\ y_k \\ \vdots \\ y_K \end{pmatrix}}_Y = \underbrace{\begin{pmatrix} 1 & x_{11} & x_{12} & \dots & x_{1n} & \dots & x_{1N} \\ \vdots & \vdots & \vdots & & \vdots & & \vdots \\ 1 & x_{k1} & & & x_{kn} & & \\ \vdots & \vdots & & & \vdots & & \vdots \\ 1 & \vdots & & & \vdots & & \vdots \\ 1 & x_{K1} & x_{K2} & \dots & x_{Kn} & \dots & x_{KN} \end{pmatrix}}_X \underbrace{\begin{pmatrix} \beta_0 \\ \beta_1 \\ \vdots \\ \beta_n \\ \vdots \\ \beta_N \end{pmatrix}}_{\beta} + \underbrace{\begin{pmatrix} \epsilon_1 \\ \epsilon_2 \\ \vdots \\ \epsilon_k \\ \vdots \\ \epsilon_K \end{pmatrix}}_{\epsilon}$$

**y:** response  
**x:** regressors  
 **$\beta$ :** factors  
 **$\beta_0$ :** intercept  
 **$\epsilon$ :** error

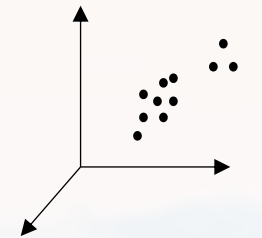
$$Y = X\beta + \epsilon$$

$k^{th}$ data point ↓	$N$						
		$x_1$	$x_2$	$x_3$	$x_4$	$x_5$	$y_k$
	Index	molecular_weight	electronegativity	bond_lengths	num_hydrogen_bonds	logP	toxicity_score
	0	341.704	2.65585	3.09407	2	9.11147	80.9281
	1	335.951	3.22262	2.89039	7	8.92848	83.4911
	2	235.203	2.44115	2.48203	1	6.49731	61.8406
	3	246.505	2.76656	2.71547	7	7.45089	57.0538
4	437.939	3.4801	3.59569	3	10.9156	131.326	



for  $K$  data points in  $N$  dimensional space

$$y_k = \beta_0 + \sum_{n=1}^N \beta_n x_n + \epsilon$$



The Math  
What is Linear?  
**Stats**  
a Python Example

$$\underbrace{\begin{pmatrix} y_1 \\ \vdots \\ y_k \\ \vdots \\ y_K \end{pmatrix}}_Y = \underbrace{\begin{pmatrix} 1 & x_{11} & x_{12} & \dots & x_{1n} & \dots & x_{1N} \\ \vdots & \vdots & \vdots & & \vdots & & \vdots \\ 1 & x_{k1} & & & x_{kn} & & \\ \vdots & \vdots & & & \vdots & & \vdots \\ 1 & \vdots & & & \vdots & & \vdots \\ 1 & x_{K1} & x_{K2} & \dots & x_{Kn} & \dots & x_{KN} \end{pmatrix}}_X \underbrace{\begin{pmatrix} \beta_0 \\ \beta_1 \\ \vdots \\ \beta_n \\ \vdots \\ \beta_N \end{pmatrix}}_{\beta} + \underbrace{\begin{pmatrix} \epsilon_1 \\ \epsilon_2 \\ \vdots \\ \epsilon_k \\ \vdots \\ \epsilon_K \end{pmatrix}}_{\epsilon}$$

$$Y = X\beta + \epsilon$$

<b>y:</b>	<b>response</b>
<b>x:</b>	<b>regressors</b>
<b><math>\beta</math>:</b>	<b>factors</b>
<b><math>\beta_0</math>:</b>	<b>intercept</b>
<b><math>\epsilon</math>:</b>	<b>error</b>

fitting: finding the best  $\beta$  in terms of minimizing the errors

$$(Y - X\beta)^T(Y - X\beta) = \sum_k \epsilon_k^2$$

the model

$$\frac{\partial}{\partial \beta} \sum_k \epsilon_k^2 = 0 \longrightarrow \beta_{best} = \hat{\beta} = (X^T X)^{-1} X^T Y \longrightarrow \hat{Y} = X\hat{\beta} = X(X^T X)^{-1} X^T Y$$

X and Y are all observables



$$Y = X\beta + \varepsilon$$

finding the best  $\beta$  in terms of minimizing the errors

$$(Y - X\beta)^T(Y - X\beta) = \sum_k \varepsilon_k^2$$

the model

$$\hat{Y} = X\hat{\beta} = X \underbrace{(X^T X)^{-1} X^T}_{\text{hat matrix } H} Y$$

hat matrix  $H$

some properties of the hat matrix:

- $H = H^T$  (symmetry)
- $HH = H \rightarrow H^n = H$  (idempotency)

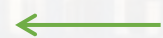
evaluating the fit:

$$\hat{\varepsilon} = Y - X\hat{\beta} = Y - \hat{Y} = (I - H)Y$$

$$\hat{\varepsilon}^T \hat{\varepsilon} = [(I - H)Y]^T (I - H)Y = Y^T (I - H)^T (I - H)Y = Y^T (I - H)Y$$

sum of squared errors (SSE)

$$\hat{\sigma}^2 = \frac{\hat{\varepsilon}^T \hat{\varepsilon}}{K - N}$$



degrees of freedom

variance or

mean of squared errors (MSE)

The Math  
What is Linear?

**Stats**  
a Python Example

y:	response
x:	regressors
$\beta$ :	factors
$\beta_0$ :	intercept
$\varepsilon$ :	error





summary:

the model:

$$Y = X\beta + \varepsilon$$

the fit:

$$\hat{Y} = X\hat{\beta} = X(X^T X)^{-1} X^T Y$$

SSE:

$$\hat{\varepsilon}^T \hat{\varepsilon} = Y^T (I - H) Y \quad (\text{after the fit})$$

variance or MSE:

$$\hat{\sigma}^2 = \frac{\hat{\varepsilon}^T \hat{\varepsilon}}{K - N} \quad (\text{after the fit})$$

often fit quality is judged by

$$R^2 := 1 - \frac{\sum_k (\hat{y}_k - y_k)^2}{\sum_k (y_k - \langle y \rangle)^2}$$

or adjusted  $R^2$

$$\bar{R}^2 := R^2 - (1 - R^2) \frac{N}{K - N - 1}$$

and it is said that the fit is good if  $R^2$  is close to one....

**...but that is not true...**

The Math  
What is Linear?

**Stats**  
a Python Example

<b>y:</b>	<b>response</b>
<b>x:</b>	<b>regressors</b>
<b><math>\beta</math>:</b>	<b>factors</b>
<b><math>\beta_0</math>:</b>	<b>intercept</b>
<b><math>\varepsilon</math>:</b>	<b>error</b>





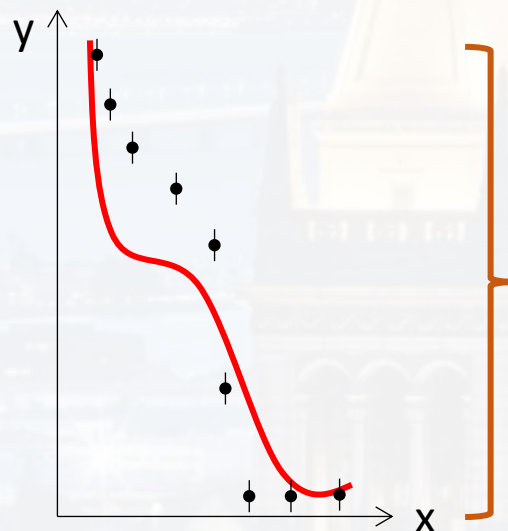
$$R^2 = 1 - \frac{\sum_{i=1}^N (y_i - \hat{y}_i)^2}{\sum_{i=1}^N (y_i - \bar{y})^2}$$

variance data vs model  
(aka residual sum of squares)

variance of the data  
(aka total sum of squares)

The Math  
What is Linear?  
**Stats**  
a Python Example

**Note:** do not confuse  $R^2$  with Pearsons coefficient:  $\rho = \frac{cov(x,y)}{\sqrt{var(x)var(y)}}$



data variance can be huge  
(i. e. exponential functions)  
→  $R^2$  could be around 1.0  
even if fit is completely off!

$\bar{y}$ : mean of the data point values

$\bar{y}$ : mean of the data point values  
 $y_i$ : measured value of data point  
 $\sigma_i$ : statistical error of  $y_i$  (often aka  $ey_i$ )  
 $\hat{y}_i$ : prediction by the model *after the fit*  
 $N$ : number of data points  
 $p$ : number of fit parameter

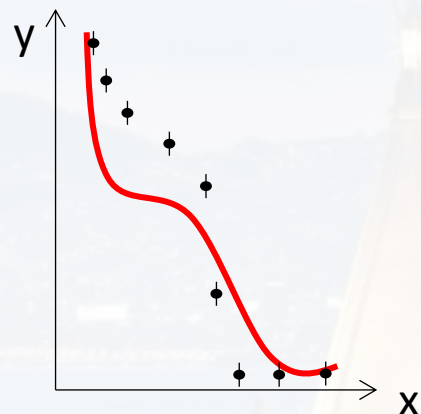


$$R^2 = 1 - \frac{\sum_{i=1}^N (y_i - \hat{y}_i)^2}{\sum_{i=1}^N (y_i - \bar{y})^2}$$

variance data vs model  
(aka residual sum of squares)

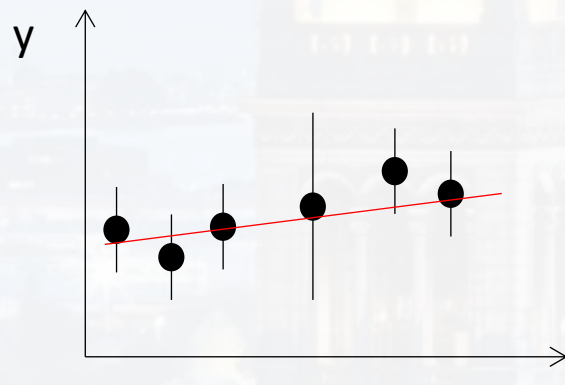
variance of the data  
(aka total sum of squares)

The Math  
What is Linear?  
**Stats**  
a Python Example



$\bar{y}$ : mean of the data point values

data variance can be huge  
(i. e. exponential functions)  
→  $R^2$  could be around 1.0  
even if fit is completely off!



variance data vs model  
(aka residual sum of squares)

variance of the data  
(aka total sum of squares)

$$\approx 1 \rightarrow R^2 = 0$$

**although the fit is good!**

$\bar{y}$ : mean of the data point values  
 $y_i$ : measured value of data point  
 $\sigma_i$ : statistical error of  $y_i$  (often aka  $ey_i$ )  
 $\hat{y}_i$ : prediction by the model *after the fit*  
 $N$ : number of data points  
 $p$ : number of fit parameter

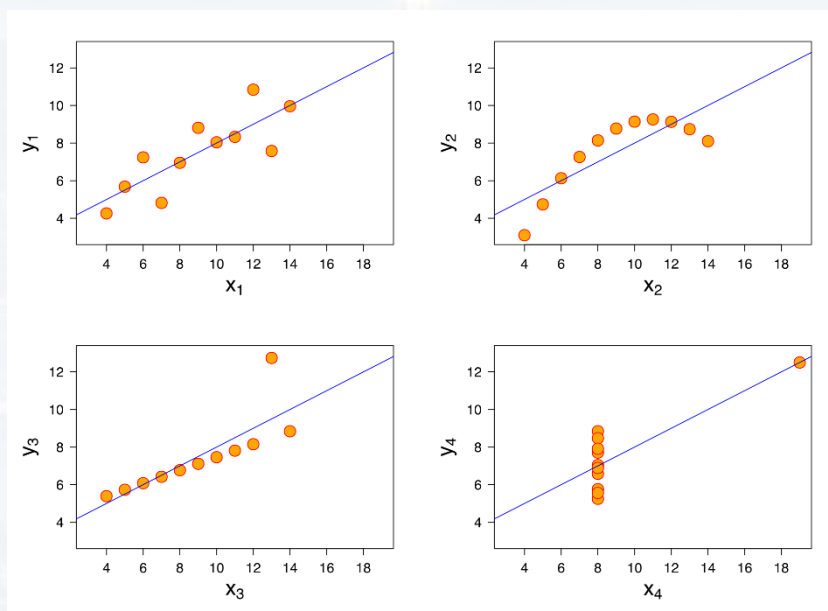


$$R^2 = 1 - \frac{\sum_{i=1}^N (y_i - \hat{y}_i)^2}{\sum_{i=1}^N (y_i - \bar{y})^2}$$

variance data vs model  
(aka residual sum of squares)

variance of the data  
(aka total sum of squares)

The Math  
What is Linear?  
**Stats**  
a Python Example



all plots: same  $R^2$

**conclusion:**

$R^2$  is **not** a measure of the fit quality (but  $\chi^2$  is)

**Given a good fit**,  $R^2$  tells how strong the dependent variable responds to the independent variable

Also, Wiki is full of examples...

...and warnings (see “caveats” therein)



see `Walk_Through_LinRegression.ipynb`

```
import numpy as np
```

```
import pandas as pd
```

```
import matplotlib.pyplot as plt
```

```
import seaborn as sns
```

```
import pylab
```

```
import scipy.stats as stats
```

```
import statsmodels.api as sm
```

```
from statsmodels.formula.api import ols
```

```
from sklearn.preprocessing import MinMaxScaler
```

The Math  
What is Linear?  
Stats  
**a Python Example**

reading .xlsx  
.csv  
.txt  
...

standard plots

fancy plots:  
here a pair-  
plot

Q-Q plot

the actual  
super tool for  
superb data  
analysis

scaling and normalizing





```
Train = pd.read_csv("molecular_train_gbc.csv")  
Test  = pd.read_csv("molecular_test_gbc.csv")
```

- 1) loading data
- 2) plotting data
- 3) scaling data
- 4) fitting model
- 5) evaluating model

	$x_1$	$x_2$	$x_3$	$x_4$	$x_5$	$y_k$
Index	molecular_weight	electronegativity	bond_lengths	num_hydrogen_bonds	logP	toxicity_score
0	341.704	2.65585	3.09407	2	9.11147	80.9281
1	335.951	3.22262	2.89039	7	8.92848	83.4911
2	235.203	2.44115	2.48203	1	6.49731	61.8406
3	246.505	2.76656	2.71547	7	7.45089	57.0538
4	437.939	3.4801	3.59569	3	10.9156	131.326

$$y_k = \beta_0 + \sum_{n=1}^N \beta_n x_n + \epsilon$$

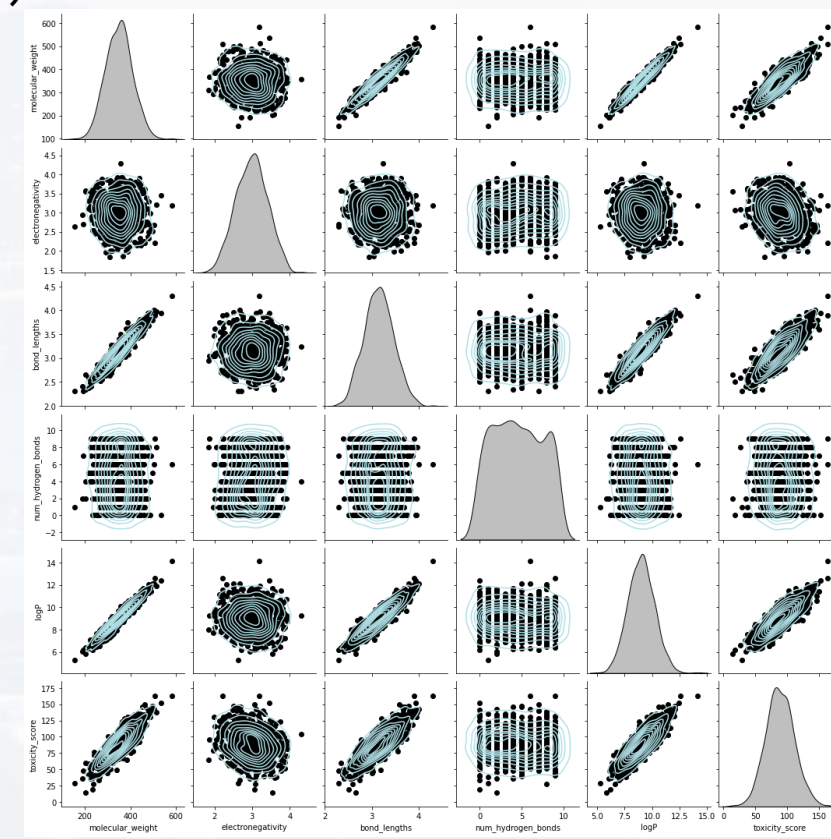
$y$ : toxicity\_score  
 $x_n$ : molecular\_weight, electronegativity,  
bond\_lengths, num\_hydrogen\_bonds, logP



```
Train = pd.read_csv("molecular_train_gbc.csv")  
Test  = pd.read_csv("molecular_test_gbc.csv")
```

- 1) loading data
- 2) plotting data
- 3) scaling data
- 4) fitting model
- 5) evaluating model

```
out = sns.pairplot(Train, kind = "kde", \  
                  plot_kws = {'color': [176/255, 224/255, 230/255]}, \  
                  diag_kws = {'color': 'black'})  
out.map_offdiag(plt.scatter, color = 'black')
```





```
Train = pd.read_csv("molecular_train_gbc.csv")  
Test  = pd.read_csv("molecular_test_gbc.csv")
```

```
out = sns.pairplot(Train, kind = "kde", \  
                  plot_kws = {'color': [176/255, 224/255, 230/255]}, \  
                  diag_kws = {'color': 'black'})  
out.map_offdiag(plt.scatter, color = 'black')
```

```
scaler = MinMaxScaler(feature_range = (0, 1))  
TrainS = scaler.fit_transform(Train)  
TestS  = scaler.transform(Test)
```

the scaler returns an `np.array`  
→ convert back to data frame

```
TrainS = pd.DataFrame(TrainS, columns = Train.columns)  
TestS  = pd.DataFrame(TestS, columns = Train.columns)
```

- 1) loading data
- 2) plotting data
- 3) scaling data
- 4) fitting model
- 5) evaluating model





```
TrainS = pd.DataFrame(TrainS, columns = Train.columns)
TestS   = pd.DataFrame(TestS,  columns = Train.columns)
```

```
equation = 'toxicity_score ~ ' + '+'.join(Train.columns[:-1])
print(equation)
```

- 1) loading data
- 2) plotting data
- 3) scaling data
- 4) fitting model
- 5) evaluating model

$$y_k = \beta_0 + \sum_{n=1}^N \beta_n x_n + \epsilon$$

```
toxicity_score ~      molecular_weight + electronegativity +
                      bond_lengths + num_hydrogen_bonds + logP
```

```
my_model = ols(equation, data = TrainS).fit()
my_model.summary()
```

**OLS** (ordinary least squares)





my\_model.summary()

- 1) loading data
- 2) plotting data
- 3) scaling data
- 4) fitting model
- 5) evaluating model

number of data points  
is much larger than  
the number of regressors  
→ degree of freedom  
approx. no of obs

OLS Regression Results

not the fit quality!

Dep. Variable:	toxicity_score	R-squared:	0.790
Model:	OLS	Adj. R-squared:	0.789
Method:	Least Squares	F-statistic:	597.5
Date:	Fri, 13 Sep 2024	Prob (F-statistic):	3.34e-266
Time:	20:57:10	Log-Likelihood:	1013.0
No. Observations:	800	AIC:	-2014.
Df Residuals:	794	BIC:	-1986.
Df Model:	5		

Covariance Type: nonrobust

p-values for factors

	coef	std err	t	P> t	[0.025
Intercept	0.1494	0.012	12.533	0.000	0.126
molecular_weight	0.7961	0.089	8.982	0.000	0.622
electronegativity	-0.1682	0.015	-11.591	0.000	-0.197
bond_lengths	0.0204	0.049	0.417	0.677	-0.076
num_hydrogen_bonds	0.0035	0.008	0.458	0.647	-0.011
logP	0.1246	0.072	1.723	0.085	-0.017

Omnibus:	2.249	Durbin-Watson:	1.984
Prob(Omnibus):	0.325	Jarque-Bera (JB):	2.240
Skew:	-0.129	Prob(JB):	0.326
Kurtosis:	2.980	Cond. No.	65.6

p-value for  
constant model

p-values for  
factors

2σ conf range of  
factors

$$y_k = \beta_0 + \sum_{n=1}^N \beta_n x_n + \epsilon$$

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.



more accurate: determining the **p-values for the factors using ANOVA** for the corresponding residuals

```
table = sm.stats.anova_lm(my_model, typ = 1)
print(table)
```

- 1) loading data
- 2) plotting data
- 3) scaling data
- 4) fitting model
- 5) evaluating model

$$y_k = \beta_0 + \sum_{n=1}^N \beta_n x_n + \epsilon$$

	df	sum_sq	mean_sq	F	PR(>F)	vs from t-test
molecular_weight	1.0	13.346285	13.346285	2847.525516	8.024085e-265	0.0000
electronegativity	1.0	0.640388	0.640388	136.631363	3.085962e-29	0.0000
bond_lengths	1.0	0.000684	0.000684	0.145954	7.025342e-01	0.6766
num_hydrogen_bonds	1.0	0.000703	0.000703	0.150055	6.985866e-01	0.6473
logP	1.0	0.013917	0.013917	2.969353	8.524510e-02	0.0852
Residual	794.0	3.721459	0.004687	NaN	NaN	

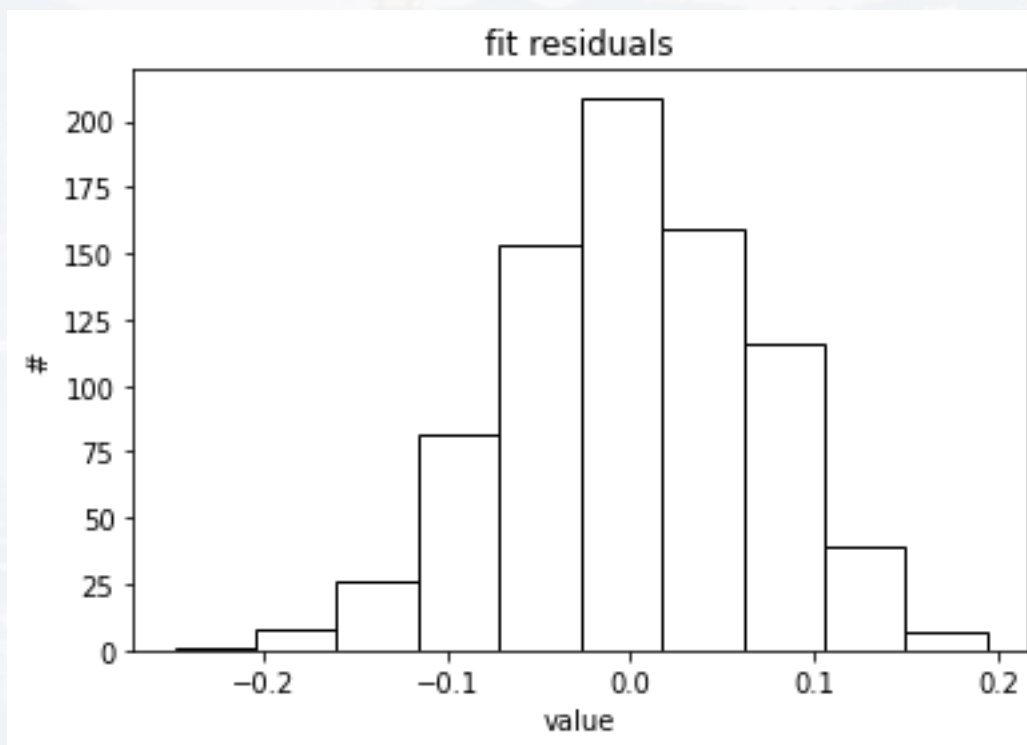


```
residuals = my_model.resid
```

```
plt.hist(residuals, color = 'w', edgecolor = 'black')  
plt.title('fit residuals')  
plt.ylabel('#')  
plt.xlabel('value')  
plt.show()
```

- 1) loading data
- 2) plotting data
- 3) scaling data
- 4) fitting model
- 5) evaluating model

$$y_k = \beta_0 + \sum_{n=1}^N \beta_n x_n + \epsilon$$



residuals approx.  
normally distributed  
around  $\mu = 0$



```
residuals = my_model.resid
```

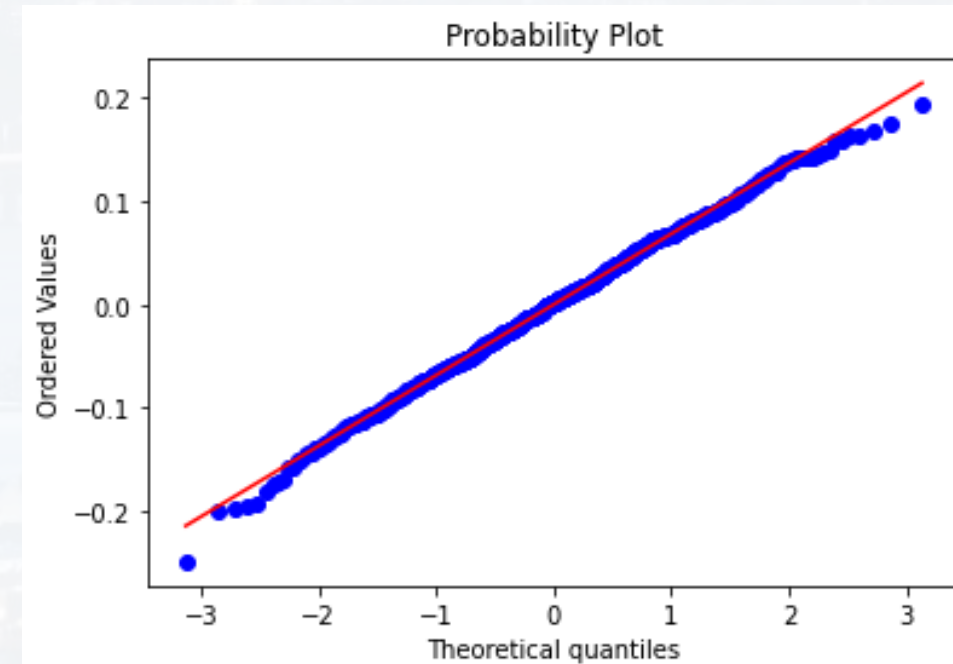
```
plt.hist(residuals, color = 'w', edgecolor = 'black')  
plt.title('fit residuals')  
plt.ylabel('#')  
plt.xlabel('value')  
plt.show()
```

```
stats.probplot(residuals, dist = "norm", plot = pylab)  
pylab.show()
```

- 1) loading data
- 2) plotting data
- 3) scaling data
- 4) fitting model
- 5) evaluating model

residuals approx.  
normally distributed  
around  $\mu = 0$

$$y_k = \beta_0 + \sum_{n=1}^N \beta_n x_n + \epsilon$$







```
Ypred = my_model.predict(TestS)
```

```
higher = np.max([Ypred, TestS.toxicity_score])
```

```
lower = np.min([Ypred, TestS.toxicity_score])
```

```
plt.plot([lower, higher], [lower, higher], c = [0, 0, 0, 0.2],\n         linewidth = 4)
```

```
plt.scatter(TestS.toxicity_score, Ypred, marker = '.', c = 'k')
```

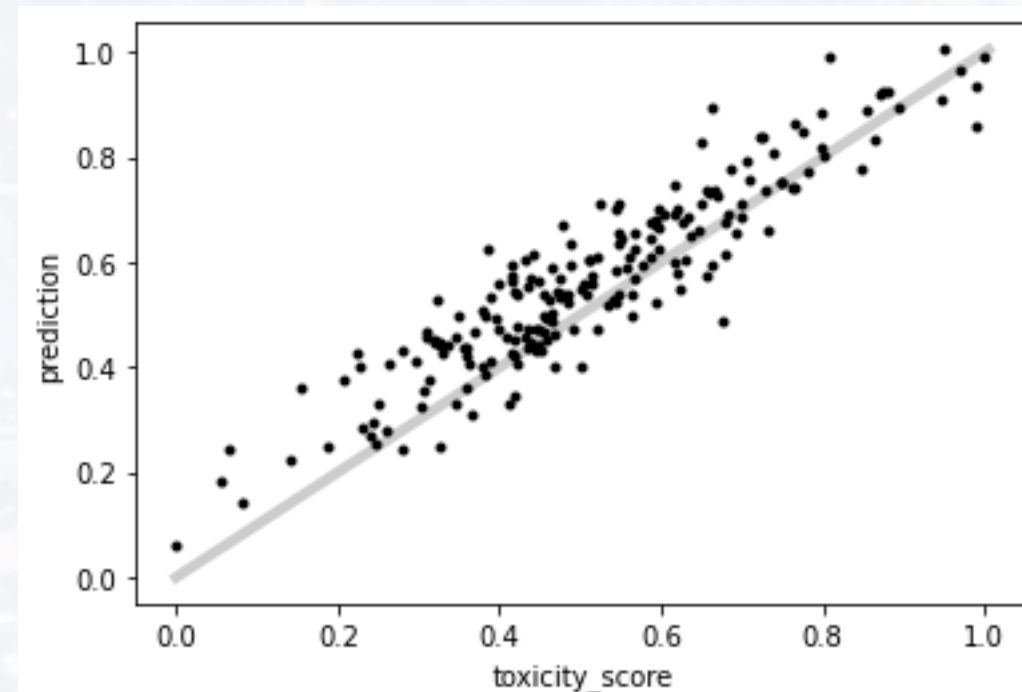
```
plt.ylabel('prediction')
```

```
plt.xlabel('toxicity score')
```

```
plt.show()
```

- 1) loading data
- 2) plotting data
- 3) scaling data
- 4) fitting model
- 5) evaluating model

$$y_k = \beta_0 + \sum_{n=1}^N \beta_n x_n + \epsilon$$





```
Ypred = my_model.predict(TestS)
```

```
higher = np.max([Ypred, TestS.toxicity_score])
```

```
lower = np.min([Ypred, TestS.toxicity_score])
```

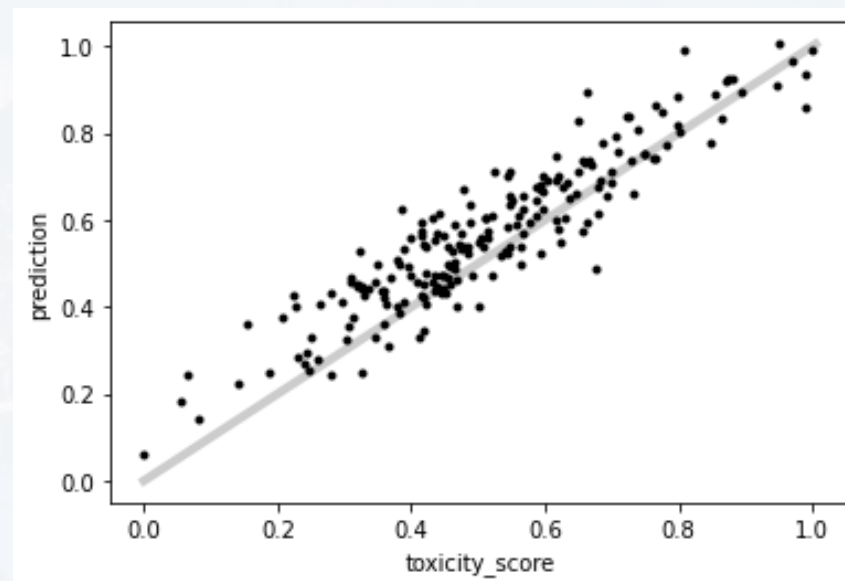
```
plt.plot([lower, higher], [lower, higher], c = [0, 0, 0, 0.2], linewidth = 4)
```

```
plt.scatter(TestS.toxicity_score, Ypred, marker = '.', c = 'k')
```

```
plt.ylabel('prediction')
```

```
plt.xlabel('toxicity score')
```

```
plt.show()
```



```
mean_dev = np.sum( abs(TestS.toxicity_score - Ypred) )/len(Ypred)  
print(mean_dev)
```

5%

- 1) loading data
- 2) plotting data
- 3) scaling data
- 4) fitting model
- 5) evaluating model

$$y_k = \beta_0 + \sum_{n=1}^N \beta_n x_n + \epsilon$$



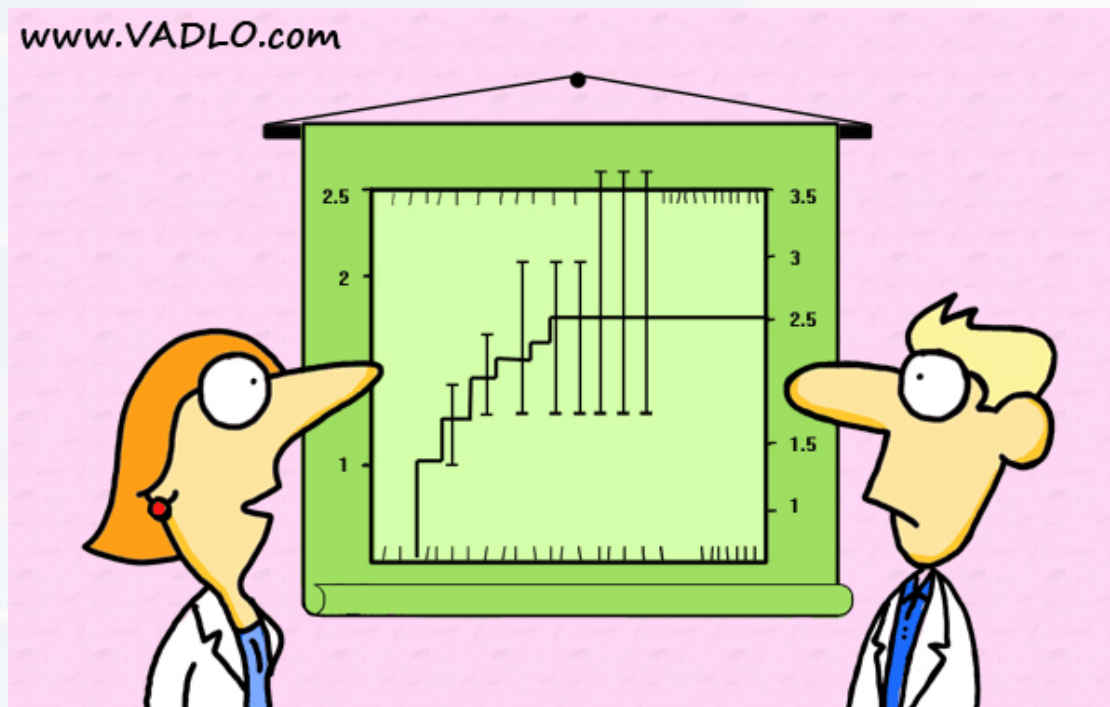
### Outline

#### Linear Regression

- The Math
- What is Linear?
- Stats
- a Python Example

#### Logistic Regression

#### Curve Fitting



“Did you really have to show the error bars?”





linear model: regressors are continuous or categorical,  
response is continuous

logistic model: response is **categorical**

$y$ :	response
$x$ :	regressors (assumed to be independent)
$\beta$ :	factors
$\beta_0$ :	intercept
$\epsilon$ :	error (stochasticity of the data, assumed to be normally dist.)

Index	molecular_weight	electronegativity	bond_lengths	num_hydrogen_bonds	logP	label
0	341.704	2.65585	3.09407	2	9.11147	Toxic
1	335.951	3.22262	2.89039	7	8.92848	Toxic
2	235.203	2.44115	2.48203	1	6.49731	Non-Toxic
3	246.505	2.76656	2.71547	7	7.45089	Non-Toxic
4	437.939	3.4801	3.59569	3	10.9156	Non-Toxic



linear model: regressors are continuous or categorical,  
response is continuous

logistic model: response is **categorical**

y:	response
x:	regressors (assumed to be independent)
$\beta$ :	factors
$\beta_0$ :	intercept
$\epsilon$ :	error (stochasticity of the data, assumed to be normally dist.)

dichotomic model: **probability** to be in state A)  $\rightarrow p$

**probability** to be in state B)  $\rightarrow 1 - p$

label
Toxic
Toxic
Non-Toxic
Non-Toxic
Non-Toxic

ansatz:

$$\log \left( \frac{p}{1-p} \right) = \beta_0 + \sum_{n=1}^N \beta_n x_n + \epsilon$$

**log odds ratio: linear model**



dichotomic model:

label
Toxic
Toxic
Non-Toxic
Non-Toxic
Non-Toxic

**probability** to be in state A)  $\rightarrow p$

**probability** to be in state B)  $\rightarrow 1 - p$

y:	response
x:	regressors (assumed to be independent)
$\beta$ :	factors
$\beta_0$ :	intercept
$\epsilon$ :	error (stochasticity of the data, assumed to be normally dist.)

ansatz:

$$\log\left(\frac{p}{1-p}\right) = \beta_0 + \sum_{n=1}^N \beta_n x_n + \epsilon$$

**log odds ratio: linear model**

$\rightarrow$  probability for being in a certain state

$$p = \frac{e^{\beta_0 + \beta_1 x_1 + \dots}}{1 + e^{\beta_0 + \beta_1 x_1 + \dots}}$$

often:

$$\text{logit}(p) = \log\left(\frac{p}{1-p}\right)$$

examples:

- probability that a gene has been mutated
- probability of being diseased (cancer, alzheimer etc) as function of age, environmental influence etc ...
- Verhulst equation:  $N(t) = N_0 \frac{e^{rt}}{C + e^{rt}}$
- activation functions in ANNs



$$\log\left(\frac{p}{1-p}\right) = \beta_0 + \sum_{n=1}^N \beta_n x_n + \epsilon$$

Note: one can derive the logit function from max. entropy too!

y:	response
x:	regressors (assumed to be independent)
$\beta$ :	factors
$\beta_0$ :	intercept
$\epsilon$ :	error (stochasticity of the data, assumed to be normally dist.)

$$p = \frac{e^{\beta_0 + \beta_1 x_1 + \dots}}{1 + e^{\beta_0 + \beta_1 x_1 + \dots}} = \frac{1}{1 + e^{-\beta_0 - \beta_1 x_1 - \dots}}$$

onset of Alzheimer's disease (AD) is a function of **age** and years spent in **education**  
(and other risk factors we ignore here for the sake of simplicity)

education:  $d = x_1$  [yrs]

age:  $a = x_2$  [yrs]

model: 
$$p_{AD} = \frac{1}{1 + e^{-\beta_0 - \beta_1 d - \beta_2 a}}$$

+ data set + fit  $\rightarrow$

$$\begin{aligned}\beta_0 &= +0.1 \\ \beta_1 &= -1.5 \\ \beta_2 &= +0.12\end{aligned}$$

- positive value  $\rightarrow$  increasing  $p$
- negative value  $\rightarrow$  decreasing  $p$
- intercept: "background" prevalence, not related to environmental/internal conditions





model: 
$$p_{AD} = \frac{1}{1 + e^{-\beta_0 - \beta_1 d - \beta_2 a}}$$

education:  $d = x_1$  [yrs]  $\beta_0 = +0.1$   
age:  $a = x_2$  [yrs]  $\beta_1 = -1.5$   
 $\beta_2 = +0.12$

example: **65yrs** old person, **8yrs** spent in education

$$\rightarrow p_{AD} = 1.6\%$$

**65yrs** old person, **13yrs** spent in education

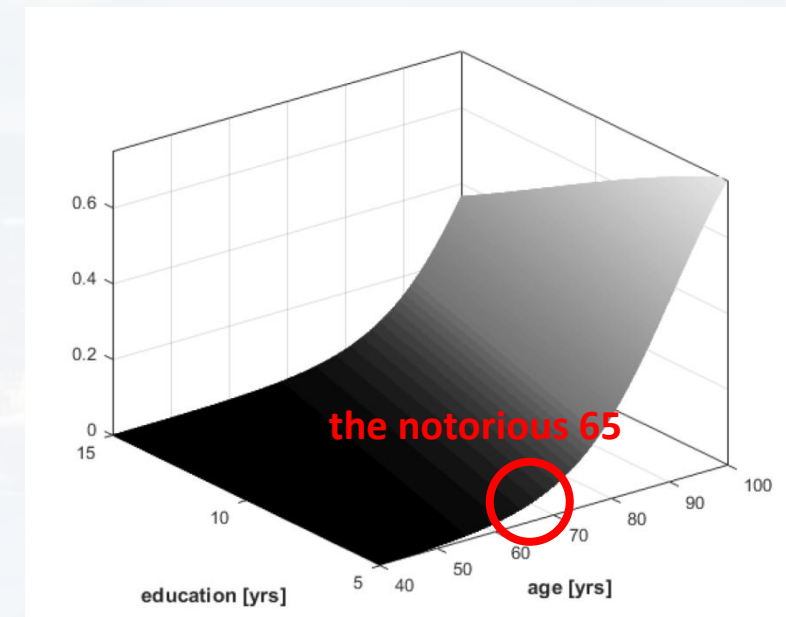
$$\rightarrow p_{AD} = 0.001\%$$

How does education compensate aging?

$$p_{AD}(d + \bar{d}, a + \bar{a}) = p_{AD}(d, a)$$

$$\rightarrow \bar{a} = 12.5 \bar{d}$$

y:	response
x:	regressors (assumed to be independent)
$\beta$ :	factors
$\beta_0$ :	intercept
$\varepsilon$ :	error (stochasticity of the data, assumed to be normally dist.)



hence, one more year prolonged education compensates  
**12.5 years of aging**

(warning: don't confuse correlation with causation here!)



model: 
$$p_{AD} = \frac{1}{1 + e^{-\beta_0 - \beta_1 d - \beta_2 a}}$$

education:  $d = x_1$  [yrs]

age:  $a = x_2$  [yrs]

$$\beta_0 = +0.1$$

$$\beta_1 = -1.5$$

$$\beta_2 = +0.12$$

y:	response
x:	regressors (assumed to be independent)
$\beta$ :	factors
$\beta_0$ :	intercept
$\varepsilon$ :	error (stochasticity of the data, assumed to be normally dist.)

How does the risk of onset changes *per year*?

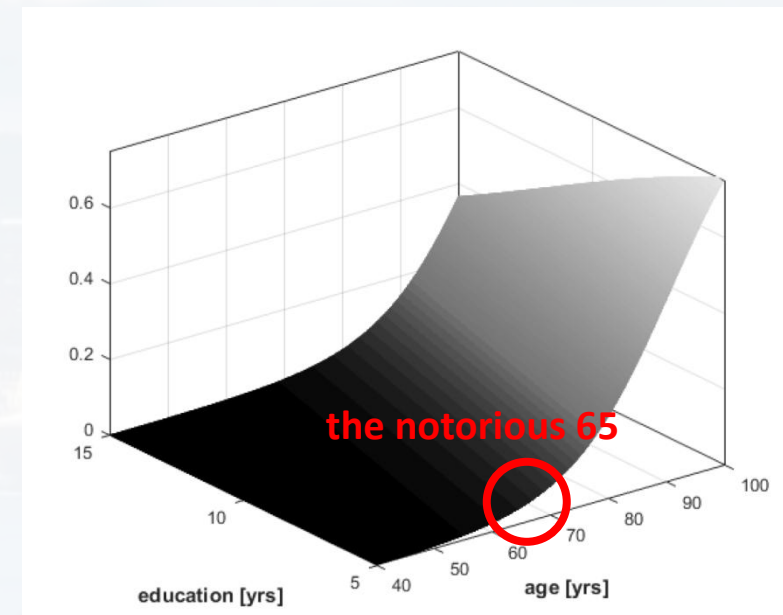
relative change:

$$\frac{p_{AD}(a+1) - p_{AD}(a)}{p_{AD}(a)} \approx e^{\beta_2} - 1 \approx 12.7\%$$

$p_{AD} \ll 1$  (hence, for small  $\Delta a$  and “young” ages, i. e. below  $\approx 80$  yrs )

**the risk of getting AD increases by 12.7% every year**

(warning: does not mean that it increases by 127% in ten yrs – we made an approximation!)





model:  $p_{AD} = \frac{1}{1+e^{-\beta_0-\beta_1d-\beta_2a}}$

education:  $d = x_1$  [yrs]

age:  $a = x_2$  [yrs]

$$\beta_0 = +0.1$$

$$\beta_1 = -1.5$$

$$\beta_2 = +0.12$$

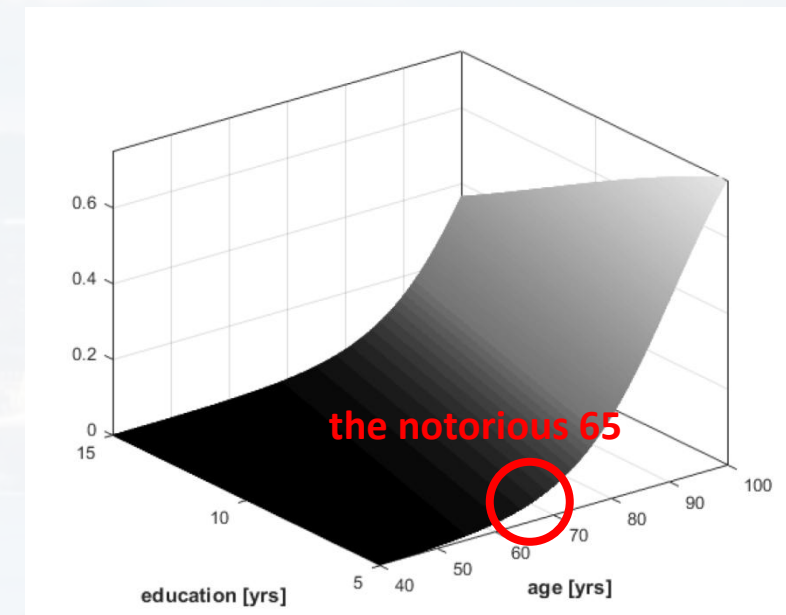
y:	response
x:	regressors (assumed to be independent)
$\beta$ :	factors
$\beta_0$ :	intercept
$\varepsilon$ :	error (stochasticity of the data, assumed to be normally dist.)

How does the risk of onset changes *per year*?

more precise: relative change of the odds ratio

$$\frac{\frac{\partial}{\partial x_i} \left( \frac{p_{AD}}{1 - p_{AD}} \right)}{\frac{p_{AD}}{1 - p_{AD}}} = \beta_i$$

$x_i$  is the desired regressor,  
for example, age again ( $x_2$ )



the factors  $\beta_i$  indicate how strong (and in which direction)  $p$  changes wrt a regressor  $x_i$





let us return to the molecule data set:

```
Train = pd.read_csv("molecular_train_gbc_cat.csv")  
Test  = pd.read_csv("molecular_test_gbc_cat.csv")
```

Index	molecular_weight	electronegativity	bond_lengths	num_hydrogen_bonds	logP	label
0	341.704	2.65585	3.09407	2	9.11147	Toxic
1	335.951	3.22262	2.89039	7	8.92848	Toxic
2	235.203	2.44115	2.48203	1	6.49731	Non-Toxic
3	246.505	2.76656	2.71547	7	7.45089	Non-Toxic
4	437.939	3.4801	3.59569	3	10.9156	Non-Toxic

```
import statsmodels.api as sm
```





it is the same data set → plotting and scaling is as before

```
X = sm.add_constant(TrainS)
```

adding the  
intercept

```
Y = pd.get_dummies(Train['Label'])
```

Python needs  
True/False  
as categorical

$$p = \frac{e^{\beta_0 + \beta_1 x_1 + \dots}}{1 + e^{\beta_0 + \beta_1 x_1 + \dots}}$$

```
In [48]: print(Y)
      Non-Toxic  Toxic
0         False   True
1         False   True
2          True  False
3          True  False
4          True  False
```

we have two  
states: toxic /  
non-toxic

```
my_model = sm.GLM(Y, X, family = sm.families.Binomial()).fit()
```

```
my_model.summary()
```

GLM: general  
linear model



it is the same data set → plotting and scaling is as before

```
X = sm.add_constant(TrainS)
Y = pd.get_dummies(Train['Label'])
```

```
my_model = sm.GLM(Y, X, family = sm.families.Binomial()).fit()
my_model.summary()
```

Generalized Linear Model Regression Results						
=====						
Dep. Variable:	['Non-Toxic', 'Toxic']		No. Observations:	800		
Model:	GLM		Df Residuals:	794		
Model Family:	Binomial		Df Model:	5		
Link Function:	Logit		Scale:	1.0000		
Method:	IRLS		Log-Likelihood:	-332.82		
Date:	Sat, 14 Sep 2024		Deviance:	665.64		
Time:	20:59:18		Pearson chi2:	1.14e+03		
No. Iterations:	6		Pseudo R-squ. (CS):	0.4243		
Covariance Type:	nonrobust		p-values for factors			
=====						
	coef	std err	z	P> z	[0.025	0.975]
-----						
const	6.1641	0.585	10.536	0.000	5.017	7.311
molecular_weight	-10.4920	3.626	-2.893	0.004	-17.599	-3.385
electronegativity	3.2874	0.599	5.492	0.000	2.114	4.461
bond_lengths	0.6736	1.913	0.352	0.725	-3.075	4.422
num_hydrogen_bonds	-0.3082	0.303	-1.018	0.309	-0.902	0.285
logP	-7.6090	2.978	-2.555	0.011	-13.447	-1.771
=====						

$$p = \frac{e^{\beta_0 + \beta_1 x_1 + \dots}}{1 + e^{\beta_0 + \beta_1 x_1 + \dots}}$$

p-value for  
constant model

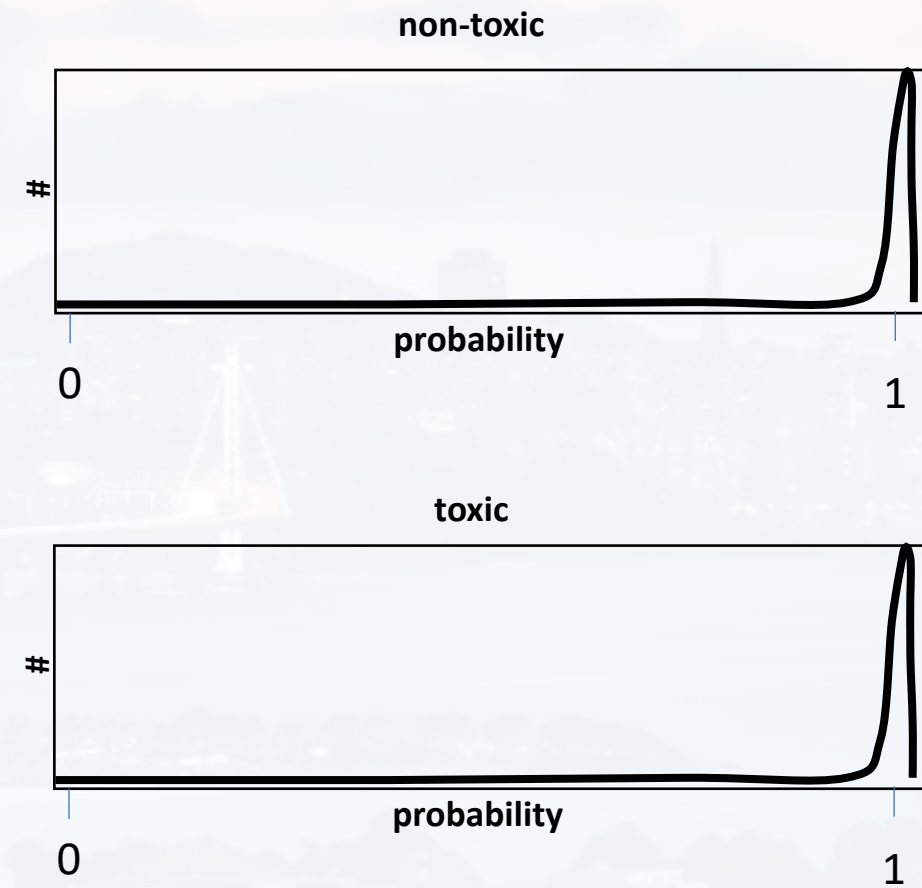
2σ conf range of  
factors



accuracy: How **often** did the model make the correct prediction.  
cross-entropy: How **certain** was the model when making the prediction.

ideal world:

true label	non-toxic	toxic
	100%	0%
non-toxic	100%	0%
toxic	0%	100%
		predicted label
		non-toxic      toxic

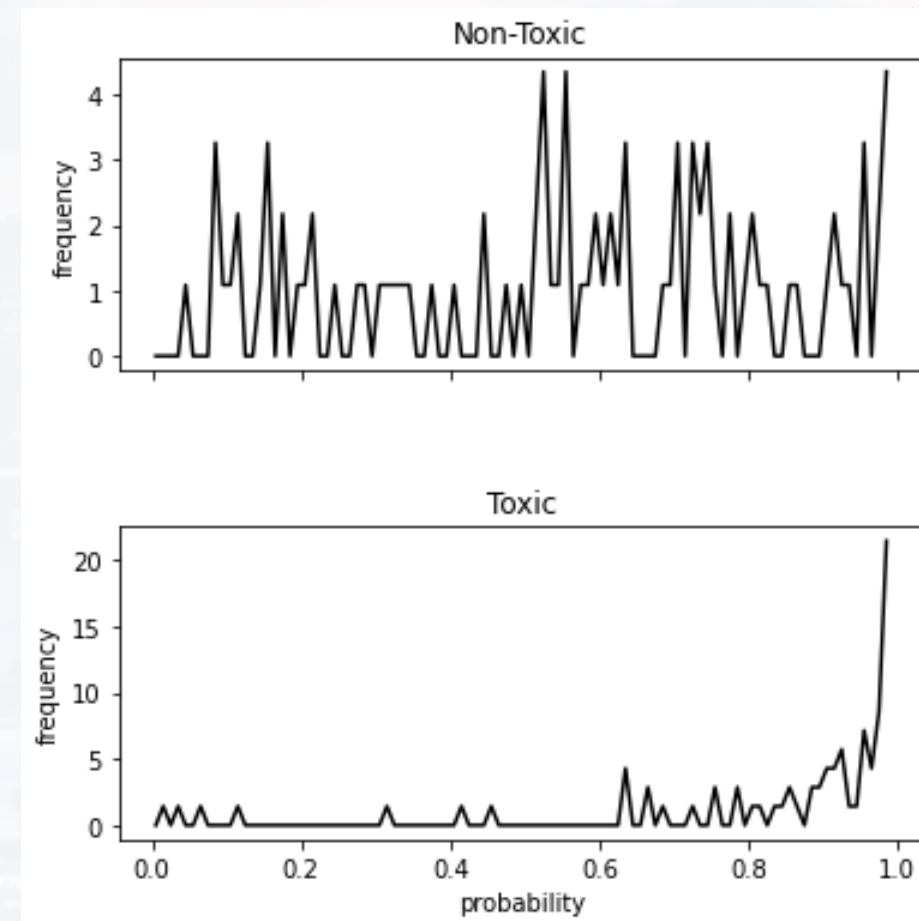
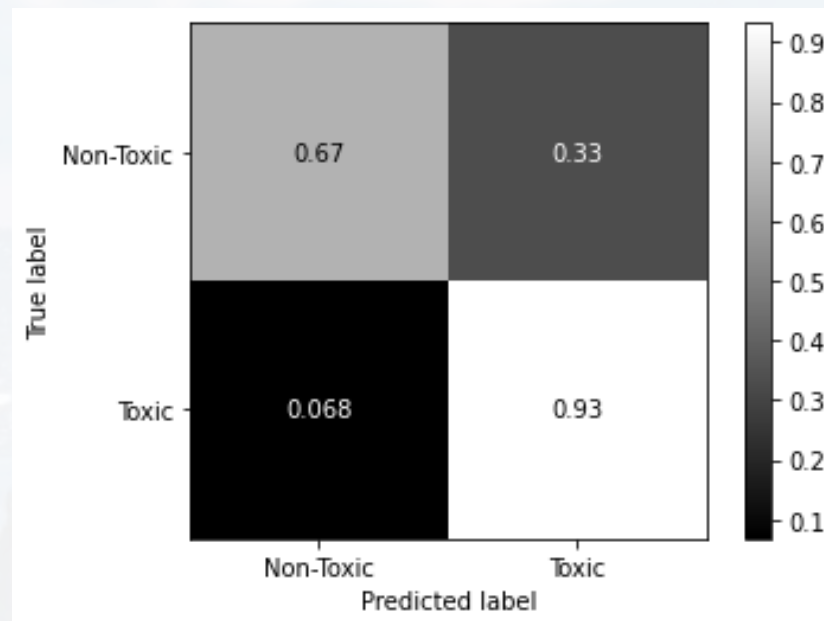


more details see [Chem 277](#)



accuracy: How **often** did the model make the correct prediction.  
cross-entropy: How **certain** was the model when making the prediction.

real world:



more details see [Chem 277](#)





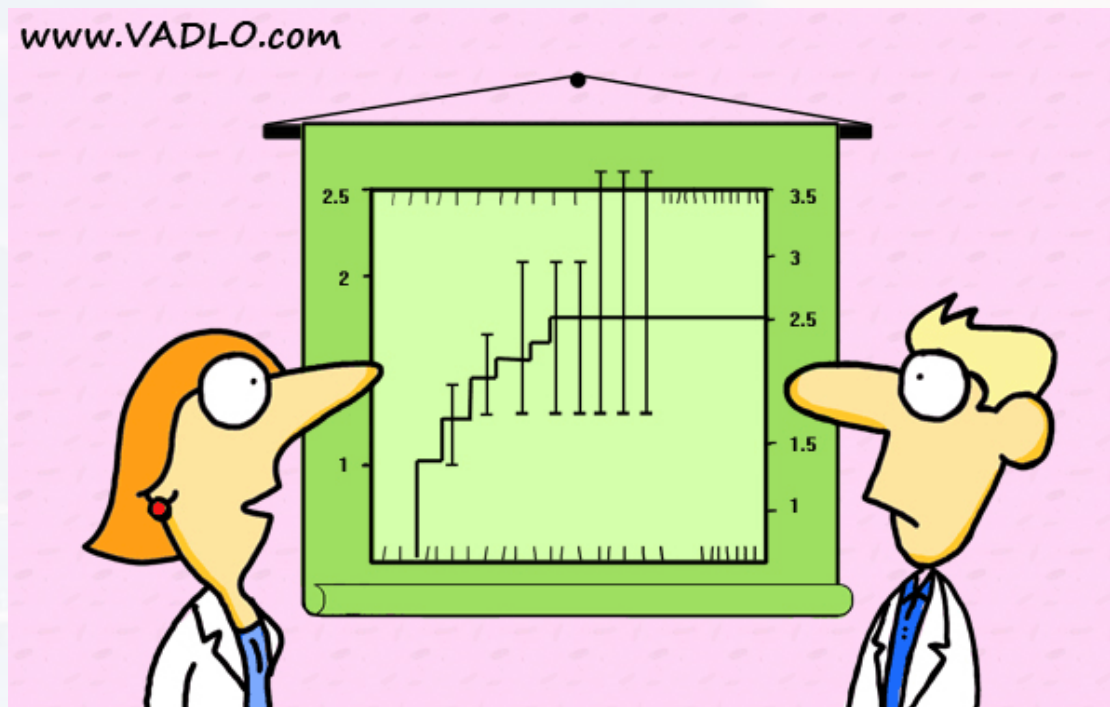
## Outline

### Linear Regression

- The Math
- What is Linear?
- Stats
- a Python Example

### Logistic Regression

### Curve Fitting

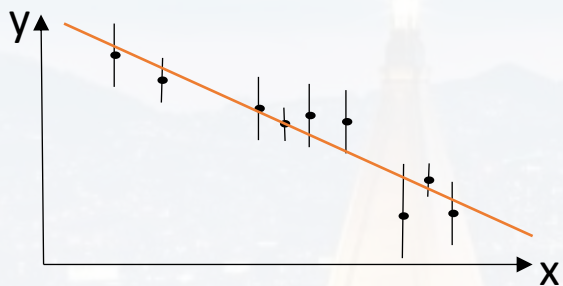


“Did you really have to show the error bars?”

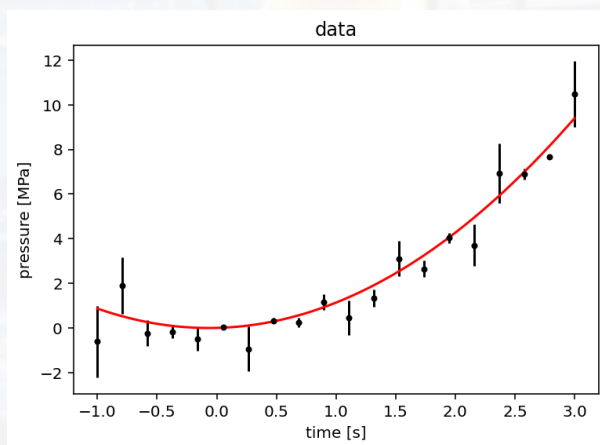


**problem:** set of data points  $y_i = f(x_i)$ , ideally each data point  $y_i$  has an error  $\sigma_i$   
we know  $f(x_i)$

**goal:** find parameters of  $f(x_i)$   
once we know the parameters  $\rightarrow$  prediction



model:  $f(x_i) = m x_i + n$   
parameter:  $m, n$

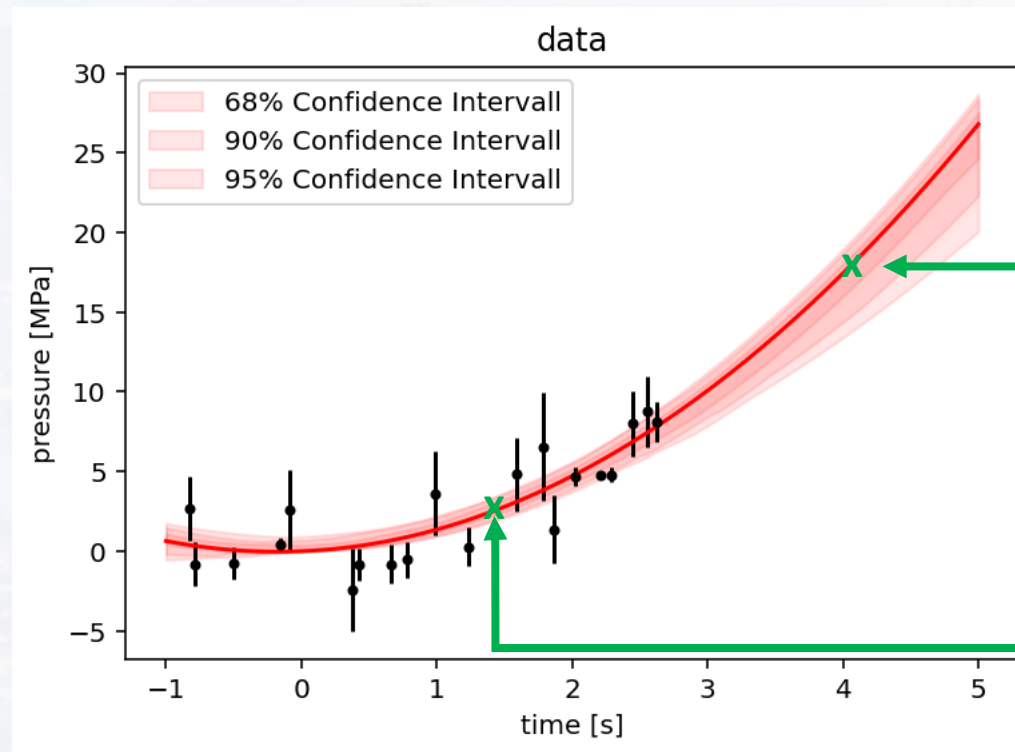


model:  $f(x_i) = a x_i^2 + b x_i + c$   
parameter:  $a, b, c$



**problem:** set of data points  $y_i = f(x_i)$ , ideally each data point  $y_i$  has an error  $\sigma_i$   
we know  $f(x_i)$

**goal:** find parameters of  $f(x_i)$   
once we know the parameters → **prediction**



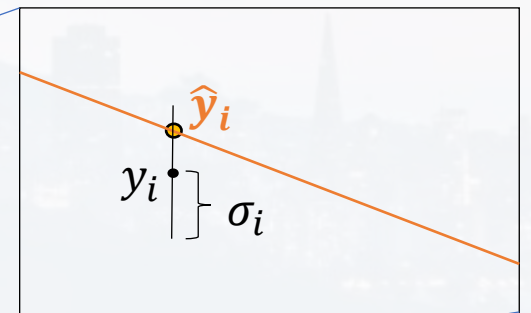
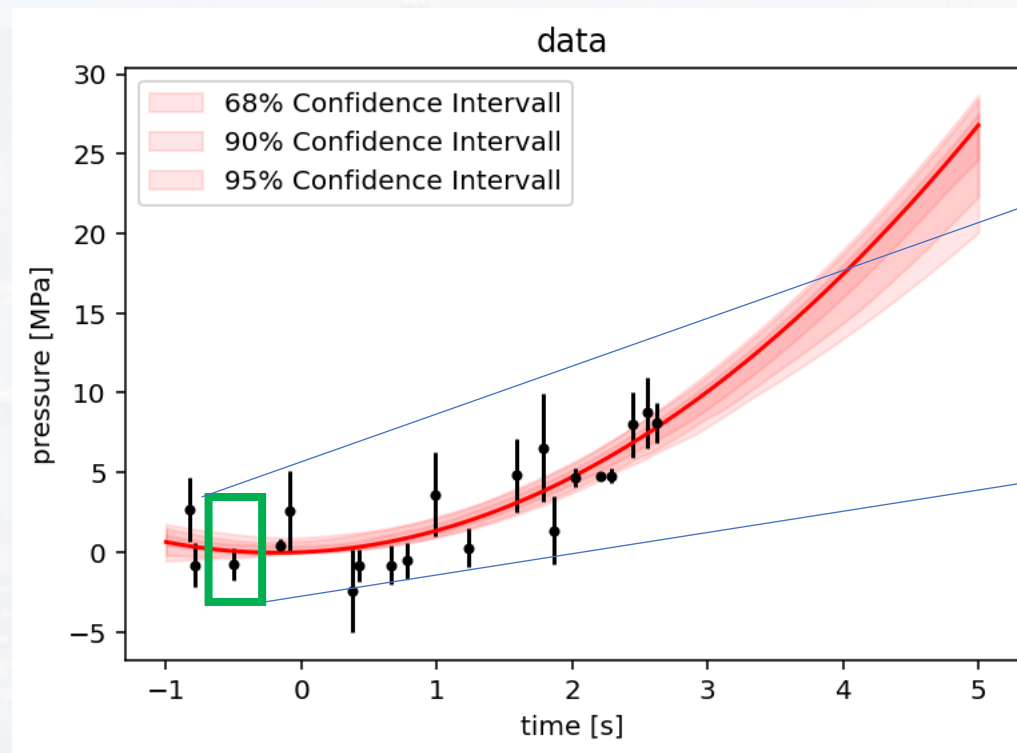
$y_i$  for data point  
**outside** known interval

$y_i$  for data point  
**within** known interval



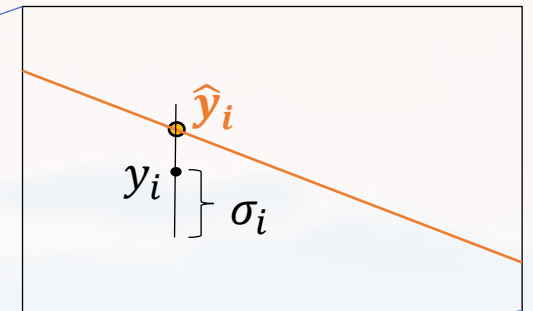
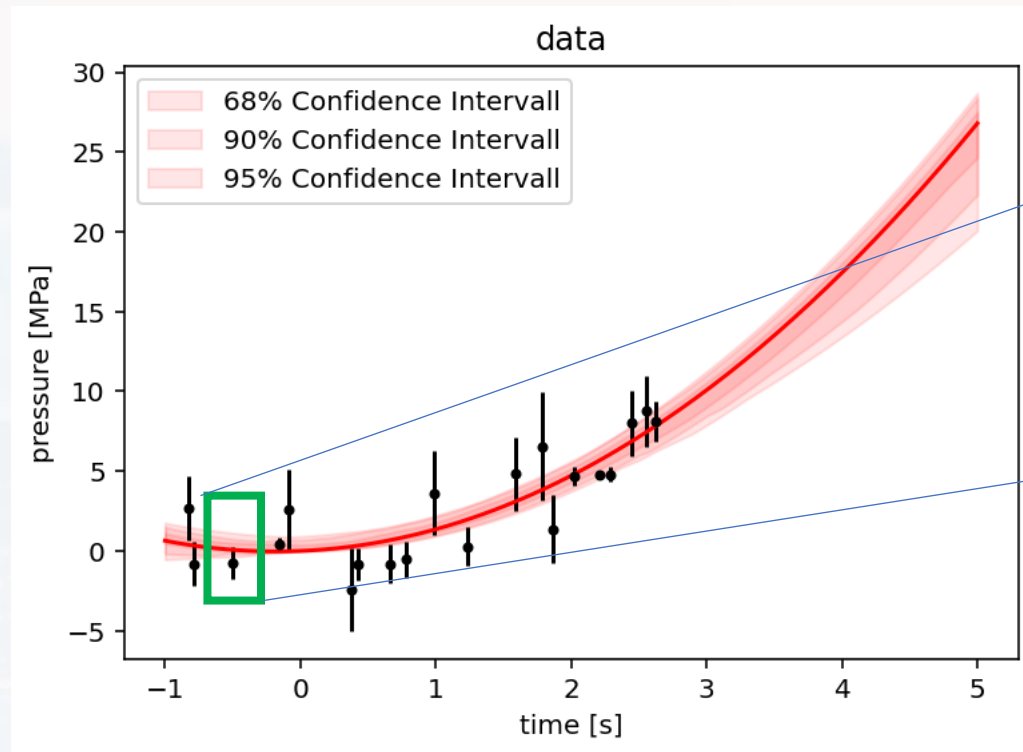
**problem:** set of data points  $y_i = f(x_i)$ , ideally each data point  $y_i$  has an error  $\sigma_i$   
we know  $f(x_i)$

**goal:** find parameters of  $f(x_i)$   
once we know the parameters  $\rightarrow$  prediction



$y_i$ : measured value of data point  
 $\sigma_i$ : statistical error of  $y_i$  (often aka  $ey_i$ )  
 $\hat{y}_i$ : prediction by the model *after the fit*  
 $N$ : number of data points  
 $p$ : number of fit parameter





finding best parameters by minimizing

$$\chi_{red}^2 = \frac{1}{df} \sum_{i=1}^N \left( \frac{y_i - \hat{y}_i}{\sigma_i} \right)^2$$

$df = N - p - 1$   
see module 8

$y_i$ : measured value of data point  
 $\sigma_i$ : statistical error of  $y_i$  (often aka  $ey_i$ )  
 $\hat{y}_i$ : prediction by the model *after the fit*  
 $N$ : number of data points  
 $p$ : number of fit parameter

or  $MSE = \frac{1}{df} \sum_{i=1}^N (y_i - \hat{y}_i)^2$

if no errors given

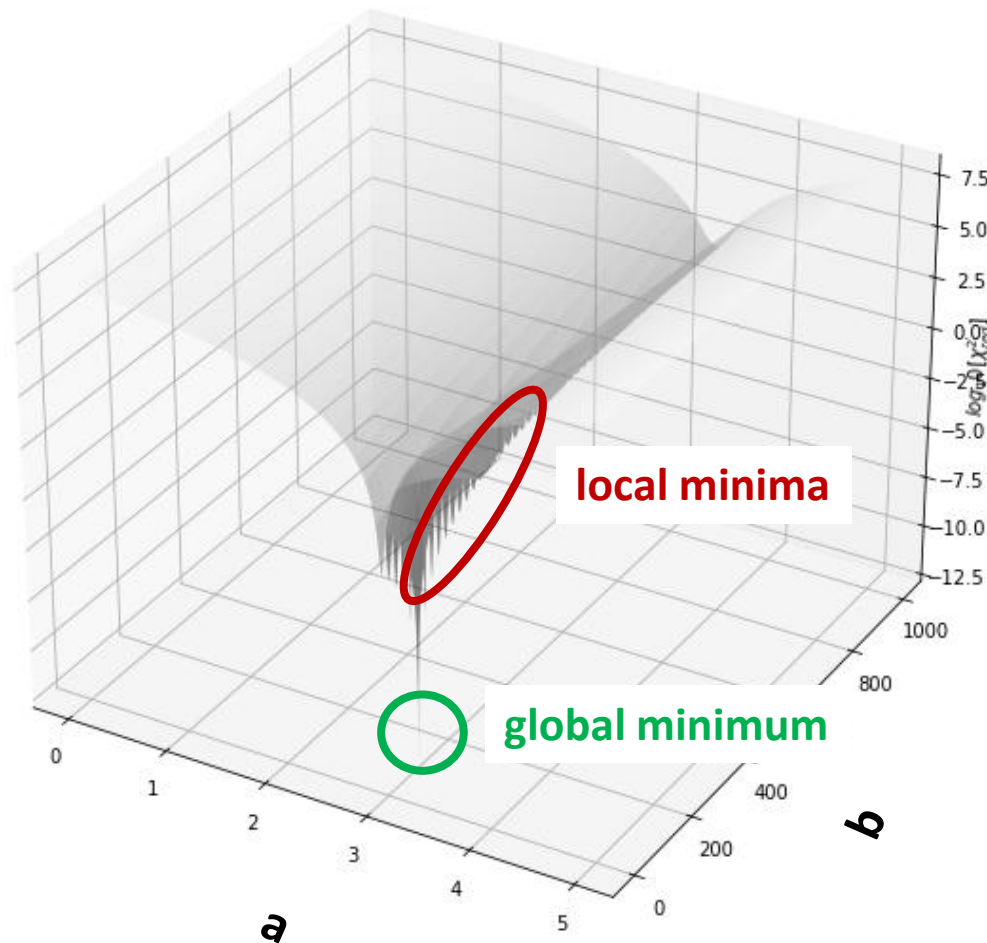


finding best parameters by minimizing

$$\chi_{red}^2 = \frac{1}{df} \sum_{i=1}^N \left( \frac{y_i - \hat{y}_i}{\sigma_i} \right)^2 \quad df = N - p - 1$$

see module 8

$\log(\chi_{red}^2)$



or

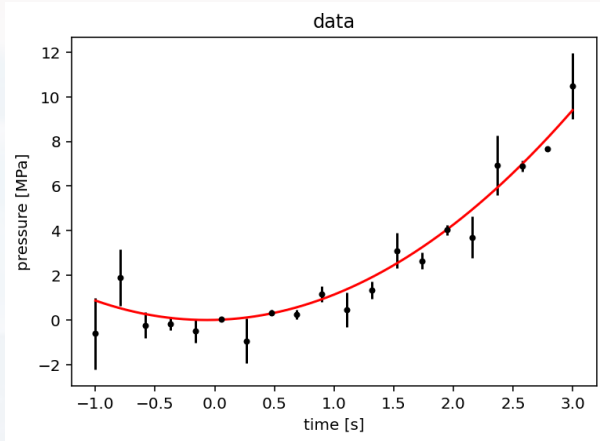
$$MSE = \frac{1}{df} \sum_{i=1}^N (y_i - \hat{y}_i)^2 \quad \text{if no errors given}$$

finding the global minimum of a higher dimensional non-analytical function

→ see also module 3



from module 8



$$\chi_{red}^2 = \frac{1}{df} \sum_{i=1}^N \left( \frac{y_i - \hat{y}_i}{\sigma_i} \right)^2$$

$$df = N - p - 1$$

**N** : number of data points

**p**: number of fit parameter (model)

given a fitted model:

$\chi_{red}^2$  is a measure of the fit quality

for large (> 50...100) N (number of data points):

≈ **2/3** of the data points should be consistent with the model within their **1σ** error bars

≈ **95%** of the data points should be consistent with the model within their **2σ** error bars

≈ **99.7%** of the data points should be consistent with the model within their **3σ** error bars

$\chi_{red}^2 \approx$

1.0 excellent fit

1.0...1.5 acceptable fit

1.5...1.7 bad fit

>2.0 not acceptable

<<1.0 suspicious, errors are overestimated!



```
from scipy.optimize import curve_fit
```

```
def fun_to_fit(x, a, b, c):  
    return a*x**2 + b*x + c
```

```
ValsBest, Cov = curve_fit(fun_to_fit, x, y)
```

```
ValsBest  
array([ 1.37347967, -0.31706189, -0.20746552])  
      a              b              c
```

```
In [16]: Cov  
Out[16]:  
array([[ 0.15053344, -0.28309719, -0.05890134],  
       [-0.28309719,  0.6630454 , -0.02965514],  
       [-0.05890134, -0.02965514,  0.33825647]])
```

$$\begin{pmatrix} \sigma_a^2 & cov(a, b) & cov(a, c) \\ cov(b, a) & \sigma_b^2 & cov(b, c) \\ cov(c, a) & cov(c, b) & \sigma_c^2 \end{pmatrix}$$





```
from scipy.optimize import curve_fit
```

scipy has the optimization  
tool cuve\_fit

```
def fun_to_fit(x, a, b, c):  
    return a*x**2 + b*x + c
```

defining the model as a  
python function

```
ValsBest, Cov = curve_fit(fun_to_fit, x, y)
```

run the curve fit

```
ValsBest  
array([ 1.37347967, -0.31706189, -0.20746552])  
      a              b              c
```

best values by minimizing MSE  
or related

```
In [16]: Cov  
Out[16]:  
array([[ 0.15053344, -0.28309719, -0.05890134],  
       [-0.28309719,  0.6630454 , -0.02965514],  
       [-0.05890134, -0.02965514,  0.33825647]])
```

$$\begin{pmatrix} \sigma_a^2 & cov(a,b) & cov(a,c) \\ cov(b,a) & \sigma_b^2 & cov(b,c) \\ cov(c,a) & cov(c,b) & \sigma_c^2 \end{pmatrix}$$

covariance matrix for the model parameters

- **non diagonal values should be  $\approx 0$**  (mutually independent)
- **diagonal: squared errors** (see module 8)



```
from scipy.optimize import curve_fit
```

```
def fun_to_fit(x, a, b, c):  
    return a*x**2 + b*x + c
```

```
ValsBest, Cov = curve_fit(fun_to_fit, x, y)
```

```
ValsBest  
array([ 1.37347967, -0.31706189, -0.20746552])  
      a              b              c
```

```
In [16]: Cov  
Out[16]:  
array([[ 0.15053344, -0.28309719, -0.05890134],  
       [-0.28309719,  0.6630454 , -0.02965514],  
       [-0.05890134, -0.02965514,  0.33825647]])
```

$a = 1.37 \pm 0.39$   
 $b = -0.32 \pm 0.81$   
 $c = -0.21 \pm 0.58$

$$\begin{pmatrix} \sigma_a^2 & \text{cov}(a,b) & \text{cov}(a,c) \\ \text{cov}(b,a) & \sigma_b^2 & \text{cov}(b,c) \\ \text{cov}(c,a) & \text{cov}(c,b) & \sigma_c^2 \end{pmatrix}$$

```
error_1sigma = np.sqrt(np.diagonal(Cov))
```



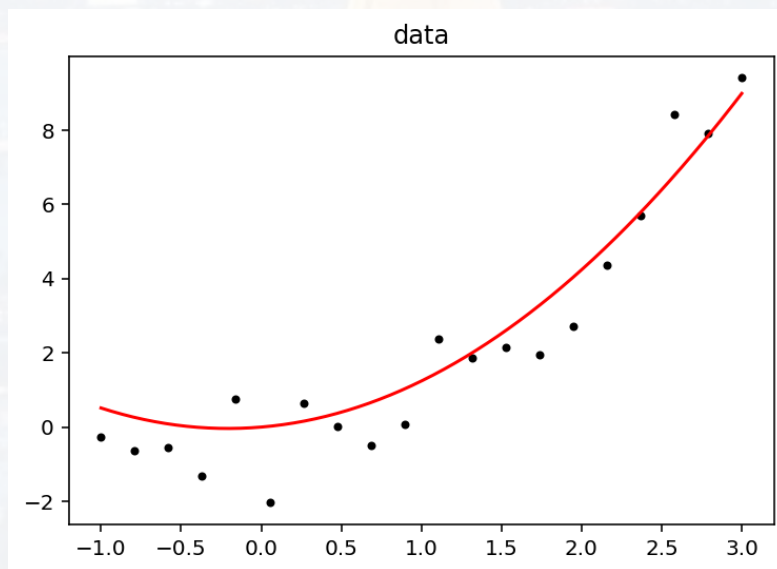
```
from scipy.optimize import curve_fit
```

```
def fun_to_fit(x, a, b, c):  
    return a*x**2 + b*x + c
```

$$\begin{aligned}a &= 1.37 \pm 0.39 \\ b &= -0.32 \pm 0.81 \\ c &= -0.21 \pm 0.58\end{aligned}$$

```
ValsBest, Cov = curve_fit(fun_to_fit, x, y)
```

$$y = 1.37 x^2 - 0.32 x - 0.21$$





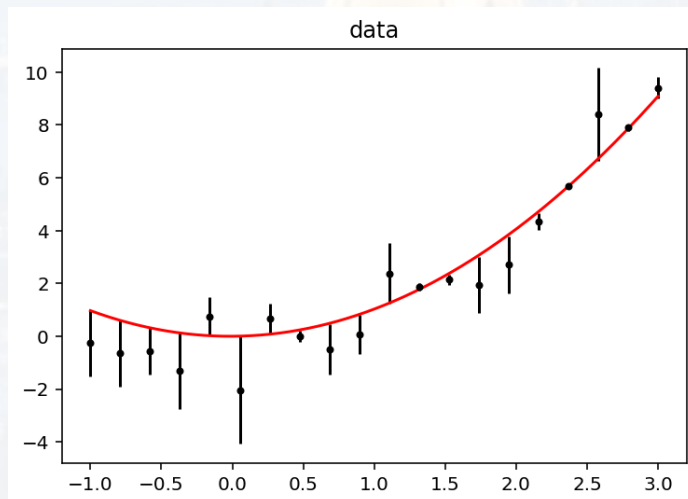
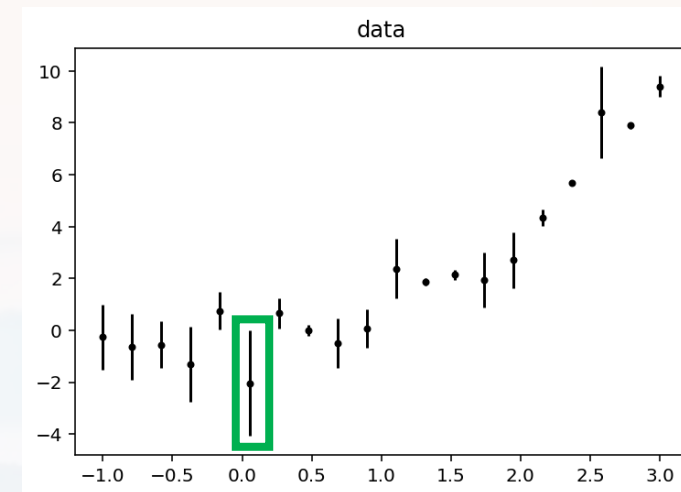
```
from scipy.optimize import curve_fit
```

```
def fun_to_fit(x, a, b, c):  
    return a*x**2 + b*x + c
```

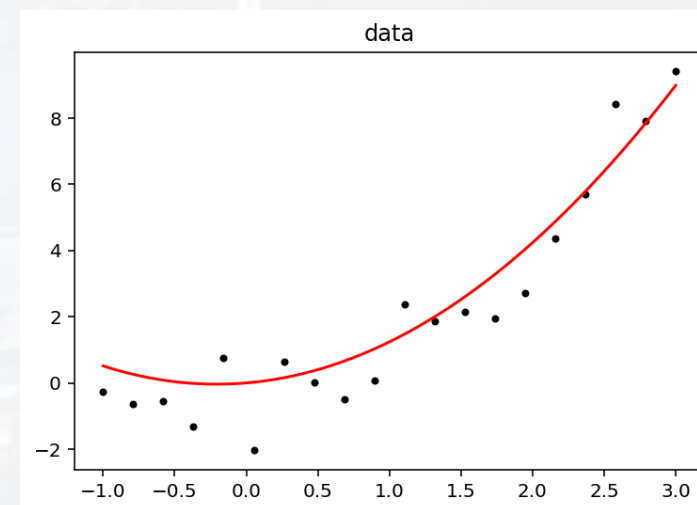
```
ValsBest, Cov = curve_fit(fun_to_fit, x, y)
```

if error bars know → **error weighted fit!**

```
ValsBest, Cov = curve_fit(fun_to_fit, x, y, sigma = err, absolute_sigma = True)
```



VS





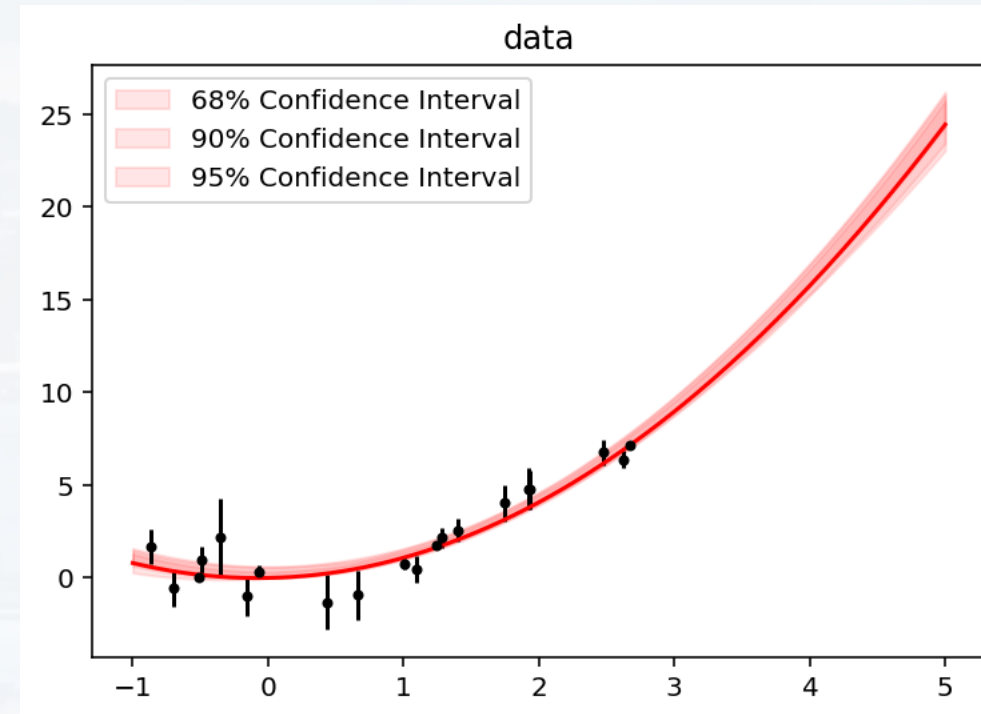
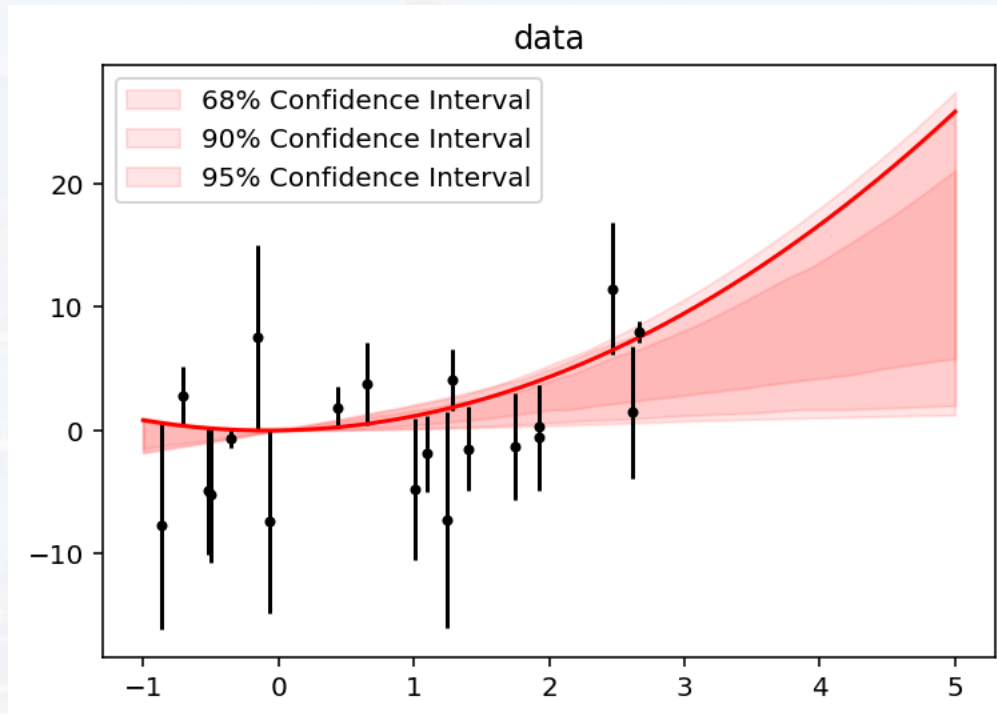


data points have errors and/or fit parameter have uncertainties too

→ implies uncertainties of the model

→ confidence band/interval

$$\begin{aligned} a &= 1.37 \pm 0.39 \\ b &= -0.32 \pm 0.81 \\ c &= -0.21 \pm 0.58 \end{aligned}$$





data points have errors and/or fit parameter have uncertainties too

- implies uncertainties of the model
- confidence band/interval
- idea: **bootstrapping!**

$$a = 1.37 \pm 0.39$$

$$b = -0.32 \pm 0.81$$

$$c = -0.21 \pm 0.58$$



data points have errors and/or fit parameter have uncertainties too

$$\begin{aligned}a &= 1.37 \pm 0.39 \\b &= -0.32 \pm 0.81 \\c &= -0.21 \pm 0.58\end{aligned}$$

→ implies uncertainties of the model

→ confidence band/interval

→ idea: **bootstrapping!**

if **no errors of  $y_i$  known**

→ assuming that fitted parameters follow a **normal distribution**, i.e.  $a = 1.37 \pm 0.39$   
where  $\mu_a = 1.37$  and  $\sigma_a = 0.39$  and so on...

→ **varying** the parameters **within their errors** using `np.random.normal( $\mu_a$ ,  $\sigma_a$ , N)`  
N times

→ for each **N**, **generating a curve fit**

→ from set of N curve fits → **calculating percentiles** for confidence band/ interval



data points have errors and/or fit parameter have uncertainties too

→ implies uncertainties of the model

→ confidence band/interval

→ idea: **bootstrapping!**

$$a = 1.37 \pm 0.39$$

$$b = -0.32 \pm 0.81$$

$$c = -0.21 \pm 0.58$$

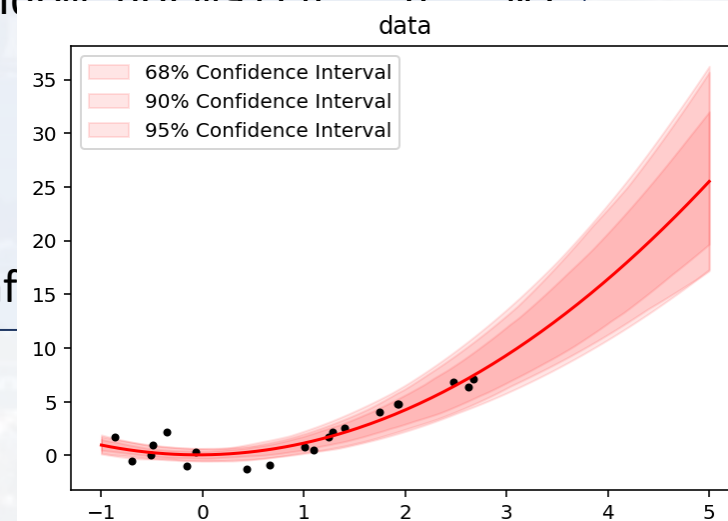
if **no errors of  $y_i$  known**

→ assuming that fitted parameters follow a **normal distribution**, i.e.  $a = 1.37 \pm 0.39$  where  $\mu_a = 1.37$  and  $\sigma_a = 0.39$  and so on...

→ **varying** the parameters **within their errors** using `np.random.normal( $\mu$ ,  $\sigma$ , N)` **N times**

→ for each **N**, **generating a curve fit**

→ from set of N curve fits → **calculating percentiles** for conf







data points have errors and/or fit parameter have uncertainties too

→ implies uncertainties of the model

→ confidence band/interval

→ idea: **bootstrapping!**

$$a = 1.37 \pm 0.39$$

$$b = -0.32 \pm 0.81$$

$$c = -0.21 \pm 0.58$$

if **errors of  $y_i$  known**

→ assuming that errors of  $y_i$  follow a **normal distribution**, i.e.  $y_i(\text{boot}) = y_i \pm \sigma_i$

→ **varying** all  $y_i$  **within their errors** using `np.random.normal( $y_i$ ,  $\sigma_i$ , N)`  
N times

→ for each **N**, **generating a curve fit**

→ from set of N curve fits → **calculating percentiles** for confidence band/ interval



data points have errors and/or fit parameter have uncertainties too

→ implies uncertainties of the model

→ confidence band/interval

→ idea: **bootstrapping!**

$$a = 1.37 \pm 0.39$$

$$b = -0.32 \pm 0.81$$

$$c = -0.21 \pm 0.58$$

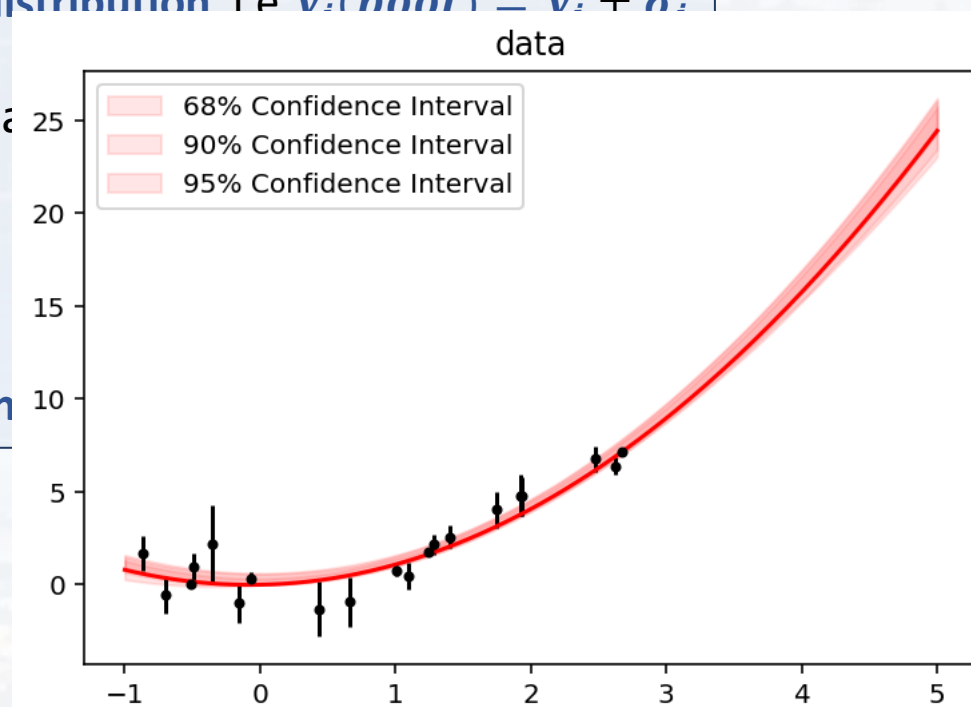
if **errors of  $y_i$  known**

→ assuming that errors of  $y_i$  follow a **normal distribution** i.e.  $v_i(\text{boot}) = v_i + \sigma_i$

→ **varying** all  $y_i$  **within their errors** using np.random.normal  
N times

→ for each **N**, **generating a curve fit**

→ from set of N curve fits → **calculating percentiles**

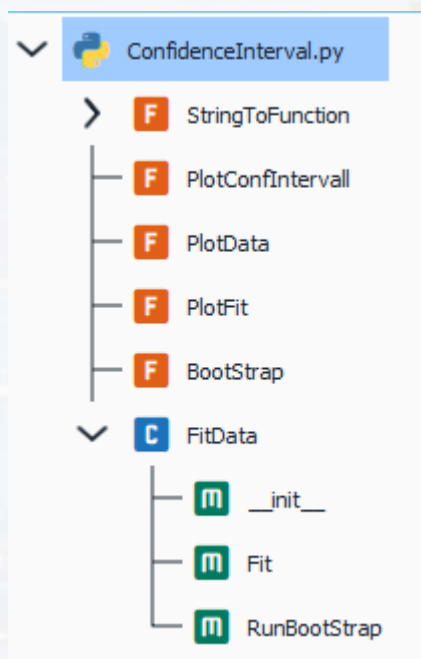




## bootstrapping

explore the .py file

ConfidenceInterval.py



### USAGE:

```

...generating a test sample:
....
....x.....=np.linspace(-1,3,20)
....err.....=np.random.normal(0,1,(len(x),))#1sigma errorbars
....y.....=x**2+err
....errorbars=abs(err)
....
....
....1).plotting data
.....
....F1=FitData(x,y)
....F2=FitData(x,y,errorbars)
....F3=FitData(x,y,errorbars,time='[s]',pressure='[MPa]')
....
....
....2).fitting data (returns best values of fitted params, 1sigma confidence and
.....reduced chi2 if errorbars given, MSE else)
.....
....res1=F1.Fit()
....res2=F2.Fit()
....res2=F3.Fit()
....
....res12=F1.Fit("a*x**2",[1],(-0.5,10))
....
....
....3).Bootstrapping (either varying within errorbars or within conf of fitted
.....params)
.....
....F1.RunBootStrap()
....F2.RunBootStrap()
....F3.RunBootStrap()
....
....F1.RunBootStrap(100,[90,95],np.linspace(-1,5,200))
....F3.RunBootStrap(100,[90,95],np.linspace(-1,5,200))
....

```



Thank you very much for your attention!

