

Lecture 11:

Data Fitting and Regression



Markus Hohle

University of California, Berkeley

**Numerical Methods for
Computational Science**

Course Map

Week 1:	Introduction to Scientific Computing and Python Libraries
Week 2:	Linear Algebra Fundamentals
Week 3:	Vector Calculus
Week 4:	Numerical Differentiation and Integration
Week 5:	Solving Nonlinear Equations
Week 6:	Probability Theory Basics
Week 7:	Random Variables and Distributions
Week 8:	Statistics for Data Science
Week 9:	Eigenvalues and Eigenvectors
Week 10:	Simulation and Monte Carlo Method
Week 11:	Data Fitting and Regression
Week 12:	Optimization Techniques
Week 13:	Machine Learning Fundamentals




Outline

Linear Regression

- The Math
- What is Linear?
- Stats
- a Python Example

Logistic Regression

Curve Fitting



“Did you really have to show the error bars?”




Outline

Linear Regression

- The Math
- What is Linear?
- Stats
- a Python Example

Logistic Regression

Curve Fitting



“Did you really have to show the error bars?”



problem: a parameter y depends on many **regressors** x_n

goal: finding a model that tells us **how** y depends on each x_n
deriving a model to **predict** y from new data points

The Math


What is Linear?

Stats

a Python Example

N

		x_1	x_2	x_3	x_4	x_5	y_k
	Index	molecular_weight	electronegativity	bond_lengths	num_hydrogen_bonds	logP	toxicity_score
k^{th} data point	0	341.704	2.65585	3.09407	2	9.11147	80.9281
	1	335.951	3.22262	2.89039	7	8.92848	83.4911
	2	235.203	2.44115	2.48203	1	6.49731	61.8406
	3	246.505	2.76656	2.71547	7	7.45089	57.0538
	4	437.939	3.4801	3.59569	3	10.9156	131.326



idea: data point y_k in N dimensional space

$$\rightarrow y_k = f(x_1, \dots, x_n, \dots, x_N) + \epsilon \quad \text{for each data point } k$$

The Math

What is Linear?

Stats

a Python Example

ansatz:

$$y_k = \beta_0 + \sum_{n=1}^N \beta_n x_n + \epsilon$$

linear combination

y : response

x : regressors (assumed to be independent)

β : factors (how a regressor contributes to the response)


β_0 : intercept

ϵ : error (stochasticity of the data, assumed to be normally dist.)

N

x_1 x_2 x_3 x_4 x_5 y_k

	Index	molecular_weight	electronegativity	bond_lengths	num_hydrogen_bonds	logP	toxicity_score
k^{th} data point ↓	0	341.704	2.65585	3.09407	2	9.11147	80.9281
	1	335.951	3.22262	2.89039	7	8.92848	83.4911
	2	235.203	2.44115	2.48203	1	6.49731	61.8406
	3	246.505	2.76656	2.71547	7	7.45089	57.0538
	4	437.939	3.4801	3.59569	3	10.9156	131.326





linear $\neq$ not curved

$$y_k = \beta_0 + \sum_{n=1}^N \beta_n x_n^n + \epsilon$$

...is still linear

just define: $\bar{x}_n := x_n^n$

$$y_k = \beta_1 x_n^{\beta_2}$$

...is still linear

$$\begin{aligned} \text{just use log: } \bar{y}_k &= \log(y_k) = \log(\beta_1) + \beta_2 \log(x_n) \\ &= \bar{\beta}_1 + \beta_2 \bar{x}_n \end{aligned}$$

As long as we can recover the linear structure by any transformation $\rightarrow$ it is linear

in part. log scaling is quite common examples:

- log fold change (DESeq/RNASeq)
- log odds ratio (comparing models, HMM)
- sound $\rightarrow$ dB is a log unit
- log incidence rates (medical studies)
- percentiles (medical studies)
-

The Math

What is Linear?

Stats

a Python Example

y:	response
x:	regressors
β :	factors
β_0 :	intercept
ϵ :	error



...what is **not** linear?



$$y_k = \beta_0 + \beta_1 x_n^{\beta_2} \quad \text{log trick does not work here}$$

general: linear refers to the **factors**


$$y_k = \beta_0 + \beta_1 x_1 + \beta_2 x_2 \quad \text{2D plane in 3D space}$$

$$y_k = \beta_0 + \beta_1 x_1^2 + \beta_2 x_2^2 \quad \text{2D parabolic}$$

$$y_k = \beta_0 + \beta_1 x_1^2 - \beta_2 x_2^2 \quad \text{2D hyperbolic}$$



...and many more...



The Math
What is Linear?
Stats
a Python Example

y:	response
x:	regressors
β :	factors
β_0 :	intercept
ϵ :	error


all linear

$$y_k = \beta_0 + \sum_{n=1}^N \beta_n x_n + \epsilon$$



for K data points in N dimensional space

$$y_k = \beta_0 + \sum_{n=1}^N \beta_n x_n + \epsilon$$



The Math
What is Linear?
Stats
a Python Example

$$\underbrace{\begin{pmatrix} y_1 \\ \vdots \\ y_k \\ \vdots \\ y_K \end{pmatrix}}_Y = \underbrace{\begin{pmatrix} 1 & x_{11} & x_{12} & \dots & x_{1n} & \dots & x_{1N} \\ \vdots & \vdots & \vdots & & \vdots & & \vdots \\ 1 & x_{k1} & & & x_{kn} & & \\ \vdots & \vdots & & & \vdots & & \vdots \\ 1 & \vdots & & & \vdots & & \vdots \\ 1 & x_{K1} & x_{K2} & \dots & x_{Kn} & \dots & x_{KN} \end{pmatrix}}_X \underbrace{\begin{pmatrix} \beta_0 \\ \beta_1 \\ \vdots \\ \beta_n \\ \vdots \\ \beta_N \end{pmatrix}}_{\beta} + \underbrace{\begin{pmatrix} \epsilon_1 \\ \epsilon_2 \\ \vdots \\ \epsilon_k \\ \vdots \\ \epsilon_K \end{pmatrix}}_{\epsilon}$$

y: response
x: regressors
 β : factors
 β_0 : intercept
 ϵ : error


$$Y = X\beta + \epsilon$$

N							
$x_1 \quad x_2 \quad x_3 \quad x_4 \quad x_5 \quad y_k$							
k^{th} data point	Index	molecular_weight	electronegativity	bond_lengths	num_hydrogen_bonds	logP	toxicity_score
	0	341.704	2.65585	3.09407	2	9.11147	80.9281
	1	335.951	3.22262	2.89039	7	8.92848	83.4911
	2	235.203	2.44115	2.48203	1	6.49731	61.8406
	3	246.505	2.76656	2.71547	7	7.45089	57.0538
	4	437.939	3.4801	3.59569	3	10.9156	131.326



for K data points in N dimensional space

$$y_k = \beta_0 + \sum_{n=1}^N \beta_n x_n + \epsilon$$



The Math
What is Linear?
Stats
a Python Example

$$\underbrace{\begin{pmatrix} y_1 \\ \vdots \\ y_k \\ \vdots \\ y_K \end{pmatrix}}_Y = \underbrace{\begin{pmatrix} 1 & x_{11} & x_{12} & \dots & x_{1n} & \dots & x_{1N} \\ \vdots & \vdots & \vdots & & \vdots & & \vdots \\ 1 & x_{k1} & & & x_{kn} & & \\ \vdots & \vdots & & & \vdots & & \vdots \\ 1 & \vdots & & & \vdots & & \vdots \\ 1 & x_{K1} & x_{K2} & \dots & x_{Kn} & \dots & x_{KN} \end{pmatrix}}_X \underbrace{\begin{pmatrix} \beta_0 \\ \beta_1 \\ \vdots \\ \beta_n \\ \vdots \\ \beta_N \end{pmatrix}}_{\beta} + \underbrace{\begin{pmatrix} \epsilon_1 \\ \epsilon_2 \\ \vdots \\ \epsilon_k \\ \vdots \\ \epsilon_K \end{pmatrix}}_{\epsilon}$$

$$Y = X\beta + \epsilon$$

y:	response
x:	regressors
β:	factors
β_0:	intercept
ϵ:	error

fitting: finding the best β in terms of minimizing the errors

$$(Y - X\beta)^T(Y - X\beta) = \sum_k \epsilon_k^2$$

the model

$$\frac{\partial}{\partial \beta} \sum_k \epsilon_k^2 = 0 \longrightarrow \beta_{best} = \hat{\beta} = (X^T X)^{-1} X^T Y \longrightarrow \hat{Y} = X\hat{\beta} = X(X^T X)^{-1} X^T Y$$

X and Y are all observables



$$Y = X\beta + \varepsilon$$

finding the best β in terms of minimizing the errors

$$(Y - X\beta)^T(Y - X\beta) = \sum_k \varepsilon_k^2$$

the model

$$\hat{Y} = X\hat{\beta} = X \underbrace{(X^T X)^{-1} X^T}_{\text{hat matrix } H} Y$$

hat matrix H

some properties of the hat matrix:

- $H = H^T$ (symmetry)
- $HH = H \rightarrow H^n = H$ (idempotency)


evaluating the fit:

$$\hat{\varepsilon} = Y - X\hat{\beta} = Y - \hat{Y} = (I - H)Y$$

$$\hat{\varepsilon}^T \hat{\varepsilon} = [(I - H)Y]^T (I - H)Y = Y^T (I - H)^T (I - H)Y = Y^T (I - H)Y$$

sum of squared errors (SSE)

$$\hat{\sigma}^2 = \frac{\hat{\varepsilon}^T \hat{\varepsilon}}{K - N}$$



degrees of freedom

variance or

mean of squared errors (MSE)

The Math
What is Linear?

Stats
a Python Example

y:	response
x:	regressors
β :	factors
β_0 :	intercept
ε :	error



summary:

the model:

$$Y = X\beta + \varepsilon$$

the fit:

$$\hat{Y} = X\hat{\beta} = X(X^T X)^{-1} X^T Y$$

SSE:

$$\hat{\varepsilon}^T \hat{\varepsilon} = Y^T (I - H) Y \quad (\text{after the fit})$$

variance or MSE:

$$\hat{\sigma}^2 = \frac{\hat{\varepsilon}^T \hat{\varepsilon}}{K - N} \quad (\text{after the fit})$$

often fit quality is judged by

$$R^2 := 1 - \frac{\sum_k (\hat{y}_k - y_k)^2}{\sum_k (y_k - \langle y \rangle)^2}$$

or adjusted R^2

$$\bar{R}^2 := R^2 - (1 - R^2) \frac{N}{K - N - 1}$$

and it is said that the fit is good if R^2 is close to one....

...but that is not true...

The Math
What is Linear?

Stats
a Python Example

y:	response
x:	regressors
β:	factors
β_0:	intercept
ε:	error




$$R^2 = 1 - \frac{\sum_{i=1}^N (y_i - \hat{y}_i)^2}{\sum_{i=1}^N (y_i - \bar{y})^2}$$

variance data vs model
(aka residual sum of squares)

variance of the data
(aka total sum of squares)

The Math
What is Linear?
Stats
a Python Example

Note: do not confuse R^2 with Pearsons coefficient: $\rho = \frac{cov(x,y)}{\sqrt{var(x)var(y)}}$



data variance can be huge
(i. e. exponential functions)
→ R^2 could be around 1.0
even if fit is completely off!

$\bar{y}$: mean of the data point values

$\bar{y}$: mean of the data point values
 y_i : measured value of data point
 σ_i : statistical error of y_i (often aka ey_i)
 $\hat{y}_i$: prediction by the model *after the fit*
 N : number of data points
 p : number of fit parameter




$$R^2 = 1 - \frac{\sum_{i=1}^N (y_i - \hat{y}_i)^2}{\sum_{i=1}^N (y_i - \bar{y})^2}$$

variance data vs model
(aka residual sum of squares)


variance of the data
(aka total sum of squares)

The Math
What is Linear?
Stats
a Python Example



$\bar{y}$: mean of the data point values

data variance can be huge
(i. e. exponential functions)
→ R^2 could be around 1.0
even if fit is completely off!



variance data vs model
(aka residual sum of squares)

variance of the data
(aka total sum of squares)

$$\approx 1 \rightarrow R^2 = 0$$

although the fit is good!

$\bar{y}$: mean of the data point values
 y_i : measured value of data point
 σ_i : statistical error of y_i (often aka ey_i)
 $\hat{y}_i$: prediction by the model *after the fit*
 N : number of data points
 p : number of fit parameter




$$R^2 = 1 - \frac{\sum_{i=1}^N (y_i - \hat{y}_i)^2}{\sum_{i=1}^N (y_i - \bar{y})^2}$$

variance data vs model
(aka residual sum of squares)

variance of the data
(aka total sum of squares)

The Math
What is Linear?
Stats
a Python Example



all plots: same R^2

conclusion:

R^2 is **not** a measure of the fit quality (but χ^2 is)

Given a good fit, R^2 tells how strong the dependent variable responds to the independent variable

Also, Wiki is full of examples...
...and warnings (see “caveats” therein)



see `Walk_Through_LinRegression.ipynb`

```
import numpy as np
```

```
import pandas as pd
```

```
import matplotlib.pyplot as plt
```

```
import seaborn as sns
```

```
import pylab
```

```
import scipy.stats as stats
```

```
import statsmodels.api as sm
```

```
from statsmodels.formula.api import ols
```

```
from sklearn.preprocessing import MinMaxScaler
```

The Math
What is Linear?
Stats
a Python Example

reading .xlsx
.csv
.txt
...

standard plots

fancy plots:
here a pair-
plot

Q-Q plot

the actual
super tool for
superb data
analysis

scaling and normalizing



```
Train = pd.read_csv("molecular_train_gbc.csv")  
Test  = pd.read_csv("molecular_test_gbc.csv")
```

- 1) loading data
- 2) plotting data
- 3) scaling data
- 4) fitting model
- 5) evaluating model

	x_1	x_2	x_3	x_4	x_5	y_k
Index	molecular_weight	electronegativity	bond_lengths	num_hydrogen_bonds	logP	toxicity_score
0	341.704	2.65585	3.09407	2	9.11147	80.9281
1	335.951	3.22262	2.89039	7	8.92848	83.4911
2	235.203	2.44115	2.48203	1	6.49731	61.8406
3	246.505	2.76656	2.71547	7	7.45089	57.0538
4	437.939	3.4801	3.59569	3	10.9156	131.326

$$y_k = \beta_0 + \sum_{n=1}^N \beta_n x_n + \epsilon$$


y: toxicity_score
 x_n : molecular_weight, electronegativity,
bond_lengths, num_hydrogen_bonds, logP



```
Train = pd.read_csv("molecular_train_gbc.csv")  
Test  = pd.read_csv("molecular_test_gbc.csv")
```

- 1) loading data
- 2) plotting data
- 3) scaling data
- 4) fitting model
- 5) evaluating model

```
out = sns.pairplot(Train, kind = "kde", \  
                  plot_kws = {'color': [176/255, 224/255, 230/255]}, \  
                  diag_kws = {'color': 'black'})  
out.map_offdiag(plt.scatter, color = 'black')
```





```
Train = pd.read_csv("molecular_train_gbc.csv")  
Test  = pd.read_csv("molecular_test_gbc.csv")
```

```
out = sns.pairplot(Train, kind = "kde", \  
                  plot_kws = {'color': [176/255, 224/255, 230/255]}, \  
                  diag_kws = {'color': 'black'})  
out.map_offdiag(plt.scatter, color = 'black')
```

```
scaler = MinMaxScaler(feature_range = (0, 1))  
TrainS = scaler.fit_transform(Train)  
TestS  = scaler.transform(Test)
```

the scaler returns an np.array
→ convert back to data frame

```
TrainS = pd.DataFrame(TrainS, columns = Train.columns)  
TestS  = pd.DataFrame(TestS,  columns = Train.columns)
```

- 1) loading data
- 2) plotting data
- 3) scaling data
- 4) fitting model
- 5) evaluating model



```
TrainS = pd.DataFrame(TrainS, columns = Train.columns)
TestS   = pd.DataFrame(TestS,  columns = Train.columns)
```

```
equation = 'toxicity_score ~ ' + '+'.join(Train.columns[:-1])
print(equation)
```

- 1) loading data
- 2) plotting data
- 3) scaling data
- 4) fitting model
- 5) evaluating model

$$y_k = \beta_0 + \sum_{n=1}^N \beta_n x_n + \epsilon$$

```
toxicity_score ~      molecular_weight + electronegativity +
                      bond_lengths + num_hydrogen_bonds + logP
```

```
my_model = ols(equation, data = TrainS).fit()
my_model.summary()
```

OLS (ordinary least squares)



my_model.summary()

- 1) loading data
- 2) plotting data
- 3) scaling data
- 4) fitting model
- 5) evaluating model

number of data points
is much larger than
the number of regressors
→ degree of freedom
approx. no of obs

OLS Regression Results

not the fit quality!

Dep. Variable:	toxicity_score	R-squared:	0.790
Model:	OLS	Adj. R-squared:	0.789
Method:	Least Squares	F-statistic:	597.5
Date:	Fri, 13 Sep 2024	Prob (F-statistic):	3.34e-266
Time:	20:57:10	Log-Likelihood:	1013.0
No. Observations:	800	AIC:	-2014.
Df Residuals:	794	BIC:	-1986.
Df Model:	5		

Covariance Type: nonrobust

p-values for factors

	coef	std err	t	P> t	[0.025
Intercept	0.1494	0.012	12.533	0.000	0.126
molecular_weight	0.7961	0.089	8.982	0.000	0.622
electronegativity	-0.1682	0.015	-11.591	0.000	-0.197
bond_lengths	0.0204	0.049	0.417	0.677	-0.076
num_hydrogen_bonds	0.0035	0.008	0.458	0.647	-0.011
logP	0.1246	0.072	1.723	0.085	-0.017

Omnibus:	2.249	Durbin-Watson:	1.984
Prob(Omnibus):	0.325	Jarque-Bera (JB):	2.240
Skew:	-0.129	Prob(JB):	0.326
Kurtosis:	2.980	Cond. No.	65.6

p-value for
constant model

p-values for
factors

2σ conf range of
factors

$$y_k = \beta_0 + \sum_{n=1}^N \beta_n x_n + \epsilon$$

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.



more accurate: determining the **p-values for the factors using ANOVA** for the corresponding residuals

```
table = sm.stats.anova_lm(my_model, typ = 1)
print(table)
```

- 1) loading data
- 2) plotting data
- 3) scaling data
- 4) fitting model
- 5) evaluating model

$$y_k = \beta_0 + \sum_{n=1}^N \beta_n x_n + \epsilon$$

	df	sum_sq	mean_sq	F	PR(>F)	vs from t-test
molecular_weight	1.0	13.346285	13.346285	2847.525516	8.024085e-265	0.0000
electronegativity	1.0	0.640388	0.640388	136.631363	3.085962e-29	0.0000
bond_lengths	1.0	0.000684	0.000684	0.145954	7.025342e-01	0.6766
num_hydrogen_bonds	1.0	0.000703	0.000703	0.150055	6.985866e-01	0.6473
logP	1.0	0.013917	0.013917	2.969353	8.524510e-02	0.0852
Residual	794.0	3.721459	0.004687	NaN	NaN	




```
residuals = my_model.resid
```

```
plt.hist(residuals, color = 'w', edgecolor = 'black')  
plt.title('fit residuals')  
plt.ylabel('#')  
plt.xlabel('value')  
plt.show()
```

- 1) loading data
- 2) plotting data
- 3) scaling data
- 4) fitting model
- 5) evaluating model

$$y_k = \beta_0 + \sum_{n=1}^N \beta_n x_n + \epsilon$$



residuals approx.
normally distributed
around $\mu = 0$



```
residuals = my_model.resid
```


```
plt.hist(residuals, color = 'w', edgecolor = 'black')  
plt.title('fit residuals')  
plt.ylabel('#')  
plt.xlabel('value')  
plt.show()
```

```
stats.probplot(residuals, dist = "norm", plot = pylab)  
pylab.show()
```

- 1) loading data
- 2) plotting data
- 3) scaling data
- 4) fitting model
- 5) evaluating model

residuals approx.
normally distributed
around $\mu = 0$

$$y_k = \beta_0 + \sum_{n=1}^N \beta_n x_n + \epsilon$$





```
Ypred = my_model.predict(TestS)
```

```
higher = np.max([Ypred, TestS.toxicity_score])
```

```
lower = np.min([Ypred, TestS.toxicity_score])
```

```
plt.plot([lower, higher], [lower, higher], c = [0, 0, 0, 0.2],\n         linewidth = 4)
```

```
plt.scatter(TestS.toxicity_score, Ypred, marker = '.', c = 'k')
```


```
plt.ylabel('prediction')
```

```
plt.xlabel('toxicity score')
```

```
plt.show()
```

- 1) loading data
- 2) plotting data
- 3) scaling data
- 4) fitting model
- 5) evaluating model

$$y_k = \beta_0 + \sum_{n=1}^N \beta_n x_n + \epsilon$$





```
Ypred = my_model.predict(TestS)
```

```
higher = np.max([Ypred, TestS.toxicity_score])
```

```
lower = np.min([Ypred, TestS.toxicity_score])
```


```
plt.plot([lower, higher], [lower, higher], c = [0, 0, 0, 0.2], linewidth = 4)
```

```
plt.scatter(TestS.toxicity_score, Ypred, marker = '.', c = 'k')
```

```
plt.ylabel('prediction')
```

```
plt.xlabel('toxicity score')
```

```
plt.show()
```



```
mean_dev = np.sum( abs(TestS.toxicity_score - Ypred) )/len(Ypred)  
print(mean_dev)
```

5%

- 1) loading data
- 2) plotting data
- 3) scaling data
- 4) fitting model
- 5) evaluating model

$$y_k = \beta_0 + \sum_{n=1}^N \beta_n x_n + \epsilon$$




Outline

Linear Regression

- The Math
- What is Linear?
- Stats
- a Python Example

Logistic Regression

Curve Fitting



“Did you really have to show the error bars?”



linear model: regressors are continuous or categorical,
response is continuous

logistic model: response is **categorical**

y :	response
x :	regressors (assumed to be independent)
β :	factors
β_0 :	intercept
ϵ :	error (stochasticity of the data, assumed to be normally dist.)

Index	molecular_weight	electronegativity	bond_lengths	num_hydrogen_bonds	logP	label
0	341.704	2.65585	3.09407	2	9.11147	Toxic
1	335.951	3.22262	2.89039	7	8.92848	Toxic
2	235.203	2.44115	2.48203	1	6.49731	Non-Toxic
3	246.505	2.76656	2.71547	7	7.45089	Non-Toxic
4	437.939	3.4801	3.59569	3	10.9156	Non-Toxic



linear model: regressors are continuous or categorical,
response is continuous

logistic model: response is **categorical**

y:	response
x:	regressors (assumed to be independent)
β :	factors
β_0 :	intercept
ϵ :	error (stochasticity of the data, assumed to be normally dist.)

dichotomic model: **probability** to be in state A) $\rightarrow p$

probability to be in state B) $\rightarrow 1 - p$

label
Toxic
Toxic
Non-Toxic
Non-Toxic
Non-Toxic

ansatz:

$$\log \left(\frac{p}{1-p} \right) = \beta_0 + \sum_{n=1}^N \beta_n x_n + \epsilon$$

log odds ratio: linear model



dichotomic model:

label
Toxic
Toxic
Non-Toxic
Non-Toxic
Non-Toxic

probability to be in state A) $\rightarrow p$

probability to be in state B) $\rightarrow 1 - p$

y:	response
x:	regressors (assumed to be independent)
β :	factors
β_0 :	intercept
ϵ :	error (stochasticity of the data, assumed to be normally dist.)

ansatz:

$$\log\left(\frac{p}{1-p}\right) = \beta_0 + \sum_{n=1}^N \beta_n x_n + \epsilon$$

log odds ratio: linear model

$\rightarrow$ probability for being in a certain state

$$p = \frac{e^{\beta_0 + \beta_1 x_1 + \dots}}{1 + e^{\beta_0 + \beta_1 x_1 + \dots}}$$

often:

$$\text{logit}(p) = \log\left(\frac{p}{1-p}\right)$$

examples:

- probability that a gene has been mutated
- probability of being diseased (cancer, alzheimer etc) as function of age, environmental influence etc ...
- Verhulst equation: $N(t) = N_0 \frac{e^{rt}}{C + e^{rt}}$
- activation functions in ANNs



$$\log\left(\frac{p}{1-p}\right) = \beta_0 + \sum_{n=1}^N \beta_n x_n + \epsilon$$

Note: one can derive the logit function from max. entropy too!

y:	response
x:	regressors (assumed to be independent)
β :	factors
β_0 :	intercept
ϵ :	error (stochasticity of the data, assumed to be normally dist.)

$$p = \frac{e^{\beta_0 + \beta_1 x_1 + \dots}}{1 + e^{\beta_0 + \beta_1 x_1 + \dots}} = \frac{1}{1 + e^{-\beta_0 - \beta_1 x_1 - \dots}}$$

onset of Alzheimer's disease (AD) is a function of **age** and years spent in **education**
(and other risk factors we ignore here for the sake of simplicity)

education: $d = x_1$ [yrs]

age: $a = x_2$ [yrs]

model:
$$p_{AD} = \frac{1}{1 + e^{-\beta_0 - \beta_1 d - \beta_2 a}}$$

+ data set + fit $\rightarrow$

$$\begin{aligned}\beta_0 &= +0.1 \\ \beta_1 &= -1.5 \\ \beta_2 &= +0.12\end{aligned}$$

- positive value $\rightarrow$ increasing p
- negative value $\rightarrow$ decreasing p
- intercept: "background" prevalence, not related to environmental/internal conditions



model:
$$p_{AD} = \frac{1}{1 + e^{-\beta_0 - \beta_1 d - \beta_2 a}}$$

education: $d = x_1$ [yrs] $\beta_0 = +0.1$
age: $a = x_2$ [yrs] $\beta_1 = -1.5$
 $\beta_2 = +0.12$

example: **65yrs** old person, **8yrs** spent in education

$$\rightarrow p_{AD} = 1.6\%$$

65yrs old person, **13yrs** spent in education


$$\rightarrow p_{AD} = 0.001\%$$

How does education compensate aging?

$$p_{AD}(d + \bar{d}, a + \bar{a}) = p_{AD}(d, a)$$

$$\rightarrow \bar{a} = 12.5 \bar{d}$$

y:	response
x:	regressors (assumed to be independent)
β :	factors
β_0 :	intercept
ε :	error (stochasticity of the data, assumed to be normally dist.)



hence, one more year prolonged education compensates
12.5 years of aging

(warning: don't confuse correlation with causation here!)



model:
$$p_{AD} = \frac{1}{1 + e^{-\beta_0 - \beta_1 d - \beta_2 a}}$$

education: $d = x_1$ [yrs]

age: $a = x_2$ [yrs]

$$\beta_0 = +0.1$$

$$\beta_1 = -1.5$$

$$\beta_2 = +0.12$$

y:	response
x:	regressors (assumed to be independent)
β :	factors
β_0 :	intercept
ε :	error (stochasticity of the data, assumed to be normally dist.)

How does the risk of onset changes *per year*?


relative change:

$$\frac{p_{AD}(a+1) - p_{AD}(a)}{p_{AD}(a)} \approx e^{\beta_2} - 1 \approx 12.7\%$$

$p_{AD} \ll 1$ (hence, for small Δa and “young” ages, i. e. below ≈ 80 yrs)

the risk of getting AD increases by 12.7% every year

(warning: does not mean that it increases by 127% in ten yrs – we made an approximation!)





model:
$$p_{AD} = \frac{1}{1 + e^{-\beta_0 - \beta_1 d - \beta_2 a}}$$

education: $d = x_1$ [yrs]

age: $a = x_2$ [yrs]

$$\beta_0 = +0.1$$

$$\beta_1 = -1.5$$

$$\beta_2 = +0.12$$


y:	response
x:	regressors (assumed to be independent)
β :	factors
β_0 :	intercept
ε :	error (stochasticity of the data, assumed to be normally dist.)

How does the risk of onset changes *per year*?

more precise: relative change of the odds ratio

$$\frac{\frac{\partial}{\partial x_i} \left(\frac{p_{AD}}{1 - p_{AD}} \right)}{\frac{p_{AD}}{1 - p_{AD}}} = \beta_i$$

x_i is the desired regressor,
for example, age again (x_2)



the factors β_i indicate how strong (and in which direction) p changes wrt a regressor x_i



let us return to the molecule data set:

```
Train = pd.read_csv("molecular_train_gbc_cat.csv")  
Test  = pd.read_csv("molecular_test_gbc_cat.csv")
```

Index	molecular_weight	electronegativity	bond_lengths	num_hydrogen_bonds	logP	label
0	341.704	2.65585	3.09407	2	9.11147	Toxic
1	335.951	3.22262	2.89039	7	8.92848	Toxic
2	235.203	2.44115	2.48203	1	6.49731	Non-Toxic
3	246.505	2.76656	2.71547	7	7.45089	Non-Toxic
4	437.939	3.4801	3.59569	3	10.9156	Non-Toxic

```
import statsmodels.api as sm
```




it is the same data set → plotting and scaling is as before

```
X = sm.add_constant(TrainS)
```

adding the
intercept

```
Y = pd.get_dummies(Train['Label'])
```

Python needs
True/False
as categorical

$$p = \frac{e^{\beta_0 + \beta_1 x_1 + \dots}}{1 + e^{\beta_0 + \beta_1 x_1 + \dots}}$$

```
In [48]: print(Y)
   Non-Toxic  Toxic
0         False   True
1         False   True
2          True  False
3          True  False
4          True  False
```

we have two
states: toxic /
non-toxic

```
my_model = sm.GLM(Y, X, family = sm.families.Binomial()).fit()
```

```
my_model.summary()
```

GLM: general
linear model



it is the same data set → plotting and scaling is as before

```
X = sm.add_constant(TrainS)
Y = pd.get_dummies(Train['Label'])
```

```
my_model = sm.GLM(Y, X, family = sm.families.Binomial()).fit()
my_model.summary()
```

Generalized Linear Model Regression Results						
=====						
Dep. Variable:	['Non-Toxic', 'Toxic']		No. Observations:	800		
Model:	GLM		Df Residuals:	794		
Model Family:	Binomial		Df Model:	5		
Link Function:	Logit		Scale:	1.0000		
Method:	IRLS		Log-Likelihood:	-332.82		
Date:	Sat, 14 Sep 2024		Deviance:	665.64		
Time:	20:59:18		Pearson chi2:	1.14e+03		
No. Iterations:	6		Pseudo R-squ. (CS):	0.4243		
Covariance Type:	nonrobust		p-values for factors			
=====						
	coef	std err	z	P> z	[0.025	0.975]

const	6.1641	0.585	10.536	0.000	5.017	7.311
molecular_weight	-10.4920	3.626	-2.893	0.004	-17.599	-3.385
electronegativity	3.2874	0.599	5.492	0.000	2.114	4.461
bond_lengths	0.6736	1.913	0.352	0.725	-3.075	4.422
num_hydrogen_bonds	-0.3082	0.303	-1.018	0.309	-0.902	0.285
logP	-7.6090	2.978	-2.555	0.011	-13.447	-1.771
=====						

$$p = \frac{e^{\beta_0 + \beta_1 x_1 + \dots}}{1 + e^{\beta_0 + \beta_1 x_1 + \dots}}$$

p-value for
constant model


2σ conf range of
factors



accuracy: How **often** did the model make the correct prediction.
cross-entropy: How **certain** was the model when making the prediction.

ideal world:

true label	non-toxic	toxic
	100%	0%
non-toxic	100%	0%
toxic	0%	100%
		predicted label
		non-toxic toxic





more details see [Chem 277](#)



accuracy: How **often** did the model make the correct prediction.
cross-entropy: How **certain** was the model when making the prediction.

real world:





more details see [Chem 277](#)



accuracy: How **often** did the model make the correct prediction.
cross-entropy: How **certain** was the model when making the prediction.

real world:



more details see [Chem 277](#)




Outline

Linear Regression

- The Math
- What is Linear?
- Stats
- a Python Example

Logistic Regression

Curve Fitting




“Did you really have to show the error bars?”




problem: set of data points $y_i = f(x_i)$, ideally each data point y_i has an error σ_i
we know $f(x_i)$

goal: find parameters of $f(x_i)$
once we know the parameters $\rightarrow$ prediction



model: $f(x_i) = m x_i + n$
parameter: m, n




model: $f(x_i) = a x_i^2 + b x_i + c$
parameter: a, b, c



problem: set of data points $y_i = f(x_i)$, ideally each data point y_i has an error σ_i
we know $f(x_i)$

goal: find parameters of $f(x_i)$
once we know the parameters → **prediction**





y_i for data point
outside known interval

y_i for data point
within known interval





problem: set of data points $y_i = f(x_i)$, ideally each data point y_i has an error σ_i
we know $f(x_i)$

goal: find parameters of $f(x_i)$
once we know the parameters $\rightarrow$ prediction



y_i : measured value of data point
 σ_i : statistical error of y_i (often aka ey_i)
 $\hat{y}_i$: prediction by the model *after the fit*
 N : number of data points
 p : number of fit parameter



finding best parameters by minimizing

$$\chi_{red}^2 = \frac{1}{df} \sum_{i=1}^N \left(\frac{y_i - \hat{y}_i}{\sigma_i} \right)^2$$

$df = N - p - 1$
see module 8

y_i : measured value of data point
 σ_i : statistical error of y_i (often aka ey_i)
 $\hat{y}_i$: prediction by the model *after the fit*
 N : number of data points
 p : number of fit parameter

or

$$MSE = \frac{1}{df} \sum_{i=1}^N (y_i - \hat{y}_i)^2$$

if no errors given




finding best parameters by minimizing

$$\chi_{red}^2 = \frac{1}{df} \sum_{i=1}^N \left(\frac{y_i - \hat{y}_i}{\sigma_i} \right)^2 \quad df = N - p - 1$$

see module 8

$\log(\chi_{red}^2)$



or


$$MSE = \frac{1}{df} \sum_{i=1}^N (y_i - \hat{y}_i)^2 \quad \text{if no errors given}$$

finding the global minimum of a higher dimensional non-analytical function

→ see also module 3



from module 8



$$\chi_{red}^2 = \frac{1}{df} \sum_{i=1}^N \left(\frac{y_i - \hat{y}_i}{\sigma_i} \right)^2$$

$$df = N - p - 1$$

N : number of data points

p: number of fit parameter (model)

given a fitted model:

χ_{red}^2 is a measure of the fit quality

for large (> 50...100) N (number of data points):

≈ **2/3** of the data points should be consistent with the model within their **1σ** error bars

≈ **95%** of the data points should be consistent with the model within their **2σ** error bars

≈ **99.7%** of the data points should be consistent with the model within their **3σ** error bars

$\chi_{red}^2 \approx$

1.0 excellent fit

1.0...1.5 acceptable fit

1.5...1.7 bad fit

>2.0 not acceptable

<<1.0 suspicious, errors are overestimated!



```
from scipy.optimize import curve_fit
```

```
def fun_to_fit(x, a, b, c):  
    return a*x**2 + b*x + c
```

```
ValsBest, Cov = curve_fit(fun_to_fit, x, y)
```

```
ValsBest  
array([ 1.37347967, -0.31706189, -0.20746552])  
      a              b              c
```

```
In [16]: Cov  
Out[16]:  
array([[ 0.15053344, -0.28309719, -0.05890134],  
       [-0.28309719,  0.6630454 , -0.02965514],  
       [-0.05890134, -0.02965514,  0.33825647]])
```

$$\begin{pmatrix} \sigma_a^2 & cov(a, b) & cov(a, c) \\ cov(b, a) & \sigma_b^2 & cov(b, c) \\ cov(c, a) & cov(c, b) & \sigma_c^2 \end{pmatrix}$$



```
from scipy.optimize import curve_fit
```

scipy has the optimization tool `cuve_fit`

```
def fun_to_fit(x, a, b, c):  
    return a*x**2 + b*x + c
```

defining the model as a python function

```
ValsBest, Cov = curve_fit(fun_to_fit, x, y)
```

run the curve fit

```
ValsBest  
array([ 1.37347967, -0.31706189, -0.20746552])  
      a              b              c
```

best values by minimizing MSE or related

```
In [16]: Cov  
Out[16]:  
array([[ 0.15053344, -0.28309719, -0.05890134],  
       [-0.28309719,  0.6630454 , -0.02965514],  
       [-0.05890134, -0.02965514,  0.33825647]])
```

$$\begin{pmatrix} \sigma_a^2 & \text{cov}(a,b) & \text{cov}(a,c) \\ \text{cov}(b,a) & \sigma_b^2 & \text{cov}(b,c) \\ \text{cov}(c,a) & \text{cov}(c,b) & \sigma_c^2 \end{pmatrix}$$

covariance matrix for the model parameters

- **non diagonal values should be ≈ 0** (mutually independent)
- **diagonal: squared errors** (see module 8)



```
from scipy.optimize import curve_fit
```

```
def fun_to_fit(x, a, b, c):  
    return a*x**2 + b*x + c
```

```
ValsBest, Cov = curve_fit(fun_to_fit, x, y)
```

```
ValsBest  
array([ 1.37347967, -0.31706189, -0.20746552])  
      a              b              c
```

```
In [16]: Cov  
Out[16]:  
array([[ 0.15053344, -0.28309719, -0.05890134],  
       [-0.28309719,  0.6630454 , -0.02965514],  
       [-0.05890134, -0.02965514,  0.33825647]])
```

$a = 1.37 \pm 0.39$
 $b = -0.32 \pm 0.81$
 $c = -0.21 \pm 0.58$

$$\begin{pmatrix} \sigma_a^2 & \text{cov}(a,b) & \text{cov}(a,c) \\ \text{cov}(b,a) & \sigma_b^2 & \text{cov}(b,c) \\ \text{cov}(c,a) & \text{cov}(c,b) & \sigma_c^2 \end{pmatrix}$$

```
error_1sigma = np.sqrt(np.diagonal(Cov))
```




```
from scipy.optimize import curve_fit
```

```
def fun_to_fit(x, a, b, c):  
    return a*x**2 + b*x + c
```

$$\begin{aligned}a &= 1.37 \pm 0.39 \\ b &= -0.32 \pm 0.81 \\ c &= -0.21 \pm 0.58\end{aligned}$$

```
ValsBest, Cov = curve_fit(fun_to_fit, x, y)
```

$$y = 1.37 x^2 - 0.32 x - 0.21$$







```
from scipy.optimize import curve_fit
```

```
def fun_to_fit(x, a, b, c):  
    return a*x**2 + b*x + c
```


```
ValsBest, Cov = curve_fit(fun_to_fit, x, y)
```

if error bars known → **error weighted fit!**

```
ValsBest, Cov = curve_fit(fun_to_fit, x, y, sigma = err, absolute_sigma = True)
```



VS







data points have errors and/or fit parameter have uncertainties too

→ implies uncertainties of the model

→ confidence band/interval

$$\begin{aligned} a &= 1.37 \pm 0.39 \\ b &= -0.32 \pm 0.81 \\ c &= -0.21 \pm 0.58 \end{aligned}$$





data points have errors and/or fit parameter have uncertainties too

- implies uncertainties of the model
- confidence band/interval
- idea: **bootstrapping!**

$$\begin{aligned}a &= 1.37 \pm 0.39 \\b &= -0.32 \pm 0.81 \\c &= -0.21 \pm 0.58\end{aligned}$$



data points have errors and/or fit parameter have uncertainties too

$$\begin{aligned} a &= 1.37 \pm 0.39 \\ b &= -0.32 \pm 0.81 \\ c &= -0.21 \pm 0.58 \end{aligned}$$

→ implies uncertainties of the model

→ confidence band/interval

→ idea: **bootstrapping!**

if **no errors of y_i known**

→ assuming that fitted parameters follow a **normal distribution**, i.e. $a = 1.37 \pm 0.39$
where $\mu_a = 1.37$ and $\sigma_a = 0.39$ and so on...

→ **varying** the parameters **within their errors** using `np.random.normal( $\mu_a$ ,  $\sigma_a$ , N)`
N times

→ for each **N**, **generating a curve fit**

→ from set of N curve fits → **calculating percentiles** for confidence band/ interval



data points have errors and/or fit parameter have uncertainties too

→ implies uncertainties of the model

→ confidence band/interval

→ idea: **bootstrapping!**

$$a = 1.37 \pm 0.39$$

$$b = -0.32 \pm 0.81$$

$$c = -0.21 \pm 0.58$$


if **no errors of y_i known**

→ assuming that fitted parameters follow a **normal distribution**, i.e. $a = 1.37 \pm 0.39$ where $\mu_a = 1.37$ and $\sigma_a = 0.39$ and so on...

→ **varying** the parameters **within their errors** using `np.random.normal( $\mu$ ,  $\sigma$ , N)` **N times**

→ for each **N**, **generating a curve fit**

→ from set of N curve fits → **calculating percentiles** for conf





data points have errors and/or fit parameter have uncertainties too

→ implies uncertainties of the model

→ confidence band/interval

→ idea: **bootstrapping!**

$$a = 1.37 \pm 0.39$$

$$b = -0.32 \pm 0.81$$

$$c = -0.21 \pm 0.58$$

if **errors of y_i known**

→ assuming that errors of y_i follow a **normal distribution**, i.e. $y_i(\text{boot}) = y_i \pm \sigma_i$

→ **varying** all y_i **within their errors** using `np.random.normal( $y_i$ ,  $\sigma_i$ , N)`
N times

→ for each **N**, **generating a curve fit**

→ from set of N curve fits → **calculating percentiles** for confidence band/ interval



data points have errors and/or fit parameter have uncertainties too

→ implies uncertainties of the model

→ confidence band/interval

→ idea: **bootstrapping!**

$$a = 1.37 \pm 0.39$$

$$b = -0.32 \pm 0.81$$

$$c = -0.21 \pm 0.58$$


if **errors of y_i known**

→ assuming that errors of y_i follow a **normal distribution** i.e. $v_i(\text{boot}) = v_i + \sigma_i$

→ **varying** all y_i **within their errors** using np.random.normal
N times

→ for each **N**, **generating a curve fit**

→ from set of N curve fits → **calculating percentiles**






bootstrapping

explore the .py file

ConfidenceInterval.py



USAGE:

```

...generating a test sample:
....
....x.....=np.linspace(-1,3,20)
....err.....=np.random.normal(0,1,(len(x),))#1sigma errorbars
....y.....=x**2+err
....errorbars=abs(err)
....
....
....1).plotting data
.....
....F1=FitData(x,y)
....F2=FitData(x,y,errorbars)
....F3=FitData(x,y,errorbars,time='[s]',pressure='[MPa]')
....
....
....2).fitting data (returns best values of fitted params, 1sigma confidence and
.....reduced chi2 if errorbars given, MSE else)
.....
....res1=F1.Fit()
....res2=F2.Fit()
....res2=F3.Fit()
....
....res12=F1.Fit("a*x**2",[1],(-0.5,10))
....
....
....3).Bootstrapping (either varying within errorbars or within conf of fitted
.....params)
.....
....F1.RunBootStrap()
....F2.RunBootStrap()
....F3.RunBootStrap()
....
....F1.RunBootStrap(100,[90,95],np.linspace(-1,5,200))
....F3.RunBootStrap(100,[90,95],np.linspace(-1,5,200))
....

```




Thank you very much for your attention!

