

Lecture 01:

Introduction to Data Science

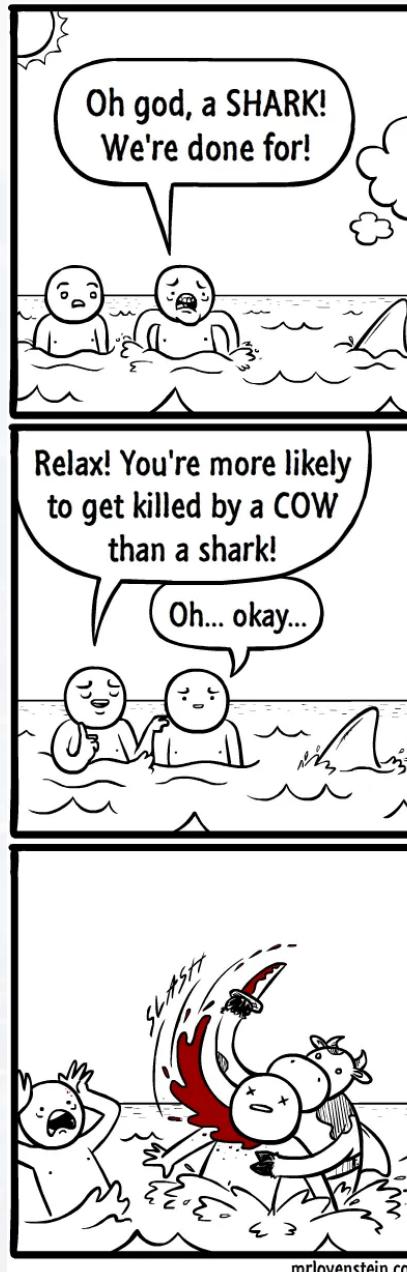


Markus Hohle

University California, Berkeley

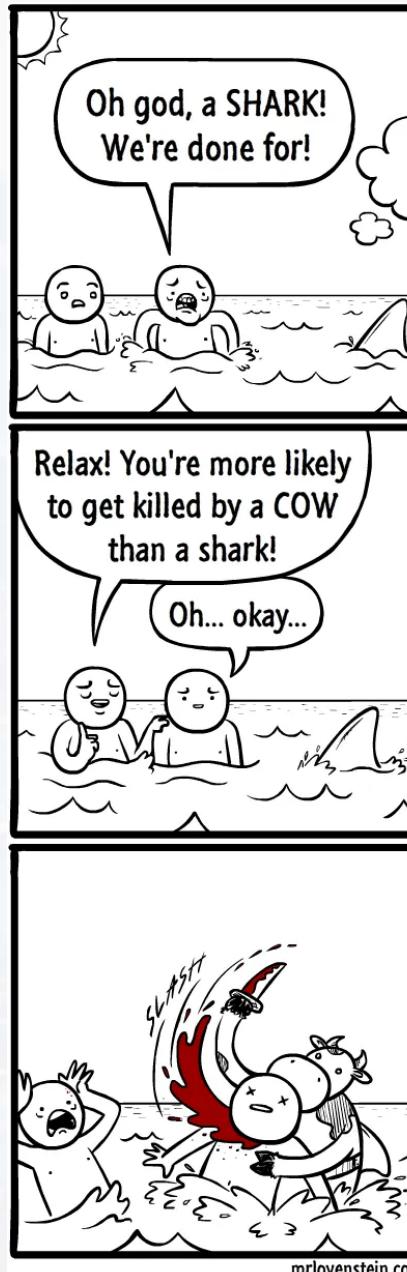
Data Science for Scientific
Computing

MSSE 277A, 3 Units



Outline

- motivation for this course
- structure
- syllabus
- guide through the recordings



Outline

- motivation for this course
- structure
- syllabus
- guide through the recordings



typical workflow:

start:

raw data

```
>GBDP50119-19|Aedes africanus|COI-5P|MF183656  
TTAAATTCGATCTGTTATAATAGTAATAGCTCCTGCTAAAAC  
ATTGCTAAA >GBMIN56476-17|Aedes albopictus|COI-5P|KY378921  
GTTTTAATTCTGATTGAACCTAGACATCCTGGTATATTATTGGAAATGA  
TACTAGGAGCCCTGATATAGCTTTCTCGAATAAAATAATAGTTT  
TCATGCTGGGGCTCAGTGATTAGCAATTTCCTTACATTAGCGG  
GTAATTACAGCTTTTATTACTCTTCTACCCGTATTAGCCGGAGC
```

data exploration (EDA)

- How many files?
- How big are they?
- How many sets?
- consistency
- homogeneity
- gaps? Missing data?

data extraction

- sampling
- filtering

feature selection II

- correlation analysis

feature engineering

- normalization
- encoding
- interpolation/extrapolation

feature selection I

- Which information is relevant?
- Which information is redundant?
- Is there missing information/biases?

model selection

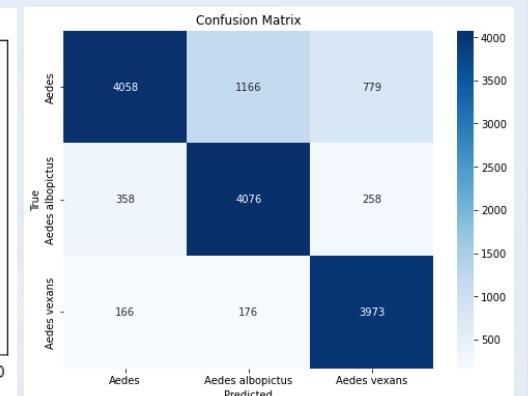
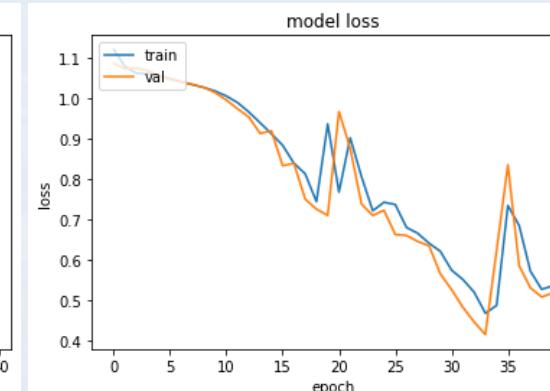
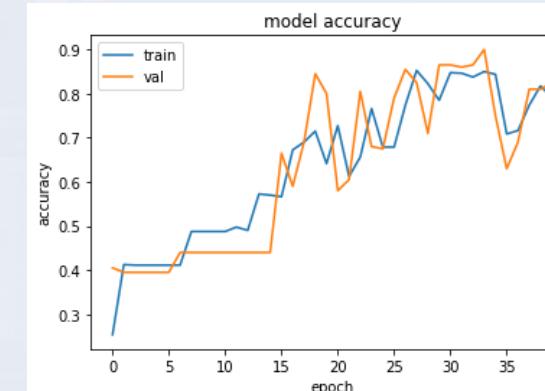
goal: classification/ regression/
prediction/ generation
structure: timeseries/images/
sequences/perm invariance

feature selection III

- significance
- “leave one out”

goal: a model that:

- well **describes** the data
- covers **relevant information** of the data set
- can be used for **predictions**
- ...





typical tools:

start:



matplotlib

model selection
goal: classification/ regression/
prediction/ generation
structure: sequences/
TensorFlow
sequences/perm invariance

NumPy

data exploration (EDA)

- How many files?
-
-
-
-
-

pandas



python

dask

data extraction

- sampling
- filtering

polars

feature selection I

- Which information is relevant?
- Which information is redundant?
- Is there missing information/
- uses?

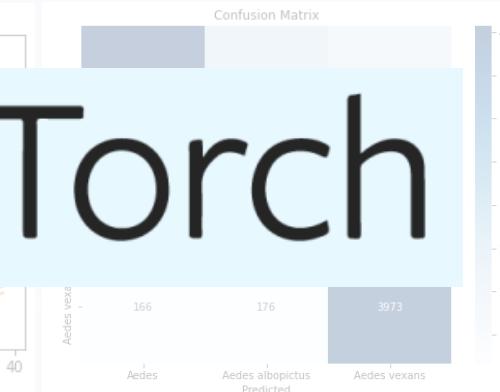
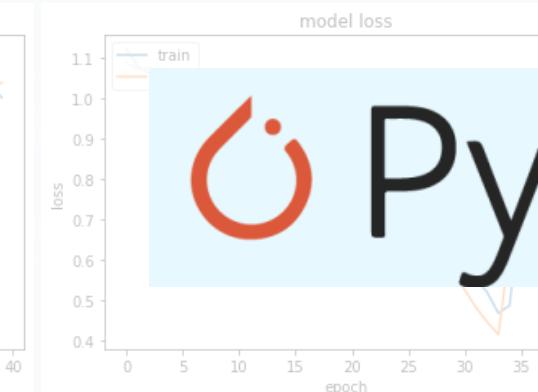
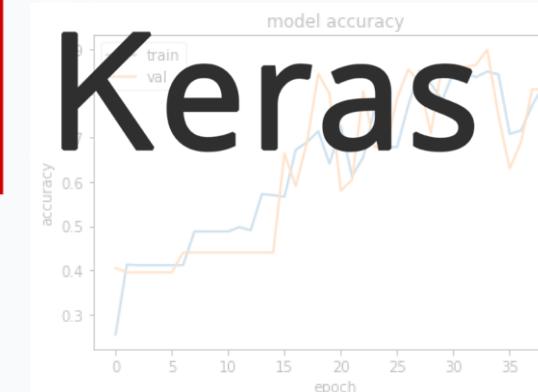
scikit learn

feature selection III

- significance
- "leave one out"

K

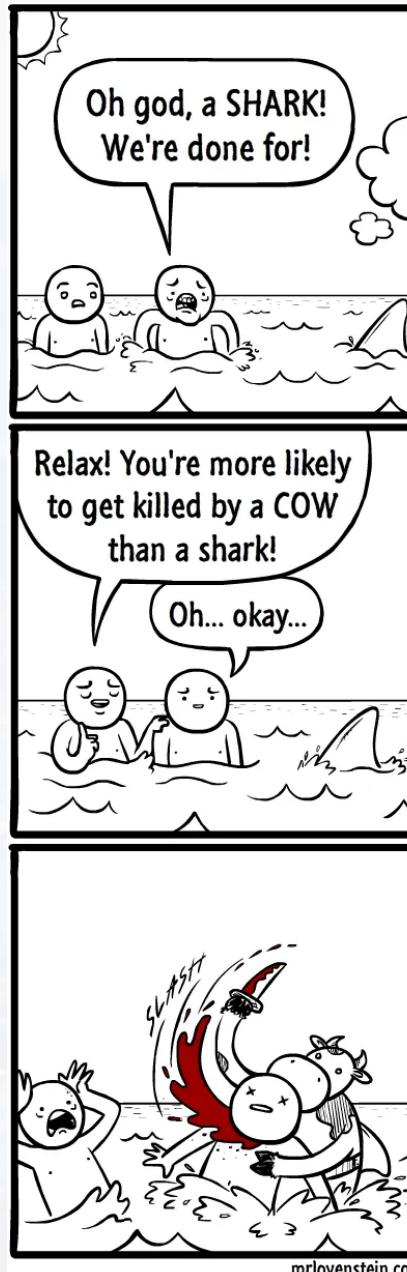
Keras



goal: a model that:

- well describes the data
- covers relevant information of the data set
- can be used for predictions
- ...

PyTorch



Outline

- motivation for this course
- structure**
- syllabus
- guide through the recordings



GSI:

Elizabeth (Lizzie) Gilson

Toxicology Data Scientist at EPA,
UC Berkeley Alumna (MSSE)



Lecturer:

Markus Hohle

Lecturer at UC Berkeley &
Data Analysis Consultant
PhD Physics





Lecture (Markus):	3 hours of asynchronous recorded lectures per week watch any time, but prior to Discussions/Lab Sessions/Homework	
Discussion (Lizzie/ Markus):	Tuesday,	5:30 – 6:30pm PT
Lab Session (Lizzie/ Markus):	every other Wednesday,	6:00 – 8:00pm PT
Office Hours (Markus):	Friday,	5:00 – 6:00pm PT

Grades:

assignment	weight
5 Problem Sets:	40%
2 Programming Projects (midterm & final project)	20%
Lecture Exercises	20%
Discussion & Lab Participation (be active! ask/answer questions!)	20%

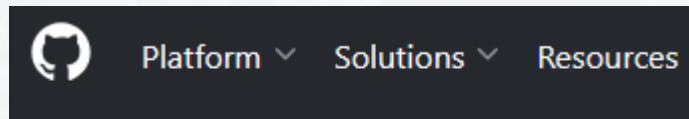


bcourses

- dates & times
- homework assignments and sample solutions
- videos/slides/notebooks/data/codes
- links to discussions/labs/office hours

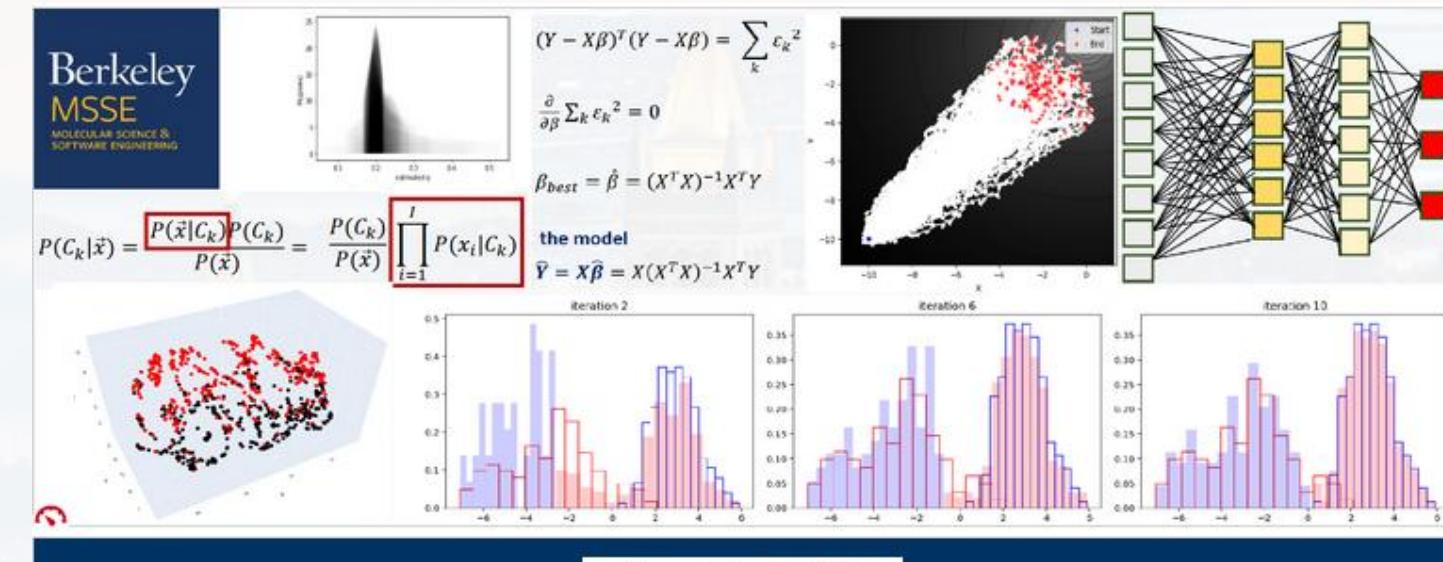
GitHub

slides/notebooks/data/codes



MarkusHohle / UC-Berkeley

Public



Chemistry 277A

Data Science for Scientific Computing

Module 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 11 | 12 | 13 | 14 | 15 | 16 | 17





bcourses

- dates & times
- homework assignments and sample solutions
- videos/slides/notebooks/data/codes
- links to discussions/labs/office hours

Course Map				
Week	Dates	Topics	Reading/Quizzes	Material
1	Jan 20th	Introduction to Data Science		Module 1
2	Jan 26th	Data Sampling and Probability, Pandas		Module 2
3	Feb 2nd	Exploratory Data Analysis (EDA), and Regex Part I		Module 3
4	Feb 9th	Exploratory Data Analysis (EDA), and Regex Part II		Module 4
5	Feb 16th	Introduction to SQL		Module 5
6	Feb 23rd	Feature Analysis, Engineering and Encoding		Module 6



bcourses

- **dates & times**
- homework assignments and sample solutions
- videos/slides/notebooks/data/codes
- **links to discussions/labs/office hours**

Course Start Date: Tuesday, January, 20th, 2026

Course End Date: Friday, May 8th, 2026

Spring Recess: Monday, March 23–Thursday, March 26, 2026

All times are PST.

Schedule

Event	Day	Time	Link
Lecture	asynchronous	3 hours prerecorded	
Discussions	Tuesday	5:30 pm - 6:30 pm	here
Lab (every other week)	Wednesday	6:00 pm - 8:00 pm	here
Office Hour GSI	TBD	TBD	here
Office Hour Lecturer	Friday	5:00 pm - 6:00 pm	here



bcourses

- dates & times
- **homework assignments and sample solutions**
- **videos/slides/notebooks/data/codes**
- links to discussions/labs/office hours

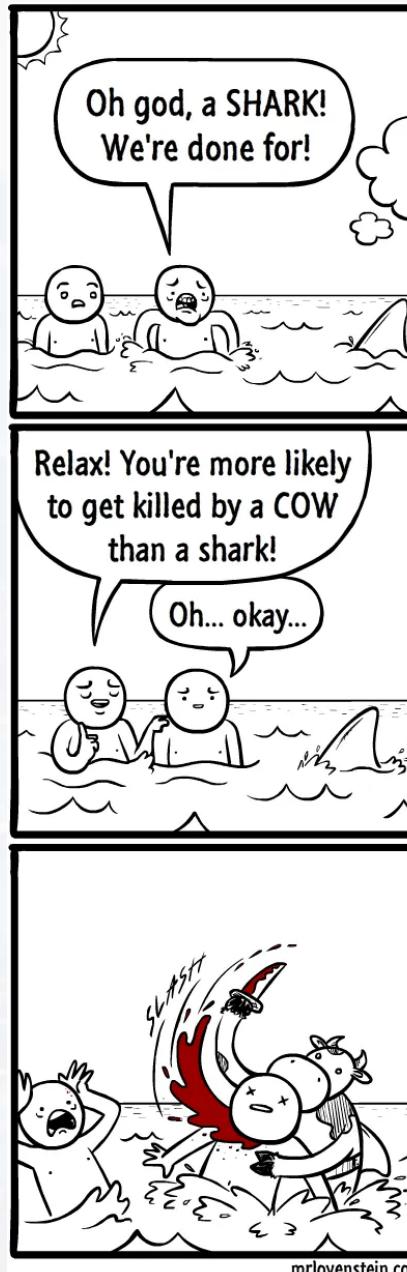
 ▾ Module 1: Introduction
 LectureExercise 01.ipynb
 LectureExercise 01 Solution.ipynb
 MessyFile.xlsx
 ▾ Module 2: Pandas and Memory Efficient Sampling
 LargerThanMemoryExample.ipynb
 LectureExercise 02.ipynb
 LectureExercise 02 Solution.ipynb
 Data_Set.txt



bcourses

- dates & times
- **homework assignments and sample solutions**
- videos/slides/notebooks/data/codes
- links to discussions/labs/office hours

- ☰ ▾ Problem Sets
- ☰ ▾ Programming Projects
- ☰ ▾ Lecture Exercises
- ☰ ▾ Discussion & Lab Participation



Outline

- motivation for this course
- structure
- syllabus**
- guide through the recordings



Lecture 1: Introduction to Data Science

Lecture 2: Data Sampling and Probability, Pandas

data acquisition and analysis

Lecture 3: Exploratory Data Analysis (EDA), and Regex Part I

Lecture 4: Exploratory Data Analysis (EDA), and Regex Part II

Lecture 5: Introduction to SQL

Lecture 6: Feature Analysis, Engineering and Encoding

feature Selection and Analysis

Lecture 7: PCA, LDA and Correlation, Dimension Reduction

Lecture 8: Advanced Visualization Tools (UMAP, T-SNE) in Python

project 1

Lecture 9: Feature Selection via Correlation Analysis

Spring Recess

Lecture 10: Introduction to linear and logistic regression; OLS

modelling

Lecture 11: Avoiding Overfitting and Regularization (L1 & L2)

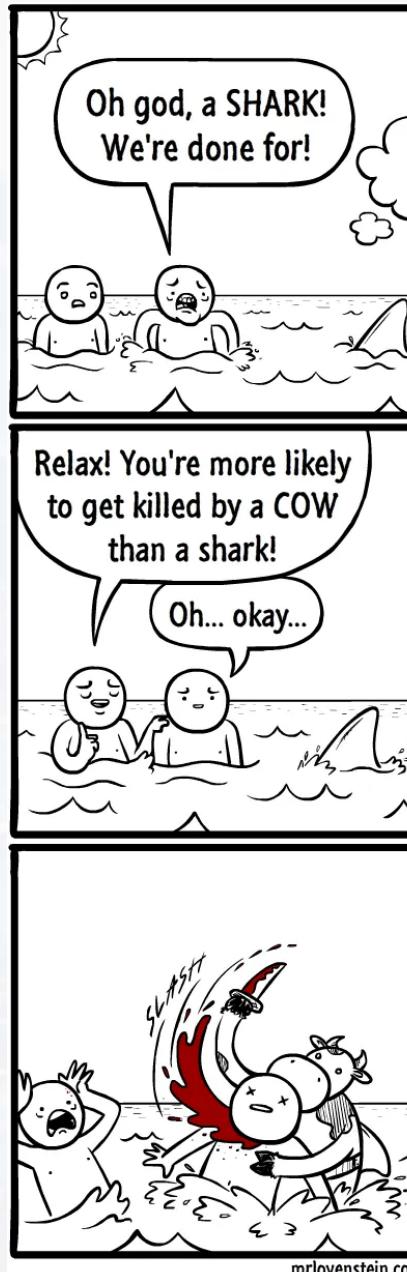
Lecture 12: Feature Selection via Regression Analysis

Lecture 13: Gradient Descent

Lecture 14: Clustering and Classification: Trees, KNN, K-Mean, GMM

project 2

Lecture 15: Feature Selection via ANNs ("leave-one-out", Accuracy Drop & Entropy)



Outline

- motivation for this course
- structure
- syllabus
- guide through the recordings**



1st: watch the lecture recordings

- Lect_5_1_Kmeans_Intro
- Lect_5_2_Kmeans_WalkThrough
- Lect_5_3_Kmeans_Summary
- Lect_5_4_GMM_Intro
- Lect_5_5_GMM_WalkThrough

lectures are ordered: *Lecture_Module_Order_Topic*

- try to understand as much as possible
- it is ok, if you don't understand everything!

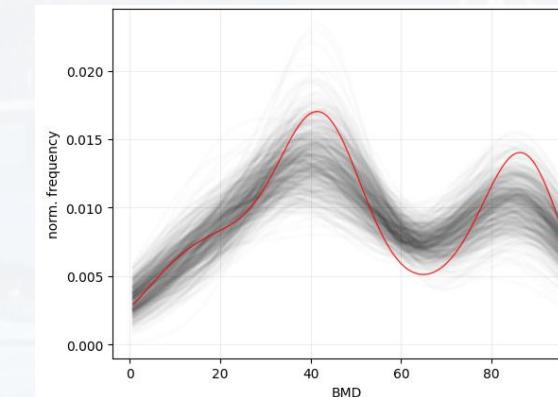
2nd: explore the jupyter notebooks and/or .py scripts

▼ Data Sampling - Methods

1) Objective

In the last examples, we deal with larger than memory files. Another approach is to load the entire data set in order to analyze it. We pick specific portions via sampling.

- Random Sampling
- Stratified Sampling
- Systematic Sampling
- Cluster Sampling
- Bootstrap Sampling
- Oversampling & Undersampling (Basic Concepts)



- follow the instructions
- try to understand each step



3rd: rewatch the lecture recordings if necessary

- Lect_5_1_Kmeans_Intro
- Lect_5_2_Kmeans_WalkThrough
- Lect_5_3_Kmeans_Summary
- Lect_5_4_GMM_Intro
- Lect_5_5_GMM_WalkThrough

lectures are ordered: *Lecture_Module_Order_Topic*

4th: write down questions for the lab and/or discussions

note: students have to actively attend the lab sessions and discussions (**asking/answering questions**) for grades!

5th: solve the homework assignments

3) Task

Read the file *Data_Set.txt* to a standard pandas dataframe.

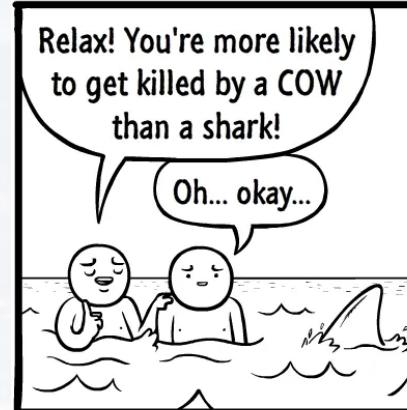
1. Write a function that is similar to "*def ReadWithAnyToolAnyMethod.py*"
 - .txt
 - .csv
 - .pkl
 - .parquetand monitor time and memory usage.
2. list the size of all these files
3. generate a table (as dataframe) that lists time, memory usage and size



Berkeley
UNIVERSITY OF CALIFORNIA

Data Science for Scientific Computing:

Introduction



Enjoy the Course ☺!