

## Lecture 06:

# Feature analysis, engineering and encoding



Markus Hohle  
University California, Berkeley

Data Science for Scientific  
Computing  
MSSE 277A, 3 Units



## Outline

**What is a feature?**

**Encodings**

**Normalization**

**Visualization**



## Outline

**What is a feature?**

Encodings

Normalization

Visualization



data analysis tasks

- classification

0.40 cat  
0.32 frog  
0.16 bird  
0.06 ship  
0.03 dog

- regression

- generation

data analysis tasks

## - classification

Index	molecular_weight	electronegativity	bond_lengths	num_hydrogen_bonds	logP	label
0	341.704	2.65585	3.09407	2	9.11147	Toxic
1	335.951	3.22262	2.89039	7	8.92848	Toxic
2	235.203	2.44115	2.48203	1	6.49731	Non-Toxic
3	246.505	2.76656	2.71547	7	7.45089	Non-Toxic
4	437.939	3.4801	3.59569	3	10.9156	Non-Toxic

## - regression

Index	molecular_weight	electronegativity	bond_lengths	num_hydrogen_bonds	logP	toxicity_score
0	341.704	2.65585	3.09407	2	9.11147	80.9281
1	335.951	3.22262	2.89039	7	8.92848	83.4911
2	235.203	2.44115	2.48203	1	6.49731	61.8406
3	246.505	2.76656	2.71547	7	7.45089	57.0538
4	437.939	3.4801	3.59569	3	10.9156	131.326

## - generation

“create a molecule with the properties XYZ,  
such that it is non-toxic”


carrier or linker proteins (drug development)



This is a **feature vector**

- tells us where the datapoint is located in **feature space**
- can be **high dimensional!**


Index	molecular_weight	electronegativity	bond_lengths	num_hydrogen_bonds	logP	label
0	341.704	2.65585	3.09407	2	9.11147	Toxic
1	335.951	3.22262	2.89039	7	8.92848	Toxic
2	235.203	2.44115	2.48203	1	6.49731	Non-Toxic
3	246.505	2.76656	2.71547	7	7.45089	Non-Toxic
4	437.939	3.4801	3.59569	3	10.9156	Non-Toxic



high dimensional feature space  
(visualization via PCA, UMAP, t-SNE etc see later)

- toxic molecules
- non-toxic molecules

classification  
boundary from  
ML-algorithm



lower dimensional projection  
(visualization via PCA, UMAP, t-SNE  
etc see later)



This is a **feature vector**

- tells us where the datapoint is located in **feature space**
- can be **high dimensional!**

Index	molecular_weight	electronegativity	bond_lengths	num_hydrogen_bonds	logP	label
0	341.704	2.65585	3.09407	2	9.11147	Toxic
1	335.951	3.22262	2.89039	7	8.92848	Toxic
2	235.203	2.44115	2.48203	1	6.49731	Non-Toxic
3	246.505	2.76656	2.71547	7	7.45089	Non-Toxic
4	437.939	3.4801	3.59569	3	10.9156	Non-Toxic

```
data = pd.read_csv('Cystfibr.txt', delimiter = '\t')
```

[source](#)

data - DataFrame

Index	age	sex	height	weight	bmp	fev1	rv	frc	tlc	pemax
0	7	0	109	13.1	68	32	258	183	137	95
1	7	1	112	12.9	65	19	449	245	134	85
2	8	0	124	14.1	64	22	441	268	147	100
3	8	1	125	16.2	67	41	234	146	124	85

↑      ↑      ↑      ↑      ↑      ↑      ↑      ↑      ↑

**feature vector**

features can be **numerical: units are arbitrary** (cm, inches etc), i.e. the **value itself is arbitrary** (depending on unit system → **normalization!**)



This is a **feature vector**

- tells us where the datapoint is located in **feature space**
- can be **high dimensional!**

Index	molecular_weight	electronegativity	bond_lengths	num_hydrogen_bonds	logP	label
0	341.704	2.65585	3.09407	2	9.11147	Toxic
1	335.951	3.22262	2.89039	7	8.92848	Toxic
2	235.203	2.44115	2.48203	1	6.49731	Non-Toxic
3	246.505	2.76656	2.71547	7	7.45089	Non-Toxic
4	437.939	3.4801	3.59569	3	10.9156	Non-Toxic

```
data = pd.read_csv('Cystfibr.txt', delimiter = '\t')
```

[source](#)

data - DataFrame

Index	age	sex	height	weight	bmp	fev1	rv	frc	tlc	pemax
0	7	0	109	13.1	68	32	258	183	137	95
1	7	1	112	12.9	65	19	449	245	134	85
2	8	0	124	14.1	64	22	441	268	147	100
3	8	1	125	16.2	67	41	234	146	124	85

feature vector



features can be **categorical**: have to be converted to numerical values → **encoding!**)



This is a **feature vector**

- tells us where the datapoint is located in **feature space**
- can be **high dimensional!**

Index	molecular_weight	electronegativity	bond_lengths	num_hydrogen_bonds	logP	label
0	341.704	2.65585	3.09407	2	9.11147	Toxic
1	335.951	3.22262	2.89039	7	8.92848	Toxic
2	235.203	2.44115	2.48203	1	6.49731	Non-Toxic
3	246.505	2.76656	2.71547	7	7.45089	Non-Toxic
4	437.939	3.4801	3.59569	3	10.9156	Non-Toxic

```
AD = pd.read_excel('AD_data.xlsx', sheet_name = 'Summary Form')
```

[source](#)

or:

```
data      = pd.ExcelFile('AD_data.xlsx')
sheet_names = data.sheet_names
```

```
AD      = data.parse(sheet_name = sheet_names[1])
```



This is a **feature vector**

- tells us where the datapoint is located in **feature space**
- can be **high dimensional!**

Index	molecular_weight	electronegativity	bond_lengths	num_hydrogen_bonds	logP	label
0	341.704	2.65585	3.09407	2	9.11147	Toxic
1	335.951	3.22262	2.89039	7	8.92848	Toxic
2	235.203	2.44115	2.48203	1	6.49731	Non-Toxic
3	246.505	2.76656	2.71547	7	7.45089	Non-Toxic
4	437.939	3.4801	3.59569	3	10.9156	Non-Toxic

```
AD = pd.read_excel('AD_data.xlsx', sheet_name = 'Summary Form')
```

[source](#)

Index	state	sex	age range(0: <75; 1: >=75)	education	economic	lifestyle	heridity	marriage	health	BIM	p
0	0	0	0	2	1	2	2	3	5	24.19	0.553777
1	0	0	1	1	4	1	1	3	5	18.05	0.375283
2	0	0	0	3	2	2	3	3	5	21.98	0.541784
3	0	0	0	1	1	4	2	3	5	26.18	0.527598

features can be actual **numerical**, but binned to **categorical**

**Note: That is no good practice! Binning leads to loss of information!**



This is a **feature vector**

- tells us where the datapoint is located in **feature space**
- can be **high dimensional!**

Index	molecular_weight	electronegativity	bond_lengths	num_hydrogen_bonds	logP	label
0	341.704	2.65585	3.09407	2	9.11147	Toxic
1	335.951	3.22262	2.89039	7	8.92848	Toxic
2	235.203	2.44115	2.48203	1	6.49731	Non-Toxic
3	246.505	2.76656	2.71547	7	7.45089	Non-Toxic
4	437.939	3.4801	3.59569	3	10.9156	Non-Toxic

```
AD = pd.read_excel('AD_data.xlsx', sheet_name = 'Summary Form')
```

[source](#)

Index	state	sex	age range(0: <75; 1: >=75)	education	economic	lifestyle	heridity	marriage	health	BIM	p
0	0	0	0	2	1	2	2	3	5	24.19	0.553777
1	0	0	1	1	4	1	1	3	5	18.05	0.375283
2	0	0	0	3	2	2	3	3	5	21.98	0.541784
3	0	0	0	1	1	4	2	3	5	26.18	0.527598

most features are **categorical**

Number of categories/states for a specific feature is called **cardinality**.



This is a **feature vector**

- tells us where the datapoint is located in **feature space**
- can be **high dimensional!**

Index	molecular_weight	electronegativity	bond_lengths	num_hydrogen_bonds	logP	label
0	341.704	2.65585	3.09407	2	9.11147	Toxic
1	335.951	3.22262	2.89039	7	8.92848	Toxic
2	235.203	2.44115	2.48203	1	6.49731	Non-Toxic
3	246.505	2.76656	2.71547	7	7.45089	Non-Toxic
4	437.939	3.4801	3.59569	3	10.9156	Non-Toxic

```
AD = pd.read_excel('AD_data.xlsx', sheet_name = 'Summary Form')
```

[source](#)

Number of categories for a specific feature is called **cardinality**.

Let's check out cardinality:

```
Categoricals = AD.columns[:-2]
```

```
for c in Categoricals:
```

```
    card = len(set(AD[c]))
```

```
    print("Cardinality of " + c + ': ' + str(card))
```

np.max won't be save:  
We don't know how  
categories have been  
enumerated



This is a **feature vector**

- tells us where the datapoint is located in **feature space**
- can be **high dimensional!**

Index	molecular_weight	electronegativity	bond_lengths	num_hydrogen_bonds	logP	label
0	341.704	2.65585	3.09407	2	9.11147	Toxic
1	335.951	3.22262	2.89039	7	8.92848	Toxic
2	235.203	2.44115	2.48203	1	6.49731	Non-Toxic
3	246.505	2.76656	2.71547	7	7.45089	Non-Toxic
4	437.939	3.4801	3.59569	3	10.9156	Non-Toxic

```
AD = pd.read_excel('AD_data.xlsx', sheet_name = 'Summary Form')
```

[source](#)

Number of categories for a specific feature is called **cardinality**.

Let's check out cardinality:

```
Categoricals = AD.columns[:-2]
```

```
for c in Categoricals:
```

```
    card = len(set(AD[c]))
```

```
    print("Cardinality of " + c + ': ' + str(card))
```

Cardinality of state: 2

Cardinality of sex: 2

Cardinality of age range(0: <75; 1: >=75): 2

Cardinality of education: 3

Cardinality of economic: 5

Cardinality of lifestyle: 5


Cardinality of heridity: 3

Cardinality of marriage: 3

Cardinality of health: 5



cardinality  $c$  of features:




four nucleotides  
 $c = 4$

;Cytochrome c1, heme protein, mitochondrial

>cow  
MAAAAATLRGAMVGPRGAGLPGARARGLLCGARPG  
LSALGMLAAGGAGLAVALHSAVSASDLELHPPSYPWS  
CSSCHSMDYVAYRHLVGVCYTEDEAKALAEEVEVQDG

twenty amino acids  
 $c = 20$



$10^4$  words  
 $c = 10^4$

**summary:**

- **features** are entries (= variables) of the **feature vector**
- the **feature vector** tells us where the datapoint is located in **feature space**
- the **feature space** can be **high dimensional**
- **features** can be numerical; values change depending on unit system → **normalization**.
- **features** can be categorical → proper **encoding**
- features need to be **equally weighted** so that any algorithm/model identifies patterns in the data and makes predictions **based on those patterns**, not based on the choice of feature encoding!

Index	molecular_weight	electronegativity	bond_lengths	num_hydrogen_bonds	logP	label
0	341.704	2.65585	3.09407	2	9.11147	Toxic
1	335.951	3.22262	2.89039	7	8.92848	Toxic
2	235.203	2.44115	2.48203	1	6.49731	Non-Toxic
3	246.505	2.76656	2.71547	7	7.45089	Non-Toxic
4	437.939	3.4801	3.59569	3	10.9156	Non-Toxic



## Outline

What is a feature?

Encodings

Normalization

Visualization



features can be categorical → proper encoding

Index	state	sex	age range(0: <75; 1: >=75)	education	economic	lifestyle	heridity	marriage	health	BIM	p
0	0	0	0	2	1	2	2	3	5	24.19	0.553777
1	0	0	1	1	4	1	1	3	5	18.05	0.375283
2	0	0	0	3	2	2	3	3	5	21.98	0.541784
3	0	0	0	1	1	4	2	3	5	26.18	0.527598

Cardinality of state: 2

Cardinality of sex: 2

Cardinality of age range(0: <75; 1: >=75): 2

Cardinality of education: 3

Cardinality of economic: 5

Cardinality of lifestyle: 5

Cardinality of heridity: 3

Cardinality of marriage: 3

Cardinality of health: 5



features can be categorical → proper encoding

Index	state	sex	age range(0: <75; 1: >=75)	education	economic	lifestyle	heridity	marriage	health	BIM	p
0	0	0	0	2	1	2	2	3	5	24.19	0.553777
1	0	0	1	1	4	1	1	3	5	18.05	0.375283
2	0	0	0	3	2	2	3	3	5	21.98	0.541784
3	0	0	0	1	1	4	2	3	5	26.18	0.527598

Cardinality of state: 2

Cardinality of sex: 2

Cardinality of age range(0: <75; 1: >=75): 2

Cardinality of education: 3

Cardinality of economic: 5

Cardinality of lifestyle: 5

Cardinality of heridity: 3

Cardinality of marriage: 3

Cardinality of health: 5

- encoding the different categories of the feature “heridity” just 1, 2, 3 ...etc **is arbitrary!**
- Why not encoding like -2, -1, 0, 1, 2?
- choice **will affect location** in feature space and therefore **influence result** for any categorization/regression/generation task
- **What is a good encoding strategy?**



**problem:** Features need to be **equally weighted** so that any algorithm/model identifies patterns in the data and makes predictions **based on those patterns**, not based on the choice of feature encoding!

**categorical features can be:**

**ordinal** = there is an **inherent order** and features can be **ranked**

example: - health status  
poor, fair, good, excellent etc

- grades  
A+, A, A-, B+, ...

Cardinality of economic: 5  
Cardinality of lifestyle: 5  
Cardinality of heridity: 3  
Cardinality of marriage: 3  
**Cardinality of health: 5**

**nominal = no inherent order**

example: - ACGT

- name of a person




### The most common encoding strategies:

- one-hot encoding
- dummy encoding
- ordinal encoding
- binary encoding
- count encoding



- **fast and simple**
- if categories are not related and **purely nominal**
- suitable if **low cardinality**
- each category is a **binary vector**

one-hot encoding  
dummy encoding  
ordinal encoding  
binary encoding  
count encoding



Cat = "ACGT"

Enc = np.identity(len(Cat))

Encoding = {c: e for c, e in zip(Cat, Enc)}

Encoding - Dictionary (4 elements) ✖

Key	Type	Size	Value
A	Array of float64	(4,)	[1. 0. 0. 0.]
C	Array of float64	(4,)	[0. 1. 0. 0.]
G	Array of float64	(4,)	[0. 0. 1. 0.]
T	Array of float64	(4,)	[0. 0. 0. 1.]



- **fast** and **simple**
- if categories are not related and **purely nominal**
- suitable if **low cardinality**
- each category is a **binary vector**

one-hot encoding  
dummy encoding  
ordinal encoding  
binary encoding  
count encoding

```
Cat = "ACGT"
```

```
Enc = np.identity(len(Cat))
```

```
Encoding = {c: e for c, e in zip(Cat, Enc)}
```

```
Sequence = 'CCTGTAATC'
```

```
Encoded = [Encoding[s] for s in Sequence]
```

```
print(np.array(Encoded))
```

```
[[0. 1. 0. 0.]  
 [0. 1. 0. 0.]  
 [0. 0. 0. 1.]  
 [0. 0. 1. 0.]  
 [0. 0. 0. 1.]  
 [1. 0. 0. 0.]  
 [1. 0. 0. 0.]  
 [0. 0. 0. 1.]  
 [0. 1. 0. 0.]]
```



- similar to one-hot encoding
- if categories are not related and **purely nominal**
- suitable if **low cardinality**
- each category is a **binary value**

one-hot encoding  
dummy encoding  
ordinal encoding  
binary encoding  
count encoding

```
AD_dummy = pd.get_dummies(AD, columns = Categoricals, dtype = int)
```

Index	BIM	p	state_0	state_1	sex_0	sex_1	(0: <75; 1: (0: <75; 1: education_1	education_2	education_3	economic_1	economic_2	economic_3	economic_4	economic_5
0	24.19	0.553777	1	0	1	0	1	0	0	1	0	1	0	0
1	18.05	0.375283	1	0	1	0	0	1	1	0	0	0	1	0
2	21.98	0.541784	1	0	1	0	1	0	0	1	0	1	0	0
3	26.18	0.527598	1	0	1	0	1	0	1	0	1	0	0	0

lifestyle_1	lifestyle_2	lifestyle_3	lifestyle_4	lifestyle_5	heridity_1	heridity_2	heridity_3	marriage_1	marriage_2	marriage_3	health_1	health_2	health_3	health_4	health_5
0	1	0	0	0	0	1	0	0	0	1	0	0	0	0	1
1	0	0	0	0	1	0	0	0	0	1	0	0	0	0	1
0	1	0	0	0	0	0	1	0	0	1	0	0	0	0	1
0	0	0	1	0	0	1	0	0	0	1	0	0	0	0	1



- similar to one-hot encoding
- if categories are not related and **purely nominal**
- suitable if **low cardinality**
- each category is a **binary value**

one-hot encoding  
dummy encoding  
ordinal encoding  
binary encoding  
count encoding

```
AD_dummy = pd.get_dummies(AD, columns = Categoricals, dtype = int)
```

Index	BIM	p	state_0	state_1	sex_0	sex_1	(0: <75; 1: (0: <75; 1:	education_1	education_2	education_3	economic_1	economic_2	economic_3	economic_4	economic_5
0	24.19	0.553777	1	0	1	0	1	0	0	1	1	0	0	0	0
1	18.05	0.375283	1	0	1	0	0	1	1	0	0	0	0	1	0
2	21.98	0.541784	1	0	1	0	1	0	0	0	1	0	1	0	0
3	26.18	0.527598	1	0	1	0	1	0	1	0	0	1	0	0	0

information is redundant:

- if i.e. `sex_0 = 1`, we know that `sex_1 = 0`
- 100% correlation, many ML methods will perform poorly

idea:

- drop one column

```
AD_dummy = pd.get_dummies(AD, columns = Categoricals, dtype = int, drop_first = True)
```



These features are purely nominal

- one-hot or dummy encoding applicable

one-hot encoding  
dummy encoding  
ordinal encoding  
binary encoding  
count encoding

Index	BIM	p	state_1	sex_1	age range(0: <75; 1: >=75)_1	education_2	education_3	economic_2	economic_3	economic_4	economic_5
0	24.19	0.553777	0	0	0	1	0	0	0	0	0
1	18.05	0.375283	0	0	1	0	0	0	0	1	0
2	21.98	0.541784	0	0	0	0	1	1	0	0	0
3	26.18	0.527598	0	0	0	0	0	0	0	0	0

lifestyle_2	lifestyle_3	lifestyle_4	lifestyle_5	heridity_2	heridity_3	marriage_2	marriage_3	health_2	health_3	health_4	health_5
1	0	0	0	1	0	0	1	0	0	0	1
0	0	0	0	0	0	0	1	0	0	0	1
1	0	0	0	0	1	0	1	0	0	0	1
0	0	1	0	1	0	0	1	0	0	0	1



These features **might be ordinal!**

- one-hot or dummy encoding **might not be applicable!**

one-hot encoding  
dummy encoding  
ordinal encoding  
binary encoding  
count encoding

Index	BIM	p	state_1	sex_1	age range(0: <75; 1: >=75)_1	education_2	education_3	economic_2	economic_3	economic_4	economic_5
0	24.19	0.553777	0	0	0	1	0	0	0	0	0
1	18.05	0.375283	0	0	1	0	0	0	0	1	0
2	21.98	0.541784	0	0	0	0	1	1	0	0	0
3	26.18	0.527598	0	0	0	0	0	0	0	0	0

lifestyle_2	lifestyle_3	lifestyle_4	lifestyle_5	heridity_2	heridity_3	marriage_2	marriage_3	health_2	health_3	health_4	health_5
1	0	0	0	1	0	0	1	0	0	0	1
0	0	0	0	0	0	0	1	0	0	0	1
1	0	0	0	0	1	0	1	0	0	0	1
0	0	1	0	1	0	0	1	0	0	0	1



- if categories are **ordinal**
- assigning numerical value in specific order

one-hot encoding  
dummy encoding  
**ordinal encoding**  
binary encoding  
count encoding

```
AD = pd.read_excel('AD_data.xlsx', sheet_name = 'Summary Form') source
```

widowed or divorced, and 3 representing married with a spouse. The education levels were divided into high (university degree or other professional qualification), middle (high school or junior high school), and low (practical qualification related to work). Economic status was divided into five categories based on the Townsend Deprivation Index (which combines information on social class, employment, cars, housing, etc). Higher scores indicate better marital status, higher education levels and better economic status, respectively.



- if categories are **ordinal**
- assigning numerical value in specific order

one-hot encoding  
dummy encoding  
**ordinal encoding**  
binary encoding  
count encoding

```
AD = pd.read_excel('AD_data.xlsx', sheet_name = 'Summary Form') source
```

widowed or divorced, and 3 representing married with a spouse. The education levels were divided into high (university degree or other professional qualification), middle (high school or junior high school), and low (practical qualification related to work). Economic status was divided into five categories based on the Townsend Deprivation Index (which combines information on social class, employment, cars, housing, etc). Higher scores indicate better marital status, higher education levels and better economic status, respectively.



- if categories are **ordinal**
- assigning numerical value in specific order

one-hot encoding  
dummy encoding  
**ordinal encoding**  
binary encoding  
count encoding

```
AD = pd.read_excel('AD_data.xlsx', sheet_name = 'Summary Form') source
```

Index	state	sex	e(0: <75; 1: >=75)	education	economic
0	0	0	0	2	1
1	0	0	1	1	4
2	0	0	0	3	2
3	0	0	0	1	1
4	0	0	0	3	2
5	0	0	0	2	4



- uses **binary system**
- suitable if **high cardinality**
- mixture between **one-hot** and **ordinal** encoding
- **memory efficient** alternative to one-hot or dummy
- **reduces dimensionality**
- often used for language models

one-hot encoding  
dummy encoding  
ordinal encoding  
**binary encoding**  
count encoding

**one hot:**    Cat        = "abcdefghijklmnopqrstuvwxyz .!?"  
                    Enc        = np.identity(len(Cat))

Encoding = {c: e for c, e in zip(Cat, Enc)}

Sequence = 'this is a sentence.'  
Encoded = [Encoding[s] for s in Sequence]  
  
print(np.array(Encoded))



one-hot encoding  
dummy encoding  
ordinal encoding  
**binary encoding**  
count encoding

- one-hot uses memory very inefficiently if high cardinality
  - dimension of encoded vector = cardinality



binary:

```
Cat      = "abcdefghijklmnopqrstuvwxyz .!?"
```

```
bytes_object = Cat.encode('utf-8')
```

```
Encoding = {c: f'{b:08b}' for c, b in zip(Cat, bytes_object)}
```

```
Sequence = 'this is a sentence.'
```

```
Encoded   = [Encoding[s] for s in Sequence]
```

```
print(np.array(Encoded))
```

one-hot encoding  
dummy encoding  
ordinal encoding  
**binary encoding**  
count encoding

turning keys into 8-bit  
bytes object

**binary:**

```
Cat = "abcdefghijklmnopqrstuvwxyz .!?"
```

```
bytes_object = Cat.encode('utf-8')
```

```
Encoding = {c: f'{b:08b}' for c, b in zip(Cat, bytes_object)}
```

```
Sequence  
Encoded = 'this is a sentence.'  
= [Encoding[s] for s in Sequence]
```

```
print(np.array(Encoded))
```

one-hot encoding  
dummy encoding  
ordinal encoding  
**binary encoding**  
count encoding

Key	Type	Size	Value
a	str	8	01100001
b	str	8	01100010
c	str	8	01100011
d	str	8	01100100
e	str	8	01100101
f	str	8	01100110
g	str	8	01100111

**binary:**

```
Cat          = "abcdefghijklmnopqrstuvwxyz .!?"  
  
bytes_object = Cat.encode('utf-8')  
  
Encoding     = {c: f"{b:08b}" for c, b in zip(Cat, bytes_object)}  
  
Sequence     = 'this is a sentence.'  
Encoded      = [Encoding[s] for s in Sequence]  
  
print(np.array(Encoded))
```

```
[ '01110100' '01101000' '01101001' '01110011' '00100000' '01101001'  
'01110011' '00100000' '01100001' '00100000' '01110011' '01100101'  
'01101110' '01110100' '01100101' '01101110' '01100011' '01100101'  
'00101110']
```

one-hot encoding  
dummy encoding  
ordinal encoding  
**binary encoding**  
count encoding



## one-hot

one-hot encoding  
dummy encoding  
ordinal encoding  
**binary encoding**  
count encoding



- assigns values based on **how often categories** appear in the dataset
- aka **frequency encoding**
- **memory efficient**
- **reduces dimensionality**
- if categories are not related and **purely nominal**
- **for large datasets: rel. frequencies are approximations of probabilities**
- problem: different values for different data sets → **hard to compare for small datasets, data leakage**

Sequence = '*this is a sentence.*'

Encoding = {s: Sequence.count(s) **for s in set(Sequence)}**}

one-hot encoding  
dummy encoding  
ordinal encoding  
binary encoding  
count encoding

Key	Type	Size	Value
	int	1	3
n	int	1	2
t	int	1	2
.	int	1	1
h	int	1	1
s	int	1	3
c	int	1	1
a	int	1	1
e	int	1	3
i	int	1	2



```
Sequence = 'this is a sentence.'
```

one-hot encoding  
dummy encoding  
ordinal encoding  
binary encoding  
**count encoding**

```
Encoding = {s: Sequence.count(s) for s in set(Sequence)}
```

```
Encoded = [Encoding[s] for s in Sequence]
```

```
print(np.array(Encoded))
```

```
[2 1 2 3 3 2 3 3 1 3 3 3 2 2 3 2 1 3 1]
```



## one-hot

one-hot encoding  
dummy encoding  
ordinal encoding  
binary encoding  
**count encoding**

## binary

## count/frequency

```
[2 1 2 3 3 2 3 3 1 3 3 3 2 2 3 2 1 3 1]
```

summary

encoding	pros	cons	applicable if
one-hot	simple, variable is binary vector	inefficient for high cardinality	categories are nominal & low cardinality
dummy	simple, variable is binary number	inefficient for high cardinality	categories are nominal & low cardinality
ordinal	efficient for high cardinality	not applicable for nominal categories	categories are ordinal, i. e. have an inherent order
binary	very efficient for high cardinality, reduces dimensionality	less interpretable for humans	recommended if high cardinality
count	simple, easy to interpret	encodes the variable with values from itself (data leakage)	recommended if high cardinality

note, there are more encoding strategies such as target or hash encoding, embeddings etc



## Outline

What is a feature?

Encodings

Normalization

Visualization



Imagine a competition like **Decathlon** with various disciplines

**Decathlon:** each athlete has to perform in a competition of **ten** combined track and field events.

100 Meters Running:

measure time

Long Jump:

measure distance

Shot Put:

measure distance

High Jump:

measure distance

400 Meters Running:

measure time

110 Meters Hurdles:

measure time

Discus Throw:

measure distance

Pole Vault:

measure distance

Javelin Throw:

measure distance

1500 Meters Running:

measure time

**feature vector**

**How do you decide who's the best athlete?**

- time: the **lower** the value the better
- distance: the **larger** the value the better
- comparing values: being 1sec faster in 100m running are eons, but are no big deal in 1500m running



Same problem: how do we compare different quantities measured in arbitrary units?

```
data = pd.read_csv('Cystfibr.txt', delimiter = '\t')
```

[source](#)

data - DataFrame

Index	age	sex	height	weight	bmp	fev1	rv	frc	tlc	pemax
0	7	0	109	13.1	68	32	258	183	137	95
1	7	1	112	12.9	65	19	449	245	134	85
2	8	0	124	14.1	64	22	441	268	147	100
3	8	1	125	16.2	67	41	234	146	124	85

feature vector

↑      ↑      ↑      ↑      ↑      ↑      ↑      ↑      ↑

features can be **numerical: units are arbitrary** (cm, inches etc), i.e. the **value itself is arbitrary** (depending on unit system → **normalization!**)



## two main normalization methods:

**min/max:**

mapping all values to a certain interval (usually 0/1 or -1/+1)

**standard:**


mapping values to standard normal ( $\mu = 0, \sigma^2 = 1$ )

less robust if outliers: squeezes too many datapoints to  $\sigma^2 = 1$

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

data = pd.read_csv('Cystfibr.txt', \
                   delimiter = '\t')

sns.violinplot(data = data)
plt.title("Cystfibr Dataset")
plt.xticks(rotation = 45)
plt.show()
```





two main normalization methods:

```
import pandas as pd  
import seaborn as sns  
import matplotlib.pyplot as plt  
  
from sklearn.preprocessing import StandardScaler, MinMaxScaler
```

scaling data to standard normal

scaling data to min/max range

```
data = pd.read_csv('Cystfibr.txt', delimiter = '\t')
```

initializing scaler

```
SMinMax = MinMaxScaler(feature_range = (0, 1))  
SStand = StandardScaler()
```

actual scaling:  
fitting data and  
transforming

```
Scaled = SMinMax.fit_transform(data)
```

```
sns.violinplot(data = pd.DataFrame(Scaled, columns = data.columns))  
plt.title("Cystfibr Dataset - min/max scaled")  
plt.xticks(rotation = 45)  
plt.show()
```

scaler returns array,  
converting back to  
dataframe



two main normalization methods:

```
import pandas as pd  
import seaborn as sns  
import matplotlib.pyplot as plt  
  
from sklearn.preprocessing import StandardScaler, MinMaxScaler
```

scaling data to standard normal

scaling data to min/max range

```
data = pd.read_csv('Cystfibr.txt', delimiter = '\t')
```

initializing scaler

```
SMinMax = MinMaxScaler(feature_range = (0, 1))  
SStand = StandardScaler()
```




actual scaling:  
fitting data and  
transforming

```
Scaled = SStand.fit_transform(data)
```

```
sns.violinplot(data = pd.DataFrame(Scaled, columns = data.columns))  
plt.title("Cystfibr Dataset - standard normal scaled")  
plt.xticks(rotation = 45)  
plt.show()
```



two main normalization methods:





`fit_transform`

`transform`

`fit`

`fit`: fits the data to the desired scaling

`transform`: performs the actual scaling, applying fitted values

`fit_transform`: performs `fit` and `transform`

### Why are there these three options?

later, for supervised learning we need a **training dataset** in order to train a model

→ run `fit_transform`

once the model has been trained, we want to evaluate its performance using a **test dataset**

→ run `transform` only, **not** `fit` in order to test the model unbiased



```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.preprocessing import StandardScaler, MinMaxScaler
```

```
data           = pd.read_csv('Cystfibr.txt', delimiter = '\t')
SStand        = StandardScaler()
```


```
Train_80       = data.sample(frac = 0.8)
Test_20        = data.drop(Train_80.index)
```

common split training/  
test data is 80%/20%  
(see random sampling)

```
Scaled_Train_80 = SStand.fit_transform(Train_80)
Scaled_Test_20   = SStand.transform(Test_20)
```

```
sns.violinplot(data = pd.DataFrame(Scaled_Train_80, columns = data.columns))
plt.title("Cystfibr Train Dataset - standard normal scaled")
plt.xticks(rotation = 45)
plt.show()
```



```
sns.violinplot(data = pd.DataFrame(Scaled_Test_20, columns = data.columns))
plt.title("Cystfibr Test Dataset - standard normal scaled")
plt.xticks(rotation = 45)
plt.show()
```





subsamples might (and will!!!) indeed be different



→ running `transform` only on the test set and evaluating the model shows **how the model deals with data variance**



→ running `fit_transform` on the test set too is **data leakage** and “**cheating**”







running `fit_transform` on the test set too would transform the test data to *its* particular parameters









### Summary:

- we need to scale features because their **numerical value is arbitrary** (depends on unit system)
- this would confuse the model (**larger numerical values dominate the optimization process**, e.g. MSE etc)
- two common scaling methods:
  - StandardScaler
  - MinMaxScaler
- `fit_transform` for **training data**
- `transform` for **test data**



### Workflow:

- 1) Encode categorical features! Pick an encoding strategy based on the following criteria:
  - Are the feature values/states **nominal** (= no inherent order such as gene names, cities etc)?
  - Are the feature values/states **ordinal** (= there is an inherent order) and the feature values can be ranked (health status, tax bracket etc)?
  - **cardinality** of the feature (= how many states/values can a feature have)
  - **memory usage**
  
- 2) Scale/normalize features using common scaling methods:  
`StandardScaler`  
`MinMaxScaler`

see also **Encoding.ipynb**



## Outline

What is a feature?

Encodings


Normalization

Visualization



We learned already about the **violine plot**:

- shows the distribution of the variables
- quick look in order to get an idea about data range, distribution and outliers





Another great visualization tool is the **pair plot**:

- shows the distribution of the variables
- quick look in order to get an idea about data range, distribution and outliers
- **shows correlations** (see later)


```
data          = pd.read_csv('Cystfibr.txt', delimiter = '\t')
SStand        = StandardScaler()
Scaled        = SStand.fit_transform(data)

Pair = sns.pairplot(pd.DataFrame(Scaled, columns = data.columns), kind = 'kde')

Pair.map_lower(sns.scatterplot, c = 'k')
Pair.map_upper(sns.scatterplot, c = 'k')
Pair.map_diag(sns.histplot, stat = 'density')
```



pair plot:



age	sex	height	weight	bmp	fev1	rv	frc	tic	pemax
7	0	109	13.1	68	32	258	183	137	95
7	1	112	12.9	65	19	449	245	134	85
8	0	124	14.1	64	22	441	268	147	100

distribution of the feature itself



Thank you very much for your attention!

