Berkeley MSSE
MOLECULAR SCIENCE & SOFTWARE ENGINEERING

Lecture 09:

Linear and Logistic Regression

Markus Hohle

University California, Berkeley

Data Science for Scientific Computing

MSSE 277A, 3 Units

Syllabus

Lecture 1: Introduction to Data Science

Lecture 2: Data Sampling and Probability, Pandas

Lecture 3: Exploratory Data Analysis (EDA), and Regex Part I

Lecture 4: Exploratory Data Analysis (EDA), and Regex Part II

Lecture 5: Introduction to SQL

data acquisition and analysis

Lecture 6: Feature Analysis, Engineering and Encoding

Lecture 7: PCA, LDA and Correlation, Dimension Reduction

Lecture 8: Advanced Visualization Tools (UMAP, T-SNE) in Python

feature Selection and Analysis

project 1

Lecture 9: Introduction to linear and logistic regression; OLS

Spring Recess

Lecture 10: Avoiding Overfitting and Regularization (L1 & L2)

Lecture 11: Feature Selection via Regression Analysis

Lecture 12: Gradient Descent

Lecture 14: Clustering and Classification: Trees, KNN, K-Mean, SVM

Lecture 14: Clustering and Classification: Trees, GMM

Lecture 15: Feature Selection via ANNs ("leave-one-out", Accuracy Drop & Entropy)

modelling

project 2

Linear and Logistic Regression

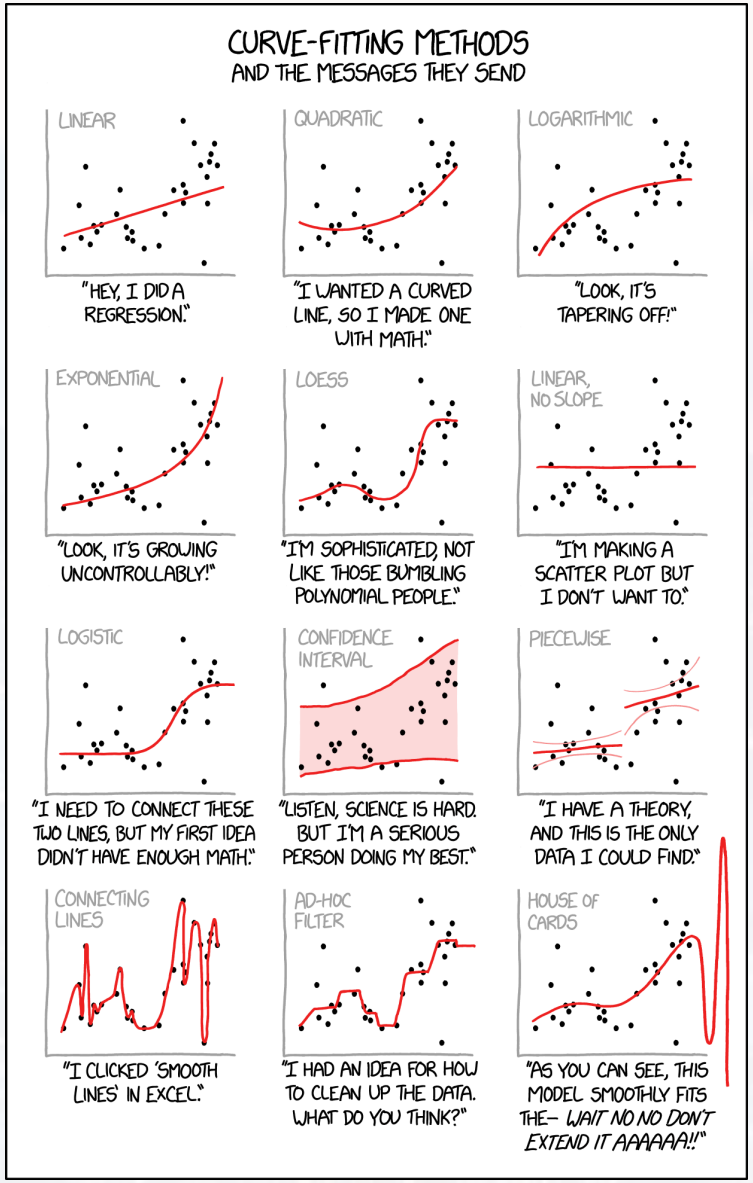

credit: CAUSEweb

Outline

Part I: **Linear Regression**

Part II: **Logistic Regression**

Linear and Logistic Regression

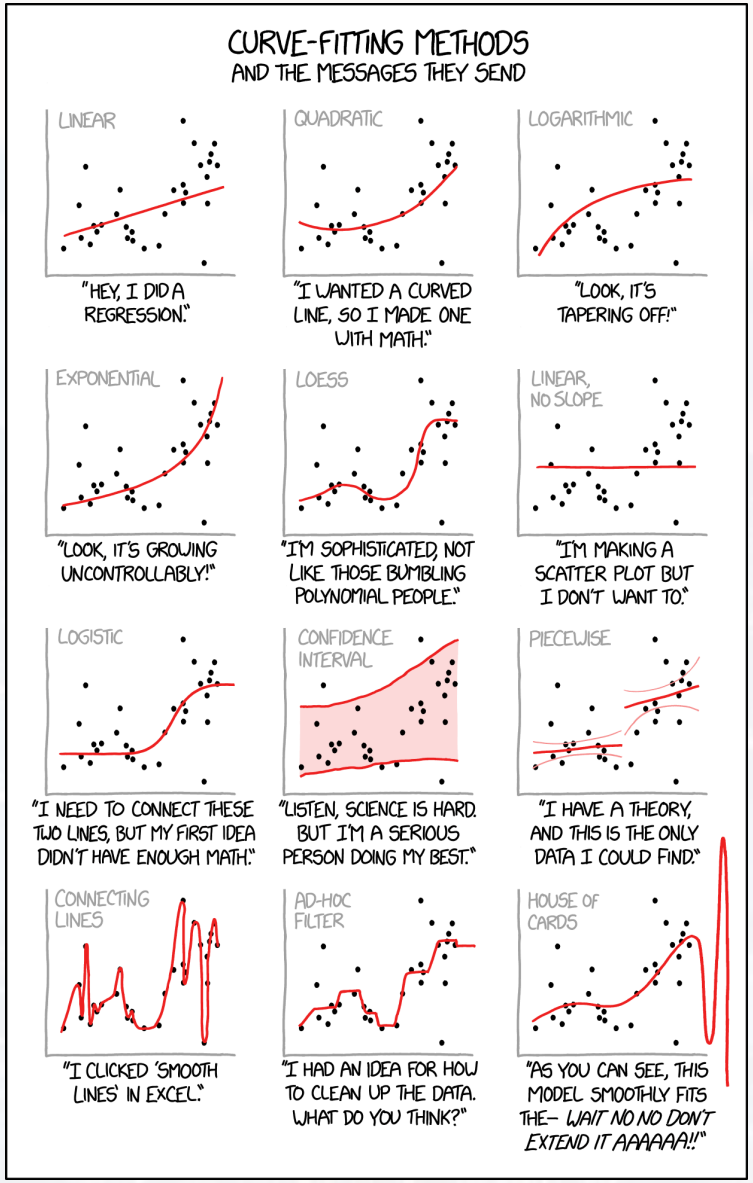

credit: CAUSEweb

Outline

Part I: **Linear Regression**

Mathematical Notation

What is Linear?

Some Statistics

a Python example

Linear and Logistic Regression

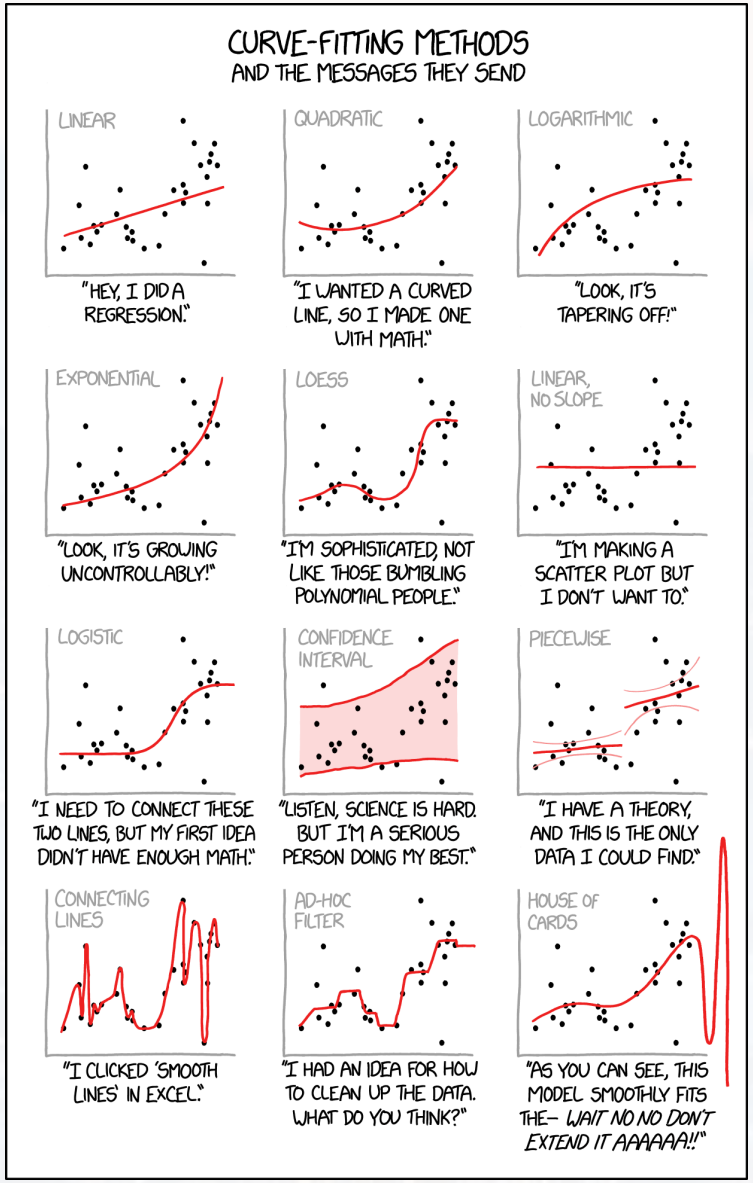

credit: CAUSEweb

Outline

Part I: **Linear Regression**

Mathematical Notation

What is Linear?

Some Statistics

a Python example

Linear Regression

Regression vs **Classification**

regression

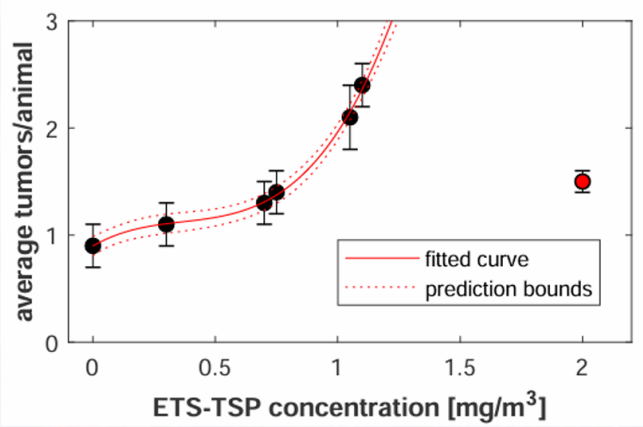

curve fit: finding model parameters by **minimizing** χ^2

$$\chi^2 = \sum_k \frac{(\hat{y}_k - y_k)^2}{\sigma_k^2}$$

fitted value $\hat{y}_k$

data point y_k

error bar σ_k

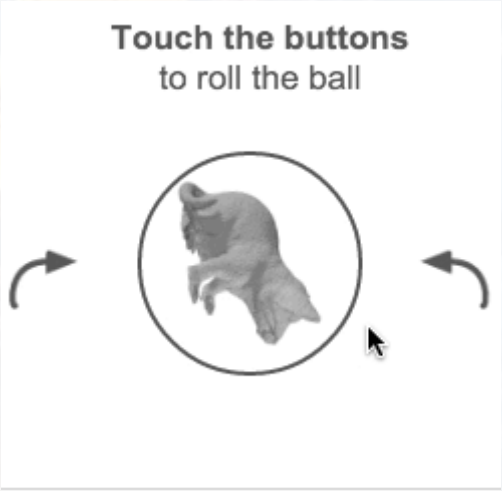

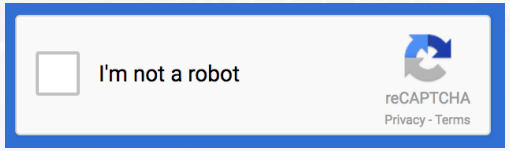

turning an image the right way:

- **maximizing** autocorrelation function
- training an AI

Regression vs Classification

regression

predicting numerical values

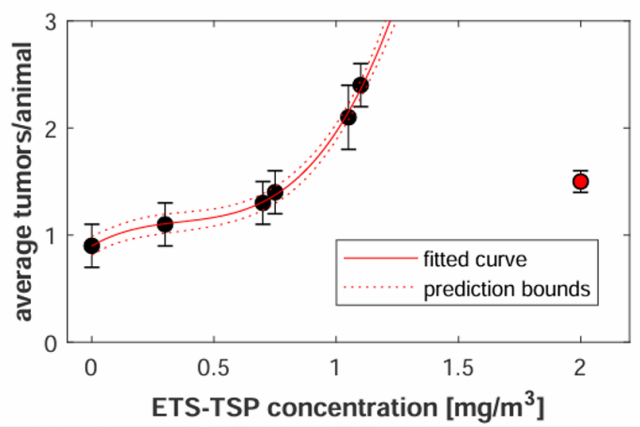

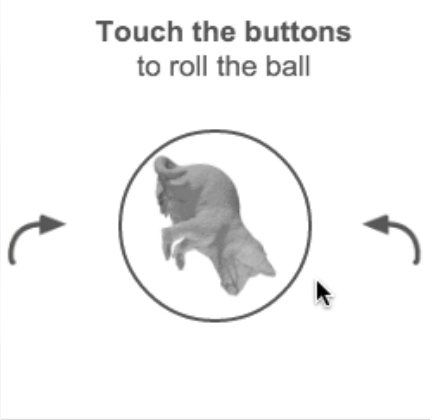

classification

predicting classes
via probabilities

cat

dog

note: we can use (non-linear) regression for classification!

Linear Regression

Mathematical Notation

idea: data point y_k in *N* dimensional space

→ $y_k = f(x_1, \dots x_n, \dots x_N) + \epsilon$ for each data point *k*

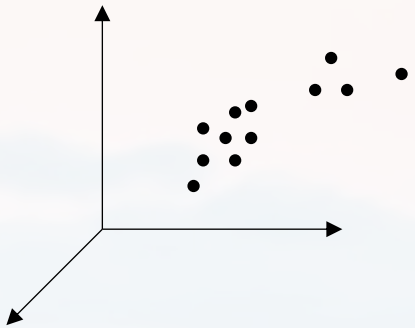

ansatz: $\boxed{y_k = \beta_0 + \sum_{n=1}^{N} \beta_n x_n + \epsilon}$ ***linear*** combination

y: response
x: regressors **(assumed to be independent)**
β: factors **(how a regressor contributes to the response)**
β_0: intercept
ε: error **(stochasticity of the data, assumed to be normally dist.)**

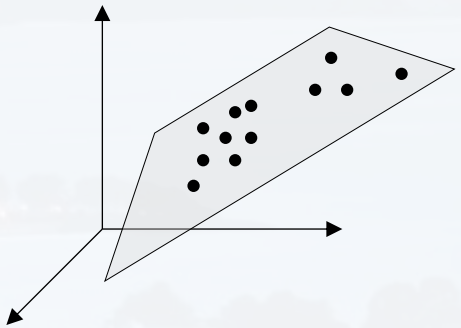

Linear Regression

Mathematical Notation

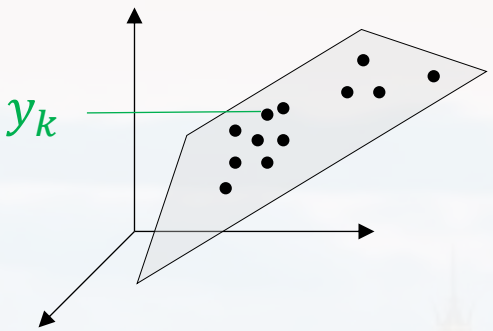

β_n: how strong does x_n affect y_k

y:	response
x:	regressors **(assumed to be independent)**
β:	factors
β_0:	intercept
ε:	error **(stochasticity of the data, assumed to be normally dist.)**

$$y_k = \beta_0 + \sum_{n=1}^{N} \beta_n x_n + \epsilon$$

	x_1	x_2	x_3	x_4	x_5	y_k
	molecular_weight	electronegativity	bond_lengths	num_hydrogen_bonds	logP	toxicity_score
0	341.704142	2.655846	3.094069	2	9.111473	80.928052
1	335.950798	3.222621	2.890393	7	8.928483	83.491053
2	235.203185	2.441153	2.482026	1	6.497307	61.840600
3	246.504930	2.766560	2.715472	7	7.450888	57.053778
4	437.938926	3.480105	3.595688	3	10.915629	131.326459

Linear and Logistic Regression

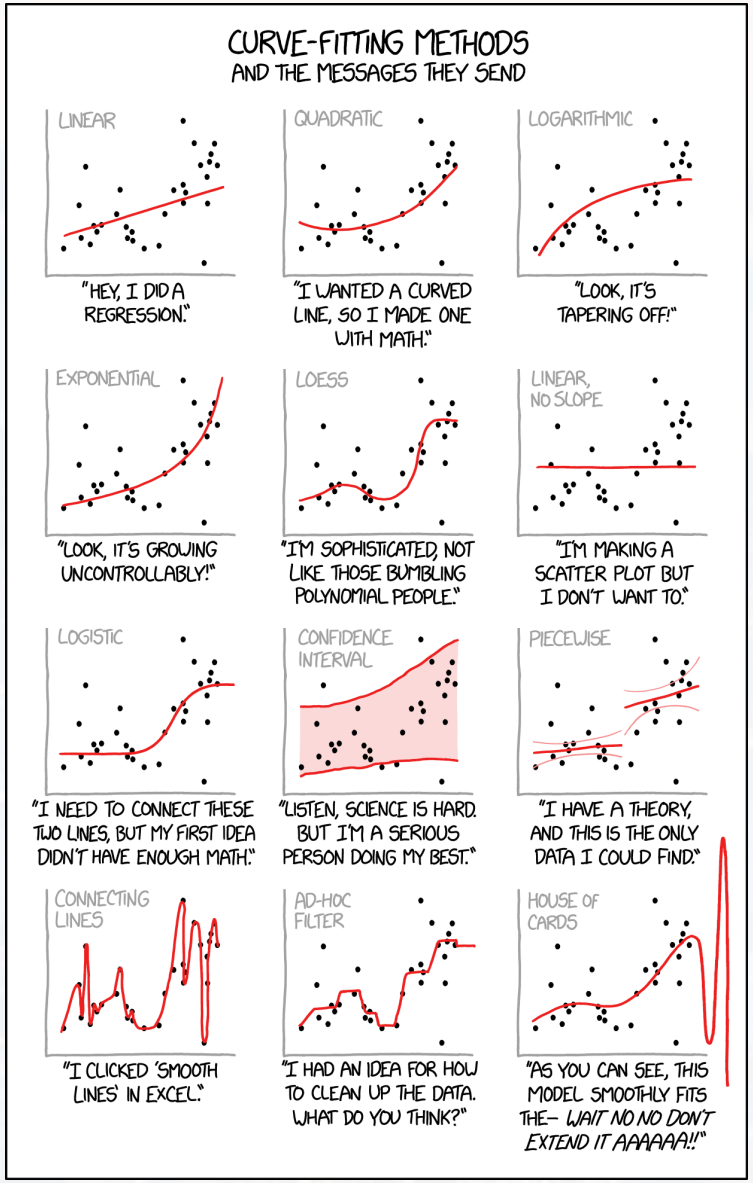

credit: CAUSEweb

Outline

Part I: **Linear Regression**

Mathematical Notation

What is Linear?

Some Statistics

a Python example

Linear Regression

What is Linear?

linear ≠ not curved

y:	response
x:	regressors **(assumed to be independent)**
β:	factors
β_0:	intercept
ε:	error **(stochasticity of the data, assumed to be normally dist.)**

$$y_k = \beta_0 + \sum_{n=1}^{N} \beta_n x_n^n + \epsilon$$

...is still linear

just define: $\bar{x}_n := x_n^n$

$$y_k = \beta_0 + \sum_{n=1}^{N} \beta_n x_n + \epsilon$$

$$y_k = \beta_1 x_n^{\beta_2}$$

...is still linear

just use log: $\bar{y}_k = \log(y_k) = \log(\beta_1) + \beta_2 \log(x_n) = \bar{\beta}_1 + \beta_2 \bar{x}_n$

note:

- as long as the linear structure $y_k = \beta_0 + \sum_{n=1}^{N} \beta_n x_n + \epsilon$ can be recovered by any transformation it is a linear model
- log transformation is pretty common (log fold change in DNA/RNAseq, intensity measured in dB etc)
- general: linear refers to the **factors**

Linear Regression

What is Linear?

general: linear refers to the **factors**

y:	response
x:	regressors **(assumed to be independent)**
β:	factors
β_0:	intercept
ε:	error **(stochasticity of the data, assumed to be normally dist.)**

$$y_k = \beta_0 + \beta_1 x_1 + \beta_2 x_2$$ 2*D* plane in 3*D* space

$$y_k = \beta_0 + \beta_1 {x_1}^2 + \beta_2 {x_2}^2$$ 2*D* parabolic

$$y_k = \beta_0 + \beta_1 {x_1}^2 - \beta_2 {x_2}^2$$ 2*D* hyperbolic

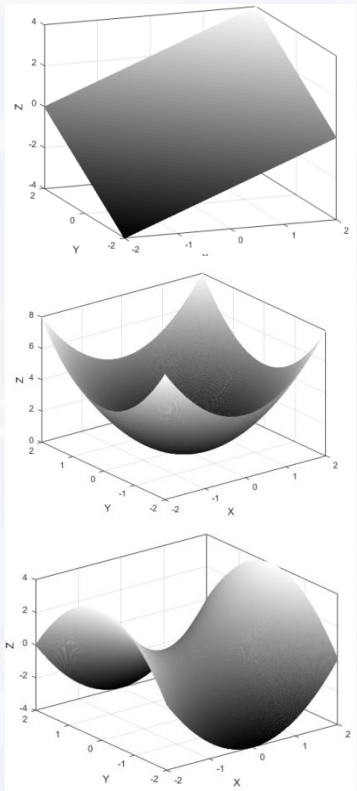

all linear

$$y_k = \beta_0 + \sum_{n=1}^{N} \beta_n x_n + \epsilon$$

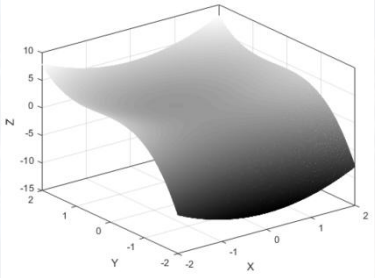

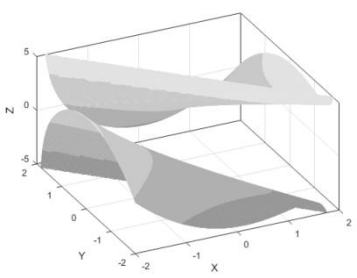

...and many more...

Linear Regression

What is Linear?

...what is **not** linear?

$$y_k = \beta_0 + \beta_1 x_n^{\beta_2}$$

log trick does not work here

y:	response
x:	regressors **(assumed to be independent)**
β:	factors
β_0:	intercept
ε:	error **(stochasticity of the data, assumed to be normally dist.)**

surfaces that can't be modelled by a linear transformation:

$$y_k = \beta_0 + \sum_{n=1}^{N} \beta_n x_n + \epsilon$$

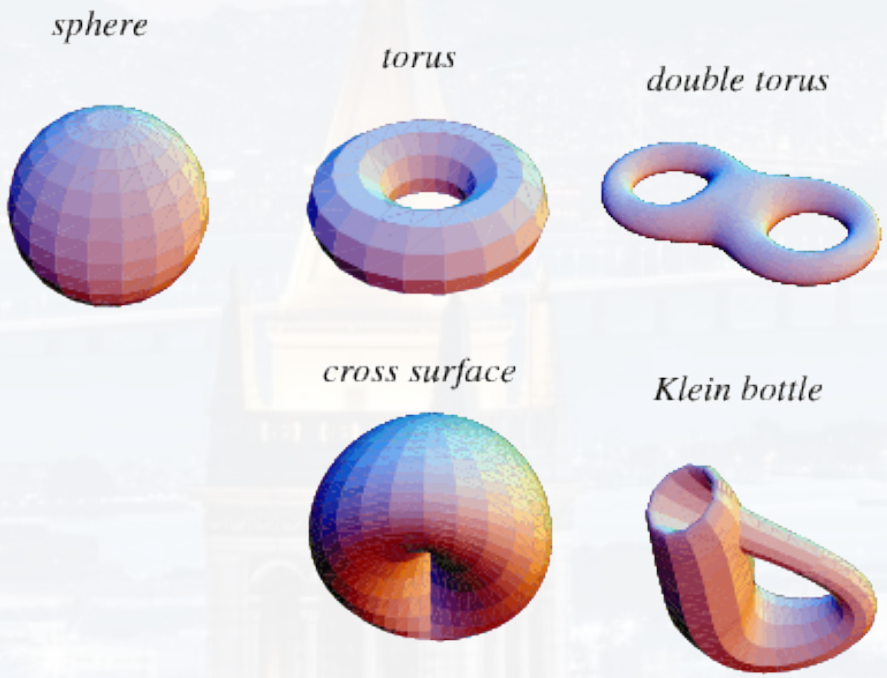

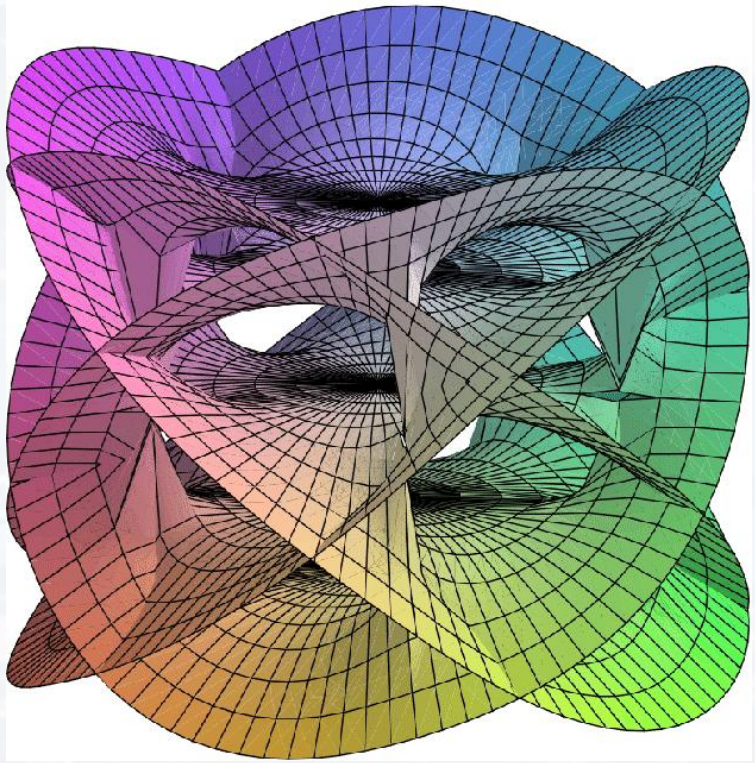

Linear and Logistic Regression

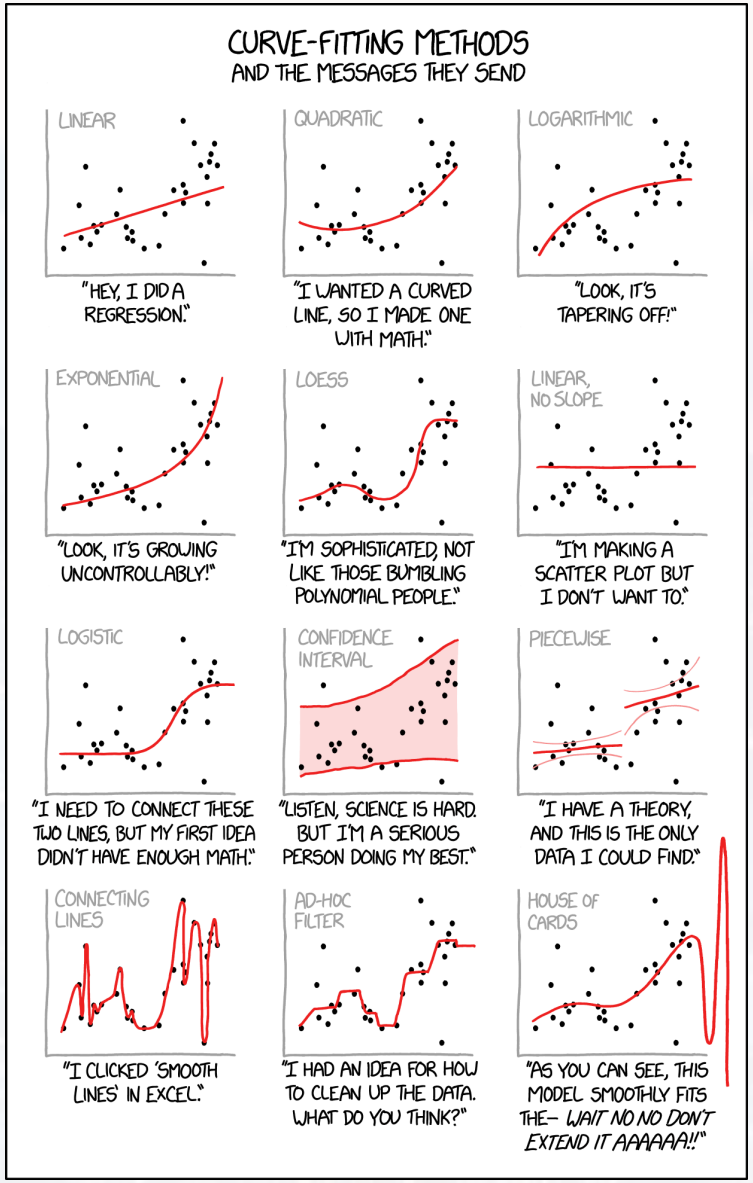

credit: CAUSEweb

Outline

Part I: **Linear Regression**

Mathematical Notation

What is Linear?

Some Statistics

a Python example

Linear Regression

y:	response
x:	regressors **(assumed to be independent)**
β:	factors
β_0:	intercept
ε:	error **(stochasticity of the data, assumed to be normally dist.)**

for *K* data points in *N* dimensional space

$$y_k = \beta_0 + \sum_{n=1}^{N} \beta_n x_n + \epsilon$$

$$\underbrace{\begin{pmatrix} y_1 \\ \dots \\ y_k \\ \dots \\ \dots \\ y_K \end{pmatrix}}_{Y} = \underbrace{\begin{pmatrix} 1 & x_{11} & x_{12} & \dots & x_{1n} & \dots & x_{1N} \\ \dots & \dots & & & \dots & & \\ 1 & x_{k1} & & & x_{kn} & & \\ 1 & \dots & & & \dots & & \\ 1 & x_{K1} & x_{K2} & \dots & x_{Kn} & \dots & x_{KN} \end{pmatrix}}_{X} \underbrace{\begin{pmatrix} \beta_0 \\ \beta_1 \\ \dots \\ \beta_n \\ \dots \\ \beta_N \end{pmatrix}}_{\beta} + \underbrace{\begin{pmatrix} \varepsilon_1 \\ \varepsilon_2 \\ \dots \\ \varepsilon_k \\ \dots \\ \varepsilon_K \end{pmatrix}}_{\varepsilon}$$

$$\boxed{Y = X\beta + \varepsilon}$$

<u>fitting:</u> finding the best β in terms of minimizing the errors

$$\hat{\beta} = \underset{\beta}{argmin} \left\{ \frac{1}{K} \|Y - X\beta\|^2 \right\} \qquad (Y - X\beta)^T (Y - X\beta) = \sum_k {\varepsilon_k}^2$$

Linear Regression

Some Statistics

$$\boxed{Y = X\beta + \varepsilon}$$

y:	response
x:	regressors **(assumed to be independent)**
β:	factors
β_0:	intercept
ε:	error **(stochasticity of the data, assumed to be normally dist.)**

<u>fitting:</u> finding the best β in by minimizing the errors

$$(Y - X\beta)^T(Y - X\beta) = \sum_k {\varepsilon_k}^2$$

$$\frac{\partial}{\partial\beta}\sum_k {\varepsilon_k}^2 = 0 \longrightarrow \beta_{best} = \hat{\beta} = (X^TX)^{-1}X^TY \longrightarrow$$

the model

$$\hat{\boldsymbol{Y}} = \boldsymbol{X}\hat{\boldsymbol{\beta}} = \underbrace{X(X^TX)^{-1}X^T}_{\text{hat matrix } \boldsymbol{H}}Y$$

<u>some properties of the hat matrix:</u>

- $H = H^T$ (symmetry)
- $HH = H \rightarrow H^n = H$ (idempotency)

$$\hat{\boldsymbol{Y}} = \boldsymbol{X}\hat{\boldsymbol{\beta}} = \boldsymbol{X}\left(\boldsymbol{X}^T\boldsymbol{X}\right)^{-1}\boldsymbol{X}^T\boldsymbol{Y}$$

all observables!

<u>evaluating the fit:</u>

$$\hat{\varepsilon} = Y - X\hat{\beta} = Y - \hat{Y} = (I - H)Y$$

$$\hat{\varepsilon}^T\hat{\varepsilon} = [(I - H)Y]^T(I - H)Y = Y^T(I - H)^T(I - H)Y = Y^T(I - H)Y$$

sum of **s**quared **e**rrors (SSE)

Linear Regression

summary:

y:	response
x:	regressors **(assumed to be independent)**
β:	factors
β_0:	intercept
ε:	error **(stochasticity of the data, assumed to be normally dist.)**
K:	number of data points
N:	number of model param

the model: $Y = X\beta + \varepsilon$

the fit: $\hat{Y} = X\hat{\beta} = X(X^T X)^{-1} X^T Y$

sum of squared errors (SSE): $\hat{\varepsilon}^T \hat{\varepsilon} = Y^T (I - H) Y$

(after the fit)

mean of squared errors (MSE): $\frac{\hat{\varepsilon}^T \hat{\varepsilon}}{K-N}$

(after the fit)

often fit quality is judged by $R^2 := 1 - \frac{\sum_k (\hat{y}_k - y_k)^2}{\sum_k (y_k - \langle y \rangle)^2}$

or adjusted R^2 $\bar{R}^2 := R^2 - (1 - R^2) \frac{K}{N-K-1}$

and it is said that the fit is good if R^2 is close to one....

...but that is not correct...

Linear Regression

Some Statistics

$$\chi^2_{red} = \frac{1}{df}\sum_{k=1}^{K}\left(\frac{y_k - \hat{y}_k}{\sigma_k}\right)^2$$

$$df = K - N - 1$$

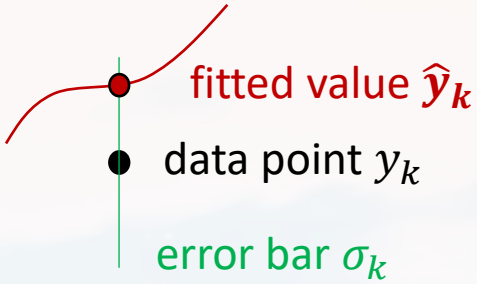

y:	response
x:	regressors **(assumed to be independent)**
β:	factors
β_0:	intercept
ε:	error **(stochasticity of the data, assumed to be normally dist.)**
K:	number of data points
N:	number of model param

$$R^2 = 1 - \frac{\sum_{i=1}^{K}(y_k - \hat{y}_k)^2}{\sum_{i=1}^{K}(y_k - \bar{y})^2}$$

note: do not confuse R^2 with Pearsons coefficient: $\rho = \frac{cov(x,y)}{\sqrt{var(x)var(y)}}$

$(y_k - \hat{y}_k)^2$ *residual sum of squares*

$(y_k - \bar{y})^2$ variance of **the data** *(aka total sum of squares)*

Linear Regression

Some Statistics

$$\chi^2_{red} = \frac{1}{df}\sum_{k=1}^{K}\left(\frac{y_k - \hat{y}_k}{\sigma_k}\right)^2$$

$$df = K - N - 1$$

- scales difference between model and data to the error bars
- can be directly translated to a p-value via the Students distribution

H0: the fitted model has in fact generated the data

$$R^2 = 1 - \frac{\sum_{k=1}^{K}(y_k - \hat{y}_k)^2}{\sum_{k=1}^{K}(y_k - \bar{y})^2}$$

data vs model / variance of the data

$\bar{y}$: mean of the data point values

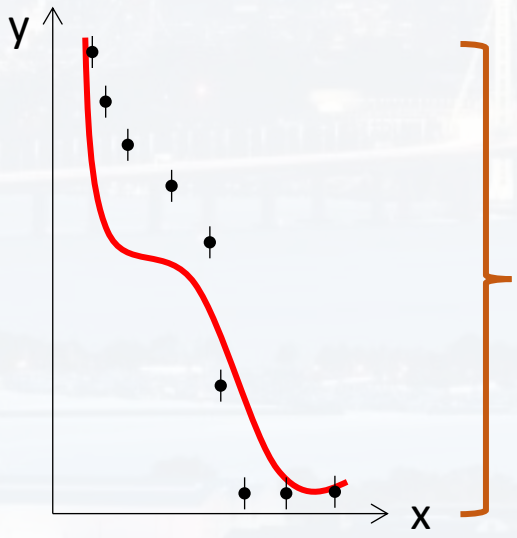

data variance can be huge (i. e. exponential functions) → R^2 could be around 1.0 even if fit is completely off!

Linear Regression

Some Statistics

$$\chi^2_{red} = \frac{1}{df}\sum_{k=1}^{K}\left(\frac{y_k - \hat{y}_k}{\sigma_k}\right)^2$$

$$df = K - N - 1$$

- scales difference between model and data to the error bars
- can be directly translated to a p-value via the Students distribution

H0: the fitted model has in fact generated the data

$$R^2 = 1 - \frac{\sum_{k=1}^{K}(y_k - \hat{y}_k)^2}{\sum_{k=1}^{K}(y_k - \bar{y})^2} \quad \frac{\text{data vs model}}{\text{variance of the data}}$$

$\bar{y}$: mean of the data point values

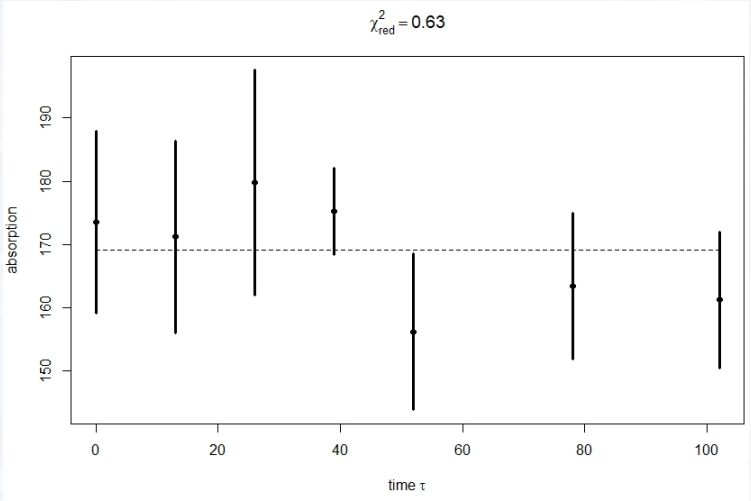

$$\frac{\text{data vs model}}{\text{variance of the data}} \approx 1 \rightarrow R^2 = 0$$

→ **although the fit is good**

Linear Regression

Some Statistics

$$\chi^2_{red} = \frac{1}{df}\sum_{k=1}^{K}\left(\frac{y_k - \hat{y}_k}{\sigma_k}\right)^2$$

$$df = K - N - 1$$

- scales difference between model and data to the error bars
- can be directly translated to a p-value via the Students distribution

H0: the fitted model has in fact generated the data

$$R^2 = 1 - \frac{\sum_{k=1}^{K}(y_k - \hat{y}_k)^2}{\sum_{k=1}^{K}(y_k - \bar{y})^2}$$

data vs model / variance of the data

$\bar{y}$: mean of the data point values

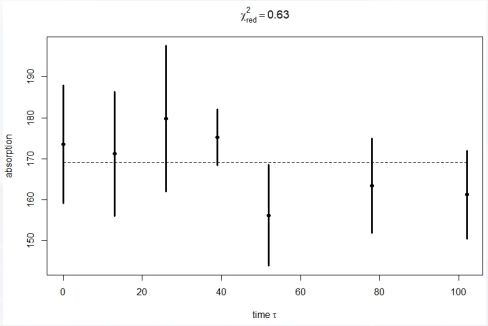

<u>conclusion:</u>

- R^2 is not a measure of the fit quality (but χ^2 is)
- error bars are important
- **given a good fit**, R^2 tells how strong the dependent variable responds to the independent variable

Also, Wiki is full of examples…

…and warnings (see "caveats" therein)

Linear and Logistic Regression

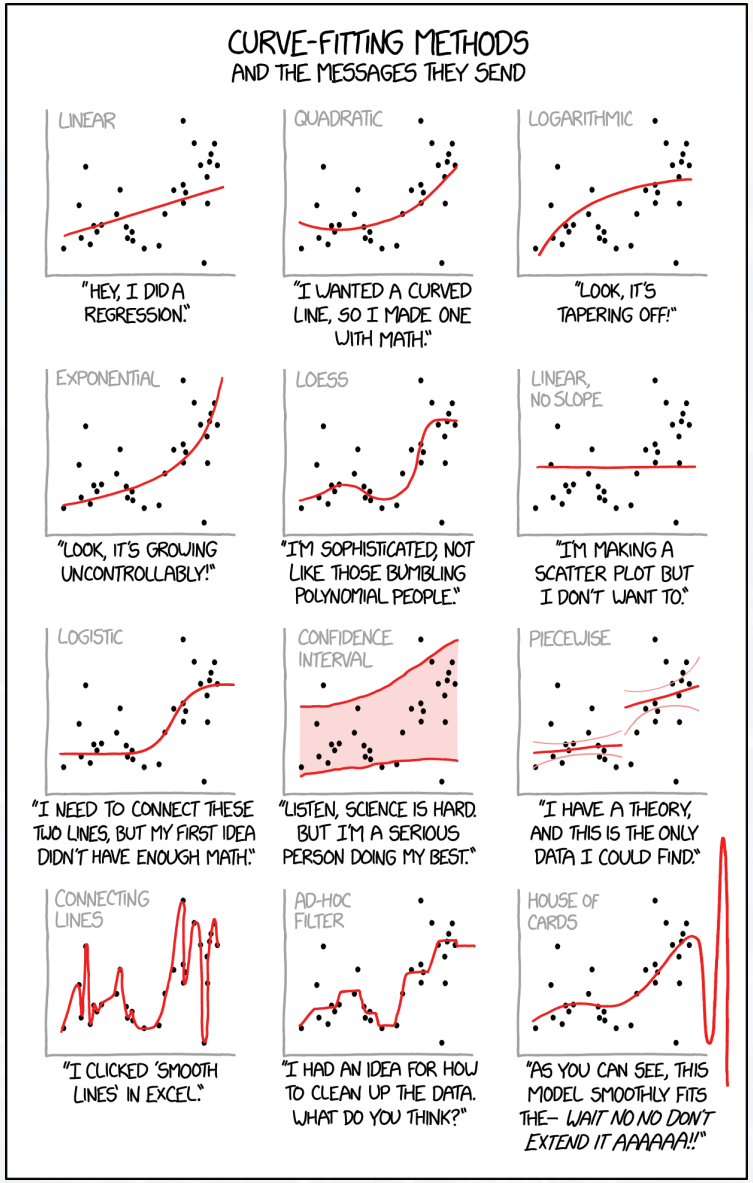

credit: CAUSEweb

Outline

Part I: **Linear Regression**

Mathematical Notation

What is Linear?

Some Statistics

a Python example

Linear Regression

A Python Example

```
import numpy as np

import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import pylab
import scipy.stats as stats
import statsmodels.api as sm

from statsmodels.formula.api import ols
from sklearn.preprocessing import MinMaxScaler
```

reading .xlsx
.csv
.txt
…

standard plots

fancy plots: here a pair-plot

Q-Q plot

the actual super tool for superb data analysis

scaling and normalizing

Linear Regression

1) loading data
2) plotting data
3) scaling data
4) fitting model
5) evaluating model

```
Train  = pd.read_csv("molecular_train_gbc.csv")
Test   = pd.read_csv("molecular_test_gbc.csv")
```

	x_1	x_2	x_3	x_4	x_5	y_k
	molecular_weight	electronegativity	bond_lengths	num_hydrogen_bonds	logP	toxicity_score
0	341.704142	2.655846	3.094069	2	9.111473	80.928052
1	335.950798	3.222621	2.890393	7	8.928483	83.491053
2	235.203185	2.441153	2.482026	1	6.497307	61.840600
3	246.504930	2.766560	2.715472	7	7.450888	57.053778
4	437.938926	3.480105	3.595688	3	10.915629	131.326459

$$y_k = \beta_0 + \sum_{n=1}^{N} \beta_n x_n + \epsilon$$

y: toxicity_score

x_n : molecular_weight, electronegativity, bond_lengths, num_hydrogen_bonds, logP

Linear Regression

1) loading data
2) plotting data
3) scaling data
4) fitting model
5) evaluating model

```
Train  = pd.read_csv("molecular_train_gbc.csv")
Test   = pd.read_csv("molecular_test_gbc.csv")


out = sns.pairplot(Train, kind = "kde", \
                   plot_kws = {'color':[176/255, 224/255, 230/255]},\
                   diag_kws = {'color': 'black'})
out.map_offdiag(plt.scatter, color = 'black')
```

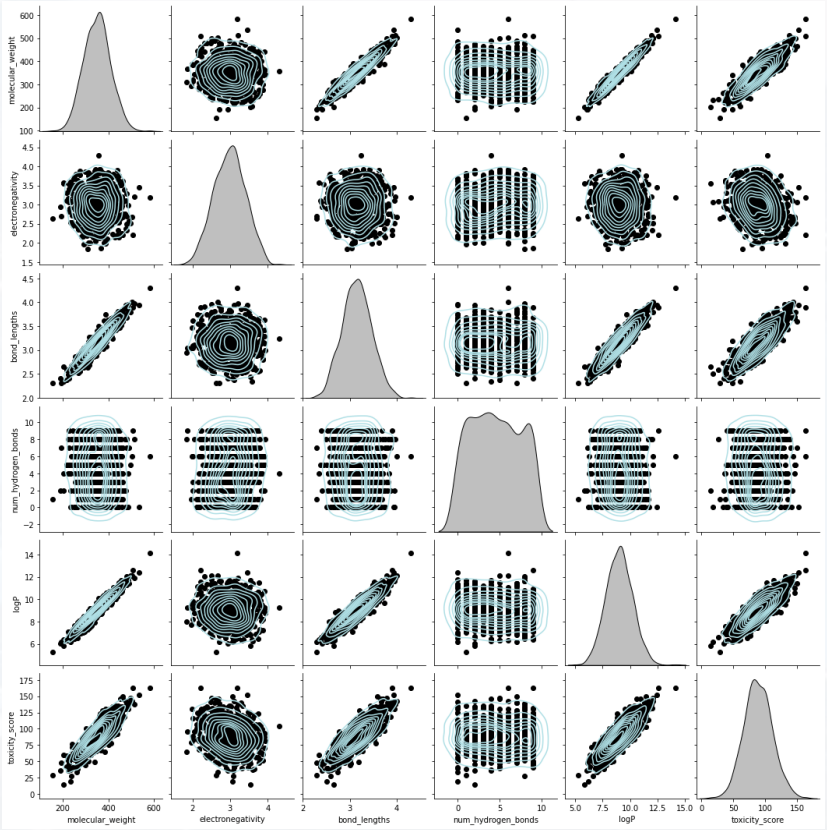

Linear Regression

1) loading data
2) plotting data
3) scaling data
4) fitting model
5) evaluating model

```
Train  = pd.read_csv("molecular_train_gbc.csv")
Test   = pd.read_csv("molecular_test_gbc.csv")
```

```
out = sns.pairplot(Train, kind = "kde", \
                   plot_kws = {'color':[176/255, 224/255, 230/255]},\
                   diag_kws = {'color': 'black'})
out.map_offdiag(plt.scatter, color = 'black')
```

```
scaler  = MinMaxScaler(feature_range = (0, 1))
TrainS  = scaler.fit_transform(Train)
TestS   = scaler.transform(Test)
```

the scaler returns an `np.array`
→ convert back to `data frame`

```
TrainS = pd.DataFrame(TrainS, columns = Train.columns)
TestS  = pd.DataFrame(TestS, columns = Train.columns)
```

Linear Regression

1) loading data
2) plotting data
3) scaling data
4) fitting model
5) evaluating model

```
TrainS = pd.DataFrame(TrainS, columns = Train.columns)
TestS  = pd.DataFrame(TestS,  columns = Train.columns)
```

```
equation = 'toxicity_score ~ ' + '+'.join(Train.columns[:-1])
print(equation)
```

$$y_k = \beta_0 + \sum_{n=1}^{N} \beta_n x_n + \epsilon$$

```
toxicity_score ~      molecular_weight + electronegativity +
                      bond_lengths + num_hydrogen_bonds + logP
```

```
my_model = ols(equation, data = TrainS).fit()
my_model.summary()
```

OLS (ordinary least squares)

Linear Regression

1) loading data
2) plotting data
3) scaling data
4) fitting model
5) evaluating model

```
my_model.summary()
```

not the fit quality!

```
                            OLS Regression Results
==============================================================================
Dep. Variable:          toxicity_score   R-squared:                       0.790
Model:                             OLS   Adj. R-squared:                  0.789
Method:                  Least Squares   F-statistic:                     597.5
Date:                 Fri, 13 Sep 2024   Prob (F-statistic):          3.34e-266
Time:                         20:57:10   Log-Likelihood:                 1013.0
No. Observations:                  800   AIC:                            -2014.
Df Residuals:                      794   BIC:                            -1986.
Df Model:                            5
Covariance Type:             nonrobust
=====================================================================================
                        coef    std err          t      P>|t|      [0.025      0.975]
-------------------------------------------------------------------------------------
Intercept             0.1494      0.012     12.533      0.000       0.126       0.173
molecular_weight      0.7961      0.089      8.982      0.000       0.622       0.970
electronegativity    -0.1682      0.015    -11.591      0.000      -0.197      -0.140
bond_lengths          0.0204      0.049      0.417      0.677      -0.076       0.116
num_hydrogen_bonds    0.0035      0.008      0.458      0.647      -0.011       0.018
logP                  0.1246      0.072      1.723      0.085      -0.017       0.267
==============================================================================
Omnibus:                        2.249   Durbin-Watson:                   1.984
Prob(Omnibus):                  0.325   Jarque-Bera (JB):                2.240
Skew:                          -0.129   Prob(JB):                        0.326
Kurtosis:                       2.980   Cond. No.                         65.6
==============================================================================

Notes:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
```

number of data points is much larger than the number of regressors
→ degree of freedom approx. no of obs

p-value for constant model

p-values for factors

2σ conf range of factors

$$y_k = \beta_0 + \sum_{n=1}^{N} \beta_n x_n + \epsilon$$

Linear Regression

1) loading data
2) plotting data
3) scaling data
4) fitting model
5) evaluating model

more accurate: determining **the p-values for the factors using ANOVA** for the corresponding residuals

```
table     = sm.stats.anova_lm(my_model, typ = 1)
print(table)
```

$$y_k = \beta_0 + \sum_{n=1}^{N} \beta_n x_n + \epsilon$$

	df	sum_sq	mean_sq	F	PR(>F)
molecular_weight	1.0	13.346285	13.346285	2847.525516	8.024085e-265
electronegativity	1.0	0.640388	0.640388	136.631363	3.085962e-29
bond_lengths	1.0	0.000684	0.000684	0.145954	7.025342e-01
num_hydrogen_bonds	1.0	0.000703	0.000703	0.150055	6.985866e-01
logP	1.0	0.013917	0.013917	2.969353	8.524510e-02
Residual	794.0	3.721459	0.004687	NaN	NaN

vs from t-test

0.0000
0.0000
0.6766
0.6473
0.0852

Linear Regression

1) loading data
2) plotting data
3) scaling data
4) fitting model
5) evaluating model

```
residuals = my_model.resid

plt.hist(residuals, color = 'w', edgecolor = 'black')
plt.title('fit residuals')
plt.ylabel('#')
plt.xlabel('value')
plt.show()
```

$$y_k = \beta_0 + \sum_{n=1}^{N} \beta_n x_n + \epsilon$$

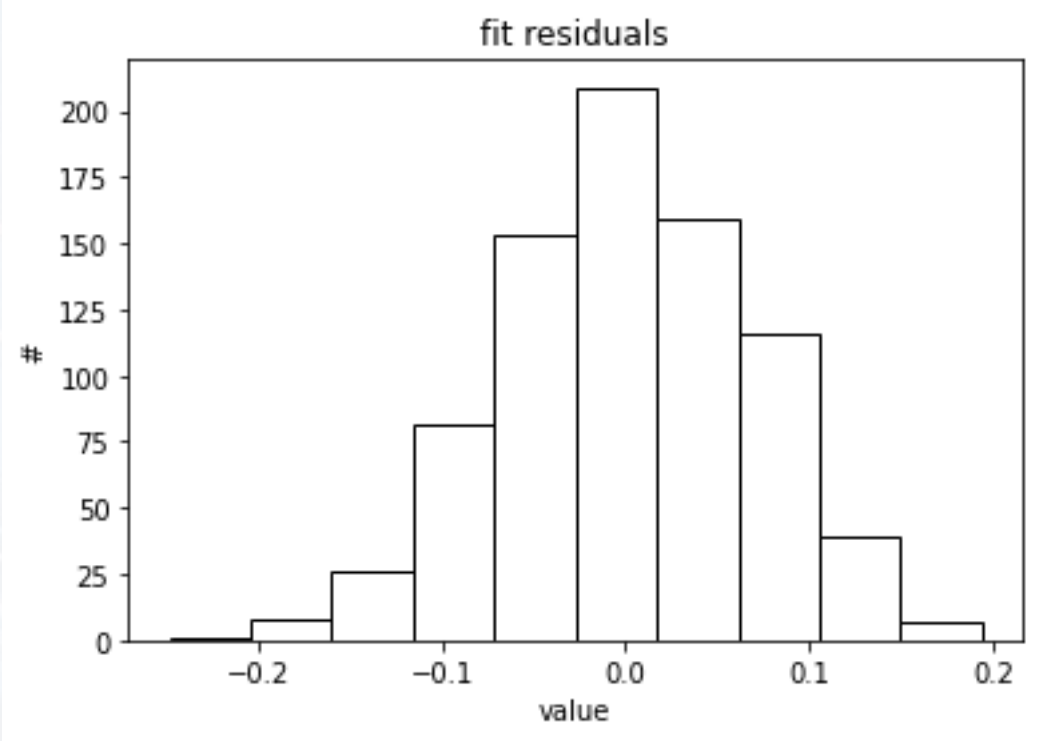

residuals approx. normally distributed around μ = 0

Linear Regression

1) loading data
2) plotting data
3) scaling data
4) fitting model
5) evaluating model

```
residuals = my_model.resid

plt.hist(residuals, color = 'w', edgecolor = 'black')
plt.title('fit residuals')
plt.ylabel('#')
plt.xlabel('value')
plt.show()

stats.probplot(residuals, dist = "norm", plot = pylab)
pylab.show()
```

residuals approx. normally distributed around μ = 0

$$y_k = \beta_0 + \sum_{n=1}^{N} \beta_n x_n + \epsilon$$

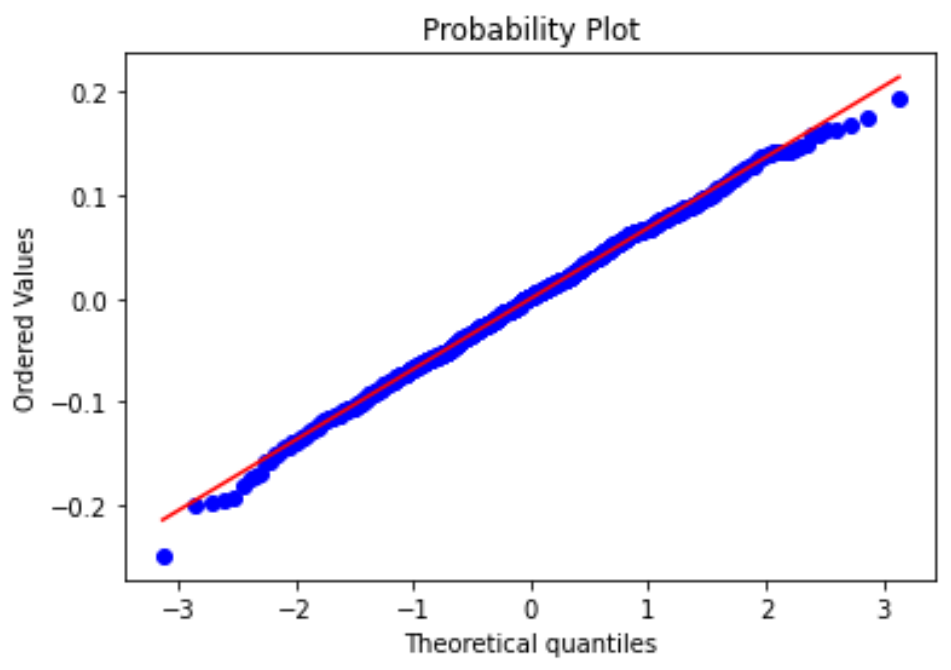

Linear Regression


```
Ypred  = my_model.predict(TestS)

higher = np.max([Ypred, TestS.toxicity_score])
lower  = np.min([Ypred, TestS.toxicity_score])

plt.plot([lower, higher], [lower, higher], c = [0, 0, 0, 0.2],\
         linewidth = 4)
plt.scatter(TestS.toxicity_score, Ypred, marker = '.', c = 'k')
plt.ylabel('prediction')
plt.xlabel('toxicity score')
plt.show()
```

1) loading data
2) plotting data
3) scaling data
4) fitting model
5) evaluating model

$$y_k = \beta_0 + \sum_{n=1}^{N} \beta_n x_n + \epsilon$$

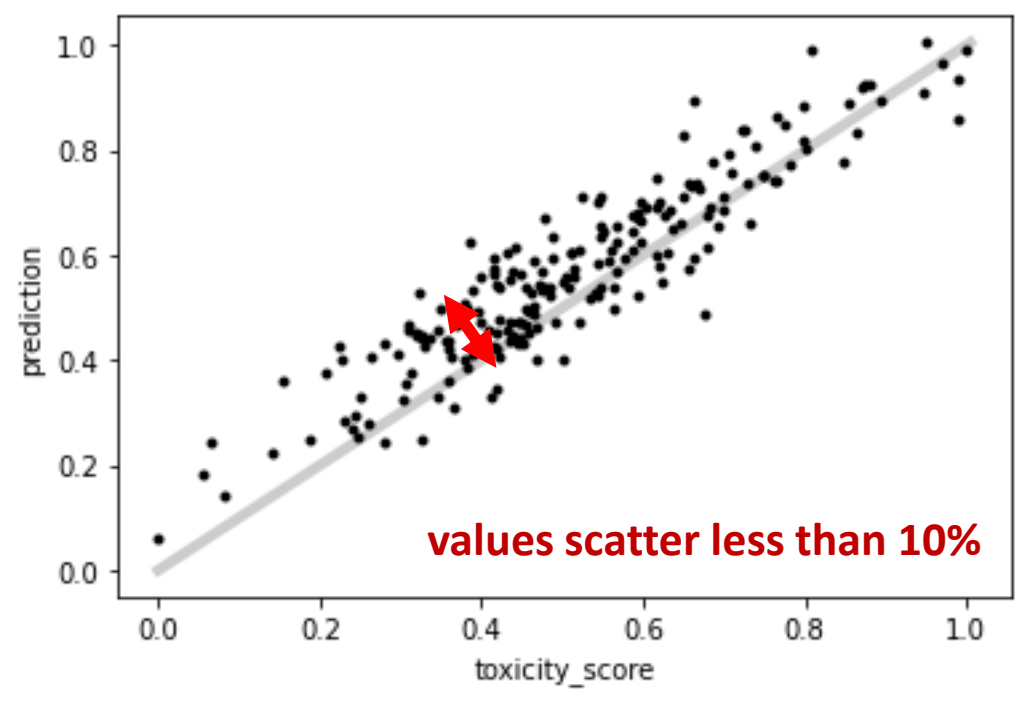

Linear Regression

1) loading data
2) plotting data
3) scaling data
4) fitting model
5) evaluating model

$$y_k = \beta_0 + \sum_{n=1}^{N} \beta_n x_n + \epsilon$$

```
Ypred  = my_model.predict(TestS)

higher = np.max([Ypred, TestS.toxicity_score])
lower  = np.min([Ypred, TestS.toxicity_score])

plt.plot([lower, higher], [lower, higher], c = [0, 0, 0, 0.2], linewidth = 4)
plt.scatter(TestS.toxicity_score, Ypred, marker = '.', c = 'k')
plt.ylabel('prediction')
plt.xlabel('toxicity score')
plt.show()
```

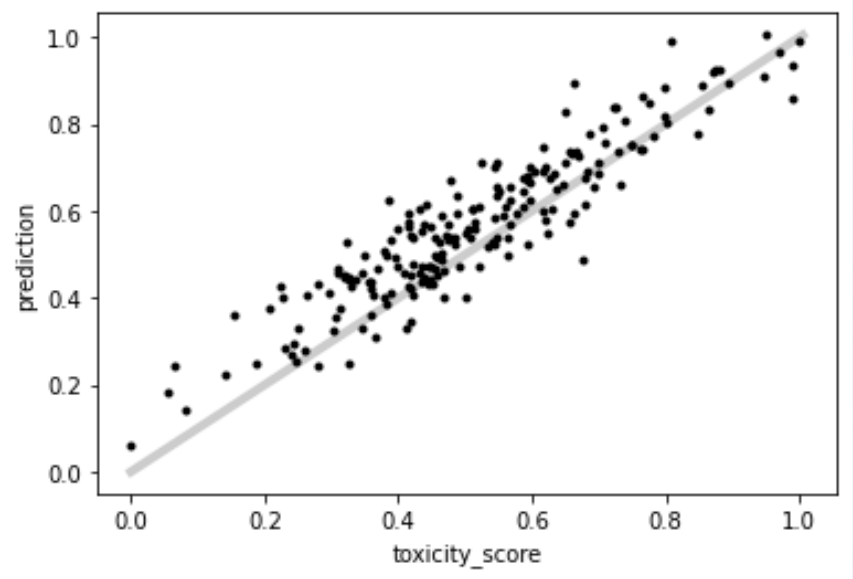


```
mean_dev = np.sum( abs(TestS.toxicity_score - Ypred) )/len(Ypred)
print(mean_dev)
```

5%

Linear Regression

note:

- we saw that many regressors are highly correlated
- run a PCA first
- so called **L1 regularization** deals with highly correlated regressors (sparicty)
- see next module!

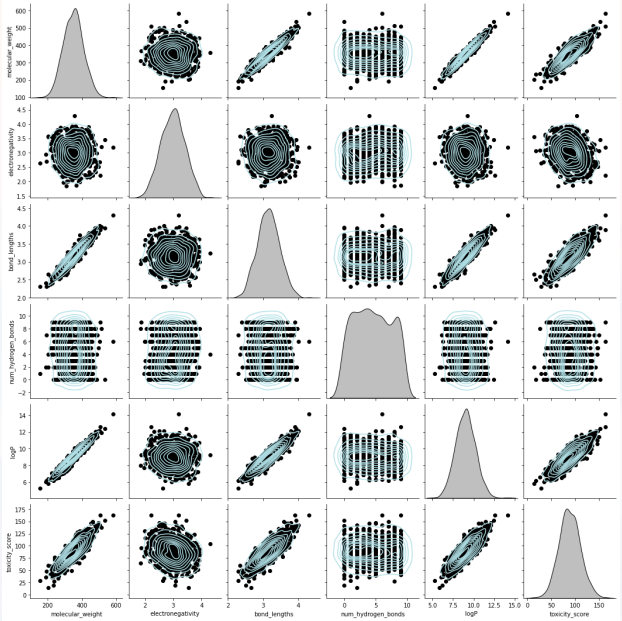

Linear and Logistic Regression

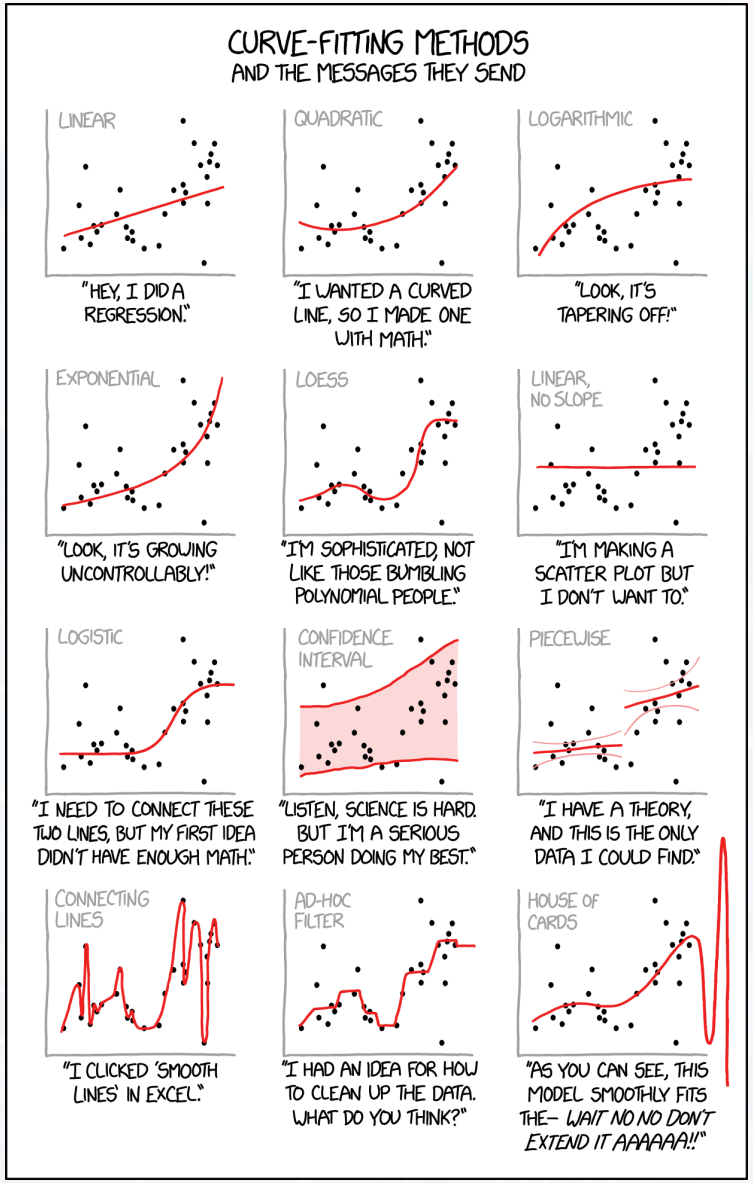

credit: CAUSEweb

Outline

Part II: **Logistic Regression**

Idea

The Factors

a Python example

Linear and Logistic Regression

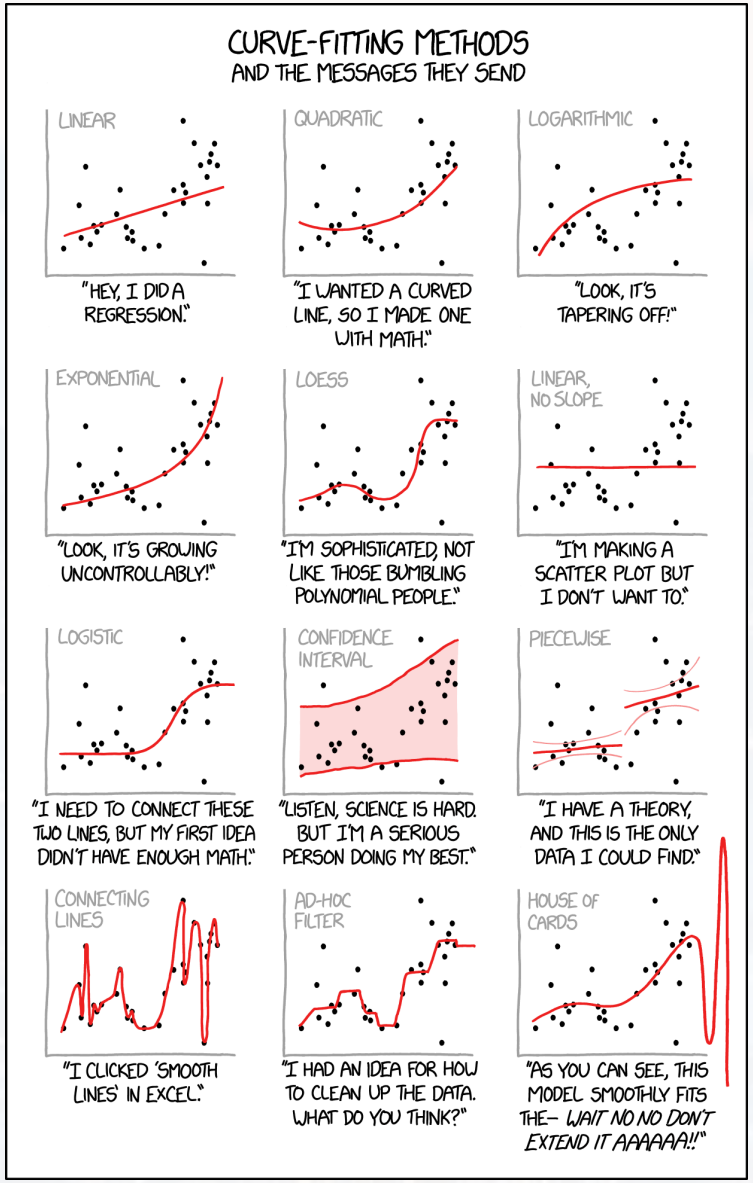

credit: CAUSEweb

Outline

Part II: **Logistic Regression**

Idea

The Factors

a Python example

Logistic Regression

Idea

linear model: regressors are continuous or categorical, **response is continuous**

logistic model: **response is categorical**

y:	response
x:	regressors **(assumed to be independent)**
β:	factors
β_0:	intercept
ε:	error **(stochasticity of the data, assumed to be normally dist.)**

	molecular_weight	electronegativity	bond_lengths	num_hydrogen_bonds	logP	label
0	341.704142	2.655846	3.094069	2	9.111473	Toxic
1	335.950798	3.222621	2.890393	7	8.928483	Toxic
2	235.203185	2.441153	2.482026	1	6.497307	Non-Toxic
3	246.504930	2.766560	2.715472	7	7.450888	Non-Toxic
4	437.938926	3.480105	3.595688	3	10.915629	Non-Toxic

Logistic Regression

Idea

linear model: regressors are continuous or categorical, **response is continuous**

logistic model: **response is categorical**

y:	response
x:	regressors **(assumed to be independent)**
β:	factors
β_0:	intercept
ε:	error **(stochasticity of the data, assumed to be normally dist.)**

dichotomic model: **probability** to be in state A) → ***p***

probability to be in state B) → ***1 - p***

label
Toxic
Toxic
Non-Toxic
Non-Toxic
Non-Toxic

<u>ansatz:</u>

$$\boxed{log\left(\frac{p}{1-p}\right) = \beta_0 + \sum_{n=1}^{N} \beta_n x_n + \epsilon}$$

log odds ratio: linear model

Logistic Regression

Idea

dichotomic model:

probability to be in state A) → ***p***

probability to be in state B) → ***1 - p***

y:	response
x:	regressors **(assumed to be independent)**
β:	factors
β_0:	intercept
ε:	error **(stochasticity of the data, assumed to be normally dist.)**

label
Toxic
Toxic
Non-Toxic
Non-Toxic
Non-Toxic

<u>ansatz:</u>

$$\boxed{log\left(\frac{p}{1-p}\right) = \beta_0 + \sum_{n=1}^{N} \beta_n x_n + \epsilon}$$

log odds ratio: linear model

→ probability for being in a certain state

$$\boxed{p = \frac{e^{\beta_0+\beta_1 x_1+\cdots}}{1+e^{\beta_0+\beta_1 x_1+\cdots}}}$$

often:

$$\boxed{logit(p) = log\left(\frac{p}{1-p}\right)}$$

examples:

- probability that a gene has been mutated
- probability of being diseased (cancer, alzheimer etc) as function of age, environmental influence etc ...
- Verhulst equation: $N(t) = N_0 \frac{e^{rt}}{C+e^{rt}}$
- activation functions in ANNs

Logistic Regression

Idea

$$log\left(\frac{p}{1-p}\right) = \beta_0 + \sum_{n=1}^{N} \beta_n x_n + \epsilon$$

y:	response
x:	regressors **(assumed to be independent)**
β:	factors
β_0:	intercept
ε:	error **(stochasticity of the data, assumed to be normally dist.)**

Note: one can derive the logit function from max. entropy too!

$$\boxed{p = \frac{e^{\beta_0+\beta_1 x_1+\cdots}}{1+e^{\beta_0+\beta_1 x_1+\cdots}} = \frac{1}{1+e^{-\beta_0-\beta_1 x_1-\cdots}}}$$

toy model: onset of Alzheimer's disease (AD) is a function of ***age*** and years spent in ***education***
(and other risk factors we ignore here for the sake of simplicity)

education: $\mathrm{d} = x_1$ [yrs]
age: $a = x_2$ [yrs]

model: $p_{AD} = \frac{1}{1+e^{-\beta_0-\beta_1 d-\beta_2 a}}$

+ data set + fit → $\beta_0 = +0.1$
$\beta_1 = -1.5$
$\beta_2 = +0.12$

- *positive* value → *increasing p*
- *negative* value → *decreasing p*
- intercept: "background" prevalence, not related to environmental/internal conditions

Logistic Regression

Idea

y:	response
x:	regressors **(assumed to be independent)**
β:	factors
β_0:	intercept
ε:	error **(stochasticity of the data, assumed to be normally dist.)**

model: $p_{AD} = \frac{1}{1+e^{-\beta_0-\beta_1 d-\beta_2 a}}$

education: $d = x_1$ [yrs] $\beta_0 = +0.1$

age: $a = x_2$ [yrs] $\beta_1 = -1.5$

$\beta_2 = +0.12$

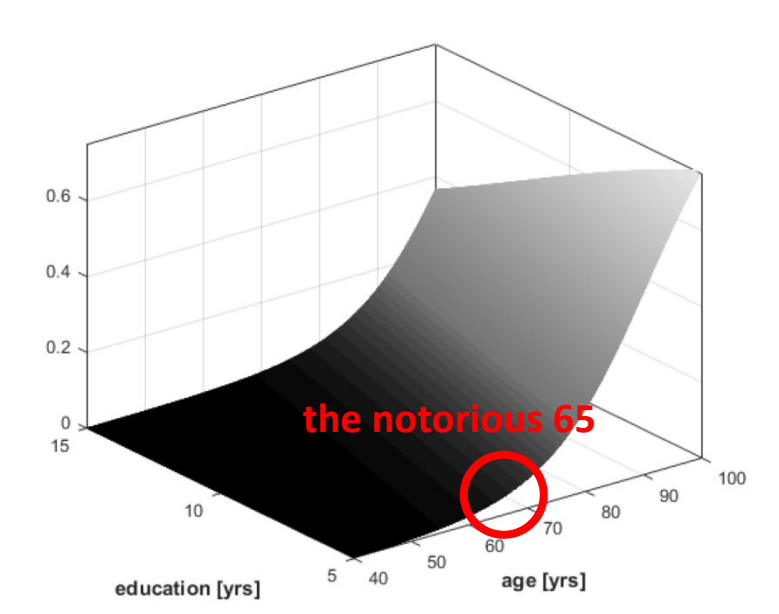

example: ***65yrs*** old person, ***8yrs*** spent in education

→ $p_{AD} = 1.6\%$

65yrs old person, ***13yrs*** spent in education

→ $p_{AD} = 0.001\%$

How does education compensate aging?

$$p_{AD}(d+\bar{d}, a+\bar{a}) = p_{AD}(d,a)$$

→ $\bar{a} = 12.5\,\bar{d}$

hence, one more year prolonged education compensates 12.5 years of aging

(warning: don't confuse correlation with causation here!)

Linear and Logistic Regression

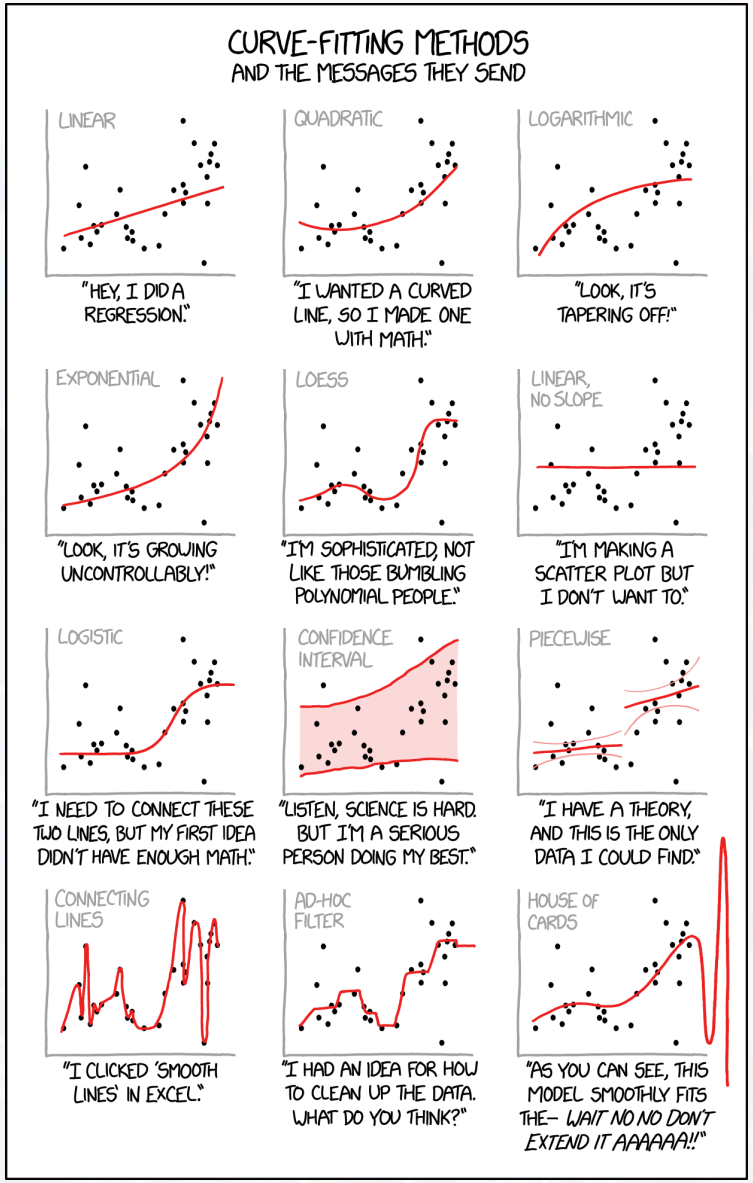

credit: CAUSEweb

Outline

Part II: **Logistic Regression**

Idea

The Factors

a Python example

Logistic Regression

model: $p_{AD} = \frac{1}{1+e^{-\beta_0-\beta_1 d-\beta_2 a}}$

education: $d = x_1$ [yrs]
age: $a = x_2$ [yrs]

$\beta_0 = +0.1$
$\beta_1 = -1.5$
$\beta_2 = +0.12$

y:	response
x:	regressors **(assumed to be independent)**
β:	factors
β_0:	intercept
ε:	error **(stochasticity of the data, assumed to be normally dist.)**

<u>How does the risk of onset changes *per year*?</u>

relative change:

$$\frac{p_{AD}(a+1) - p_{AD}(a)}{p_{AD}(a)} \approx e^{\beta_2} - 1 \approx 12.7\%$$

for $p_{AD} \ll 1$ (hence, for small Δa and "young" ages, i. e. below ≈ 80yrs)

"The risk of getting AD increases by 12.7% every year." **(warning: does not mean that it increases by 127% in ten yrs – we made an approximation!)**

more precise:
relative change of
the **odds ratio**

$$\frac{\frac{\partial}{\partial x_i}\left(\frac{p_{AD}}{1-p_{AD}}\right)}{\frac{p_{AD}}{1-p_{AD}}} = \beta_i$$

x_i is the desired regressor, for example, age again (x_2)

The factors β_i indicate how strong (and in which direction) *p* changes wrt a regressor x_i

Linear and Logistic Regression

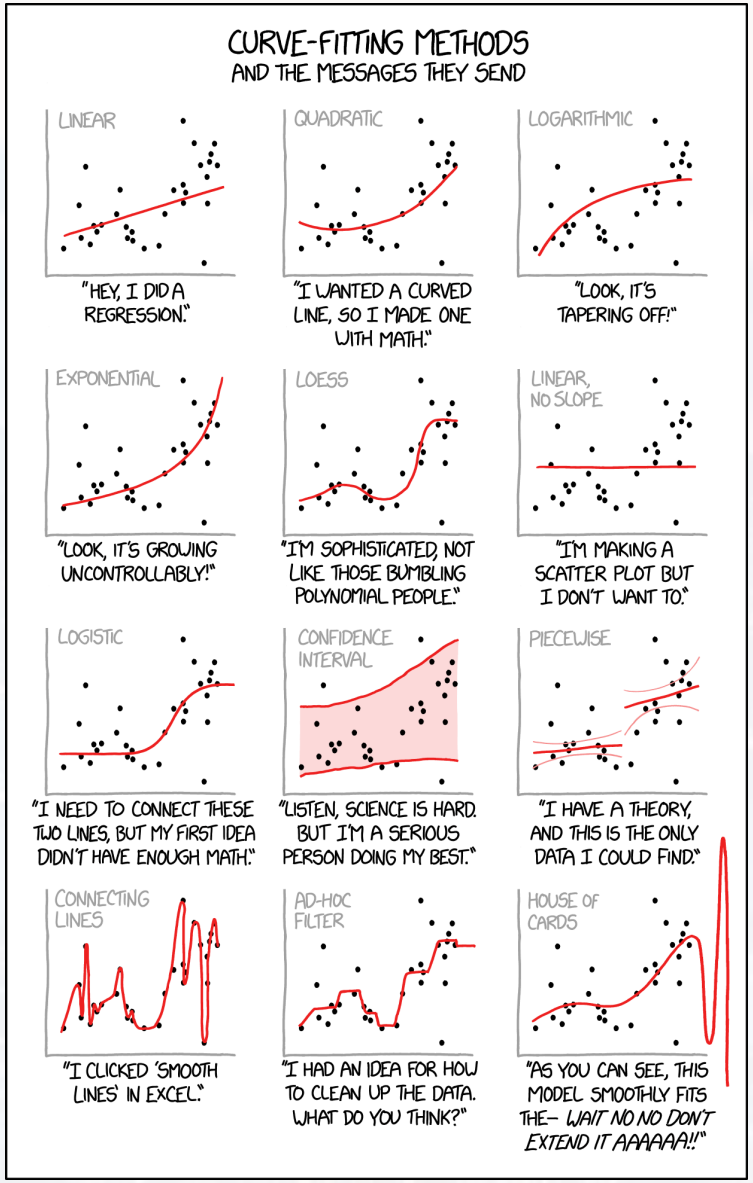

credit: CAUSEweb

Outline

Part II: **Logistic Regression**

Idea

The Factors

a Python example

Logistic Regression

1) loading data
2) plotting data
3) scaling data
4) fitting model
5) evaluating model

let us return to the molecule data set:

```
Train  = pd.read_csv("molecular_train_gbc_cat.csv")
Test   = pd.read_csv("molecular_test_gbc_cat.csv")
```

	molecular_weight	electronegativity	bond_lengths	num_hydrogen_bonds	logP	label
0	341.704142	2.655846	3.094069	2	9.111473	Toxic
1	335.950798	3.222621	2.890393	7	8.928483	Toxic
2	235.203185	2.441153	2.482026	1	6.497307	Non-Toxic
3	246.504930	2.766560	2.715472	7	7.450888	Non-Toxic
4	437.938926	3.480105	3.595688	3	10.915629	Non-Toxic

```
import statsmodels.api as sm
```

Logistic Regression

1) loading data
2) plotting data
3) scaling data
4) fitting model
5) evaluating model

it is the same data set → plotting and scaling is as before

```
X = sm.add_constant(TrainS)
```

adding the intercept

```
Y = pd.get_dummies(YTrain)
```

Python needs True/False as categorical

$$p = \frac{e^{\beta_0+\beta_1 x_1+\cdots}}{1+e^{\beta_0+\beta_1 x_1+\cdots}}$$

```
In [48]: print(Y)
     Non-Toxic  Toxic
0        False   True
1        False   True
2         True  False
3         True  False
4         True  False
```

we have two states: toxic / non-toxic

```
my_model = sm.GLM(Y, X, family = sm.families.Binomial()).fit()

my_model.summary()
```

GLM: general linear model

Logistic Regression

1) loading data
2) plotting data
3) scaling data
4) fitting model
5) evaluating model

it is the same data set → plotting and scaling is as before

```
X = sm.add_constant(TrainS)
Y = pd.get_dummies(YTrain)

my_model = sm.GLM(Y, X, family = sm.families.Binomial()).fit()
my_model.summary()
```

$$p = \frac{e^{\beta_0+\beta_1 x_1+\cdots}}{1+e^{\beta_0+\beta_1 x_1+\cdots}}$$

```
                 Generalized Linear Model Regression Results
==============================================================================
Dep. Variable:     ['Non-Toxic', 'Toxic']   No. Observations:                800
Model:                                GLM   Df Residuals:                    794
Model Family:                    Binomial   Df Model:                          5
Link Function:                      Logit   Scale:                        1.0000
Method:                              IRLS   Log-Likelihood:               -332.82
Date:                   Sat, 14 Sep 2024   Deviance:                      665.64
Time:                           20:59:18   Pearson chi2:                 1.14e+03
No. Iterations:                        6   Pseudo R-squ. (CS):            0.4243
Covariance Type:                nonrobust
=====================================================================================
                        coef    std err          z      P>|z|      [0.025      0.975]
-------------------------------------------------------------------------------------
const                  6.1641      0.585     10.536      0.000       5.017       7.311
molecular_weight     -10.4920      3.626     -2.893      0.004     -17.599      -3.385
electronegativity      3.2874      0.599      5.492      0.000       2.114       4.461
bond_lengths           0.6736      1.913      0.352      0.725      -3.075       4.422
num_hydrogen_bonds    -0.3082      0.303     -1.018      0.309      -0.902       0.285
```

p-value for constant model

p-values for factors

2σ conf range of factors

Logistic Regression

accuracy: How ***often*** did the model make the correct prediction.
cross-entropy: How ***certain*** was the model when making the prediction.

1) loading data
2) plotting data
3) scaling data
4) fitting model
5) evaluating model

Logistic Regression

accuracy: How ***often*** did the model make the correct prediction.

cross-entropy: How ***certain*** was the model when making the prediction.

1) loading data
2) plotting data
3) scaling data
4) fitting model
5) evaluating model

```
predProbs   = my_model.predict(sm.add_constant(TestS))

Pred        = np.round(predProbs).astype(int)
predictions = ['Non-Toxic' if i==1 else 'Toxic' for i in Pred]
```

```
Dep. Variable:      ['Non-Toxic', 'Toxic']
```

```
In [51]: predictions
Out[51]:
['Toxic',
 'Toxic',
 'Non-Toxic',
 'Non-Toxic',
 'Toxic',
 'Toxic',
 'Toxic',
```

```
YTest       = Test['label']
accuracy    = 100*(YTest == predictions).sum()/len(predictions)
print(f'accuracy = {accuracy: .2f}%')
```

```
accuracy =  84.00%
```

Logistic Regression

accuracy: How ***often*** did the model make the correct prediction.

cross-entropy: How ***certain*** was the model when making the prediction.

1) loading data
2) plotting data
3) scaling data
4) fitting model
5) evaluating model

accuracy is ≈ 84% But does it depend on the class? → **confusion matrix**

ideal world:

true label \ predicted label	non-toxic	toxic
non-toxic	100%	0%
toxic	0%	100%

```
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
```

Logistic Regression

accuracy: How ***often*** did the model make the correct prediction.

cross-entropy: How ***certain*** was the model when making the prediction.

1) loading data
2) plotting data
3) scaling data
4) fitting model
5) evaluating model

accuracy is ≈ 84% But does it depend on the class? → **confusion matrix**

two labels

```
L = ['Non-Toxic', 'Toxic']

cm   = confusion_matrix(TestY, predictions, labels = L, normalize = 'true')
disp = ConfusionMatrixDisplay(confusion_matrix = cm, display_labels = L)
disp.plot(cmap = 'gray')
plt.show()
```

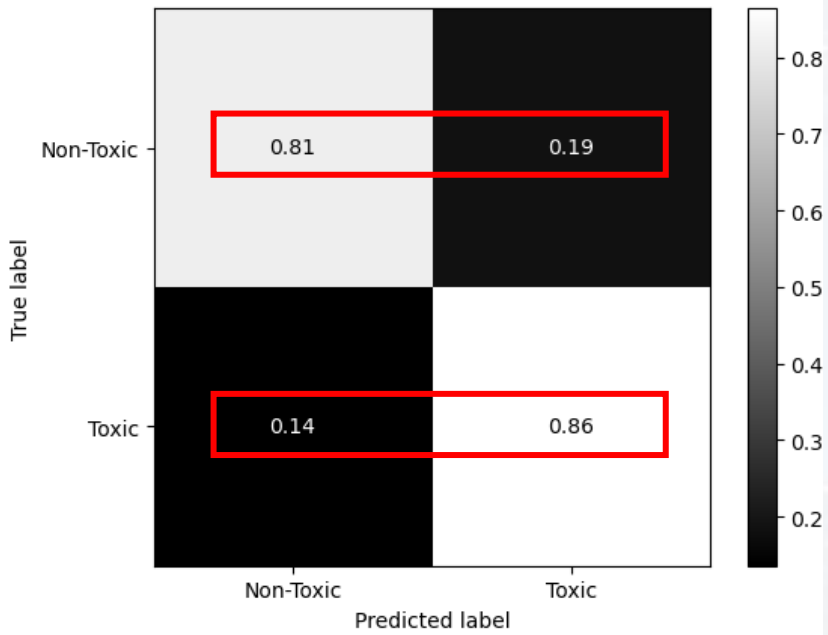

Logistic Regression

accuracy: How ***often*** did the model make the correct prediction.

cross-entropy: How ***certain*** was the model when making the prediction.

1) loading data
2) plotting data
3) scaling data
4) fitting model
5) evaluating model

accuracy is ≈ 84% But does it depend on the class? → **confusion matrix**

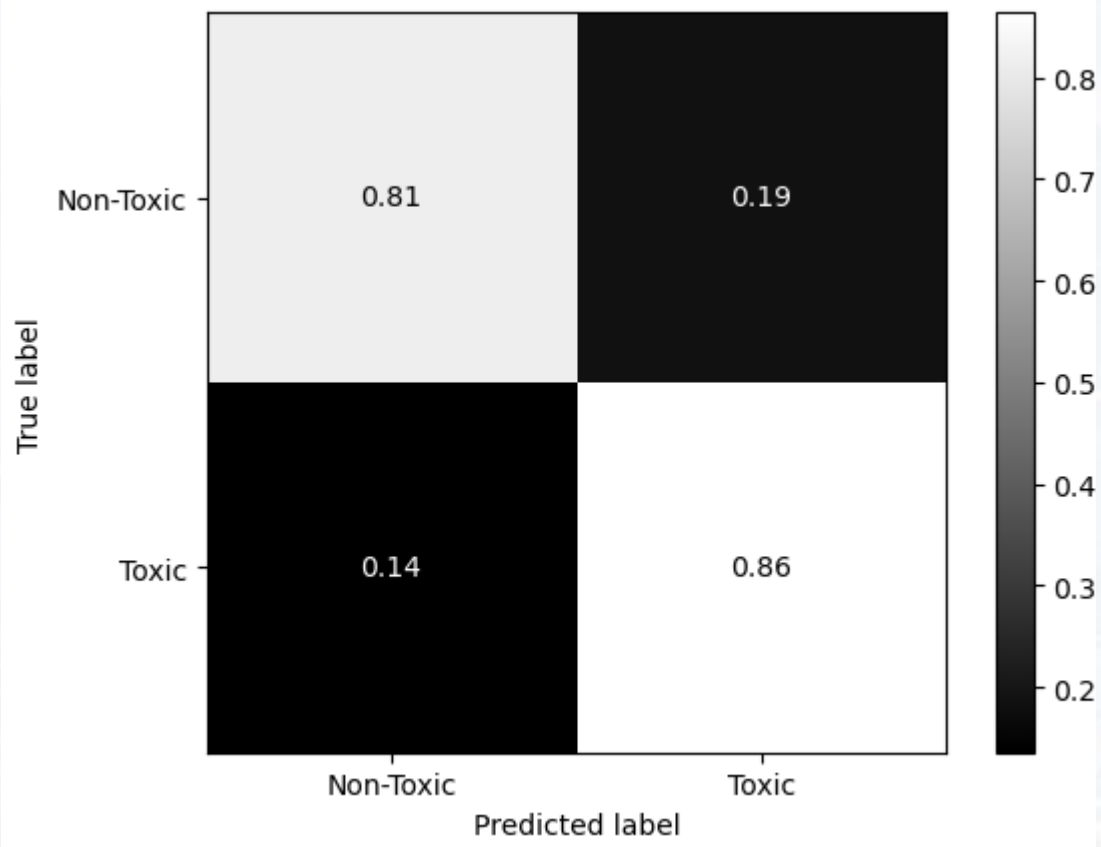

true label \ predicted label	non-toxic	toxic
non-toxic	100%	0%
toxic	0%	100%

Logistic Regression

accuracy: How ***often*** did the model make the correct prediction.

cross-entropy: How ***certain*** was the model when making the prediction.

1) loading data
2) plotting data
3) scaling data
4) fitting model
5) evaluating model

ideal world:

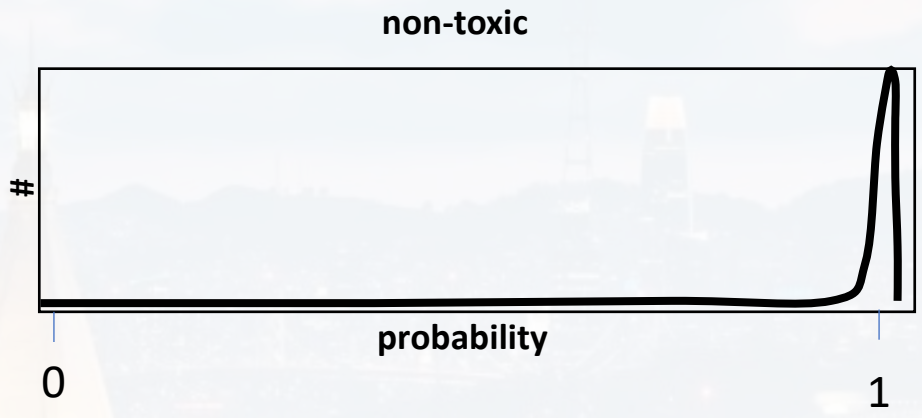

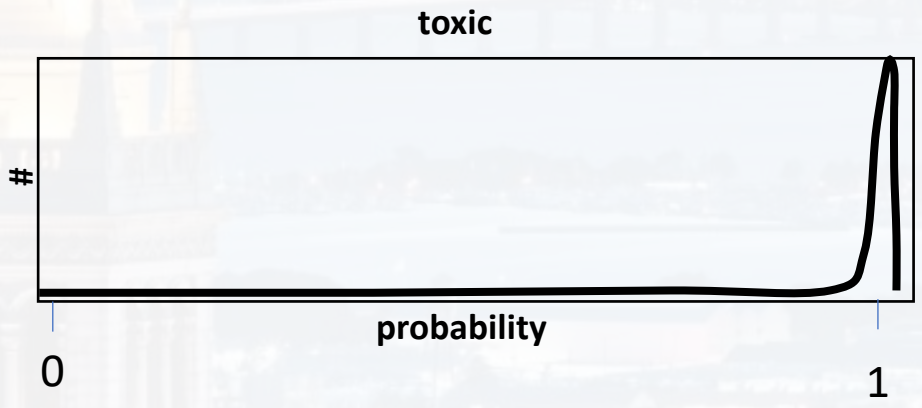

Logistic Regression

accuracy: How ***often*** did the model make the correct prediction.

cross-entropy: How ***certain*** was the model when making the prediction.

1) loading data
2) plotting data
3) scaling data
4) fitting model
5) evaluating model

```
PredProbs = np.vstack((predProbs, 1 - predProbs))

fig, ax = plt.subplots(len(L), 1, sharex = True)
fig.set_figheight(6)
fig.subplots_adjust(hspace = 0.5)
fig.suptitle('entropy')
for i, l in enumerate(L):
    idx = [k for k, y in enumerate(TestY) if y == l]
    idx = np.array(idx)
    (value, where) = np.histogram(PredProbs[i,idx],\
                                  bins = np.arange(0, 1, 0.01),\
                                  density = True)
    w = 0.5*(where[1:] + where[:-1])
    ax[i].plot(w, value, 'k-')
    ax[i].set_ylabel('frequency')
    ax[i].set_title(l)
ax[len(L)-1].set_xlabel('probability')
plt.show()
```

Logistic Regression

accuracy: How ***often*** did the model make the correct prediction.

cross-entropy: How ***certain*** was the model when making the prediction.

1) loading data
2) plotting data
3) scaling data
4) fitting model
5) evaluating model

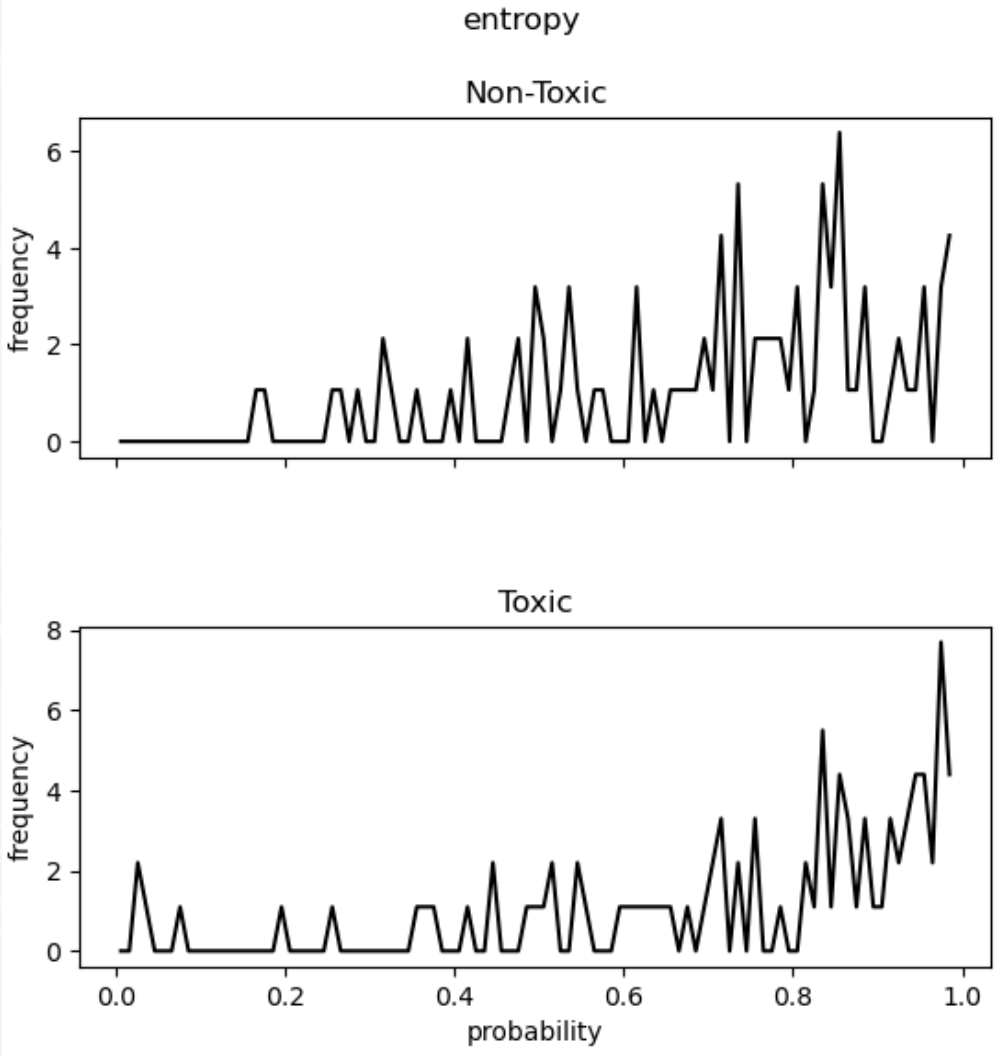

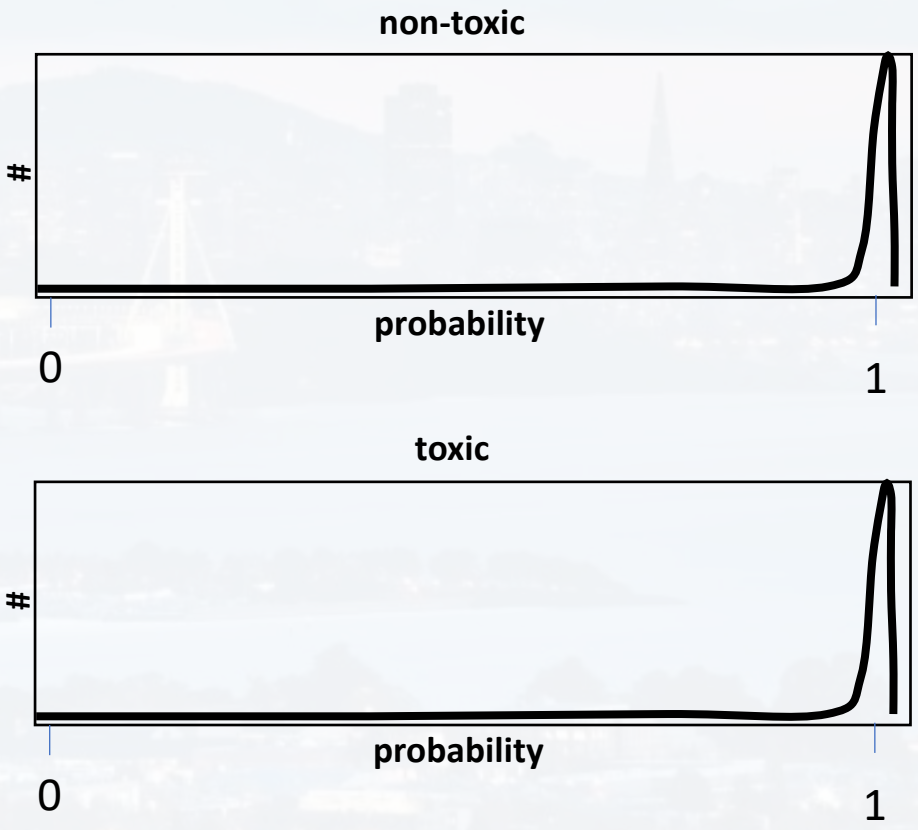

Linear and Non-Linear Regression

Thank you very much for your attention!