

Lecture 05:

Unsupervised Learning



Markus Hohle

University California, Berkeley

Machine Learning Algorithms

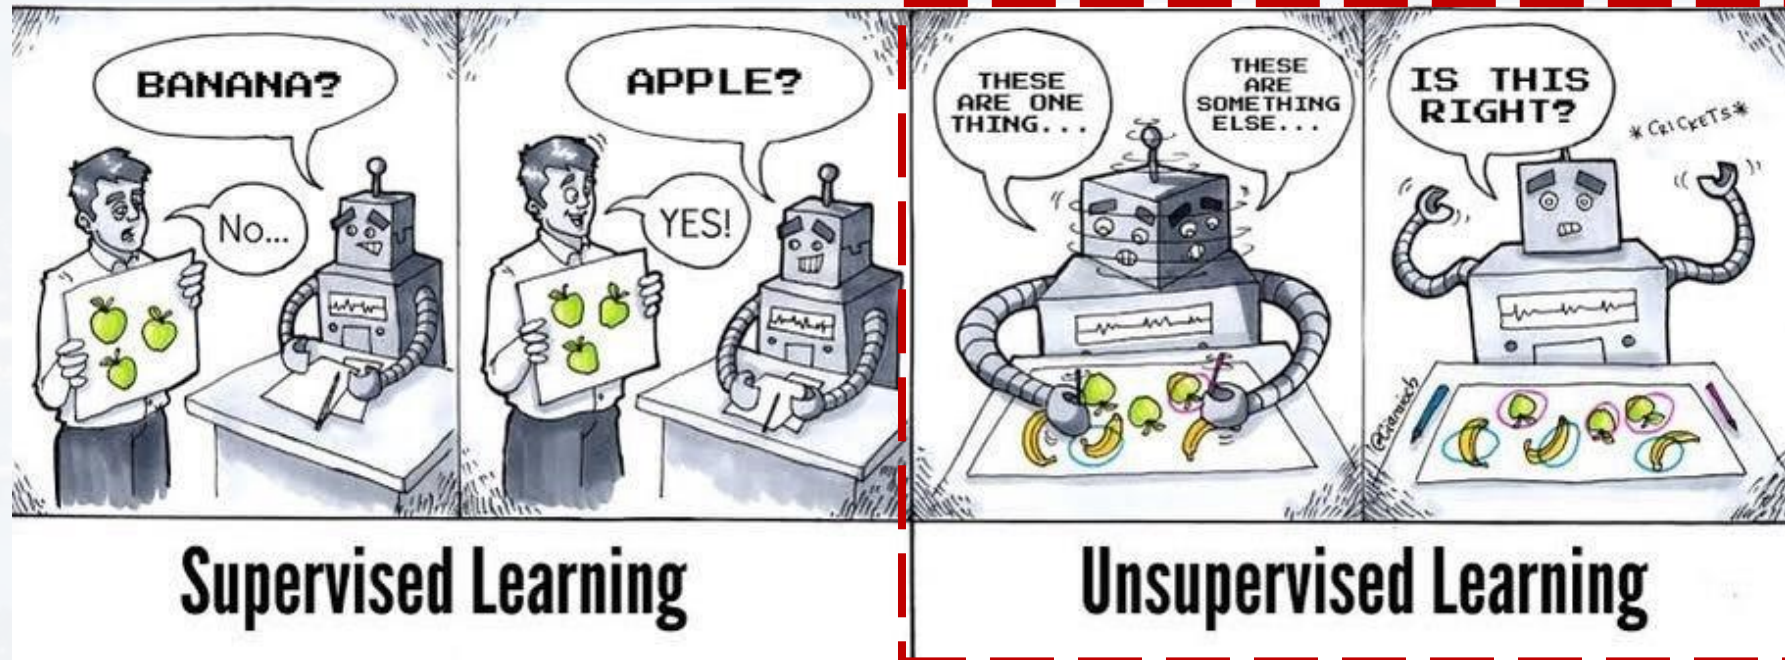
MSSE 277B, 3 Units

Fall 2024



So far, there has been a **training** data set and a **test** data set...

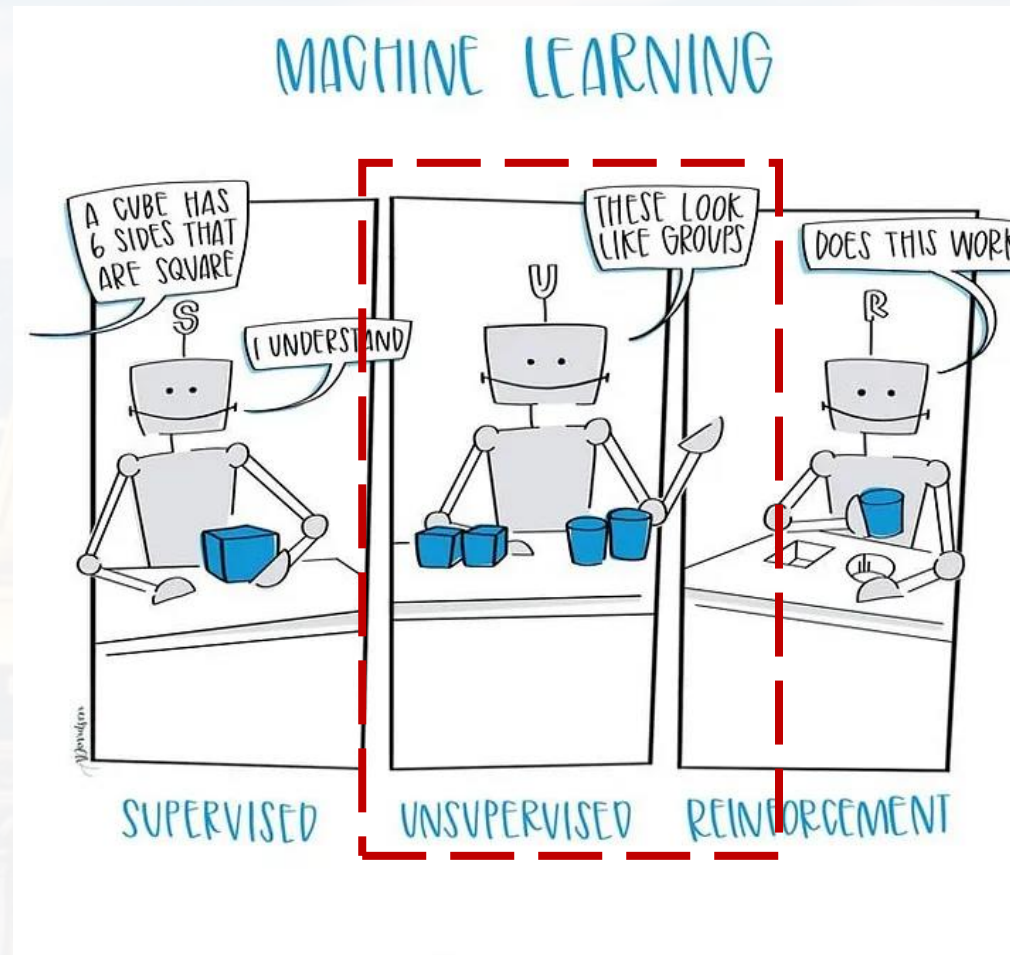
... but maybe there are ways to learn *without* training data

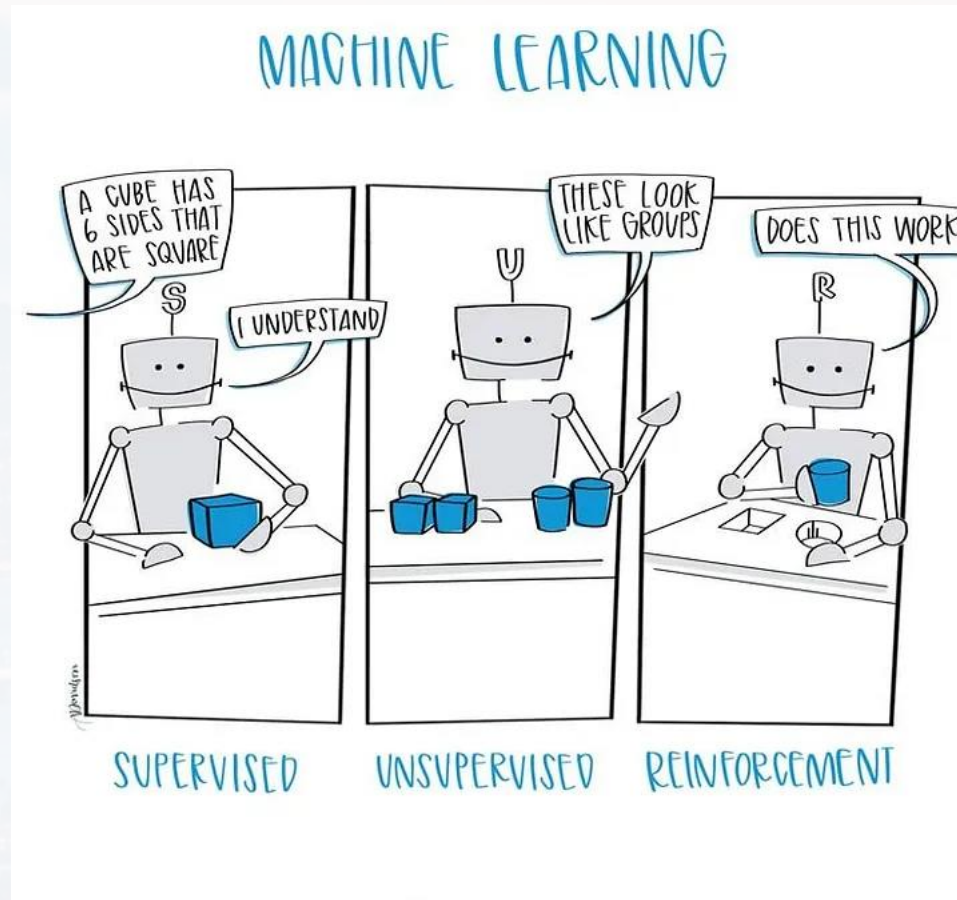




So far, there has been a **training** data set and a **test** data set...

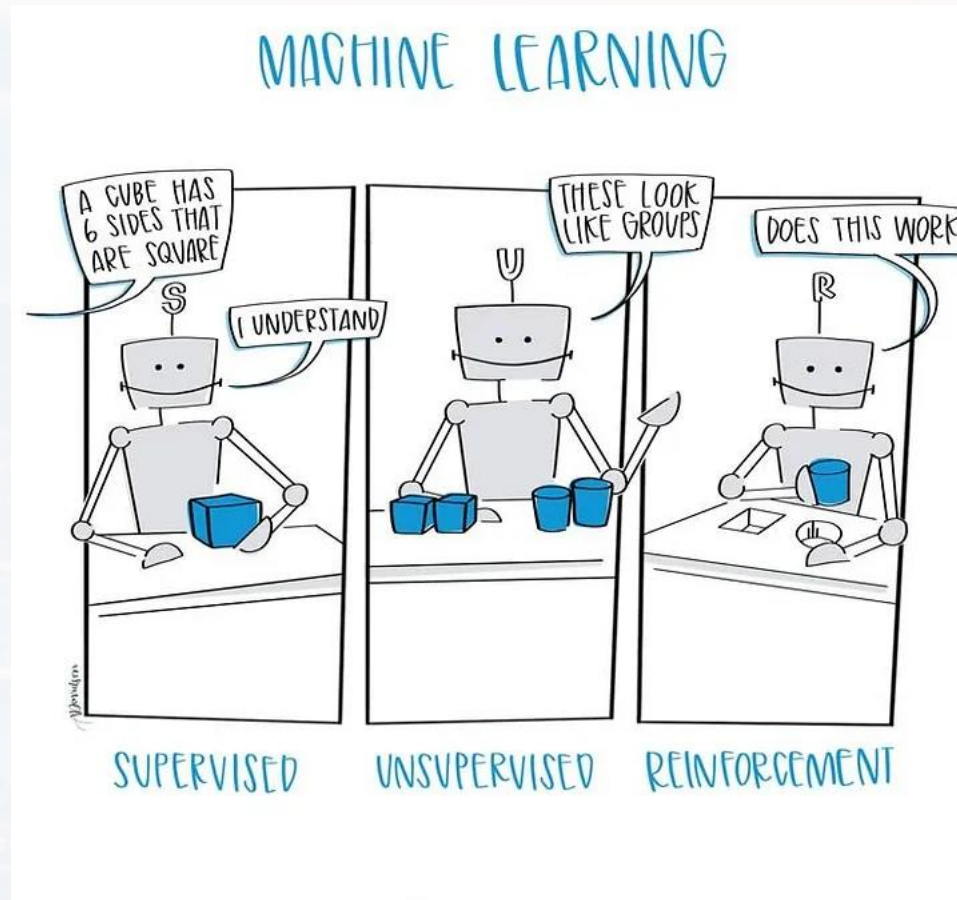
... but maybe there are ways to learn *without* training data





Outline

- K - means
- GMM
- trees



Outline

- K - means

- GMM

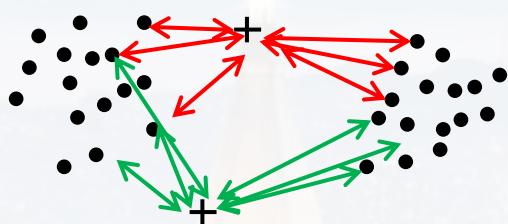
- trees



idea:



a) assign k means randomly



b) calculate *distance* from each point to each mean



c) assign each point to its closest mean

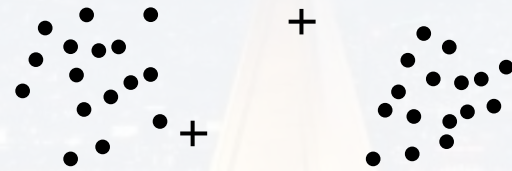


d) update the means accordingly

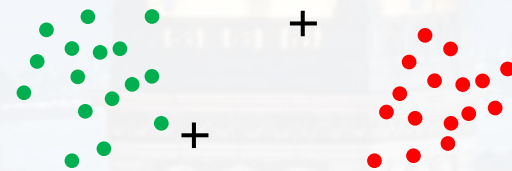
idea:



d) update the means accordingly



e) go back to b)





problem: $k = \text{number of cluster}$, is a hyperparameter. How do I know the correct value for k ?

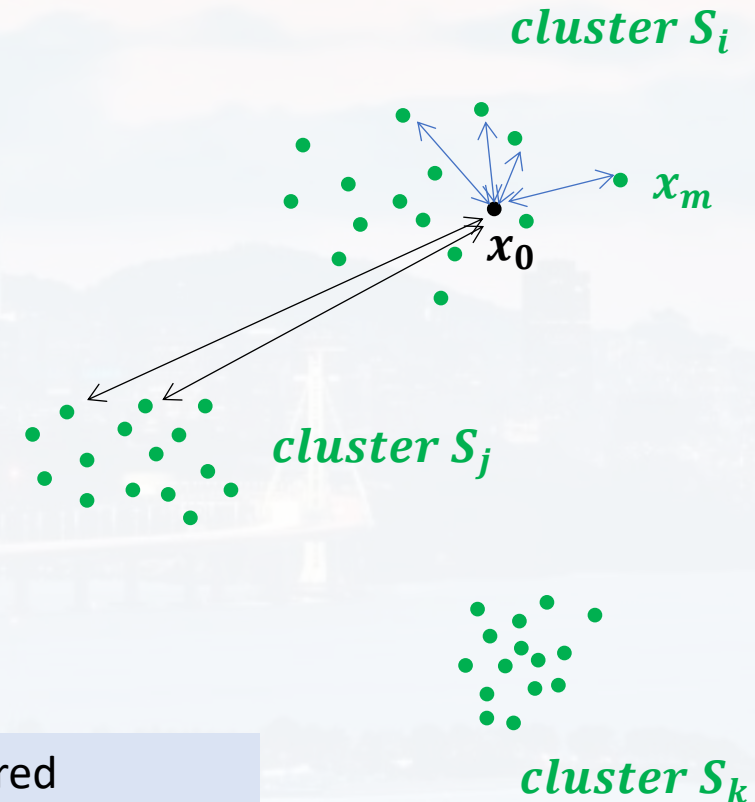
→ silhouette Ψ

- distance d_1 of a data point x_0 to *its assigned cluster* S_i
vs distance d_2 to *closest cluster (here S_j)*

$$\Psi(x_0) = \begin{cases} 0 & \text{if } d_1 = 0 \\ \frac{d_2 - d_1}{\max[d_1; d_2]} \end{cases}$$

- average over all points → ψ_{tot}

if	$\psi_{tot} = 0.75 \dots 1.00$	→ well clustered
	$\psi_{tot} = 0.50 \dots 0.75$	→ medium clustered
	$\psi_{tot} = 0.25 \dots 0.50$	→ poorly clustered
	$\psi_{tot} < 0.25$	→ data has no structure





problem: k is a hyperparameter. How do I know the correct value for k ?

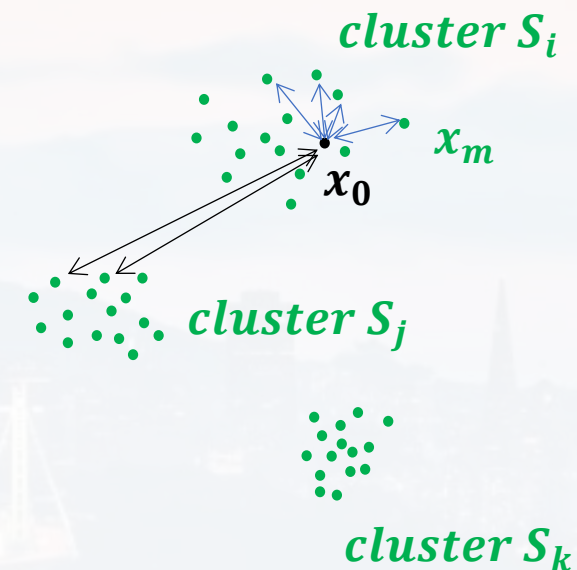
→ silhouette Ψ

- distance d_1 of a data point x_0 to **its assigned cluster** S_i
vs distance d_2 to **closest cluster (here S_j)**

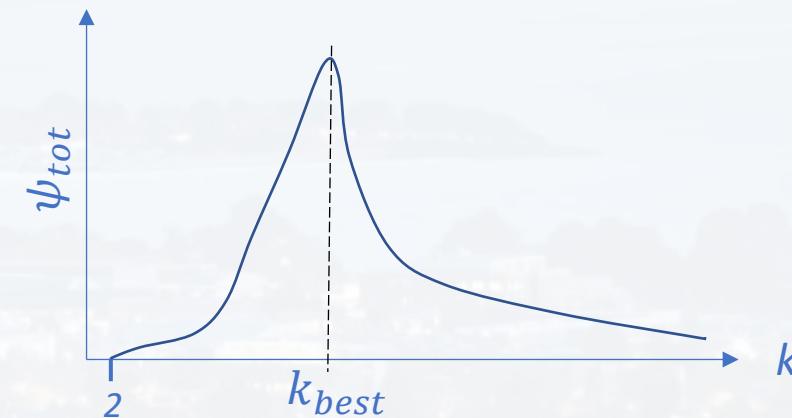
$$\Psi(x_0) = \begin{cases} 0 & \text{if } d_1 = 0 \\ \frac{d_2 - d_1}{\max[d_1; d_2]} \end{cases}$$

- average over all points → ψ_{tot}

$\psi_{tot} = 0.75 \dots 1.00$	→ well clustered
$\psi_{tot} = 0.50 \dots 0.75$	→ medium clustered
$\psi_{tot} = 0.25 \dots 0.50$	→ poorly clustered
$\psi_{tot} < 0.25$	→ data has no structure



ideal world →





```
import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns
```

our standard libraries

for having different
distances available

```
from pyclustering.utils.metric import *
```

```
from nltk.cluster.kmeans import KMeansClusterer
```

```
from sklearn.metrics import silhouette_samples, silhouette_score
```

```
from sklearn import datasets
```

calling the famous "iris" data set

calculating silhouette
coefficient for different k

performing k-means



```
import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns

from pyclustering.utils.metric import *
from nltk.cluster.kmeans import KMeans(
from sklearn.metrics import silhouette_
from sklearn import datasets
```

```
iris = datasets.load_iris()
```

```
iris.DESCR
```

Iris plants dataset

****Data Set Characteristics:****

:Number of Instances: 150 (50 in each of three classes)
:Number of Attributes: 4 numeric, predictive attributes and the class
:Attribute Information:

- sepal length in cm
- sepal width in cm
- petal length in cm
- petal width in cm
- class:

- Iris-Setosa
- Iris-Versicolour
- Iris-Virginica

ideal world: three distinct cluster

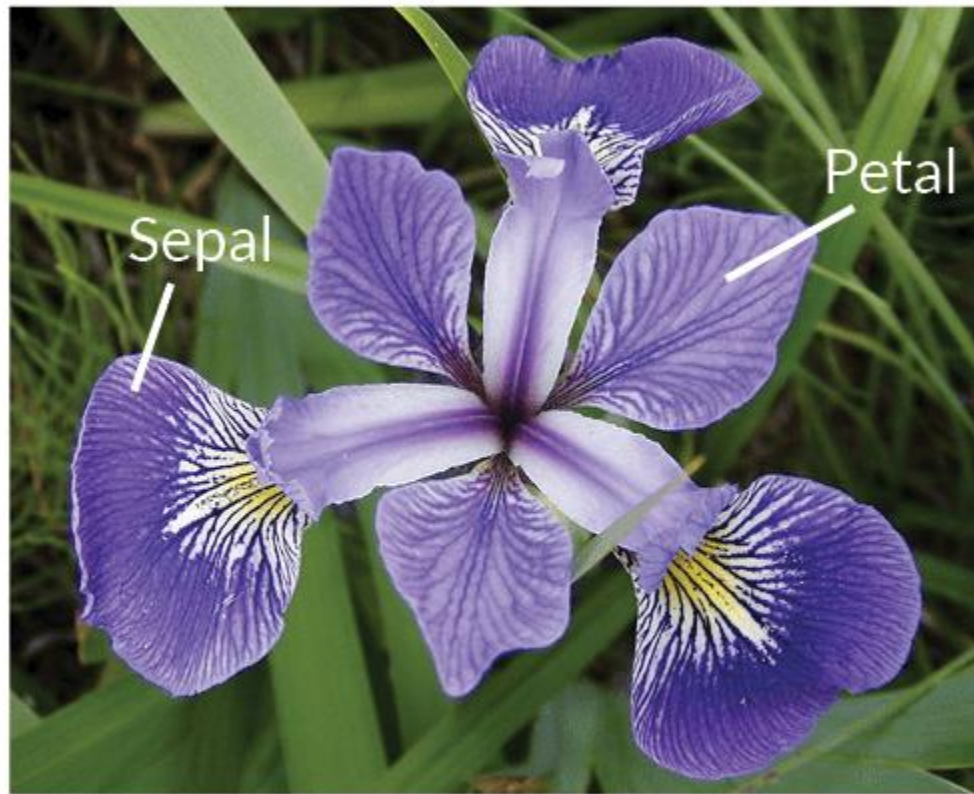
:Summary Statistics:

	Min	Max	Mean	SD	Class Correlation
sepal length:	4.3	7.9	5.84	0.83	0.7826
sepal width:	2.0	4.4	3.05	0.43	-0.4194
petal length:	1.0	6.9	3.76	1.76	0.9490 (high!)
petal width:	0.1	2.5	1.20	0.76	0.9565 (high!)



```
iris = datasets.load_iris()
```

```
iris.DESCR
```



Iris Versicolor



Iris Setosa



Iris Virginica



loading & exploring the data:

```
iris = datasets.load_iris()
```

```
iris.DESCR
```

```
iris.feature_names
```

```
iris.target_names
```

```
['sepal length (cm)',  
'sepal width (cm)',  
'petal length (cm)',  
'petal width (cm)']
```

four features → 4D

```
array(['setosa', 'versicolor', 'virginica'])
```

check out the Jupyter Notebook [Walk_Through_Kmeans](#)

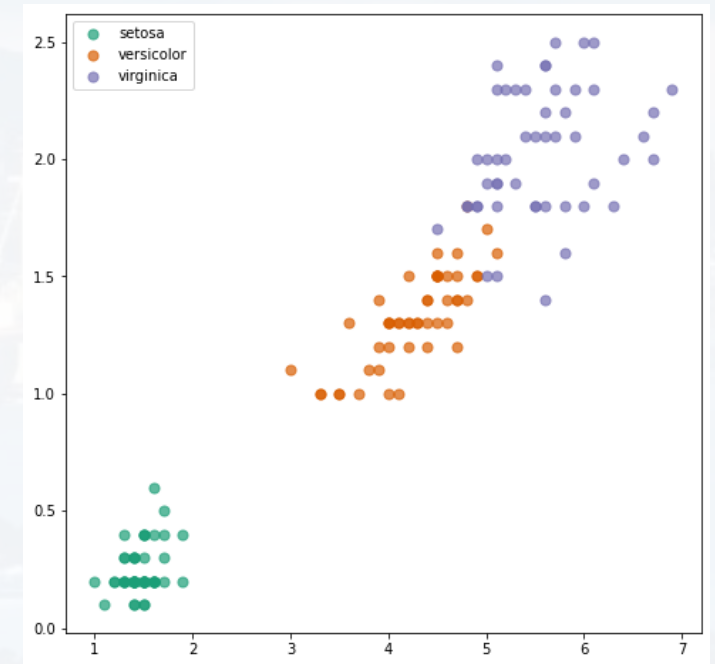
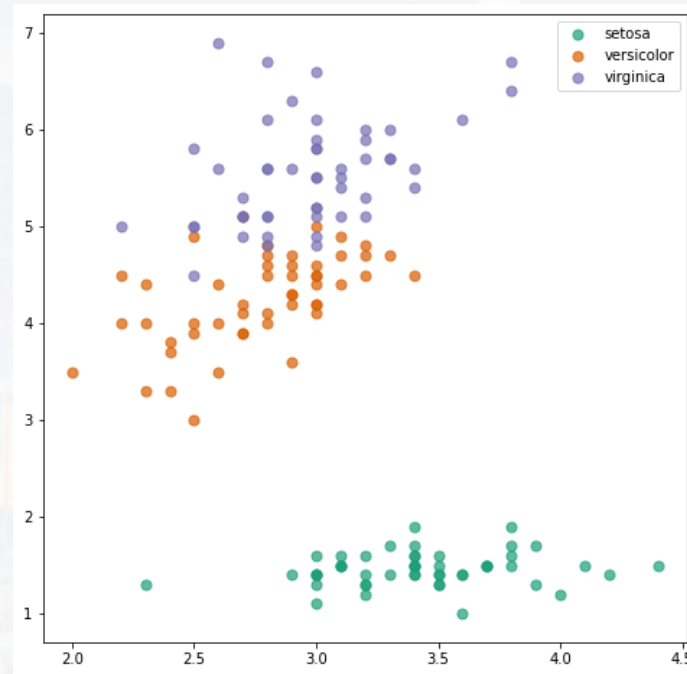
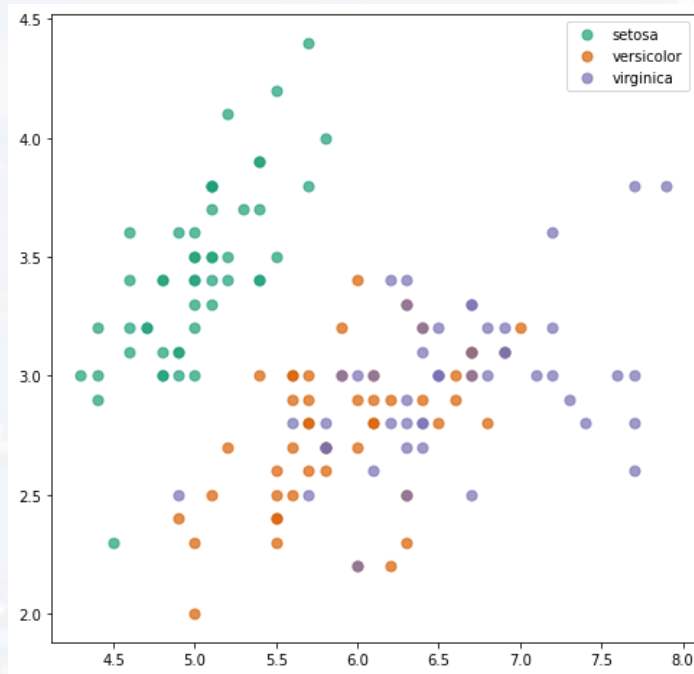
- plotting the data
- running k-means
- evaluating the result



```
['sepal length (cm)',  
'sepal width (cm)',  
'petal length (cm)',  
'petal width (cm)']
```

4D dataset → plotting two components

- **plotting the data**
- running k-means
- evaluating the result





```
nClust    = 3  
rep       = 25  
dist      = distance_metric(type_metric.EUCLIDEAN)
```

```
my_model  = KMeansClusterer(nClust, distance = dist,\n                             repeats = rep,\n                             avoid_empty_clusters = True)
```

```
PredLabels = my_model.cluster(X2D,\n                              assign_clusters = True)
```

```
Center    = my_model.means()
```

- plotting the data
- **running k-means**
- evaluating the result

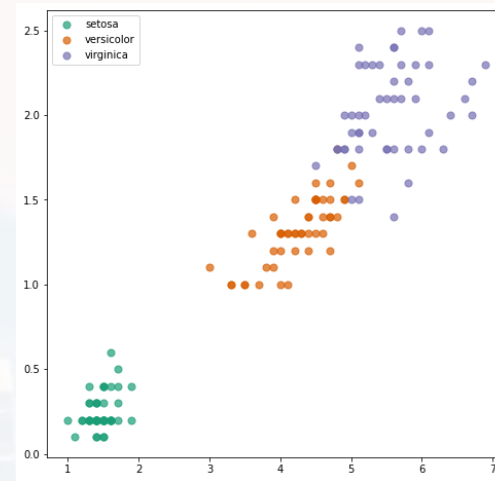
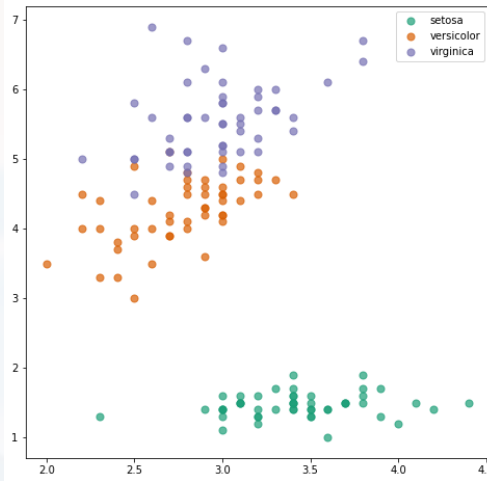
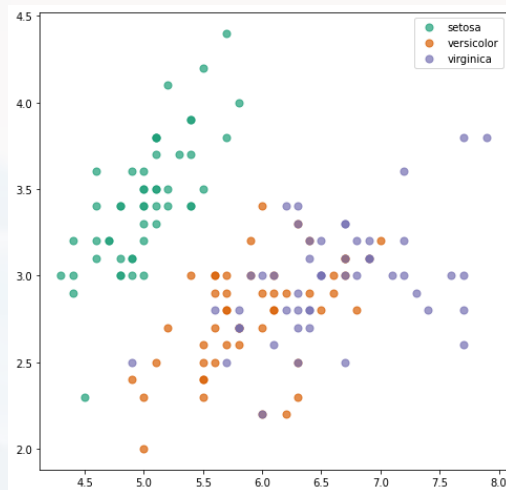
we need to “guess” the number of cluster

the **initial** means are assigned randomly.
→ repeat the procedure 25 times
→ avoiding local minimum,

The features are measured in cm, i. e. the correct distance to pick here is Euclidean

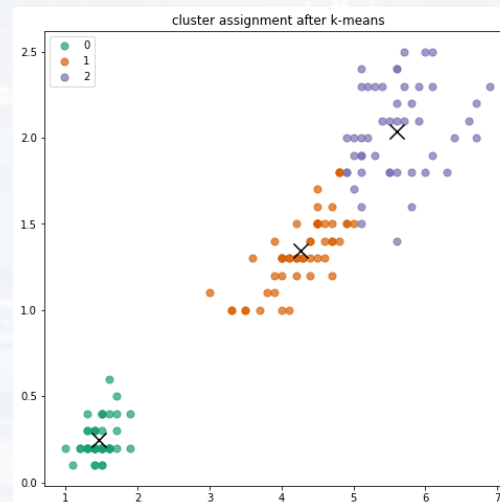
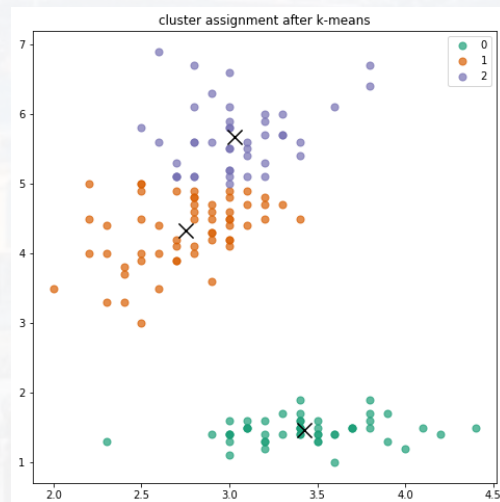
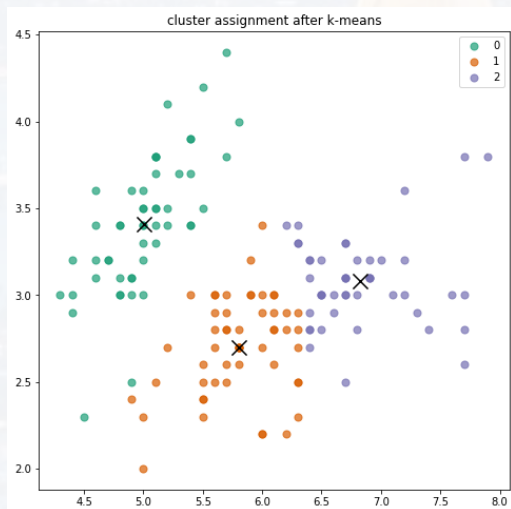


true classes



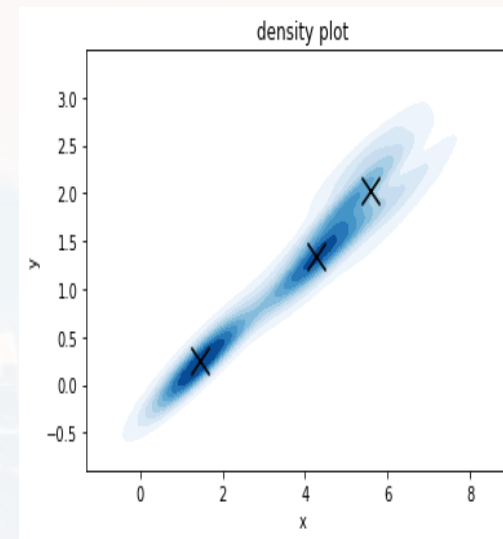
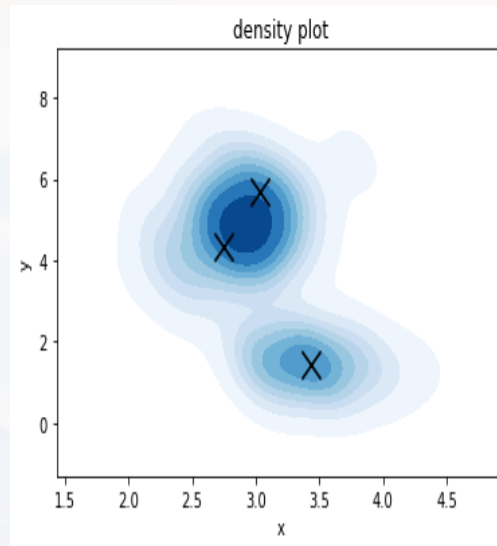
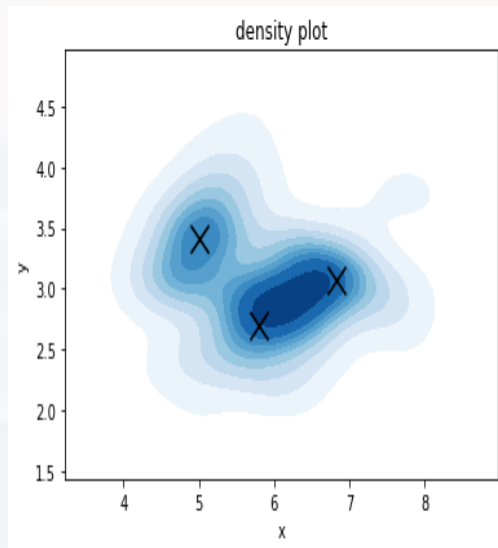
- plotting the data
- **running k-means**
- evaluating the result

assigned classes



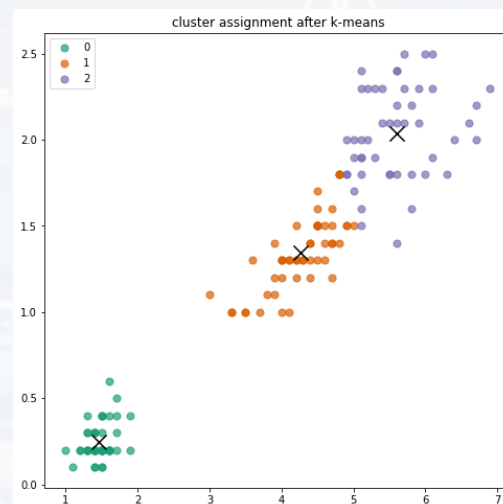
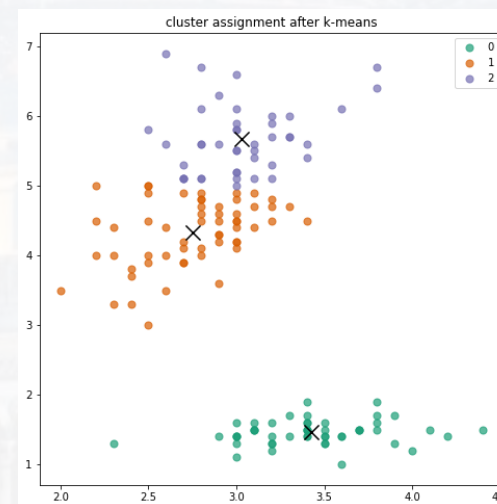
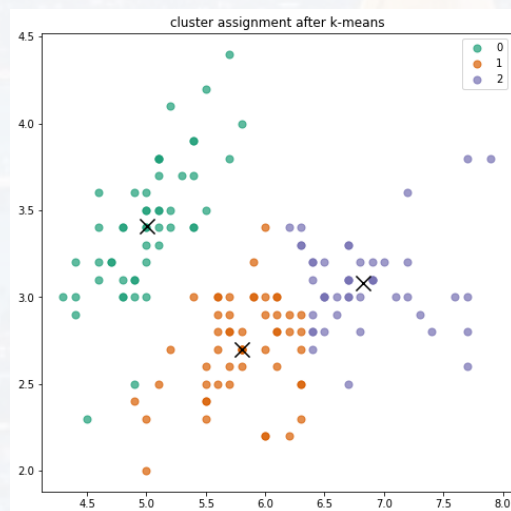


density



- plotting the data
- **running k-means**
- evaluating the result

assigned classes

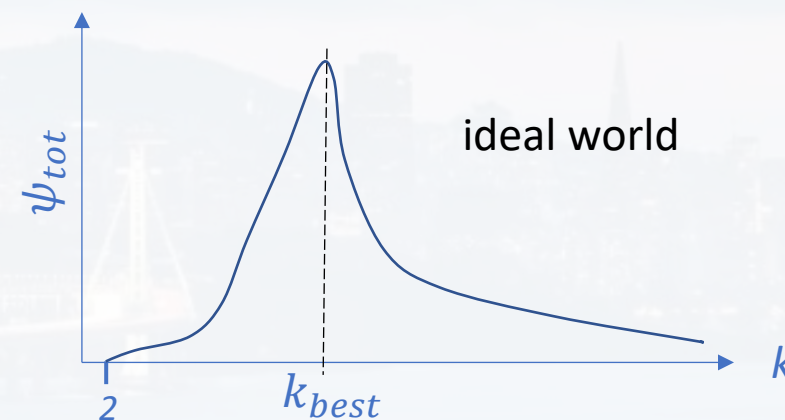
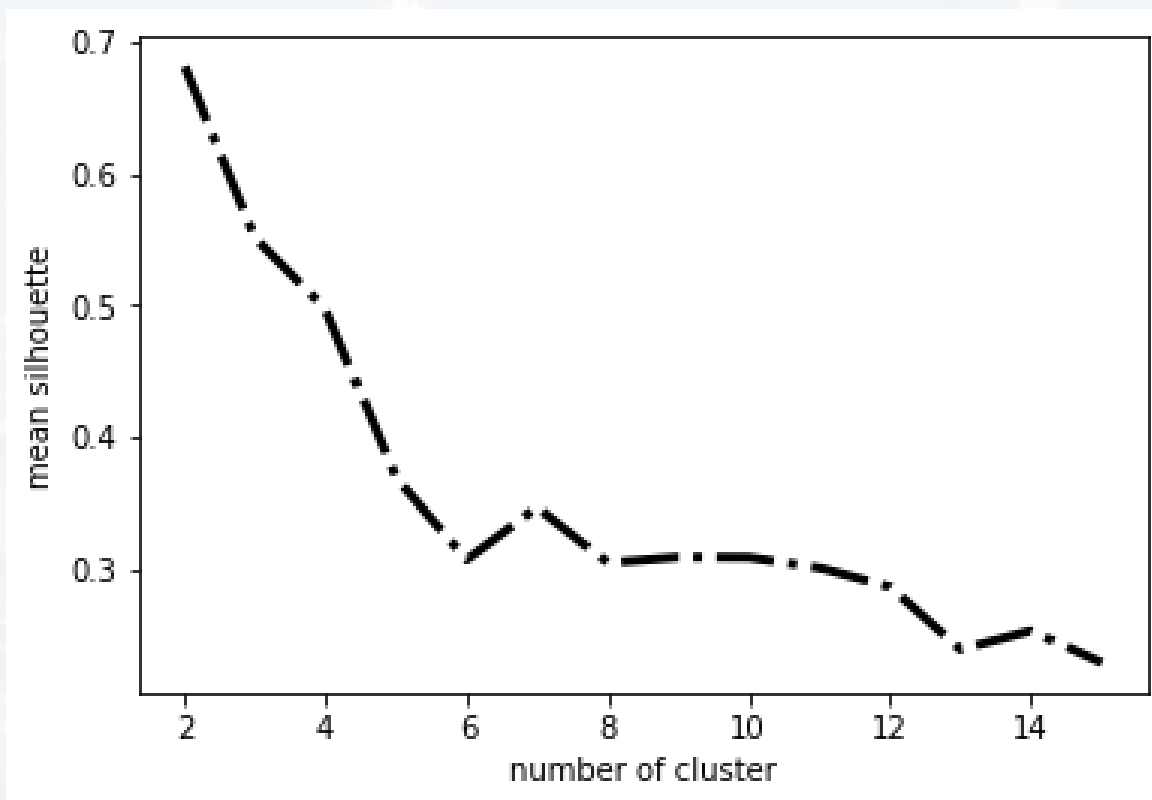




we run k-means now for the **full 4D** dataset
+ evaluate clustering with silhouette

- plotting the data
- running k-means
- **evaluating the result**

`silhouette_score(X, Labels)`

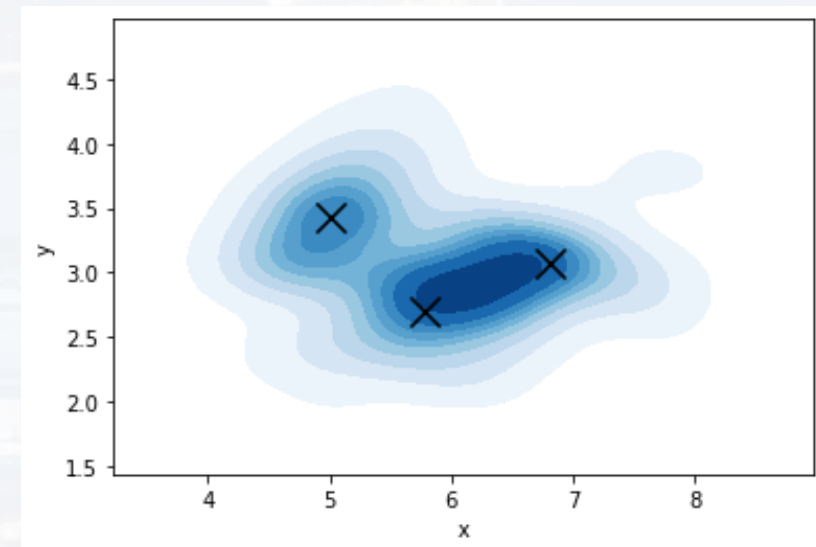
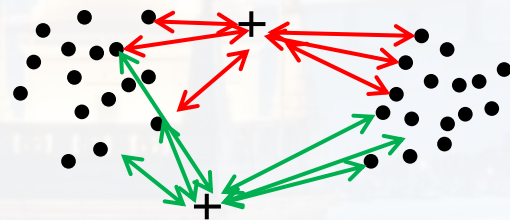


accuracy for 4D $k = 3$: 90%



summary:

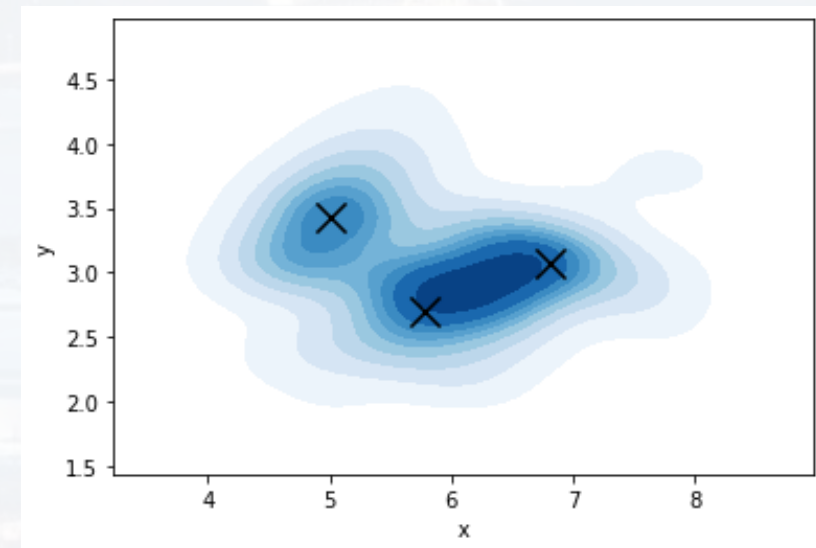
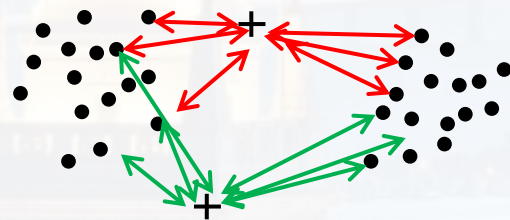
- simple and fast
- unsupervised
- k has to be given \rightarrow silhouette for determining best k
- problems if cluster have unusual shapes
(elongated, hollow inside, scattered)

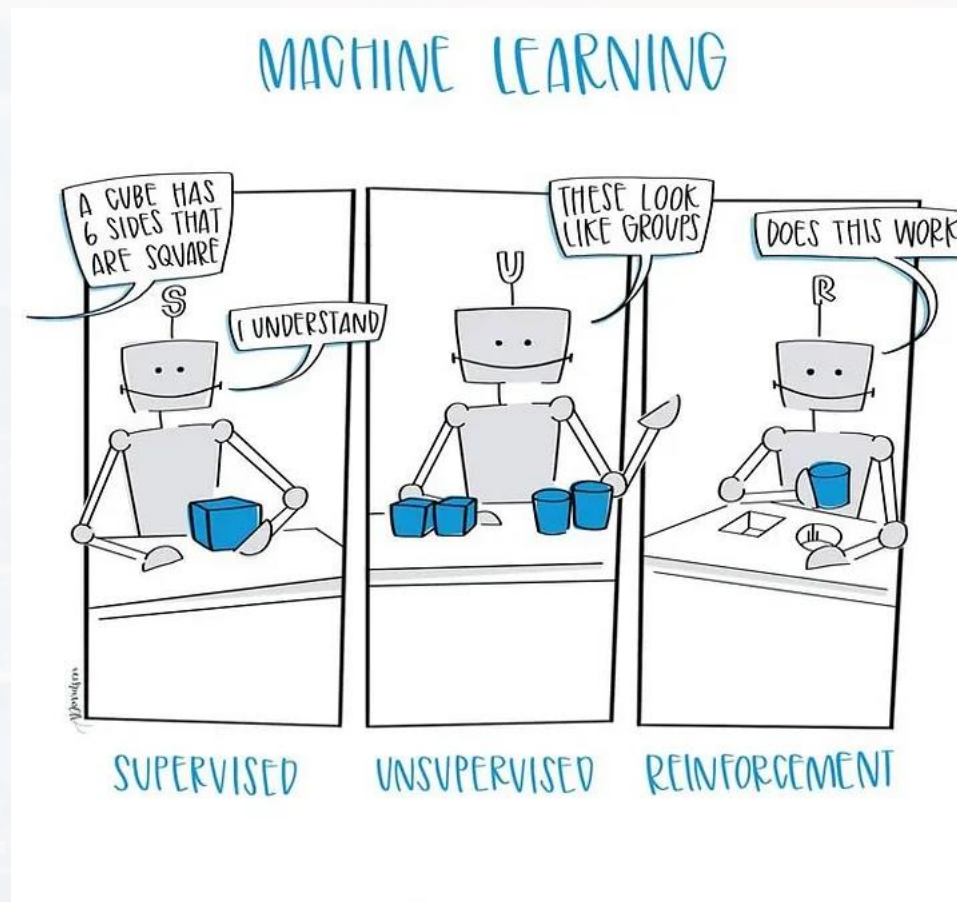




topics for the discussion/office hour:

- What is a *distance*?
- Which are different distances?
- When to use which distance?





Outline

- K - means

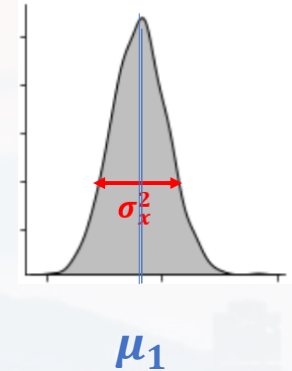
- **GMM**

- trees

Gaussian Mixture Models

one feature

$$N_1(x_1) = \frac{1}{\sqrt{2\pi \sigma_{x1}^2}} \exp -\frac{1}{2} \left(\frac{x_1 - \mu_1}{\sigma_{x1}} \right)^2$$

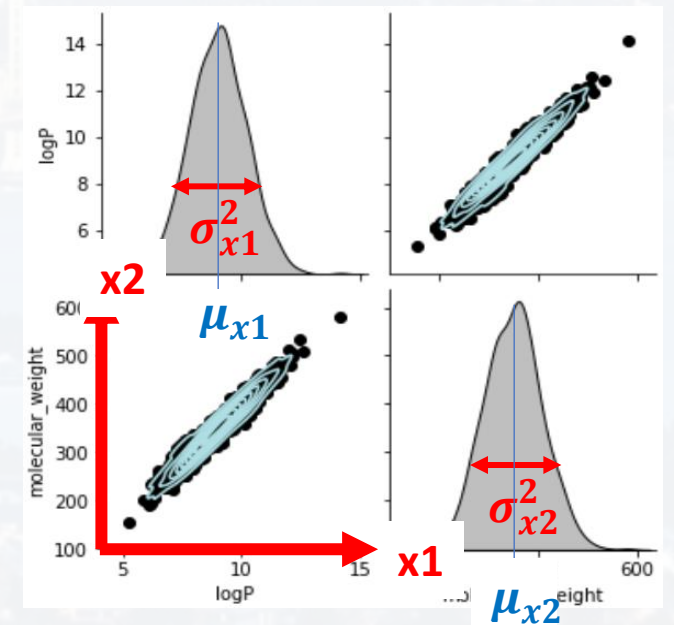


two features

$$\Sigma = \begin{pmatrix} \sigma_{x1}^2 & cov(x_1, x_2) \\ cov(x_2, x_1) & \sigma_{x2}^2 \end{pmatrix} \quad \text{covariance matrix}$$

$$\begin{pmatrix} x_1 - \mu_{x1} \\ x_2 - \mu_{x2} \end{pmatrix}^T \Sigma^{-1} \begin{pmatrix} x_1 - \mu_{x1} \\ x_2 - \mu_{x2} \end{pmatrix} \quad \text{see PCA lecture}$$

$$N_2(x_1, x_2) = \frac{1}{2\pi \det(\Sigma)^{1/2}} \exp -\frac{1}{2} \left[\begin{pmatrix} x_1 - \mu_{x1} \\ x_2 - \mu_{x2} \end{pmatrix}^T \Sigma^{-1} \begin{pmatrix} x_1 - \mu_{x1} \\ x_2 - \mu_{x2} \end{pmatrix} \right]$$

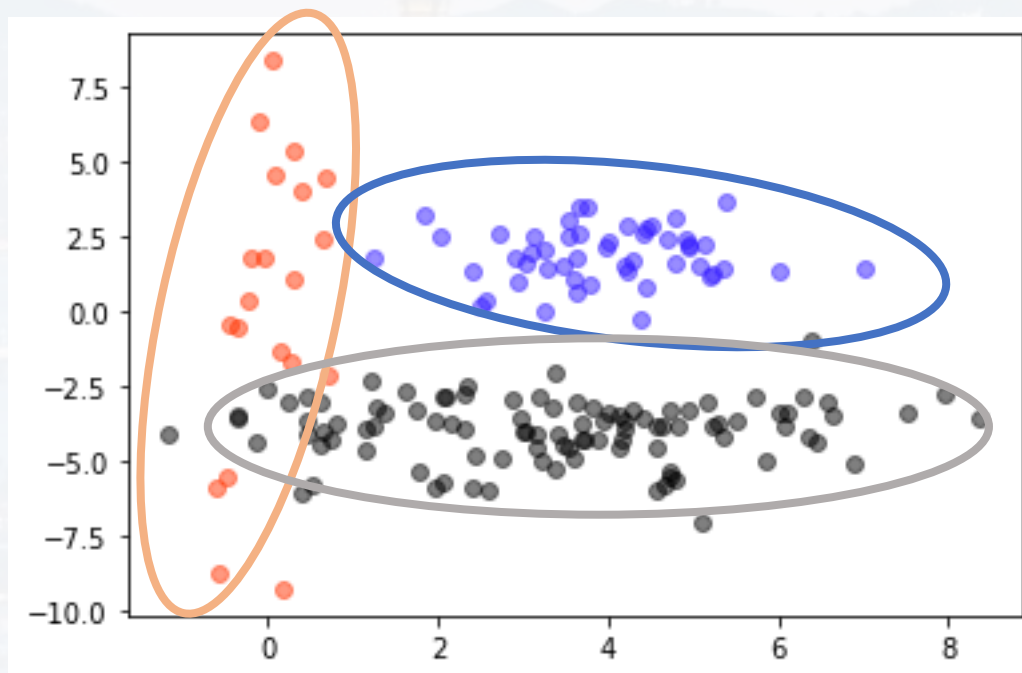


Gaussian Mixture Models

n features

$$N_k(x_1, \dots, x_n) = \frac{1}{(2\pi)^{n/2} \det(\Sigma)^{1/2}} \exp \left[-\frac{1}{2} (x - \mu)^T \Sigma^{-1} (x - \mu) \right]$$

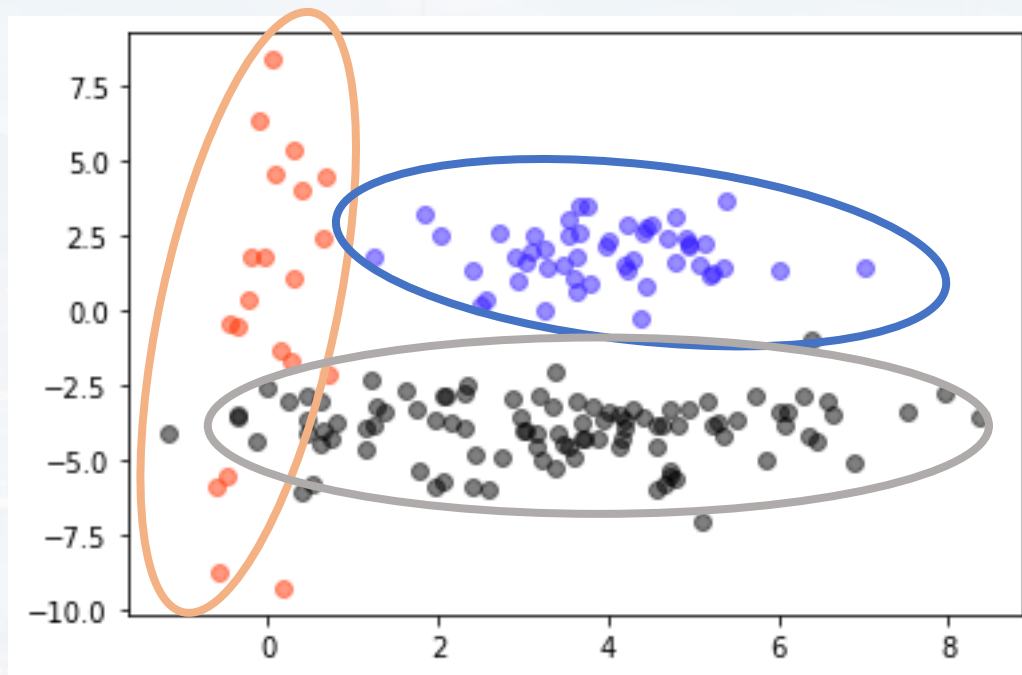
vectors x and μ



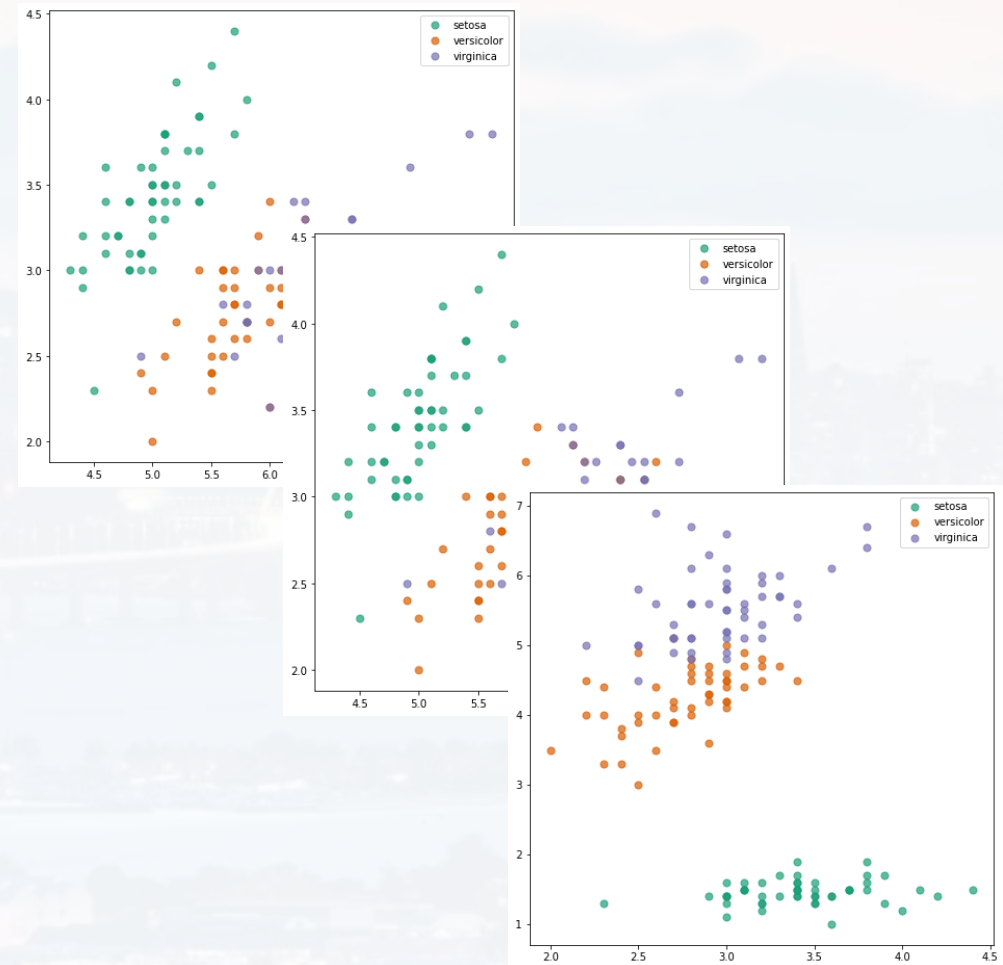
two features, $k=3$ components

Gaussian Mixture Models

two features, $k=3$ components



four features, $k=3$ components



Gaussian Mixture Models

idea: fitting the data to a GMM → analytical functions to **calculate** probabilities for labels

different algorithms:

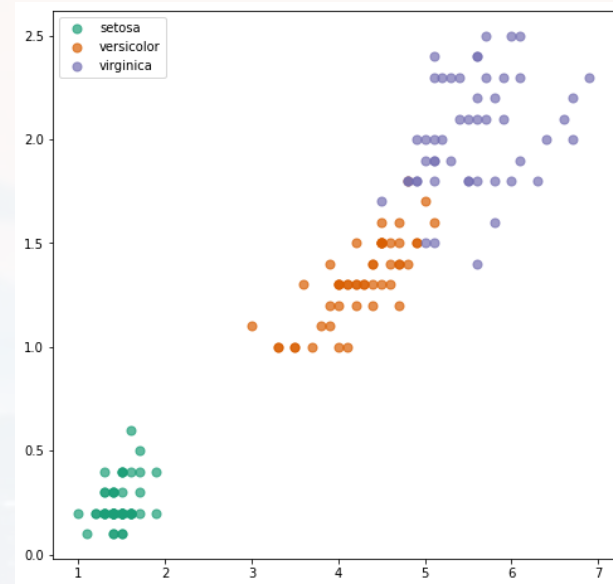
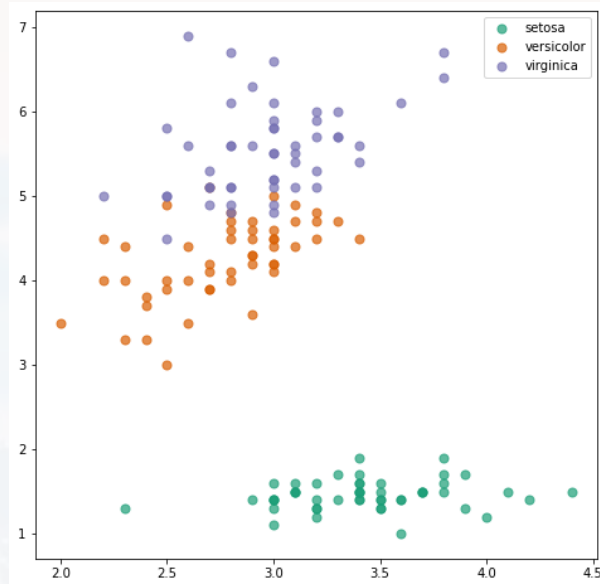
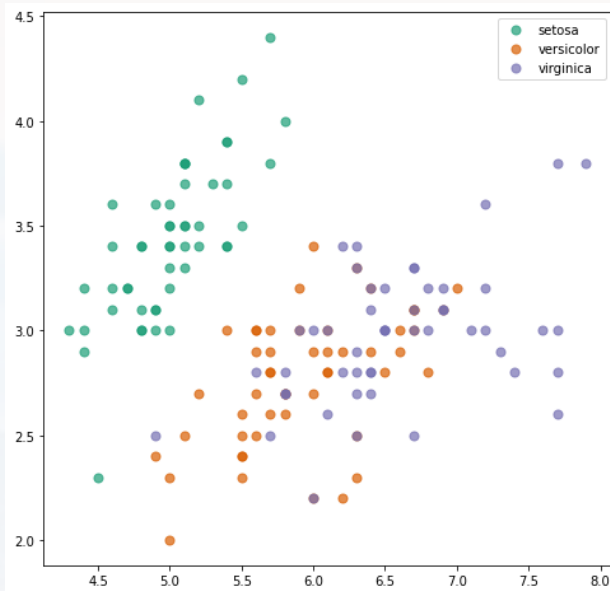
- Bayesian
- **Expectation Maximization**
- ...

```
my_model = GaussianMixture(n_components = k, random_state = 0).fit(X)
```

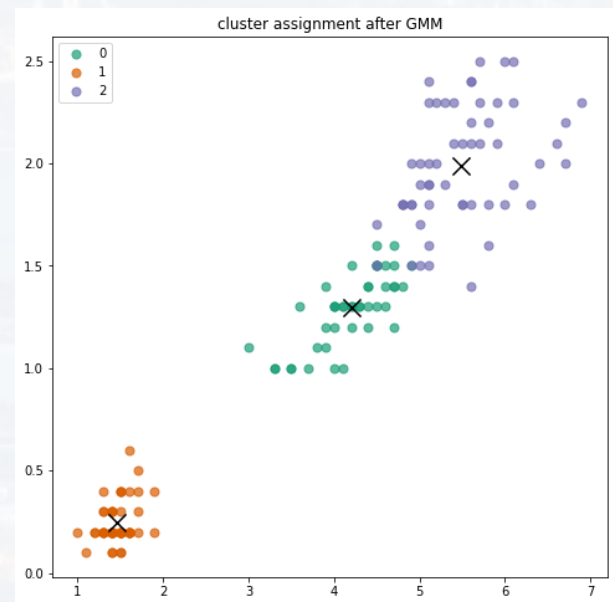
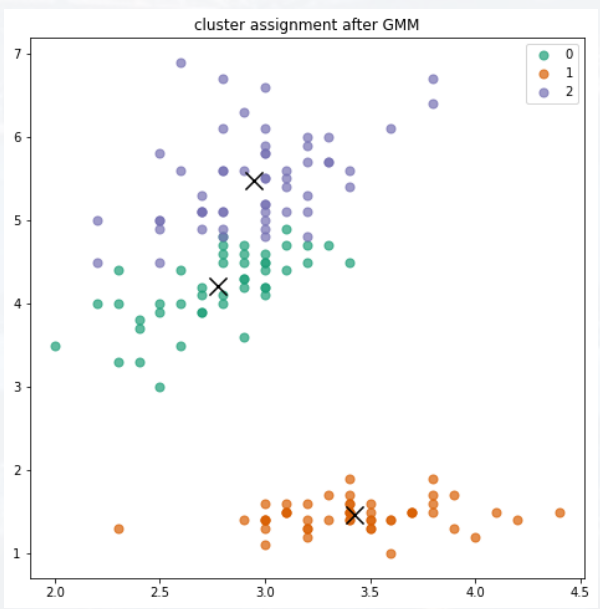
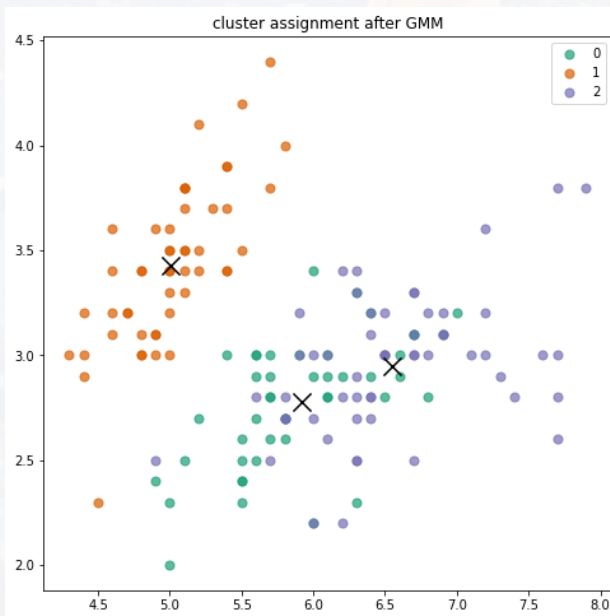
```
Center = my_model.means_  
PredLabels = my_model.predict(X)
```

setting initial labels
randomly

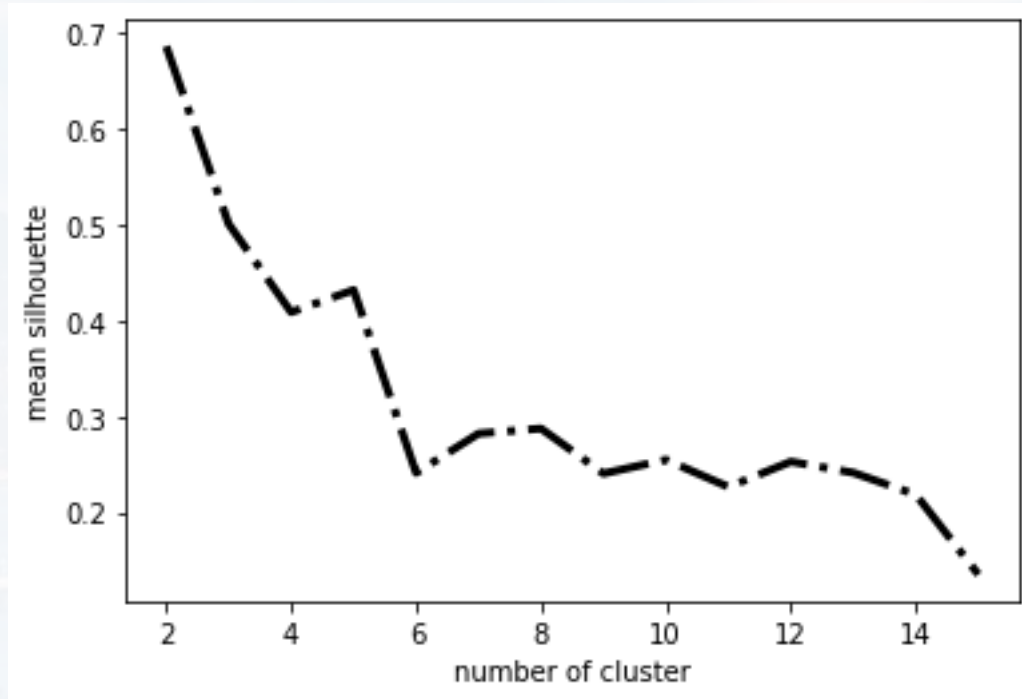
true classes



assigned classes



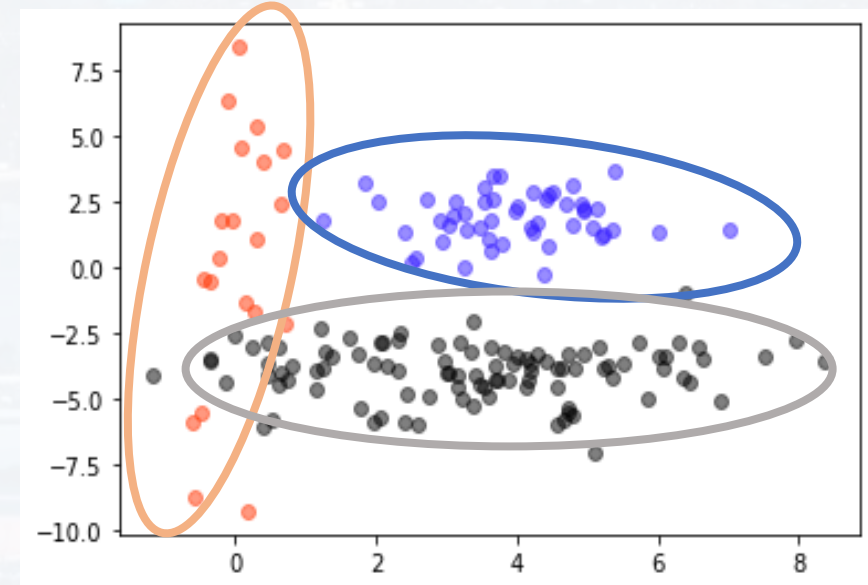
we run k-means now for the **full 4D** dataset
+ evaluate clustering with silhouette

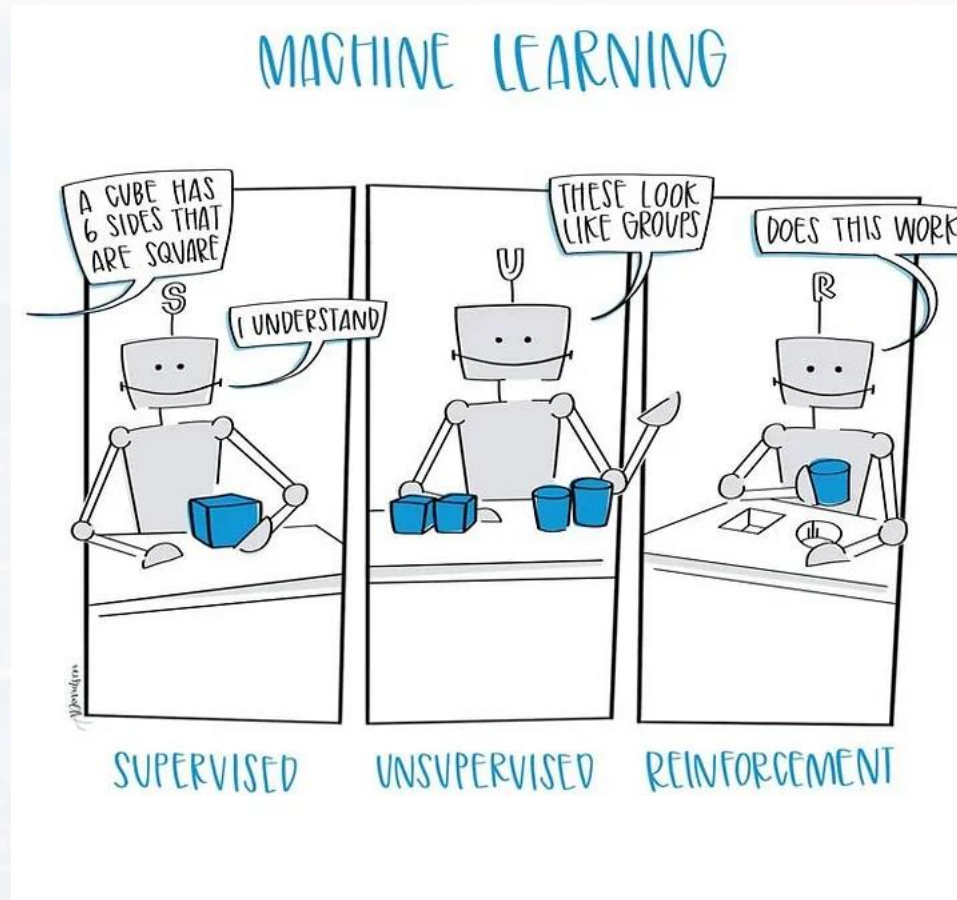


accuracy for 4D $k = 3$: 93%

topics for the discussion/office hour:

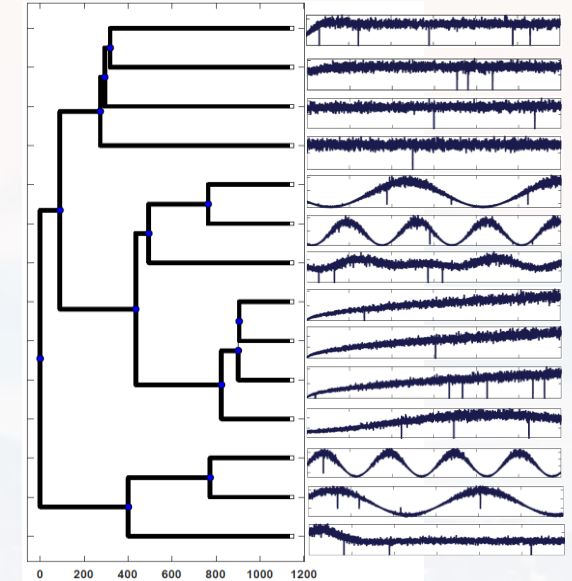
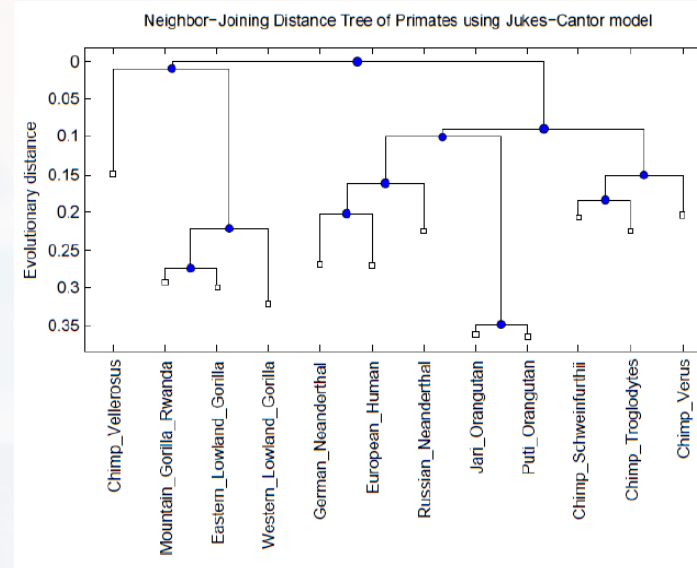
- EM algorithm
- mean, variance and covariance in more detail



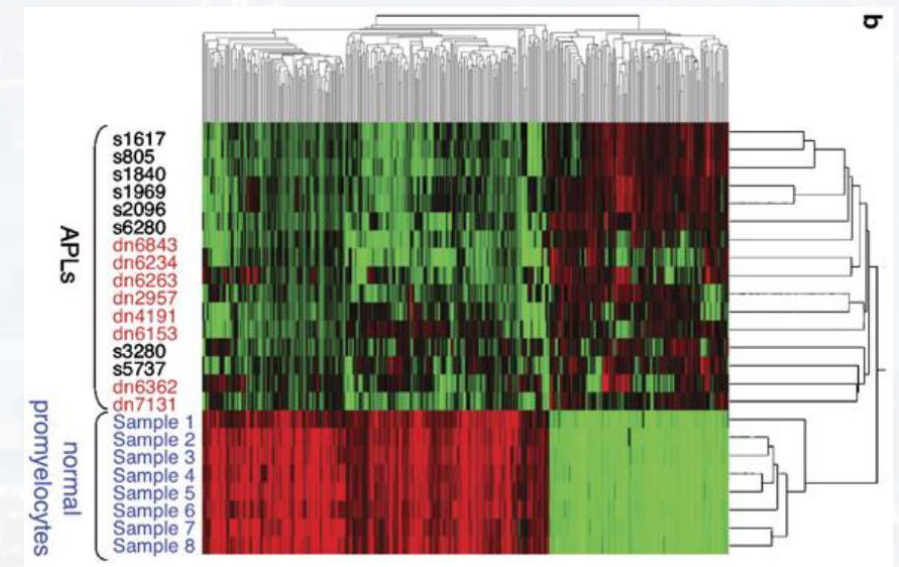


Outline

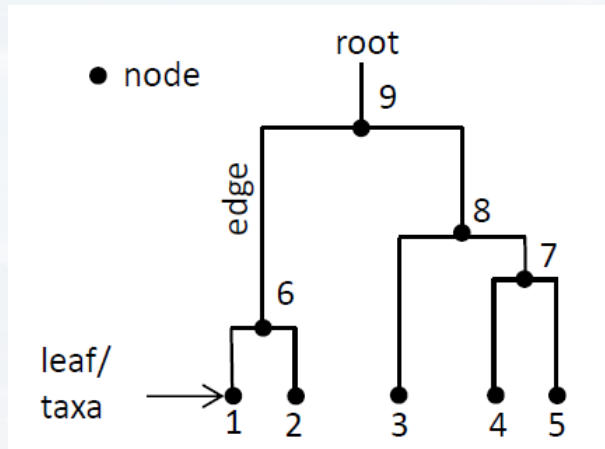
- K - means
- GMM
- **trees**



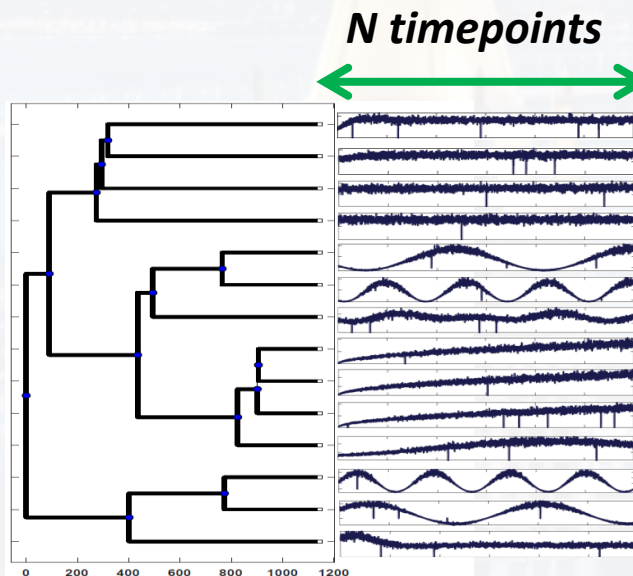
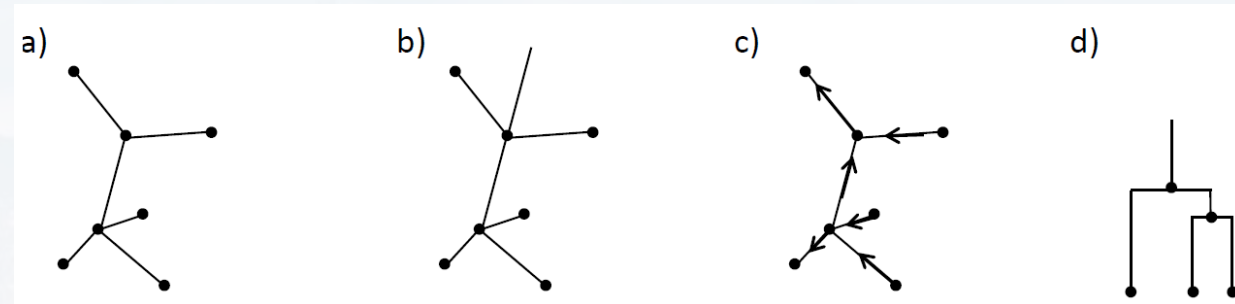
- What is a tree?
- Different kinds of trees...?
- How to build a tree?
- Why do we need trees?
- Examples...



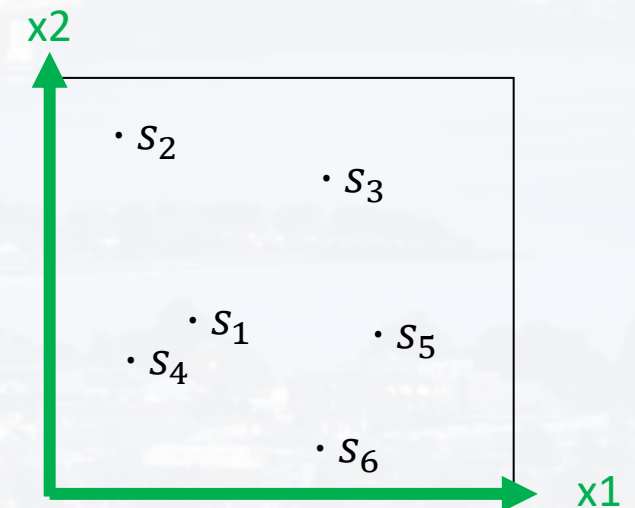
trees are a subclass of graphs (but: not fully connected \rightarrow “hierarchy”, no loops):



- a) unrooted, undirected multinary tree
- b) rooted, undirected multinary tree
- c) unrooted, directed multinary tree
- d) rooted, undirected binary tree

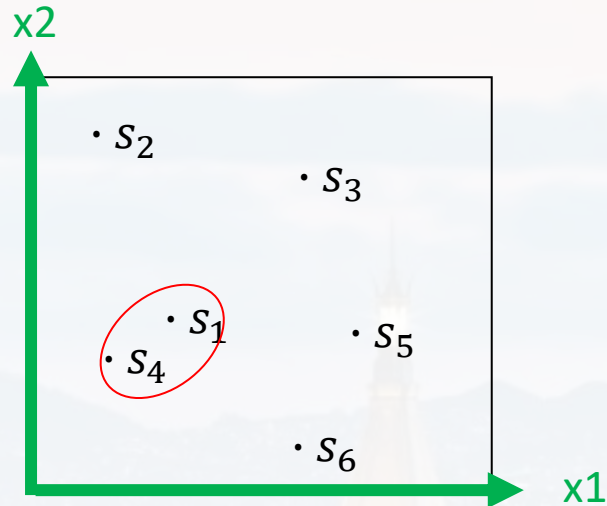


*each sample s is
a vector of N rows,
hence, a data point
in N -D*



constructing trees:

- calculating a distance between each pair of samples



	s_1	s_2	s_3	s_4	s_5	s_6
s_1	0	$d(s_1, s_2)$	$d(s_1, s_3)$	$d(s_1, s_4)$	$d(s_1, s_5)$	$d(s_1, s_6)$
s_2		0	$d(s_2, s_3)$	$d(s_2, s_4)$	$d(s_2, s_5)$	$d(s_2, s_6)$
s_3			0	$d(s_3, s_4)$	$d(s_3, s_5)$	$d(s_3, s_6)$
s_4				0	$d(s_4, s_5)$	$d(s_4, s_6)$
s_5					0	$d(s_5, s_6)$
s_6						0

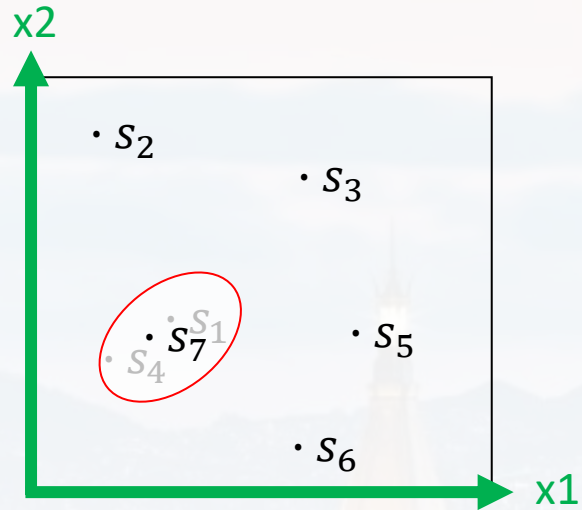
Question: What could be a proper distance definition here?

...once a distance has been defined...

→ find the closest pair

$$\begin{array}{c}
 t_4 \quad \left[\begin{array}{c} \text{---} \\ | \quad | \\ \bullet \quad \bullet \\ s_4 \quad s_1 \end{array} \right]
 \end{array}
 \quad
 t_1 = t_4 = \frac{1}{2} d(s_1, s_4)$$

constructing trees:



- calculating a distance between each pair of samples

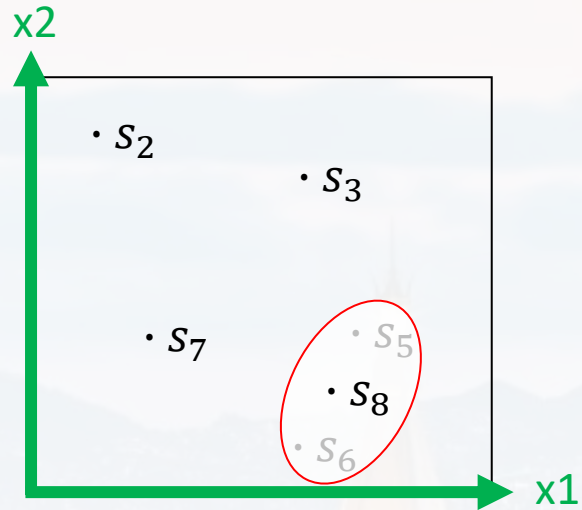
	s_1	s_2	s_3	s_4	s_5	s_6
s_1	0	$d(s_1, s_2)$	$d(s_1, s_3)$	$d(s_1, s_4)$	$d(s_1, s_5)$	$d(s_1, s_6)$
s_2		0	$d(s_2, s_3)$	$d(s_2, s_4)$	$d(s_2, s_5)$	$d(s_2, s_6)$
s_3			0	$d(s_3, s_4)$	$d(s_3, s_5)$	$d(s_3, s_6)$
s_4				0	$d(s_4, s_5)$	$d(s_4, s_6)$
s_5					0	$d(s_5, s_6)$
s_6						0

- treat it as a new cluster $s_{1,4}$
- use average of distance from cluster elements

$$\left[\begin{array}{c} \bullet \\ | \\ \bullet \end{array} \right] \begin{array}{c} s_7 \\ s_4 \end{array} \quad t_1 = t_4 = \frac{1}{2} d(s_1, s_4)$$

constructing trees:

- calculating a distance between each pair of samples

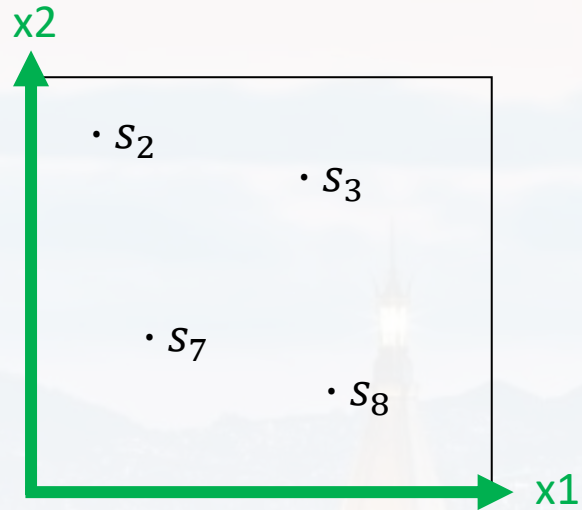


	s_1	s_2	s_3	s_4	s_5	s_6
s_1	0	$d(s_1, s_2)$	$d(s_1, s_3)$	$d(s_1, s_4)$	$d(s_1, s_5)$	$d(s_1, s_6)$
s_2		0	$d(s_2, s_3)$	$d(s_2, s_4)$	$d(s_2, s_5)$	$d(s_2, s_6)$
s_3			0	$d(s_3, s_4)$	$d(s_3, s_5)$	$d(s_3, s_6)$
s_4				0	$d(s_4, s_5)$	$d(s_4, s_6)$
s_5					0	$d(s_5, s_6)$
s_6						0

→ find the closest pair
(now including the cluster)



constructing trees:

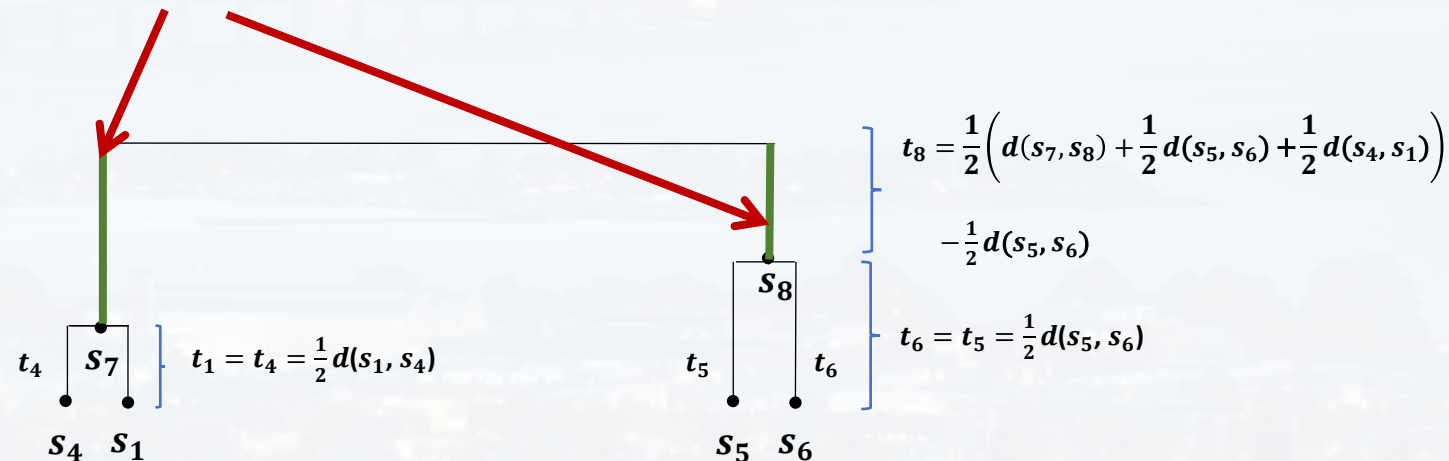


- calculating a distance between each pair of samples

	s_1	s_2	s_3	s_4	s_5	s_6
s_1	0	$d(s_1, s_2)$	$d(s_1, s_3)$	$d(s_1, s_4)$	$d(s_1, s_5)$	$d(s_1, s_6)$
s_2		0	$d(s_2, s_3)$	$d(s_2, s_4)$	$d(s_2, s_5)$	$d(s_2, s_6)$
s_3			0	$d(s_3, s_4)$	$d(s_3, s_5)$	$d(s_3, s_6)$
s_4				0	$d(s_4, s_5)$	$d(s_4, s_6)$
s_5					0	$d(s_5, s_6)$
s_6						0

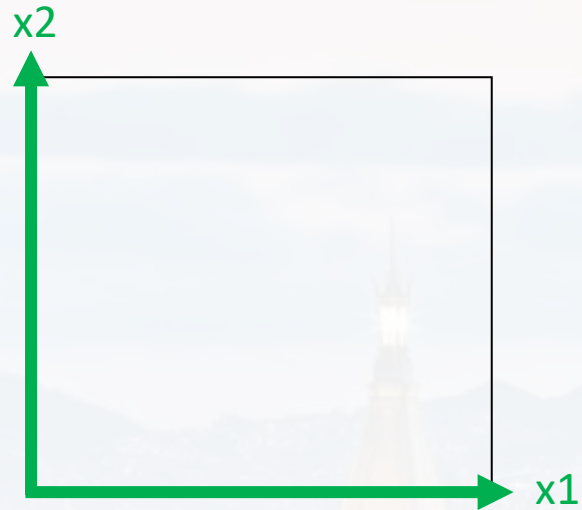
→ and so on....

$$d(s_7, s_8) = \frac{d(s_5, s_7) + d(s_6, s_7)}{2} = \frac{d(s_1, s_5) + d(s_4, s_5) + d(s_1, s_6) + d(s_4, s_6)}{4}$$



constructing trees:

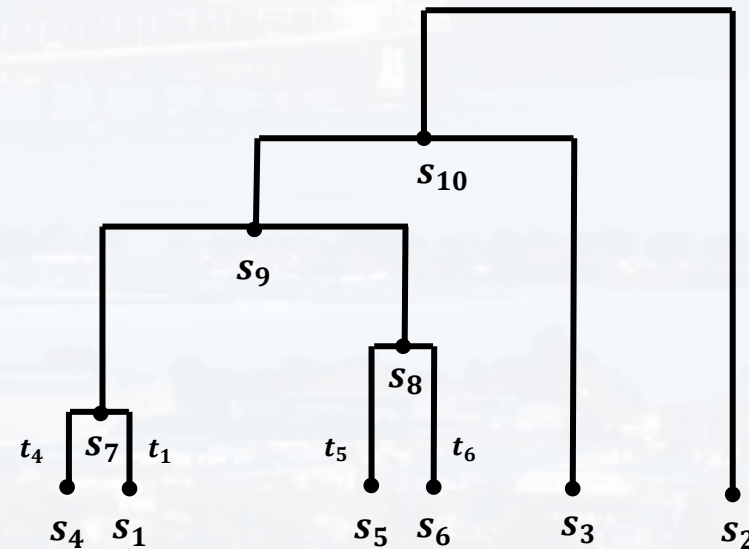
- calculating a distance between each pair of samples



	s_1	s_2	s_3	s_4	s_5	s_6
s_1	0	$d(s_1, s_2)$	$d(s_1, s_3)$	$d(s_1, s_4)$	$d(s_1, s_5)$	$d(s_1, s_6)$
s_2		0	$d(s_2, s_3)$	$d(s_2, s_4)$	$d(s_2, s_5)$	$d(s_2, s_6)$
s_3			0	$d(s_3, s_4)$	$d(s_3, s_5)$	$d(s_3, s_6)$
s_4				0	$d(s_4, s_5)$	$d(s_4, s_6)$
s_5					0	$d(s_5, s_6)$
s_6						0

....finally

→ **Unweighted Pair Group Method**
Using **Arithmetic Averages (UPGMA)**



Python libraries:

libraries from the *Bio* package

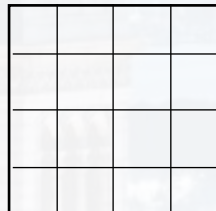
```
from Bio.Phylo.TreeConstruction import DistanceCalculator, DistanceTreeConstructor  
from Bio import Phylo  
from scipy import spatial
```

for plotting a
phylogenetic tree

for rearranging a
distance matrix

general idea:

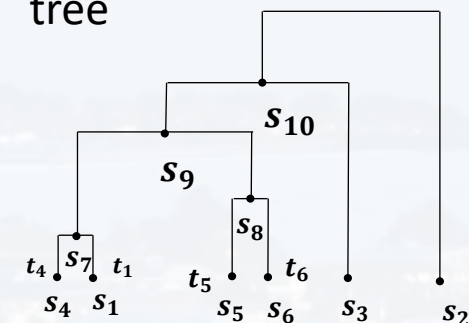
distance matrix



linkage algorithm (e.g. UPGMA)



tree



Python libraries:

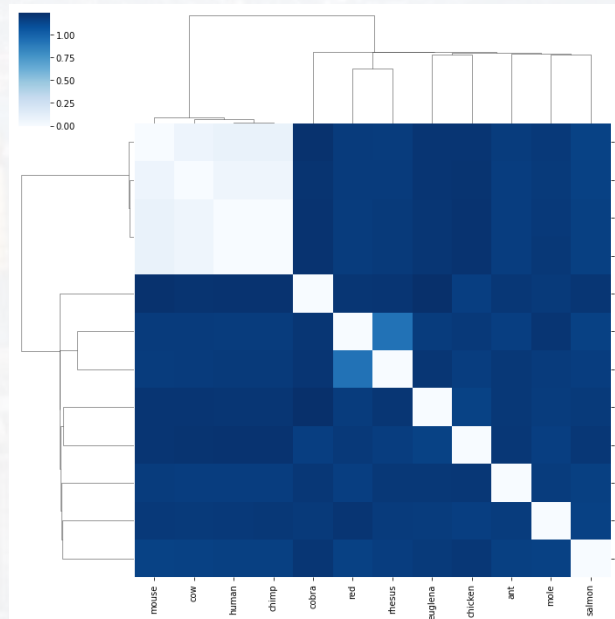
also most heatmap tools have some abilities to construct trees:

`sns.clustermap`

```
sns.clustermap(D_df, cmap = 'Blues', \
               row_cluster = True, col_cluster = True, \
               metric = 'euclidean', method = 'average', \
               yticklabels = True)
```

`plt.show()`

similar to UPGMA

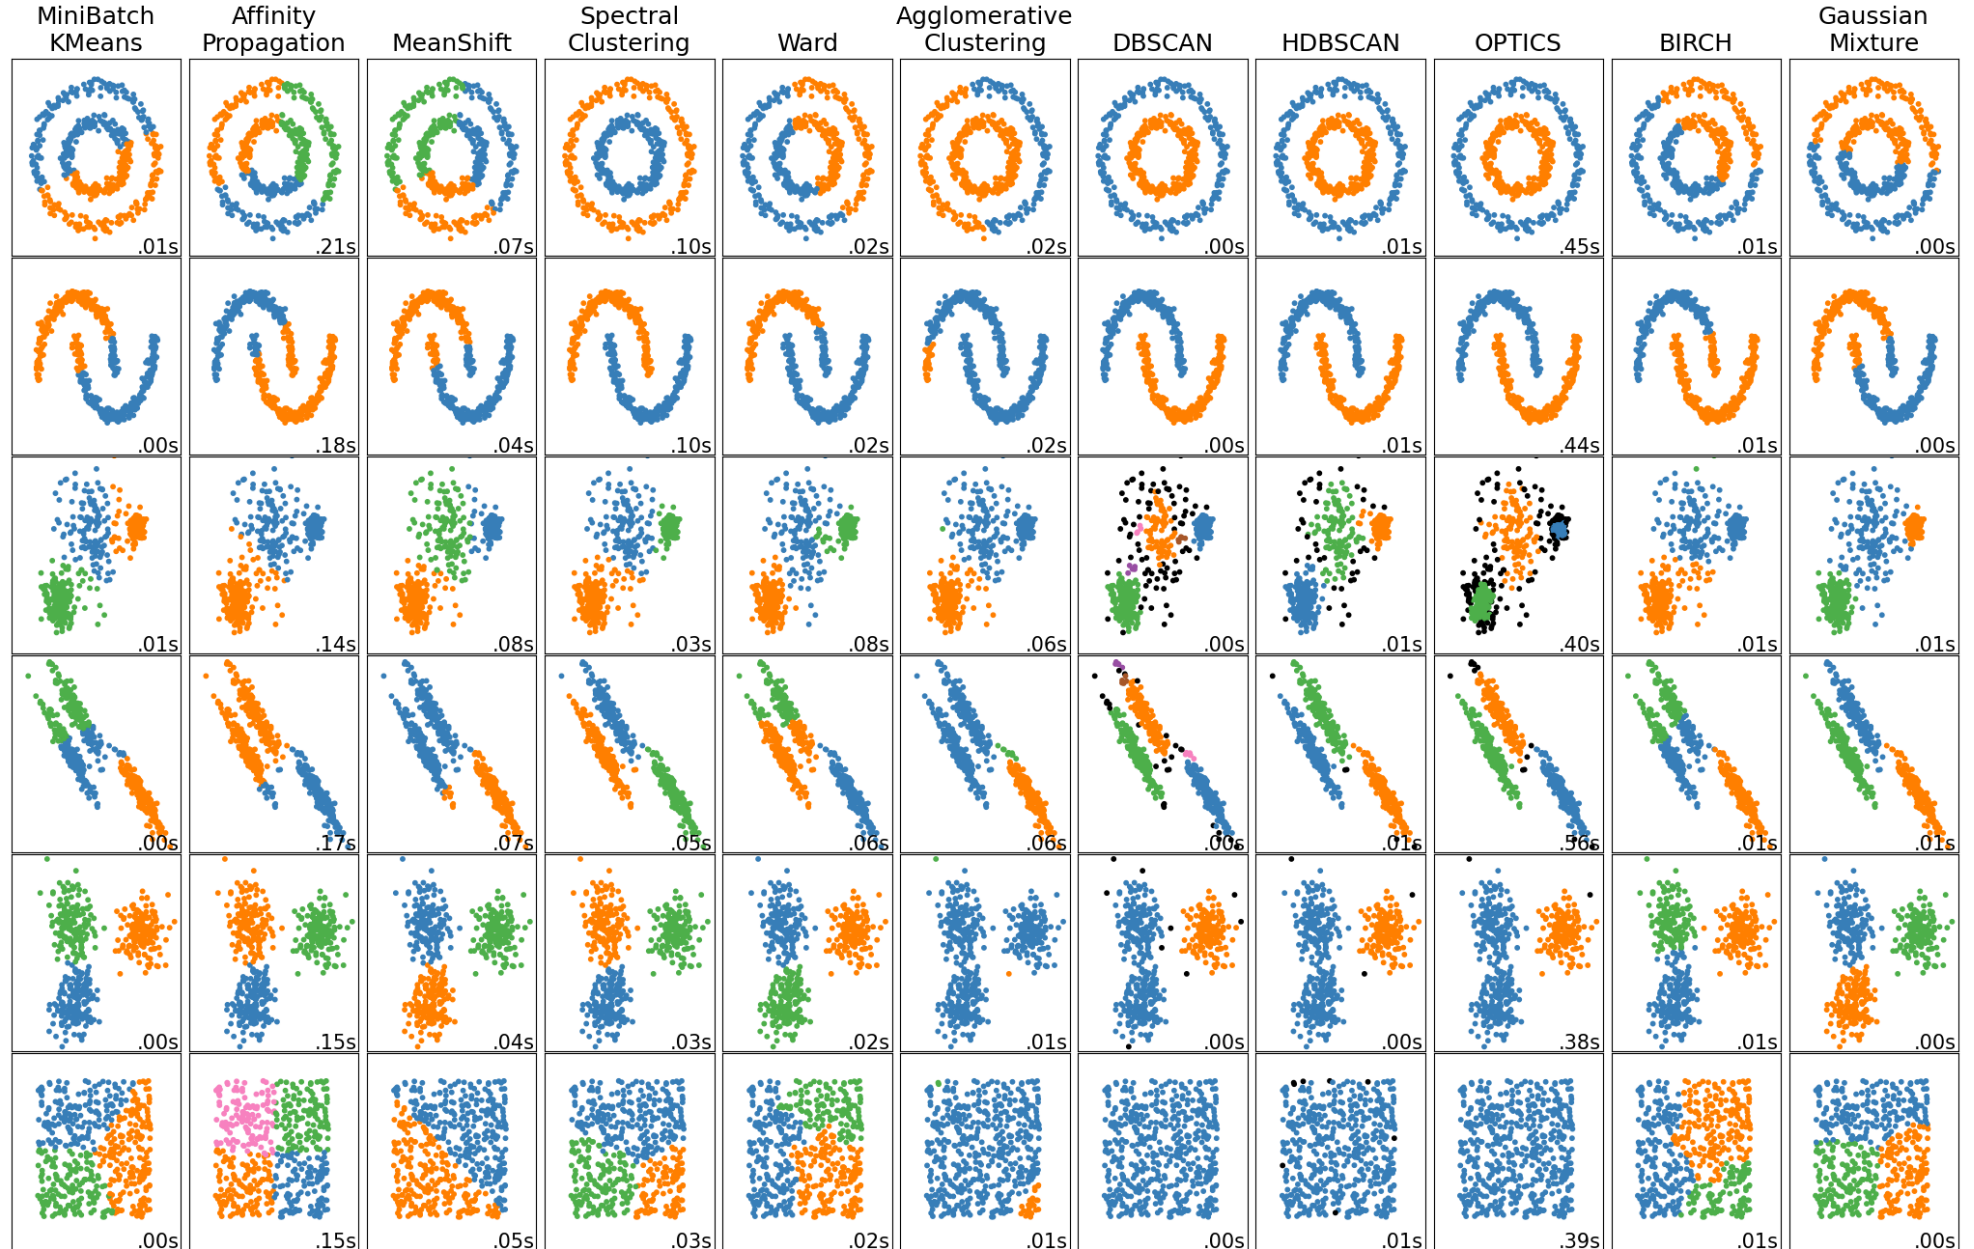


topics for the discussion/office hour:

- distances
- random forest
- graphs



there is a lot more...



Thank you very much for your attention!

