

## Lecture 06:

## Optimization



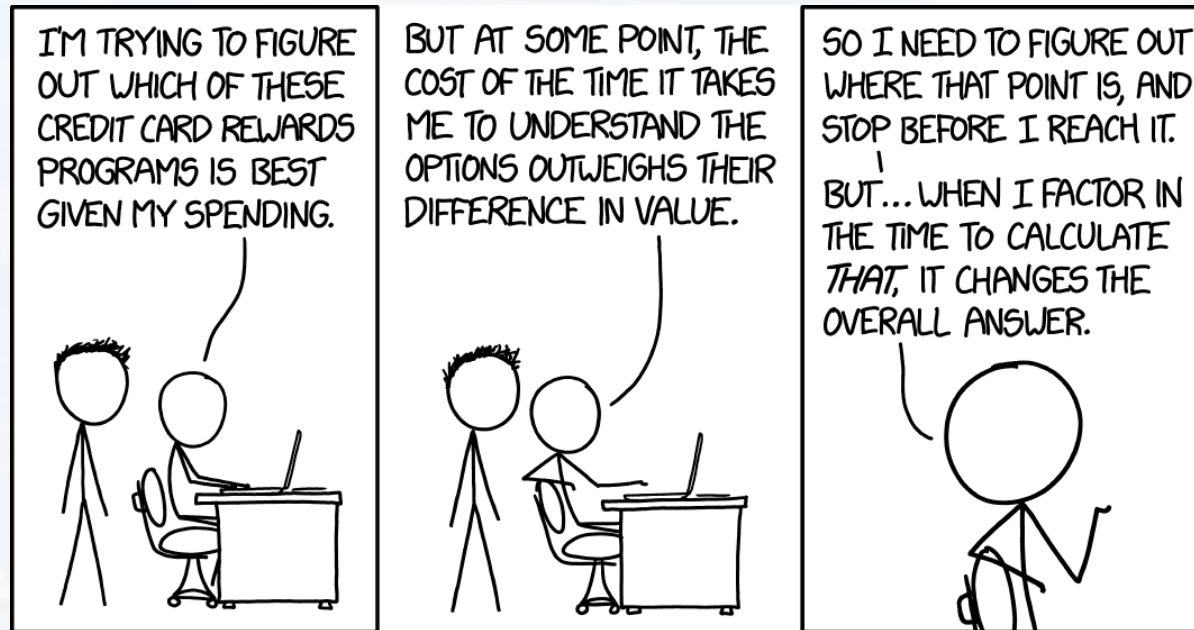
Markus Hohle

University California, Berkeley

**Machine Learning Algorithms**

MSSE 277B, 3 Units

Fall 2024

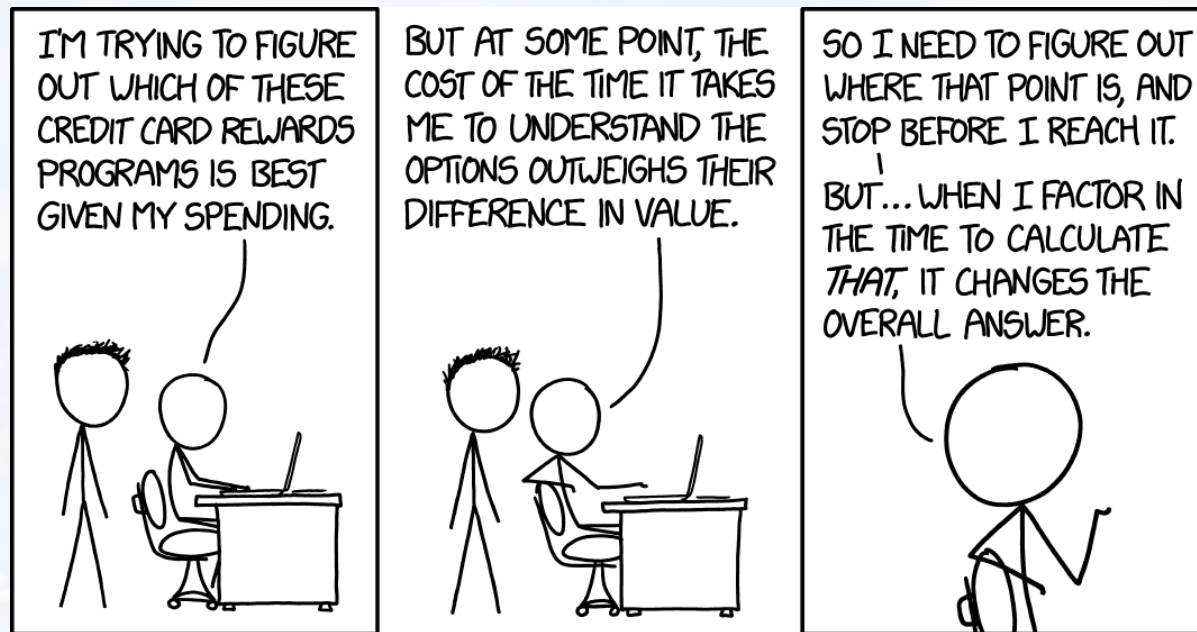


## Outline

### - The Problem

### - Gradient Descent

- Vanilla
- Learning Rate Schedule
- Momentum
- L1 and L2
- More Finetuning



## Outline

### - The Problem


### - Gradient Descent

- Vanilla
- Learning Rate Schedule
- Momentum
- L1 and L2
- More Finetuning





Any algorithm needs a “goal” aka **objective function** that has to be **optimized** (finding an **extreme**)

regression, e. g.  
curve fitting



minimize: 
$$\chi_{red}^2 = \frac{1}{N - p - 1} \sum_{i=1}^N \frac{(\bar{y}(model)_i - y_i)^2}{\sigma_i^2}$$

classification



maximize: accuracy


minimize: cross entropy

$$S = - \sum_i p(true)_i \cdot \ln p(model)_i$$





Any algorithm needs a “goal” aka **objective function** that has to be **optimized** (finding an **extreme**)



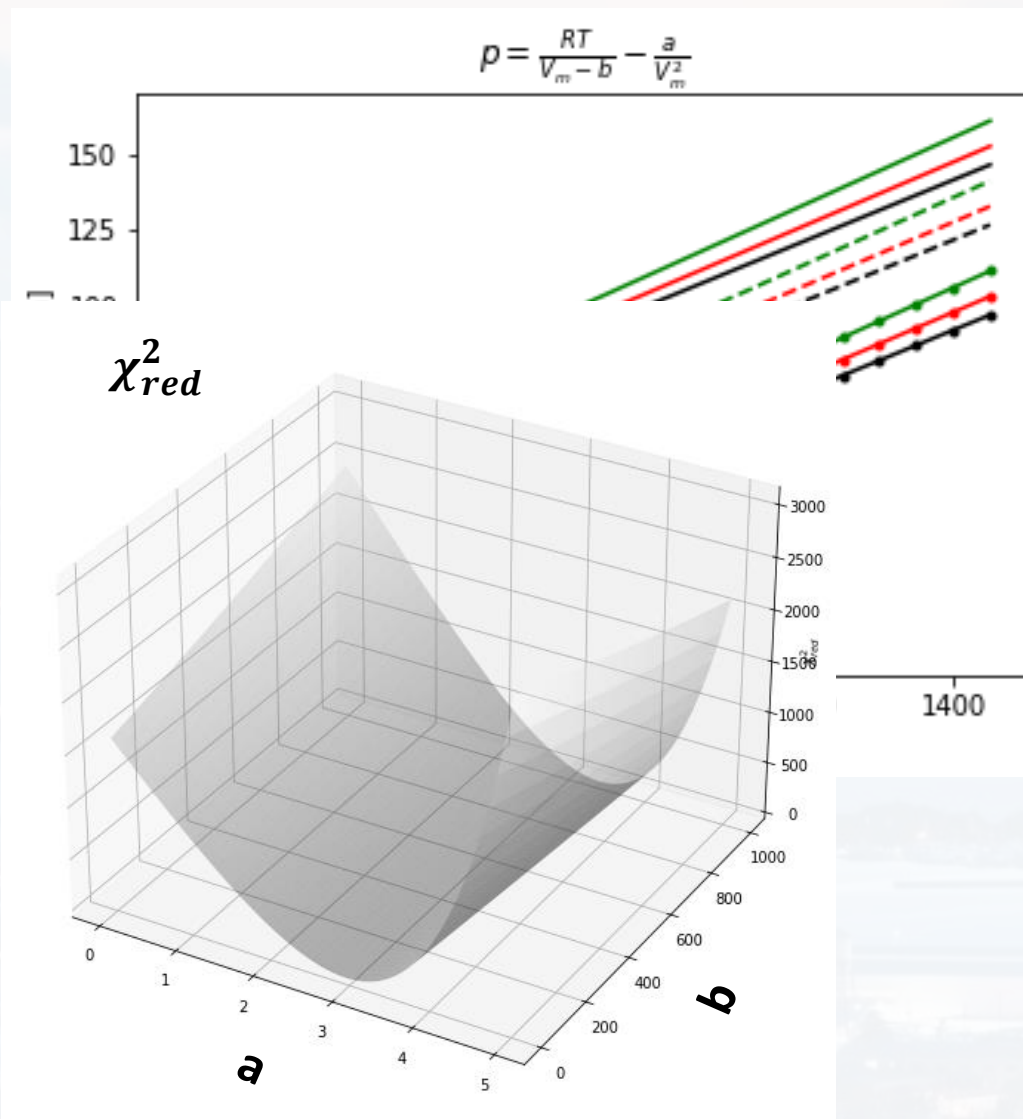
finding ***a*** and ***b*** of  
a van-der-Waals gas

if critical points are not  
accessible

→ fitting curve, finding ***a*** and ***b***




Any algorithm needs a “goal” aka **objective function** that has to be *optimized* (finding an **extreme**)



$\log(\chi_{red}^2)$

finding **a** and **b** of





Any algorithm needs a “goal” aka **objective function** that has to be **optimized** (finding an **extreme**)

Often, the extreme of the objective function is subject to **constraints**

cross entropy

$$S = - \sum_i p(\text{true})_i \cdot \ln p(\text{model})_i$$

constrain:  $\sum_i p_i = 1$

→ Lagrangian Multipliers and variational calculus

→ mathematically: *Free Energy like term = Energy like term – Entropy term*

examples:

- KL divergence
- Lasso method (linear regression)
- actual energy → Boltzmann distribution


etc



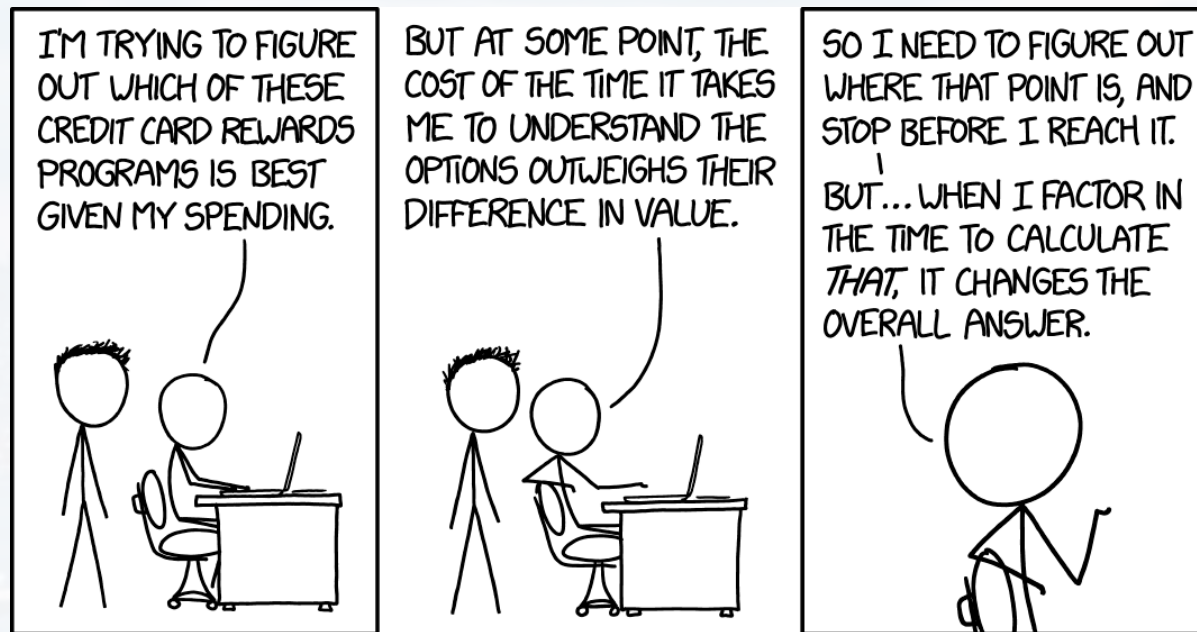


Any algorithm needs a “goal” aka **objective function** that has to be *optimized* (finding an **extreme**)

These functions are very complicated, not analytical at all







## Outline

- The Problem

- **Gradient Descent**

- Vanilla

- Learning Rate Schedule

- Momentum

- L1 and L2

- More Finetuning



main application: **ANN!**









Vanilla

$$\left. \frac{dy}{dx} \right|_{x_0} \approx \frac{y(x_0 + \Delta x) - y(x_0 - \Delta x)}{2\Delta x}$$



$$x(t=2) = x(t=1) - \varepsilon \left. \frac{dy}{dx} \right|_{x(t=1)}$$




$\varepsilon > 0$




Vanilla

$$\left. \frac{dy}{dx} \right|_{x_0} \approx \frac{y(x_0 + \Delta x) - y(x_0 - \Delta x)}{2\Delta x}$$



$$x(t=3) = x(t=2) - \epsilon \left. \frac{dy}{dx} \right|_{x(t=2)}$$





$\epsilon > 0$





Vanilla

$$\left. \frac{dy}{dx} \right|_{x_0} \approx \frac{y(x_0 + \Delta x) - y(x_0 - \Delta x)}{2\Delta x}$$




$$x(t=4) = x(t=3) - \varepsilon \left. \frac{dy}{dx} \right|_{x(t=3)}$$

$\varepsilon > 0$




Vanilla

$$\left. \frac{dy}{dx} \right|_{x_0} \approx \frac{y(x_0 + \Delta x) - y(x_0 - \Delta x)}{2\Delta x}$$



$$x(t=4) = x(t=3) - \varepsilon \left. \frac{dy}{dx} \right|_{x(t=3)}$$


$\varepsilon > 0$






Vanilla

$$\left. \frac{dy}{dx} \right|_{x_0} \approx \frac{y(x_0 + \Delta x) - y(x_0 - \Delta x)}{2\Delta x}$$



$$x(t=2) = x(t=1) - \epsilon \left. \frac{dy}{dx} \right|_{x(t=1)}$$



$\epsilon > 0$







Vanilla

$$\left. \frac{dy}{dx_1} \right|_{x_1(0)} \approx \frac{y(x_1(0) + \Delta x_1) - y(x_1(0) - \Delta x_1)}{2\Delta x_1}$$

$$\left. \frac{dy}{dx_2} \right|_{x_2(0)} \approx \frac{y(x_2(0) + \Delta x_2) - y(x_2(0) - \Delta x_2)}{2\Delta x_2}$$







Vanilla

$$\left. \frac{dy}{dx_1} \right|_{x_1(0)} \approx \frac{y(x_1(0) + \Delta x_1) - y(x_1(0) - \Delta x_1)}{2\Delta x_1}$$

$$\left. \frac{dy}{dx_2} \right|_{x_2(0)} \approx \frac{y(x_2(0) + \Delta x_2) - y(x_2(0) - \Delta x_2)}{2\Delta x_2}$$

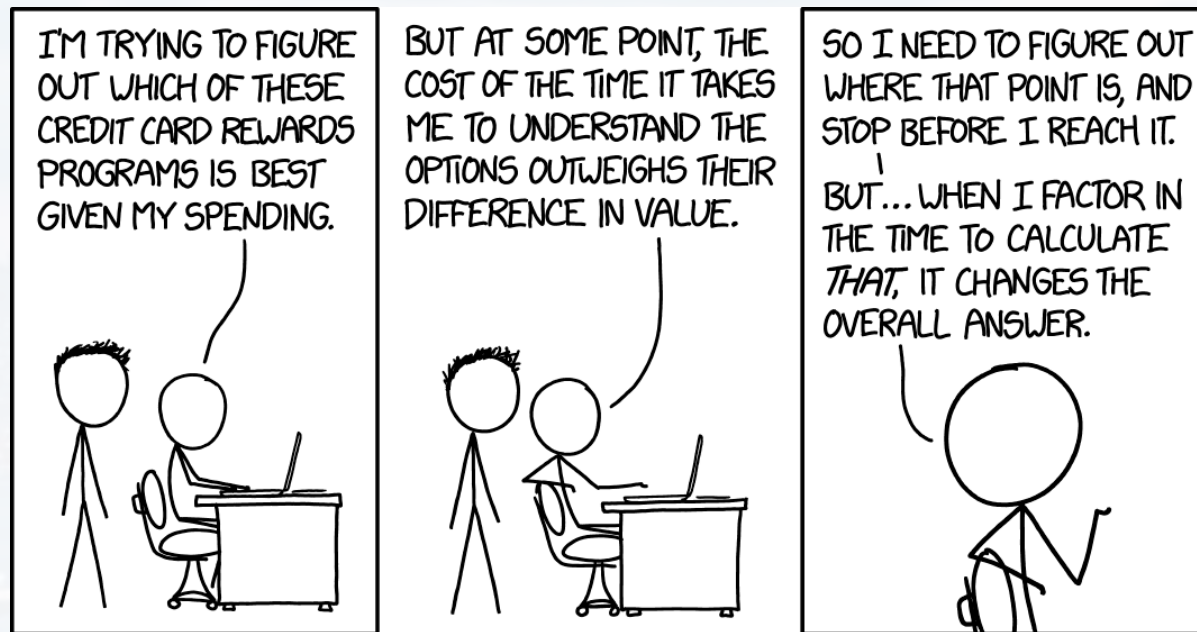
⋮

$$\left. \frac{dy}{dx_i} \right|_{x_i(0)} \approx \frac{y(x_i(0) + \Delta x_i) - y(x_i(0) - \Delta x_i)}{2\Delta x_i}$$

⋮

$$\left. \frac{dy}{dx_N} \right|_{x_N(0)} \approx \frac{y(x_N(0) + \Delta x_N) - y(x_N(0) - \Delta x_N)}{2\Delta x_N}$$

$$\begin{pmatrix} \left. \frac{dy}{dx_1} \right|_{x_1(0)} \\ \vdots \\ \left. \frac{dy}{dx_i} \right|_{x_i(0)} \\ \vdots \\ \left. \frac{dy}{dx_N} \right|_{x_N(0)} \end{pmatrix} = \text{grad}(y)_x$$



## Outline

- The Problem

- **Gradient Descent**

- Vanilla

- Learning Rate Schedule

- Momentum

- L1 and L2

- More Finetuning



### Learning Rate Schedule

$$\left. \frac{dy}{dx} \right|_{x_0} \approx \frac{y(x_0 + \Delta x) - y(x_0 - \Delta x)}{2\Delta x}$$


$$x(t+1) = x(t) - \epsilon \left. \frac{dy}{dx} \right|_{x(t)}$$

$$\epsilon > 0$$

called *learning rate*

$$\Delta x = - \epsilon \left. \frac{dy}{dx} \right|_{x(t)}$$

defines how large  
the leap  $\Delta x$  is





### Learning Rate Schedule

$$\left. \frac{dy}{dx} \right|_{x_0} \approx \frac{y(x_0 + \Delta x) - y(x_0 - \Delta x)}{2\Delta x}$$


$$x(t+1) = x(t) - \epsilon \left. \frac{dy}{dx} \right|_{x(t)}$$

$$\epsilon > 0$$

called *learning rate*

$$\Delta x = - \epsilon \left. \frac{dy}{dx} \right|_{x(t)}$$

defines how large  
the leap  $\Delta x$  is







### Learning Rate Schedule

$$\left. \frac{dy}{dx} \right|_{x_0} \approx \frac{y(x_0 + \Delta x) - y(x_0 - \Delta x)}{2\Delta x}$$

$$x(t+1) = x(t) - \epsilon \left. \frac{dy}{dx} \right|_{x(t)}$$

$$\epsilon > 0$$


called *learning rate*

$$\Delta x = - \epsilon \left. \frac{dy}{dx} \right|_{x(t)}$$

defines how large the leap  $\Delta x$  is

... and so on...

→ smaller  $\epsilon$  ?





### Learning Rate Schedule

$$\left. \frac{dy}{dx} \right|_{x_0} \approx \frac{y(x_0 + \Delta x) - y(x_0 - \Delta x)}{2\Delta x}$$

$$x(t+1) = x(t) - \epsilon \left. \frac{dy}{dx} \right|_{x(t)}$$

$$\epsilon > 0$$

called *learning rate*


$$\Delta x = - \epsilon \left. \frac{dy}{dx} \right|_{x(t)}$$

defines how large the leap  $\Delta x$  is

... and so on...

→ smaller  $\epsilon$  ?

Takes too long!





### Learning Rate Schedule

$$\left. \frac{dy}{dx} \right|_{x_0} \approx \frac{y(x_0 + \Delta x) - y(x_0 - \Delta x)}{2\Delta x}$$

$$x(t+1) = x(t) - \epsilon \left. \frac{dy}{dx} \right|_{x(t)}$$

learning rate as function of t:

$$\epsilon > 0$$

called *learning rate*

$$\epsilon(t) = \frac{\epsilon_0}{1 + \kappa t} \quad \text{decay rate } \kappa$$

$$\Delta x = - \epsilon \left. \frac{dy}{dx} \right|_{x(t)}$$

defines how large the leap  $\Delta x$  is





### Learning Rate Schedule

$$\left. \frac{dy}{dx} \right|_{x_0} \approx \frac{y(x_0 + \Delta x) - y(x_0 - \Delta x)}{2\Delta x}$$

$$x(t+1) = x(t) - \epsilon \left. \frac{dy}{dx} \right|_{x(t)}$$

learning rate as function of t:

$$\epsilon > 0$$

called *learning rate*

$$\epsilon(t) = \frac{\epsilon_0}{1 + \kappa t} \quad \text{decay rate } \kappa$$

$$\Delta x = - \epsilon \left. \frac{dy}{dx} \right|_{x(t)}$$

defines how large the leap  $\Delta x$  is

can also be a stepwise function (learning rate schedule)





### Learning Rate Schedule


learning rate as function of t:

$$\epsilon(t) = \frac{\epsilon_0}{1 + \kappa t} \quad \text{decay rate } \kappa$$

$$\Delta x = -\epsilon \left. \frac{dy}{dx} \right|_{x(t)}$$

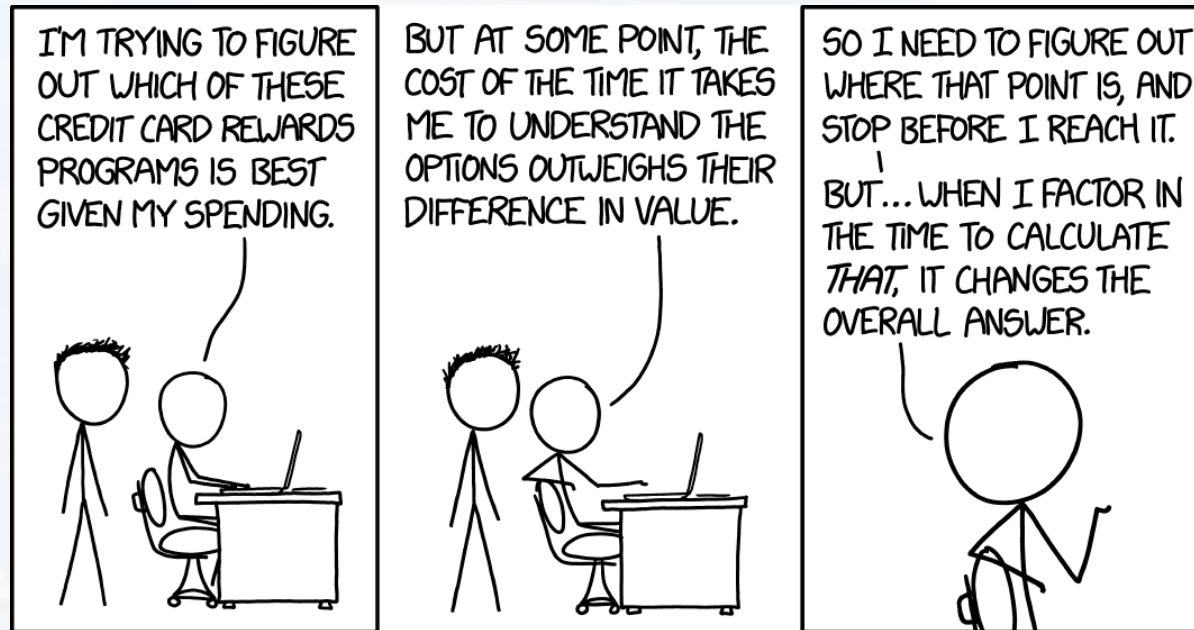
defines how large the leap  $\Delta x$  is

can also be a stepwise function (learning rate schedule)



$$\epsilon \rightarrow \frac{\epsilon}{\sqrt{\text{grad}(y)_x}}$$

adaptive gradient, aka **AdaGrad**



## Outline

- The Problem

- **Gradient Descent**

- Vanilla

- Learning Rate Schedule

- Momentum

- L1 and L2




- More Finetuning



### Momentum

even with AdaGrad and learning rate schedule  
→ would get stuck in local minimum


need to roll over → **momentum**







### Momentum



taking the **average** of  $N$  previous gradients


$$\langle grad(y)_{x(t)} \rangle = \frac{1}{N} [grad(y)_{x(t-1)} + grad(y)_{x(t-2)} + \dots + grad(y)_{x(t-N)}]$$

but we want more recent gradients to contribute more than older gradients

→ **weighted average** with weighting factor  $\mu_k$

$$\langle grad(y)_{x(t)} \rangle = \sum_{k=t-N}^{t-1} \mu_k \cdot grad(y)_{x(k)}$$

Finding a clever way to adjust  $\mu_k$  during every iteration  $t$




**weighted average** with weighting factor  $\mu_k$

**Momentum**

Finding a clever way to adjust  $\mu_k$  during every iteration  $t$

$$\langle grad(y)_{x(0)} \rangle = grad(y)_{x(0)} \quad \mu_0 = (0,1)$$

$$\langle grad(y)_{x(1)} \rangle = grad(y)_{x(1)} + \mu_0 \cdot grad(y)_{x(0)}$$



**weighted average** with weighting factor  $\mu_k$

**Momentum**

Finding a clever way to adjust  $\mu_k$  during every iteration  $t$

$$\langle grad(y)_{x(0)} \rangle = grad(y)_{x(0)} \quad \mu_0 = (0,1)$$

$$\langle grad(y)_{x(1)} \rangle = grad(y)_{x(1)} + \mu_0 \cdot grad(y)_{x(0)}$$

$$\langle grad(y)_{x(2)} \rangle = grad(y)_{x(2)} + \boxed{\mu_0} [grad(y)_{x(1)} + \boxed{\mu_0} grad(y)_{x(0)}]$$

$$\mu_{k=2} = \mu_0 \mu_0 = \mu_0^2$$

$$\langle grad(y)_{x(3)} \rangle = grad(y)_{x(3)} + \boxed{\mu_0} [grad(y)_{x(2)} + \boxed{\mu_0} [grad(y)_{x(1)} + \boxed{\mu_0} \cdot grad(y)_{x(0)}]]$$

... and so on...






**weighted average** with weighting factor  $\mu_k$

$\mu_0 = (0,1)$  called "momentum"

$$\langle \text{grad}(y)_{x(3)} \rangle = \text{grad}(y)_{x(3)} +$$

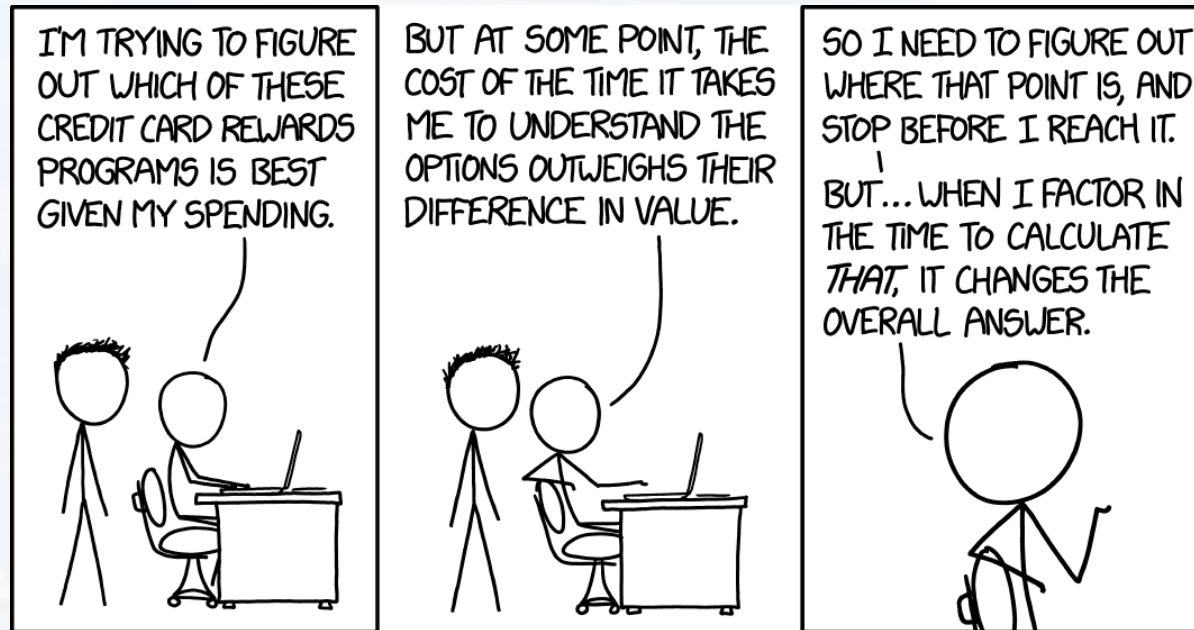
$$\mu_0 \left[ \text{grad}(y)_{x(2)} + \mu_0 \left[ \text{grad}(y)_{x(1)} + \mu_0 \cdot \text{grad}(y)_{x(0)} \right] \right] \quad \dots \text{and so on...}$$



**Momentum**

```
class Optimizer:
```

```
    def __init__(self, learning_rate = 0.1, decay = 0, momentum = 0):
        self.learning_rate      = learning_rate
        self.decay               = decay
        self.current_learning_rate = learning_rate
        self.iterations         = 0
        self.momentum            = momentum
```



## Outline

- The Problem

- **Gradient Descent**

- Vanilla
- Learning Rate Schedule
- Momentum
- L1 and L2
- More Finetuning



### L1 and L2

Any algorithm needs a “goal” aka **objective function** that has to be **optimized** (finding an **extreme**)

Often, the extreme of the objective function is subject to **constraints**

sometimes we have some **prior knowledge** about the **independent variables**

recall: linear regression

finding best  $\beta$  by

$$\min_{\beta} \left\{ \frac{1}{N} \|Y - X\beta\|^2 \right\}$$

now:

constrain: **encourages sparsity of  $\beta$**

$$\min_{\beta} \left\{ \frac{1}{N} \|Y - X\beta\|^2 + \lambda \|\beta\|^1 \right\}$$

$\lambda$  Lagrangian Multiplier

called **L1 regularization**, or LASSO






### L1 and L2

Any algorithm needs a “goal” aka **objective function** that has to be **optimized** (finding an **extreme**)

Often, the extreme of the objective function is subject to **constraints**

sometimes we have some **prior knowledge** about the **independent variables**

### L1 regularization



We often have even hard constraints based on the laws of physics!



Any algorithm needs a “goal” aka **objective function** that has to be **optimized** (finding an **extreme**)

Often, the extreme of the objective function is subject to **constraints**

sometimes we have some **prior knowledge** about the **independent variables**

recall: linear regression

finding best  $\beta$  by

$$\min_{\beta} \left\{ \frac{1}{N} \|Y - X\beta\|^2 \right\}$$

now:

$$\hat{\beta} = (X^T X + \lambda I)^{-1} X^T Y \longrightarrow \min_{\beta} \left\{ \frac{1}{N} \|Y - X\beta\|^2 + \lambda \|\beta\|^2 \right\}$$


$\lambda$  Lagrangian Multiplier

called **L2 regularization**, or RIDGE penalizes large  $\beta$



### L1 and L2 regularization

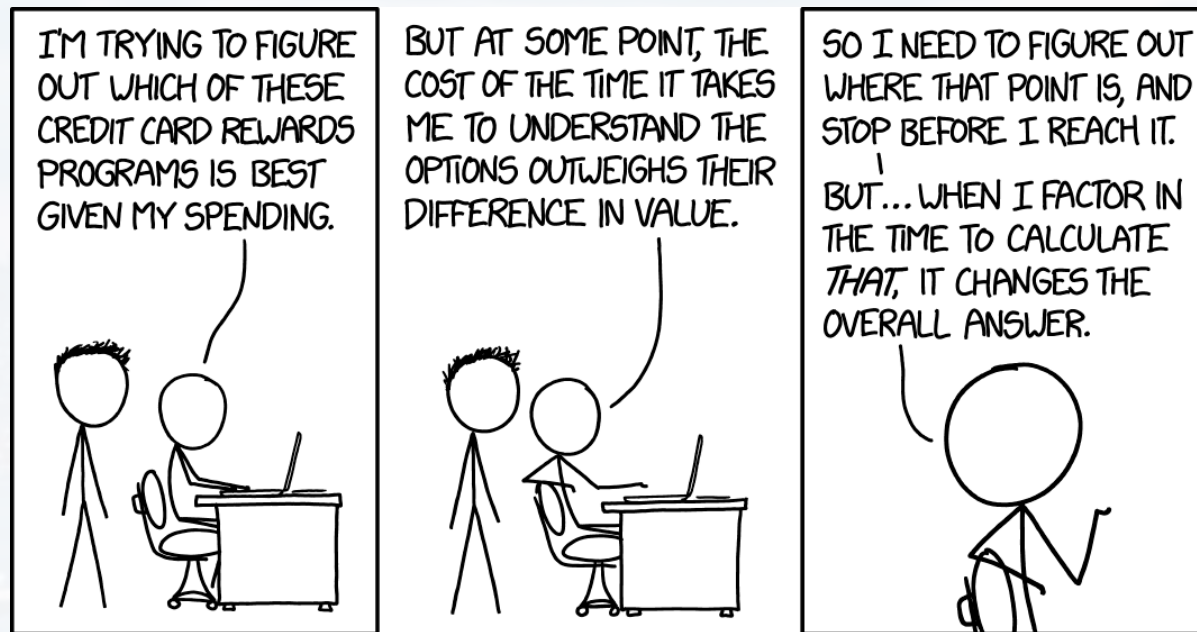
### L1 and L2



$$x(t=2) = x(t=1) - \varepsilon \frac{d[y + \lambda_1 \|x\|^1 + \lambda_2 \|x\|^2]}{dx} \Big|_{x(t=1)}$$

- gradient descent does not stop if values for  $x$  are too large and prefers sparsity
- note: the derivative of  $\|x\|^1$  returns the sign (i. e. direction)
- usually  $\lambda \ll \|x\|^n$
- will be important for ANNs later





## Outline

- The Problem

- **Gradient Descent**

- Vanilla
- Learning Rate Schedule
- Momentum
- L1 and L2
- More Finetuning



Vanilla Gradient Descent → **S**tochastic **G**radient **D**escent

More Fine Tuning

Learning Rate Schedule, L1, L2

Momentum

$$\epsilon \rightarrow \frac{\epsilon}{\sqrt{\text{grad}(y)_x}}$$

adaptive gradient, aka **AdaGrad**

different scaling for all different directions

Adding a decay factor to the sum of gradient squared (similar to momentum),  
aka **Root Mean Square Propagation RMSProp**


all combined:  
**Adaptive Moment Estimation**  
aka **Adam**



Lili Jiang

[TowardsDataScience](#)

More Fine Tuning



gradient descent (cyan),  
momentum (magenta),  
AdaGrad (white),  
RMSProp (green),  
Adam (blue)






Lili Jiang

[TowardsDataScience](#)

More Fine Tuning



gradient descent (cyan),  
momentum (magenta),  
AdaGrad (white),  
RMSProp (green),  
Adam (blue)

Thank you very much for your attention!

