

## Lecture 12:

# Long Short-Term Memory Networks (LSTMs) – Part II



Markus Hohle

University California, Berkeley

**Machine Learning Algorithms**

MSSE 277B, 3 Units

Fall 2024



<https://www.analyticsvidhya.com>



## Outline

- LSTM for Classification
- Bidirectional LSTMs
- Stacked LSTMs
- LSTM + CNN

<https://www.analyticsvidhya.com>



## Outline

- LSTM for Classification
- Bidirectional LSTMs
- Stacked LSTMs
- LSTM + CNN





minimal model:

```
[N_samples, LengthSeq, N_features] = X.shape  
[N_samples, N classes] = Y.shape
```

Y is one-hot encoded

```
model = Sequential()
```

```
model.add(LSTM(n_neurons, activation = 'tanh',\  
               input_shape = (LengthSeq, N_features)))
```


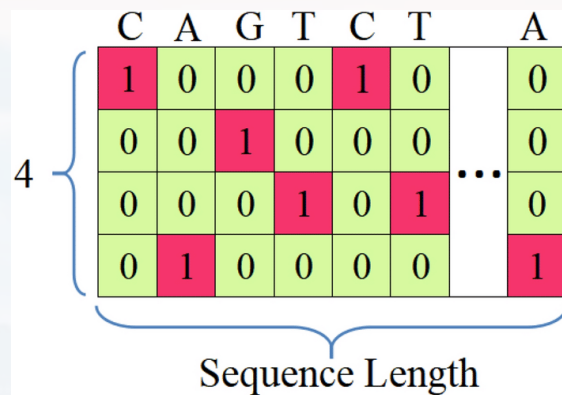
```
model.add(Dense(N classes, activation = 'softmax'))
```

```
opt = optimizers.Adam()  
model.compile(loss = 'categorical_crossentropy', optimizer = opt,\  
              metrics = ['accuracy'])
```

```
model.summary()
```





barcode example





for computational reasons:

- **three** classes
- **1k** samples total
- sequences cut to length **500**







barcode example



for computational reasons:

- **three** classes
- **1k** samples total
- sequences cut to length **500**



<https://www.analyticsvidhya.com>



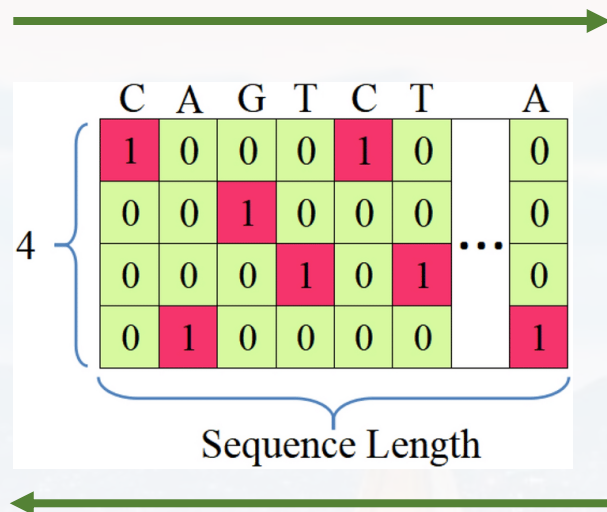
## Outline

- LSTM for Classification
- **Bidirectional LSTMs**
- Stacked LSTMs
- LSTM + CNN





sometimes, sequences can be read from two directions :



```
from keras.layers import Bidirectional
```

```
model = Sequential()
model.add(Bidirectional(LSTM(n_neurons, activation = 'tanh'),\
                           input_shape = (dt_past, n_features)))
model.add(Dense(dt_futu))
```

```
opt = optimizers.Adam()
model.compile(loss = 'mean_squared_error', optimizer = opt)
```



```
model.summary()
```







```
model = Sequential()
model.add(Bidirectional(LSTM(n_neurons, activation = 'tanh'), input_shape = (dt_past, n_features)))
model.add(Dense(dt_futu))
opt = optimizers.Adam()
model.compile(loss = 'mean_squared_error', optimizer = opt)
model.summary()
```

vanilla



bidirectional



| Layer (type)                        | Output Shape | Param # |
|-------------------------------------|--------------|---------|
| bidirectional (Bidirectional)       | (None, 800)  | 1286400 |
| dense_1 (Dense)                     | (None, 8)    | 6408    |
| Total params: 1292808 (4.93 MB)     |              |         |
| Trainable params: 1292808 (4.93 MB) |              |         |
| Non-trainable params: 0 (0.00 Byte) |              |         |

<https://www.analyticsvidhya.com>

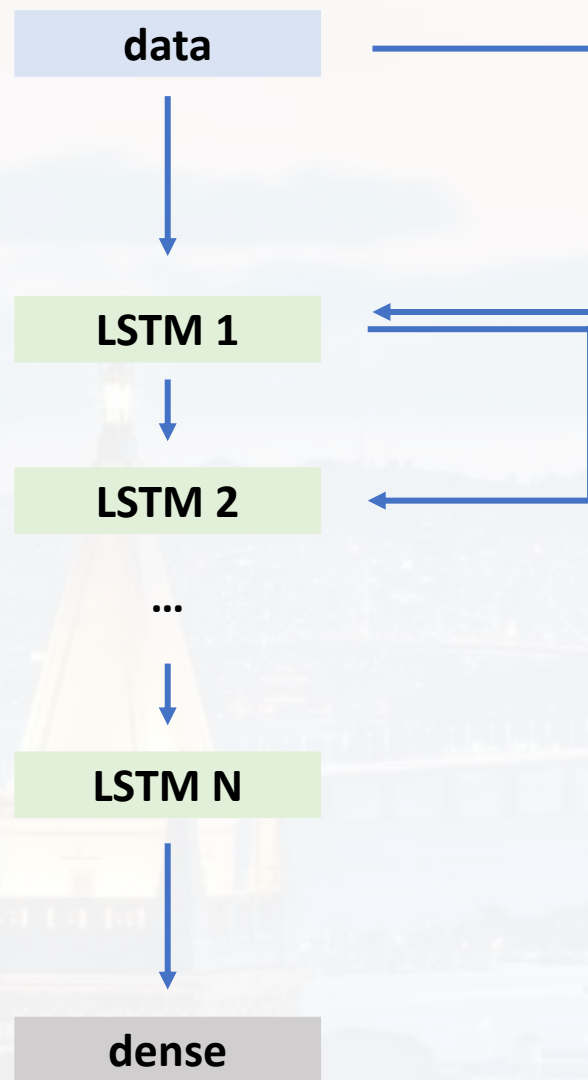


## Outline

- LSTM for Classification
- Bidirectional LSTMs
- **Stacked LSTMs**
- LSTM + CNN



idea:



shape: ( $\text{len}(y(t)) - dt\_past - dt\_futu + 1$ )  
x  $dt\_past$  x features

shape: ( $\text{len}(y(t)) - dt\_past - dt\_futu + 1$ )  
x  $dt\_past$  x hidden state

return\_sequences = True





```
model = Sequential()

model.add(LSTM(n_neurons, activation = 'tanh',\
               return_sequences = True, input_shape = (dt_past, n_features)))

model.add(LSTM(2*n_neurons, activation = 'relu',\
               return_sequences = True))

model.add(LSTM(n_neurons, activation = 'relu'))

model.add(Dense(dt_futu))

opt = optimizers.Adam()
model.compile(loss = 'mean_squared_error', optimizer = opt)



model.summary()
```

all LSTMs, **except the last** stack needs  
return\_sequences = True



| Layer (type)    | Output Shape    | Param # |
|-----------------|-----------------|---------|
| lstm_2 (LSTM)   | (None, 20, 400) | 643200  |
| lstm_3 (LSTM)   | (None, 20, 800) | 3843200 |
| lstm_4 (LSTM)   | (None, 400)     | 1921600 |
| dense_2 (Dense) | (None, 8)       | 3208    |

=====  
Total params: 6411208 (24.46 MB)  
Trainable params: 6411208 (24.46 MB)  
Non-trainable params: 0 (0.00 Byte)



<https://www.analyticsvidhya.com>



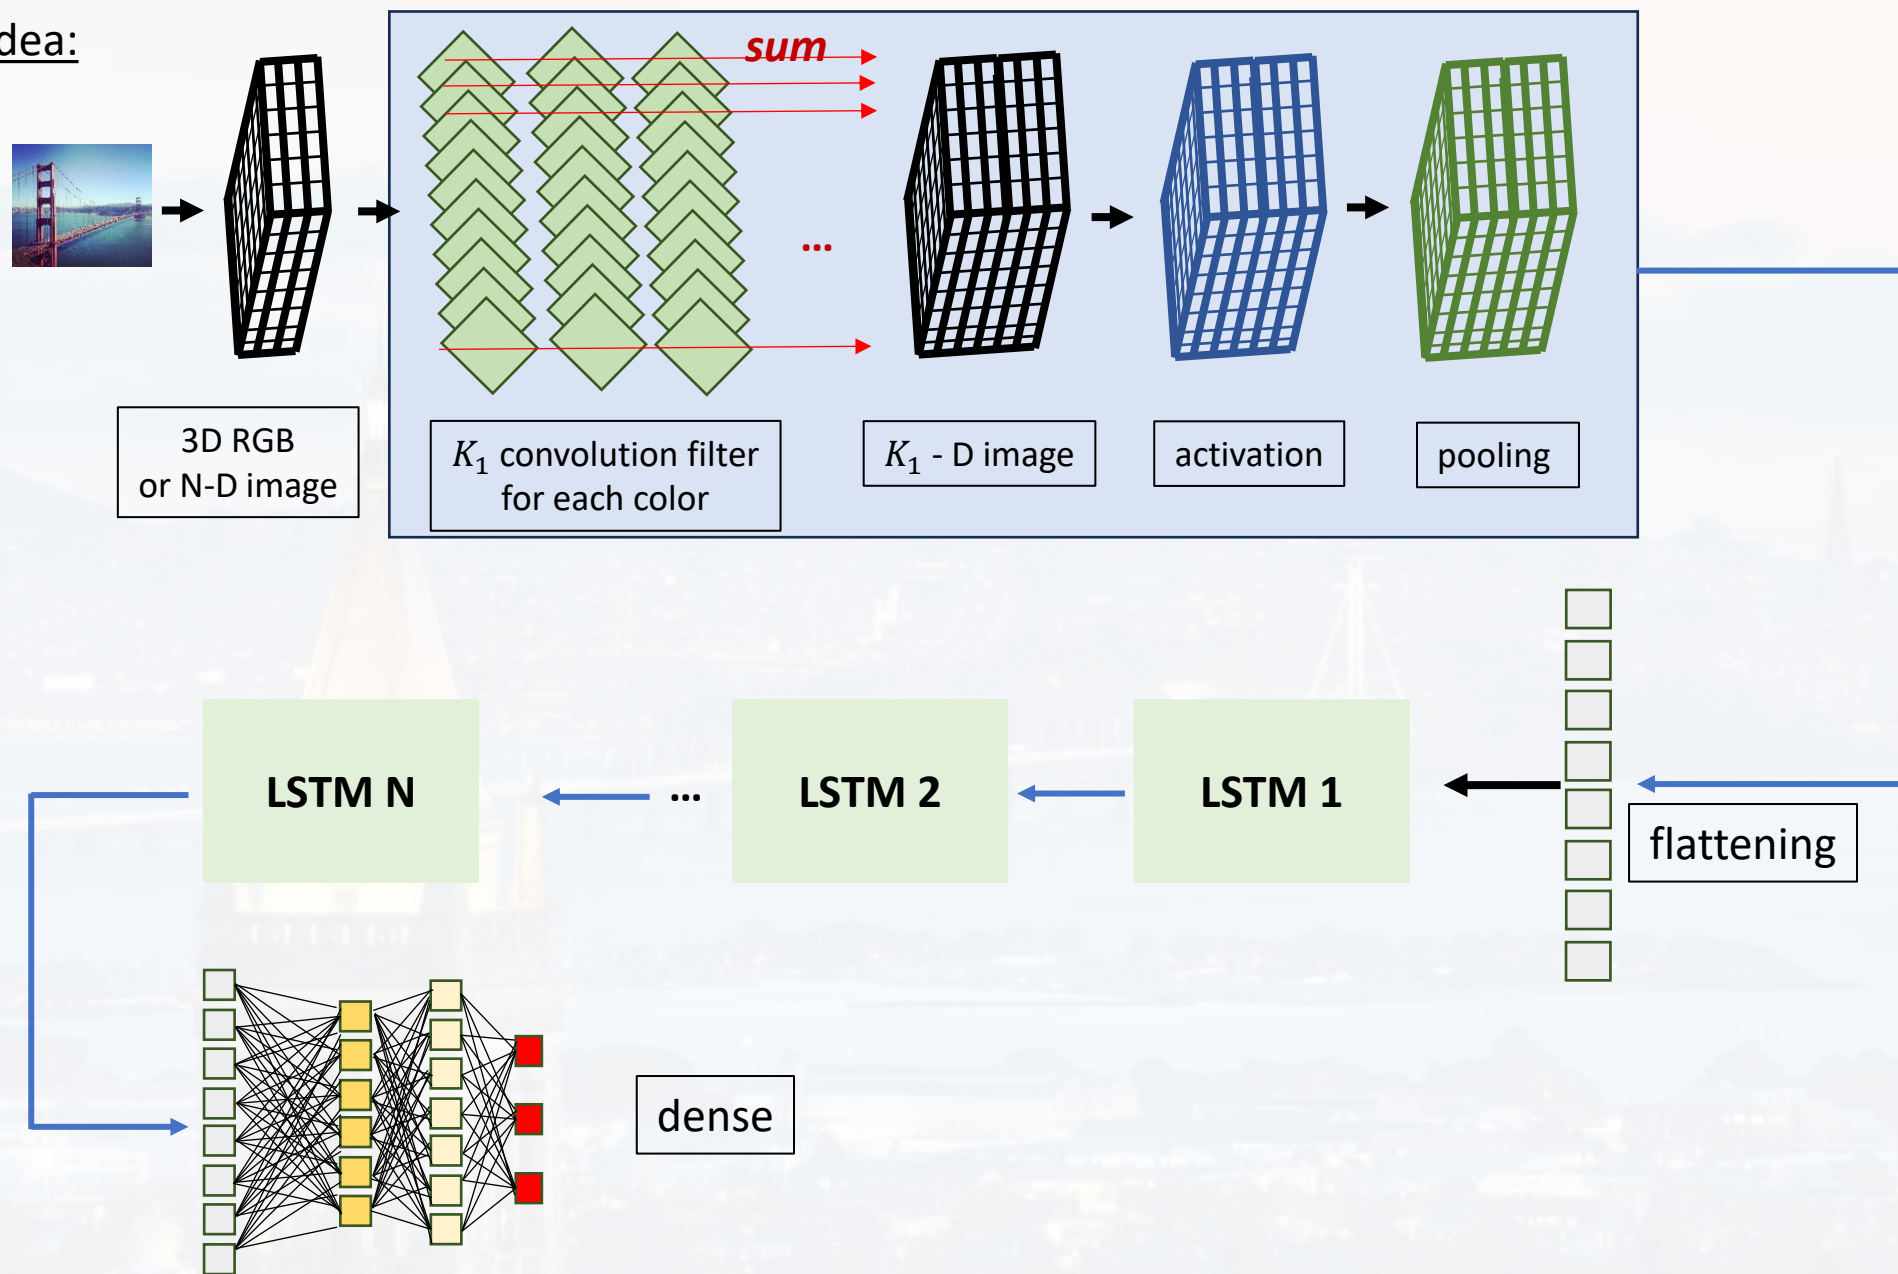
## Outline

- LSTM for Classification
- Bidirectional LSTMs
- Stacked LSTMs
- **LSTM + CNN**





idea:





idea:

input expected by CNN (images):

$(N\_images, N\_x, N\_y, N\_color)$



input expected by CNN (videos):

$(N\_videos, N\_t, N\_x, N\_y, N\_color)$

input expected by LSTM (sequences):


$(N\_sequences, N\_t, N\_feature)$

None







$(N\_images, N\_x, N\_y, N\_color)$



sample 1

sample 2


sample n









$(N\_images, N\_x, N\_y, N\_color)$



sample 1

sample 2


sample n



for **one** timepoint  $t$





$(N\_images, N\_x, N\_y, N\_color)$



sample 1

sample 2



sample n



for **one** timepoint  $t$



$(N\_images, N\_x, N\_y, N\_color)$



for **one** timepoint  $t$

regression: **one** sample of  $N\_features$  and  $dt\_past$

```
X = X.reshape((X.shape[0], N_samples, dt_past, n_features))
```





```
X = X.reshape((X.shape[0], N_samples, dt_past, n_features))
```

```
model = Sequential()  
model.add(TimeDistributed(Conv1D(filters = 64, kernel_size = 3,\br/>                                activation = 'relu'),\br/>                                input_shape = (None, dt_past, n_features)))
```

1D filter along time coordinate

```
model.add(TimeDistributed(MaxPooling1D(pool_size = 2)))
```

```
model.add(TimeDistributed(Flatten()))
```

takes care of  
maintaining  
matrix orientation

```
model.add(LSTM(n_neurons, input_shape = (dt_past, n_features),\br/>              activation = 'tanh'))
```

```
model.add(Dense(dt_futu))
```

```
opt = optimizers.Adam()  
model.compile(loss = 'mean_squared_error', optimizer = opt)
```



```
model.summary()
```

actual input is (None, None, dt\_past, n\_features)



| Layer (type)                         | Output Shape         | Param # |
|--------------------------------------|----------------------|---------|
| =====                                |                      |         |
| time_distributed (TimeDistributed)   | (None, None, 18, 64) | 256     |
| time_distributed_1 (TimeDistributed) | (None, None, 9, 64)  | 0       |
| time_distributed_2 (TimeDistributed) | (None, None, 576)    | 0       |
| lstm_5 (LSTM)                        | (None, 400)          | 1563200 |
| dense_3 (Dense)                      | (None, 8)            | 3208    |
| =====                                |                      |         |
| Total params: 1566664 (5.98 MB)      |                      |         |
| Trainable params: 1566664 (5.98 MB)  |                      |         |
| Non-trainable params: 0 (0.00 Byte)  |                      |         |


actual input is (None, None, dt\_past, n\_features)









vanilla




bidirectional



stacked



LSTM + CNN





classification: N samples of N\_features and dt\_past = LengthSeq

```
[N_sample, LengthSeq, N_features] = X.shape
```

1D filter along time coordinate  
= LengthSeq

```
model = Sequential()  
model.add(Conv1D(filters = 64, kernel_size = 3, activation = 'relu',\  
                 input_shape = (LengthSeq, N_features)))
```

```
model.add(MaxPooling1D(pool_size = 2))
```

```
model.add(LSTM(n_neurons, activation = 'tanh'))
```

```
model.add(Dense(Nclass, activation = 'softmax'))
```

```
opt = optimizers.Adam()  
model.compile(loss = 'categorical_crossentropy', optimizer = opt,\  
              metrics = ['accuracy'])
```

```
model.summary()
```



classification: **N** samples of **N\_features** and **dt\_past** = **Length\_Seq**


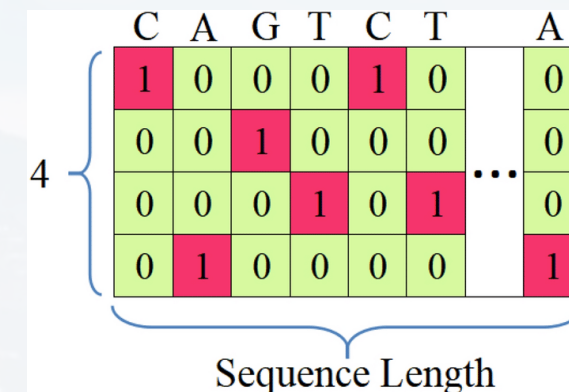
**[N\_sample, LengthSeq, N\_features]** = **X.shape**

| Layer (type)                        | Output Shape    | Param # |
|-------------------------------------|-----------------|---------|
| conv1d_8 (Conv1D)                   | (None, 498, 64) | 832     |
| max_pooling1d_7 (MaxPooling1D)      | (None, 249, 64) | 0       |
| lstm_7 (LSTM)                       | (None, 100)     | 66000   |
| dense_4 (Dense)                     | (None, 3)       | 303     |
| Total params: 67135 (262.25 KB)     |                 |         |
| Trainable params: 67135 (262.25 KB) |                 |         |
| Non-trainable params: 0 (0.00 Byte) |                 |         |

for computational reasons:

- **three** classes
- **1k** samples total
- sequences cut to length **500**





barcode example











LSTM



LSTM+CNN



Thank you for your attention!

