


Outline:

- What is LeNet?
- Forward Part *Spelled Out*
- Backward Part *Spelled Out*
- LeNet Keras TensorFlow



Outline:

- What is LeNet?
- Forward Part *Spelled Out*
- Backward Part *Spelled Out*
- LeNet Keras TensorFlow




- LeNet:
- Yann LeCun, Leon Bottou, Yoshua Bengio and Patrick Haffner, 1998
 - one of the 1st CNN that was able to categorize images
 - MNIST data set



- only **seven layers** in total
- different versions
- modern CNNs (google, ResNet etc have **100 or more** layers)



LeNet: - Yann LeCun, Leon Bottou, Yoshua Bengio and Patrick Haffner, 1998



- ***tanh*** as activation function
- **three different** filter
- from Conv1 → Conv2 not all the **channels are combined**
- **average** pool




We know already how the forward part works:

```
for c in range(numChans):# loop over channels
    for y in range(yOutput):# loop over y axis of output
        for x in range(xOutput):# loop over x axis of output

            # finding corners of the current "slice"
            y_start = y*stride
            y_end    = y*stride + yK
            x_start = x*stride
            x_end    = x*stride + xK

            #selecting the current part of the image
            current_slice = imagePadded[x_start:x_end,\
                                         y_start:y_end, c]

            #the actual convolution part
            s              = np.multiply(current_slice, K)
            output[x,y,c] = np.sum(s)
```



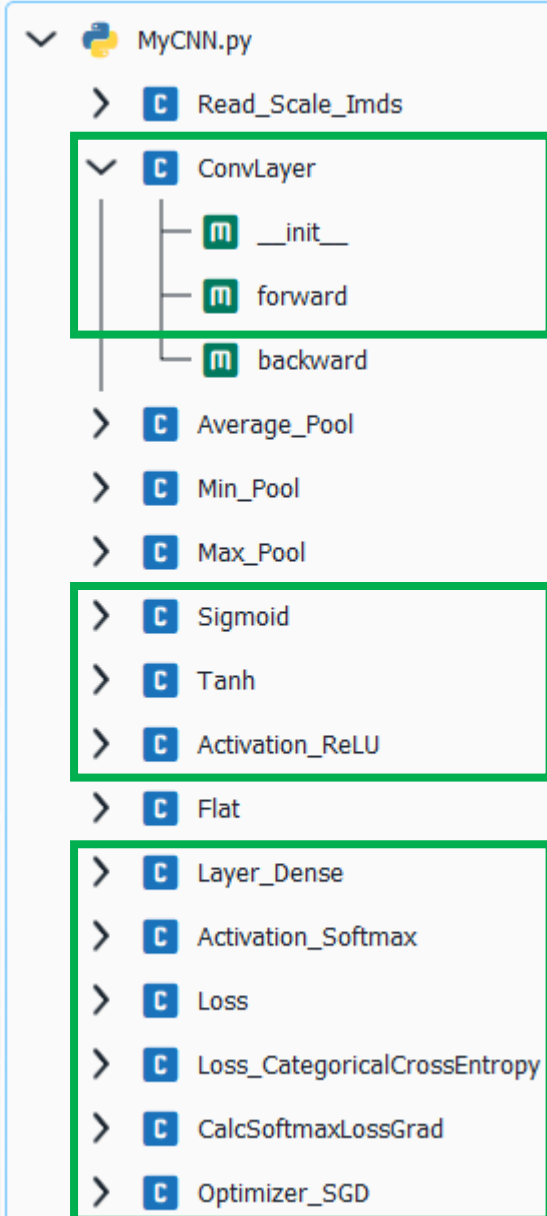
Outline:

- What is LeNet?
- **Forward Part *Spelled Out***
- Backward Part *Spelled Out*
- LeNet Keras TensorFlow



see MyCNN.py

we know these
parts already



```
import matplotlib.pyplot as plt
from MyCNN import *
```

```
read_and_scale = Read_Scale_Imds(5, [250, 250])
[I, _] = read_and_scale.Read_Scale()
```



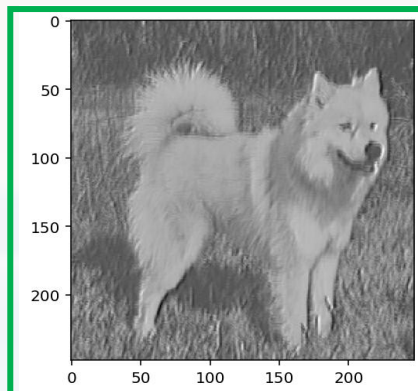
```
Conv1 = ConvLayer(3, 3, 5)
Conv2 = ConvLayer(5, 5, 4)
Conv3 = ConvLayer(2, 2, 4)
```

initializing convolution layer
(kernel size x kernel size,
number of kernels)

```
Conv1.forward(I, 0, 1)
Conv2.forward(Conv1.output, 2, 4)
Conv3.forward(Conv2.output, 0, 2)
```

passing image through
convolution layer
(padding, stride length)

```
plt.imshow(Conv1.output[:, :, 0, 4], cmap = 'gray')
plt.show()
plt.imshow(Conv2.output[:, :, 0, 4], cmap = 'gray')
plt.show()
plt.imshow(Conv3.output[:, :, 0, 4], cmap = 'gray')
plt.show()
```



```
import matplotlib.pyplot as plt  
from MyCNN import *
```

```
read_and_scale = Read_Scale_Imds(5, [250,250])  
[I, _] = read_and_scale.Read_Scale()
```

```
Conv1 = ConvLayer(3,3,5)  
Conv2 = ConvLayer(5,5,4)  
Conv3 = ConvLayer(2,2,4)
```

```
Conv1.forward(I,0,1)  
Conv2.forward(Conv1.output,2,4)  
Conv3.forward(Conv2.output,0,2)
```

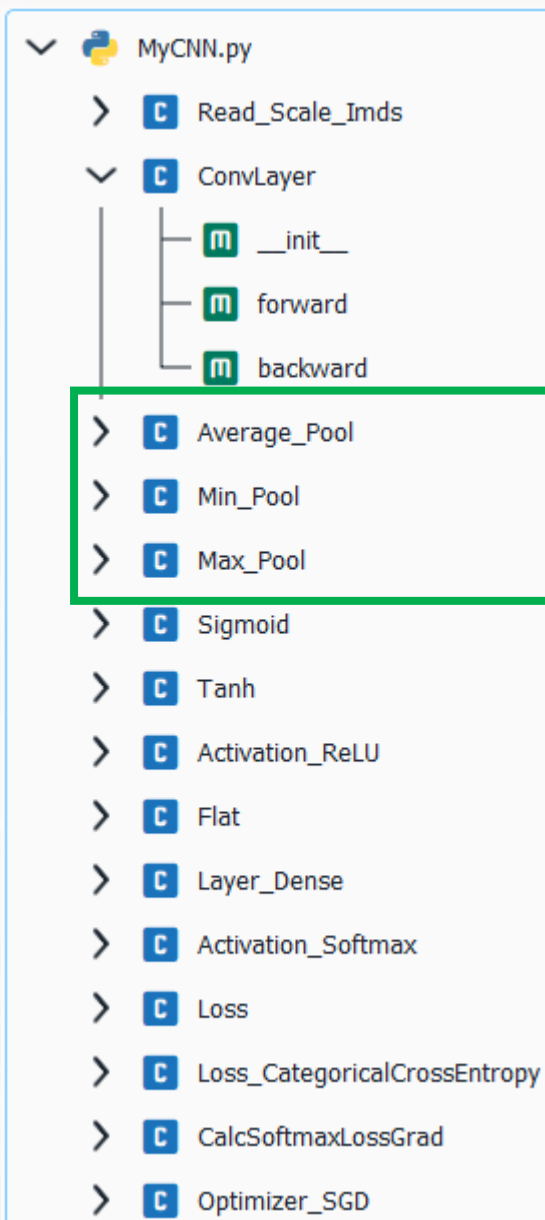
```
plt.imshow(Conv1.output[:, :, 0, 4], cmap = 'gray')  
plt.show()
```

```
plt.imshow(Conv2.output[:, :, 0, 4], cmap = 'gray')  
plt.show()
```

```
plt.imshow(Conv3.output[:, :, 0, 4], cmap = 'gray')  
plt.show()
```




see MyCNN.py




```
y_start = y * stride
y_end = y_start + yK
x_start = x * stride
x_end = x_start + xK
```

```
sx = slice(x_start, x_end)
sy = slice(y_start, y_end)
```

```
current_slice = currentIm pad[sx, sy, c]
slice_max = float(current_slice.max())
output[x, y, c, i] = slice_max
```


here: max pool



```
MP1 = Max_Pool()  
MP2 = Max_Pool()  
MP3 = Max_Pool()
```

initializing ax pool layer

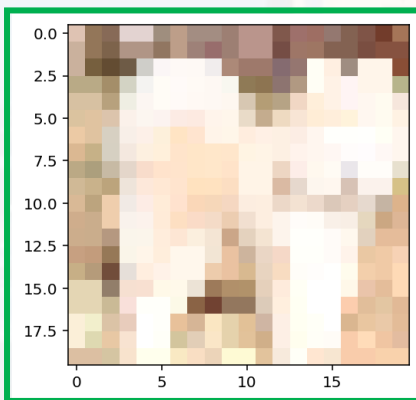
```
MP1.forward(I, 2, 3)  
MP2.forward(MP1.output, 3, 3)  
MP3.forward(MP2.output, 2, 3)
```




```
plt.imshow(MP1.output[:, :, :, 4].astype(int))  
plt.show()
```

```
plt.imshow(MP2.output[:, :, :, 4].astype(int))  
plt.show()
```

```
plt.imshow(MP3.output[:, :, :, 4].astype(int))  
plt.show()
```



passing image through max
pool layer (stride length,
kernel size)



Outline:

- What is LeNet?
- Forward Part *Spelled Out*
- Backward Part *Spelled Out*
- LeNet Keras TensorFlow



last time for dense layer:

```
class Layer_Dense:
```

```
    def __init__(self, n_inputs, n_neurons):  
        self.weights = np.random.rand(n_inputs, n_neurons)  
        self.biases   = np.zeros((1, n_neurons))
```

```
    def forward(self, inputs):  
  
        self.output = np.dot(inputs, self.weights) + self.biases  
        self.inputs  = inputs
```

```
    def backward(self, dvalues):  
  
        self.dweights = np.dot(self.inputs.T, dvalues)  
        self.dinputs  = np.dot(dvalues, self.weights.T)  
        self.dbiases  = np.sum(dvalues, axis = 0, keepdims = True)
```




outer derivative




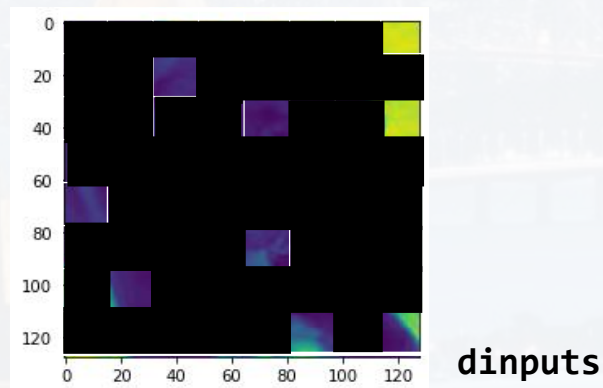
max pool

for **backpropagation**, we need to **up-sample** the images!

forward:



backward:



We need **to track, where the max came from** for each patch: creating a **mask** in the forward part
For those pixel: $dinputs = dvalues$

max pool

for **backpropagation**, we need to **up-sample** the images!

Input I stride = 3; xK = yK = 5

| | 0 | 1 | 2 | 3 | 4 | 5 | 6 |
|---|---|-----|-----|-----|-----|-----|-----|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 1 | 0 | 538 | 538 | 538 | 538 | 541 | 541 |
| 2 | 0 | 538 | 538 | 538 | 538 | 541 | 541 |
| 3 | 0 | 535 | 535 | 538 | 538 | 541 | 541 |
| 4 | 0 | 535 | 535 | 538 | 538 | 541 | 541 |
| 5 | 0 | 535 | 535 | 538 | 538 | 538 | 538 |
| 6 | 0 | 535 | 535 | 538 | 538 | 538 | 538 |
| 7 | 0 | 535 | 535 | 538 | 538 | 541 | 541 |
| 8 | 0 | 535 | 535 | 538 | 538 | 541 | 541 |
| 9 | 0 | 535 | 535 | 538 | 538 | 538 | 544 |



max pool

output

| | 0 | 1 | 2 | 3 | 4 | 5 | 6 |
|---|-----|-----|-----|-----|-----|-----|-----|
| 0 | 538 | 544 | 545 | 551 | 543 | 357 | 364 |
| 1 | 538 | 544 | 547 | 550 | 544 | 369 | 361 |
| 2 | 538 | 544 | 548 | 554 | 554 | 369 | 361 |
| 3 | 544 | 547 | 548 | 554 | 554 | 361 | 358 |
| 4 | 544 | 544 | 499 | 364 | 362 | 368 | 375 |
| 5 | 424 | 363 | 355 | 355 | 350 | 363 | 375 |
| 6 | 352 | 354 | 358 | 358 | 356 | 356 | 357 |
| 7 | 352 | 353 | 354 | 359 | 368 | 368 | 350 |
| 8 | 353 | 353 | 362 | 362 | 360 | 359 | 356 |
| 9 | 357 | 356 | 367 | 367 | 366 | 365 | 356 |

mask

and so on...

| | 0 | 1 | 2 | 3 | 4 | 5 | 6 |
|---|---|-----|-----|-----|-----|---|-----|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 1 | 0 | 538 | 538 | 538 | 538 | 0 | 0 |
| 2 | 0 | 538 | 538 | 538 | 538 | 0 | 0 |
| 3 | 0 | 0 | 0 | 538 | 538 | 0 | 0 |
| 4 | 0 | 0 | 0 | 538 | 538 | 0 | 0 |
| 5 | 0 | 0 | 0 | 538 | 538 | 0 | 0 |
| 6 | 0 | 0 | 0 | 538 | 538 | 0 | 0 |
| 7 | 0 | 0 | 0 | 538 | 538 | 0 | 0 |
| 8 | 0 | 0 | 0 | 538 | 538 | 0 | 0 |
| 9 | 0 | 0 | 0 | 538 | 538 | 0 | 544 |



max pool

for **backpropagation**, we need to **up-sample** the images!

dinputs and mask have to look the same (but different values)!

```
Conv1.forward(I,2,1)
MP1.forward(Conv1.output,3,5)
MP1.backward(MP1.output)
```





```
plt.imshow(Conv1.output[:, :, 0, 4], cmap = 'gray')
plt.show()
```

```
plt.imshow(MP1.output[:, :, 0, 4], cmap = 'gray')
plt.show()
```

```
plt.imshow(MP1.dinputs[:, :, 0, 4], cmap = 'gray_r')
plt.show()
```

```
plt.imshow(MP1.mask[:, :, 0, 4], cmap = 'gray_r')
plt.show()
```

```
Diff = MP1.mask[:, :, 0, 4]/np.max(MP1.mask[:, :, 0, 4]) - \
      MP1.dinputs[:, :, 0, 4]/np.max(MP1.dinputs[:, :, 0, 4])
```





max pool

for **backpropagation**, we need to **up-sample** the images!

dinputs and mask have to look the same (but different values)!

```
Conv1.forward(I,2,1)
MP1.forward(Conv1.output,3,5)
MP1.backward(MP1.output)
```

```
plt.imshow(Conv1.output[:, :, 0, 4], cmap = 'gray')
plt.show()
```

```
plt.imshow(MP1.output[:, :, 0, 4], cmap = 'gray')
plt.show()
```

```
plt.imshow(MP1.dinputs[:, :, 0, 4], cmap = 'gray_r')
plt.show()
```





```
plt.imshow(MP1.mask[:, :, 0, 4], cmap = 'gray_r')
plt.show()
```

```
Diff = MP1.mask[:, :, 0, 4] /
      MP1.dinputs[:, :, 0,
```

In [20]: Diff

Out[20]:

```
array([[ 0.,  0.,  0., ...,  0.,  0., -0.],
       [ 0.,  0.,  0., ...,  0.,  0., -0.],
       [ 0.,  0.,  0., ...,  0.,  0., -0.],
       ...,
       [ 0.,  0.,  0., ...,  0.,  0., -0.],
       [ 0.,  0.,  0., ...,  0.,  0., -0.],
       [ 0.,  0.,  0., ...,  0., -0., -0.]])
```





convolution

for **backpropagation**, we need to **up-sample** the images!

same for convolution layer now:

backward **dinputs = dvalues * weights** (we don't need for first conv layer!)

forward:






image = inputs




$$\sum_i I_i w_i + b$$




convolved image

backward:




dinputs

weights



dvalues





convolution

for **backpropagation**, we need to **up-sample** the images!

same for convolution layer now:

backward $dinputs = dvalues * weights$ (we don't need for first conv layer!)

forward:







image = input

backward:



```
for i in range(numImds): # loop over number of images
    ... currentIm_pad = imagePadded[:, :, :, i] # Select ith padded image
    ... for k in range(NK): # loop over kernels (= #filters)
        ... for c in range(numChan): # loop over channels of incoming data
            ...
            ... if filt[c,k] != 1:
                ...
                ... for y in range(yd): # loop over axis of output
                    ... for x in range(xd): # loop over axis of output
                        ...
                        ... # finding corners of the current "slice" (~4 lines)
                        ... y_start = y*stride
                        ... y_end = y_start + yK
                        ... x_start = x*stride
                        ... x_end = x_start + xK
                        ...
                        ... sx = slice(x_start, x_end)
                        ... sy = slice(y_start, y_end)
                        ...
                        ... current_slice = currentIm_pad[sx, sy, c]
                        ...
                        ... dweights[:, :, k] += current_slice * dvalues[x, y, k, i]
                        ... dinputs[sx, sy, c, i] += weights[:, :, k] * dvalues[x, y, k, i]
                        ...
                        ... dbiases[0, k] += np.sum(np.sum(dvalues[:, :, k, i], axis=0), axis=0)
```



convolution



for **backpropagation**, we need to **up-sample** the images!

same for convolution layer now:

backward

$dweights = dvalues * inputs$

backward:



dvalues

```
.....
.....# finding corners of the current "slice" (~4 lines)
.....y_start = y*stride
.....y_end = y_start + yK
.....x_start = x*stride
.....x_end = x_start + xK
.....
.....sx = slice(x_start,x_end)
.....sy = slice(y_start,y_end)
.....
.....current_slice = currentIm_pad[sx,sy,c]
```

```
dweights[:, :, k] += current_slice * dvalues[x, y, k, i]
```

```
.....dinputs[sx,sy,c,i] += weights[:, :, k] * dvalues[x, y, k, i]
.....
.....
.....dbiases[0, k] += np.sum(np.sum(dvalues[:, :, k, i], axis=0), axis=0)
```



convolution

for **backpropagation**, we need to **up-sample** the images!

same for convolution layer now:

backward

dbiases = dvalues


backward:

```
.....
.....# finding corners of the current "slice" (~4 lines)
.....y_start = y*stride
.....y_end = y_start + yK
.....x_start = x*stride
.....x_end = x_start + xK
.....
.....sx = slice(x_start,x_end)
.....sy = slice(y_start,y_end)
.....
.....current_slice = currentIm_pad[sx,sy,c]
.....
.....dweights[:, :, k] += current_slice * dvalues[x,y,k,i]
.....dinputs[sx,sy,c,i] += weights[:, :, k] * dvalues[x,y,k,i]
.....
.....
.....dbiases[0,k] += np.sum(np.sum(dvalues[:, :, k,i], axis=0), axis=0)
```

```
dbiases[0,k] += np.sum(np.sum(dvalues[:, :, k,i], axis=0), axis=0)
```




We are done now
and can explore
our self-made
LeNet






```
#pip install keras
#pip install tensorflow
```

```
from keras.datasets import mnist
```



```
(train_x, train_y), (test_x, test_y) = mnist.load_data()
```



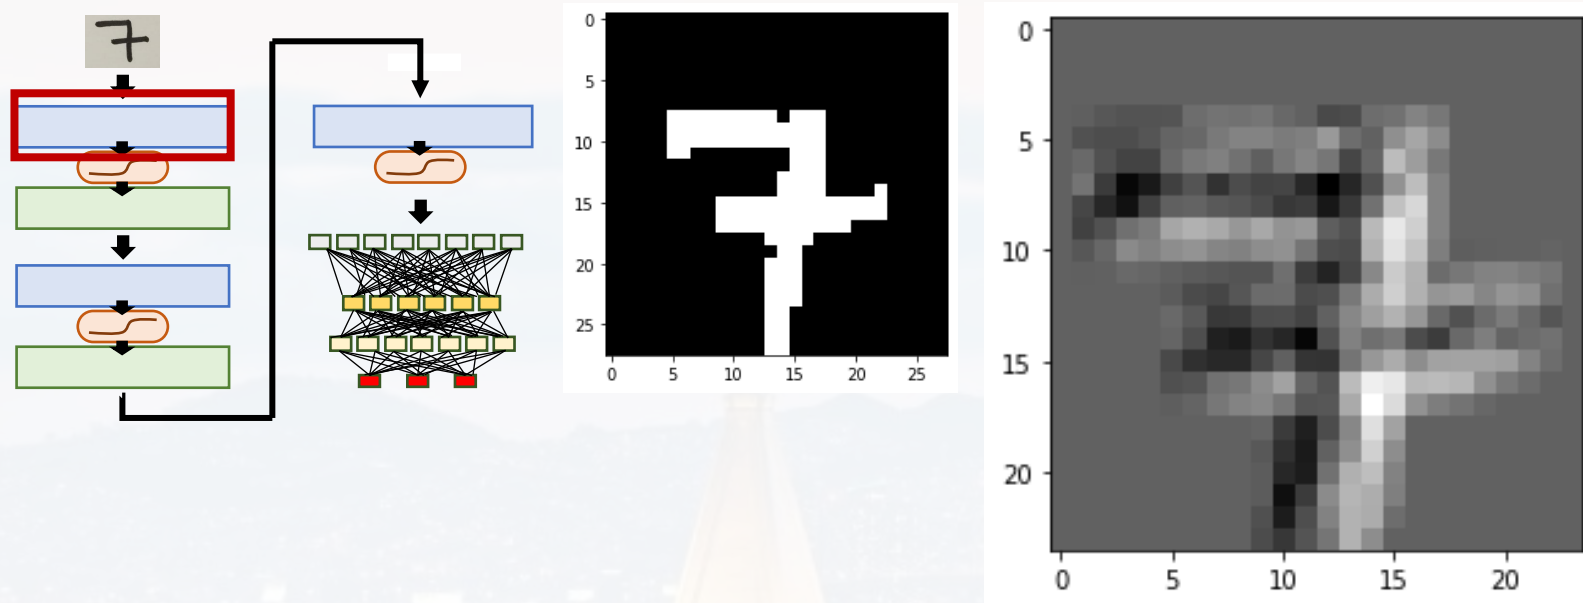
```
plt.imshow(train_x[1, :, :], cmap = 'gray')
```







exploring self-made LeNet:








exploring self-made LeNet:









exploring self-made LeNet:










exploring self-made LeNet:











exploring self-made LeNet:



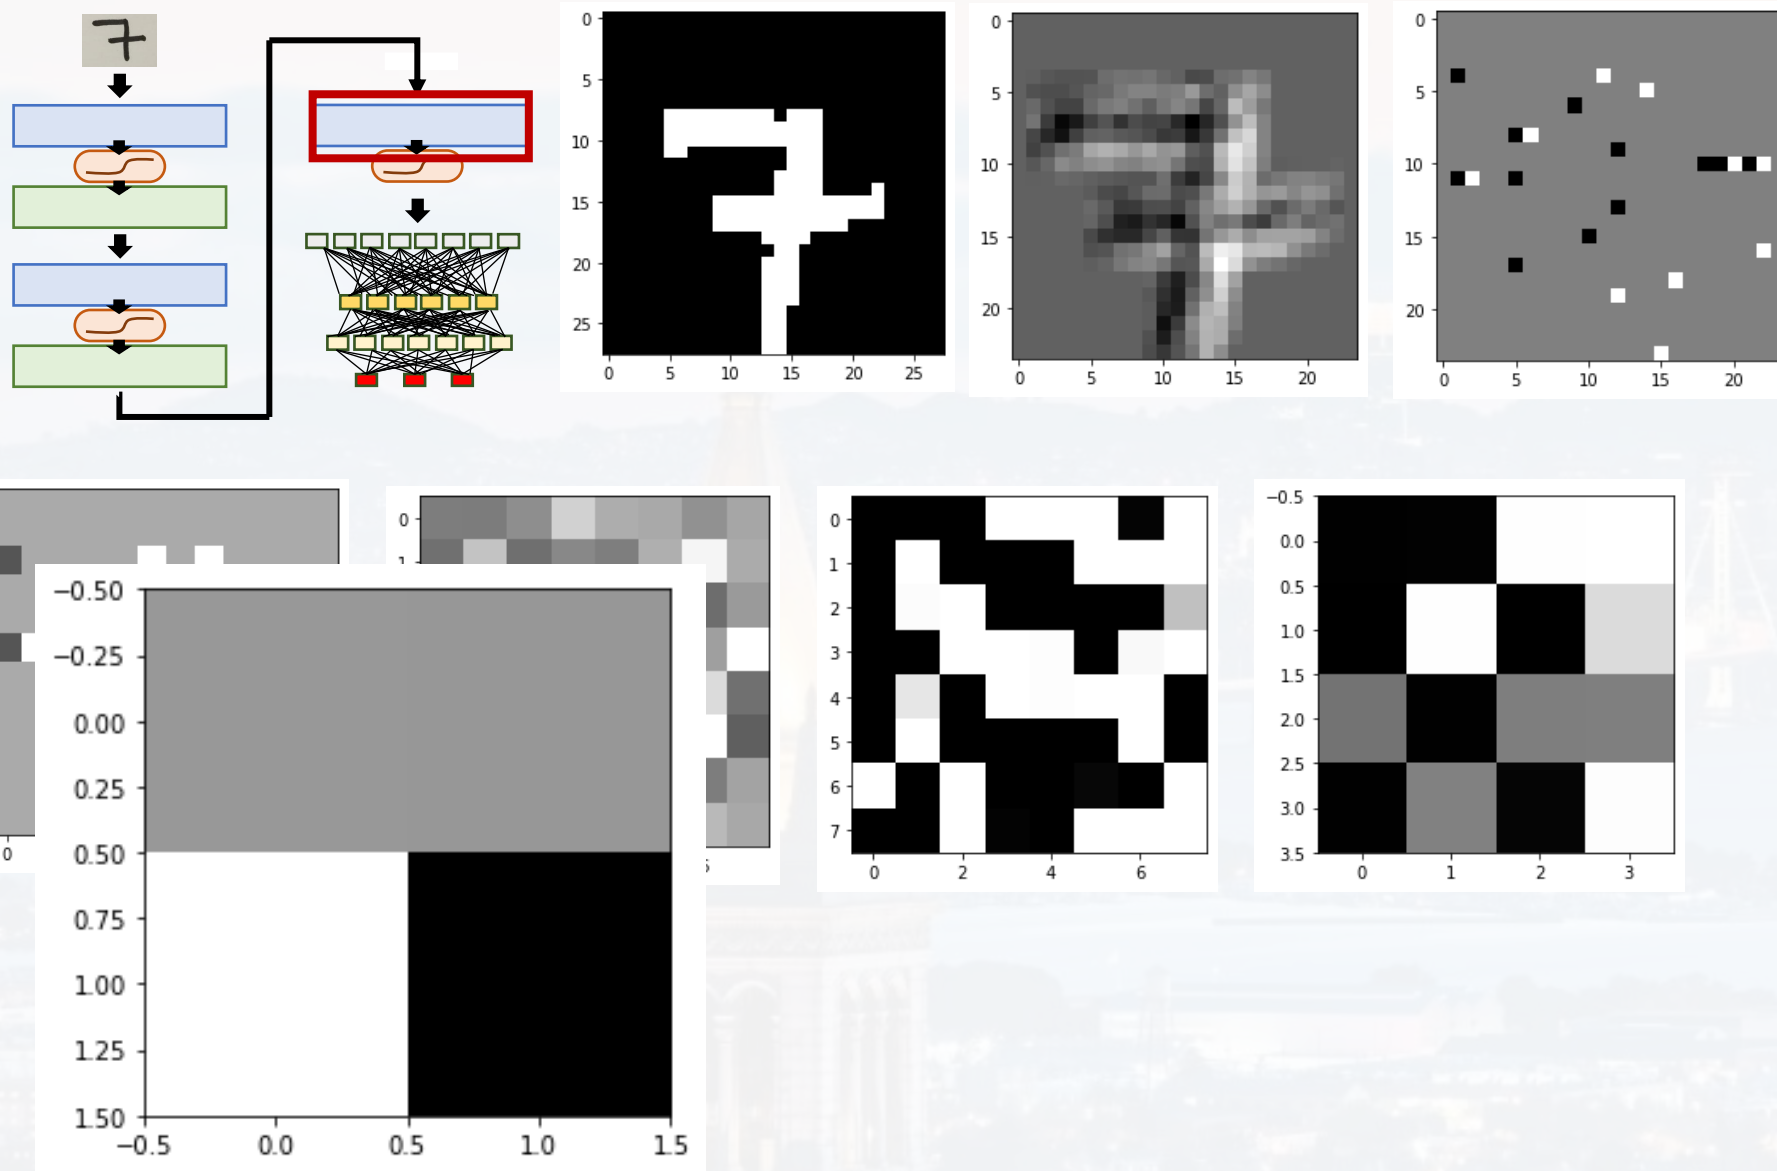
exploring self-made LeNet:












exploring self-made LeNet:



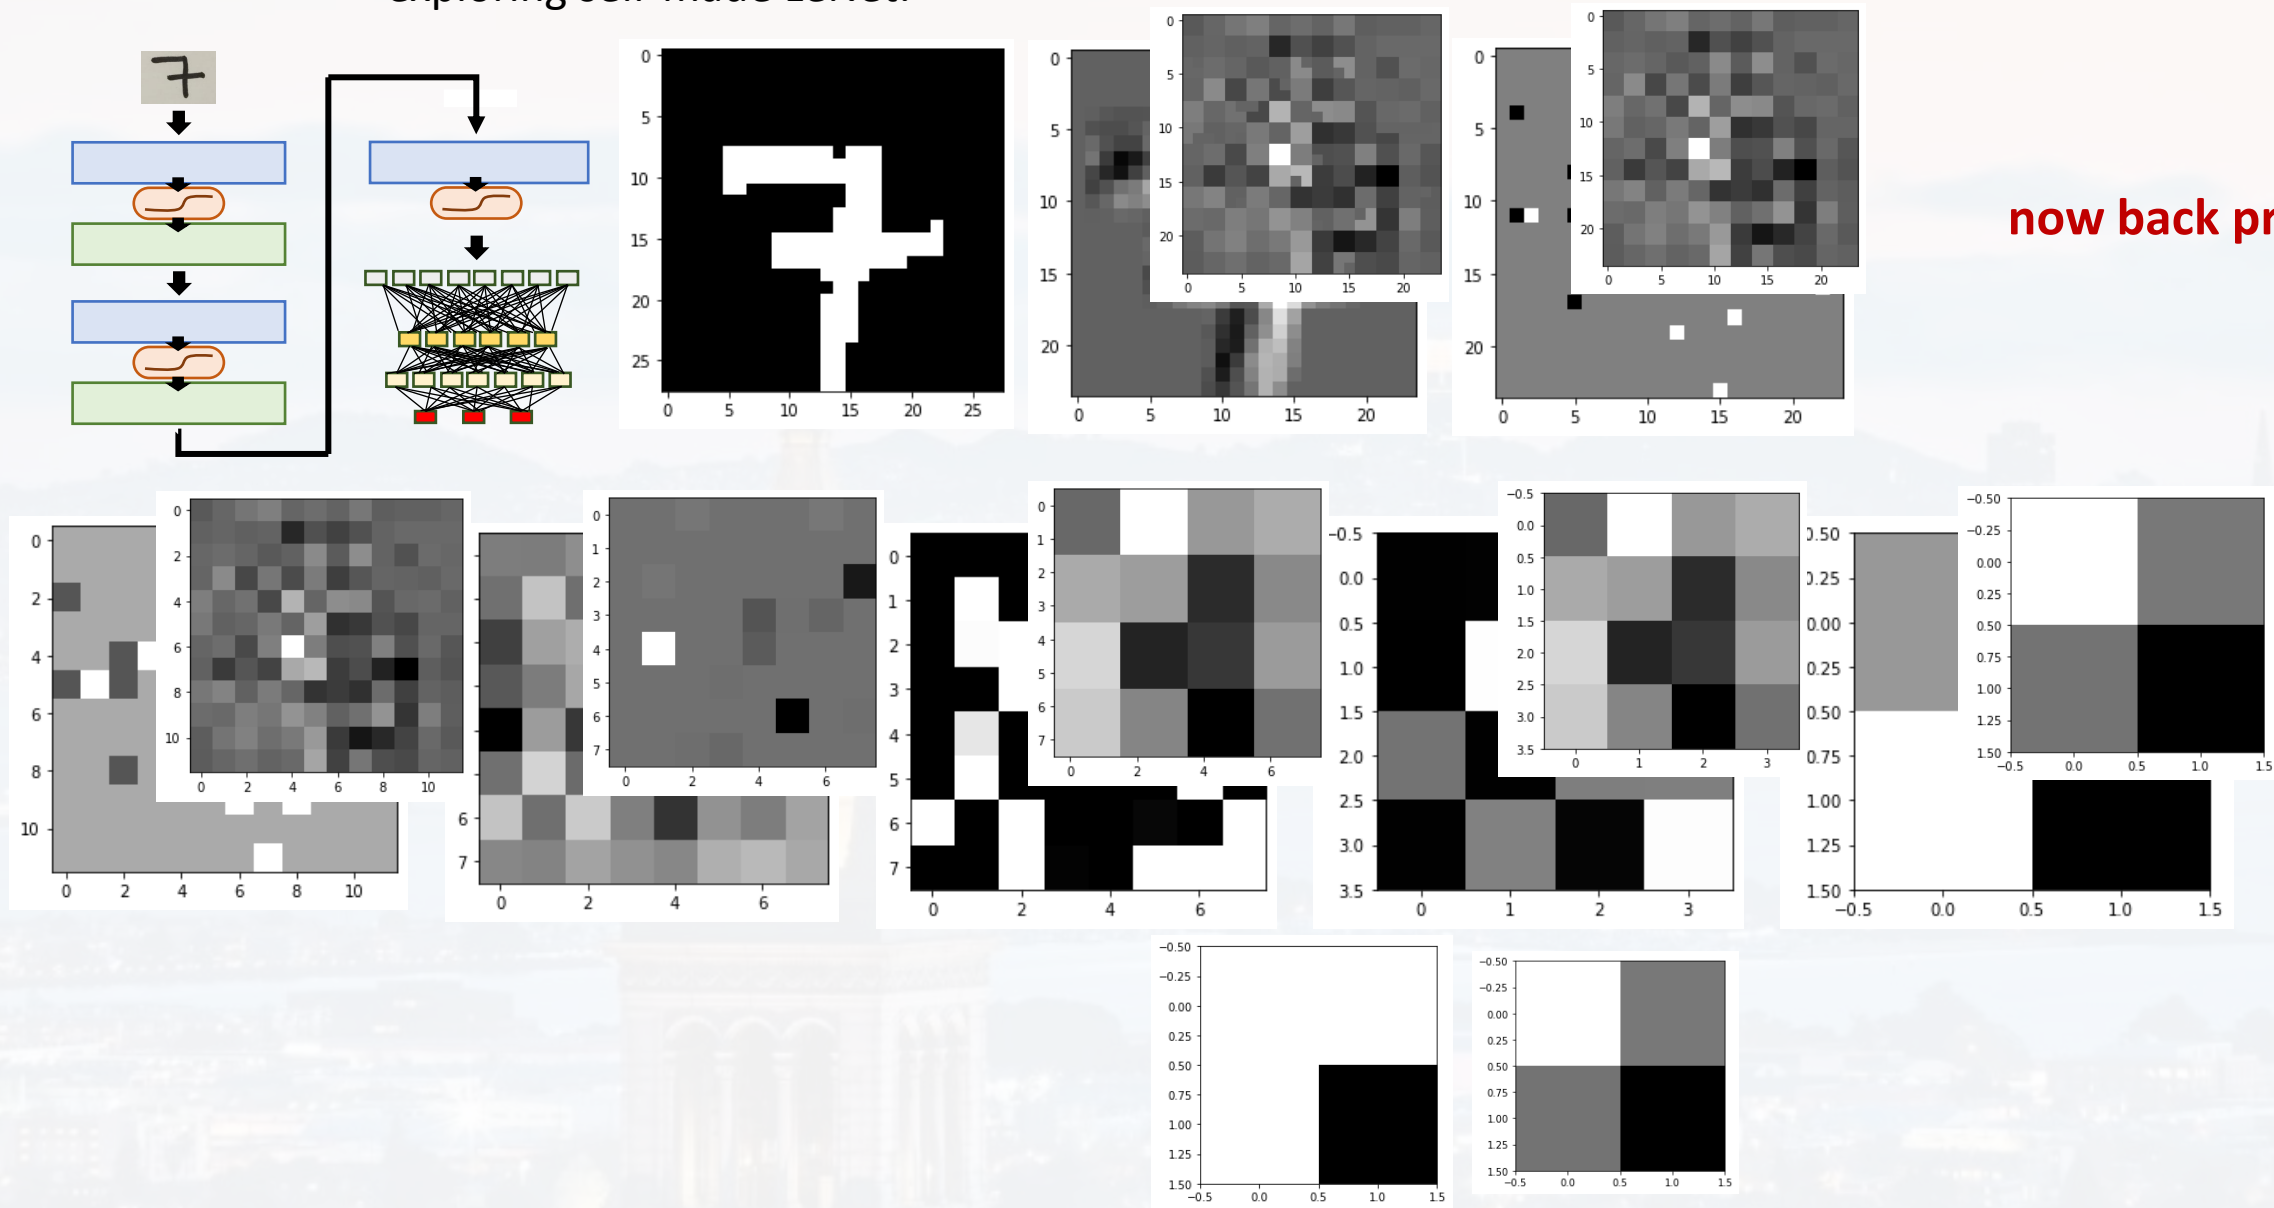
exploring self-made LeNet:



exploring self-made LeNet:

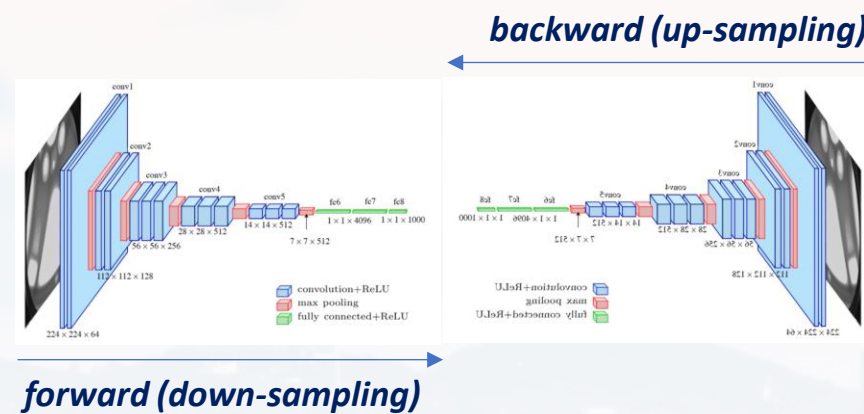



exploring self-made LeNet:

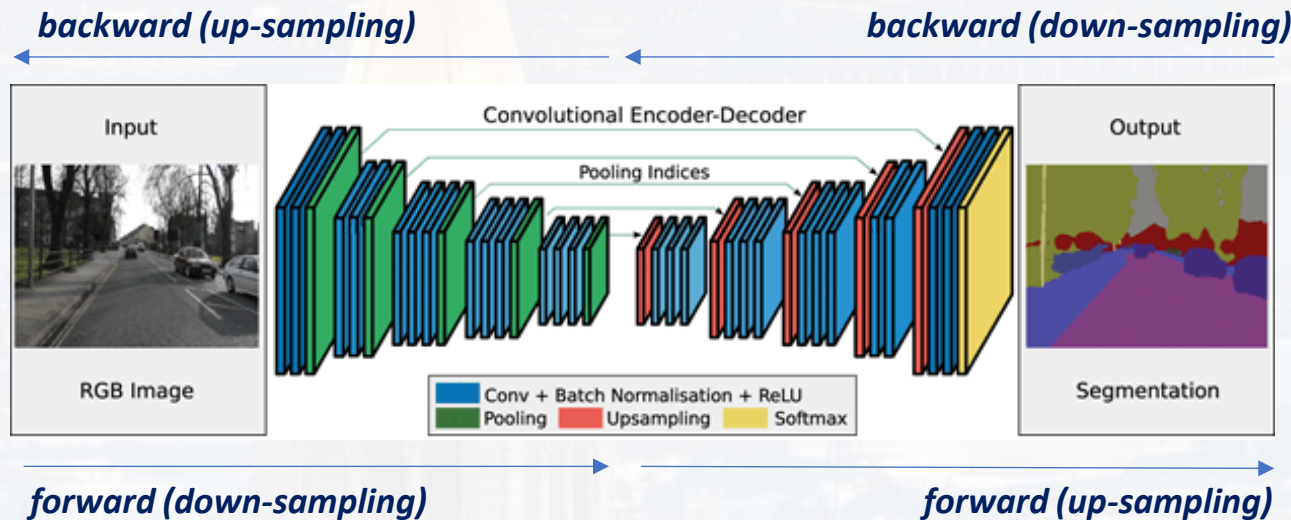


CNNs have an “hourglass” structure

classification:



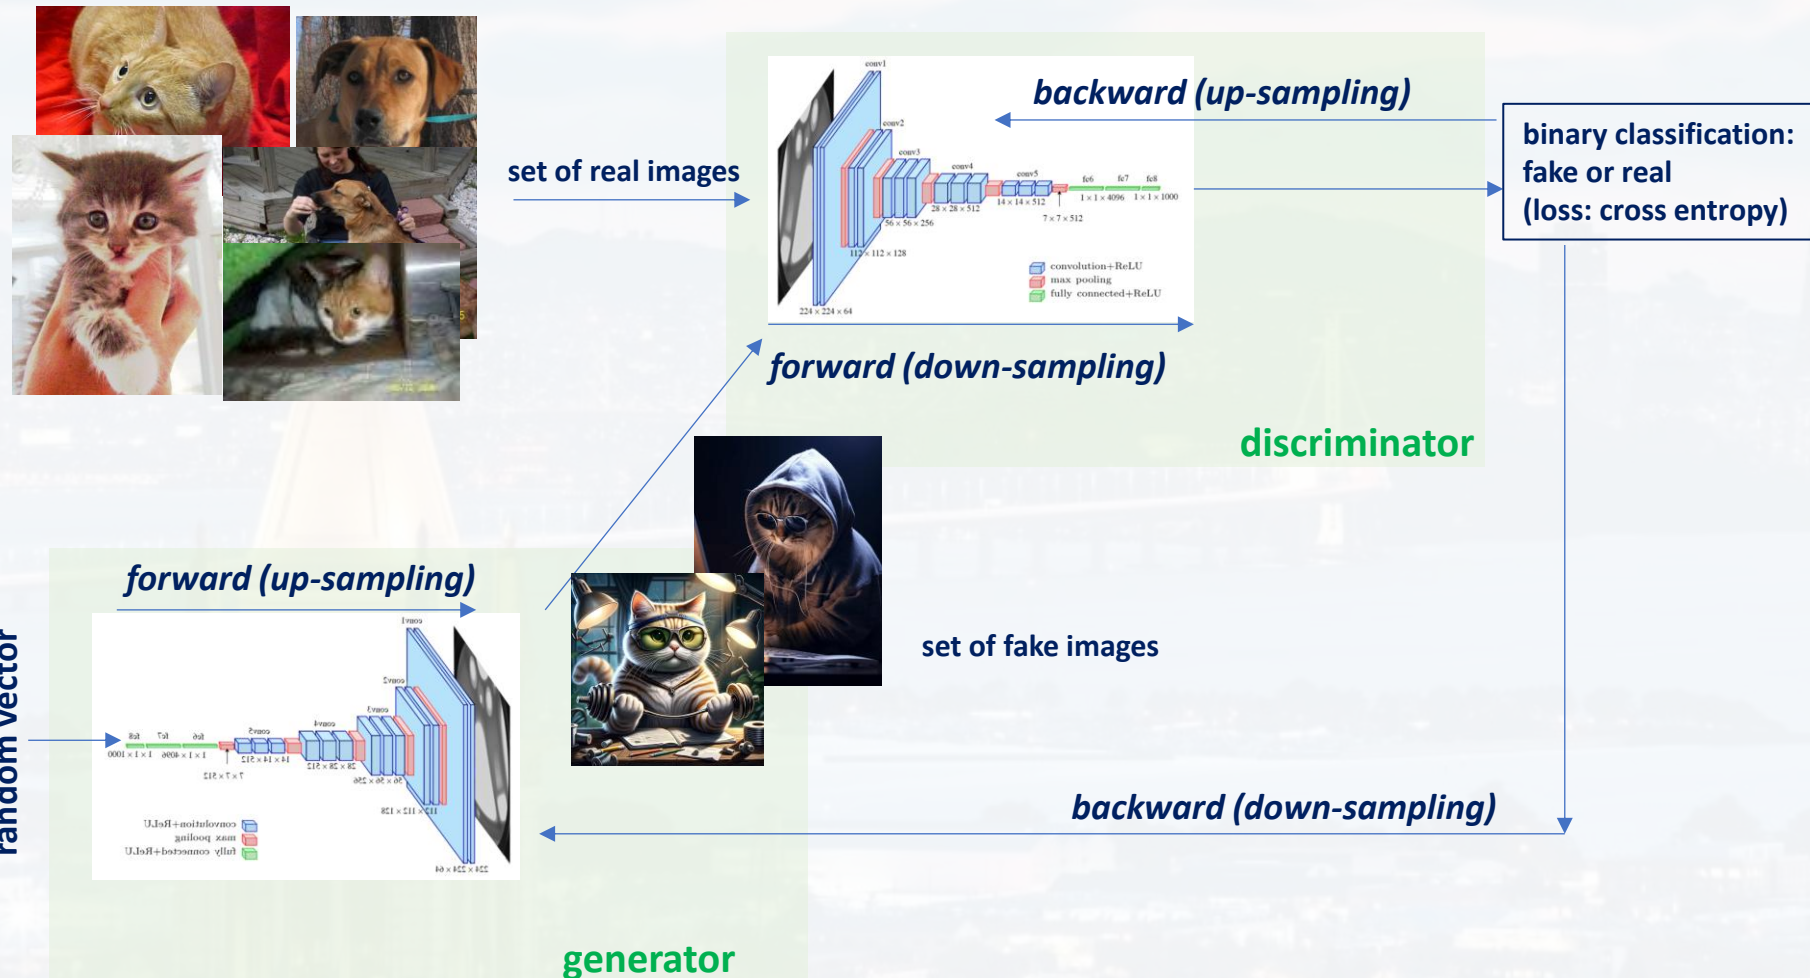
segmentation:






CNNs have an “hourglass” structure

Generative Adversarial Network (GAN):





Outline:

- What is LeNet?
- Forward Part *Spelled Out*
- Backward Part *Spelled Out*
- LeNet Keras TensorFlow



see LeNetTF.py and LeNetTF.ipynb

```
from tensorflow.keras.models import Sequential
from tensorflow.keras.losses import categorical_crossentropy
from tensorflow.keras.layers import Dense, Flatten, Conv2D, AveragePooling2D
```

```
class MyLeNet(Sequential):
```

```
    def __init__(self, input_shape, num_classes):
        super().__init__()
```

```
#building LeNet #####
```

```
    #Note padding: string, either "valid" or "same" (case-insensitive). "valid" means no padding. "same"
```

```
    #more info: https://keras.io/api/layers/convolution\_layers/convolution2d/
```

```
    self.add(Conv2D(6, kernel_size = (5, 5), strides = (1, 1), activation = 'tanh', input_shape = input_shape, padding = 'same'))
    self.add(AveragePooling2D(pool_size = (2, 2), strides = (2, 2), padding = 'valid'))
    self.add(Conv2D(16, kernel_size = (5, 5), strides = (1, 1), activation = 'tanh', padding = 'valid'))
    self.add(AveragePooling2D(pool_size = (2, 2), strides = (2, 2), padding = 'valid'))
    self.add(Conv2D(120, kernel_size = (5, 5), strides = (3, 3), activation = 'tanh', padding = 'valid'))
    self.add(Flatten())
    self.add(Dense(84, activation = 'tanh'))
    self.add(Dense(num_classes, activation = 'softmax'))
    #####
```

```
#building the optimizer #####
```

```
    lr_schedule = tf.keras.optimizers.schedules.ExponentialDecay(initial_learning_rate = 1e-2, decay_steps = 10000, decay_rate = 0.98)
    opt = tf.keras.optimizers.SGD(learning_rate = lr_schedule, momentum = 0.9)
```

```
    self.compile(optimizer = opt, loss = categorical_crossentropy, metrics = ['accuracy'])
```




see LeNetTF.py and LeNetTF.ipynb

Model: "my_le_net_1"

| Layer (type) | Output Shape | Param # |
|--|--------------------|---------|
| conv2d_3 (Conv2D) | (None, 28, 28, 6) | 156 |
| average_pooling2d_2 (AveragePooling2D) | (None, 14, 14, 6) | 0 |
| conv2d_4 (Conv2D) | (None, 10, 10, 16) | 2416 |
| average_pooling2d_3 (AveragePooling2D) | (None, 5, 5, 16) | 0 |
| conv2d_5 (Conv2D) | (None, 1, 1, 120) | 48120 |
| flatten_1 (Flatten) | (None, 120) | 0 |
| dense_2 (Dense) | (None, 84) | 10164 |
| dense_3 (Dense) | (None, 10) | 850 |

input = (N images, 28 x 28 pixel, 3 colors)
output = (N images, 28 x 28 pixel, 6 Conv filter)

6 Conv filter * shape (5, 5) = 150
plus 1 bias for each filter → total 156

Each image is represented by a vector of length 120

output layer: one neuron for each class

Total params: 61706 (241.04 KB)
Trainable params: 61706 (241.04 KB)
Non-trainable params: 0 (0.00 Byte)



see LeNetTF.py and LeNetTF.ipynb

```
running model...
Epoch 1/20
94/94 [=====] - 9s 92ms/step - loss: 1.1172 - accuracy: 0.7160 - val_loss: 0.4217 - val_accuracy: 0.8838
Epoch 2/20
94/94 [=====] - 8s 88ms/step - loss: 0.3577 - accuracy: 0.8990 - val_loss: 0.3096 - val_accuracy: 0.9109
Epoch 3/20
94/94 [=====] - 7s 73ms/step - loss: 0.2876 - accuracy: 0.9163 - val_loss: 0.2606 - val_accuracy: 0.9231
Epoch 4/20
94/94 [=====] - 8s 90ms/step - loss: 0.2444 - accuracy: 0.9287 - val_loss: 0.2238 - val_accuracy: 0.9333
Epoch 5/20
94/94 [=====] - 7s 69ms/step - loss: 0.2113 - accuracy: 0.9376 - val_loss: 0.1944 - val_accuracy: 0.9423
```

epoch: passing the entire dataset through the network



60,000 images / batch size = **512**

= **117** iterations per epoch

= **117 * 80%** for training = **94 iterations per epoch**





see `LeNetTF.py` and `LeNetTF.ipynb`





see `LeNetTF.py` and `LeNetTF.ipynb`



training loss should \approx validation loss

if validation loss \gg training loss \rightarrow overfitting

- too many parameter
- too few images in batch
- too specific/unique batch)

Thank you very much for your attention!

