

## Lecture 11:

# Long Short-Term Memory Networks (LSTMs) – Part I



Markus Hohle  
University California, Berkeley

Machine Learning Algorithms  
MSSE 277B, 3 Units



## Lecture 1: Course Overview and Introduction to Machine Learning

Lecture 2: Bayesian Methods in Machine Learning

classic ML tools & algorithms

Lecture 3: Dimensionality Reduction: Principal Component Analysis

Lecture 4: Linear and Non-linear Regression and Classification

Lecture 5: Unsupervised Learning: K-Means, GMM, Trees

Lecture 6: Adaptive Learning and Gradient Descent Optimization Algorithms

Lecture 7: Introduction to Artificial Neural Networks - The Perceptron

ANNs/AI/Deep Learning

Lecture 8: Introduction to Artificial Neural Networks - Building Multiple Dense Layers

Lecture 9: Convolutional Neural Networks (CNNs) - Part I

Lecture 10: CNNs - Part II

## **Lecture 11: Recurrent Neural Networks (RNNs) and Long Short-Term Memory (LSTMs)**

Lecture 12: Combining LSTMs and CNNs

Lecture 13: Running Models on GPUs and Parallel Processing

Lecture 14: Project Presentations

Lecture 15: Transformer

Lecture 16: GNN



<https://www.analyticsvidhya.com>



## Outline

- Idea and classic RNNs
- LSTMs
- *BackPropagation Through Time (BPTT)*
- Syntax and some examples



<https://www.analyticsvidhya.com>

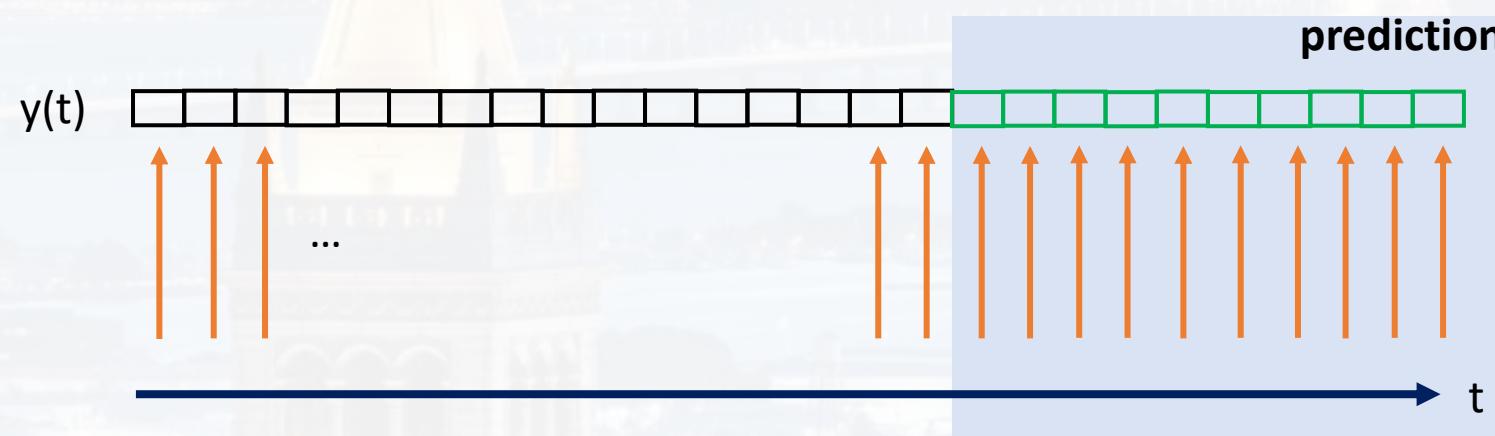


## Outline

- Idea and classic RNNs
- LSTMs
- BackPropagation Through Time (BPTT)
- Syntax and some examples

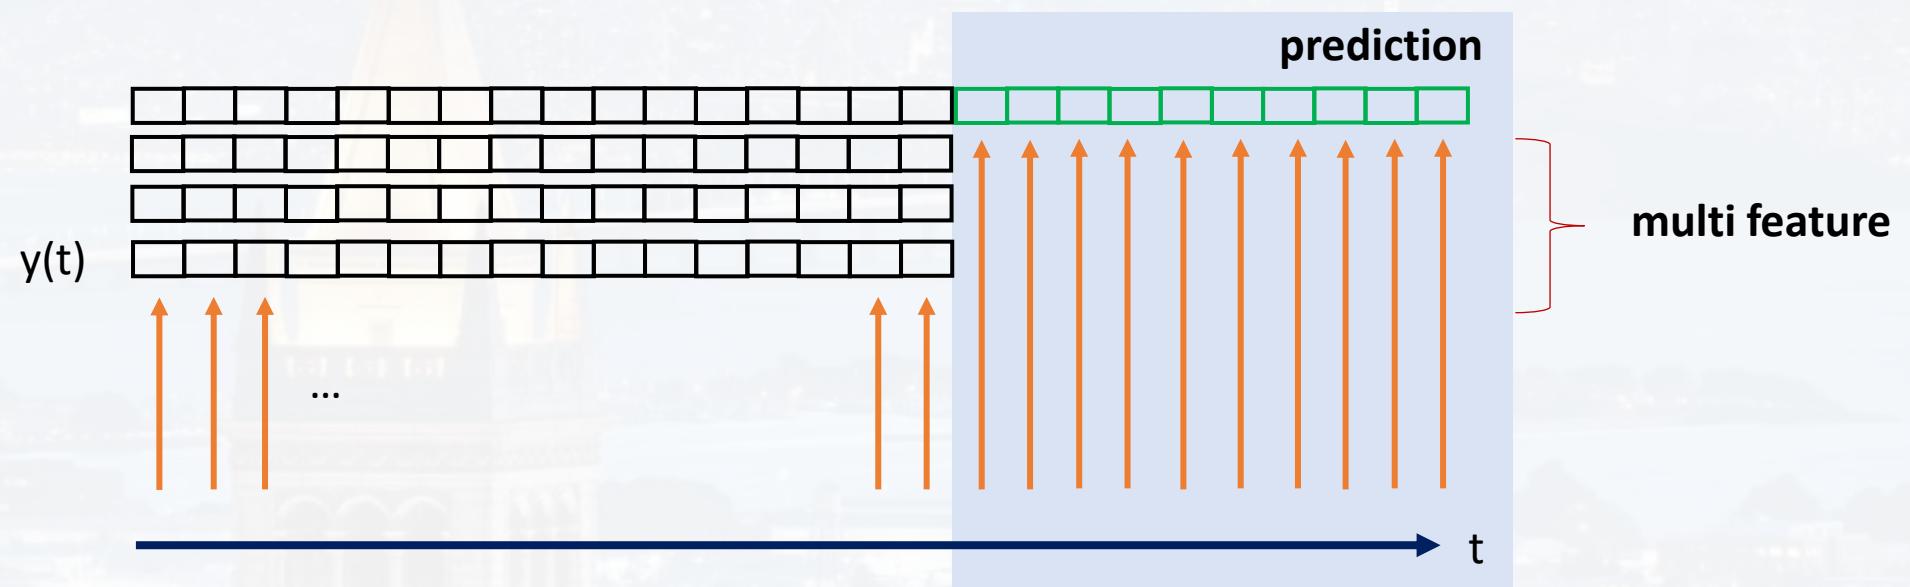


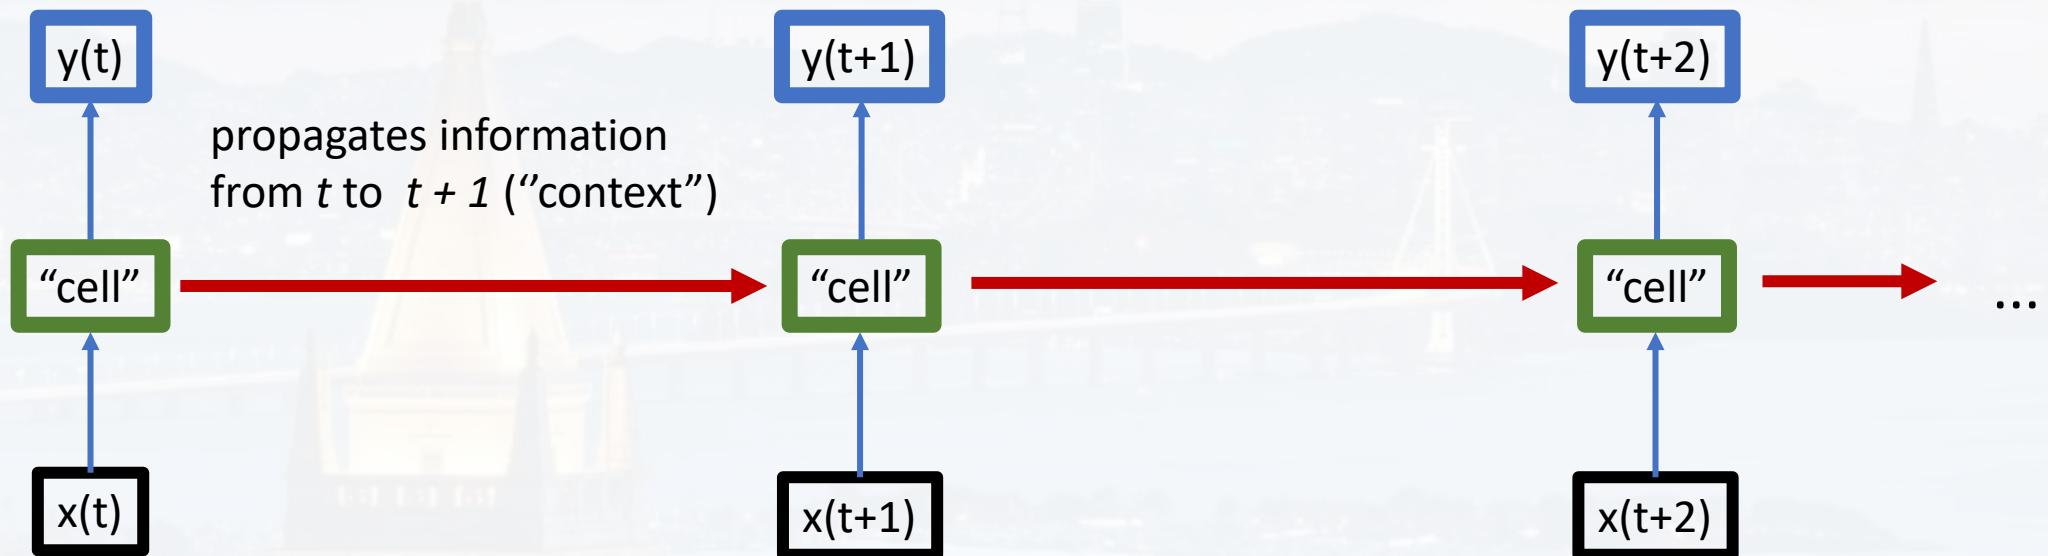
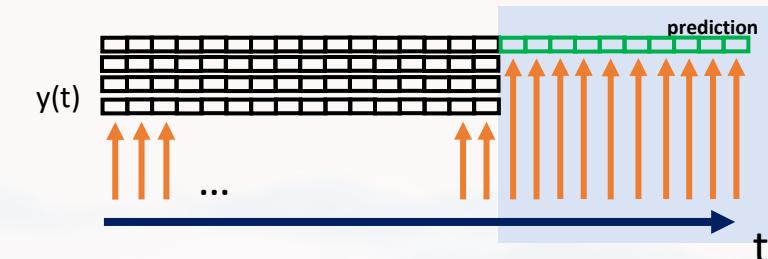
- Recurrent Neural Network
  - time series analysis regression (**prediction and forecasting**)
  - first step towards GenAI
  - time series analysis classification
  - early speech recognition
  - handwriting
  - “precursor” of LSTMs
  - invented by **Shun'ichi Amari** in 1972





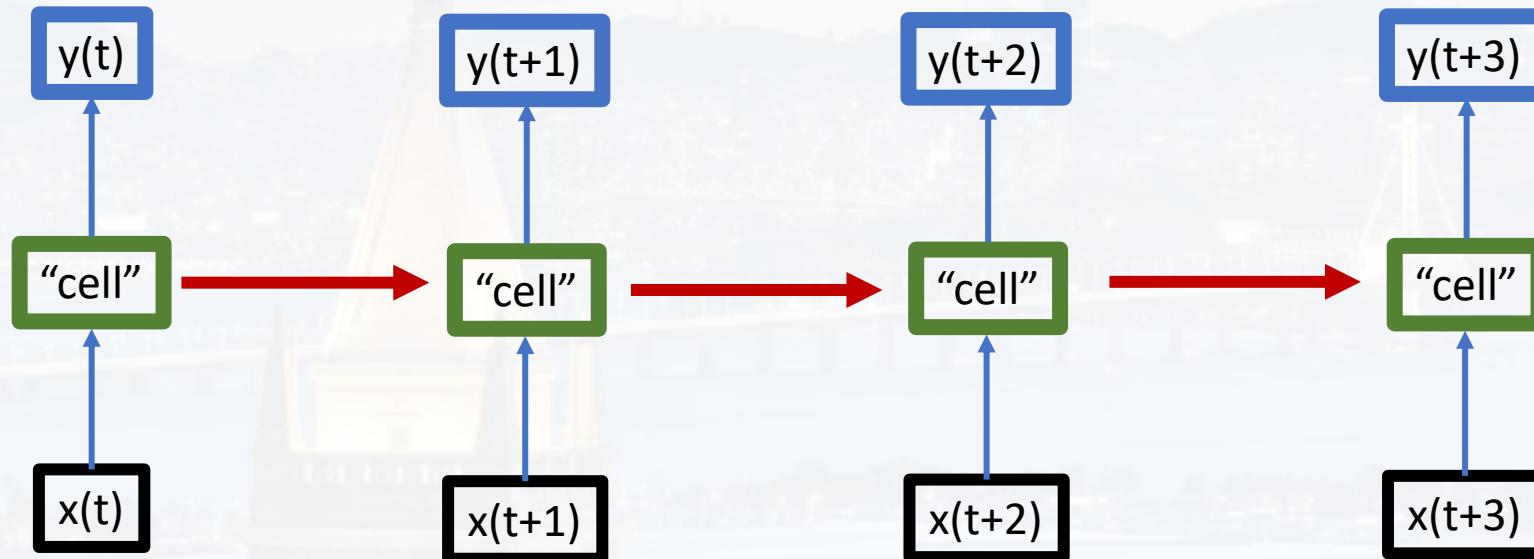
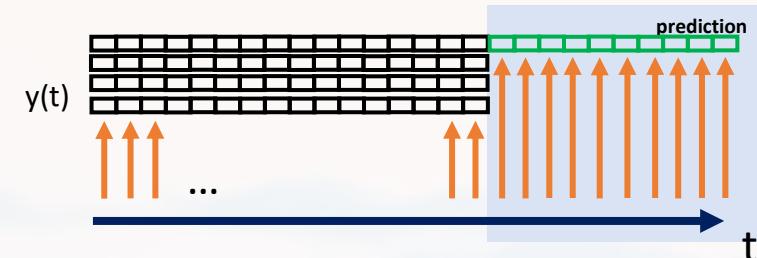
- Recurrent Neural Network
  - time series analysis regression (**prediction and forecasting**)
  - first step towards GenAI
  - time series analysis classification
  - early speech recognition
  - handwriting
  - “precursor” of LSTMs
  - invented by **Shun'ichi Amari** in 1972





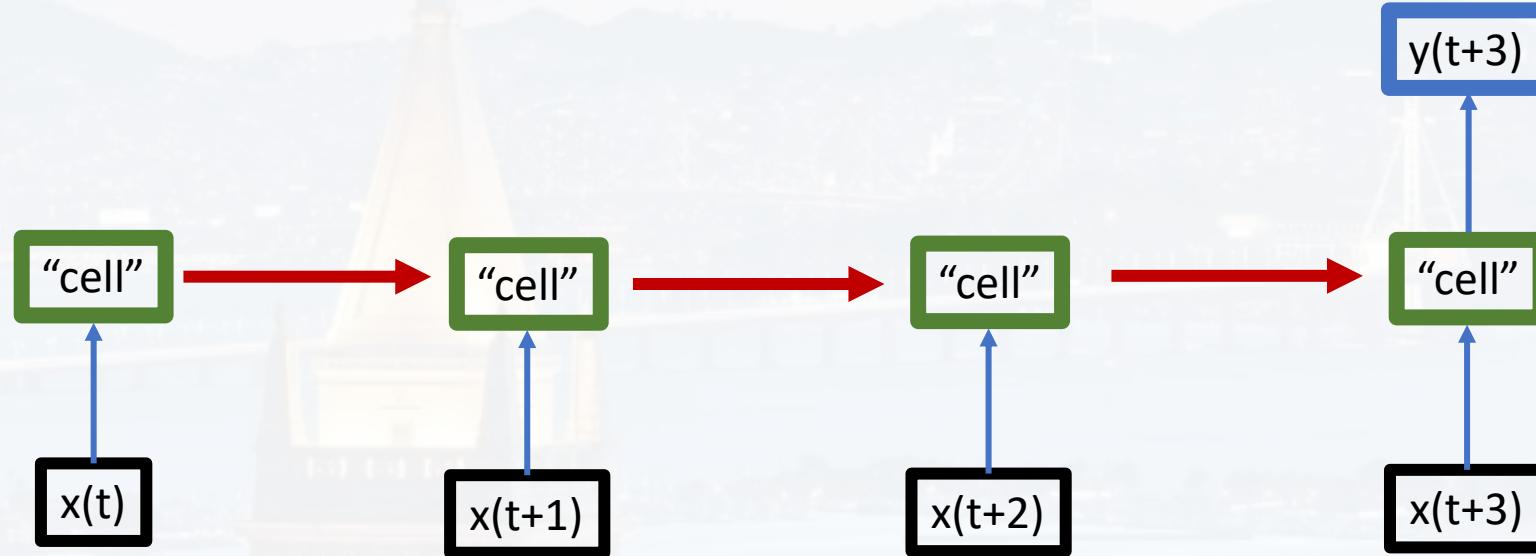
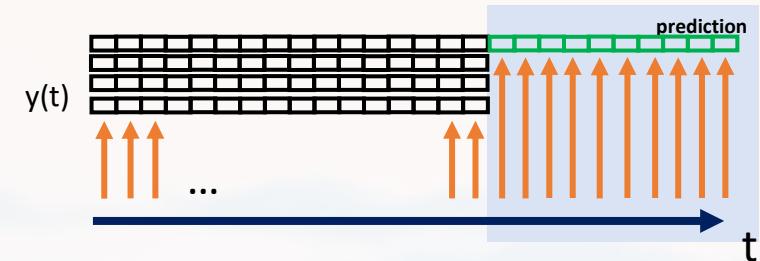


“many to many”



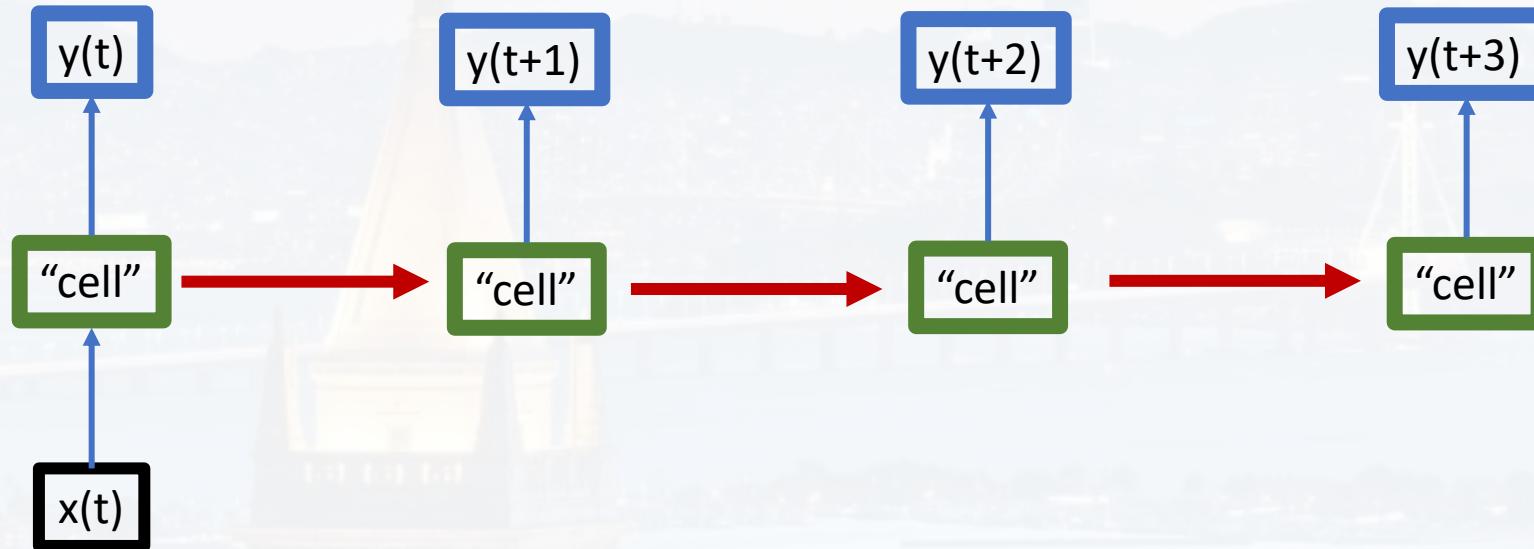
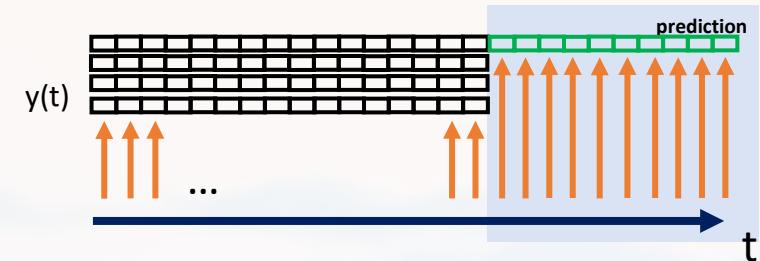


**“many to one”**  
(classification)





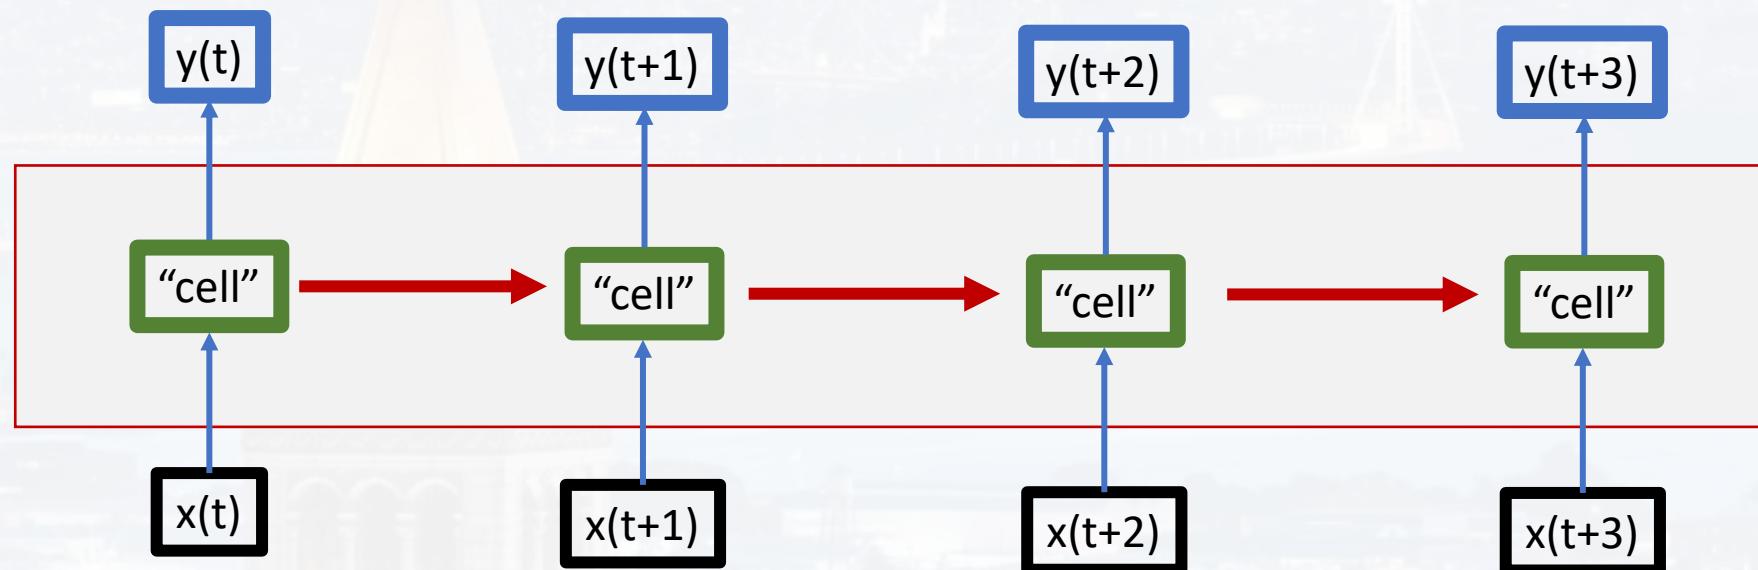
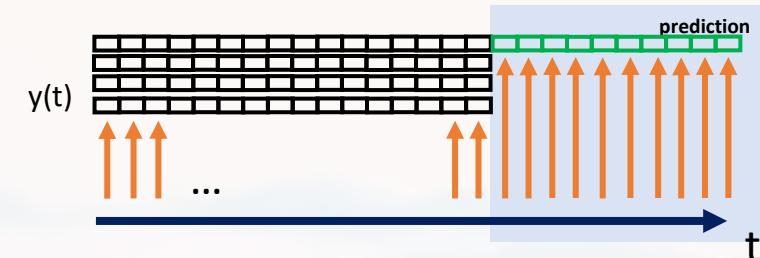
**“one to many”**  
(image capturing)





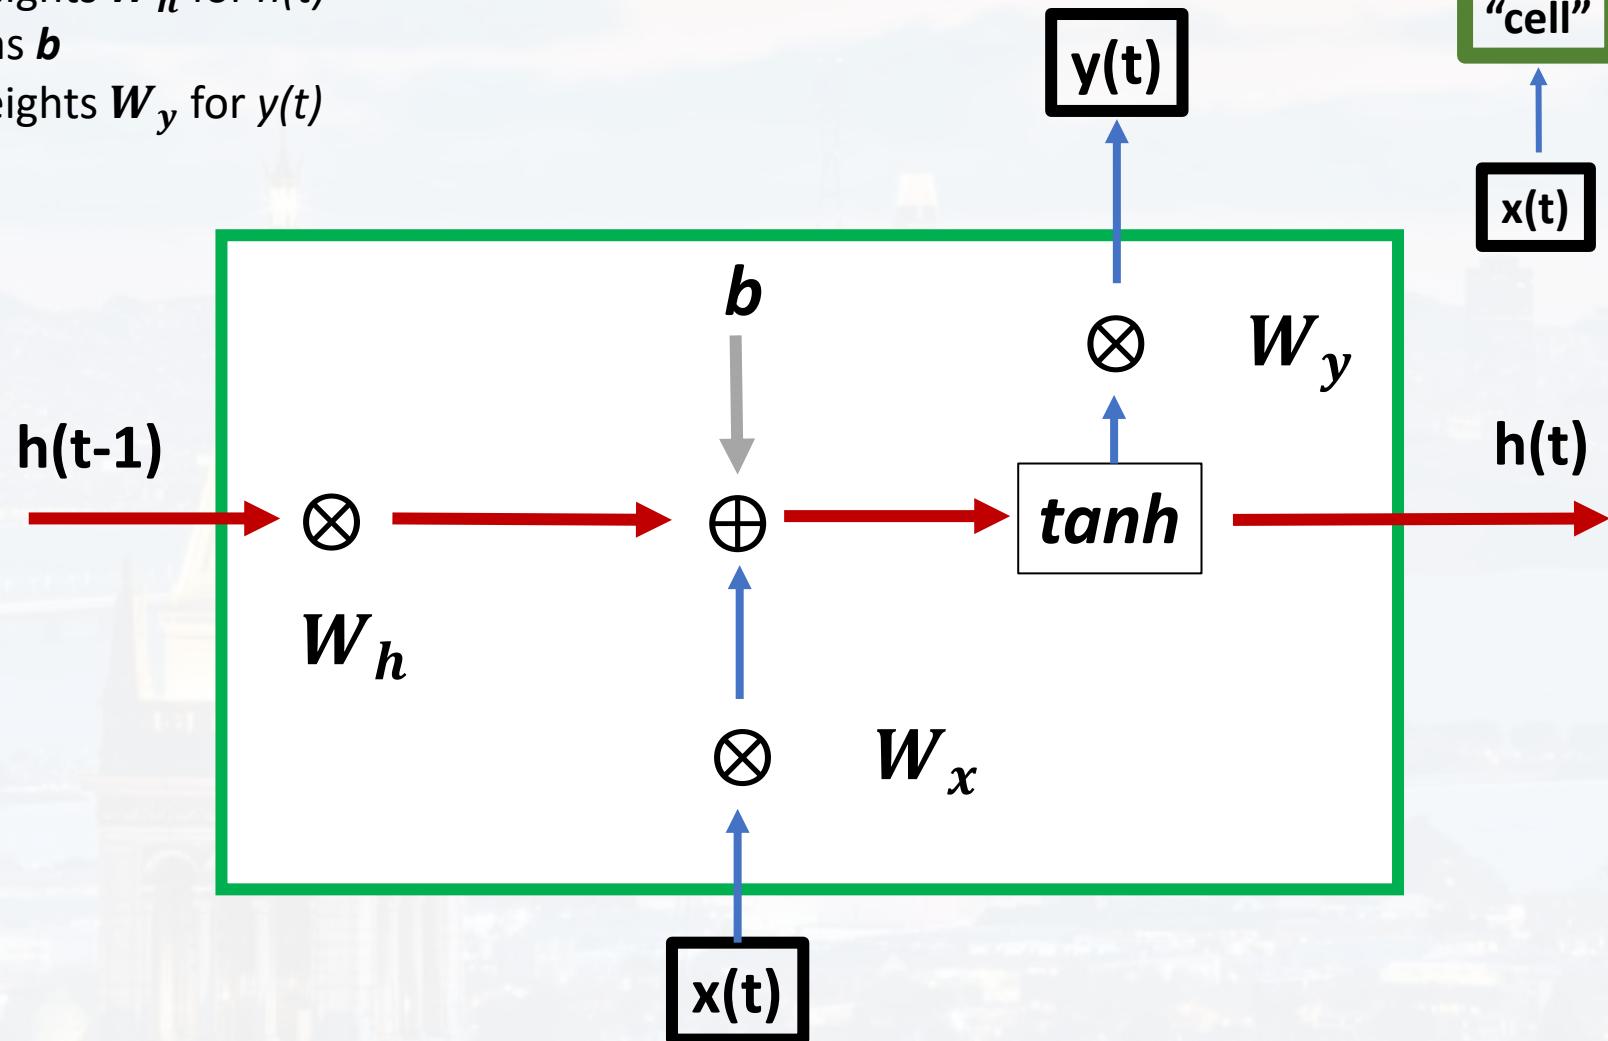
Applying the **identical cell recursively!**

- easy to implement
- direction (arrow of time, see later)
- exploding/vanishing gradients
- classic RNN has a “short memory”



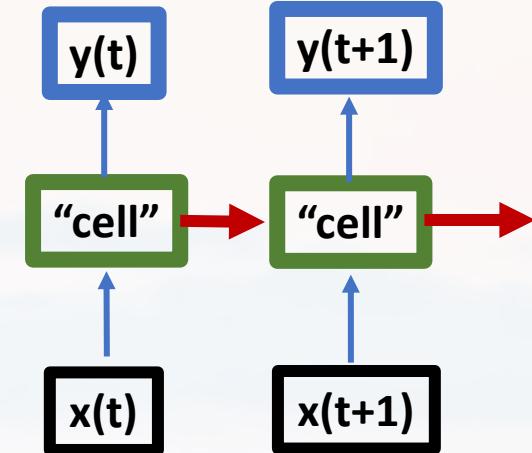
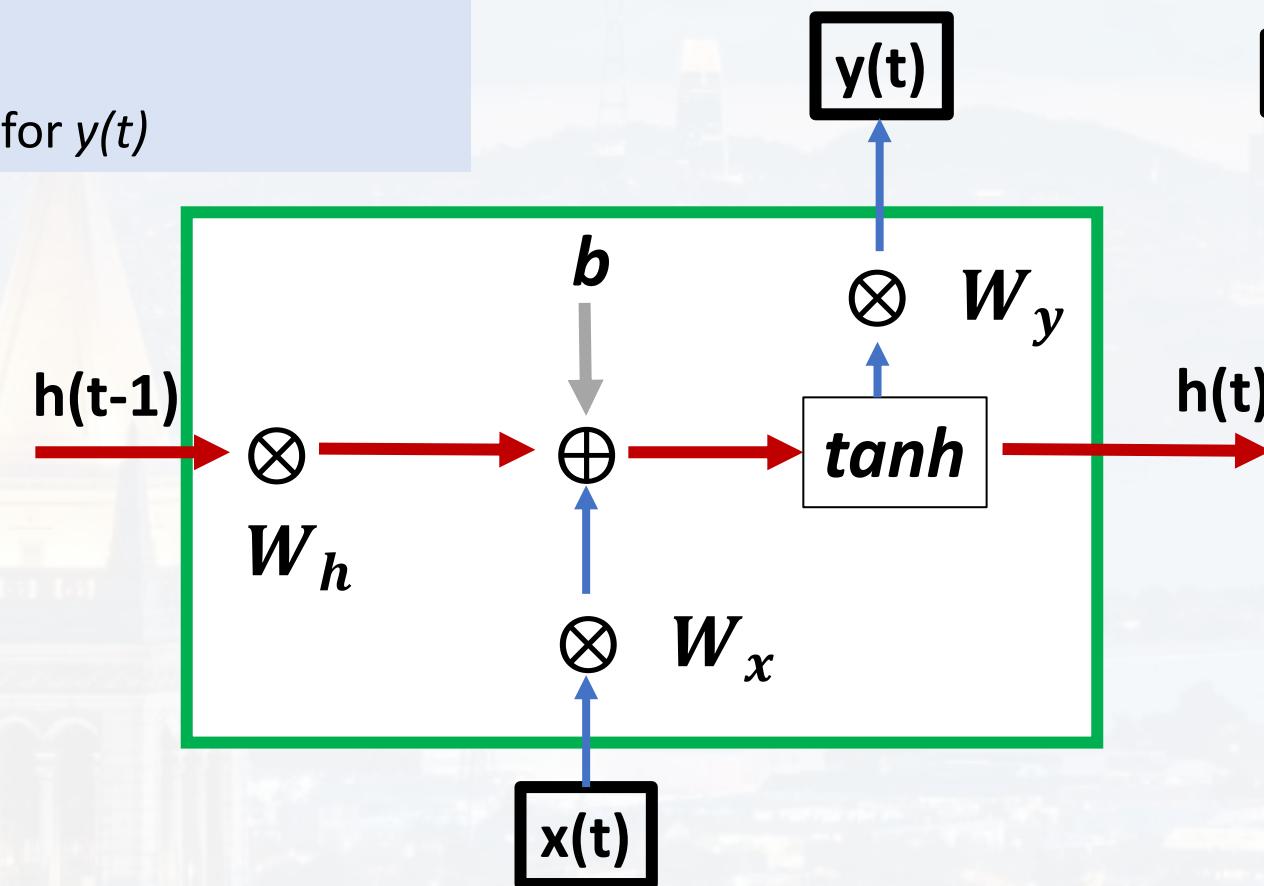


- state vector  $h(t)$ , i. e. the “memory” the system has along time
- weights  $W_x$  for  $x(t)$
- weights  $W_h$  for  $h(t)$
- bias  $b$
- weights  $W_y$  for  $y(t)$





- state vector  $h(t)$ , i. e. the “memory” the system has along time
- weights  $W_x$  for  $x(t)$  **learnable**
- weights  $W_h$  for  $h(t)$
- bias  $b$
- weights  $W_y$  for  $y(t)$





$$h(t) = \tanh[x(t) * W_x + W_h * h(t-1) + b]$$

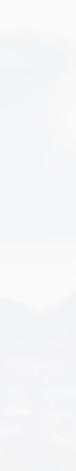


$n \times 1$

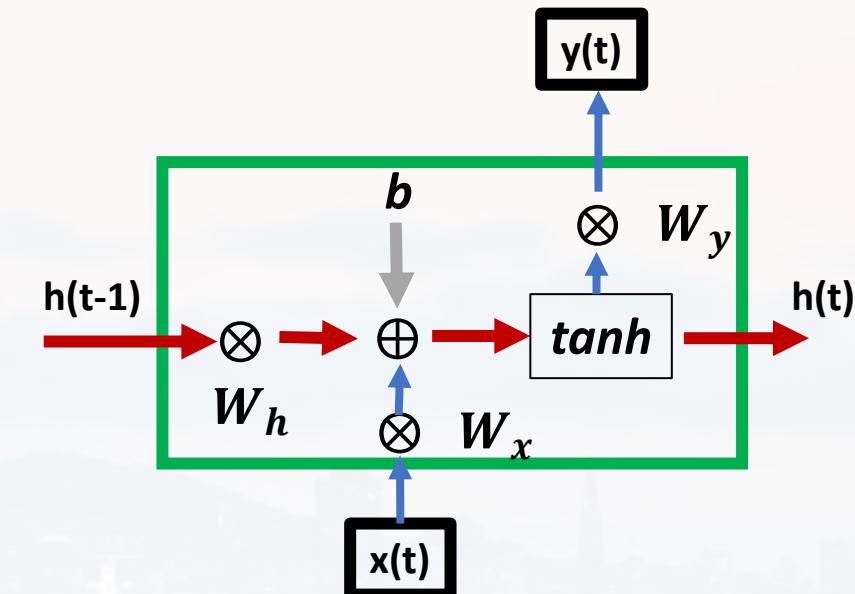
$1 \times 1$



$n \times 1$



$$y(t) = h(t) * W_y$$



$n$ : number of states/ neurons

$n \times n$



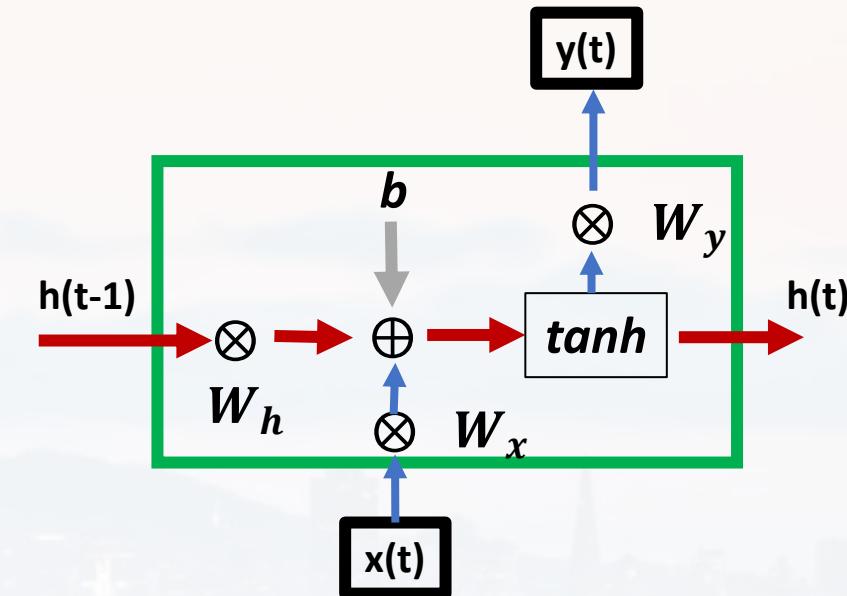
$$h(t) = \tanh[x(t) * W_x + W_h * h(t-1) + b]$$

$$y(t) = h(t) * W_y$$

**1 x 1**

**n x 1**

**1 x n**



**n:** number of states/ neurons



$$h(t) = \tanh[x(t) * W_x + W_h * h(t-1) + b]$$

$$y(t) = h(t) * W_y$$

usually,  $x(t)$  comes in **batches of size  $B$  and of length  $T$**   
and has  **$F$  features** (see also later)

for **each time point  $t$ :**

$$x(t) * W_x^T + h(t-1) * W_h + b$$

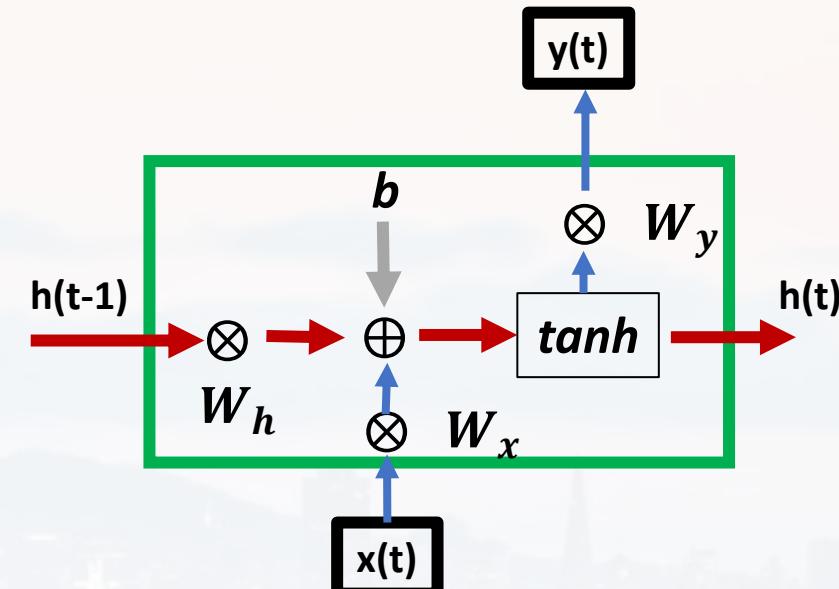
shape:  $(B \times 1 \times F) * (F \times n)^T + (B \times 1 \times n) * (n \times n) + B \times 1 \times n$

$$B \times 1 \times n$$

$$B \times 1 \times n$$

$$y(t) = h(t) * W_y^T$$

shape:  $(B \times 1 \times F) * (B \times 1 \times n) * (F \times n)^T$



**n:** number of states/ neurons



<https://www.analyticsvidhya.com>



## Outline

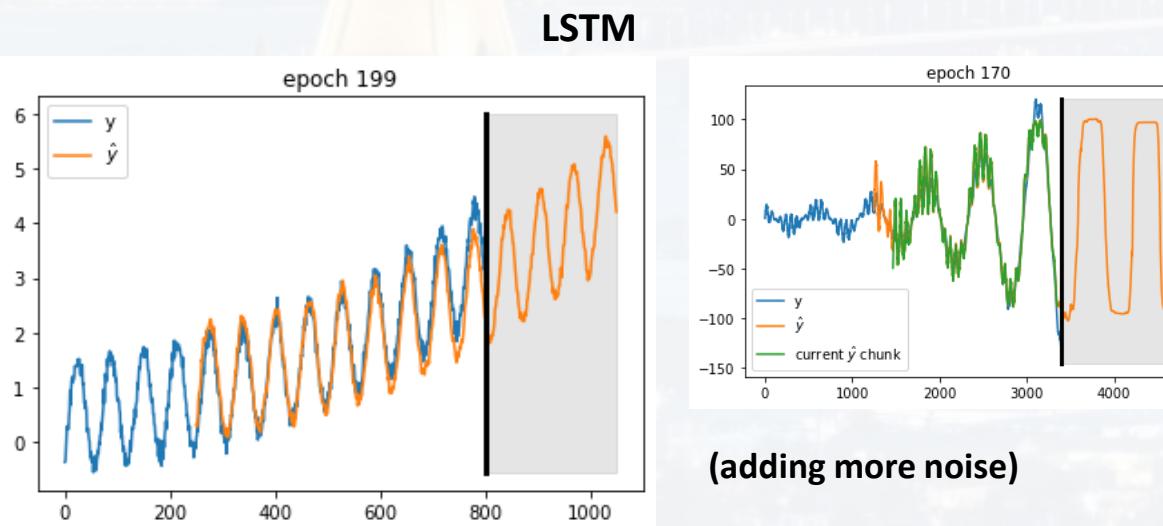
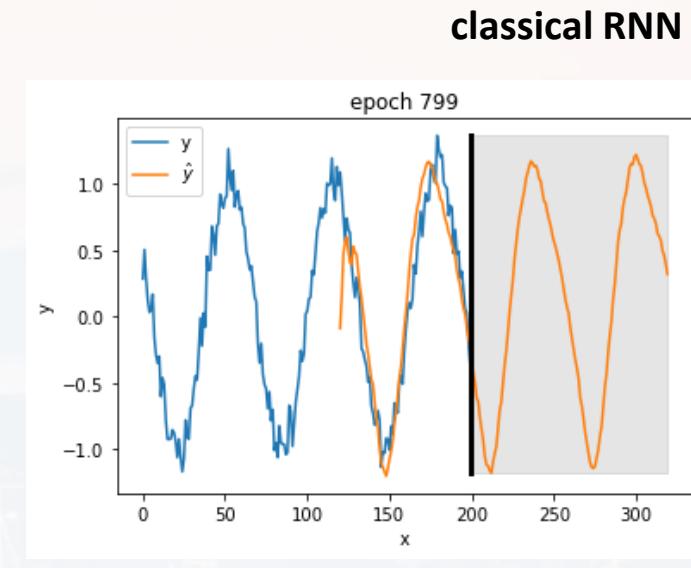
- Idea and classic RNNs
- **LSTMs**
- BackPropagation Through Time (BPTT)
- Syntax and some examples

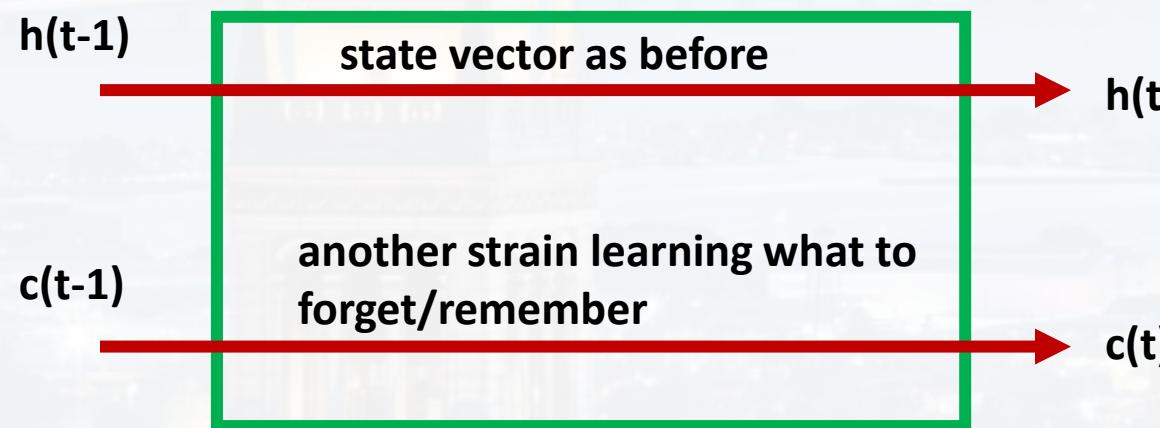
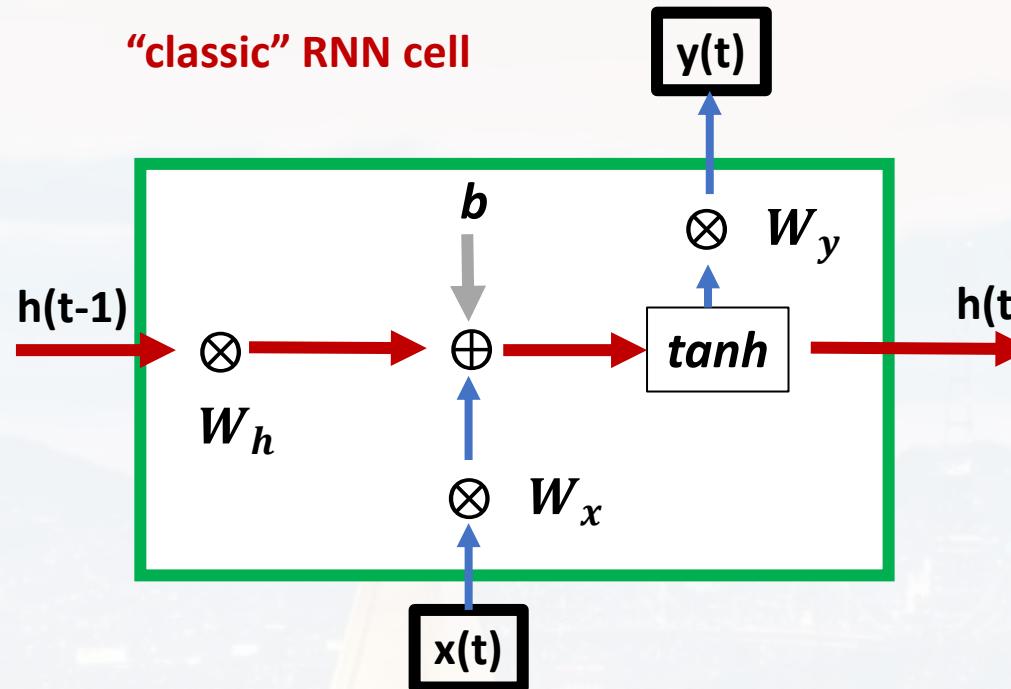


- **Long- Short Term Memory**

new:

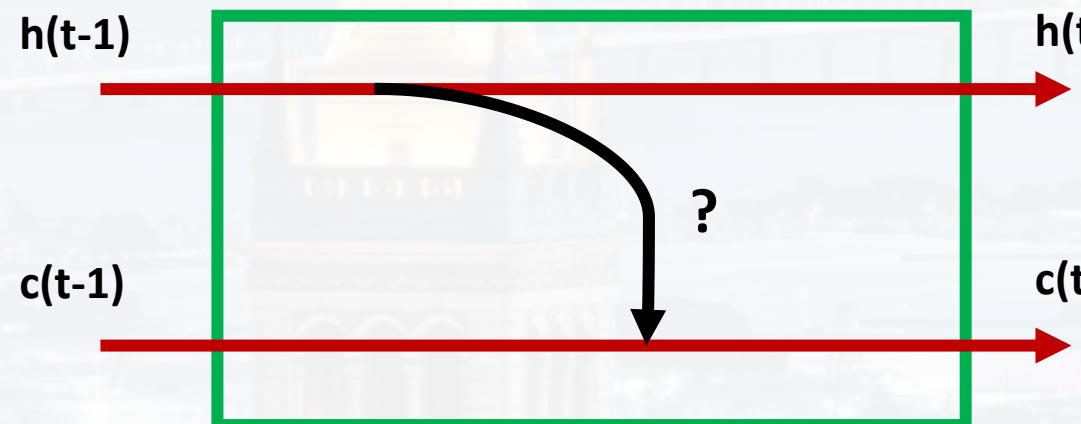
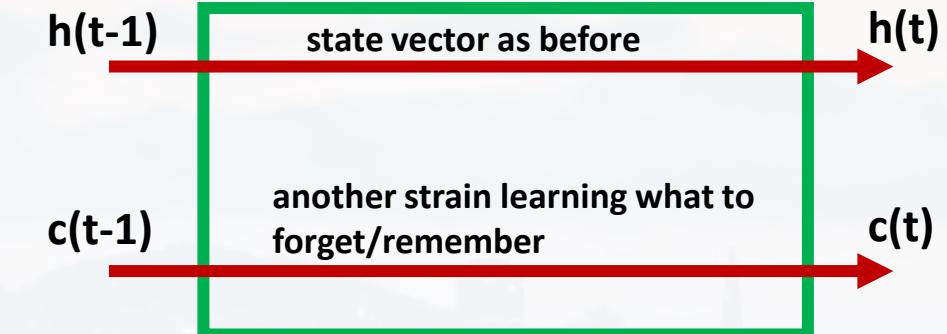
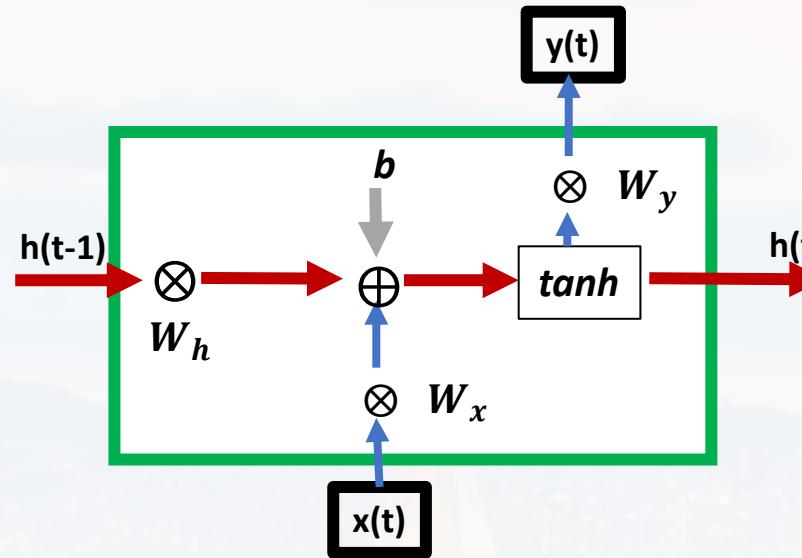
- long-term and short-term memory
- dealing with vanishing/exploding gradient
- invented 1997 by Sepp Hochreiter und Jürgen Schmidhuber

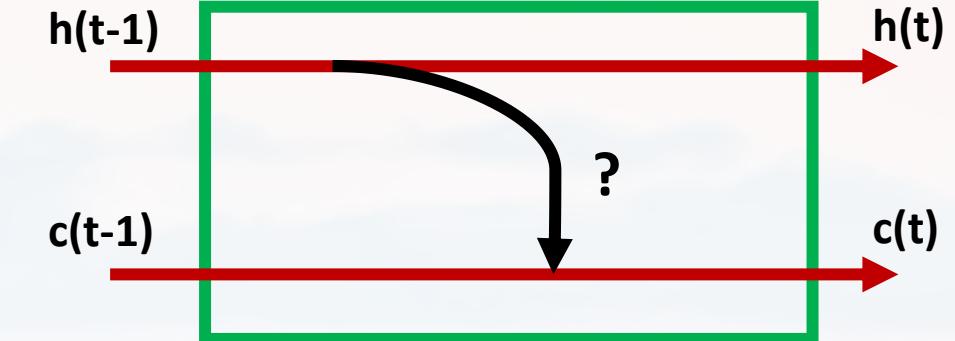
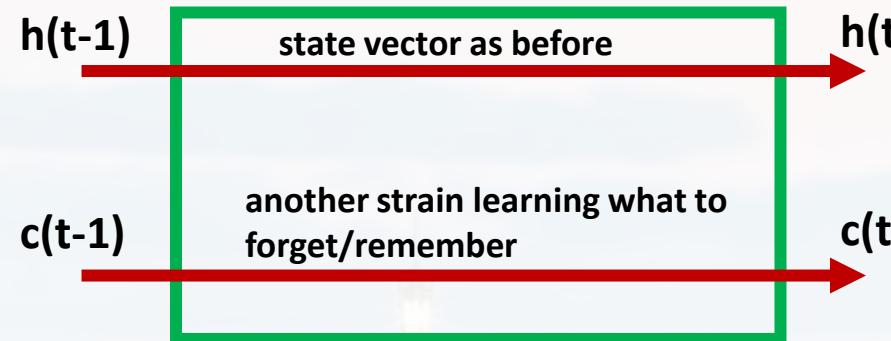




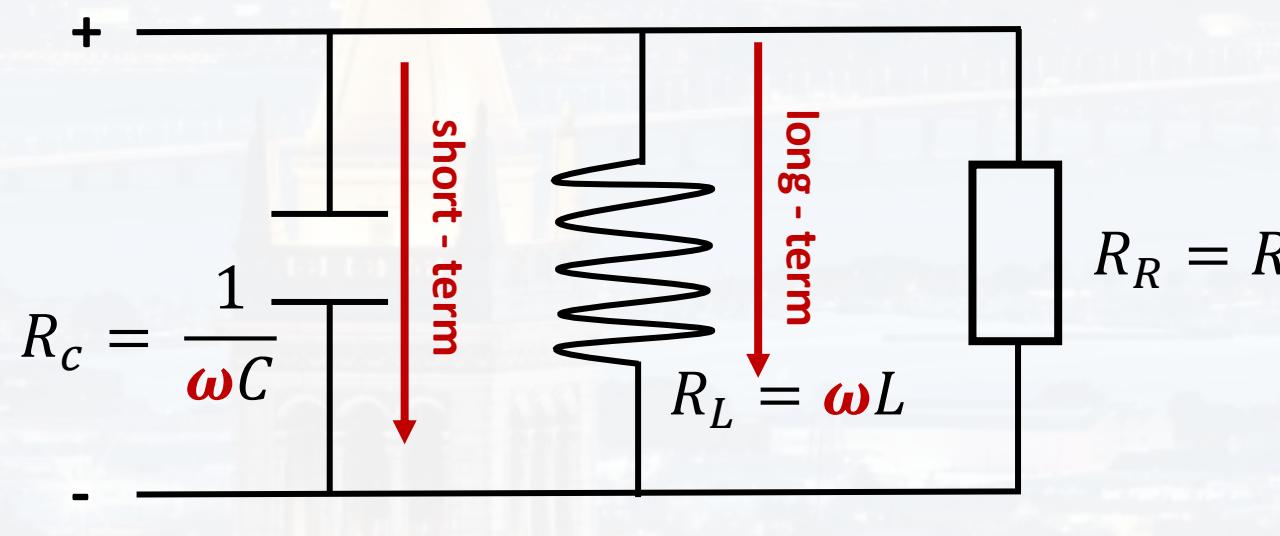


"classic" RNN cell





electrical circuits:

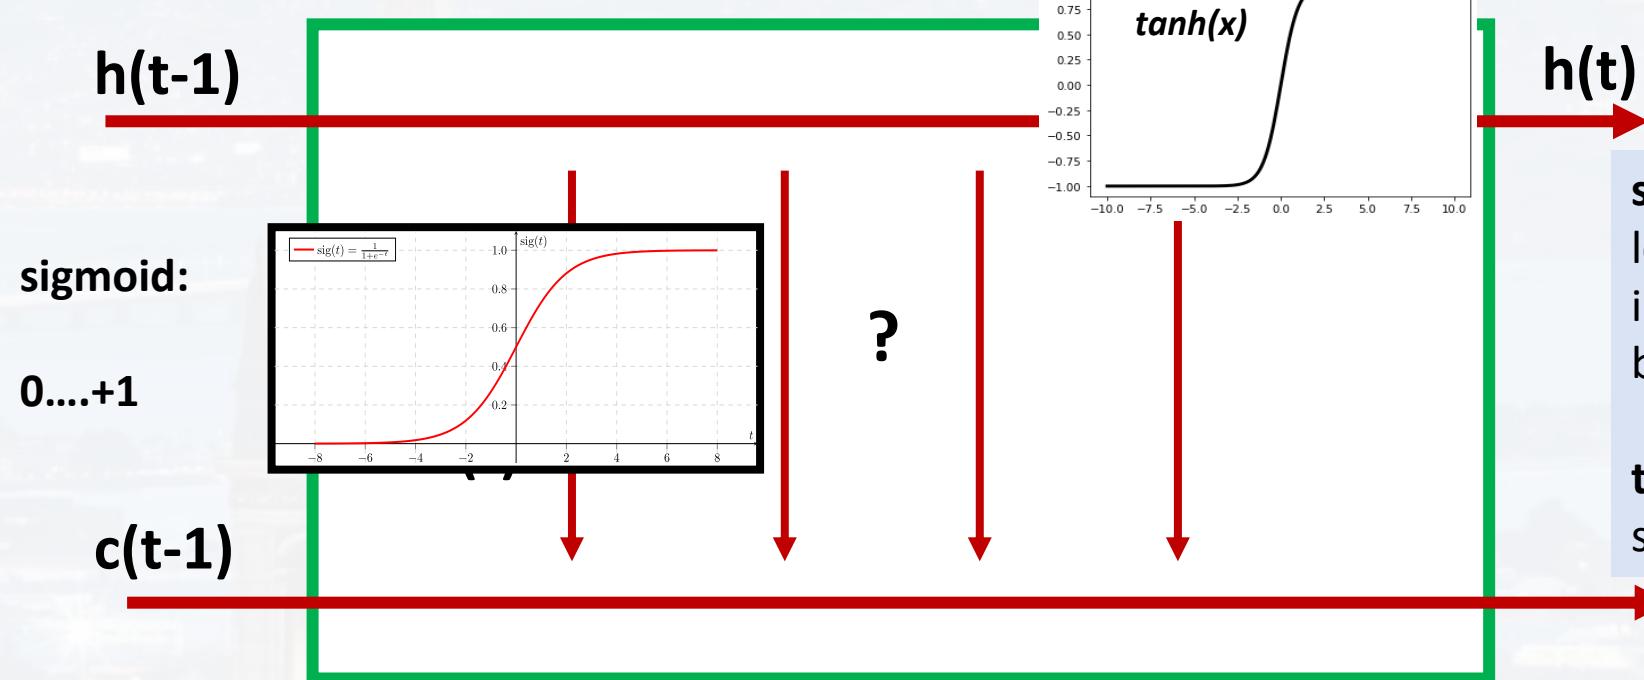
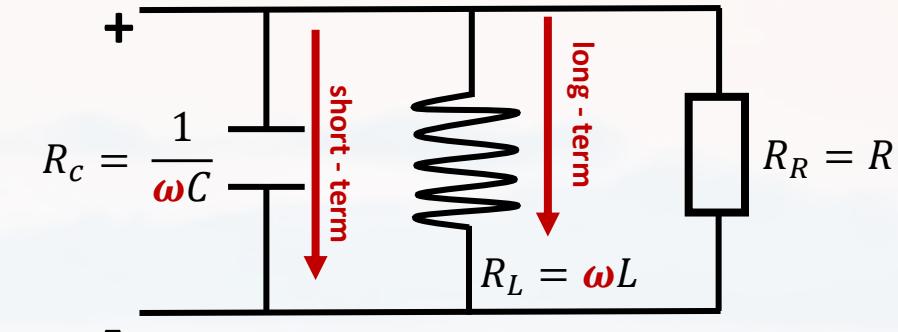
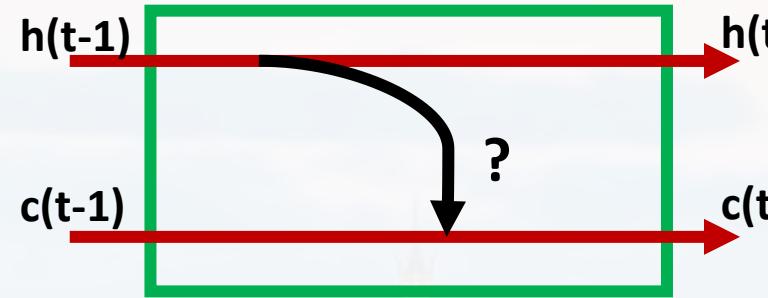


$$\text{AC: } I(t) = I_0 e^{i(\omega t + \varphi)}$$

$R_c$ : passes **short** -term changes

$R_L$ : passes **long** -term changes

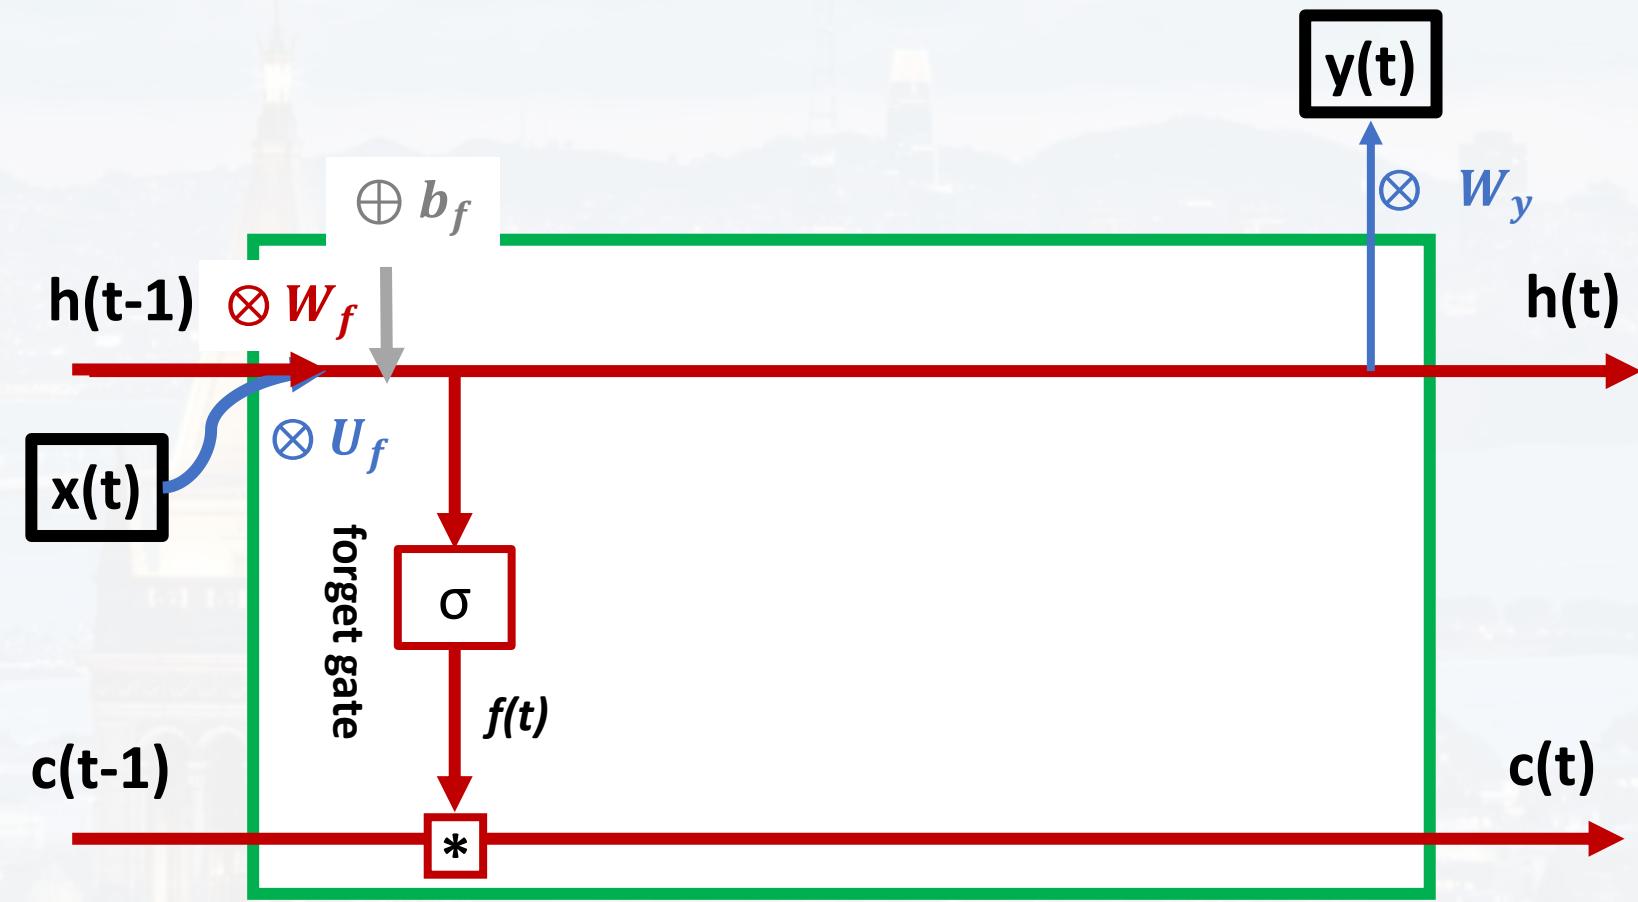
$$\frac{1}{R_{tot}} = \frac{1}{R_R} + \frac{1}{R_C} + \frac{1}{R_L}$$





$$f(t) = \sigma(U_f \otimes x(t) + W_f \otimes h(t-1) + b_f)$$

\* element-wise multiplication



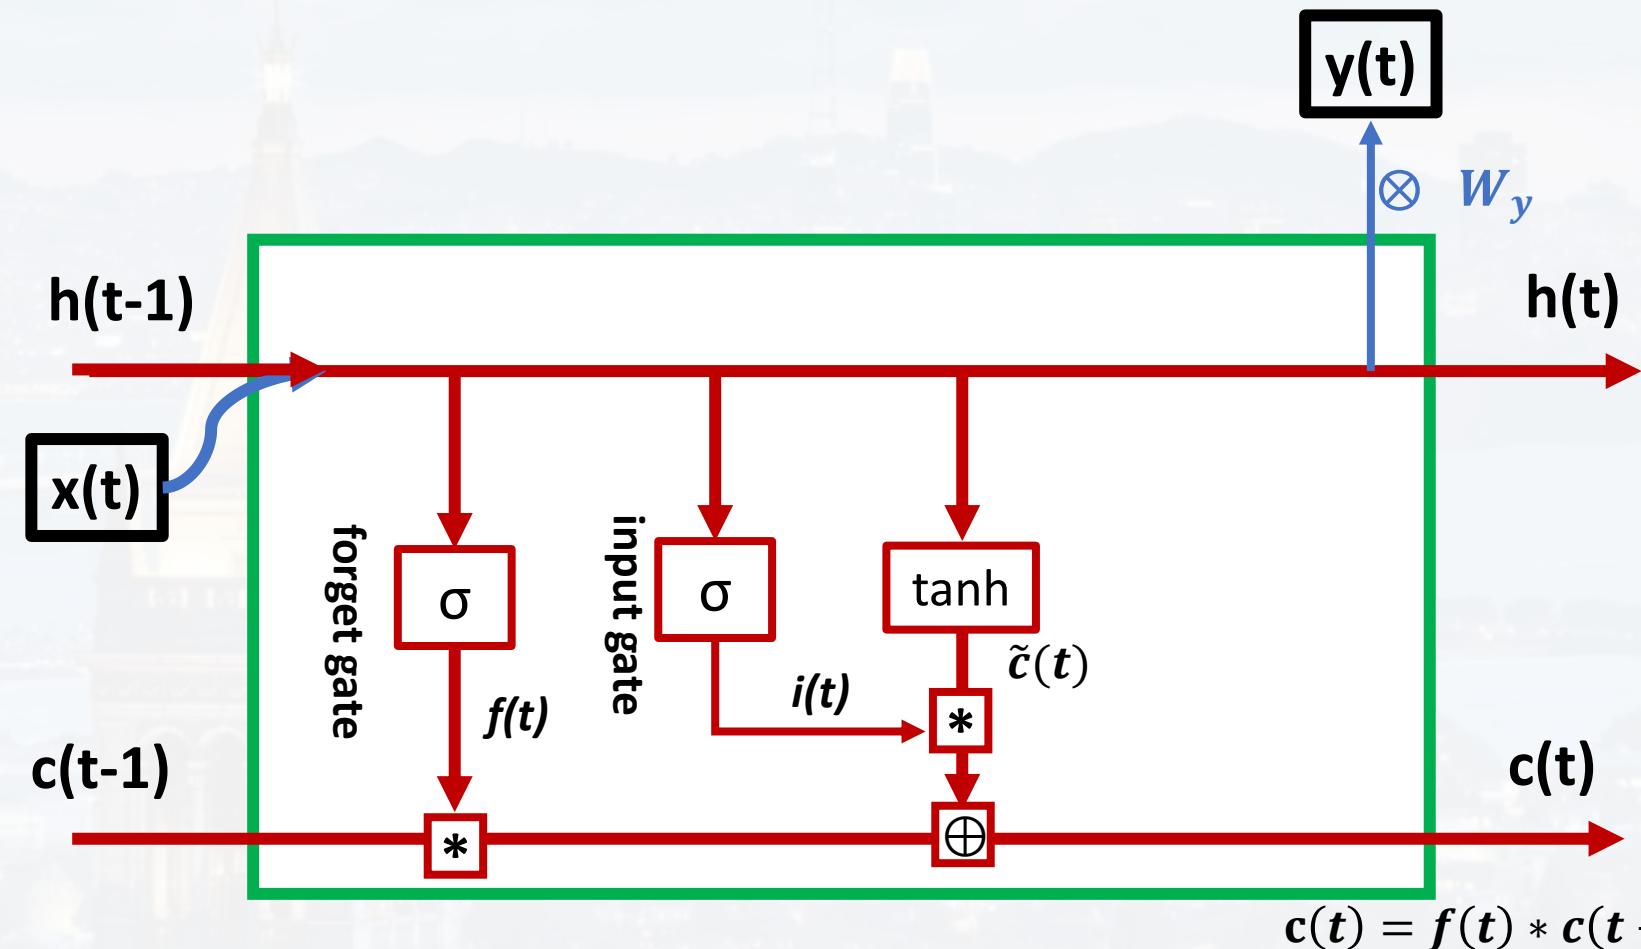


$$f(t) = \sigma(U_f \otimes x(t) + W_f \otimes h(t-1) + b_f)$$

\* element-wise multiplication

$$i(t) = \sigma(U_i \otimes x(t) + W_i \otimes h(t-1) + b_i)$$

$$\tilde{c}(t) = \tanh(U_g \otimes x(t) + W_g \otimes h(t-1) + b_g)$$





$$f(t) = \sigma(U_f \otimes x(t) + W_f \otimes h(t-1) + b_f)$$

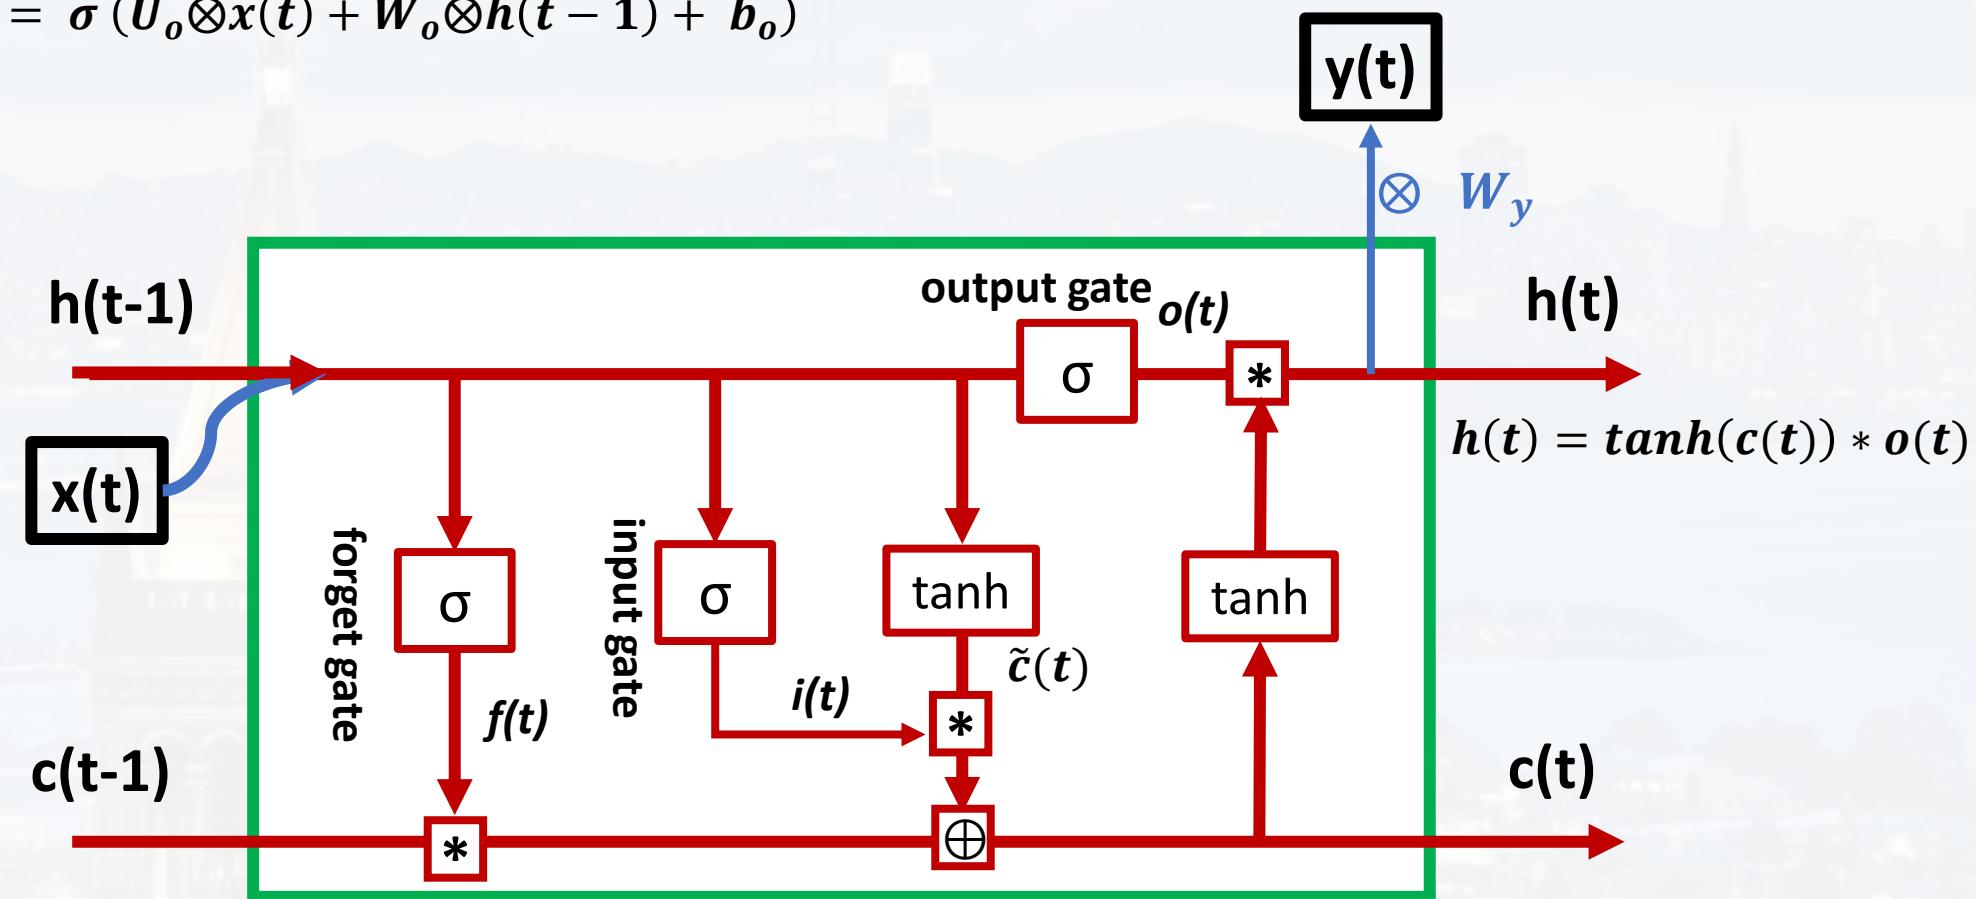
\* element-wise multiplication

$$i(t) = \sigma(U_i \otimes x(t) + W_i \otimes h(t-1) + b_i)$$

$$\tilde{c}(t) = \tanh(U_g \otimes x(t) + W_g \otimes h(t-1) + b_g)$$

$$c(t) = f(t) * c(t-1) + i(t) * \tilde{c}(t)$$

$$o(t) = \sigma(U_o \otimes x(t) + W_o \otimes h(t-1) + b_o)$$





$$f(t) = \sigma(U_f \otimes x(t) + W_f \otimes h(t-1) + b_f)$$

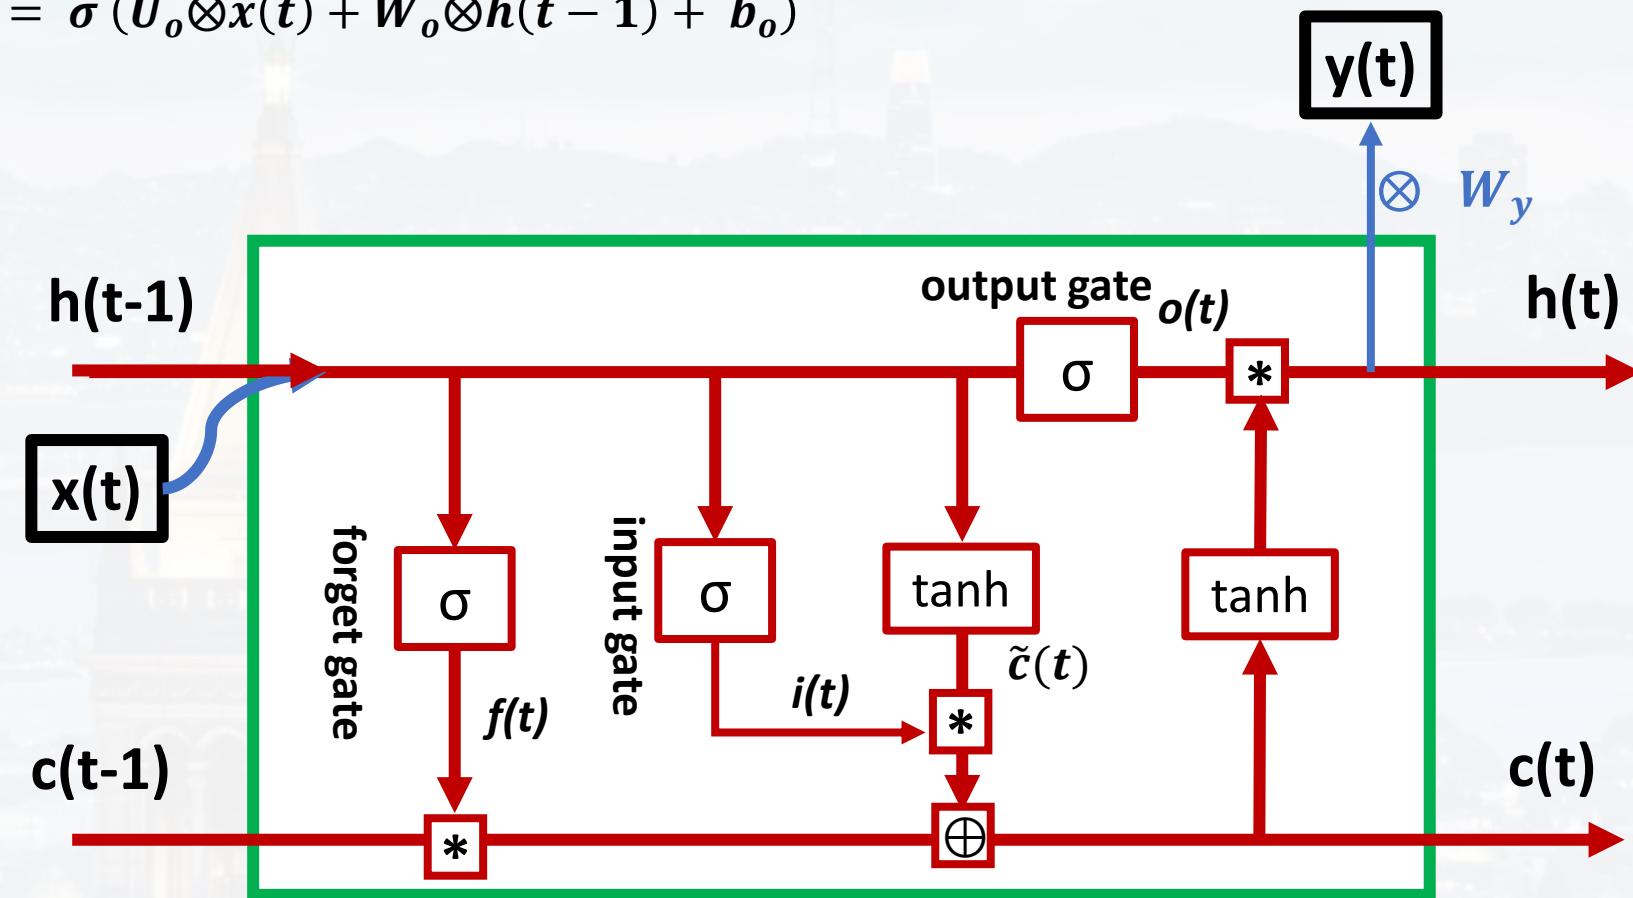
\* element-wise multiplication

$$i(t) = \sigma(U_i \otimes x(t) + W_i \otimes h(t-1) + b_i)$$

$$\tilde{c}(t) = \tanh(U_g \otimes x(t) + W_g \otimes h(t-1) + b_g)$$

$$c(t) = f(t) * c(t-1) + i(t) * \tilde{c}(t)$$

$$o(t) = \sigma(U_o \otimes x(t) + W_o \otimes h(t-1) + b_o)$$

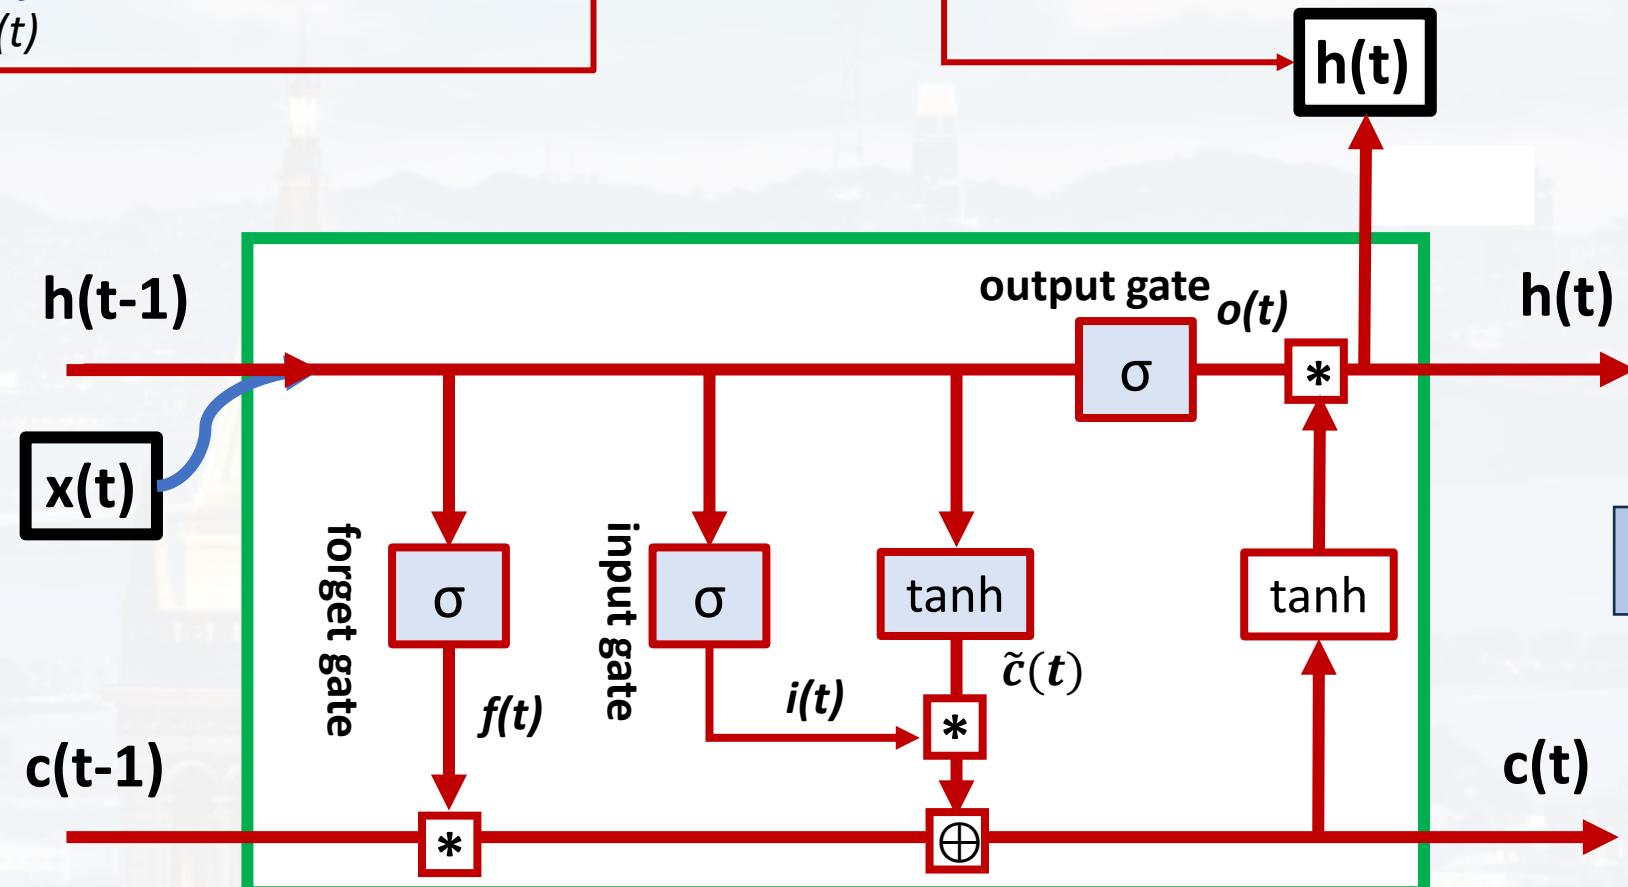




There is one more thing:

\* element – wise multiplication

we will add a **dense layer** instead of  $W_y$  at the end to convert  $h(t)$  to  $y(t)$





<https://www.analyticsvidhya.com>

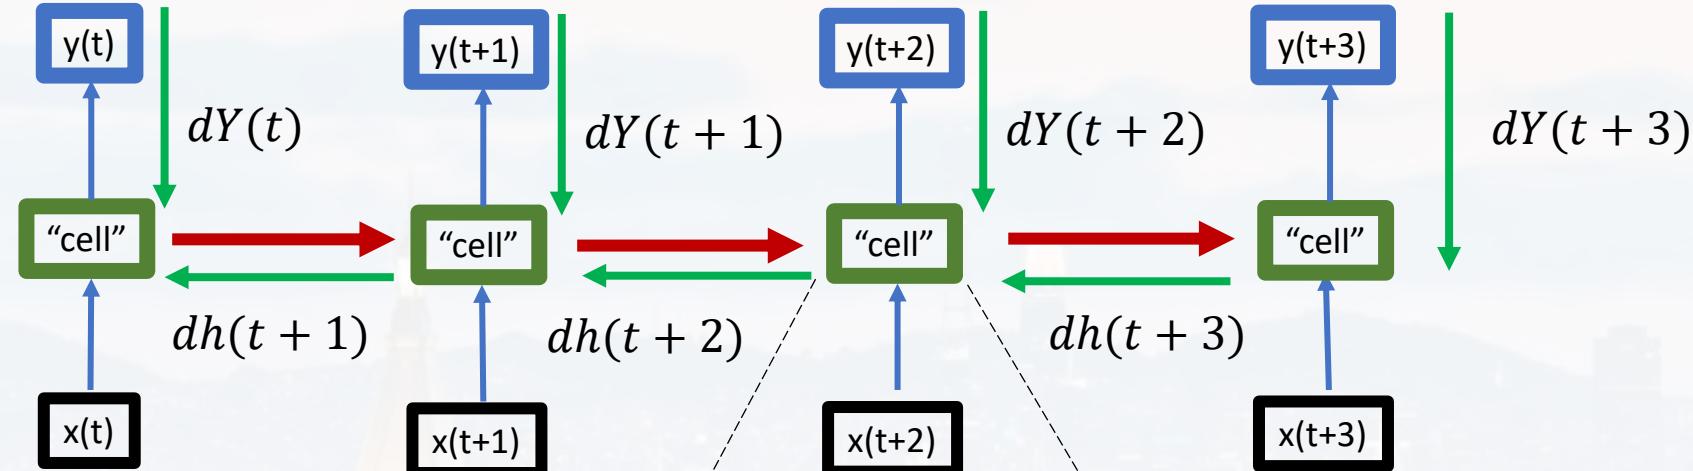


## Outline

- Idea and classic RNNs
- LSTMs
- BackPropagation Through Time (BPTT)**
- Syntax and some examples



because of the RNN/LSTM architecture, backpropagation works a bit different:



backpropagation  
through time  
(BPTT)

more details here:



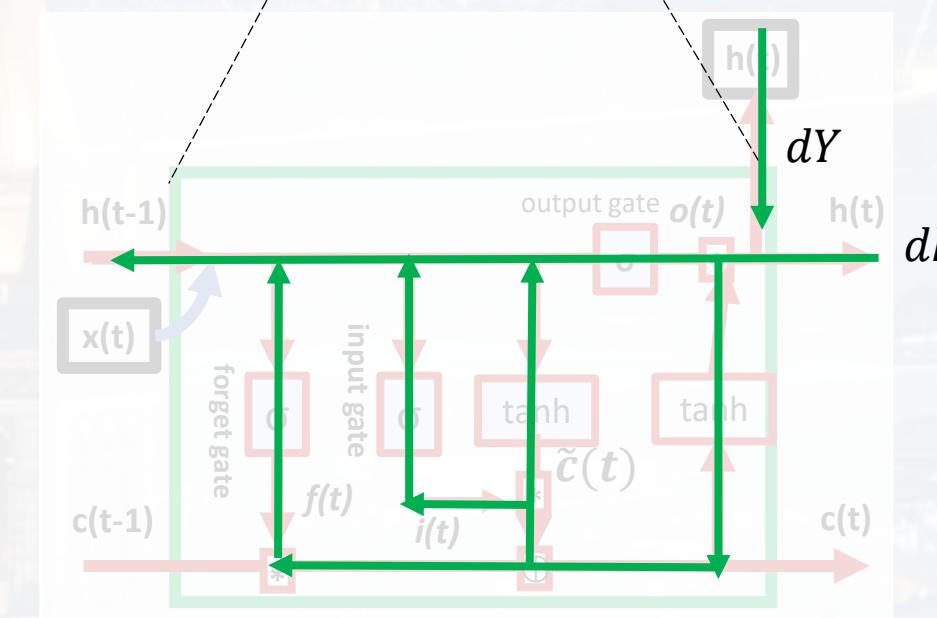
**Dr Fridolin**

@DrFridolin · 126 subscribe

Dr Fridolin shows you how to

[github.com/DrBigMau](https://github.com/DrBigMau)

Subscribe



note: there is no  $dc$  since  $c(t)$  is not a learnable!



<https://www.analyticsvidhya.com>



## Outline

- Idea and classic RNNs
- LSTMs
- BackPropagation Through Time (BPTT)
- Syntax and some examples



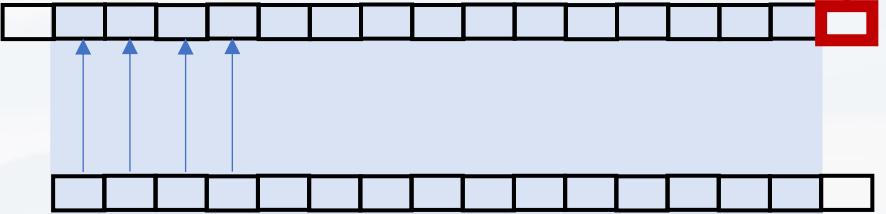
Let us first understand the logic:

$y(t)$

no data to compare with

predicting **one** step in the future  
by **one** step from the past

$$dt_{futu} = 1$$
$$dt_{past} = 1$$



length of training data is: `len[y(t)] - dtfutu - dtpast + 1`

length of training data

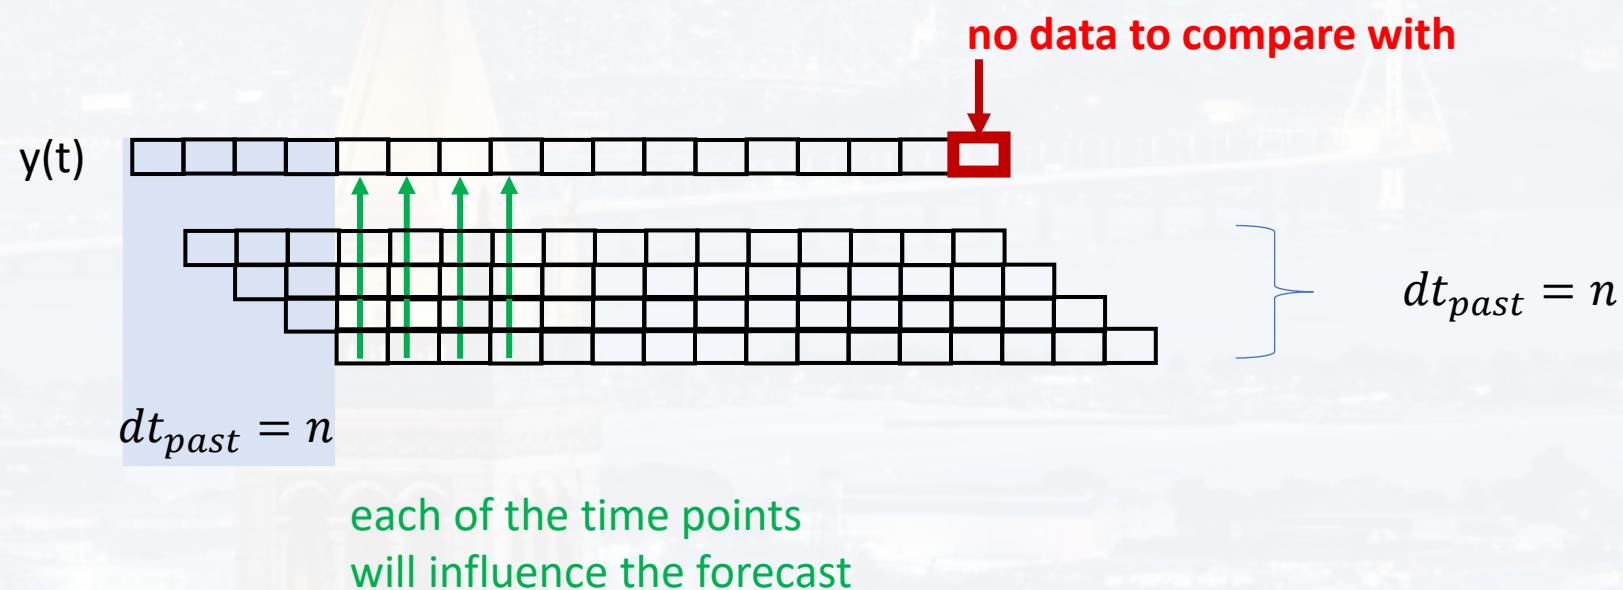


Let us first understand the logic:

length of training data is: `len[y(t)] - dtfutu - dtpast + 1`

predicting **m** steps in the future  
by **n** steps from the past

$$dt_{futu} = m$$
$$dt_{past} = n$$





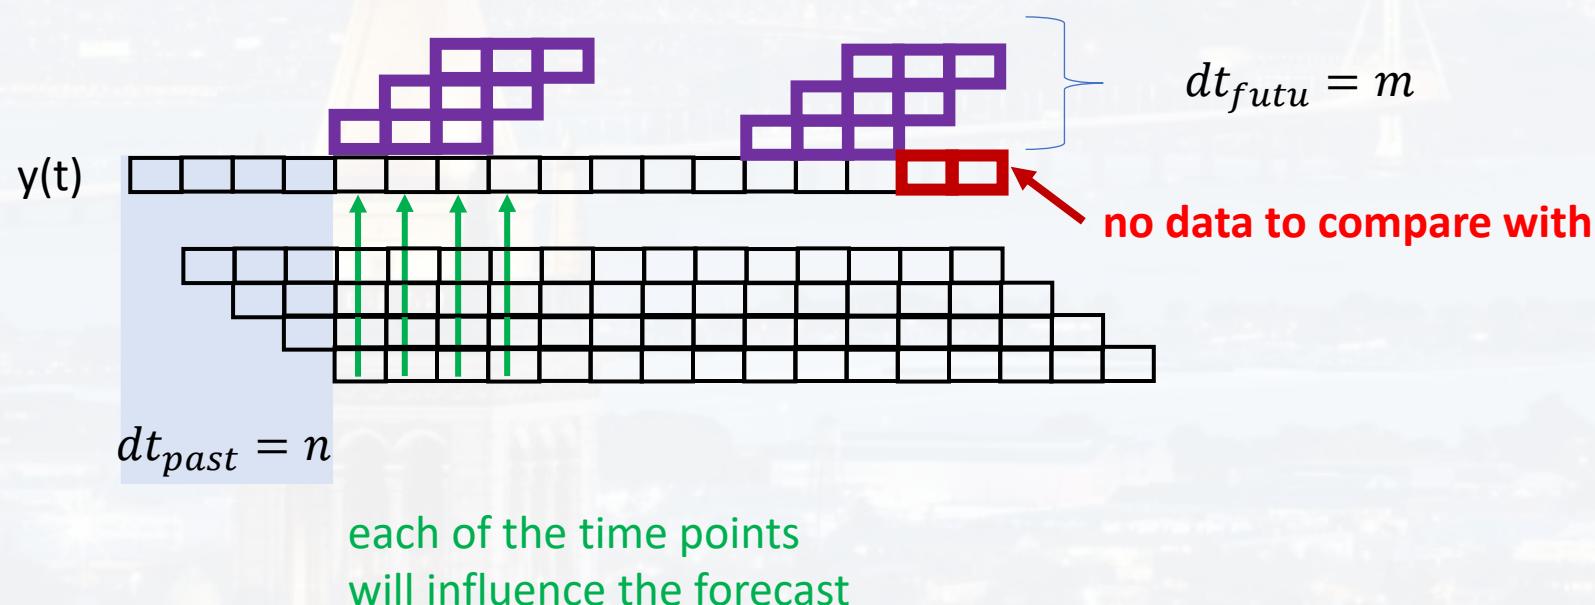
Let us first understand the logic:

length of training data is: `len[y(t)] - dtfutu - dtpast + 1`

predicting **m** steps in the future  
by **n** steps from the past

$$dt_{futu} = m$$
$$dt_{past} = n$$

predicting **m** steps of the future



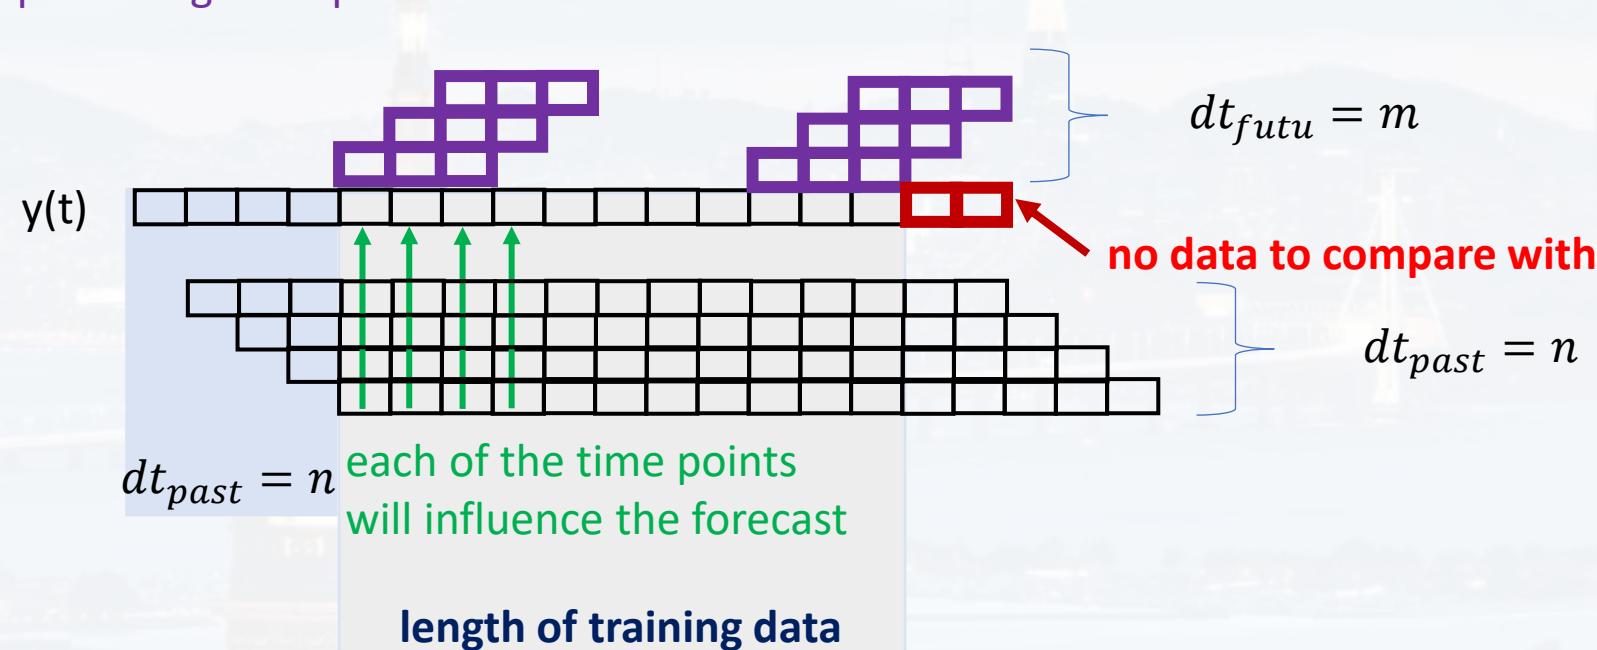


Let us first understand the logic:

predicting **m** steps in the future  
by **n** steps from the past

$$dt_{futu} = m$$
$$dt_{past} = n$$

predicting **m** steps of the future



```
X.shape = (len[y(t)] - dtfutu - dtpast + 1) x dtpast x nfeature(X)  
Y.shape = (len[y(t)] - dtfutu - dtpast + 1) x dtfutu x nfeature(Y)
```



predicting  $m$  steps of the future

$$dt_{futu} = m$$

$y(t)$

$dt_{past} = n$  each of the time points  
will influence the forecast

length of training data

predicting **m** steps in the future  
by **n** steps from the past

$$dt_{futu} = m$$

$$dt_{past} = n$$

no data to compare with

$$dt_{past} = n$$

$$X.\text{shape} = (\text{len}[y(t)] - dt_{futu} - dt_{past} + 1) \times dt_{past} \times n_{feature}(X)$$

$$Y.\text{shape} = (\text{len}[y(t)] - dt_{futu} - dt_{past} + 1) \times dt_{futu} \times n_{feature}(Y)$$

X	$dt_{past}$
	$\text{len}[y(t)] - dt_{futu} - dt_{past} + 1$
[0.23364871, 0.25531086, 0.29226308, 0.30477917, 0.34526381]	
[0.25531086, 0.29226308, 0.30477917, 0.34526381, 0.32876229]	
[0.29226308, 0.30477917, 0.34526381, 0.32876229, 0.34967038]	
[0.30477917, 0.34526381, 0.32876229, 0.34967038, 0.32374534]	
[0.34526381, 0.32876229, 0.34967038, 0.32374534, 0.34168462]	
[0.32876229, 0.34967038, 0.32374534, 0.34168462, 0.27602807]	
[0.34967038, 0.32374534, 0.34168462, 0.27602807, 0.2313527 ]	
[0.32374534, 0.34168462, 0.27602807, 0.2313527 , 0.20877584]	
[0.34168462, 0.27602807, 0.2313527 , 0.20877584, 0.16455034]	
[0.27602807, 0.2313527 , 0.20877584, 0.16455034, 0.11714726]	

Y	$dt_{futu}$
	$\text{len}[y(t)] - dt_{futu} - dt_{past} + 1$
[0.05263142, 0.10779498, 0.12263184],	
[0.10779498, 0.12263184, 0.12821065],	
[0.12263184, 0.12821065, 0.20806335],	
[0.12821065, 0.20806335, 0.2518744 ],	
[0.20806335, 0.2518744 , 0.28025766],	
[0.2518744 , 0.28025766, 0.27699119],	
[0.28025766, 0.27699119, 0.30965494],	
[0.27699119, 0.30965494, 0.37666627],	
[0.30965494, 0.37666627, 0.37879347],	
[0.37666627, 0.37879347, 0.36811853]]	



Let us first understand the logic:

Once, we have fitted the model: how do we apply the prediction?

```
PredY = model.predict(TestX)
```

```
(TestX.shape[0], dt_futu) = PredY.shape
```

X	$dt_{past}$
[0.23364871]	0.25531086, 0.29226308, 0.30477917, 0.34526381]
[0.25531086]	0.29226308, 0.30477917, 0.34526381, 0.32876229]
[0.29226308]	0.30477917, 0.34526381, 0.32876229, 0.34967038]
[0.30477917]	0.34526381, 0.32876229, 0.34967038, 0.32374534]
[0.34526381]	0.32876229, 0.34967038, 0.32374534, 0.34168462]
[0.32876229]	0.34967038, 0.32374534, 0.34168462, 0.27602807]
[0.34967038]	0.32374534, 0.34168462, 0.27602807, 0.2313527 ]
[0.32374534]	[0.34168462, 0.27602807, 0.2313527 , 0.20877584]
[0.34168462]	[0.27602807, 0.2313527 , 0.20877584, 0.16455034]
[0.27602807]	[0.2313527 , 0.20877584, 0.16455034, 0.11714726]

Y	$dt_{futu}$
[0.05263142]	0.10779498, 0.12263184]
[0.10779498]	0.12263184, 0.12821065]
[0.12263184]	0.12821065, 0.20806335]
[0.12821065]	[0.20806335, 0.2518744 ]
[0.20806335]	[0.2518744 , 0.28025766]
[0.2518744 ]	[0.28025766, 0.27699119]
[0.28025766]	[0.27699119, 0.30965494]
[0.27699119]	[0.30965494, 0.37666627]
[0.30965494]	[0.37666627, 0.37879347]
[0.37666627]	[0.37879347, 0.36811853]]

TestX[0,:,:0] should predict TestY[0,:,:0]

TestX[1,:,:0] should predict TestY[1,:,:0] etc



### Let us explore LSTM1.ipynb

```
n_neurons = 400  
batch_size = 128
```

```
model = Sequential()  
model.add(LSTM(n_neurons, input_shape = (dt_past, n_features),  
               activation = 'tanh'))  
model.add(Dense(dt_futu))
```

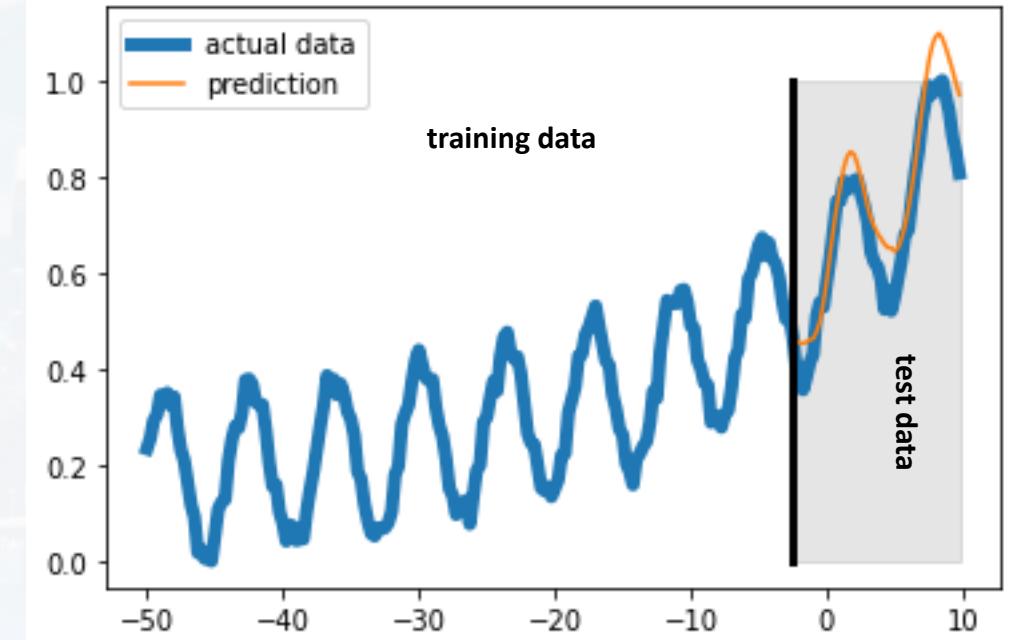
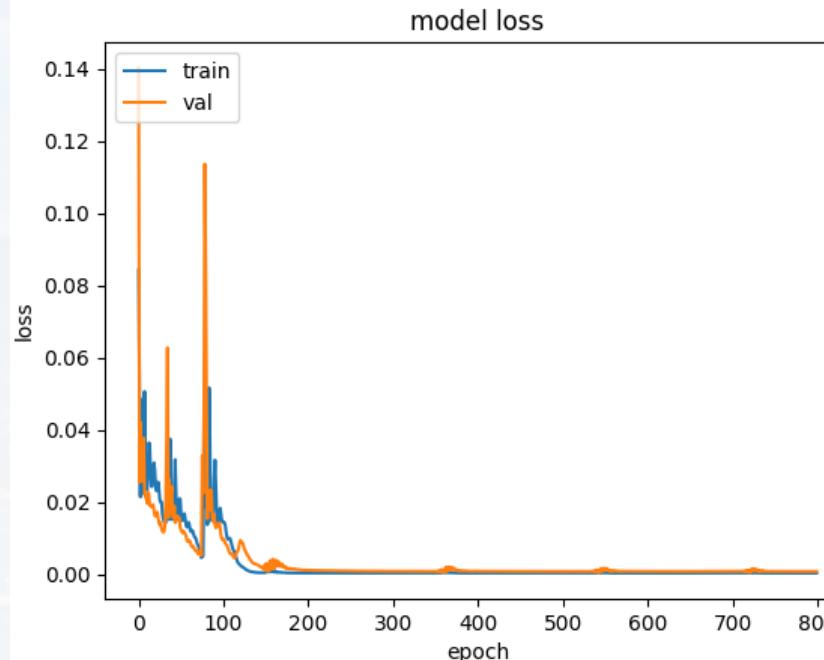
```
opt = optimizers.Adam()  
model.compile(loss = 'mean_squared_error', optimizer = opt)
```

```
model.summary()
```

Model: "sequential"		
Layer (type)	Output Shape	Param #
lstm (LSTM)	(None, 400)	643200
dense (Dense)	(None, 8)	3208
<hr/>		
Total params: 646408 (2.47 MB)		
Trainable params: 646408 (2.47 MB)		
Non-trainable params: 0 (0.00 Byte)		

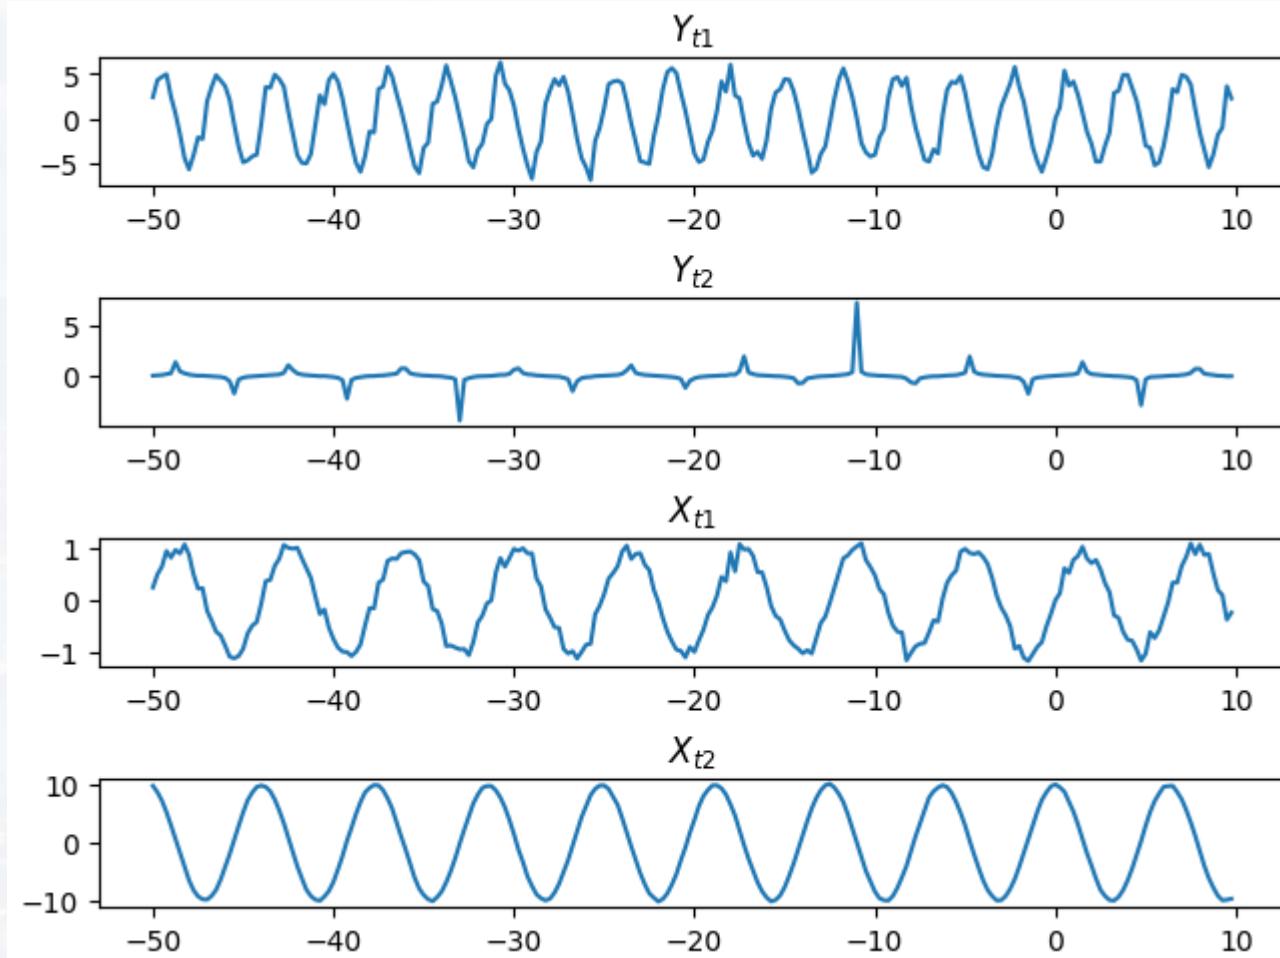


Let us explore LSTM1.ipynb



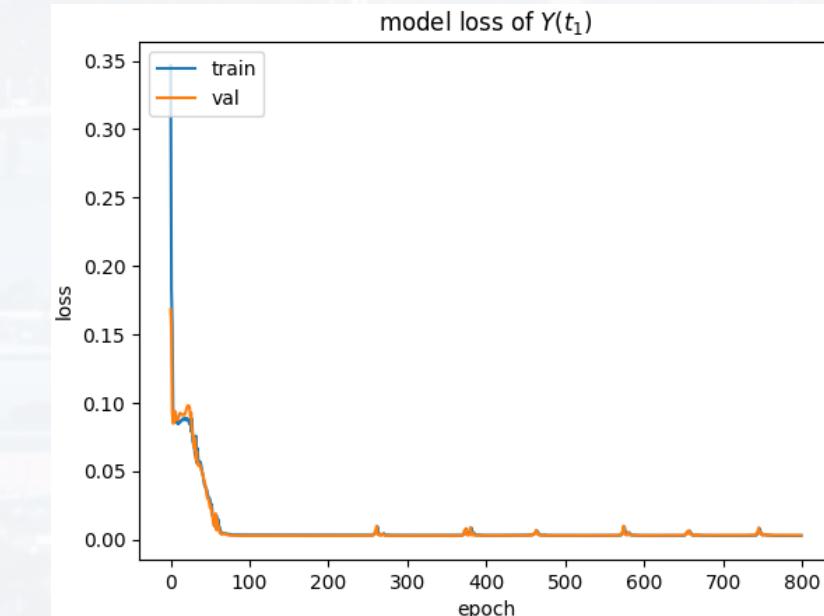
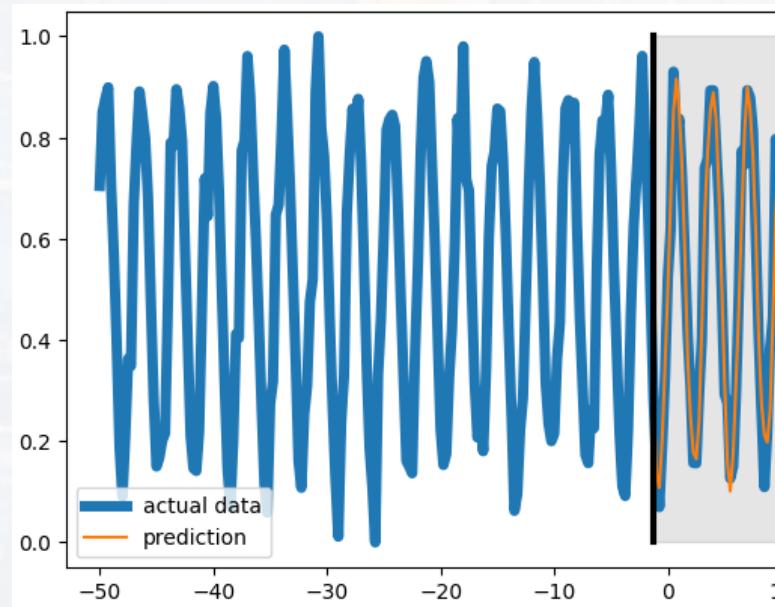
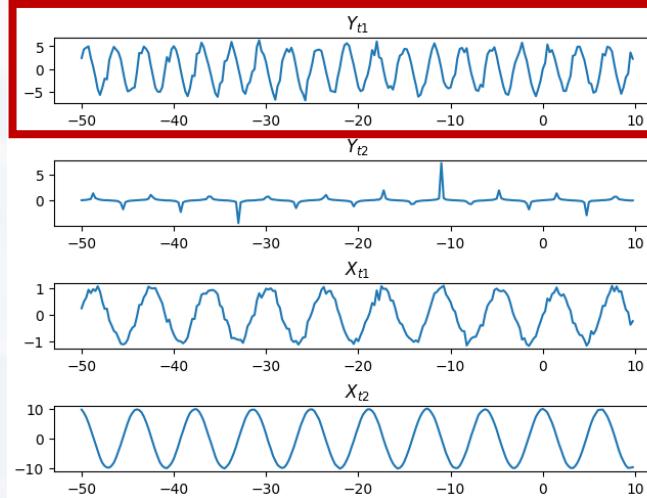


Explore LSTMII.ipynb for a multivariate, multi feature time series:



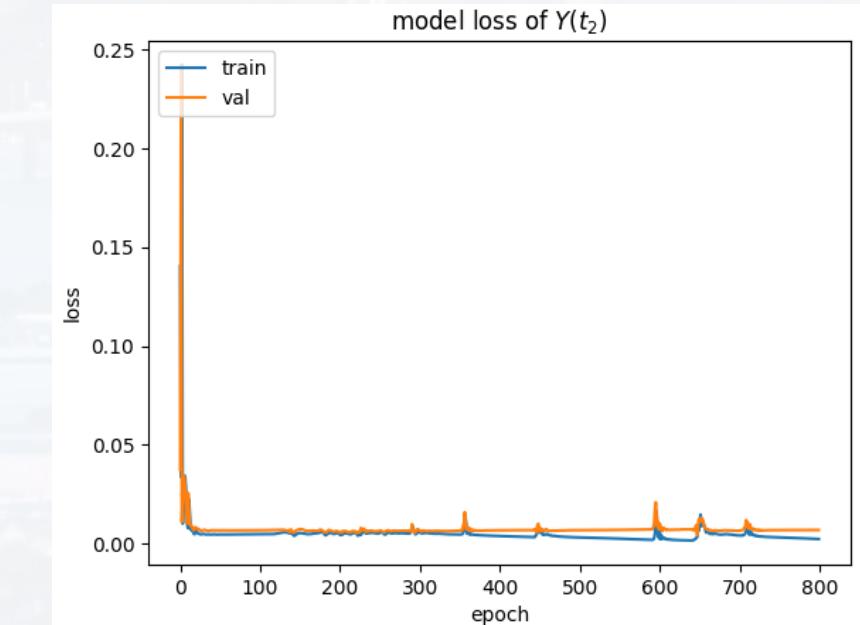
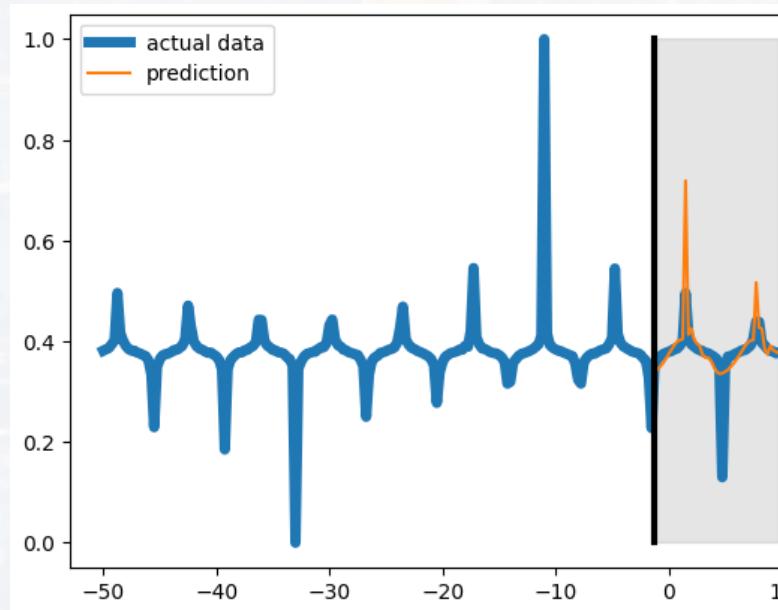
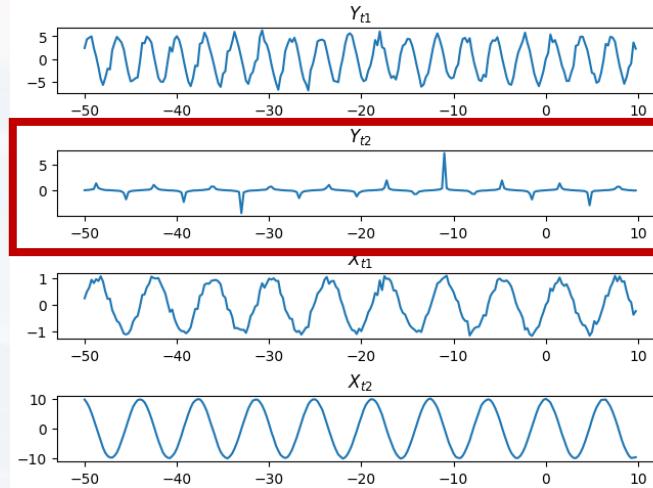


Explore LSTMII.ipynb for a multivariate, multi feature time series:





Explore LSTMII.ipynb for a multivariate, multi feature time series:





Thank you very much for your attention!



<https://www.analyticsvidhya.com>