

Lecture 12:

Long Short-Term Memory Networks (LSTMs) – Part II



Markus Hohle
University California, Berkeley

Machine Learning Algorithms
MSSE 277B, 3 Units



Lecture 1: Course Overview and Introduction to Machine Learning

Lecture 2: Bayesian Methods in Machine Learning

classic ML tools & algorithms

Lecture 3: Dimensionality Reduction: Principal Component Analysis

Lecture 4: Linear and Non-linear Regression and Classification

Lecture 5: Unsupervised Learning: K-Means, GMM, Trees

Lecture 6: Adaptive Learning and Gradient Descent Optimization Algorithms

Lecture 7: Introduction to Artificial Neural Networks - The Perceptron

ANNs/AI/Deep Learning

Lecture 8: Introduction to Artificial Neural Networks - Building Multiple Dense Layers

Lecture 9: Convolutional Neural Networks (CNNs) - Part I

Lecture 10: CNNs - Part II

Lecture 11: Recurrent Neural Networks (RNNs) and Long Short-Term Memory (LSTMs)

Lecture 12: Combining LSTMs and CNNs

Lecture 13: Running Models on GPUs and Parallel Processing

Lecture 14: Project Presentations

Lecture 15: Transformer

Lecture 16: GNN



<https://www.analyticsvidhya.com>



Outline

- LSTM for Classification
- Bidirectional LSTMs
- Stacked LSTMs
- LSTM + CNN



<https://www.analyticsvidhya.com>



Outline

- LSTM for Classification
- Bidirectional LSTMs
- Stacked LSTMs
- LSTM + CNN



minimal model:

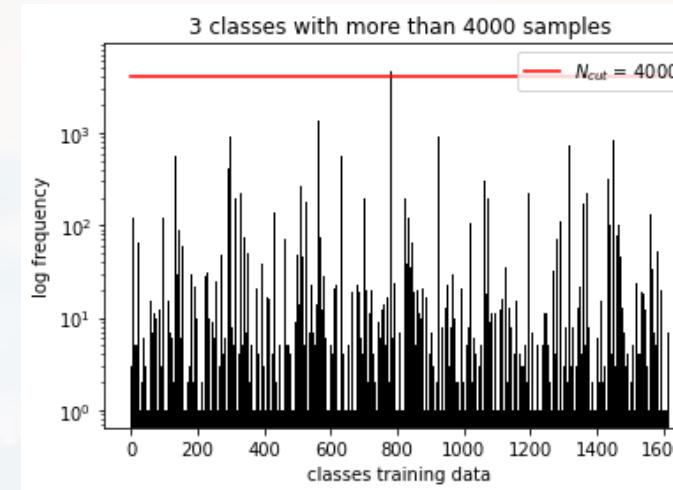
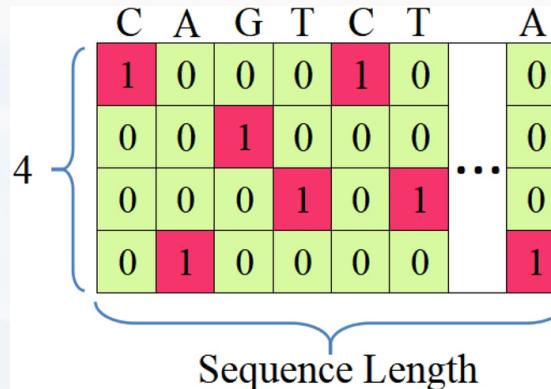
```
[N_samples, LengthSeq, N_features] = X.shape  
[N_samples, N_classes] = Y.shape
```

Y is one-hot encoded

```
model = Sequential()  
  
model.add(LSTM(n_neurons, activation = 'tanh',  
               input_shape = (LengthSeq , N_features)))  
  
model.add(Dense(N_classes, activation = 'softmax'))  
  
opt = optimizers.Adam()  
model.compile(loss = 'categorical_crossentropy', optimizer = opt,  
              metrics = [ 'accuracy' ])  
  
model.summary()
```

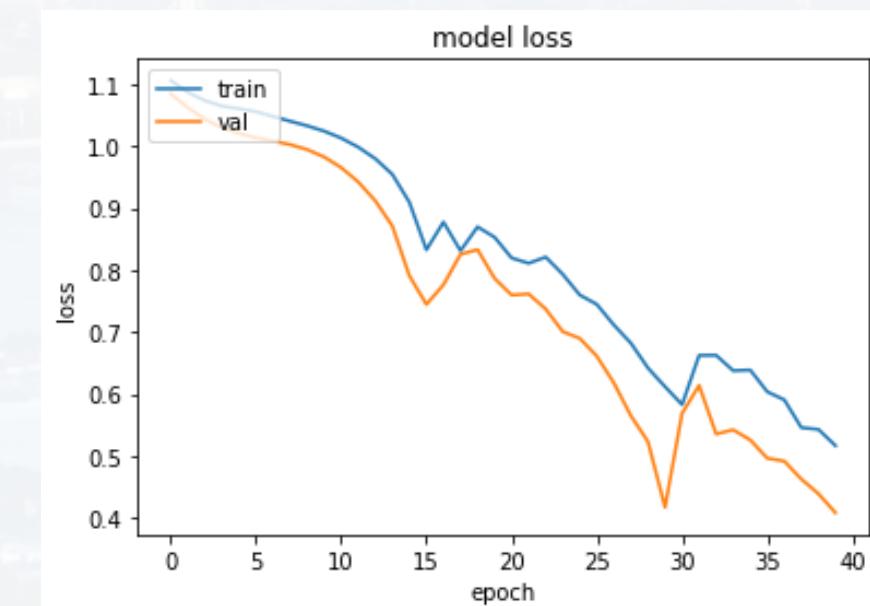
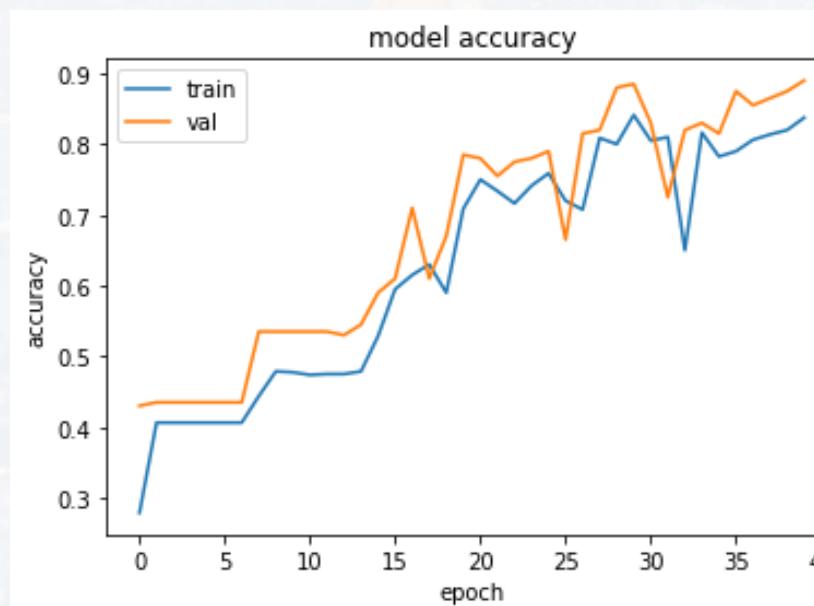


barcode example



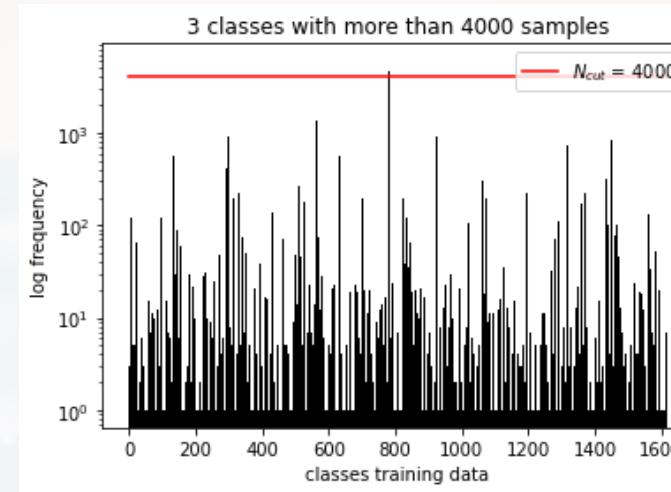
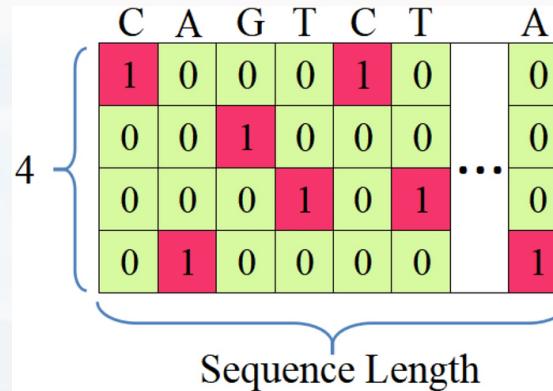
for computational reasons:

- **three** classes
- **1k** samples total
- sequences cut to length **500**



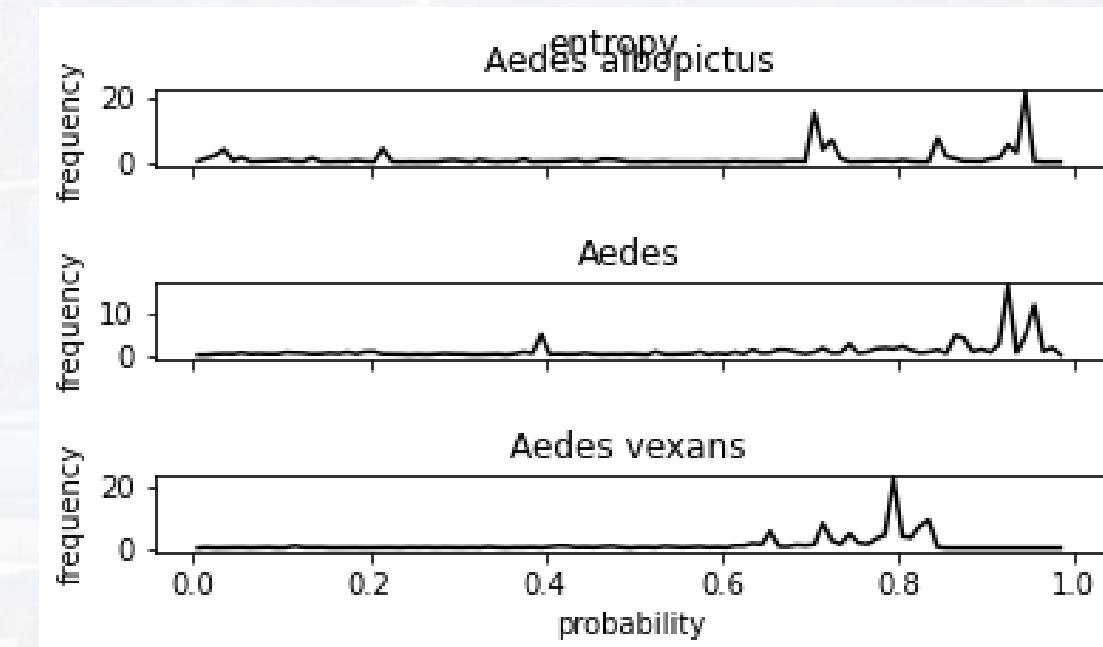
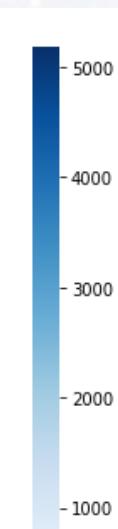
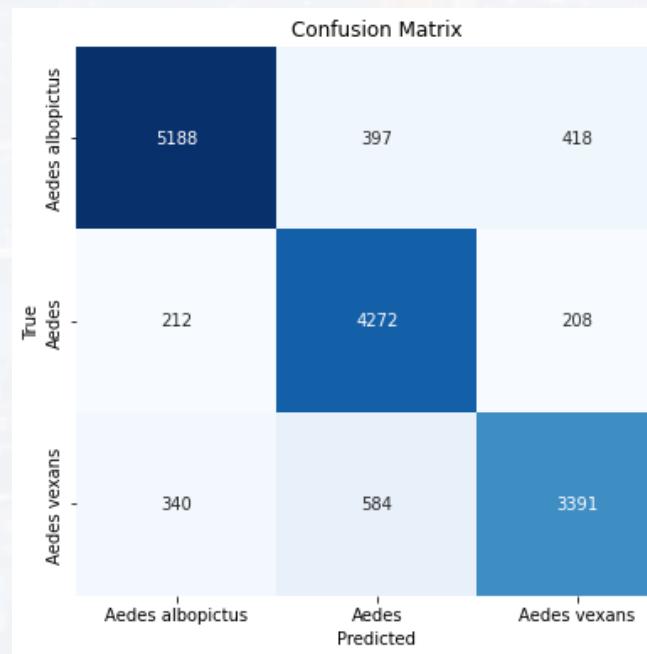


barcode example



for computational reasons:

- **three** classes
- **1k** samples total
- sequences cut to length **500**





<https://www.analyticsvidhya.com>

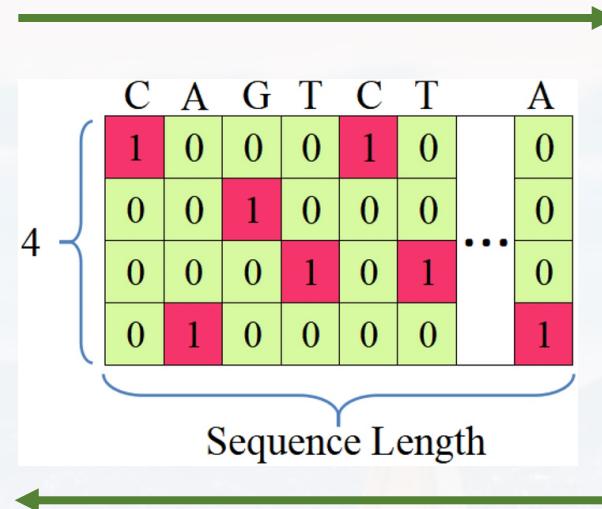


Outline

- LSTM for Classification
- Bidirectional LSTMs
- Stacked LSTMs
- LSTM + CNN



sometimes, sequences can be read from two directions :



```
from keras.layers import Bidirectional
```

```
model = Sequential()
model.add(Bidirectional(LSTM(n_neurons, activation = 'tanh'),\
                      input_shape = (dt_past, n_features)))
model.add(Dense(dt_futu))

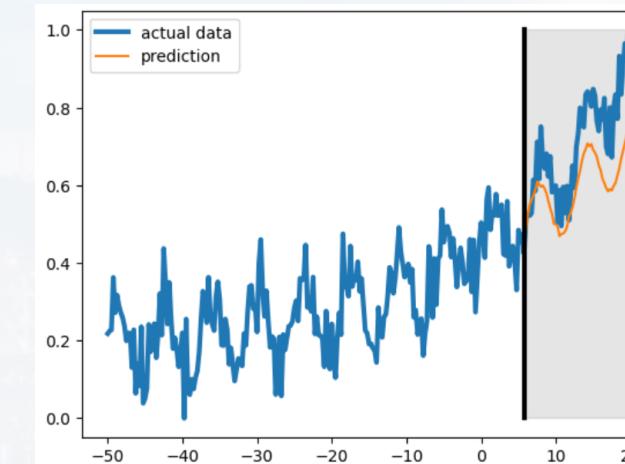
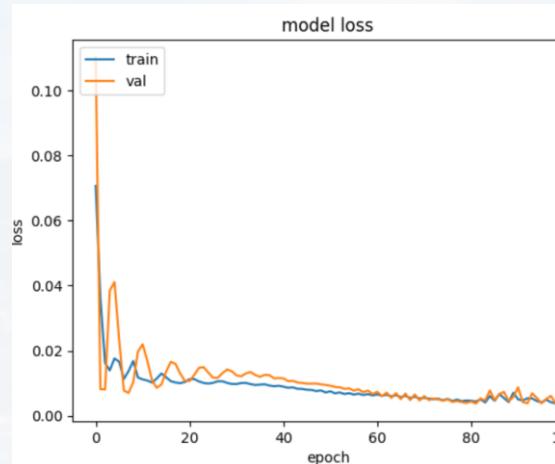
opt = optimizers.Adam()
model.compile(loss = 'mean_squared_error', optimizer = opt)

model.summary()
```

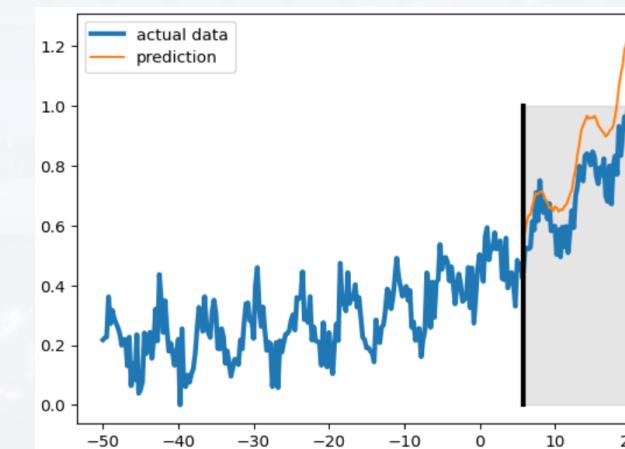
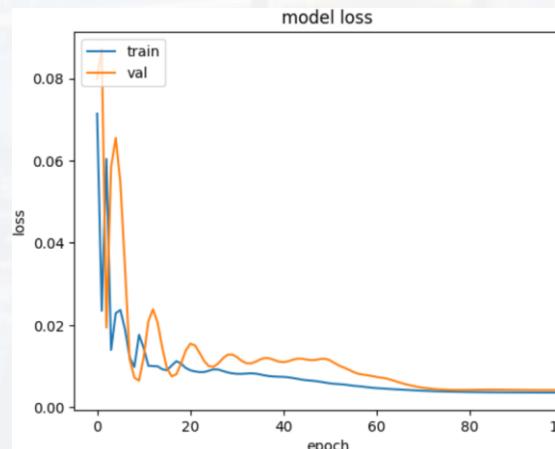


```
model = Sequential()
model.add(Bidirectional(LSTM(n_neurons, activation = 'tanh'), input_shape = (dt_past, n_features)))
model.add(Dense(dt_futu))
opt = optimizers.Adam()
model.compile(loss = 'mean_squared_error', optimizer = opt)
model.summary()
```

vanilla



bidirectional



Layer (type)	Output Shape	Param #
<hr/>		
bidirectional (Bidirectional)	(None, 800)	1286400
<hr/>		
dense_1 (Dense)	(None, 8)	6408
<hr/>		
Total params:	1292808 (4.93 MB)	
Trainable params:	1292808 (4.93 MB)	
Non-trainable params:	0 (0.00 Byte)	



<https://www.analyticsvidhya.com>

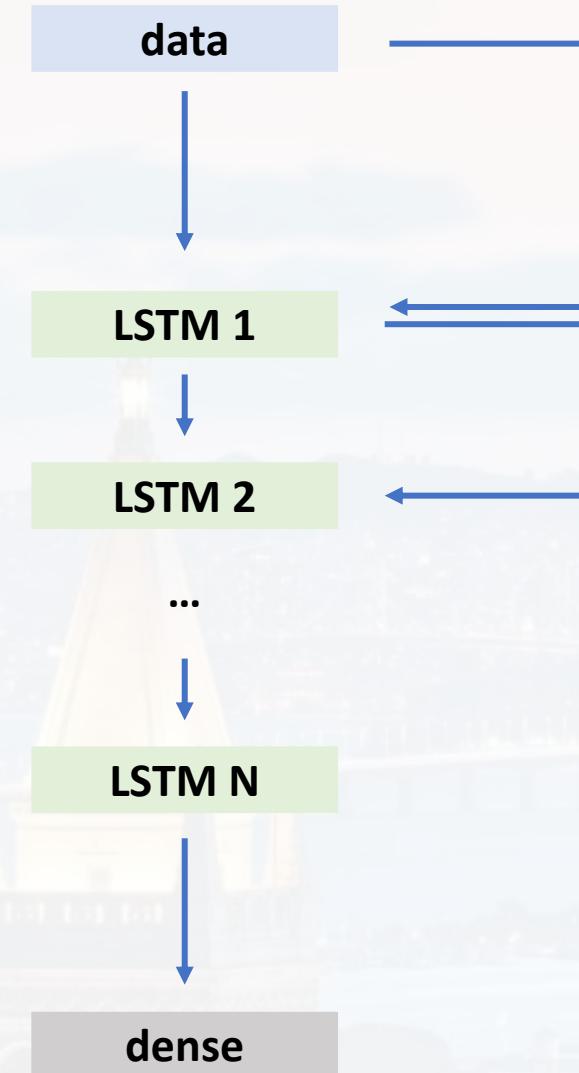


Outline

- LSTM for Classification
- Bidirectional LSTMs
- Stacked LSTMs
- LSTM + CNN



idea:



shape: (`len(y(t)) - dt_past - dt_futu + 1`
x `dt_past` x features)

shape: (`len(y(t)) - dt_past - dt_futu + 1`
x `dt_past` x hidden state)

`return_sequences = True`



data



LSTM 1



LSTM 2

...

LSTM N

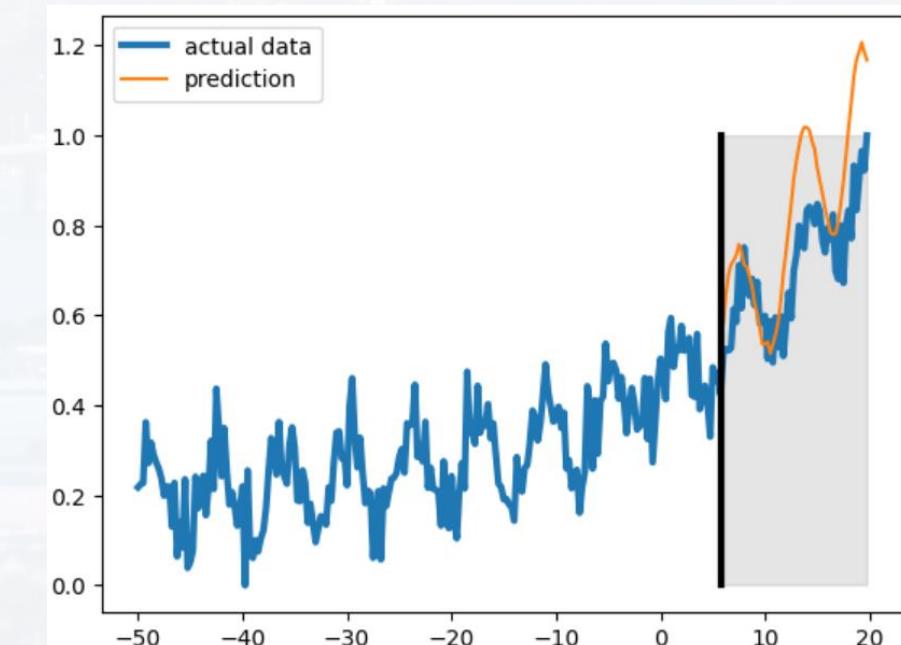
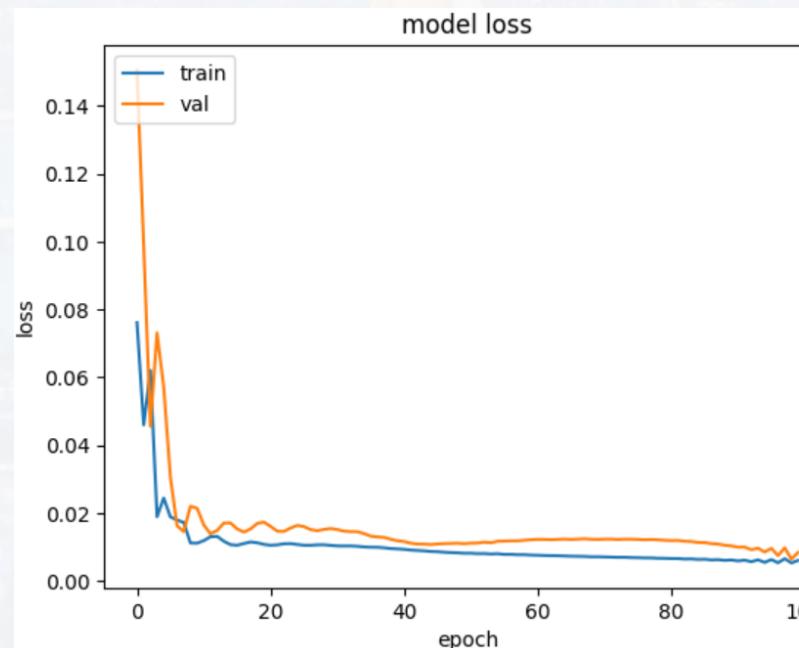


```
model = Sequential()  
  
model.add(LSTM(n_neurons, activation = 'tanh',\  
               return_sequences = True, input_shape = (dt_past, n_features)))  
  
model.add(LSTM(2*n_neurons, activation = 'relu',\  
               return_sequences = True))  
  
model.add(LSTM(n_neurons, activation = 'relu'))  
  
model.add(Dense(dt_futu))  
  
opt = optimizers.Adam()  
model.compile(loss = 'mean_squared_error', optimizer = opt)  
  
model.summary()
```

all LSTMs, **except the last** stack needs
return_sequences = True



Layer (type)	Output Shape	Param #
=====		
lstm_2 (LSTM)	(None, 20, 400)	643200
lstm_3 (LSTM)	(None, 20, 800)	3843200
lstm_4 (LSTM)	(None, 400)	1921600
dense_2 (Dense)	(None, 8)	3208
=====		
Total params: 6411208 (24.46 MB)		
Trainable params: 6411208 (24.46 MB)		
Non-trainable params: 0 (0.00 Byte)		





<https://www.analyticsvidhya.com>

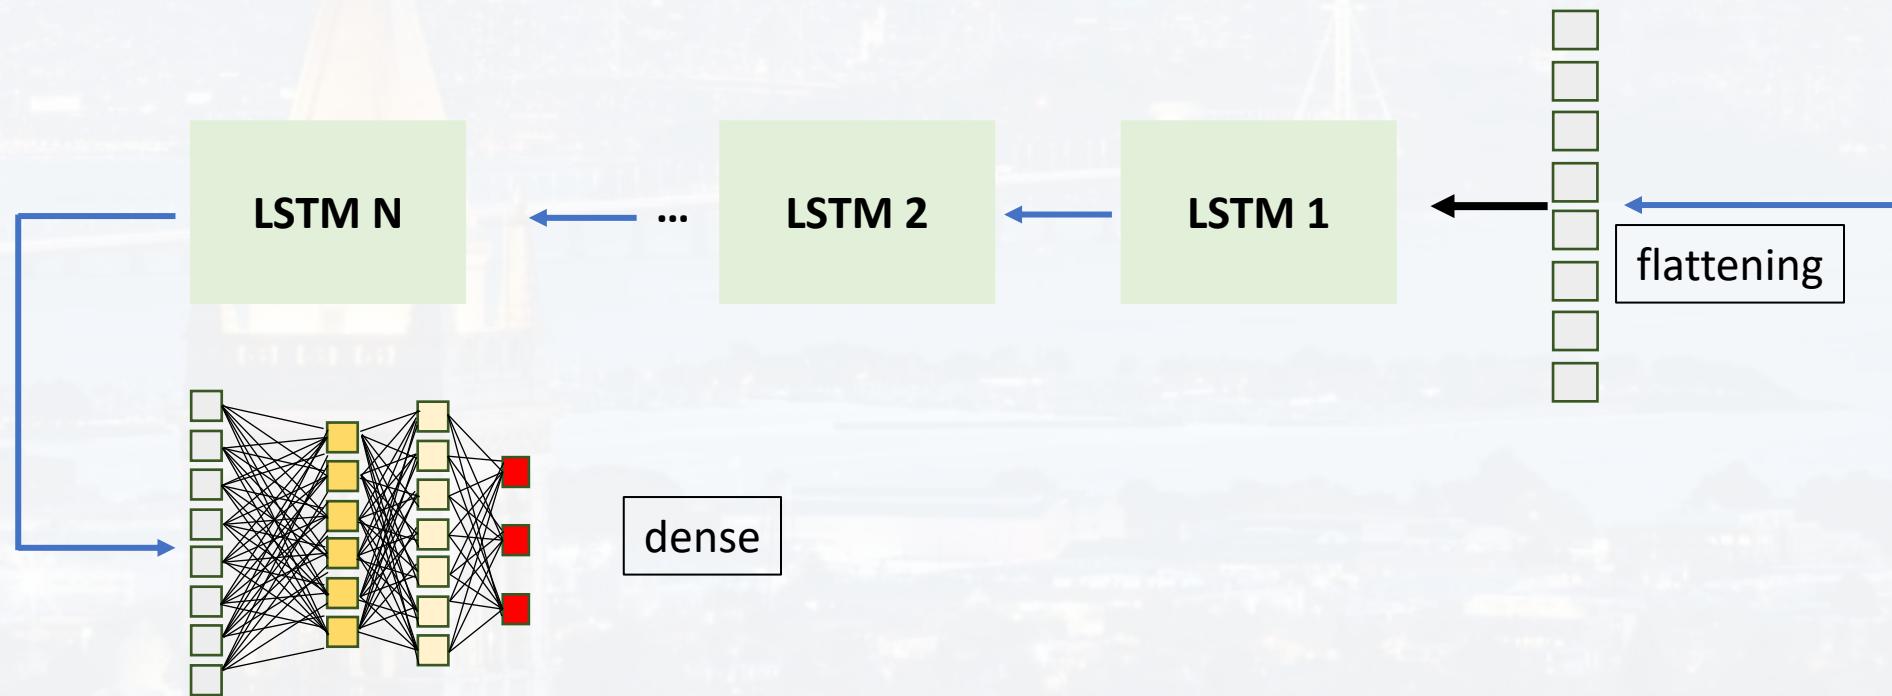
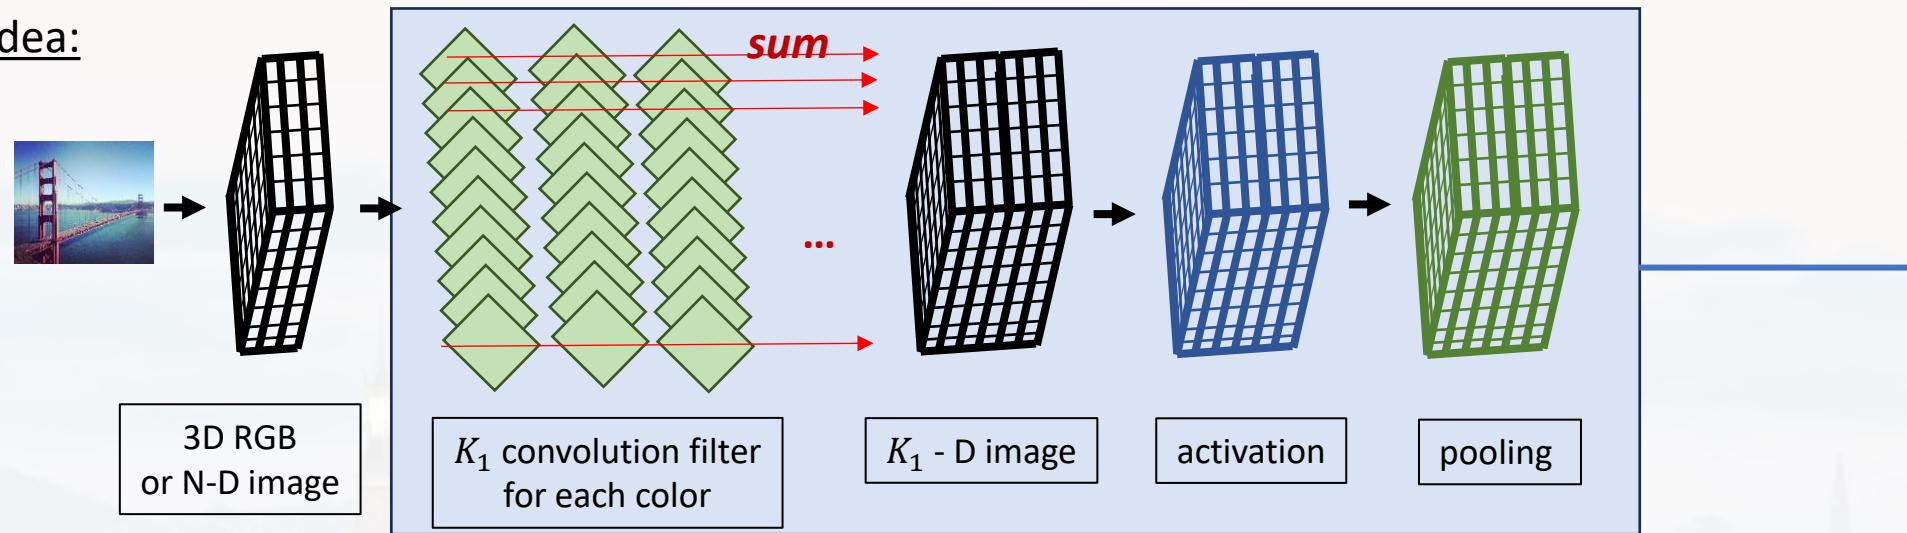


Outline

- LSTM for Classification
- Bidirectional LSTMs
- Stacked LSTMs
- LSTM + CNN



idea:

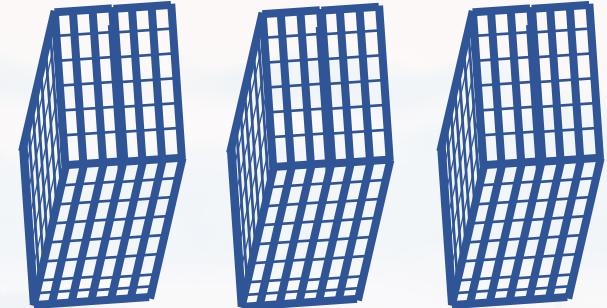




idea:

input expected by CNN (images):

$(N_images, N_x, N_y, N_color)$

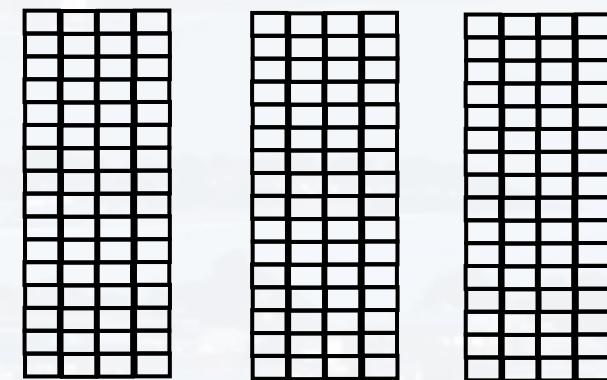


input expected by CNN (videos):

$(N_videos, N_t, N_x, N_y, N_color)$

input expected by LSTM (sequences):

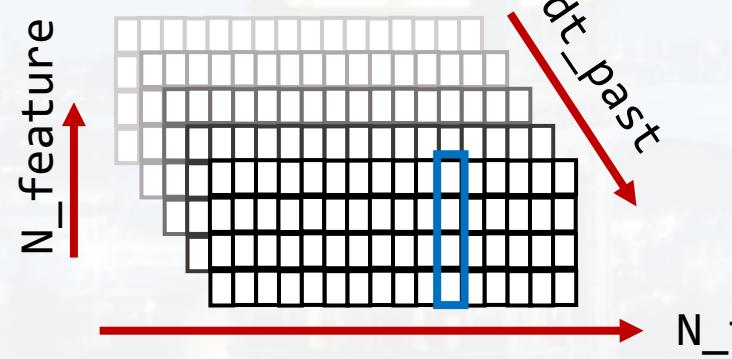
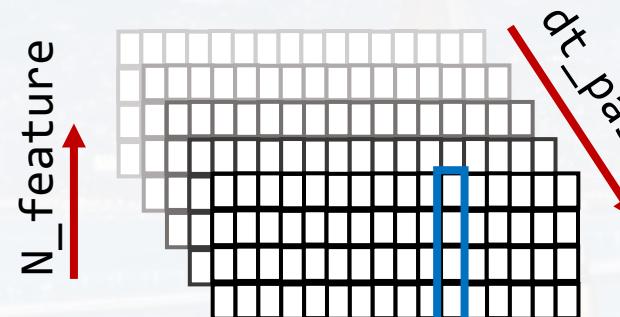
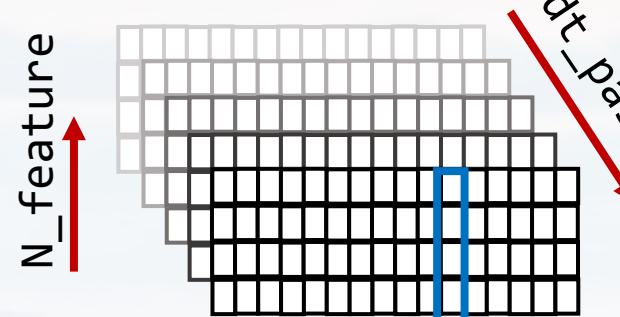
$(N_sequences, N_t, N_feature)$



None



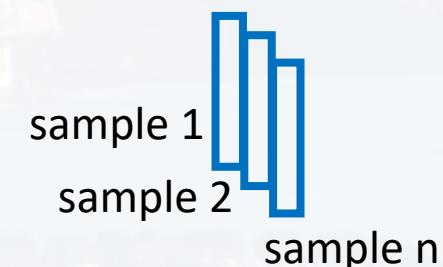
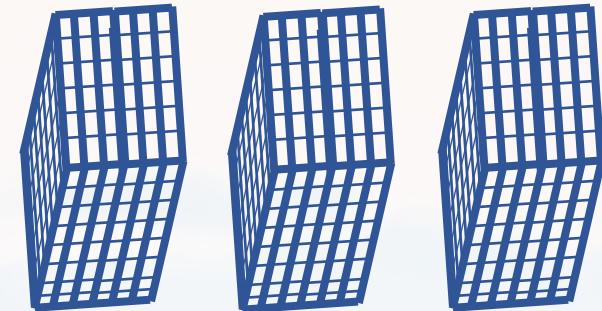
$(N_{\text{images}}, N_x, N_y, N_{\text{color}})$



sample 1

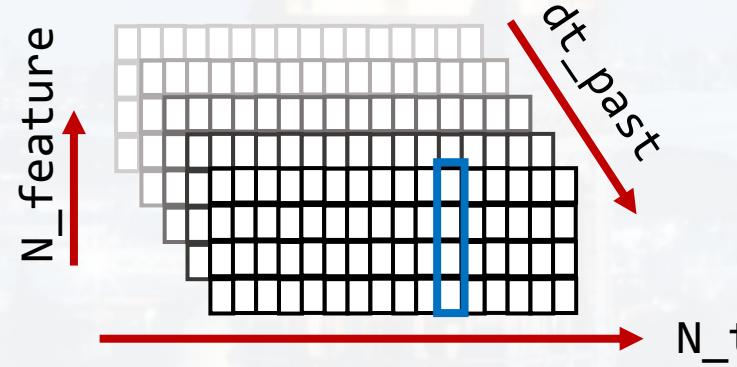
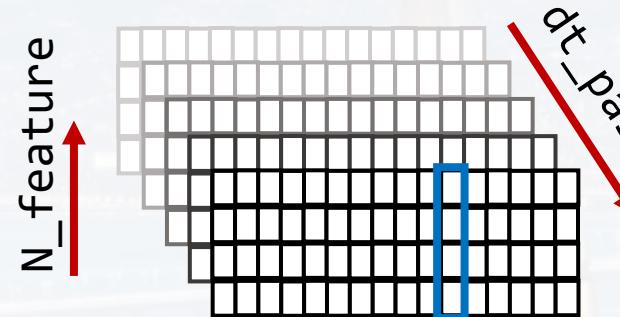
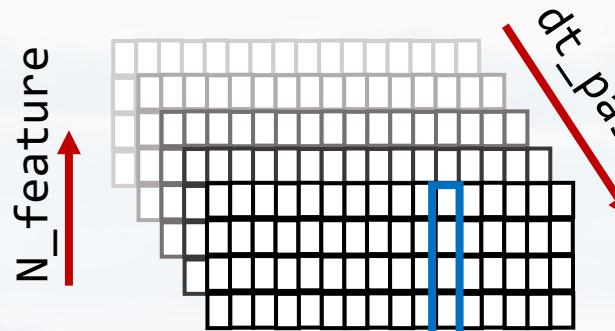
sample 2

sample n





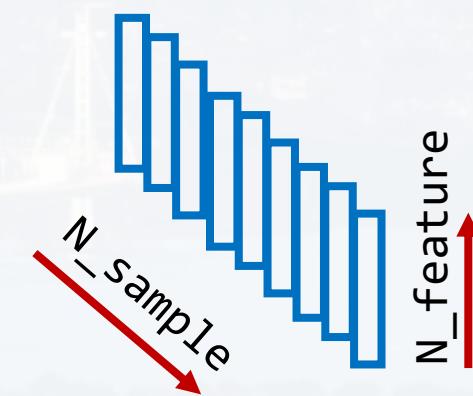
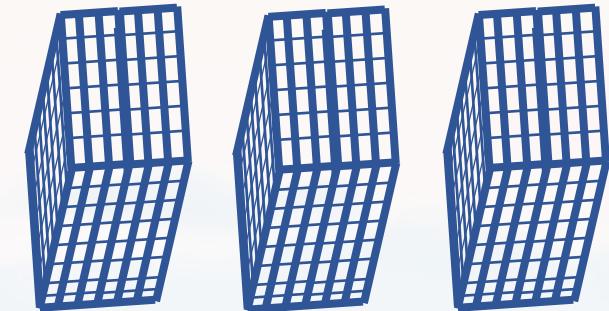
$(N_{\text{images}}, N_x, N_y, N_{\text{color}})$



sample 1

sample 2

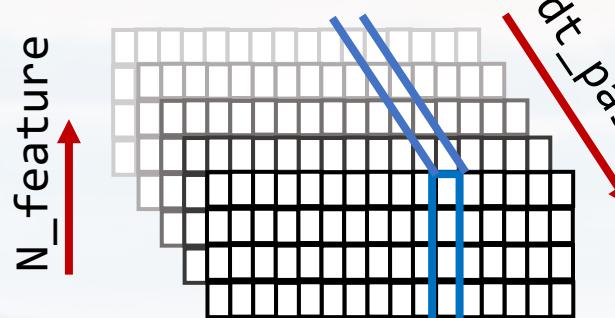
sample n



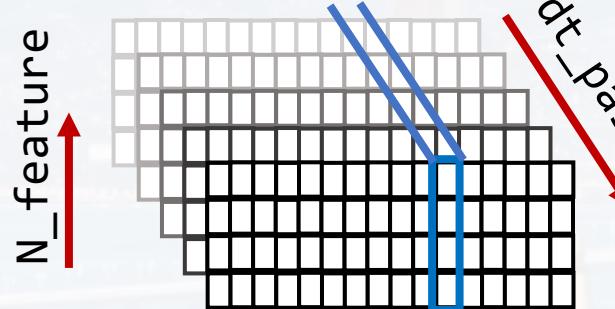
for **one** timepoint t



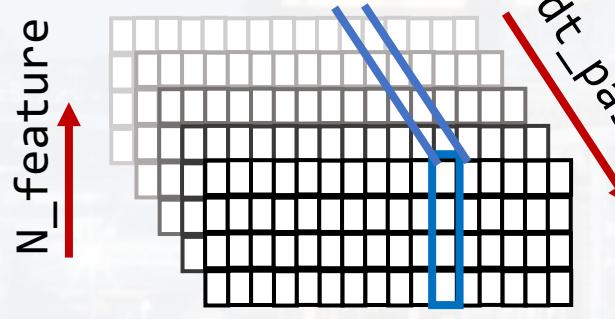
$(N_{\text{images}}, N_x, N_y, N_{\text{color}})$



sample 1

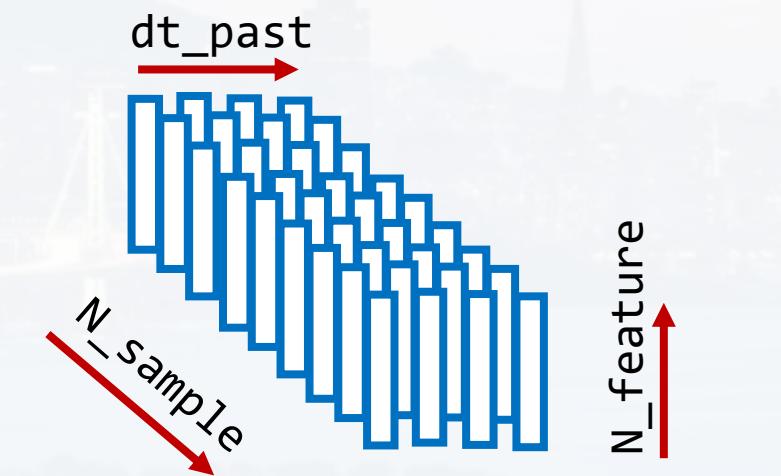
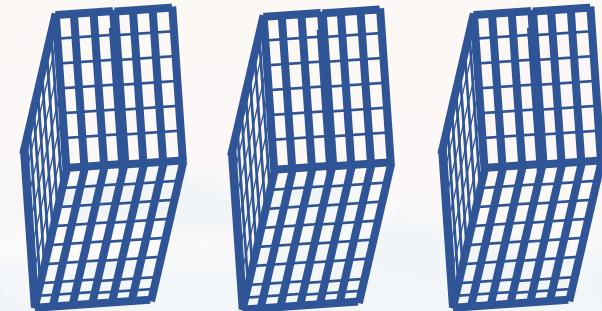


sample 2

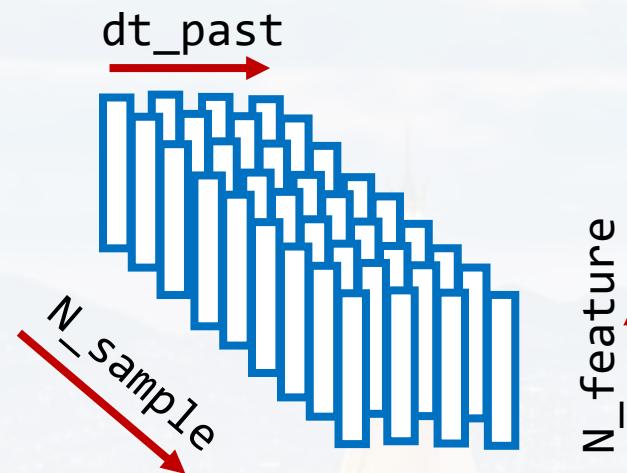


N_t

sample n

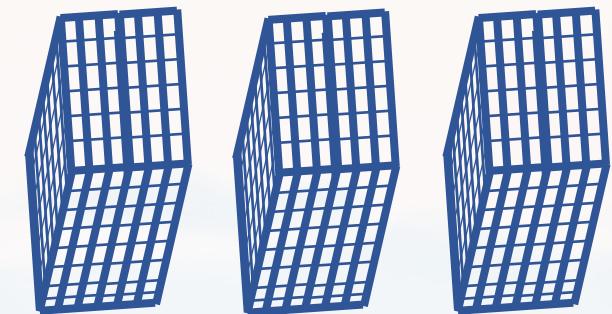


for **one** timepoint t



$(N_images, N_x, N_y, N_color)$

for **one** timepoint t



regression: one sample of $N_features$ and dt_past

```
X = X.reshape((X.shape[0], N_samples, dt_past, n_features))
```



```
X = X.reshape((X.shape[0], N_samples, dt_past, n_features))
```

```
model = Sequential()
```

```
model.add(TimeDistributed(Conv1D(filters = 64, kernel_size = 3,\n                           activation = 'relu').,\n                           input_shape = (None, dt_past, n_features))))
```

```
model.add(TimeDistributed(MaxPooling1D(pool_size = 2)))
```

```
model.add(TimeDistributed(Flatten()))
```

```
model.add(LSTM(n_neurons, input_shape = (dt_past, n_features),\n               activation = 'tanh'))
```

```
model.add(Dense(dt_futu))
```

```
opt = optimizers.Adam()
```

```
model.compile(loss = 'mean_squared_error', optimizer = opt)
```

```
model.summary()
```

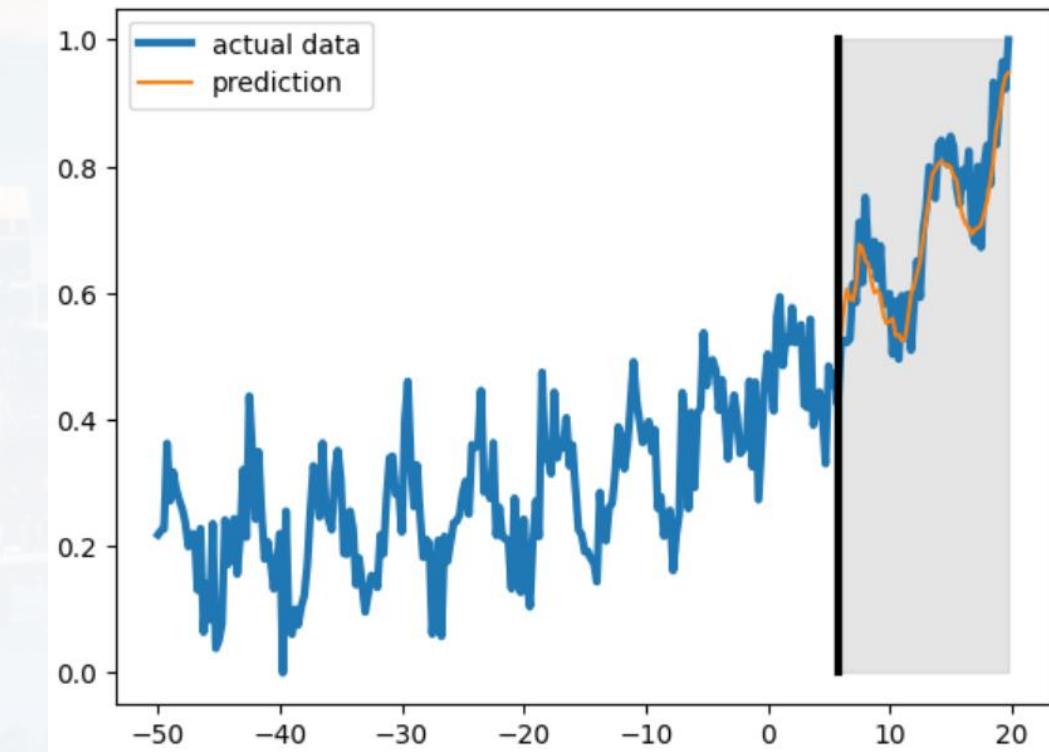
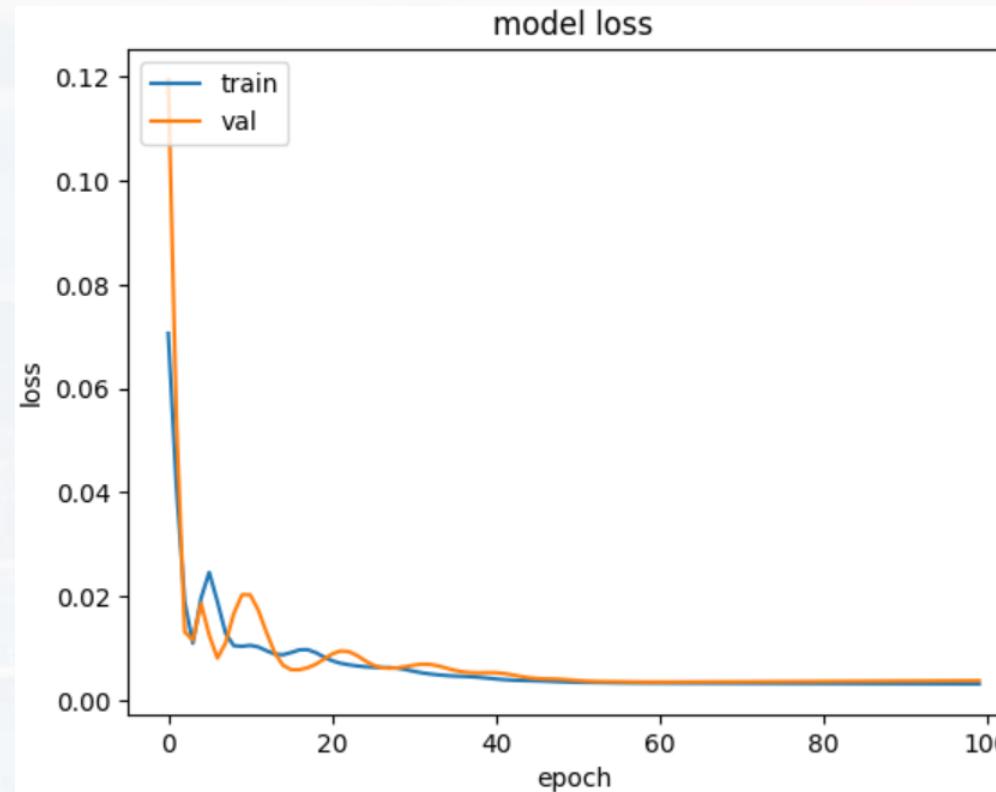
1D filter along time coordinate

takes care of
maintaining
matrix orientation

actual input is (None, None, dt_past, n_features)

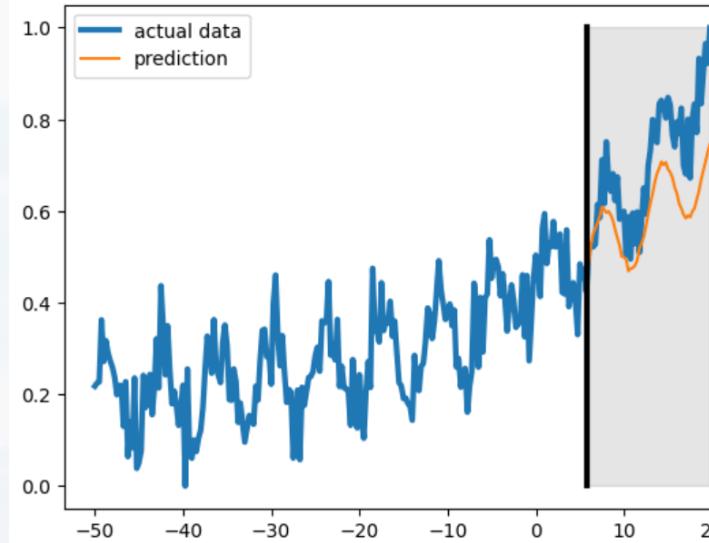


Layer (type)	Output Shape	Param #
=====		
time_distributed (TimeDistributed)	(None, None, 18, 64)	256
		actual input is (None, None, dt_past, n_features)
time_distributed_1 (TimeDistributed)	(None, None, 9, 64)	0
time_distributed_2 (TimeDistributed)	(None, None, 576)	0
lstm_5 (LSTM)	(None, 400)	1563200
dense_3 (Dense)	(None, 8)	3208
=====		
Total params: 1566664 (5.98 MB)		
Trainable params: 1566664 (5.98 MB)		
Non-trainable params: 0 (0.00 Byte)		

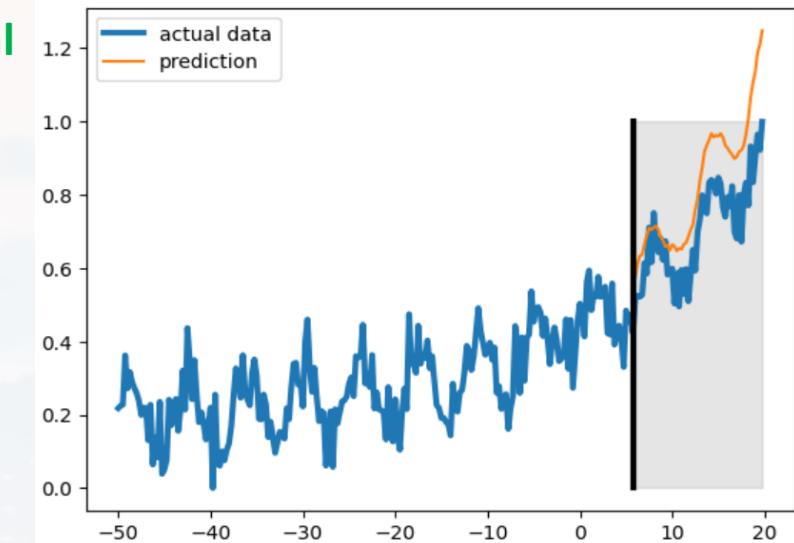




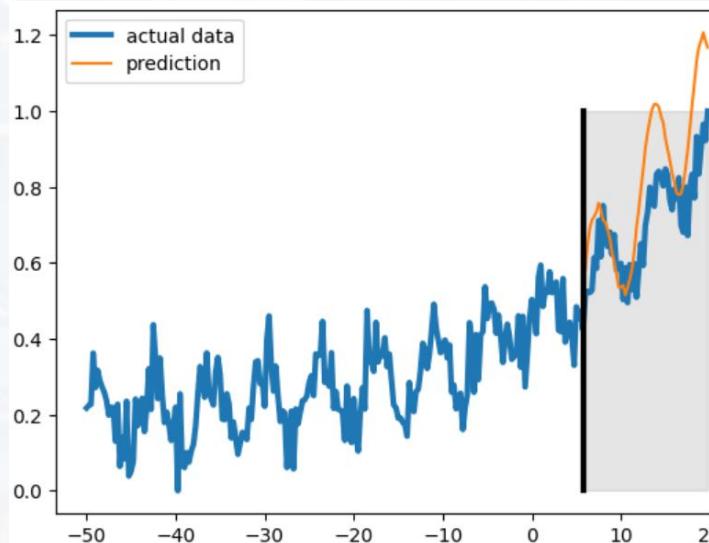
vanilla



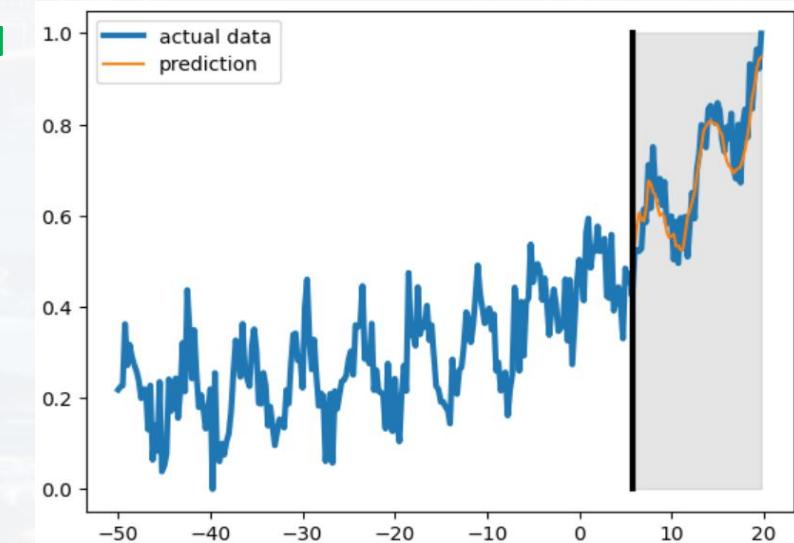
bidirectional



stacked



LSTM + CNN





classification: N samples of N_features and dt_past = Length_Seq

```
[N_sample, LengthSeq, N_features] = X.shape  
model = Sequential()  
model.add(Conv1D(filters = 64, kernel_size = 3, activation = 'relu', \  
                 input_shape = (LengthSeq, N_features)))
```

1D filter along time coordinate
= LengthSeq

```
model.add(MaxPooling1D(pool_size = 2))
```

```
model.add(LSTM(n_neurons, activation = 'tanh'))
```

```
model.add(Dense(Nclass, activation = 'softmax'))
```

```
opt = optimizers.Adam()  
model.compile(loss = 'categorical_crossentropy', optimizer = opt, \  
              metrics = ['accuracy'])
```

```
model.summary()
```



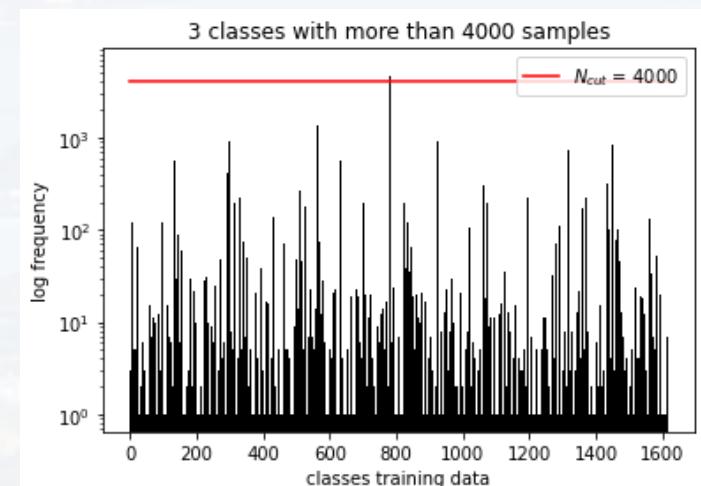
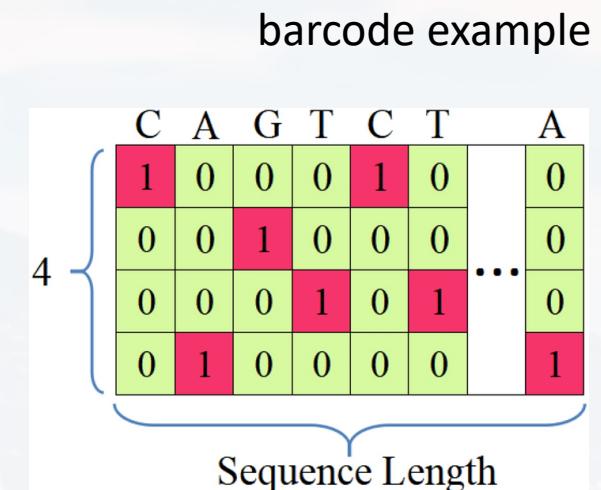
classification: N samples of N_features and dt_past = Length_Seq

[N_sample, LengthSeq, N_features] = X.shape

Layer (type)	Output Shape	Param #
<hr/>		
conv1d_8 (Conv1D)	(None, 498, 64)	832
max_pooling1d_7 (MaxPooling1D)	(None, 249, 64)	0
lstm_7 (LSTM)	(None, 100)	66000
dense_4 (Dense)	(None, 3)	303
<hr/>		
Total params: 67135 (262.25 KB)		
Trainable params: 67135 (262.25 KB)		
Non-trainable params: 0 (0.00 Byte)		

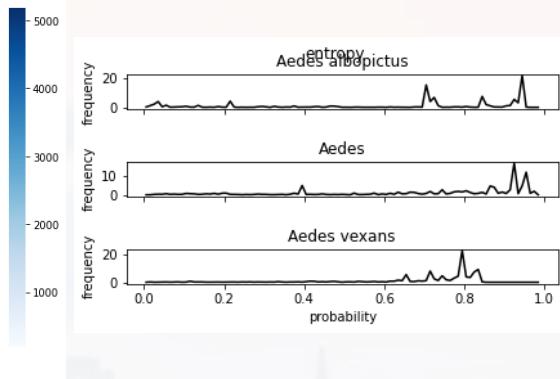
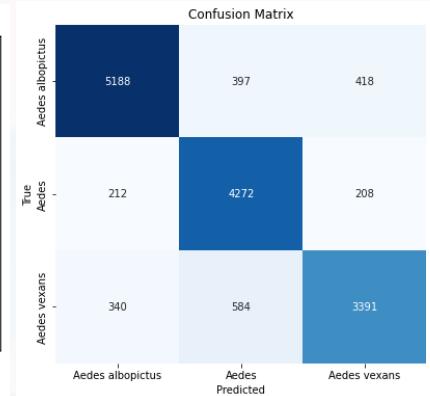
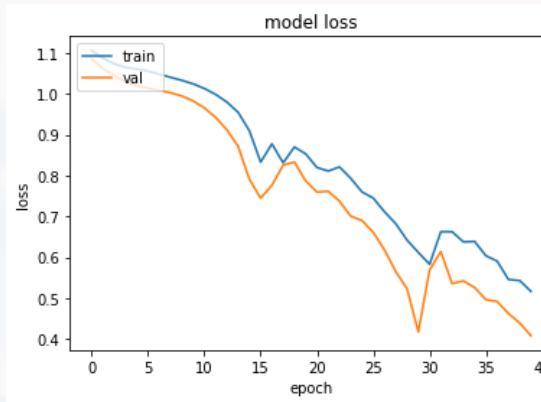
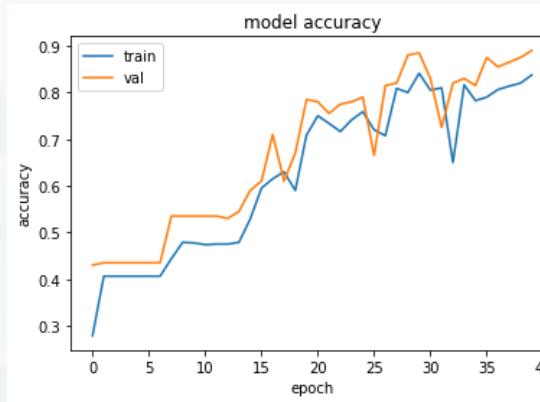
for computational reasons:

- **three** classes
- **1k** samples total
- sequences cut to length **500**

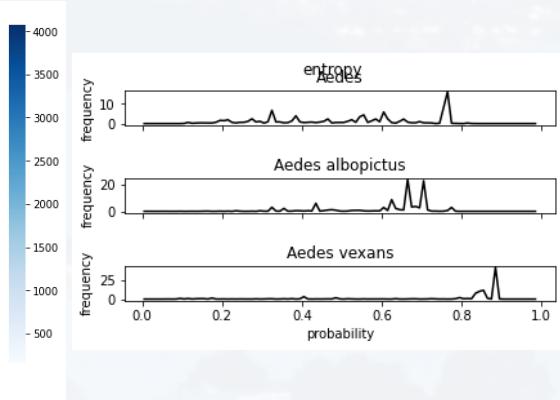
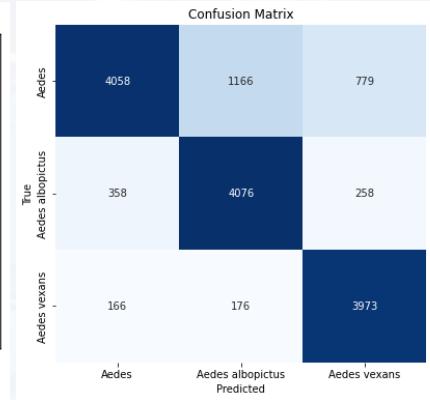
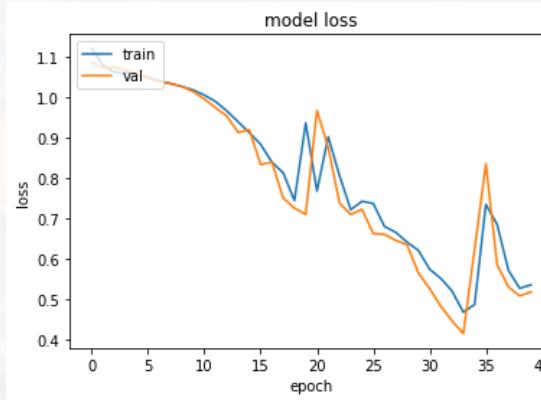
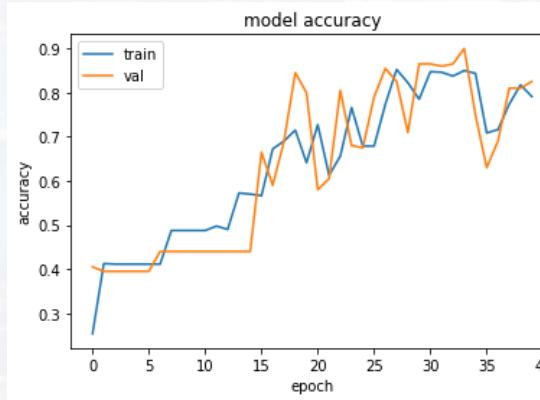




LSTM



LSTM+CNN





Thank you for your attention!

