

## Lecture 15:

# Language Models and Transformer



Markus Hohle  
University California, Berkeley

Machine Learning Algorithms  
MSSE 277B, 3 Units



## Lecture 1: Course Overview and Introduction to Machine Learning

Lecture 2: Bayesian Methods in Machine Learning

classic ML tools & algorithms

Lecture 3: Dimensionality Reduction: Principal Component Analysis

Lecture 4: Linear and Non-linear Regression and Classification

Lecture 5: Unsupervised Learning: K-Means, GMM, Trees

Lecture 6: Adaptive Learning and Gradient Descent Optimization Algorithms

Lecture 7: Introduction to Artificial Neural Networks - The Perceptron

ANNs/AI/Deep Learning

Lecture 8: Introduction to Artificial Neural Networks - Building Multiple Dense Layers

Lecture 9: Convolutional Neural Networks (CNNs) - Part I

Lecture 10: CNNs - Part II

Lecture 11: Recurrent Neural Networks (RNNs) and Long Short-Term Memory (LSTMs)

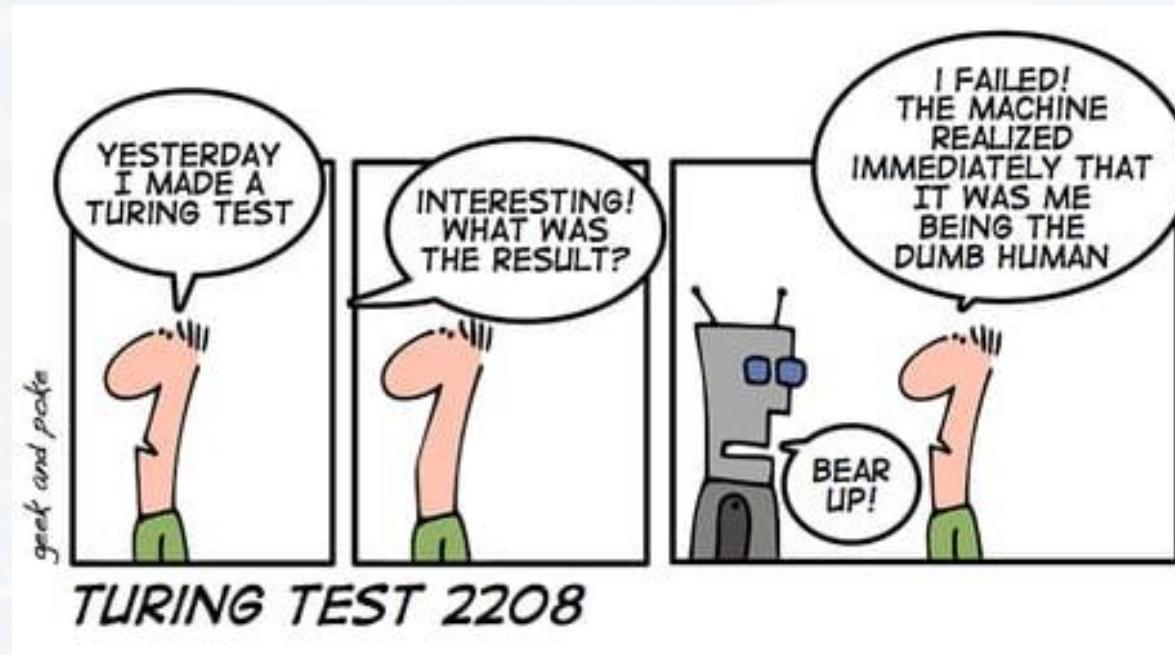
Lecture 12: Combining LSTMs and CNNs

Lecture 13: Running Models on GPUs and Parallel Processing

Lecture 14: Project Presentations

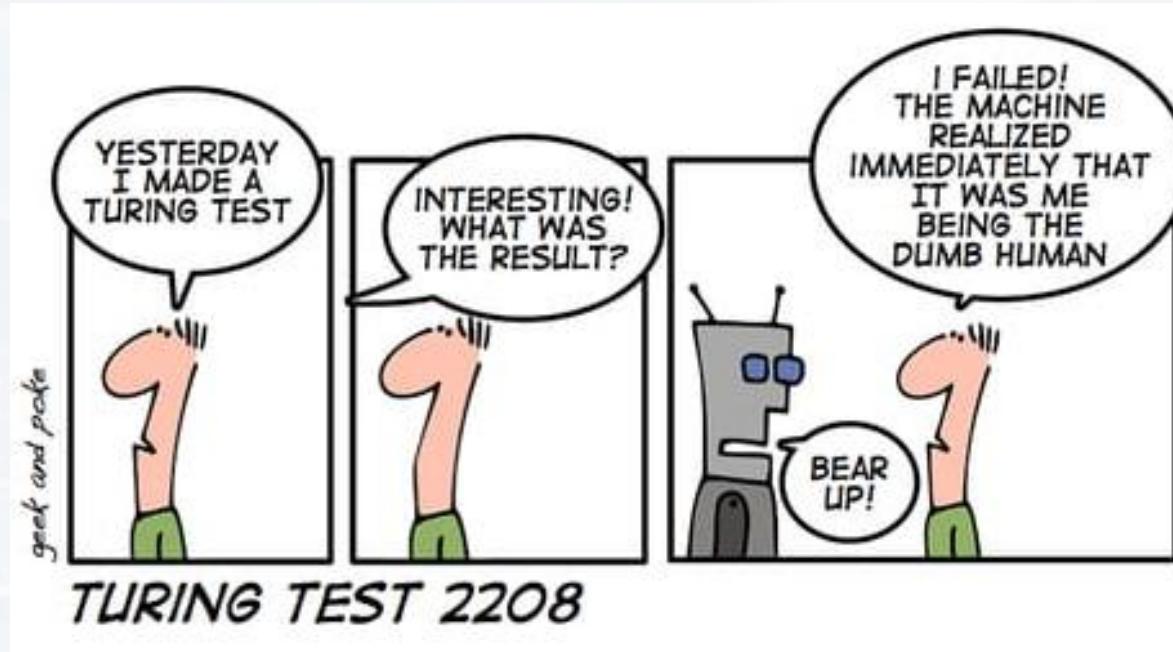
**Lecture 15: Transformer**

Lecture 16: GNN



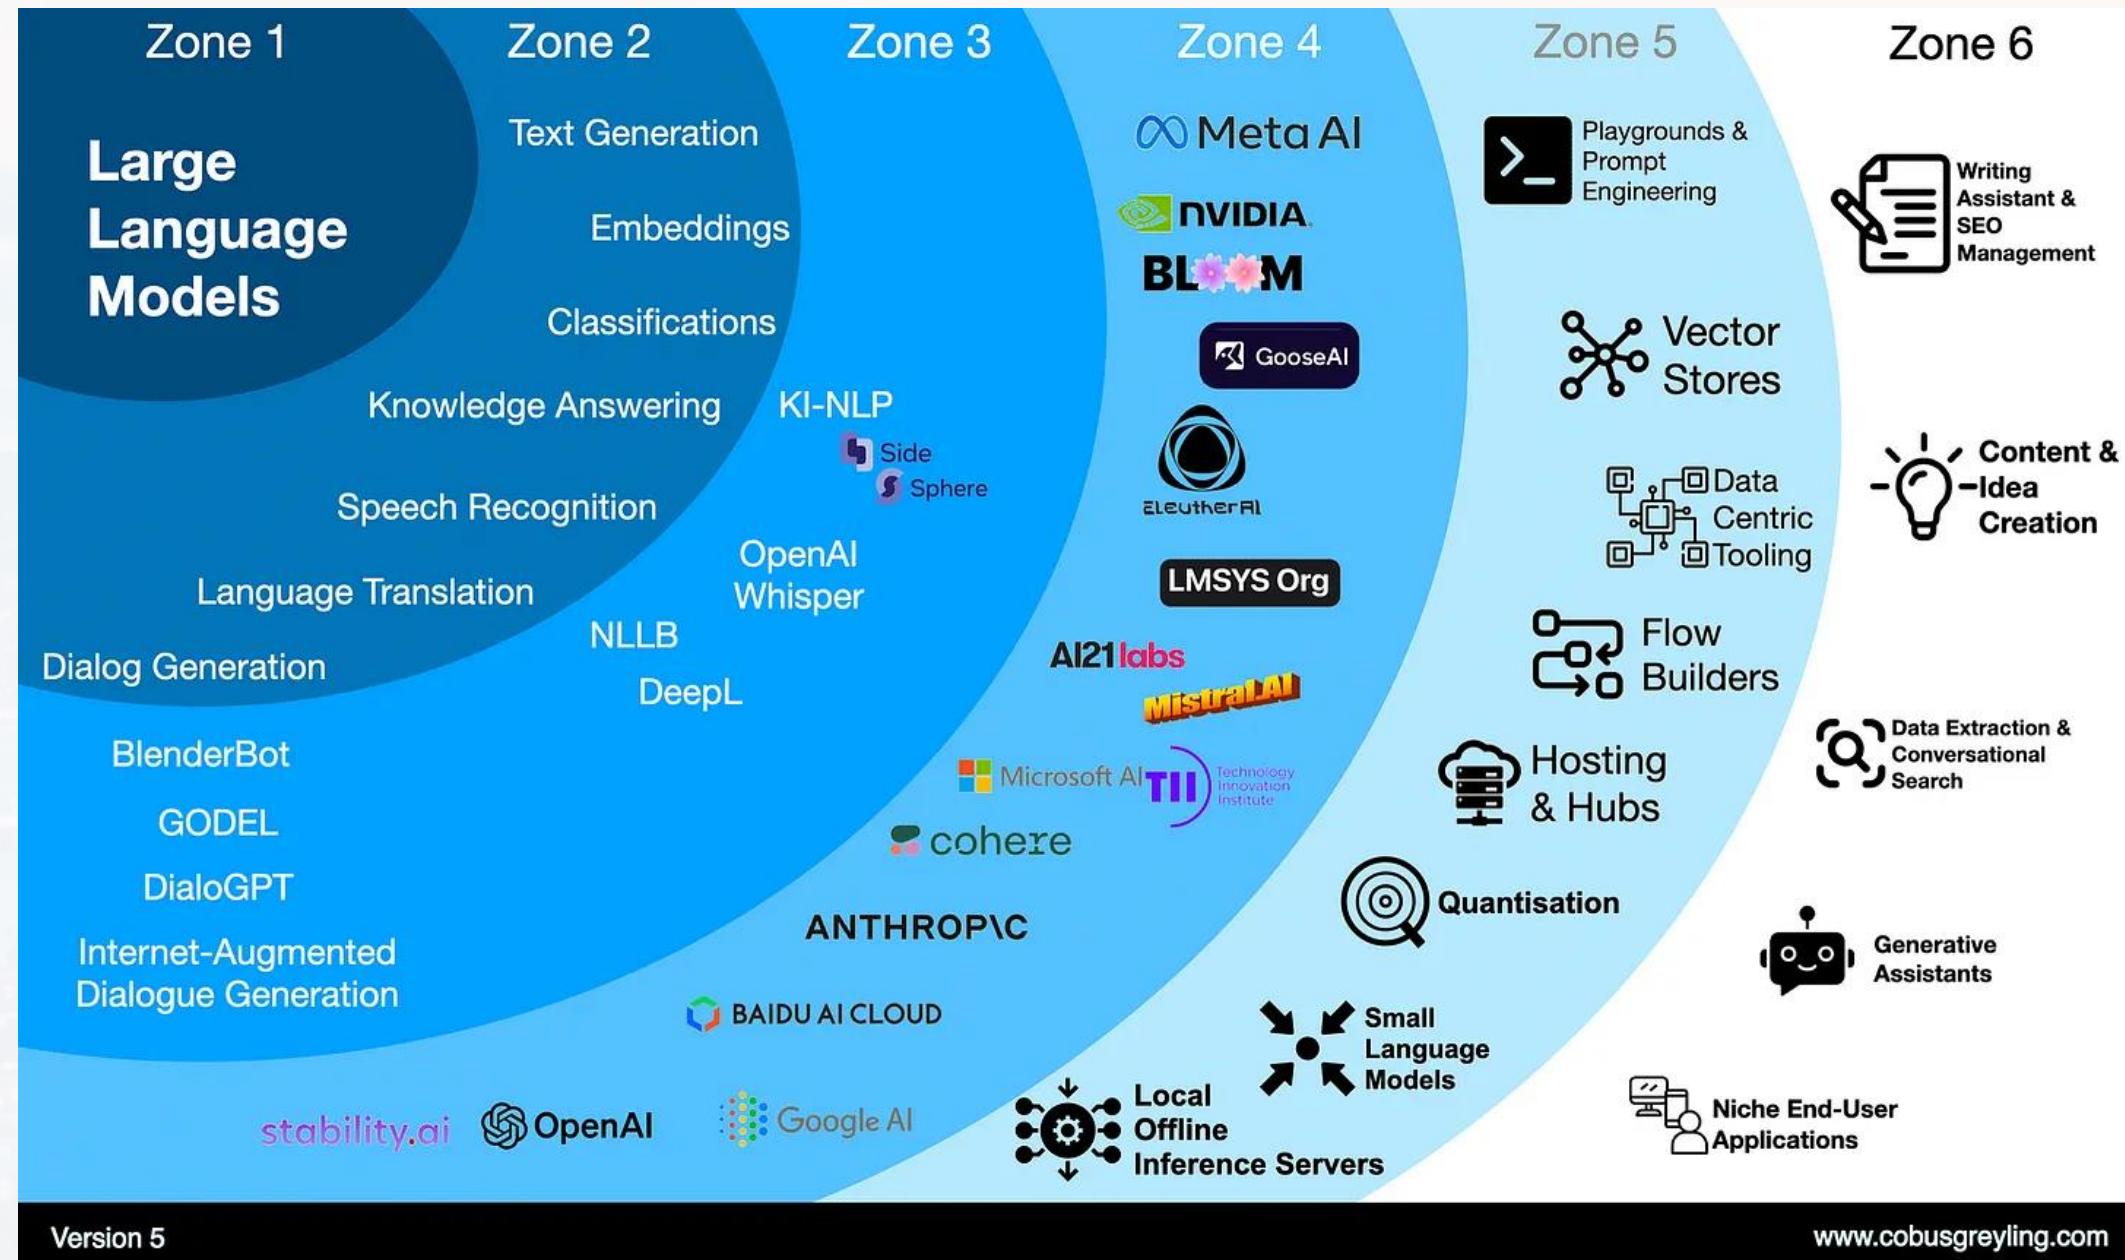
## Outline

- Introduction
- Bigram and MAP
- Positional Encoding
- Word Embedding
- Attention
- Transformer Architecture



## Outline

- Introduction
- Bigram and MAP
- Positional Encoding
- Word Embedding
- Attention
- Transformer Architecture





**corpus:** (large, representative) data set containing sequences of a language

**token:** individual, independent entity of a language

**alphabet/vocabulary:** set of tokens

### token

- letters in a word
- words in a sentence

(upper/lower case, cases, gender, tenses, conjugations)

- amino acids in a protein sequence
- nucleotides in a DNA/RNA sequence
- motifs in a DNA/RNA sequence

### size of alphabet

- $10^2$
- $10^4 \dots 10^6$
- 21
- 4
- $10^4$

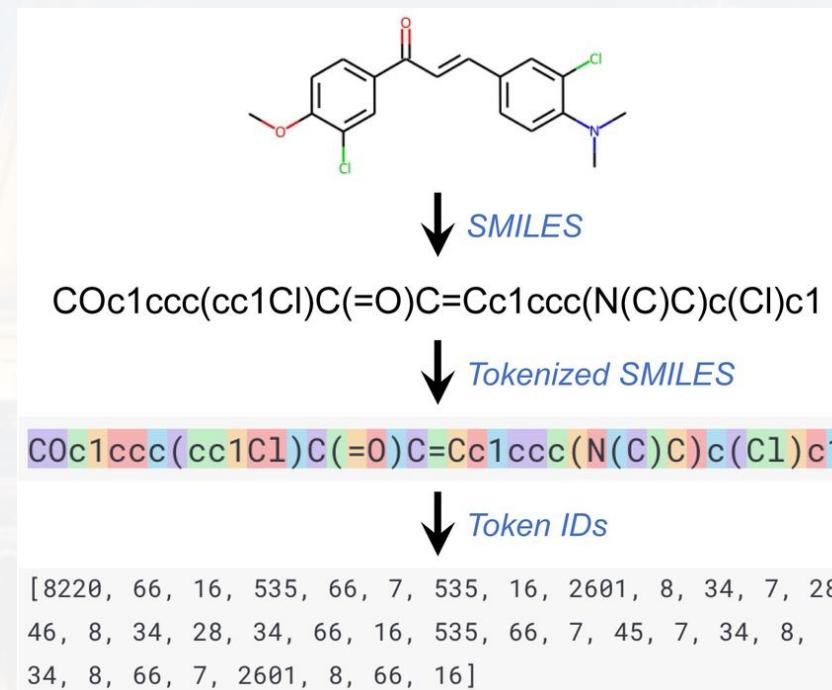


**corpus:** (large, representative) data set containing sequences of a language

**token:** individual, independent entity of a language

**alphabet/vocabulary:** set of tokens

### tokenization



**token:** - single atom vs...  
- ...functional group



**corpus:** (large, representative) data set containing sequences of a language

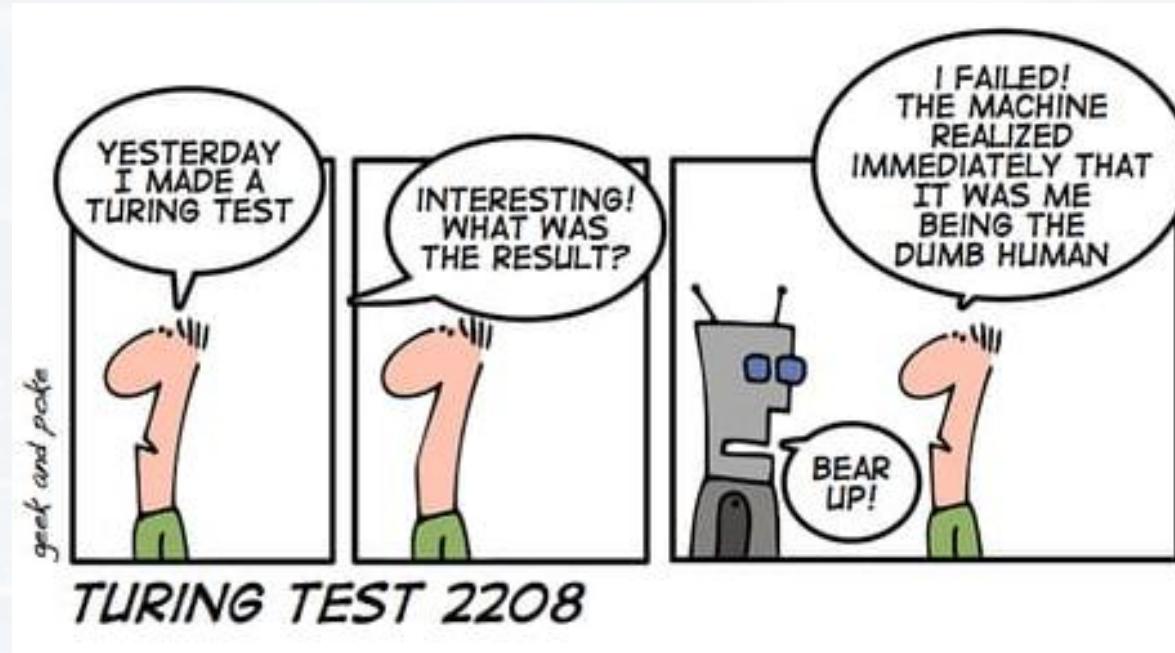
**token:** individual, independent entity of a language

**alphabet/vocabulary:** set of tokens

**note:** language models don't know grammar as we do, but they don't need to anyway...

**three things make context** (details: see later):

- **word embedding** (relation between similar/different token)
- **positional encoding** (location of token in a sequence)
- **attention** (relation between token within a sequence)



## Outline

- Introduction
- Bigram and MAP
- Positional Encoding
- Word Embedding
- Attention
- Transformer Architecture



$X_1 X_2 X_3 X_4 X_5 \dots X_n$

sequence of  $n$  token  $X$

actually:

$$P(X_1 X_2 X_3 X_4 X_5 \dots X_n) = P(X_n | X_{n-1} \dots X_1) P(X_{n-1} | X_{n-2} \dots X_1) \dots P(X_1)$$

bigram (1<sup>st</sup> order Markov Chain, see e.g. first WhatsApp versions):

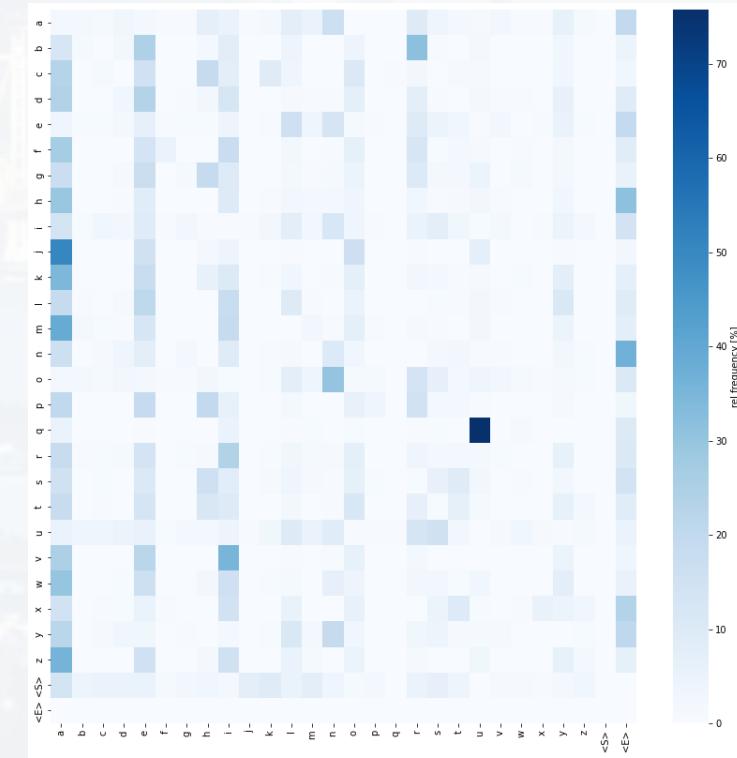
$$P(X_1 X_2 X_3 X_4 X_5 \dots X_n) = P(X_n | X_{n-1}) P(X_{n-1} | X_{n-2}) \dots P(X_1)$$



**P(i|j):** that token  $i$  is generated after token  $j$

→ N x N transition matrix from frequencies

→ “bigram” = “two words”



frequency matrix of letters in common names



### bigram (1<sup>st</sup> order Markov Chain):

let's build our own bigram model: **generate new names** based on a corpus of names

```
In [15]: words[0:12]
Out[15]:
['emma',
'olivia',
'ava',
'isabella',
'sophia',
'charlotte',
'mia',
'amelia',
'harpers',
'evelyn',
'abigail',
'emily']
```

see **Andrej Karpathy's GitHub** repository

$$P(X_1 X_2 X_3 X_4 X_5 \dots X_n) = P(X_n|X_{n-1})P(X_{n-1}|X_{n-2}) \dots P(X_1)$$

we only need to count **how often** a letter is followed by another

we also need to indicate when a name has **started** and **ended**

[ '<S>' ] + [ 'olivia' ] + [ '<E>' ]

→ **alphabet:** 26 letters + the two special “letters”

let's create a dictionary first (will help for counting):



## bigram (1<sup>st</sup> order Markov Chain):

let's build our own bigram model: **generate new names** based on a corpus of names

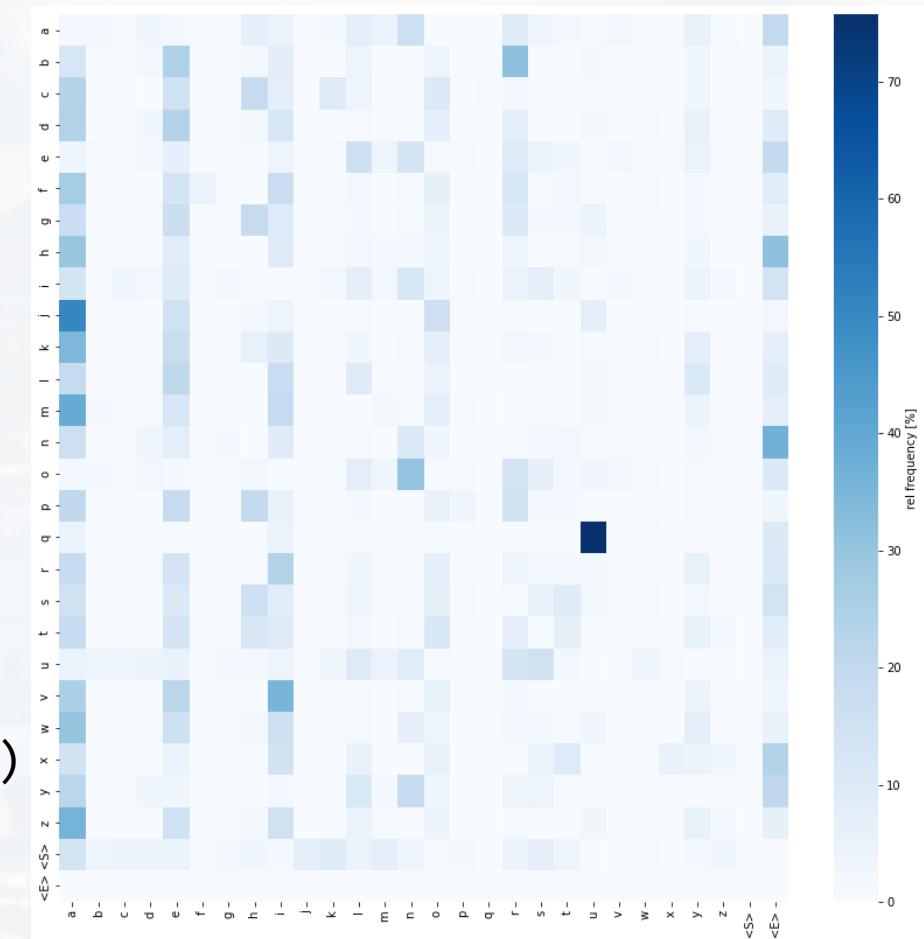
- 1) dictionary first (will help for counting)
- 2) count how often letter  $i$  is followed by letter  $j$

→ bigram matrix  $N$

- 3) normalize  $N$  accordingly
- 4) begin with a start token
- 5) draw a letter randomly based on  $N$ , using

```
np.argmax(np.random.multinomial(1,p))
```

- 6) if next token is stop token → stop



bigram (1<sup>st</sup> order Markov Chain):

let's build our own bigram model: **generate new names** based on a corpus of names  
check out **Bigram.ipynb**

**B.SampleNames(15)** vs totally random **B.SampleNames(15, False)**

some names  
are gibberish

some names  
sound real

some names  
are real

In [295]: **B.SampleNames(15)**

keesa  
ann  
ja  
jon  
nma  
malynojana  
sall  
daha  
drvah  
lzaxi  
tyunusthun  
jorrwro  
ja  
asoow  
s

In [296]: **B.SampleNames(15, False)**

mtkg  
yufexhviov  
morrhqvik  
bbbjxebpxwure  
jaqlzzuwuanx  
mmomhr<S>  
uhb  
xlmusadjfdzxadaotd  
ik<S>vdtydvxev  
taselkykcfbamceprtv1  
zyr<S>inzoerobz  
wuovx  
eg<S>pbdvikf<S>  
tomcnkfsjay<S>  
rikatnaykizszciv  
pds<S>zj  
kh<S>y<S>ualzugqgak  
akeubjbasc  
bblupnibtqmyl<S>  
vyobf  
kybs  
rznjgpml  
tnhoxuckkjzbwmj<S>  
vshkycicf<S>  
kowskphy  
rxodh  
jvswmzw  
jzpcfnpbg



Note, there is no conceptual difference between applying our model to *letters in a word vs words in a sentence*

caveats:

- the bigram model derives  $P(X_n)$  from **observed** frequencies  
→ essentially **MLE** (problematic if a letter hasn't appeared in the sequence yet  
→  $P(X_n)$  assumed to be zero!)

```
Nsam = N/np.sum(N+0.0001, axis = 1, keepdims = True)  
S_bi += np.sum(-N[:,i]*np.log(N[:,i]+1e-16))
```

- can we implement something that is closer to:

$$P(X_1 X_2 X_3 X_4 X_5 \dots X_n) = P(X_n|X_{n-1} \dots X_1)P(X_{n-1}|X_{n-2} \dots X_1) \dots P(X_1) \quad ?$$

binomial process

$$P(k|n, q) = \binom{n}{k} q^k (1 - q)^{n-k}$$

$q - 1$



$q$



$q = ?$





$$P(X_1 X_2 X_3 X_4 X_5 \dots X_n) = P(X_n|X_{n-1} \dots X_1)P(X_{n-1}|X_{n-2} \dots X_1) \dots P(X_1)$$

Bayesian  
Parameter  
Estimation\*)



$$P(k|n, q) = \binom{n}{k} q^k (1-q)^{n-k} \quad q = ?$$

likelihood function (here: binomial)

$$P(q|data\ set) = \frac{P(data\ set|q)P(q)}{P(data\ set)} \text{ prior } (\sim P(X_n|X_{n-1} \dots X_1))$$

evidence (const wrt q)

$q = const$   
before 1<sup>st</sup> data point  
(max entropy!)

$$= \frac{1}{\int_0^1 P(q|data\ set) dq} (1-q)^{n-k} q^k$$

$q = conjugate\ prior$   
after n<sup>th</sup> data point

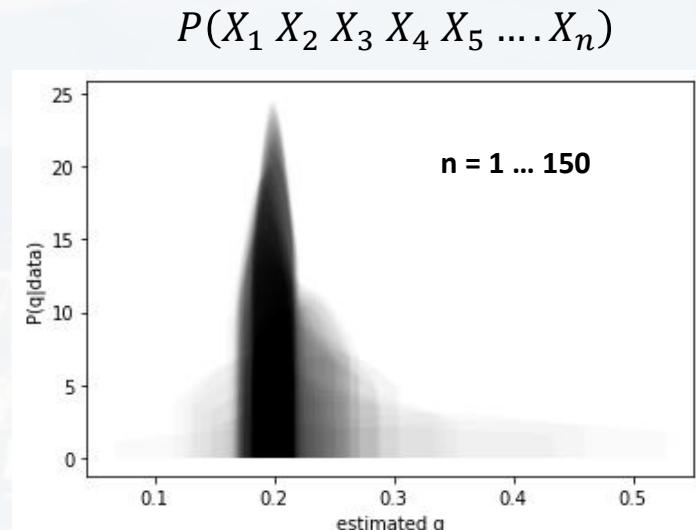
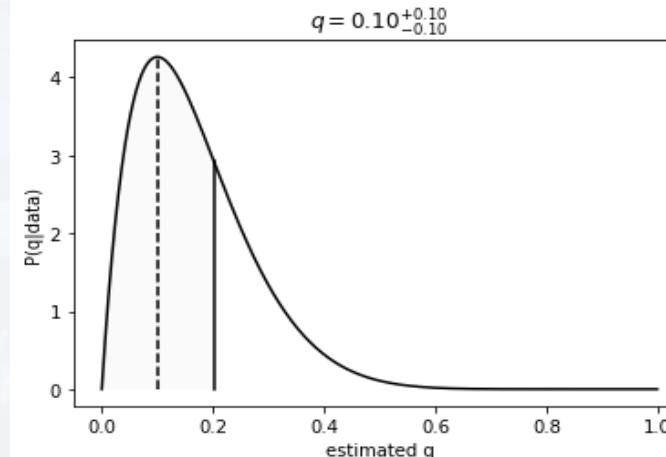
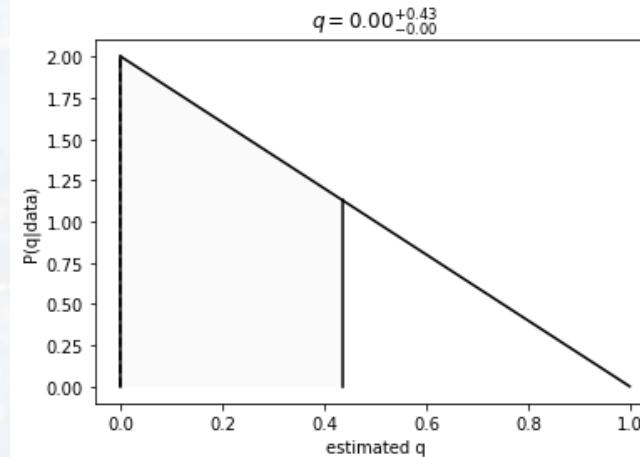
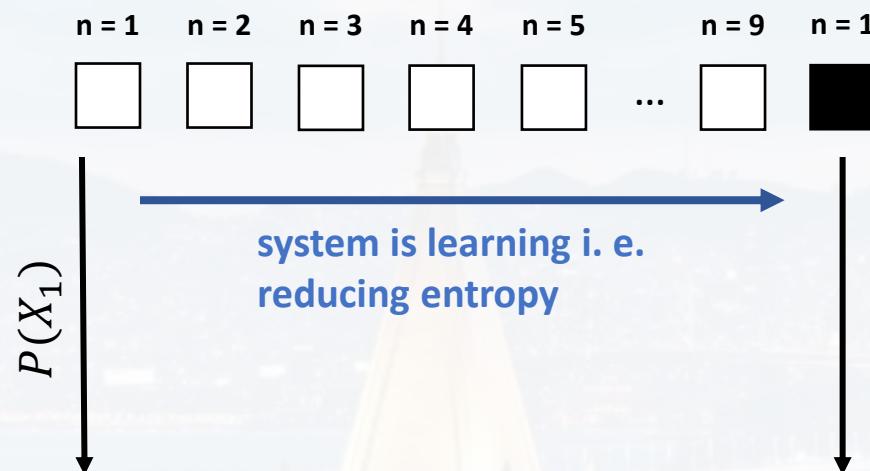
$$= \frac{q^{k+\alpha-1} (1-q)^{n-k+\beta-1}}{\int_0^1 q^{k+\alpha-1} (1-q)^{n-k+\beta-1} dq}$$



$$P(X_1 X_2 X_3 X_4 X_5 \dots X_n) = P(X_n|X_{n-1} \dots X_1)P(X_{n-1}|X_{n-2} \dots X_1) \dots P(X_1)$$

$$P(k|n, q) = \binom{n}{k} q^k (1-q)^{n-k} \quad q = ?$$

Bayesian  
Parameter  
Estimation\*)



\*) see lecture 2



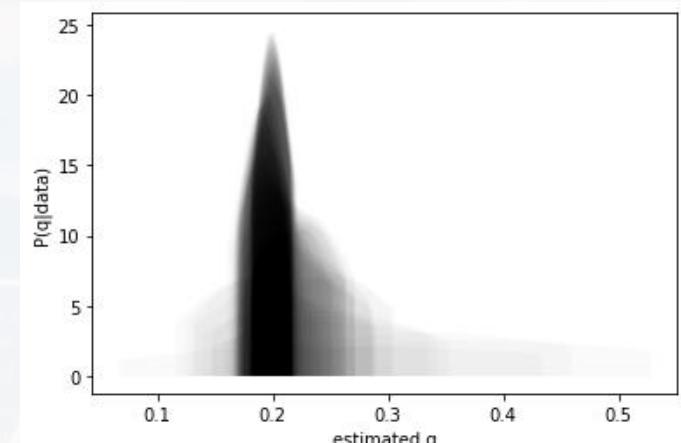
$$P(X_1 X_2 X_3 X_4 X_5 \dots X_n) = P(X_n|X_{n-1} \dots X_1)P(X_{n-1}|X_{n-2} \dots X_1) \dots P(X_1)$$

Bayesian  
Parameter  
Estimation\*)

$$P(k|n, q) = \binom{n}{k} q^k (1-q)^{n-k}$$

$$P(q|data\ set) = \frac{q^{k+\alpha-1}(1-q)^{n-k+\beta-1}}{\int_0^1 q^{k+\alpha-1}(1-q)^{n-k+\beta-1} dq}$$

Beta function



more general, we want to learn the probability  $P_j(a)$  of letter  $a$  at position  $j$

$$q \rightarrow P_j(a)$$

→ multinomial problem

→ conjugate prior is the **Dirichlet distribution**

$$P(sequence) \sim \prod_j \prod_a P(a)_j^{\alpha(a)-1}$$

equivalent to what was  $P(q|data\ set)$  earlier

\*) see lecture 2

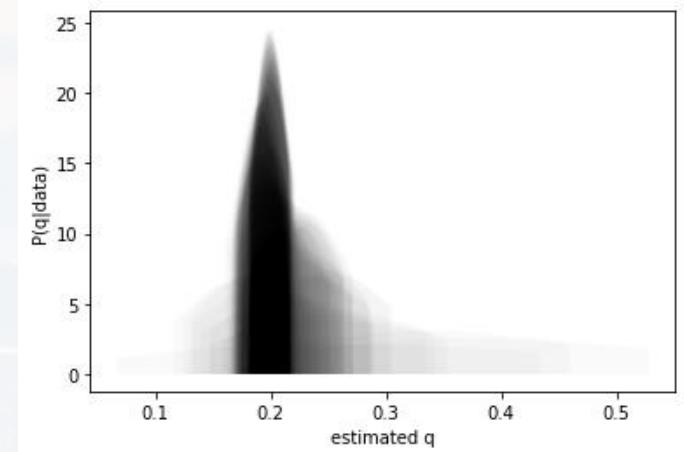


$$P(X_1 X_2 X_3 X_4 X_5 \dots X_n) = P(X_n|X_{n-1} \dots X_1)P(X_{n-1}|X_{n-2} \dots X_1) \dots P(X_1)$$

### Dirichlet distribution

$$P(\text{sequence}) \sim \prod_j \prod_a P(a)_j^{\alpha(a)-1}$$

note:  $\sum_{\text{over all } a} P(a)_j = 1 \rightarrow N - \text{dim simplex}$



- note:
- we don't need to extract  $P(a)$  from the maximum of the pdf given by the BPE posterior
  - we can directly derive the maximum of  $P(a)$  from  $P(\text{sequence})$  given the constrain  $\sum_{\text{over all } a} P(a)_j = 1$  (**Lagrangian multipliers**)
  - **Maximum a-posteriori (MAP)** approach → see XXmotif (Siebert & Soeding, 2016)

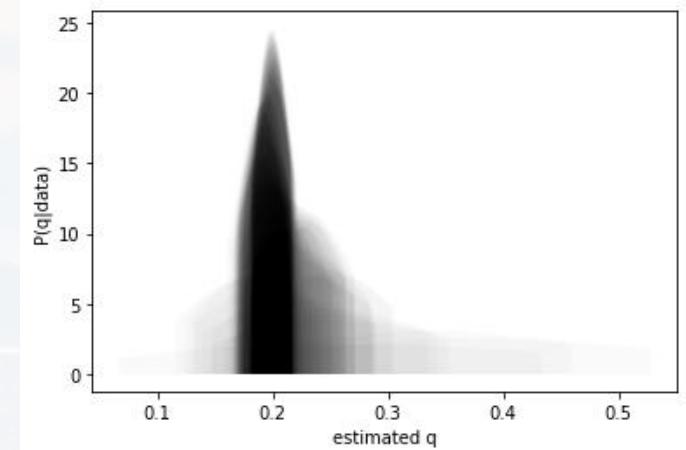


$$P(X_1 X_2 X_3 X_4 X_5 \dots X_n) = P(X_n|X_{n-1} \dots X_1)P(X_{n-1}|X_{n-2} \dots X_1) \dots P(X_1)$$

### Dirichlet distribution

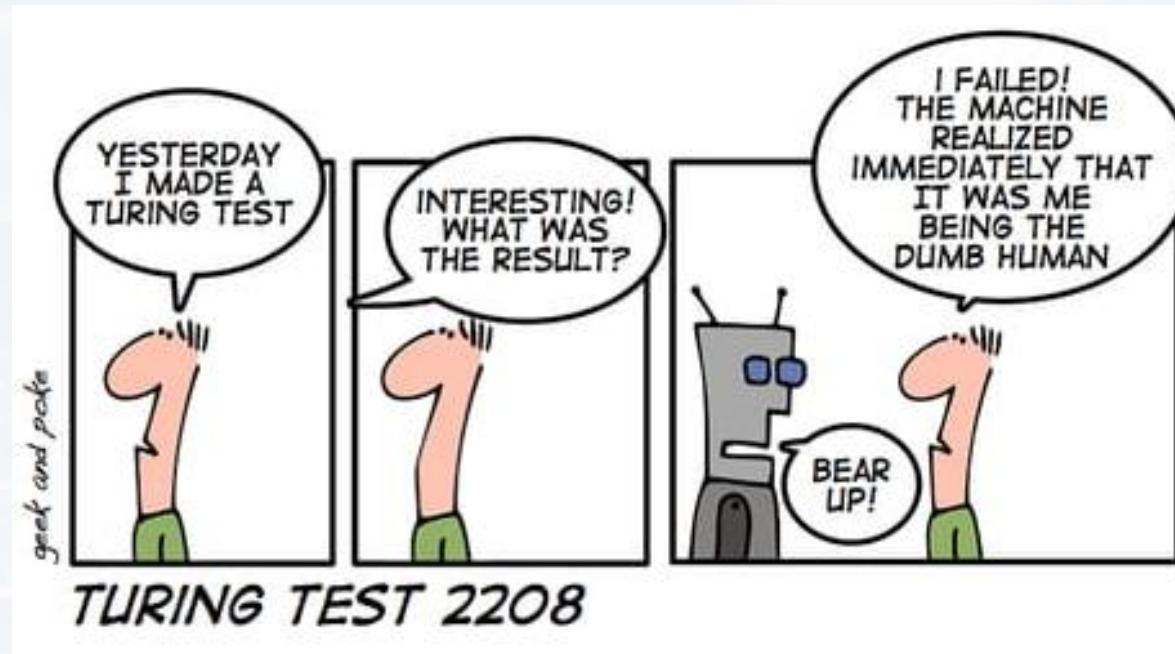
$$P(\text{sequence}) \sim \prod_j \prod_a P(a)_j^{\alpha(a)-1}$$

note:  $\sum_{\text{over all } a} P(a)_j = 1 \rightarrow \text{N - dim simplex}$



### Maximum a-posteriori (MAP)

- XXmotif (Siebert & Soeding, 2016) significantly outperformed PWMs
  - it struggled however with related motifs which where **physically located far apart** from each other
- solution see later: attention  
→ older solutions: LSTMs



## Outline

- Introduction
- Bigram and MAP
- Positional Encoding
- Word Embedding
- Attention
- Transformer Architecture



three things make context:

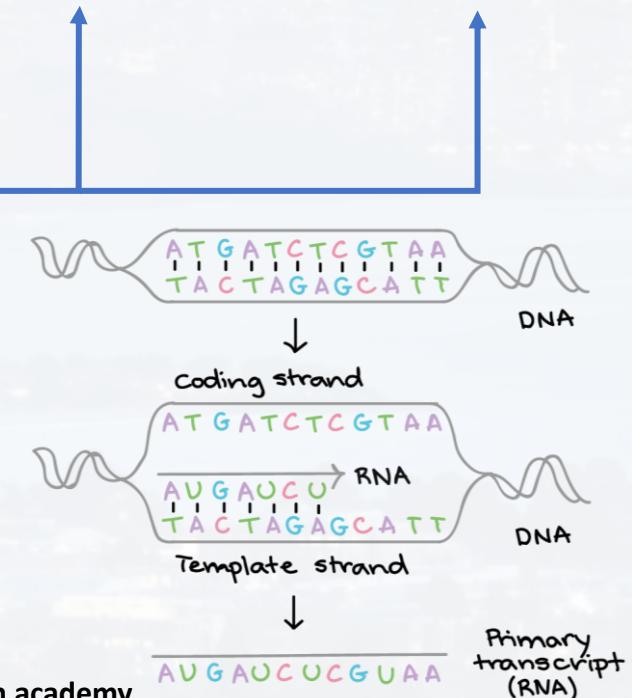
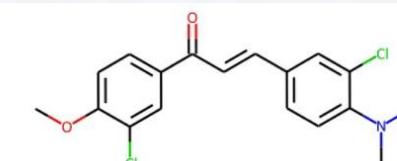
- **positional encoding** (location of token in a sequence)
- **word embedding** (relation between similar/different token)
- **attention** (relation between token within a sequence)

*"The cat jumped on the roof."*

order matters!:

- 1<sup>st</sup>: article
- 2<sup>nd</sup>: noun/subject
- 3<sup>rd</sup>: verb
- 4<sup>th</sup> : noun/object (in English)

→ positional encoding





goal: find a positional encoding that is

- reasonably simple
- independent from the length of the sequence
- somehow normalized

one idea: n-bit binary encoding

position code

76543210 ← 8bit i.e. eight dimensions

1	00000001
2	00000010
3	00000011
4	00000100
5	00000101
6	00000110
7	00000111
8	00001000
9	00001001
10	00001010
11	00001011
12	00001100
13	00001101
14	00001110
15	00001111
16	00010000

depending on dimensions (bit)

→ different frequencies

bit	frequency
0	1/2
1	1/4
2	1/8
...	

Does that look familiar?  
→ like Fourier Series



even dimensions:  $E(p, 2k) = \sin\left(\frac{p}{10,000} \frac{2k}{d}\right)$

odd dimensions:  $E(p, 2k + 1) = \cos\left(\frac{p}{10,000} \frac{2k}{d}\right)$

$p$ : position in sequence

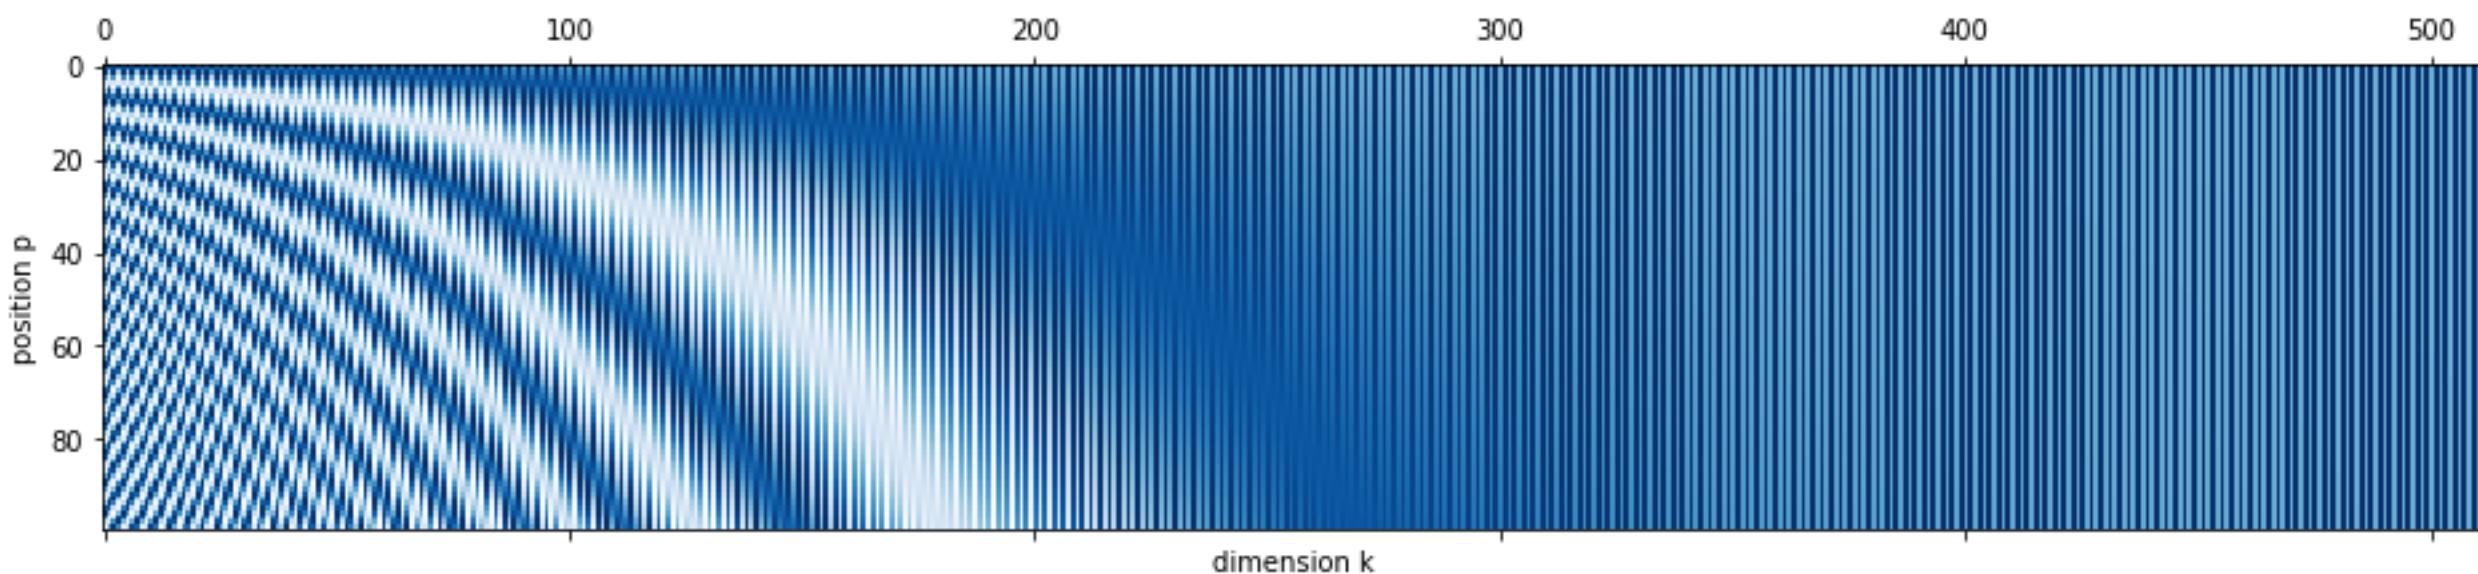
$k$ : dimension index

$d$ : number of dimensions

10,000: an arbitrary number      Vaswani et al., 2017

run **PlotPositionEncoding.py**

more info [here](#)





$$\text{even dimensions: } E(p, 2k) = \sin\left(\frac{p}{10,000 \frac{2k}{d}}\right)$$

$$\text{odd dimensions: } E(p, 2k + 1) = \cos\left(\frac{p}{10,000^2 k/d}\right)$$

$p$ : position in sequence

$k$ : dimension index

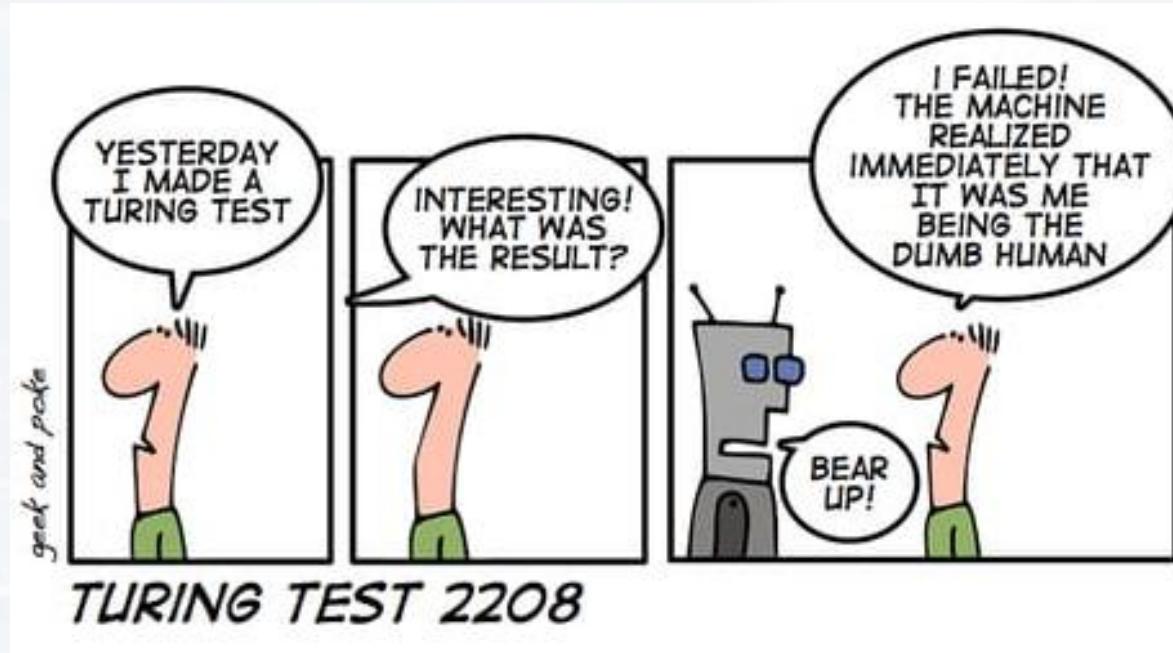
$d$ : number of dimensions

10,000: an arbitrary number      Vaswani et al., 2017

run **PlotPositionEncoding.py**

### note:

- easier to handle numerically vs discrete encoding
- $\cos(x + 2\pi k) = \cos(x)$  and  $\sin(x + 2\pi k) = \sin(x)$   
→ absolute position not relevant but **relative** position



## Outline

- Introduction
- Bigram and MAP
- Positional Encoding
- Word Embedding
- Attention
- Transformer Architecture



### three things make context:

- **positional encoding** (location of token in a sequence)
- **word embedding** (relation between similar/different token)
- **attention** (relation between token within a sequence)

### problem: turning token (words/letters) into numbers

#### single letters:

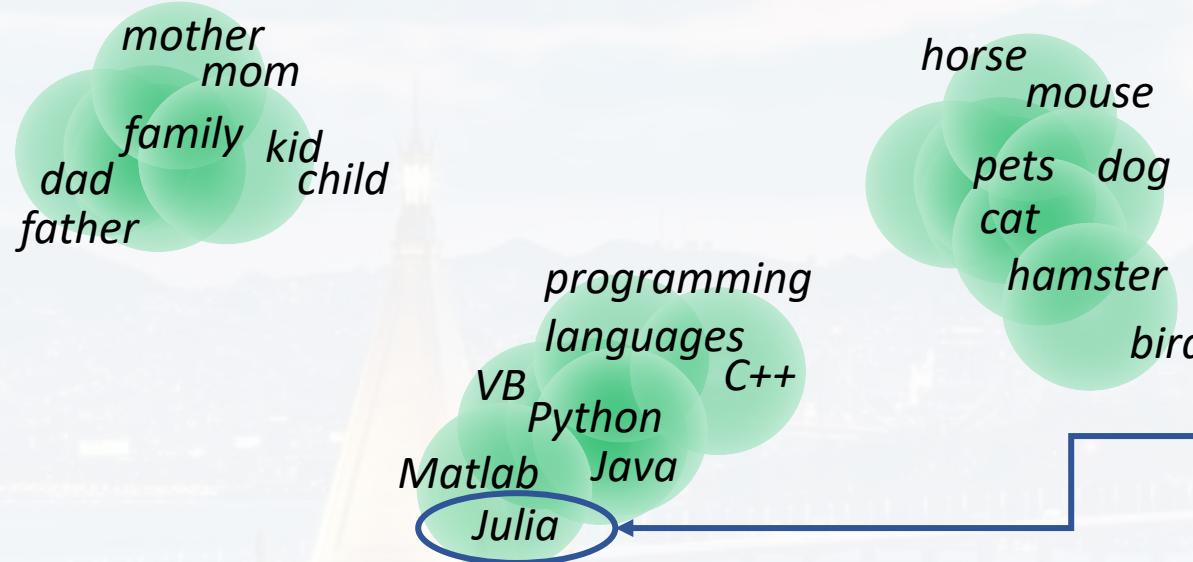
- |         |  |
|---------|--|
| ACGT    | - one – hot works perfectly (four different token)                                 |
| abcd... | - lower/upper case, special characters (50 different token), one – hot is fine too |

#### words:

- |              |  |
|--------------|--|
| actual words | - $10^4 \dots 10^6$ (upper/lower case, cases, gender, tenses, conjugations)<br>→ one – hot doesn't work (matrices would be too large)<br>→ some words have a <b>similar meaning</b> , should be <b>close</b> in parameter space ( <b>cluster</b> ) |
|--------------|--|



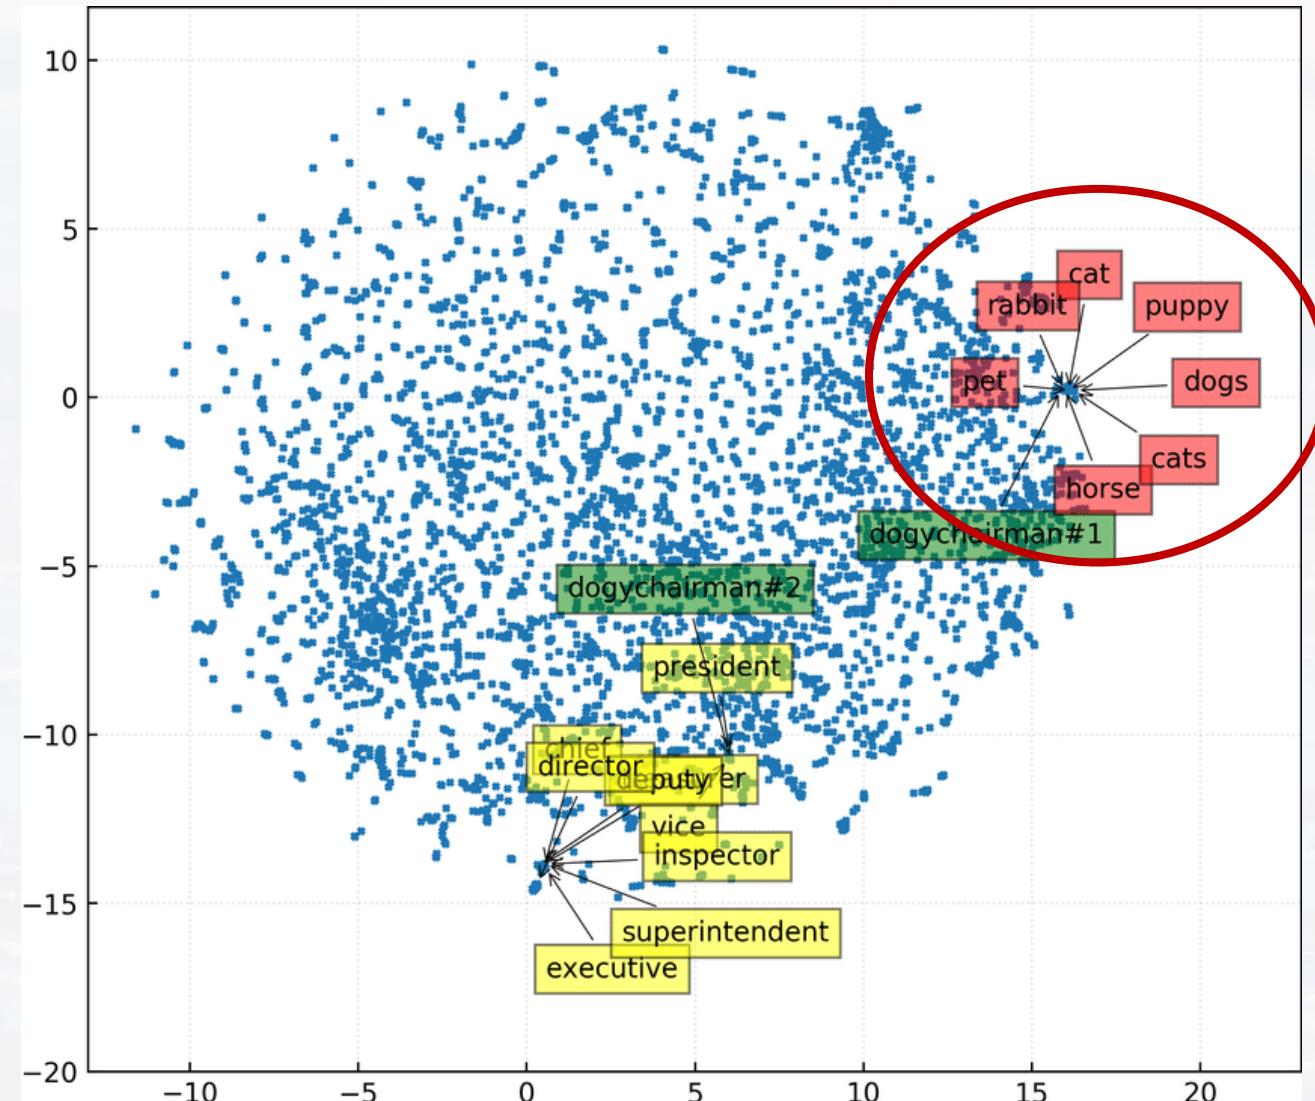
words with a **similar meaning** should form **cluster**



- embedding, instead of one – hot encoding
- from experience: **N = 30 – 300** dim vector for each token (**which is a lot less than  $10^4 \dots 10^6$** ) is sufficient
- as a result: token with **similar meaning are close** in the vector space!



words with a **similar meaning** should form **cluster**





common training set: recorded speeches from the European Parliament:

*...It seems absolutely disgraceful that we pass legislation and do not adhere to it ourselves. Mrs Lynne, you are quite right and I shall check whether this has actually not been done. I shall also refer the matter to the College of Quaestors, and I am certain that they will be keen to ensure that we comply with the regulations we ourselves vote on.*

*Madam President, Mrs Díez González and I had tabled questions on certain opinions of the Vice-President, Mrs de Palacio, which appeared in a Spanish newspaper.*

*The competent services have not included them in the agenda on the grounds that they had been answered in a previous part-session.*

*I would ask that they reconsider, since this is not the case....*

words of similar meaning should appear in similar environment

→ target token within a window

Two common algorithms are **Continuous Bag Of Words** and **skip gram**



### Continuous Bag Of Words

$n$ : number of unique token from corpus  
 $m$ : desired number of dimensions for embedding

*The competent services have not included them in the agenda on the grounds that they...*

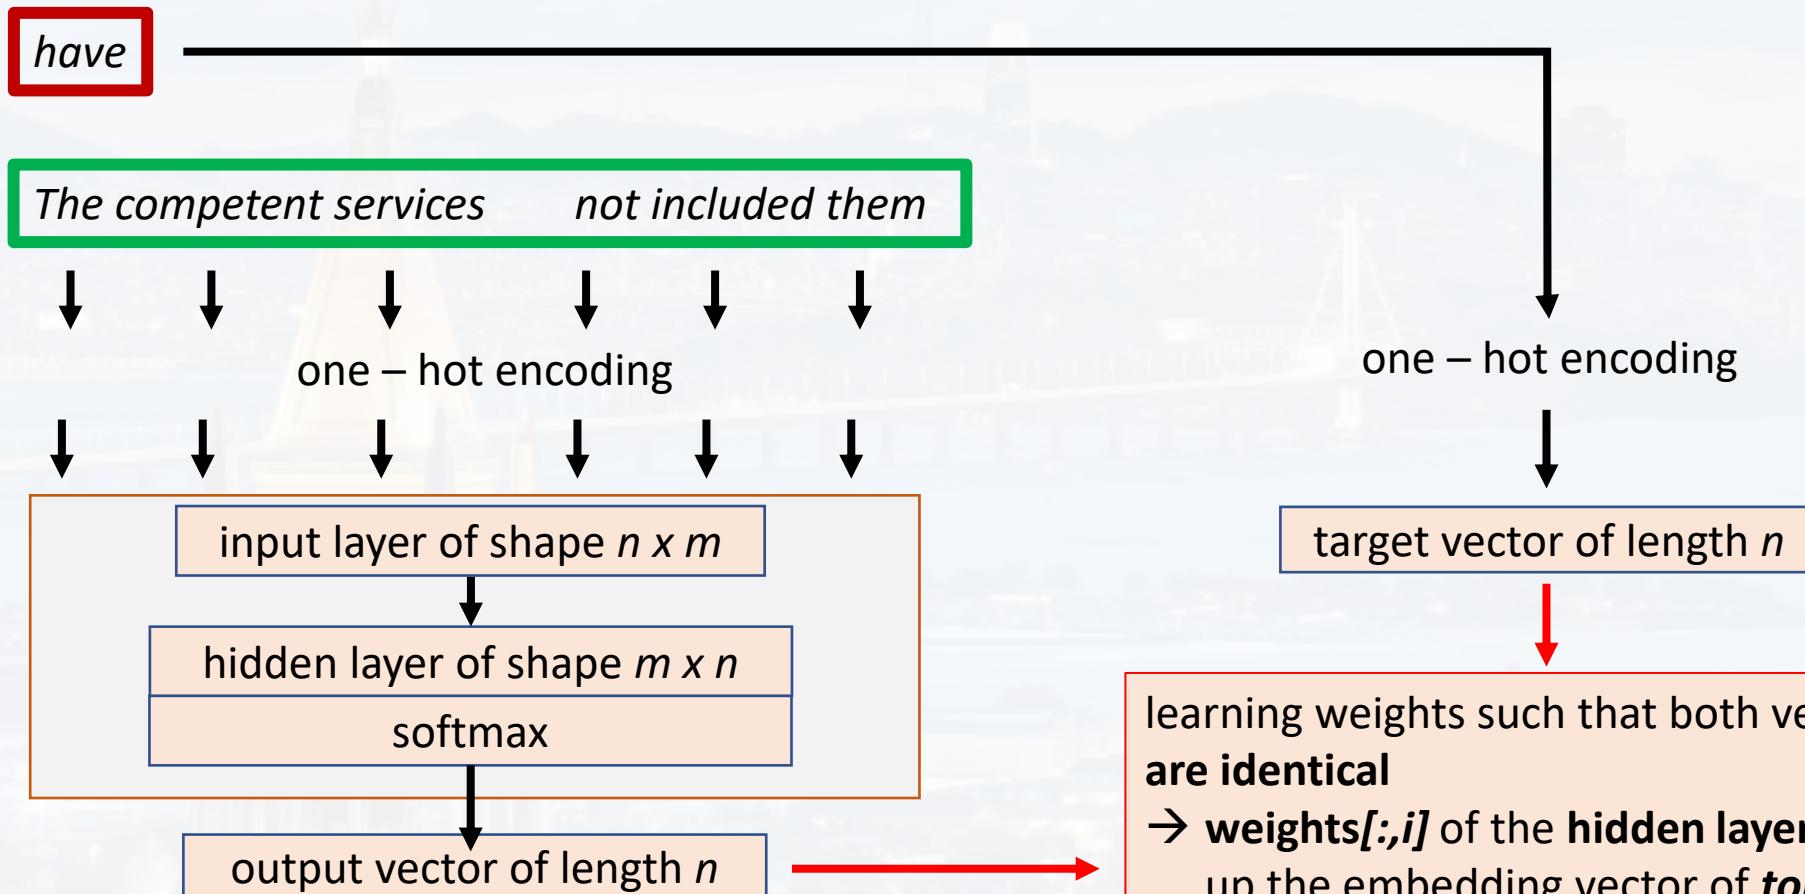
target

**have**

context window

*The competent services      not included them*

shallow ANN





### Continuous Bag Of Words

$n$ : number of unique token from corpus  
 $m$ : desired number of dimensions for embedding

The **competent services have not included them in the agenda on the grounds that they...**

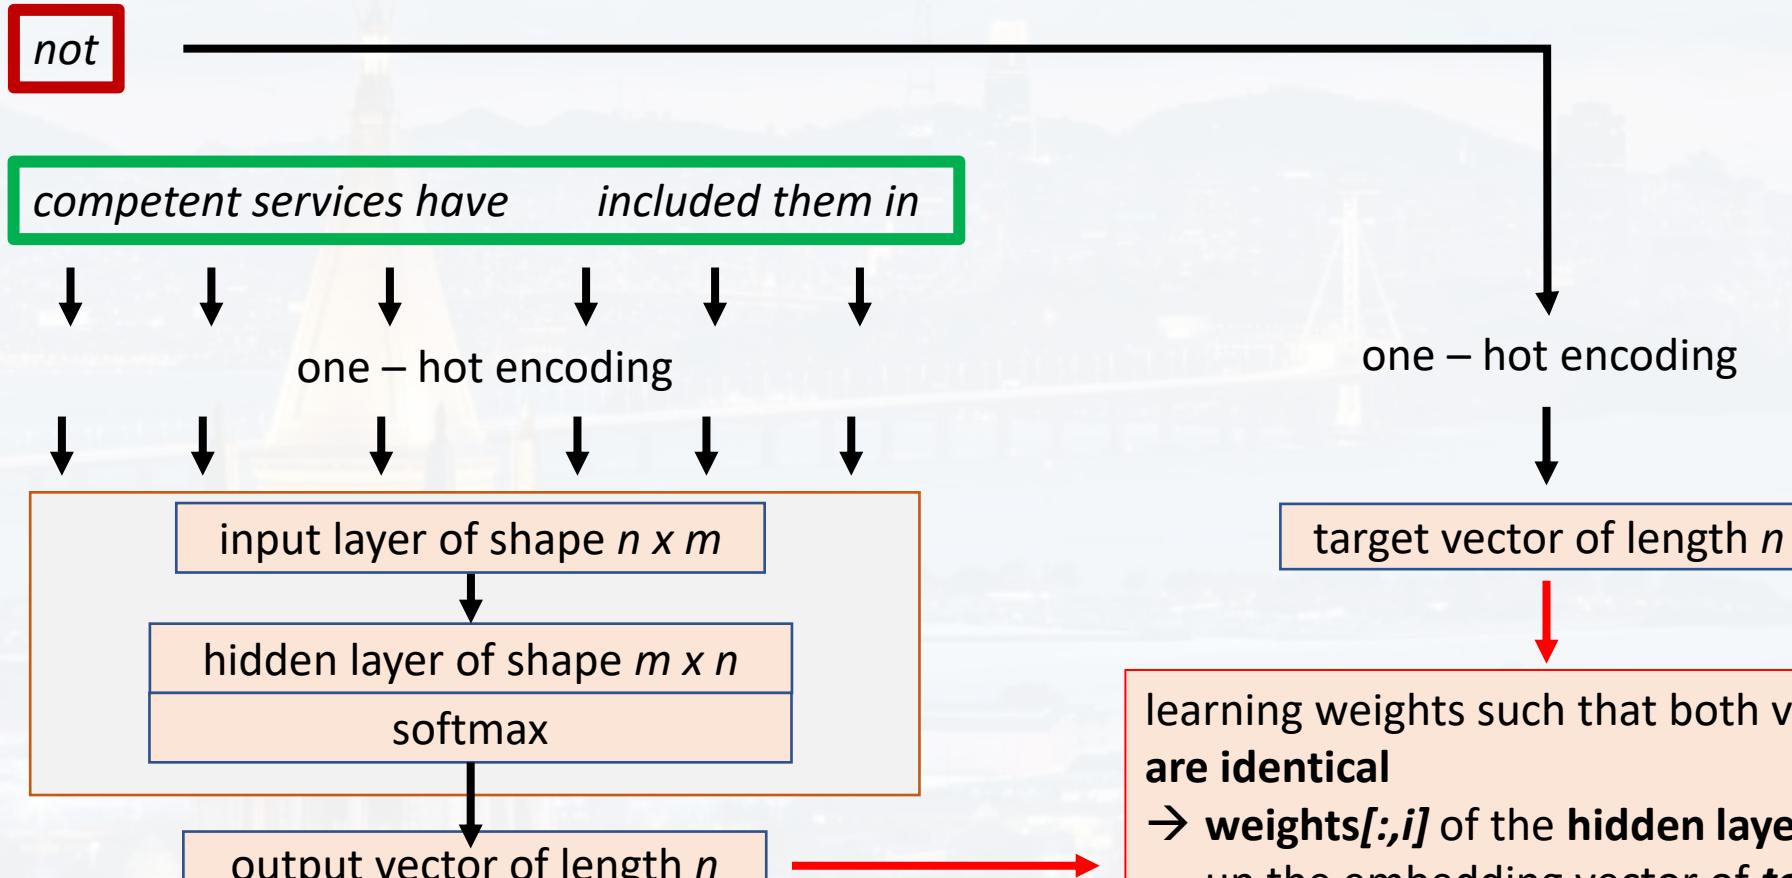
target

**not**

context window

**competent services have      included them in**

shallow ANN





### Continuous Bag Of Words

**n:** number of unique token from corpus  
**m:** desired number of dimensions for embedding

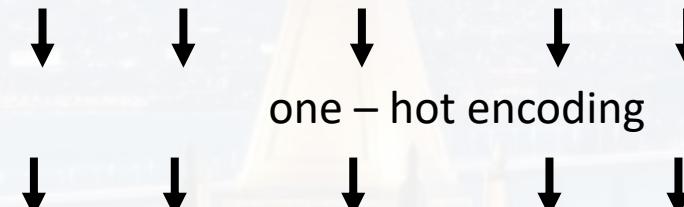
*The competent services have not included them in the agenda on the grounds that they...*

target

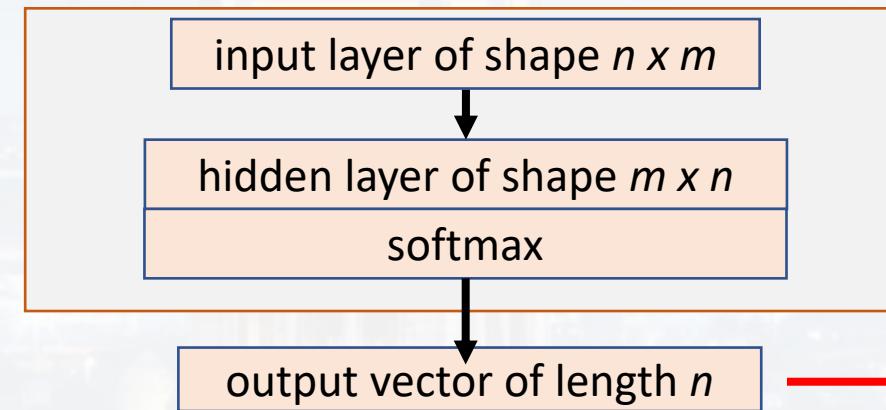
**included**

context window

**services have not      them in the**



shallow ANN

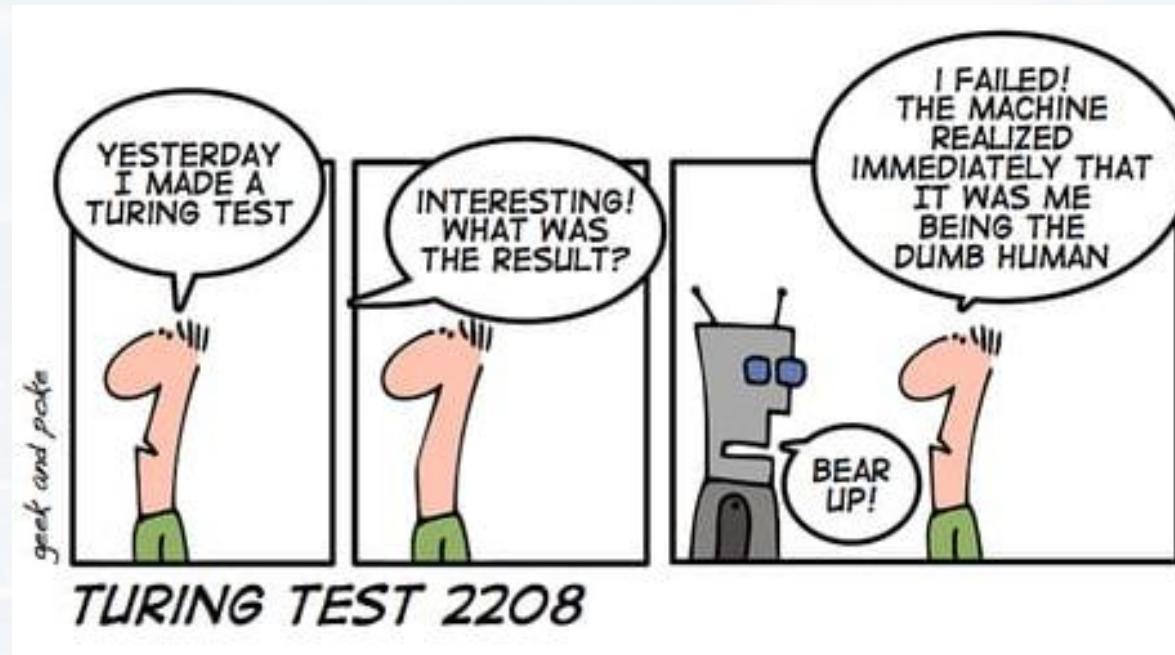


*... and so on....*

one – hot encoding

target vector of length  $n$

learning weights such that both vectors are **identical**  
→ **weights $[:,i]$**  of the **hidden layer** make up the embedding vector of **token i**



## Outline

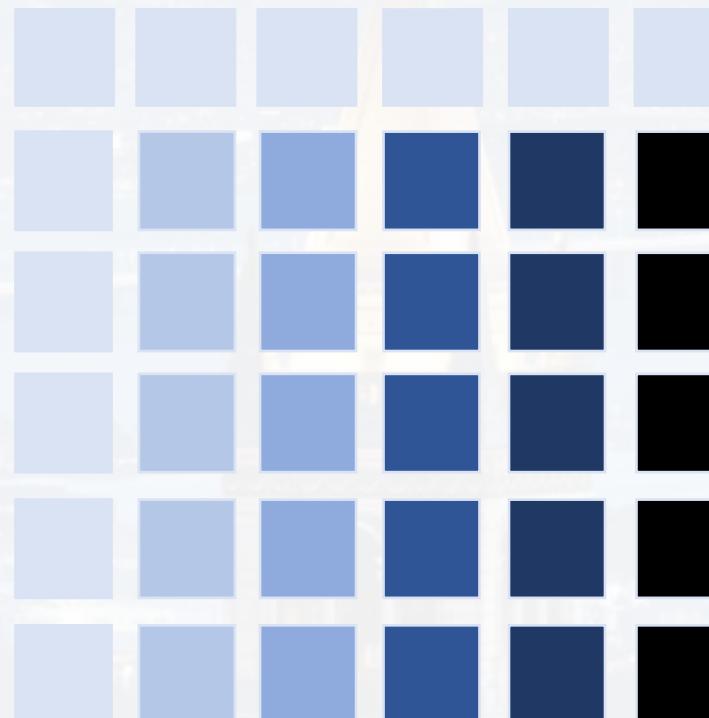
- Introduction
- Bigram and MAP
- Positional Encoding
- Word Embedding
- Attention
- Transformer Architecture



### three things make context:

- **positional encoding** (location of token in a sequence)
- **word embedding** (relation between similar/different token)
- **attention** (relation between token within a sequence)

*"The cat jumped on the roof."*

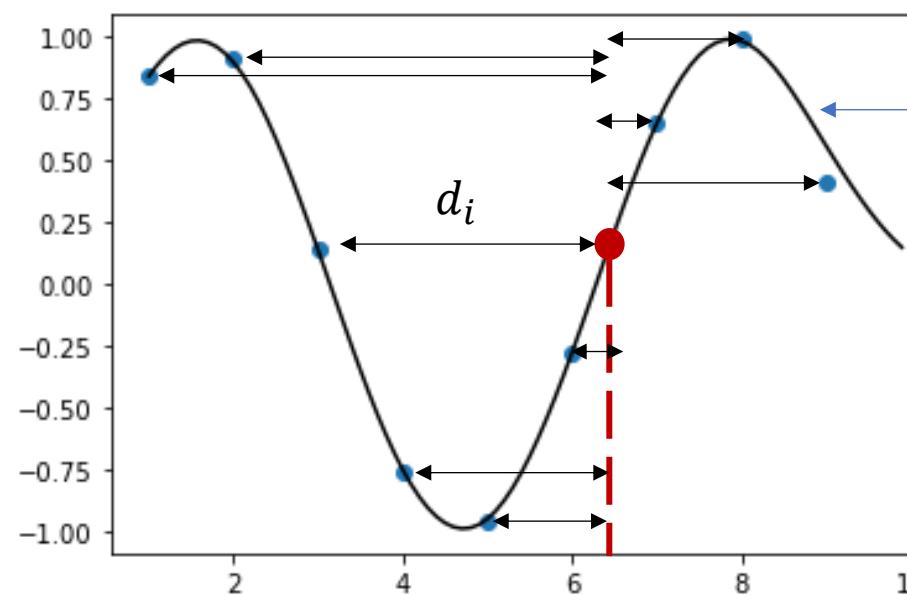


how the first token influences all other token



how the second token influences all other token

.... and so on



We want to interpolate between the blue dots  
→ generating the black line  
→ **no curve fitting!**

- idea:
- select a point for which we want the interpolation for
  - calculate distance  $d_i$  to every other point
  - each data point should influence the value of the interpolated point
  - the closer, the stronger the influence → weighted mean

$$y_{int} \sim \sum_{i=1}^I w_i \frac{1}{d_i} y_i$$

calculating distance

```
D = np.tile(V, (1, len(V))) - np.tile(L.transpose(), (len(V), 1))
```

- each data point should influence the value of the interpolated point
- the closer, the stronger the influence → weighted mean

$$y_{int} \sim \sum_{i=1}^I w_i y_i$$

Gaussian kernel

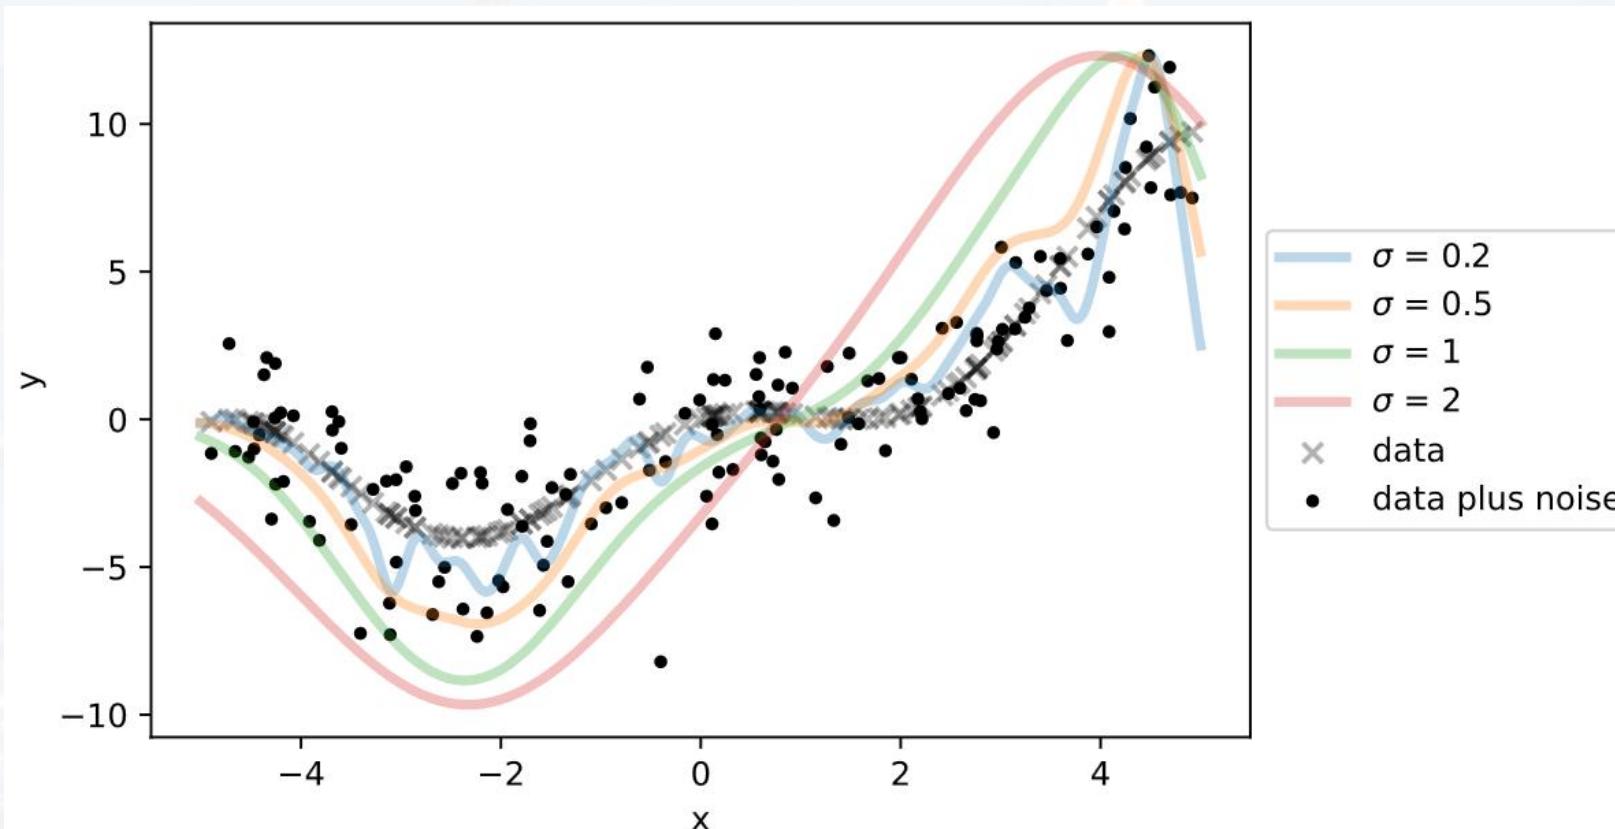
```
W      = np.exp(-(D**2)/(sigma))
W      = W/np.sum(W + 1e-16, axis = 0)
yint  = np.dot(W.transpose(), y)
```



```
D = np.tile(V, (1, len(V))) - np.tile(L.transpose(), (len(V), 1))
```

Gaussian kernel

```
W      = np.exp(-(D**2)/(sigma))
W      = W/np.sum(W + 1e-16, axis = 0)
yint  = np.dot(W.transpose(), y)
```

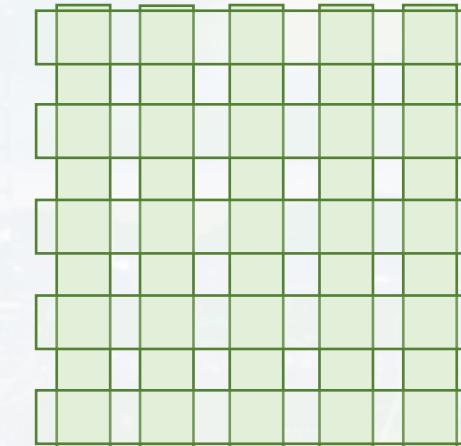
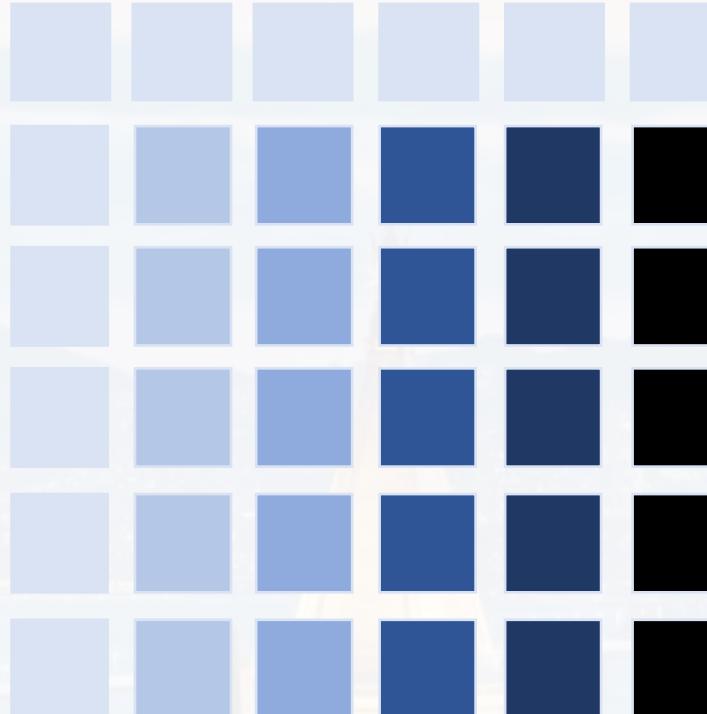


check out:

[SmoothGaussKernel.py](#)  
[SmoothExamples.py](#)



*"The cat jumped on the roof."*



Gaussian kernel	$W$	$= \text{np.exp}(-(\text{D}**2)/(\sigma))$
	$W$	$= W / \text{np.sum}(W + 1e-16, \text{axis} = 0)$
	$y_{int}$	$= \text{np.dot}(W^T, y)$

actual attention:  
these weights are learnable,  
no kernel assumed!



### self attention

imagine you want to built & train a movie GenAI **that creates movies based on queries.**

### training data

*"Thrilling horror science fiction movie, plays in space in a distant future"*

$$X = X_1, X_2, \dots, X_N$$

each token is a vector  $X_n$  of length  $E$

*"Entertaining science fiction movie, plays in space in a distant past"*

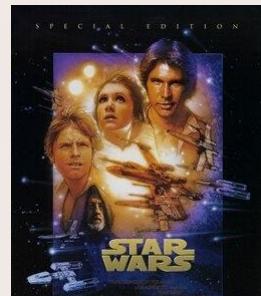
*"Boring love story, plays on a ship in the past"*

key

Alien



Star Wars



Titanic



value



### self attention

imagine you want to built & train a movie GenAI **that creates movies based on queries.**

### training data

each token is a vector  $X_n$  of length  $E$

$$X = X_1, X_2, \dots, X_N$$

key

Alien



value

*"Boring love story, plays in space."*

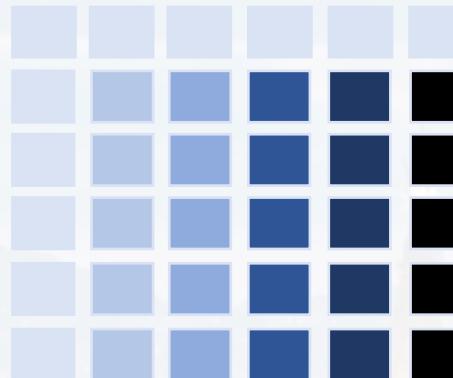
query





## self attention

key    value    query

*"Boring love story, plays on a ship in the past"*weights  $w_{nm}$ each token is a vector  $X_n$  of length  $E$  $X = X_1, X_2, \dots, X_N$ 

output:

$$Y_n = \sum_{m=1}^N w_{nm} X_m$$

- weights should be trainable
- weight = 0,  $X_m$  has no influence on output
- weights should be positive so that neg weights don't counteract positive weights
- normalization:  $\sum_{m=1}^N w_{nm} = 1$

for returning the best suggestion:

comparing key vector  $\textcolor{teal}{X}_n$  to query vector  $\textcolor{blue}{X}_m$  via dot product

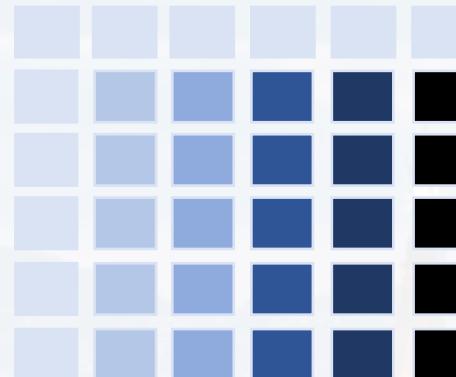
$$w_{nm} = \frac{\exp(\textcolor{teal}{X}_n \circ \textcolor{blue}{X}_m)}{\sum_{\mu=1}^N \exp(\textcolor{teal}{X}_n \circ \textcolor{blue}{X}_{\mu})}$$

softmax



## self attention

key    value    query

*"Boring love story, plays on a ship in the past"*weights  $w_{nm}$ each token is a vector  $X_n$  of length  $E$  $X = X_1, X_2, \dots, X_N$ 

output:

$$Y_n = \sum_{m=1}^N w_{nm} X_m$$

- weights should be trainable
- weight = 0,  $X_m$  has no influence on output
- weights should be positive so that neg weights don't counteract positive weights
- normalization:  $\sum_{m=1}^N w_{nm} = 1$

for returning the best suggestion:

comparing key vector  $X_n$  to query vector  $X_m$  via dot product

$$w_{nm} = \frac{\exp(X_n W(K) \circ X_m W(Q))}{\sum_{\mu=1}^N \exp(X_n W(K) \circ X_\mu W(Q))}$$

softmax



### self attention

key    value    query

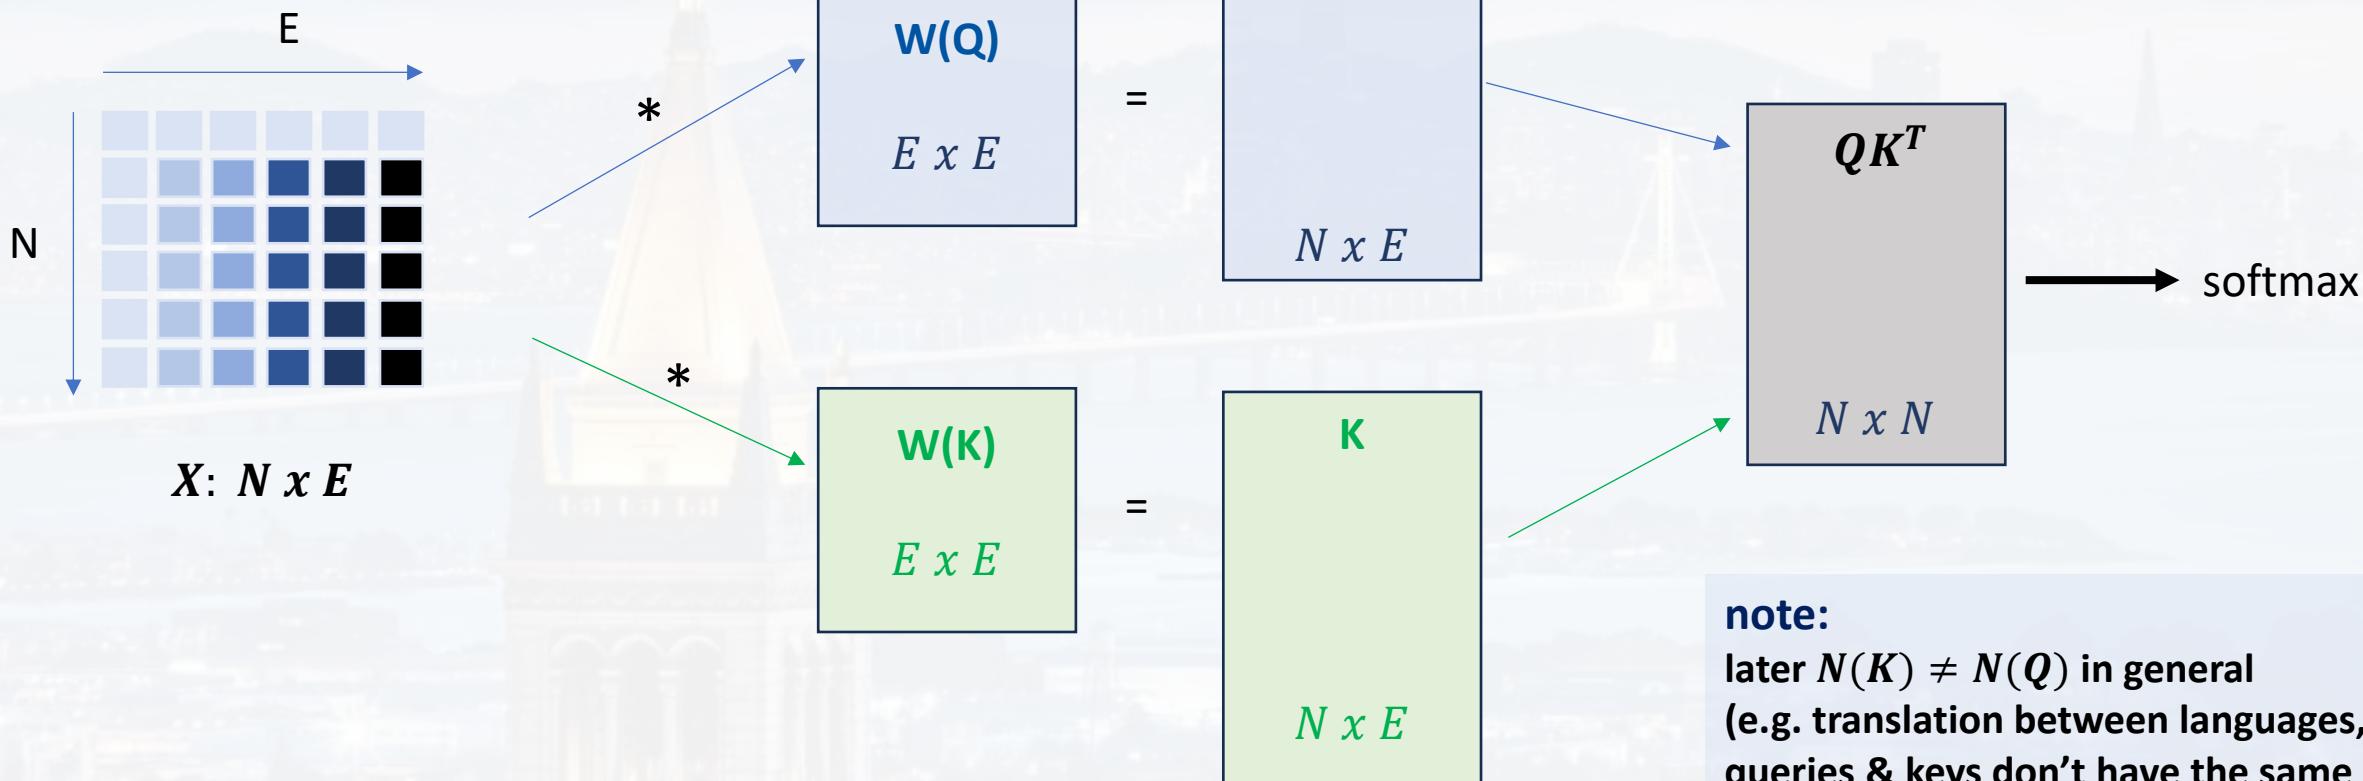
$$w_{nm} = \frac{\exp(X_n W(K) \circ X_m W(Q))}{\sum_{\mu=1}^N \exp(X_n W(K) \circ X_\mu W(Q))}$$

N: number of token

E: number of embedding dimensions

output:  $Y_n = \sum_{m=1}^N w_{nm} X_m$

"Boring love story, plays on a ship in the past"



### note:

later  $N(K) \neq N(Q)$  in general  
(e.g. translation between languages,  
queries & keys don't have the same  
number of tokens in general etc)



self attention

key    value    query

$$w_{nm} = \frac{\exp(\mathbf{X}_n \mathbf{W}(\mathbf{K}) \circ \mathbf{X}_m \mathbf{W}(\mathbf{Q}))}{\sum_{\mu=1}^N \exp(\mathbf{X}_n \mathbf{W}(\mathbf{K}) \circ \mathbf{X}_{\mu} \mathbf{W}(\mathbf{Q}))}$$

N: number of token

E: number of embedding dimensions

output:  $Y_n = \sum_{m=1}^N w_{nm} X_m = \sum_{m=1}^N \frac{\exp(\mathbf{X}_n \mathbf{W}(\mathbf{K}) \circ \mathbf{X}_m \mathbf{W}(\mathbf{Q}))}{\sum_{\mu=1}^N \exp(\mathbf{X}_n \mathbf{W}(\mathbf{K}) \circ \mathbf{X}_{\mu} \mathbf{W}(\mathbf{Q}))} X_m$

$$= \sum_{m=1}^N \text{softmax}(\mathbf{Q}_n \mathbf{K}_m^T) X_m$$

$$\rightarrow \sum_{m=1}^N \text{softmax}(\mathbf{Q}_n \mathbf{K}_m^T) X_m \mathbf{W}(\mathbf{V}) \rightarrow \mathbf{Y} = \text{softmax}(\mathbf{Q} \mathbf{K}^T) \mathbf{V}$$

↑  
summarizing the characteristics of the movie by using the movie itself

The output would be a movie, generated by weighted contributions of those movies (values), where the keys match well with the query



self attention

key    value    query

$$w_{nm} = \frac{\exp(\mathbf{X}_n \mathbf{W}(K) \circ \mathbf{X}_m \mathbf{W}(Q))}{\sum_{\mu=1}^N \exp(\mathbf{X}_n \mathbf{W}(K) \circ \mathbf{X}_{\mu} \mathbf{W}(Q))}$$

N: number of token

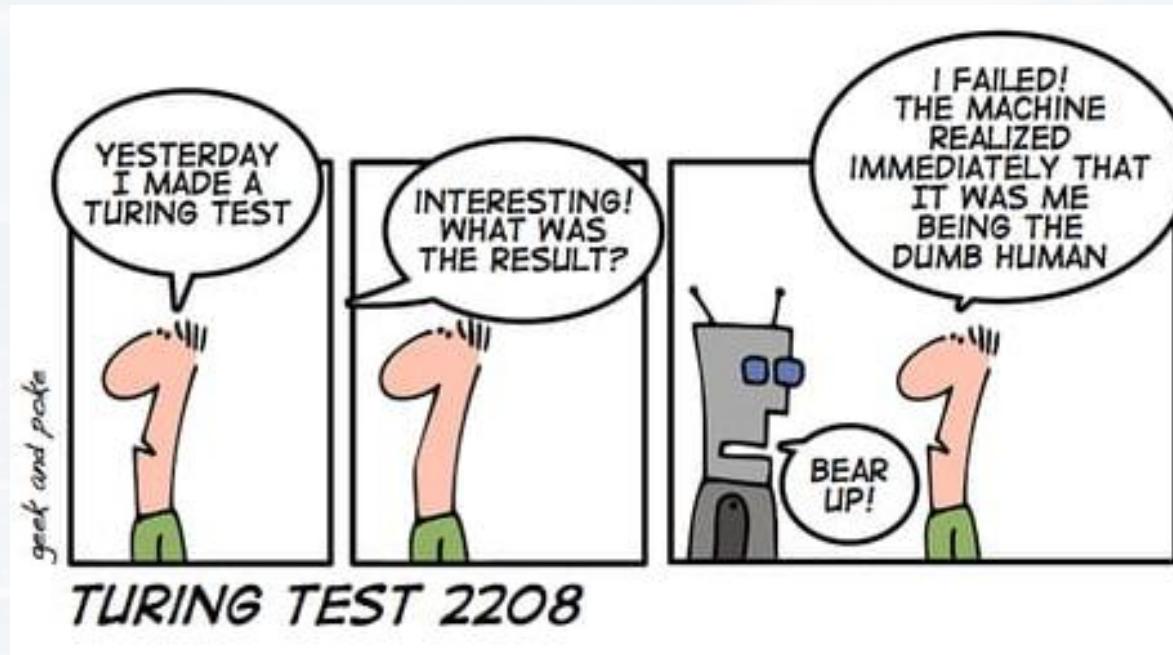
E: number of embedding dimensions

output:  $Y_n = \sum_{m=1}^N w_{nm} X_m$

$$\rightarrow \mathbf{Y} = \text{softmax}(\mathbf{Q} \mathbf{K}^T) \mathbf{V}$$

note:

- dot product scales with  $E$ , therefore normalization:  $\text{softmax}\left(\frac{\mathbf{Q}\mathbf{K}^T}{\sqrt{E}}\right)$
- cross attention:
  - eg. key a phrase in language A, query in language B
  - **encoder/decoder** structure, see next slides
- we want to recognize many underlying pattern → **multiple attention** layers in parallel → transformer



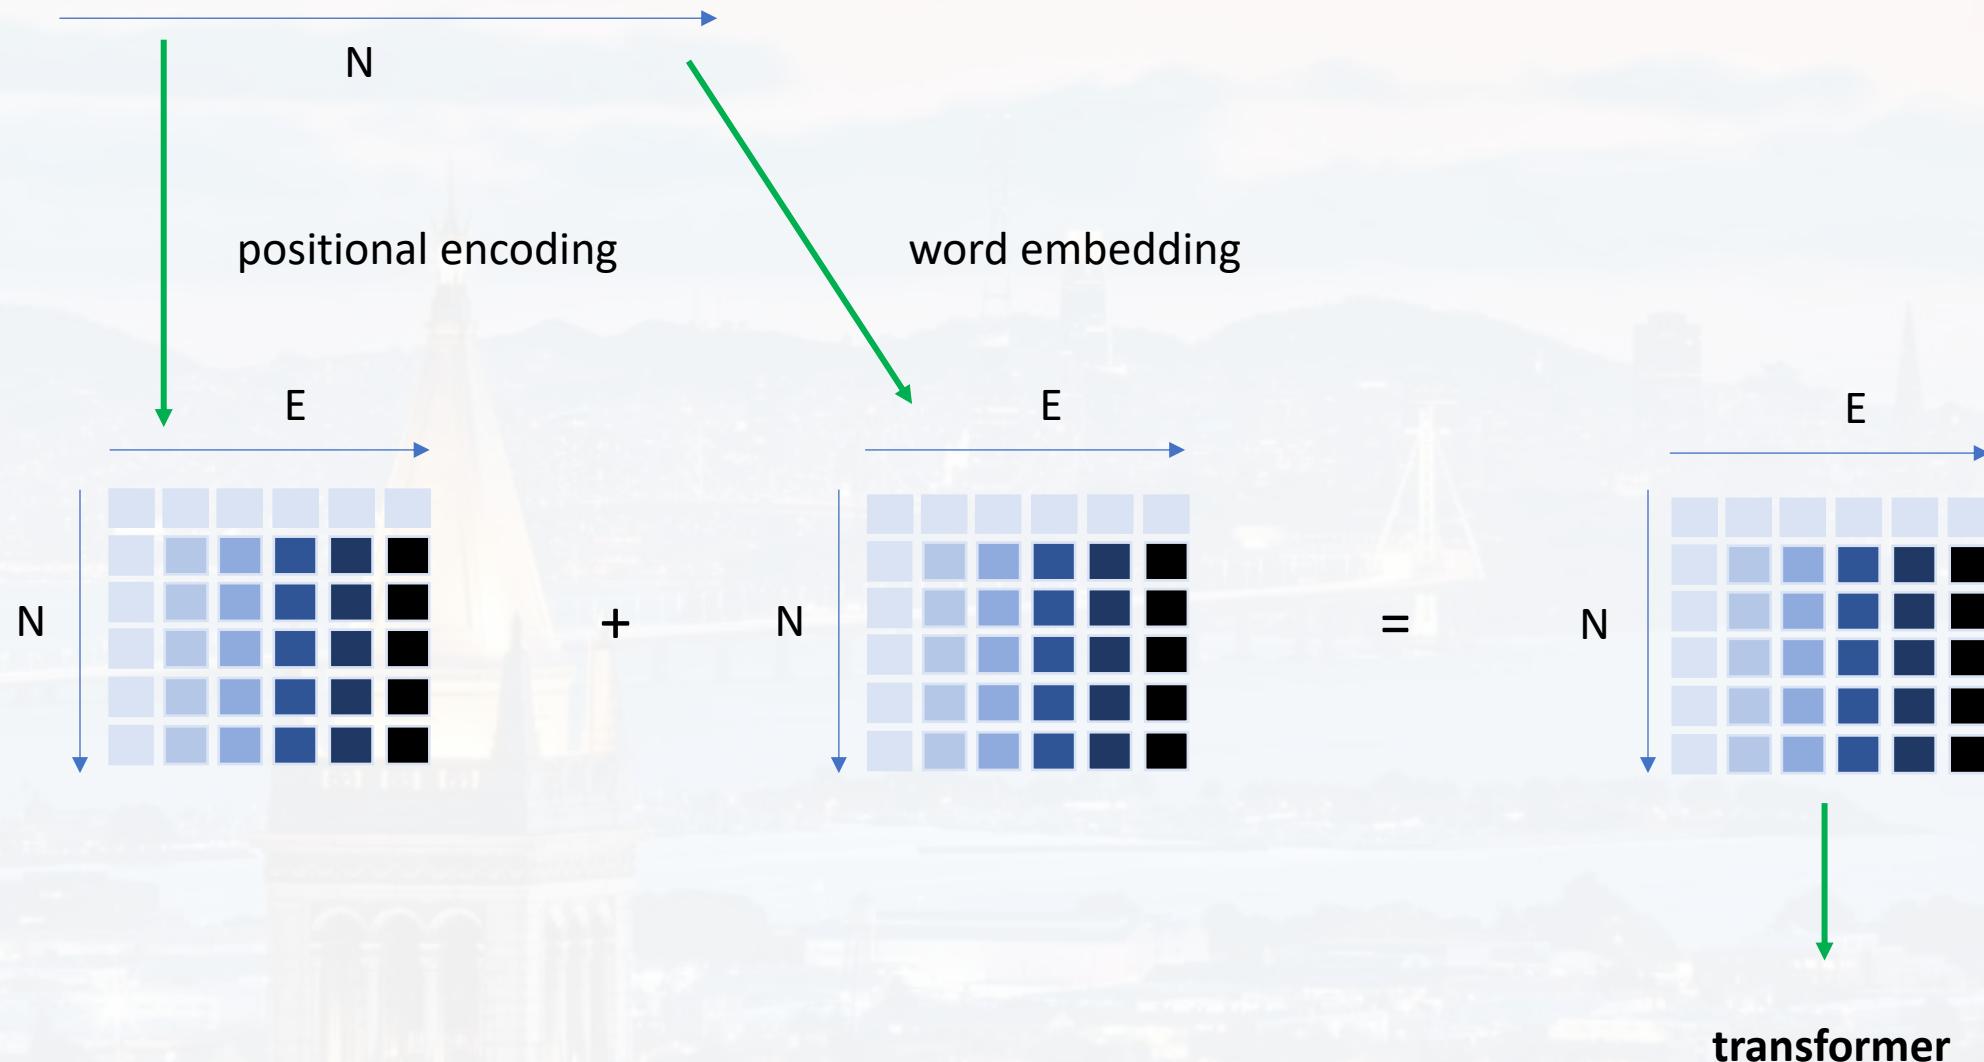
## Outline

- Introduction
- Bigram and MAP
- Positional Encoding
- Word Embedding
- Attention
- Transformer Architecture



*"The cat jumped on the roof."*

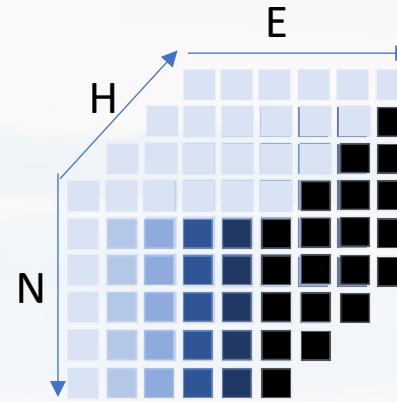
N: number of token  
E: number of embedding dimensions





within the transformer:

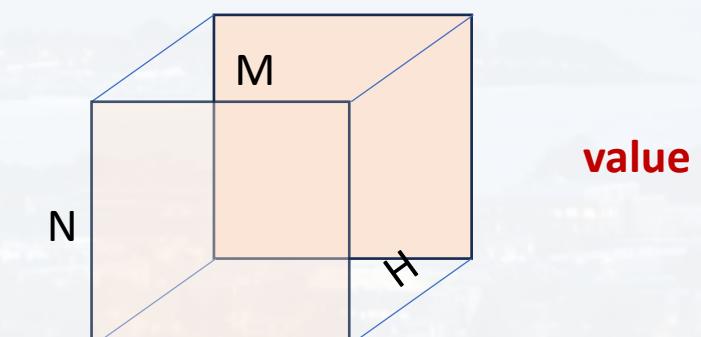
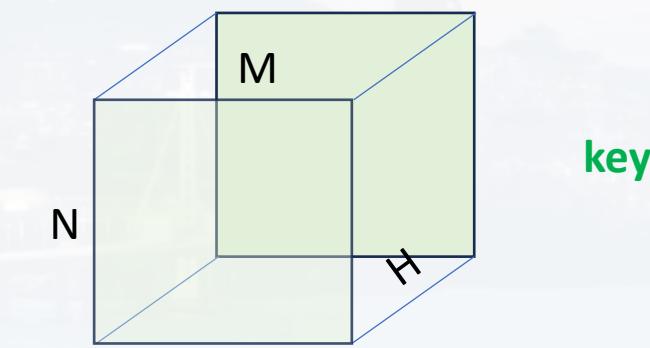
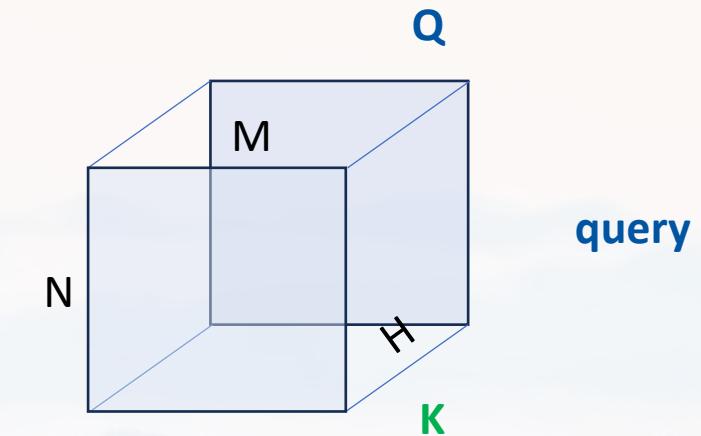
attention:



$$E \times W(Q) = Q$$

$$E \times W(K) = K$$

$$E \times W(V) = V$$



$W(Q)$ ,  $W(K)$ ,  $W(V)$ : learnable

N: number of token

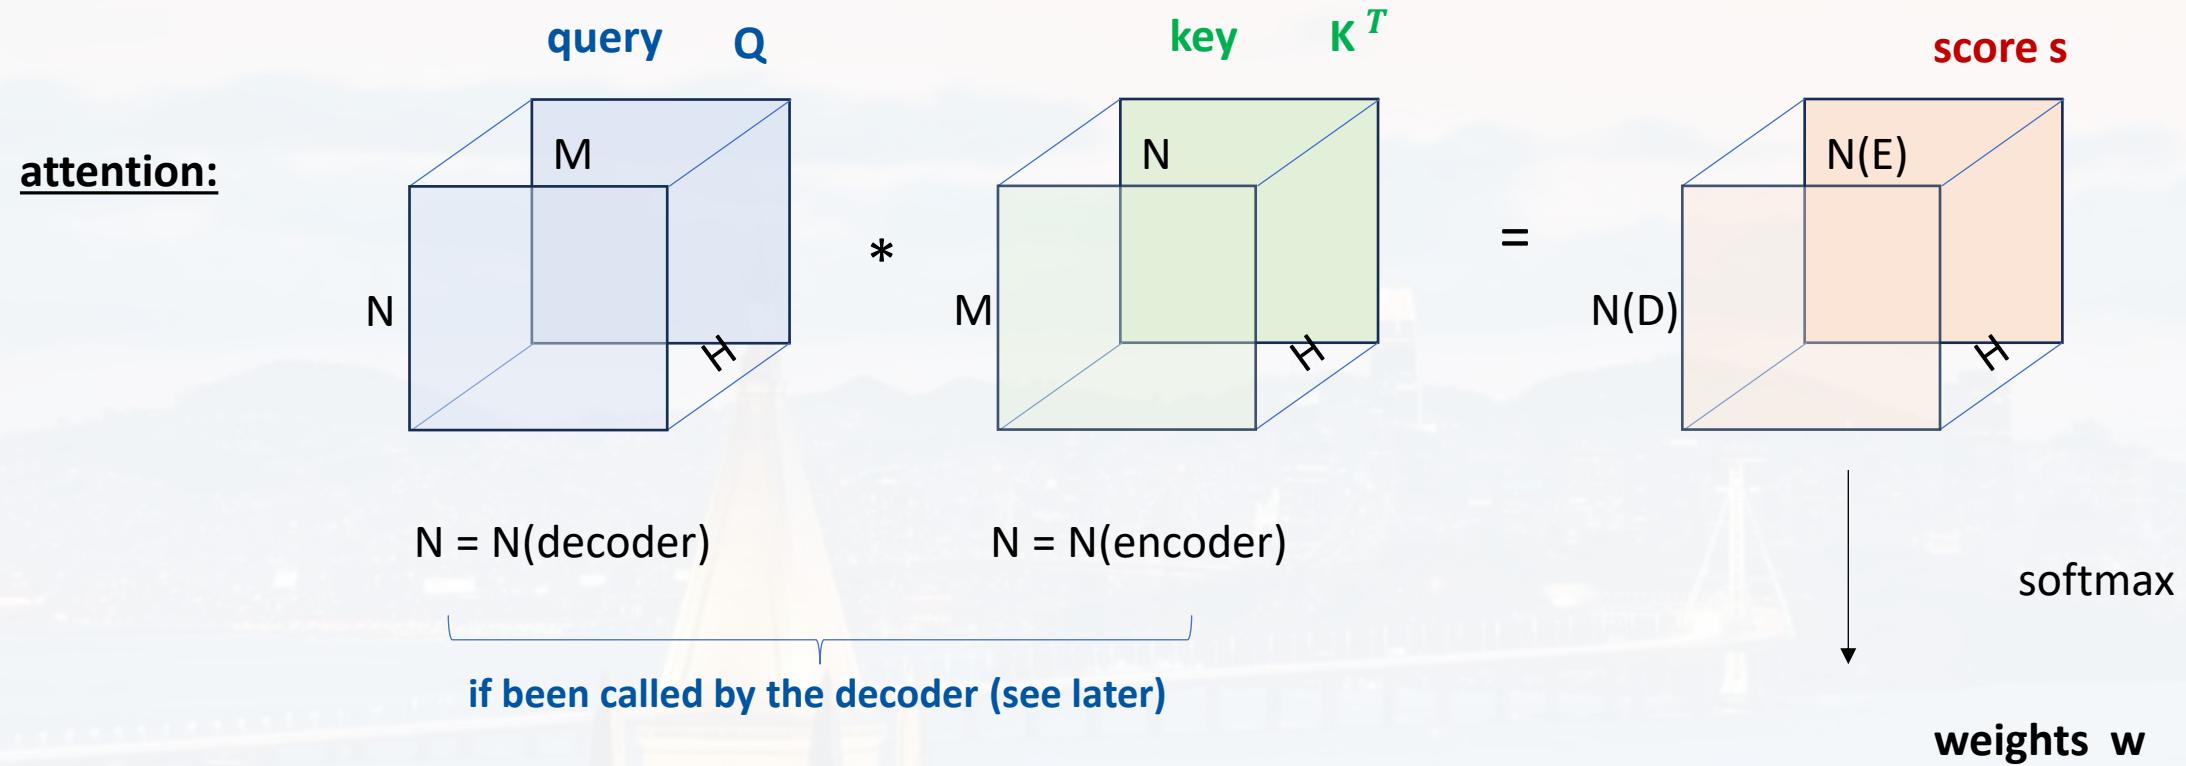
E: number of embedding dimensions

H: number of heads (= 8)

M: head size (= 64)

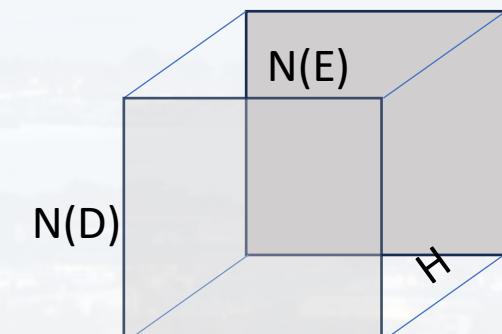


within the transformer:



$W(Q)$ ,  $W(K)$ ,  $W(V)$ : learnable

- N: number of token
- E: number of embedding dimensions
- H: number of heads (= 8)
- M: head size (= 64)

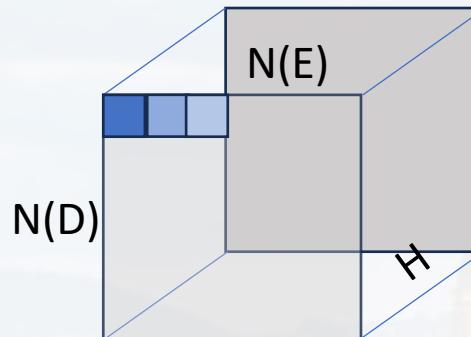




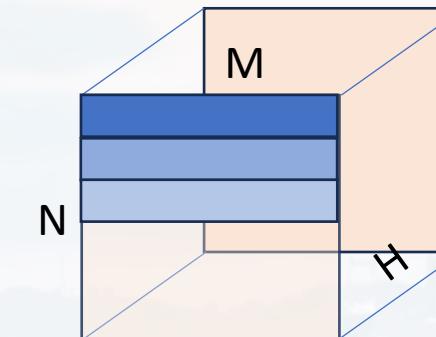
within the transformer:

attention:

weights w



value v



N:

number of token

E:

number of embedding dimensions

H:

number of heads (= 8)

M:

head size (= 64)

N(E)  
N(D)

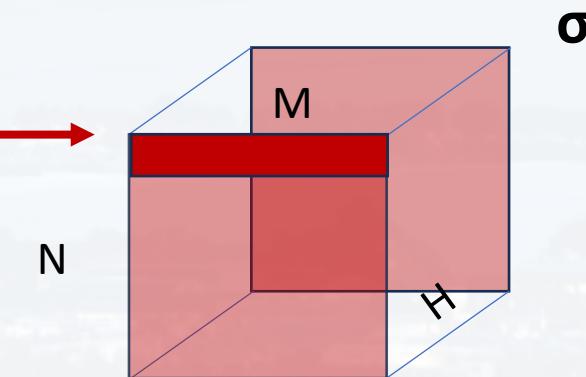
= N(encoder)

= N(decoder) , [see later](#)

$$Y_n = \sum_{m=1}^N w_{nm} X_m$$

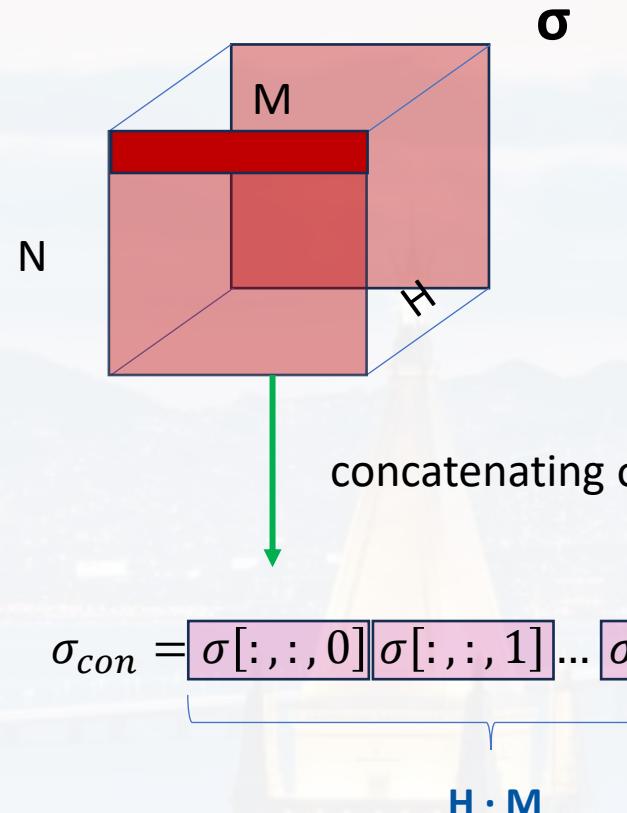
A mathematical expression showing the calculation of  $Y_n$ . It consists of a sum of terms:  $w_{n1} * X_1 + w_{n2} * X_2 + w_{n3} * X_3 + \dots$ . Each term is represented by a blue bar multiplied by a smaller blue square. Red brackets group the first three terms and the ellipsis, indicating they are summed. A red arrow points from the ellipsis to a 3D diagram representing the output.

$$\sigma[i,:,k] = \sum_j w[i,j,k] * v[j,:,k]$$





**within the transformer:**



N: number of token  
E: number of embedding dimensions  
H: number of heads (= 8)  
M: head size (= 64)

## output:

The diagram illustrates the output of a learnable  $Q$  function. On the left, a large dark red rectangle represents the input  $E$ . A blue bracket labeled  $H \cdot M$  indicates the dimension of the intermediate representation. To the right of the equals sign ( $=$ ) is a blue arrow pointing down, labeled  $N$ , indicating the final dimension of the output. The output is shown as a grid of colored squares, ranging from light blue to black, representing a probability distribution over  $N$  actions.

$E$

$H \cdot M$

$=$

$N$

$E$

$N = N(\text{decoder})$ ,

if been called by

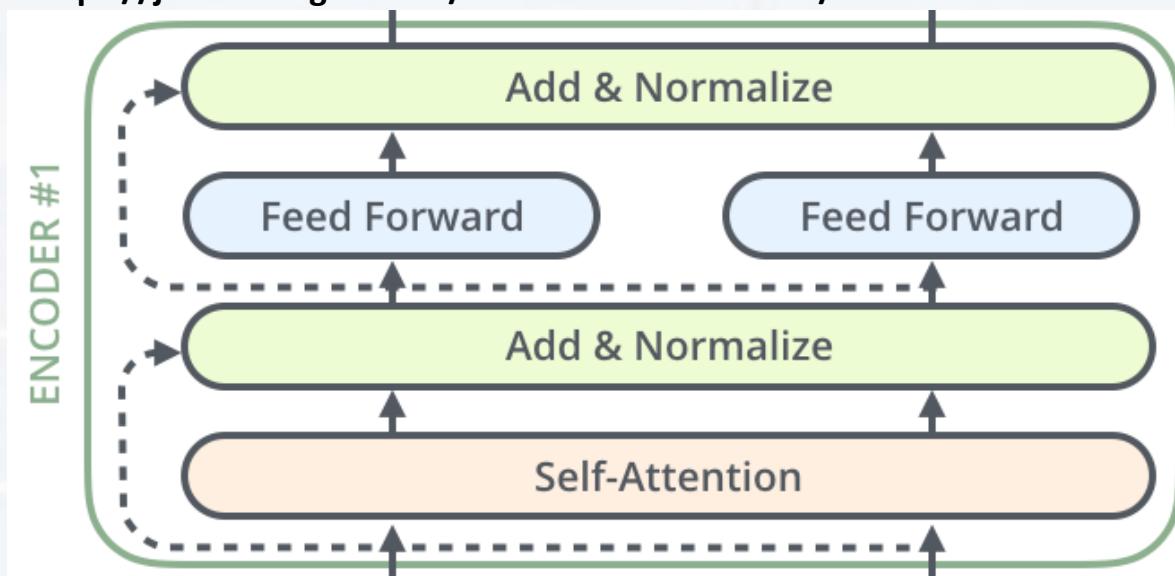
the decoder (see lat)

$N = N(\text{decoder})$ ,  
if been called by  
the decoder (see later)



that was attention → now: encoder

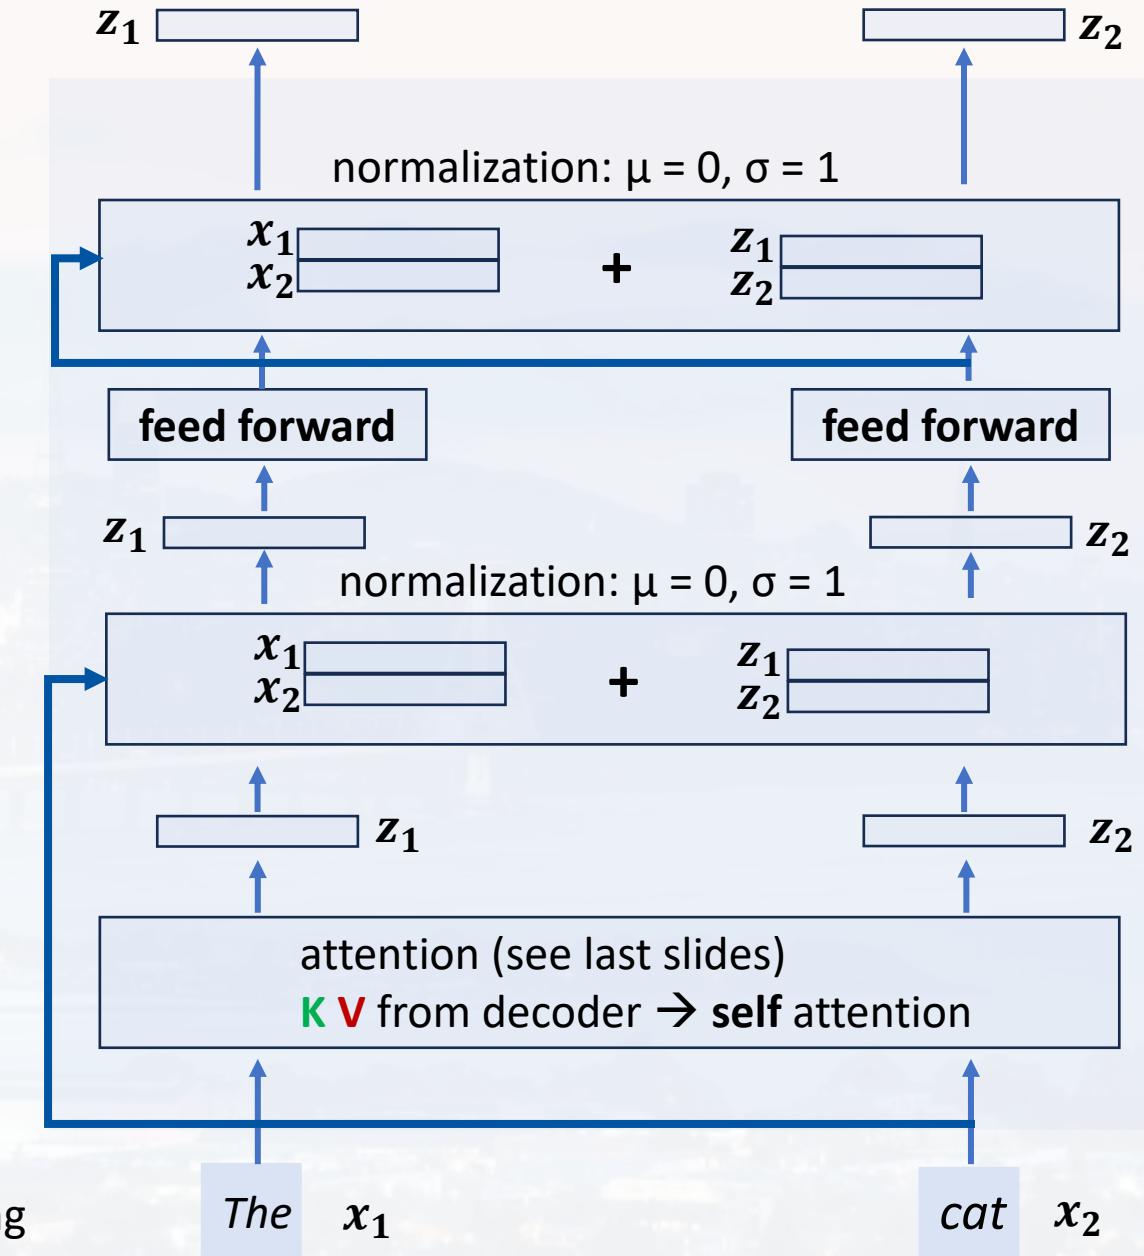
<https://jalammar.github.io/illustrated-transformer/>

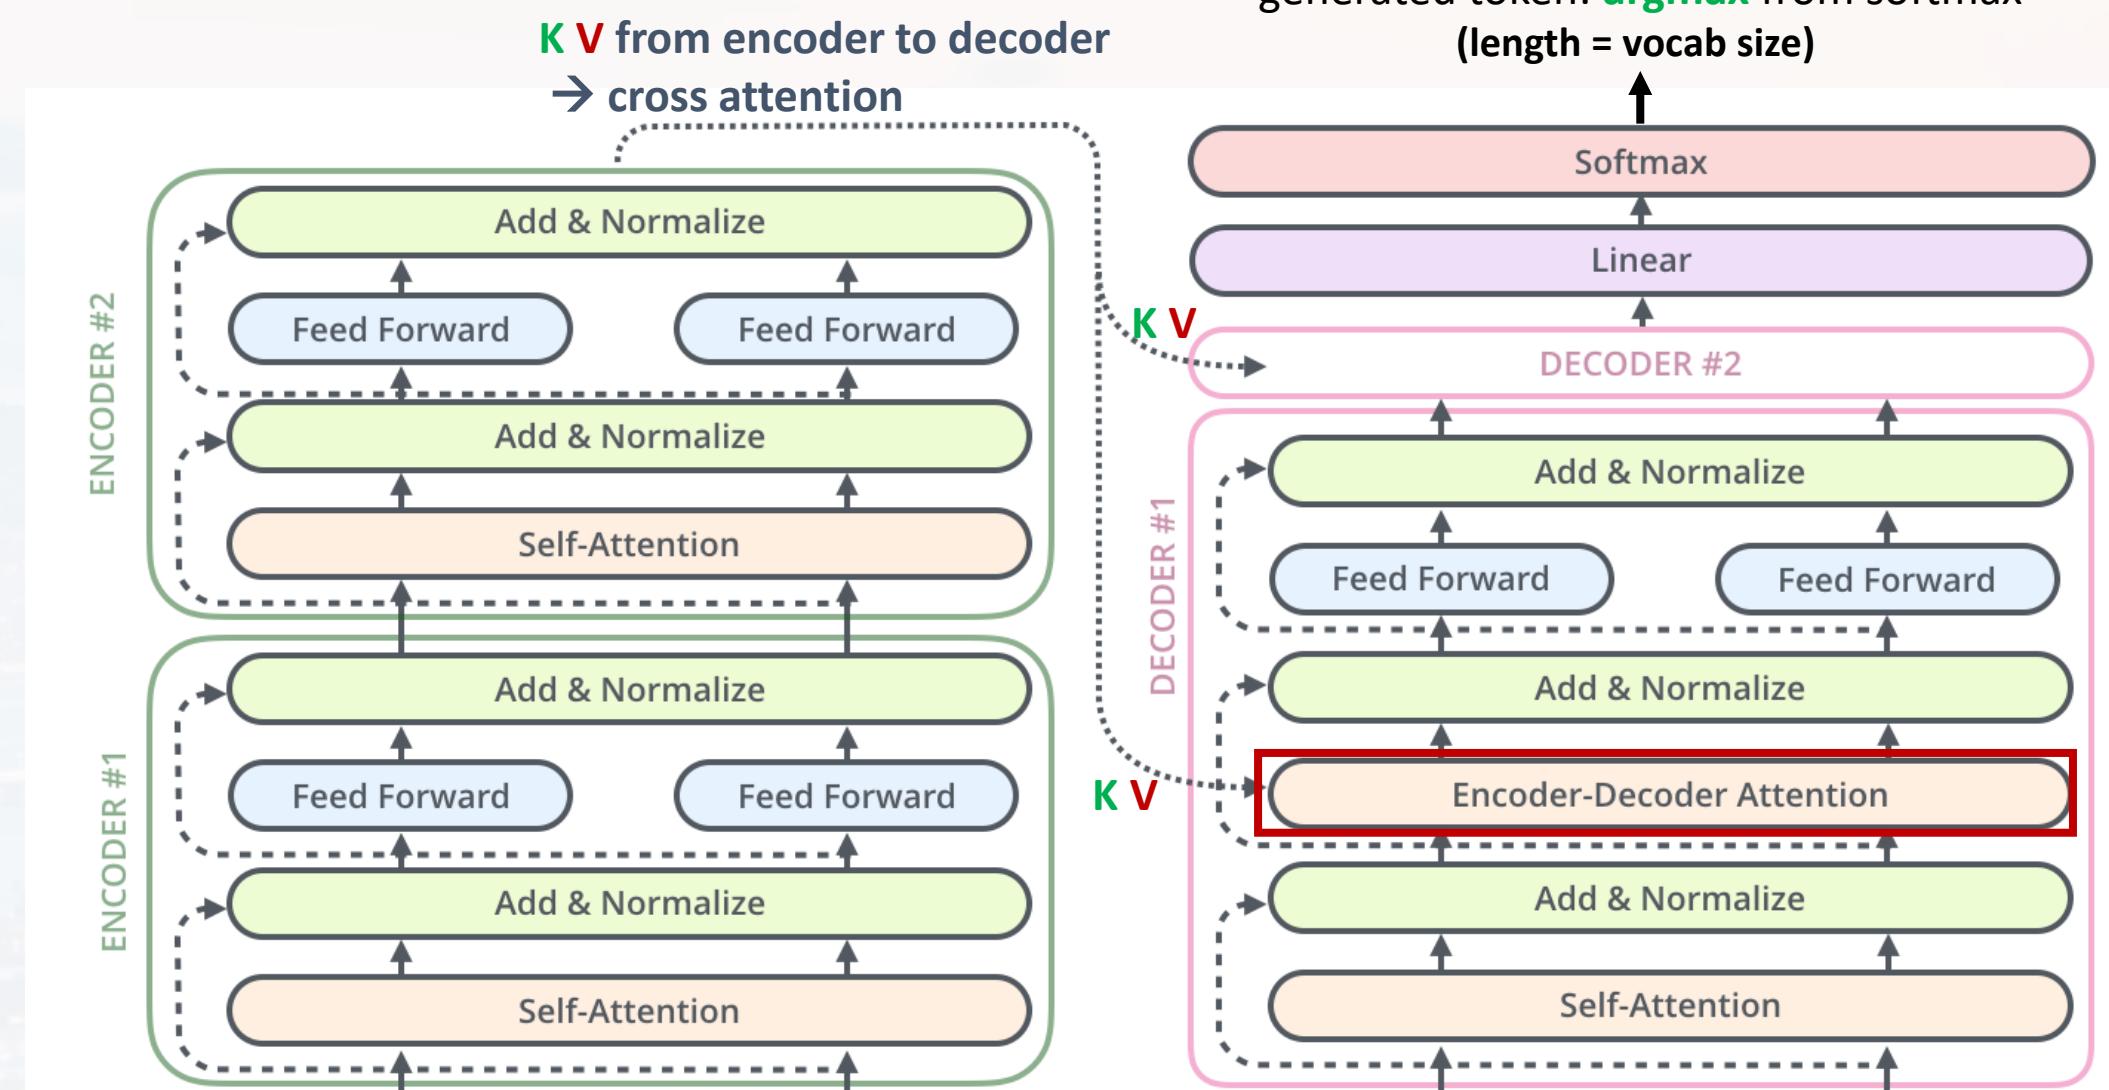


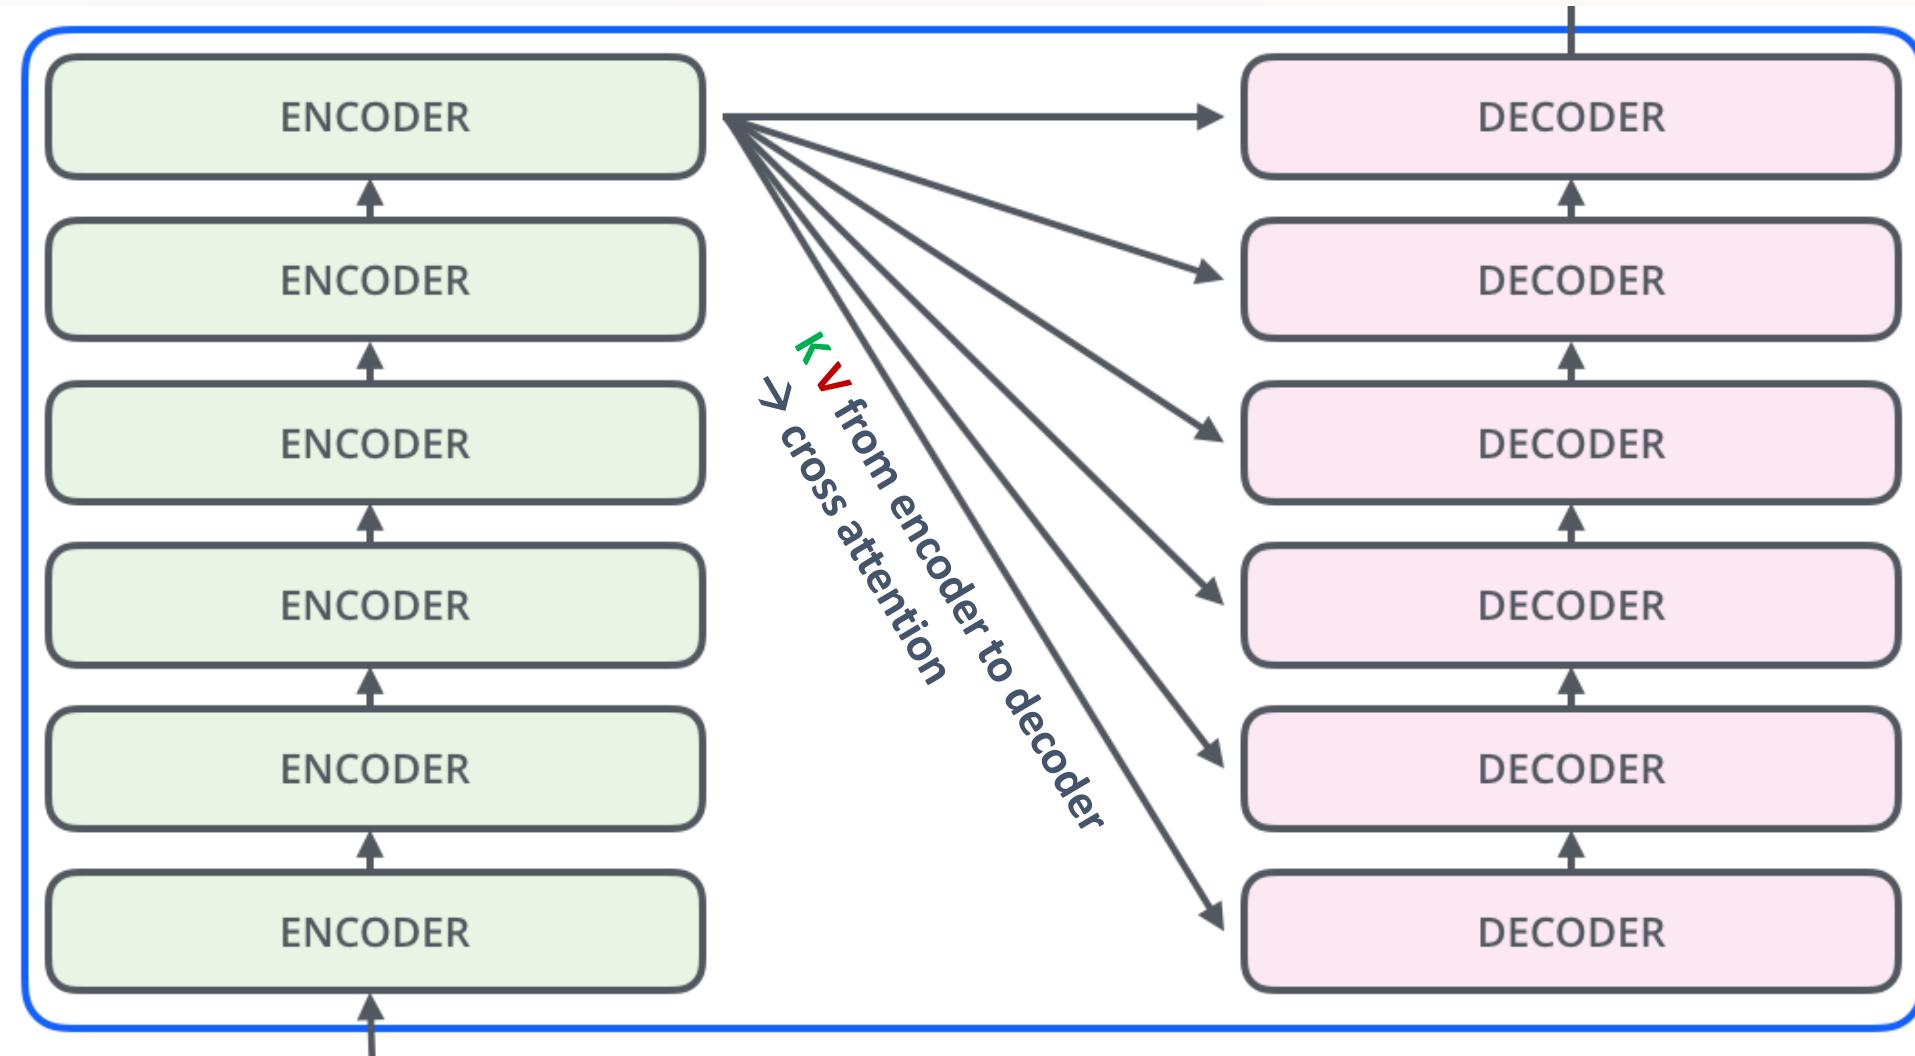
positional encoding

+

word embedding

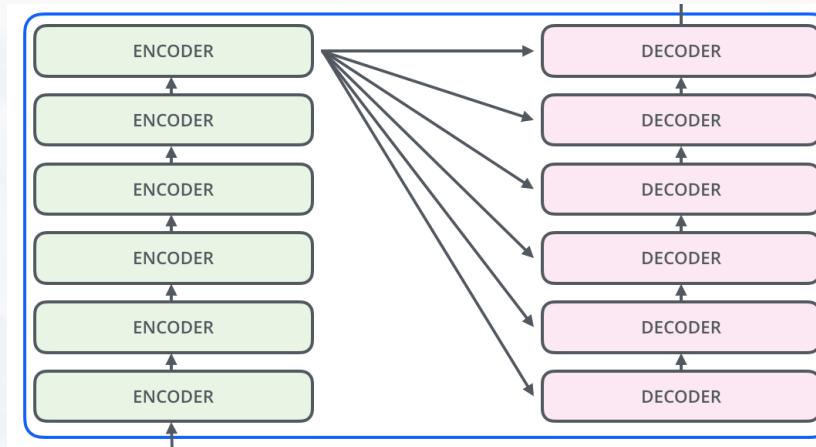








**K V** from encoder to decoder  
→ cross attention



**general:**  $N(E) \neq N(D)$

English: Let there be light.

$N(E) = 4$

German: Es werde Licht.

$N(D) = 3$

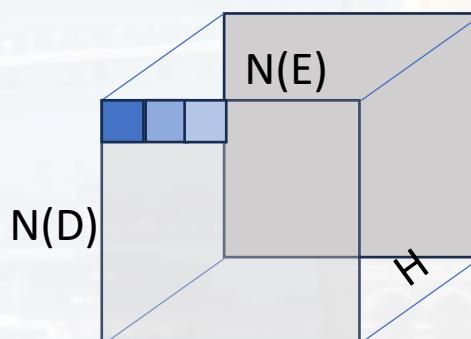
Latin: Fiat lux.

$N(D) = 2$

Hebrew: yehi ,or

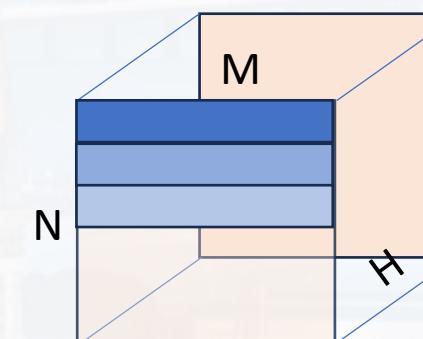
$N(D) = 2$

**weights w**

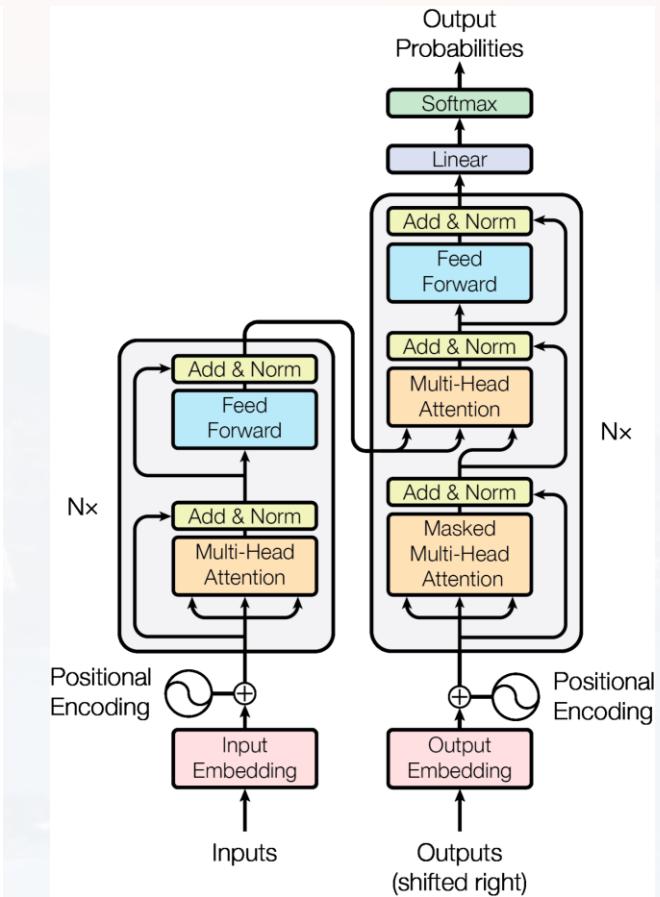
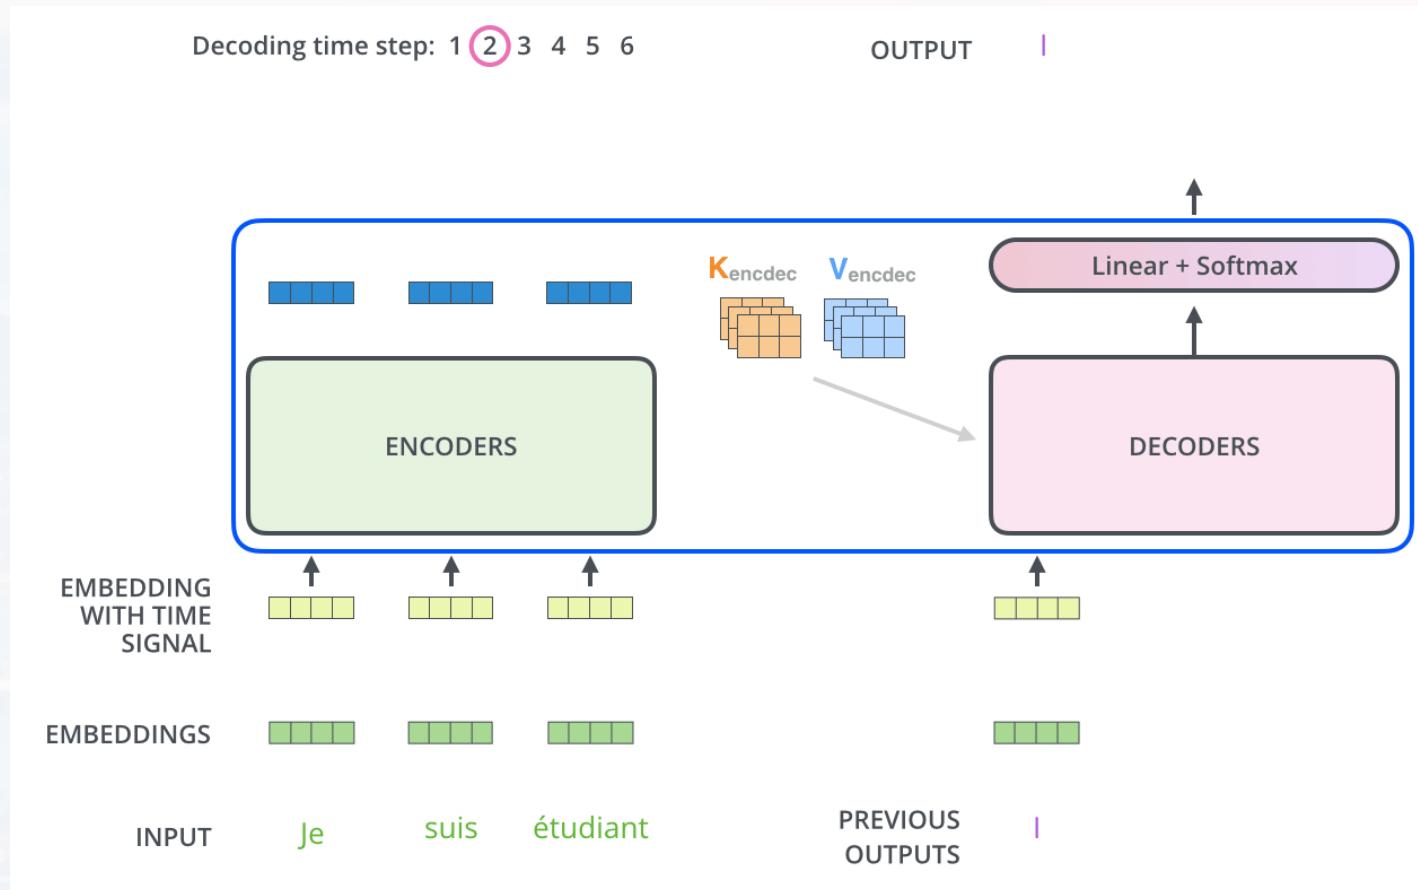


**attention:**

**value V**



$N = N(\text{encoder})$  , if been called by the decoder



improvements:

masked attention

$$\mathbf{Y} = \text{softmax}(\mathbf{QK}^T)\mathbf{V} \rightarrow \mathbf{Y} = \text{softmax}(\mathbf{QK}^T + \mathcal{M})\mathbf{V}$$

mask:  $\mathcal{M}$ **causal masking:**  $w_{nn+k} = 0 \forall k > 0$ 

$$\begin{bmatrix} 1 & 0 & 0 & 0 & 0 & 0 \\ 1 & 1 & 0 & 0 & 0 & 0 \\ 1 & 1 & 1 & 0 & 0 & 0 \\ 1 & 1 & 1 & 1 & 0 & 0 \\ 1 & 1 & 1 & 1 & 1 & 0 \\ 1 & 1 & 1 & 1 & 1 & 1 \end{bmatrix}$$

token at position  $t$  should only consider **previous token** for predicting token at  $t + 1$  (natural languages)

**padding masking:** batches of sequences might have different lengths,  
→ shorter sequences are padded with special tokens.  
→ model learns to ignore padding tokens  
→ for inference

**individual masking:** we often know that some token can't appear after each other (natural languages)



### improvements:

sampling strategies

**vanilla:**

returning most **probable token**, from which we calculate the probabilities for the next token and return the most likely one etc

**beam search:**

we store  $b$  (= **beam width**) sequences of length  $n$  and then return the **most likely sequence**

$$P(X_1 X_2 X_3 X_4 X_5 \dots X_n) = P(X_n | X_{n-1} \dots X_1) P(X_{n-1} | X_{n-2} \dots X_1) \dots P(X_1)$$

**top K-sampling:**

consider  **$K$  most probable token**  
→ renormalize their probabilities  
→ draw randomly from these  $K$  token

$$p_k = \frac{\exp(\pi_k/T)}{\sum_{k=1}^K \exp(\pi_k/T)}$$

$\pi_k$ :  
 $p_k$ :  
 $T$ :

**probability**  
**renormalized probability**  
**"softening" parameter**

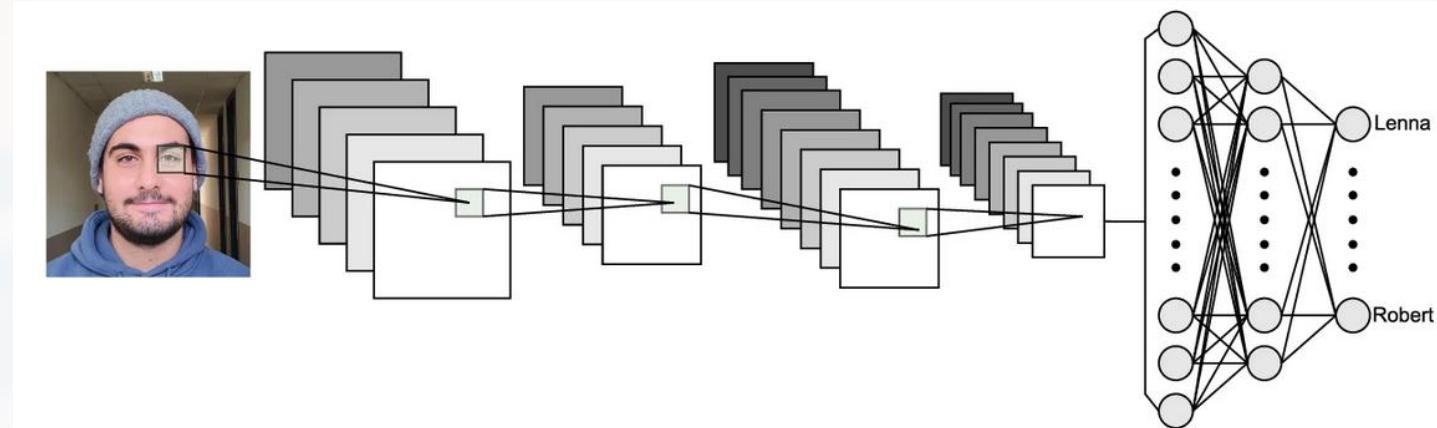
**top P-sampling:**

like top  $K$ , but for sequences (see beam search)

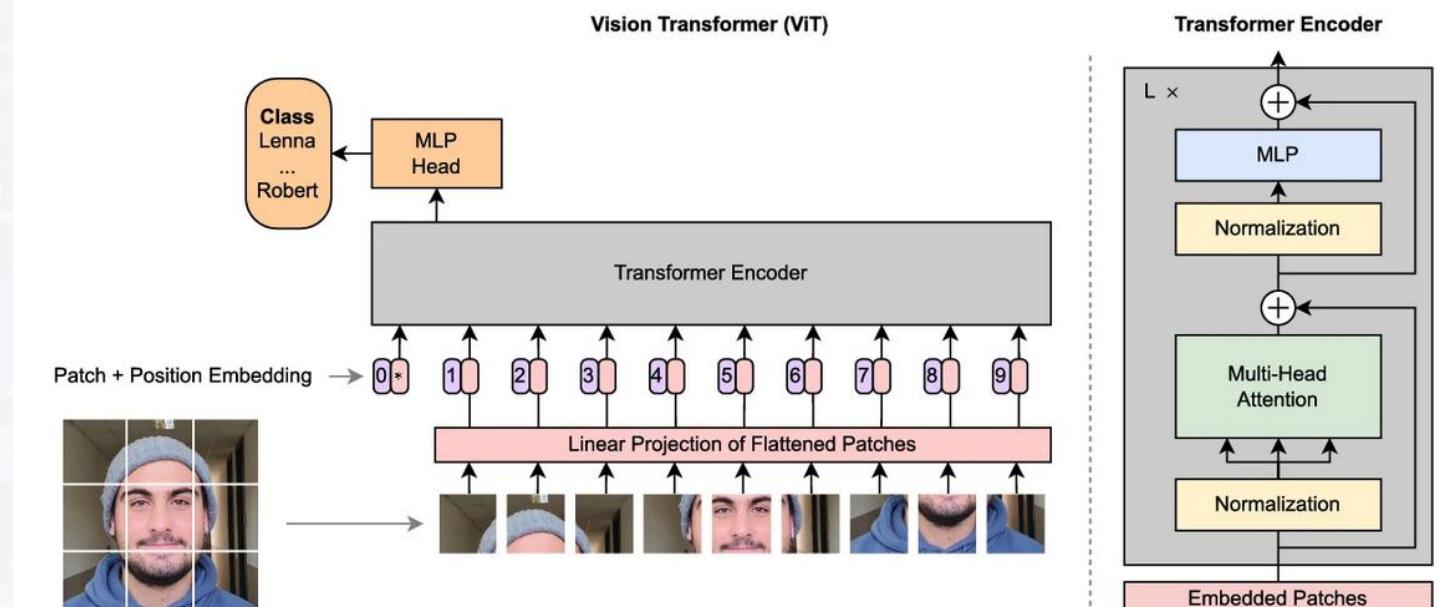


### improvements:

vision transformer



(a) Common CNN architecture





### more about transformers:

[Jay Alammar](#)

[Interactive Visualization](#)

[transformers intro](#)



### Misra Turp

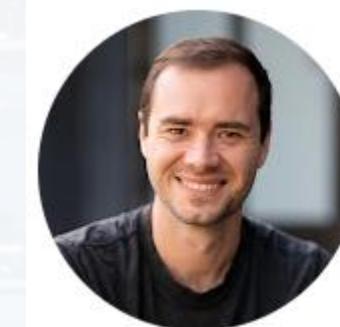
@misraturp · 40.3K subscribers · 163 videos

Here is where we learn! This is a space to take it

[misraturp.com/roadmap](#) and 3 more links

[Subscribe](#)

[building GPT from scratch!](#)



### Andrej Karpathy

@AndrejKarpathy · 451K subscribers · 14 videos

[FAQ ...more](#)

[karpathy.ai](#) and 2 more links

[Subscribe](#)



Thank you very much for your

