

Lecture 15:

Graph Neural Networks (GNN)



Markus Hohle
University California, Berkeley

Machine Learning Algorithms
MSSE 277B, 3 Units



Lecture 1: Course Overview and Introduction to Machine Learning

Lecture 2: Bayesian Methods in Machine Learning

classic ML tools & algorithms

Lecture 3: Dimensionality Reduction: Principal Component Analysis

Lecture 4: Linear and Non-linear Regression and Classification

Lecture 5: Unsupervised Learning: K-Means, GMM, Trees

Lecture 6: Adaptive Learning and Gradient Descent Optimization Algorithms

Lecture 7: Introduction to Artificial Neural Networks - The Perceptron

ANNs/AI/Deep Learning

Lecture 8: Introduction to Artificial Neural Networks - Building Multiple Dense Layers

Lecture 9: Convolutional Neural Networks (CNNs) - Part I

Lecture 10: CNNs - Part II

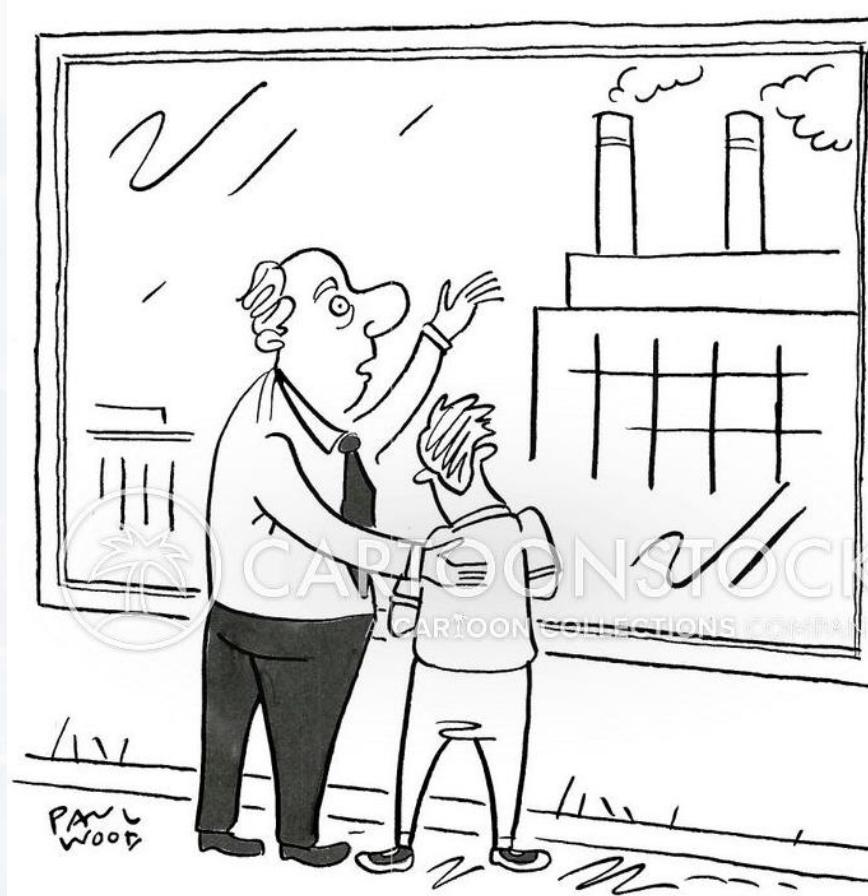
Lecture 11: Recurrent Neural Networks (RNNs) and Long Short-Term Memory (LSTMs)

Lecture 12: Combining LSTMs and CNNs

Lecture 15: Transformer

Lecture 14: Project Presentations

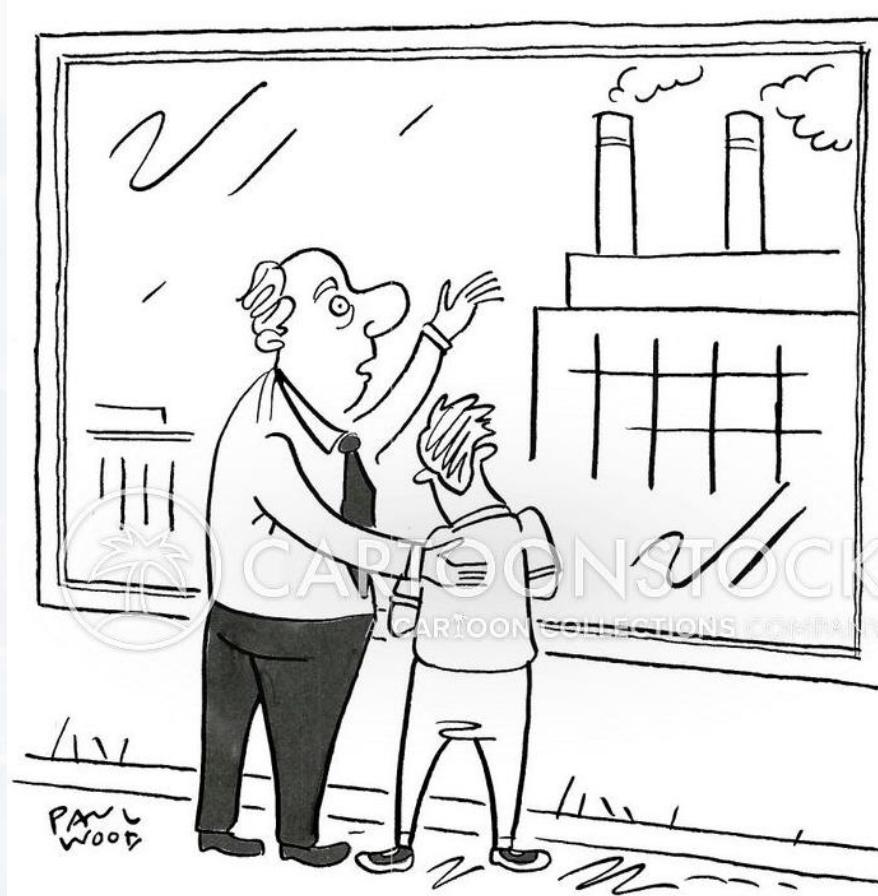
Lecture 15: GNN



ONE DAY SON, ALL THIS
WILL BE RUN BY ROBOTS

Outline

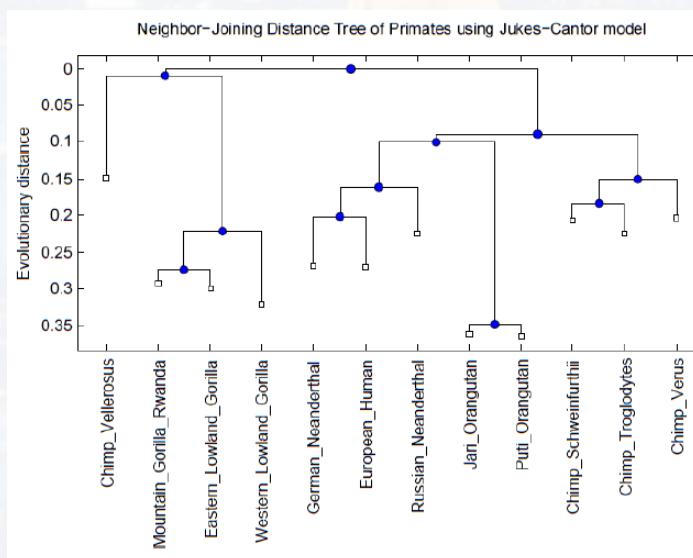
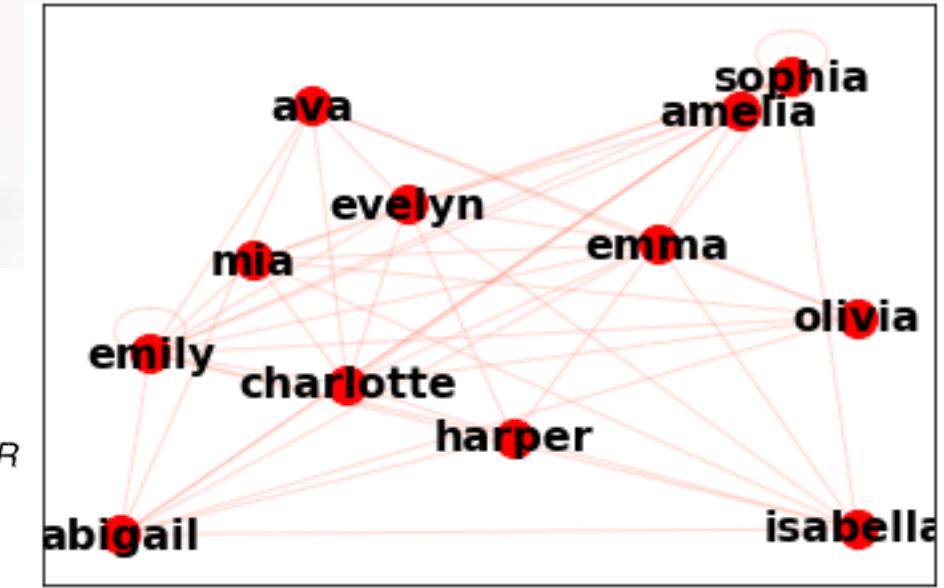
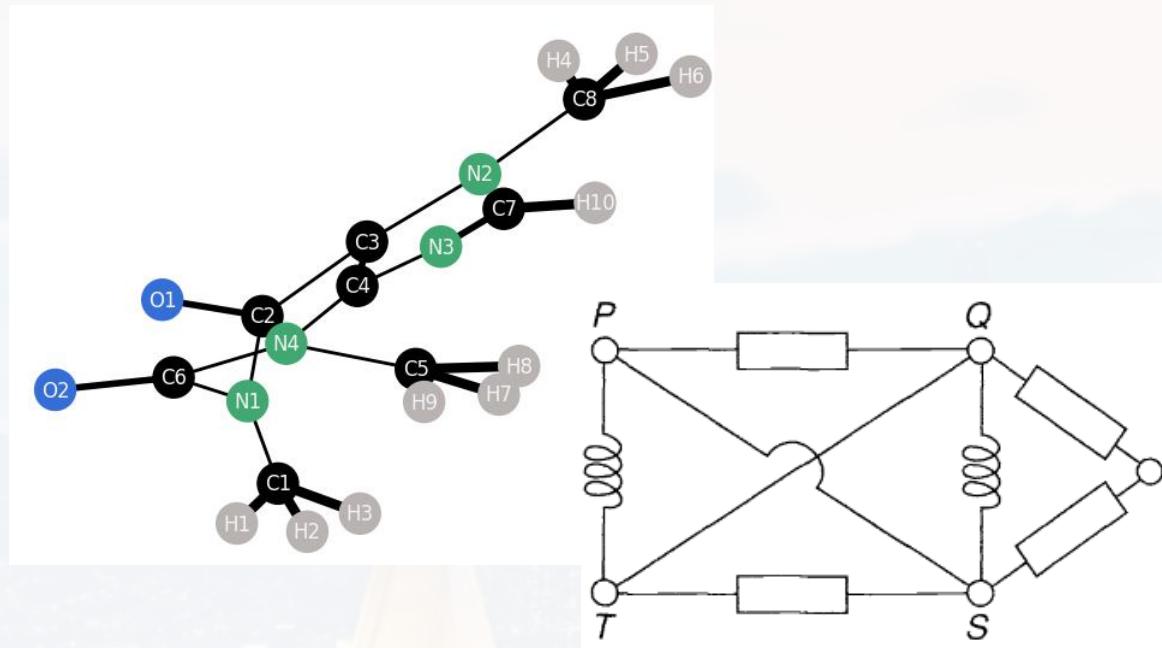
- What is a Graph
- The ANN Part
- PyTorch Example



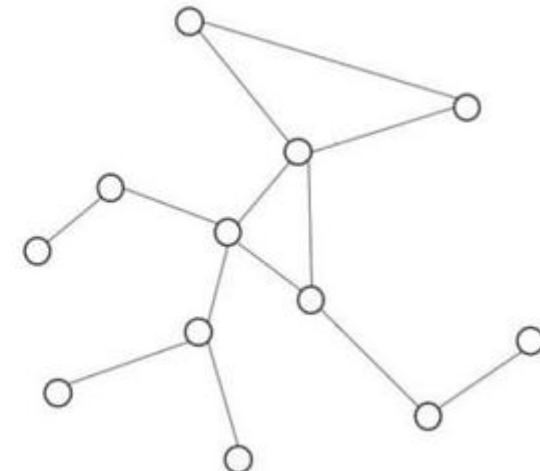
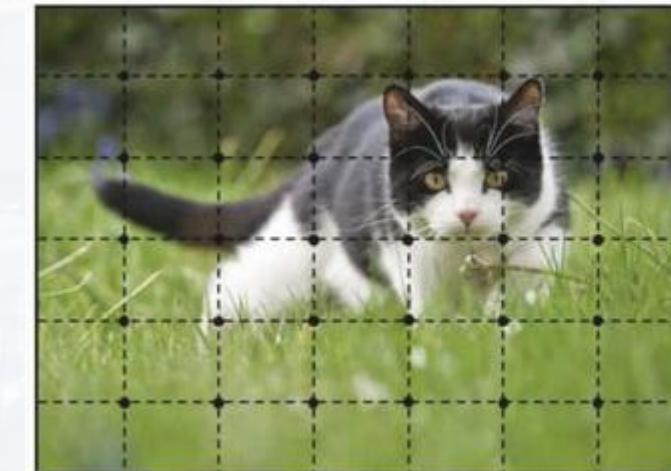
ONE DAY SON, ALL THIS
WILL BE RUN BY ROBOTS

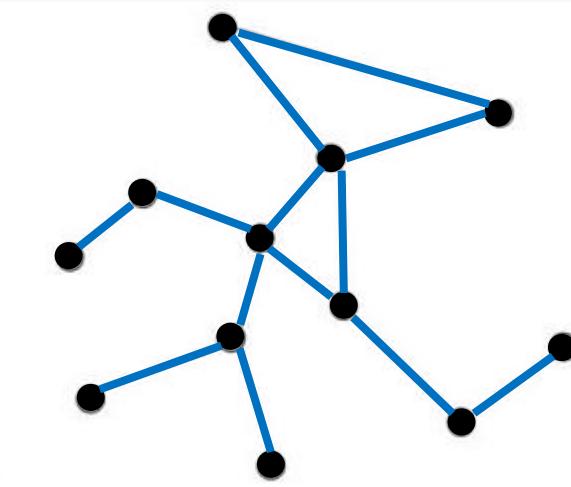
Outline

- What is a Graph
- The ANN Part
- PyTorch Example



<https://doi.org/10.1016/j.aiopen.2021.01.001>





Graph G

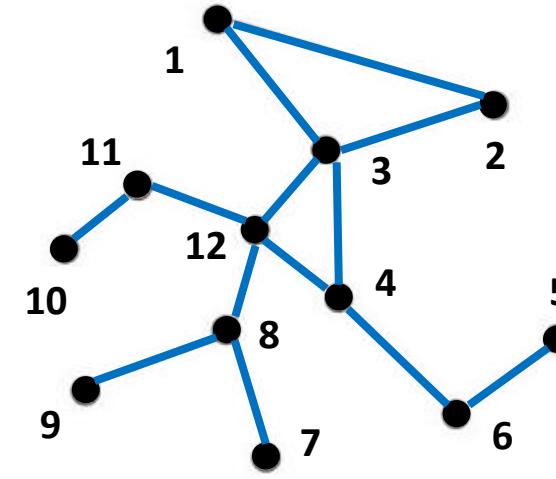
nodes N (vertices V)

edges E

$G = G(N, E)$

- social networks
- street maps
- workflows/planning
- biological signal pathways
- image processing

- nodes can have **features**
molecules: mass/ electronegativity
people: age, income, sex, ...
- edges can have **attributes**
molecules: bond length/strength
people: relations (work, friend, family)



structural information: **adjacency matrix A**

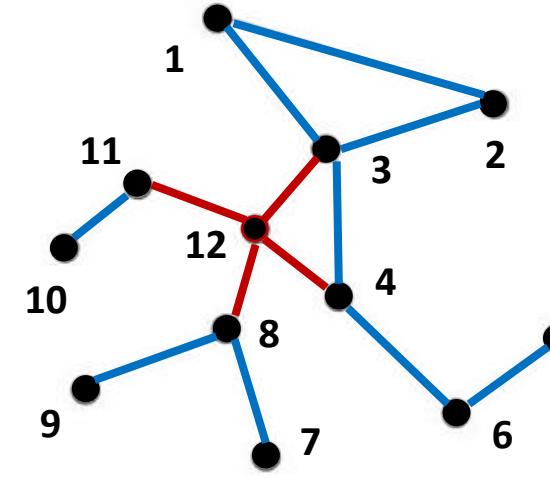
$A_{ij} = 1$ if $(n_i, n_j) \in E$
(nodes n_i and n_j have a common edge)

$A_{ij} = 0$ else

Graph $G = G(N, E)$

nodes N (vertices V)
edges E

$$A = \begin{pmatrix} 0 & 1 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 1 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 1 & 0 & 0 & 1 & 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 1 & 1 & 0 & 0 & 0 & 1 & 0 & 0 & 1 \end{pmatrix}$$



structural information: **adjacency matrix A**

$A_{ij} = 1$ if $(n_i, n_j) \in E$
(nodes n_i and n_j have a common edge)

$A_{ij} = 0$ else

node 12 has four first degree neighbors

degree d of a node

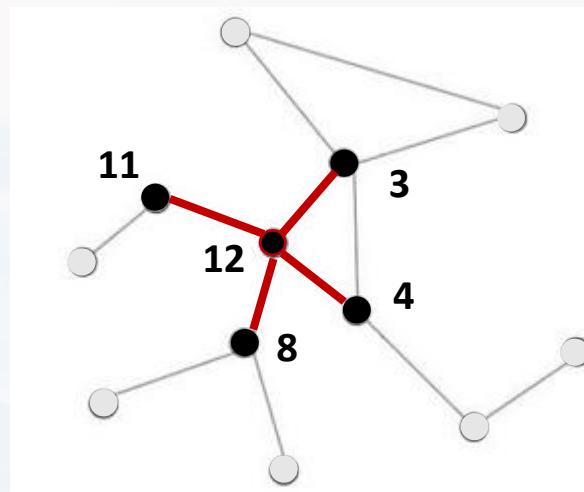
$$d(n_i) = \sum_j A_{ij}$$

$$A = \begin{pmatrix} 0 & 1 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 1 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 1 & 0 & 0 & 1 & 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 \end{pmatrix}$$

 $\boxed{\begin{matrix} 0 & 0 & \textcolor{red}{1} & \textcolor{red}{1} & 0 & 0 & 0 & \textcolor{red}{1} & 0 & 0 & \textcolor{red}{1} & 0 \end{matrix}}$

Graph $G = G(N, E)$

nodes N (vertices V)
edges E



structural information: **adjacency matrix A**

$A_{ij} = 1$ if $(n_i, n_j) \in E$
(nodes n_i and n_j have a common edge)

$A_{ij} = 0$ else

node 12 has four first degree neighbors

degree d of a node

$$d(n_i) = \sum_j A_{ij}$$

first degree neighborhood \mathcal{N}

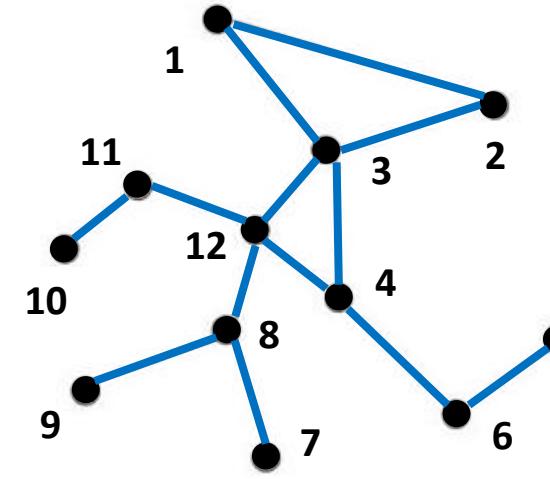
$$\mathcal{N}(n_i) = \{n_j \in N: (n_i, n_j) \in E\}$$

$$A = \begin{pmatrix} 0 & 1 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 1 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 1 & 0 & 0 & 1 & 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 1 \end{pmatrix}$$

$$\begin{matrix} 0 & 0 & \boxed{1} & \boxed{1} & 0 & 0 & 0 & \boxed{1} & 0 & 0 & \boxed{1} & 0 \end{matrix}$$

Graph $G = G(N, E)$

nodes N (vertices V)
edges E



A graph can have **loops**

structural information: **adjacency matrix A**

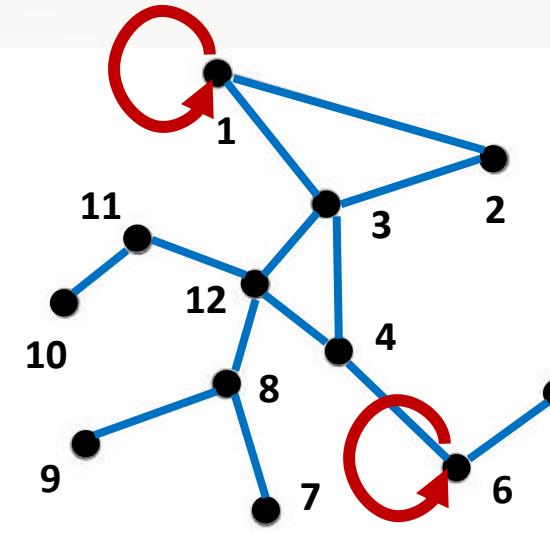
$A_{ij} = 1$ if $(n_i, n_j) \in E$
(nodes n_i and n_j have a common edge)

$A_{ij} = 0$ else

$$A = \begin{pmatrix} 0 & 1 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 1 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 1 & 0 & 0 & 1 & 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 \end{pmatrix}$$

Graph $G = G(N, E)$

nodes N (vertices V)
edges E



A graph can have **loops**

structural information: **adjacency matrix A**

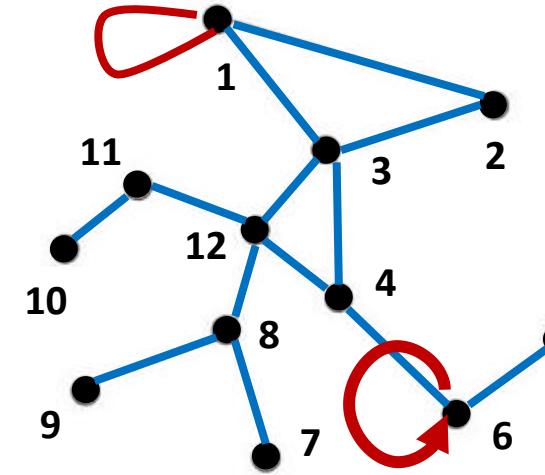
$A_{ij} = 1$ if $(n_i, n_j) \in E$
(nodes n_i and n_j have a common edge)

$A_{ij} = 0$ else

$$A = \begin{pmatrix} 1 & 1 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 1 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 1 & 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 1 & 1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 1 \\ 0 & 0 & 1 & 1 & 0 & 0 & 0 & 1 & 0 & 0 & 1 & 0 \end{pmatrix}$$

Graph $G = G(N, E)$

nodes N (vertices V)
edges E



A graph can have **loops**

note:

$d(n_1) = 4$, since loop is **undirected** and hits the node twice!

structural information: **adjacency matrix A**

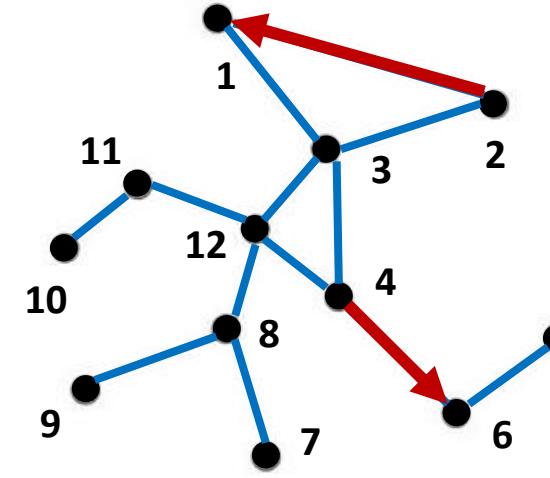
$A_{ij} = 1$ if $(n_i, n_j) \in E$
(nodes n_i and n_j have a common edge)

$A_{ij} = 0$ else

Graph $G = G(N, E)$

nodes N (vertices V)
edges E

$$A = \begin{pmatrix} 2 & 1 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 1 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 1 & 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 1 & 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 1 & 1 & 0 & 0 & 0 & 1 & 0 & 0 & 1 & 0 \end{pmatrix}$$



A graph can be **directed**

structural information: **adjacency matrix A**

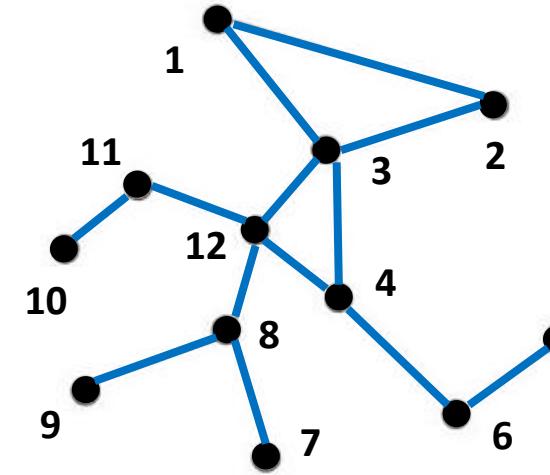
$A_{ij} = 1$ if $(n_i, n_j) \in E$
(nodes n_i and n_j have a common edge)

$A_{ij} = 0$ else

$$A = \begin{pmatrix} 0 & \textcolor{red}{0} & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ \textcolor{red}{1} & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 1 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 1 & 0 & 0 & 0 & \textcolor{red}{1} & 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & \textcolor{red}{0} & 1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 1 & 1 & 0 & 0 & 0 & 1 & 0 & 0 & 1 & 0 \end{pmatrix}$$

Graph $G = G(N, E)$

nodes N (vertices V)
edges E



structural information: **adjacency matrix A**

$A_{ij} = 1$ if $(n_i, n_j) \in E$
(nodes n_i and n_j have a common edge)

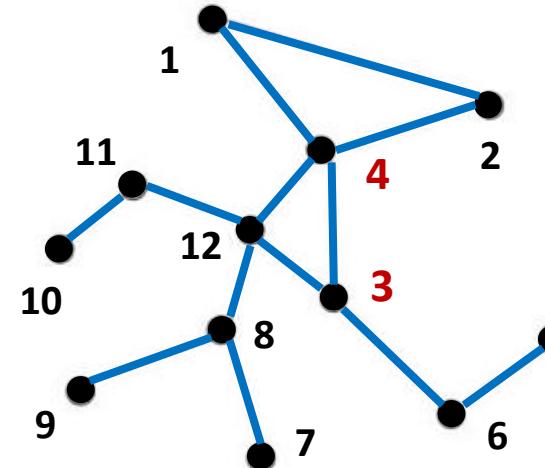
$A_{ij} = 0$ else

The order of counting the nodes is not relevant!
(permutation invariance)

$$A = \begin{pmatrix} 0 & 1 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 1 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 1 & 0 & 0 & 1 & 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 1 & 1 & 0 & 0 & 0 & 1 & 0 & 0 & 1 \end{pmatrix}$$

Graph $G = G(N, E)$

nodes N (vertices V)
edges E



structural information: **adjacency matrix A**

$A_{ij} = 1$ if $(n_i, n_j) \in E$
(nodes n_i and n_j have a common edge)

$A_{ij} = 0$ else

The order of counting the nodes is not relevant!
(permutation invariance)

Each graph can be represented by $N!$
adjacency matrices!

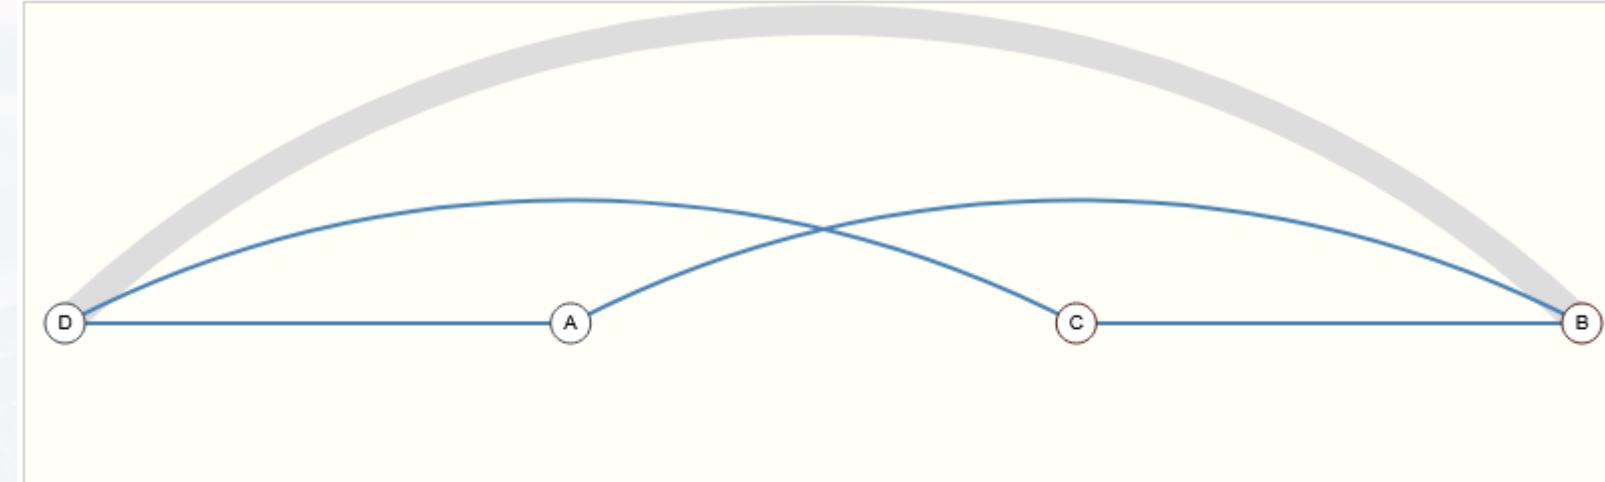
$$A = \begin{pmatrix} 0 & 1 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & \textcolor{red}{1} & 0 & 0 & \textcolor{red}{1} & 0 & 0 & 0 & 0 & \textcolor{red}{1} \\ \textcolor{red}{1} & \textcolor{red}{1} & \textcolor{red}{1} & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 & 1 & 0 & 0 & 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 \end{pmatrix}$$

Graph $G = G(N, E)$

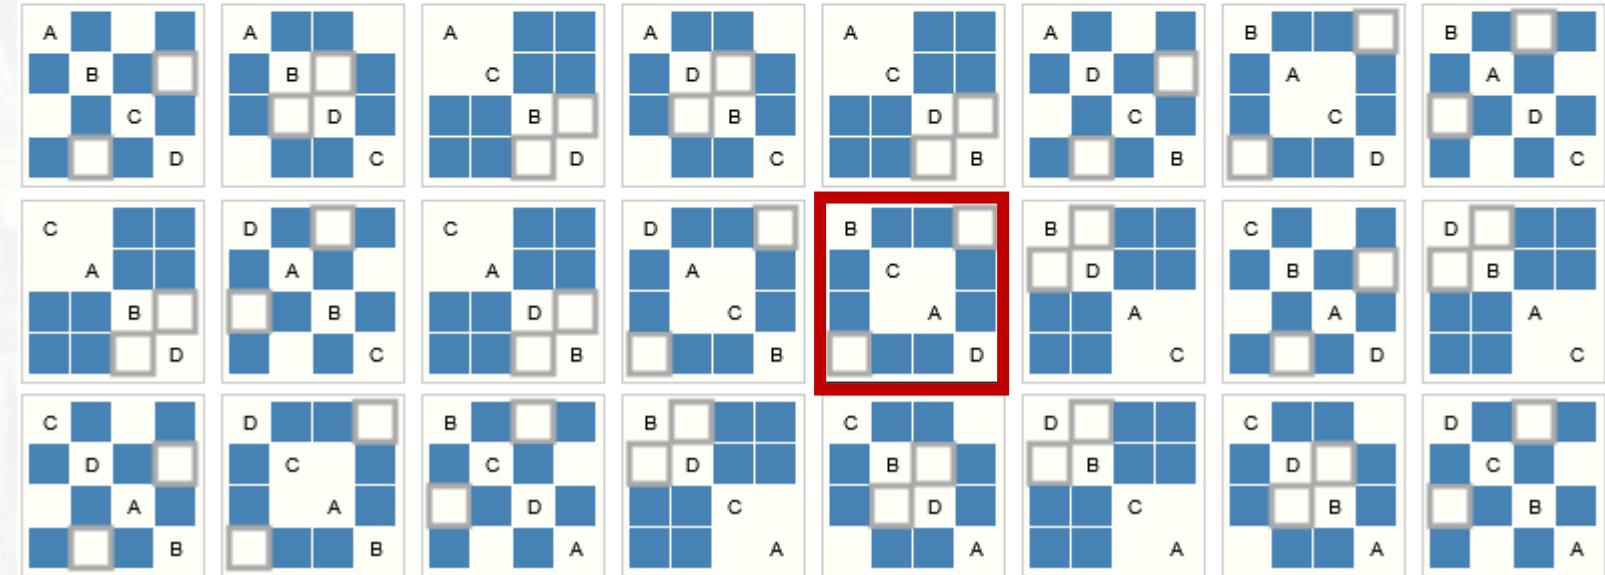
nodes N (vertices V)
edges E



Each graph can be represented by $N!$ adjacency matrices!



animation here

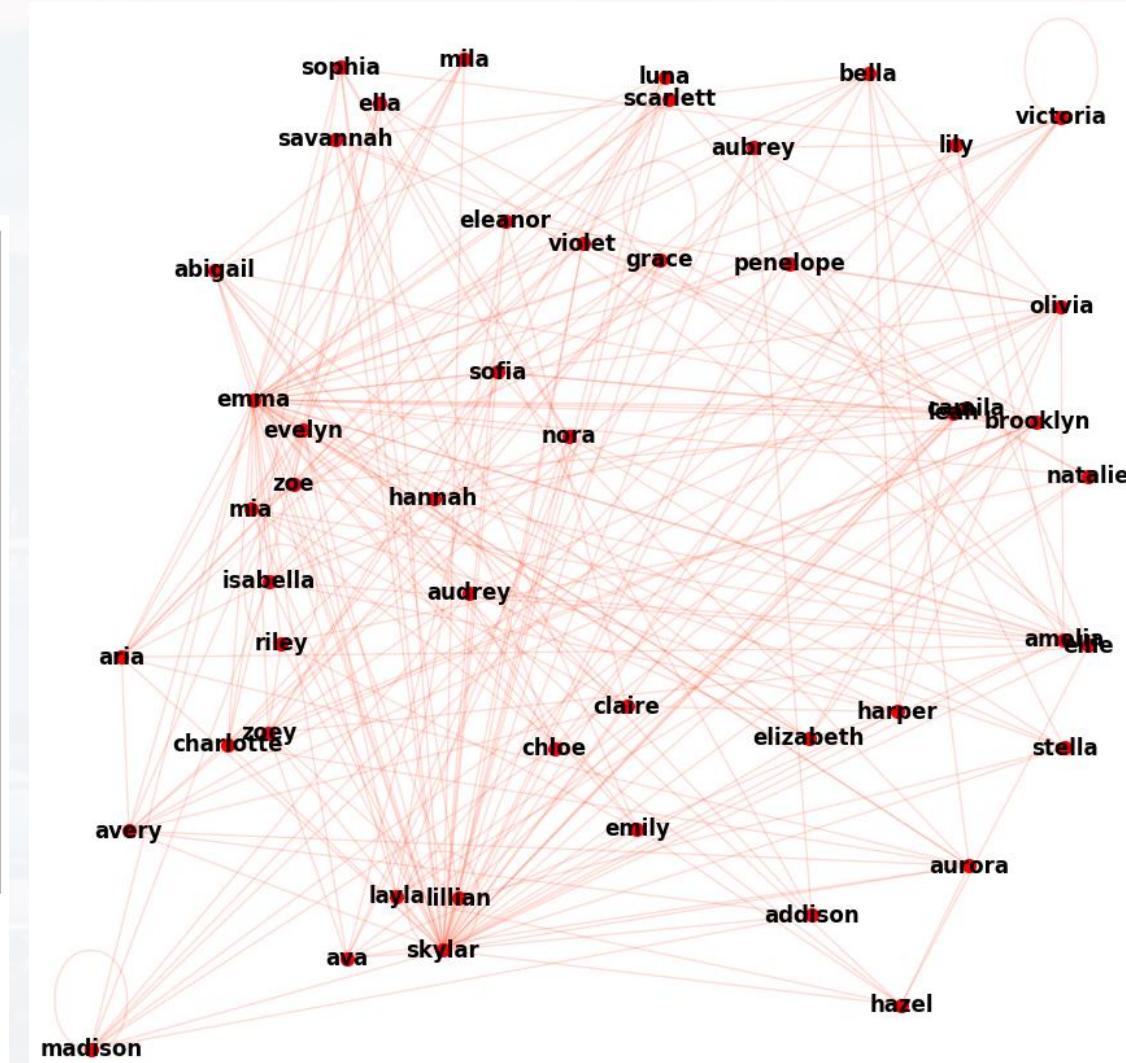
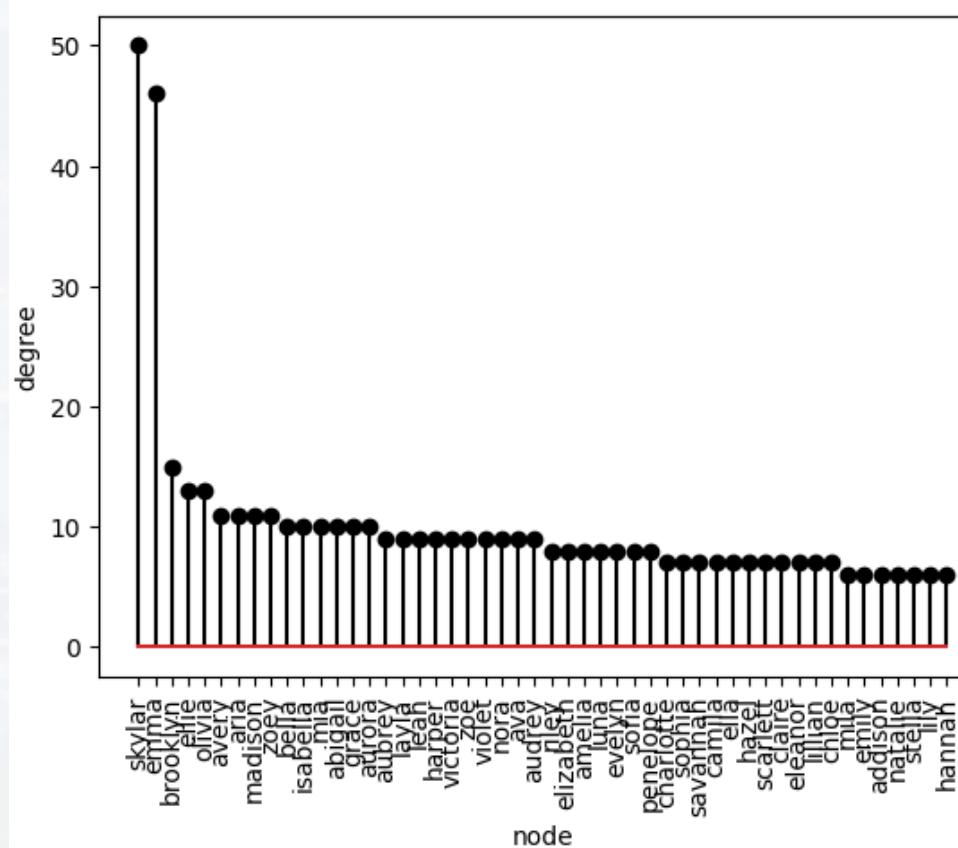




visualizing a graph:

```
import networkx as nx #pip install networkx
```

see: Graph_I.ipynb

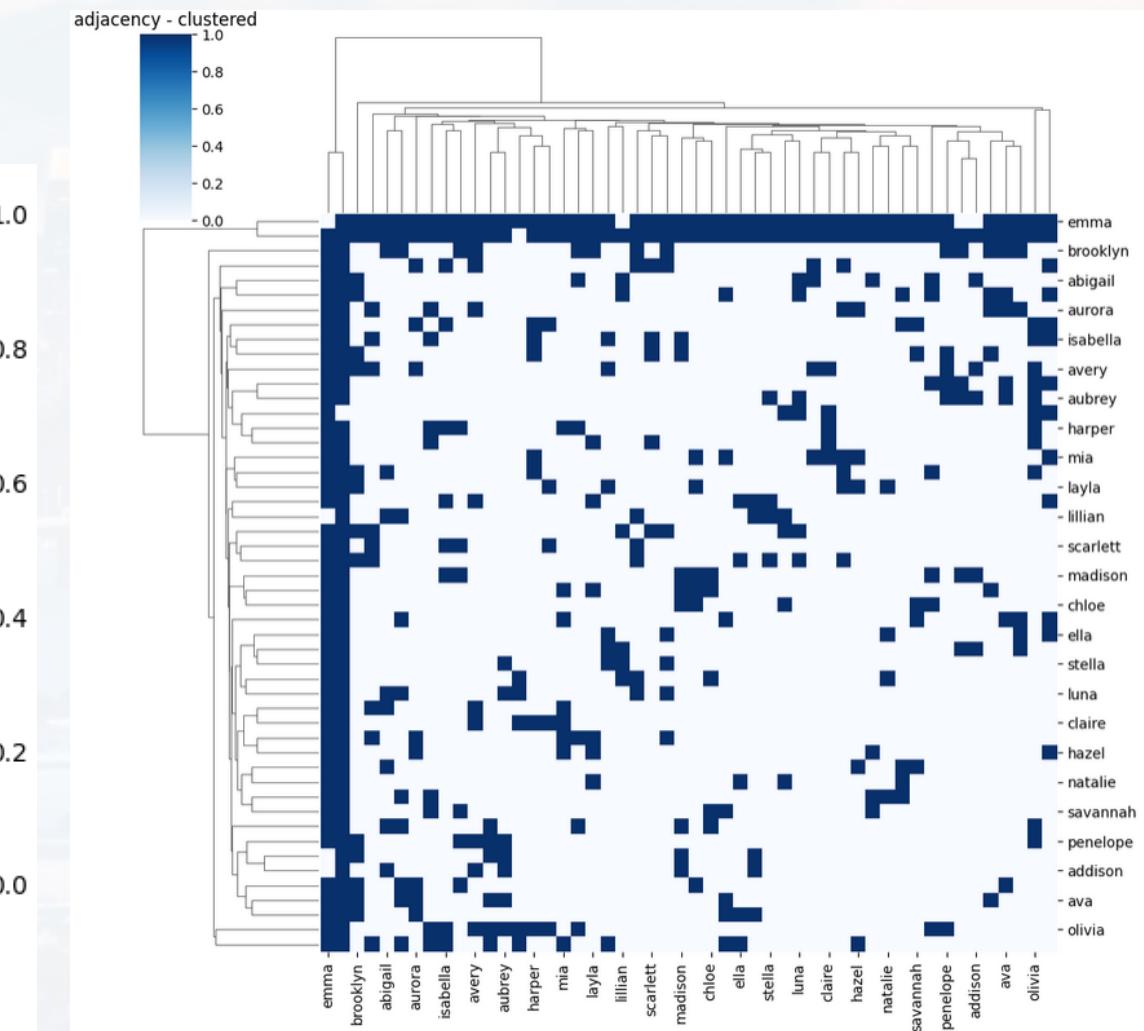
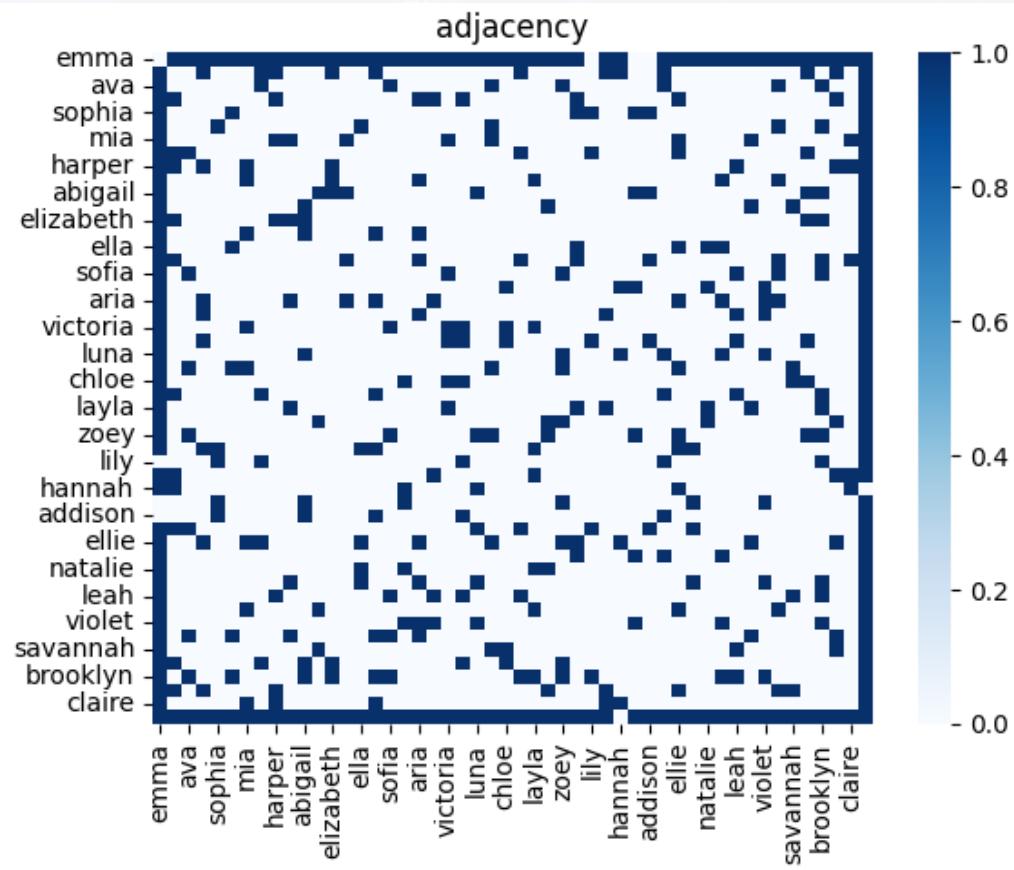




visualizing a graph:

```
import networkx as nx #pip install networkx
```

see: [Graph_I.ipynb](#)



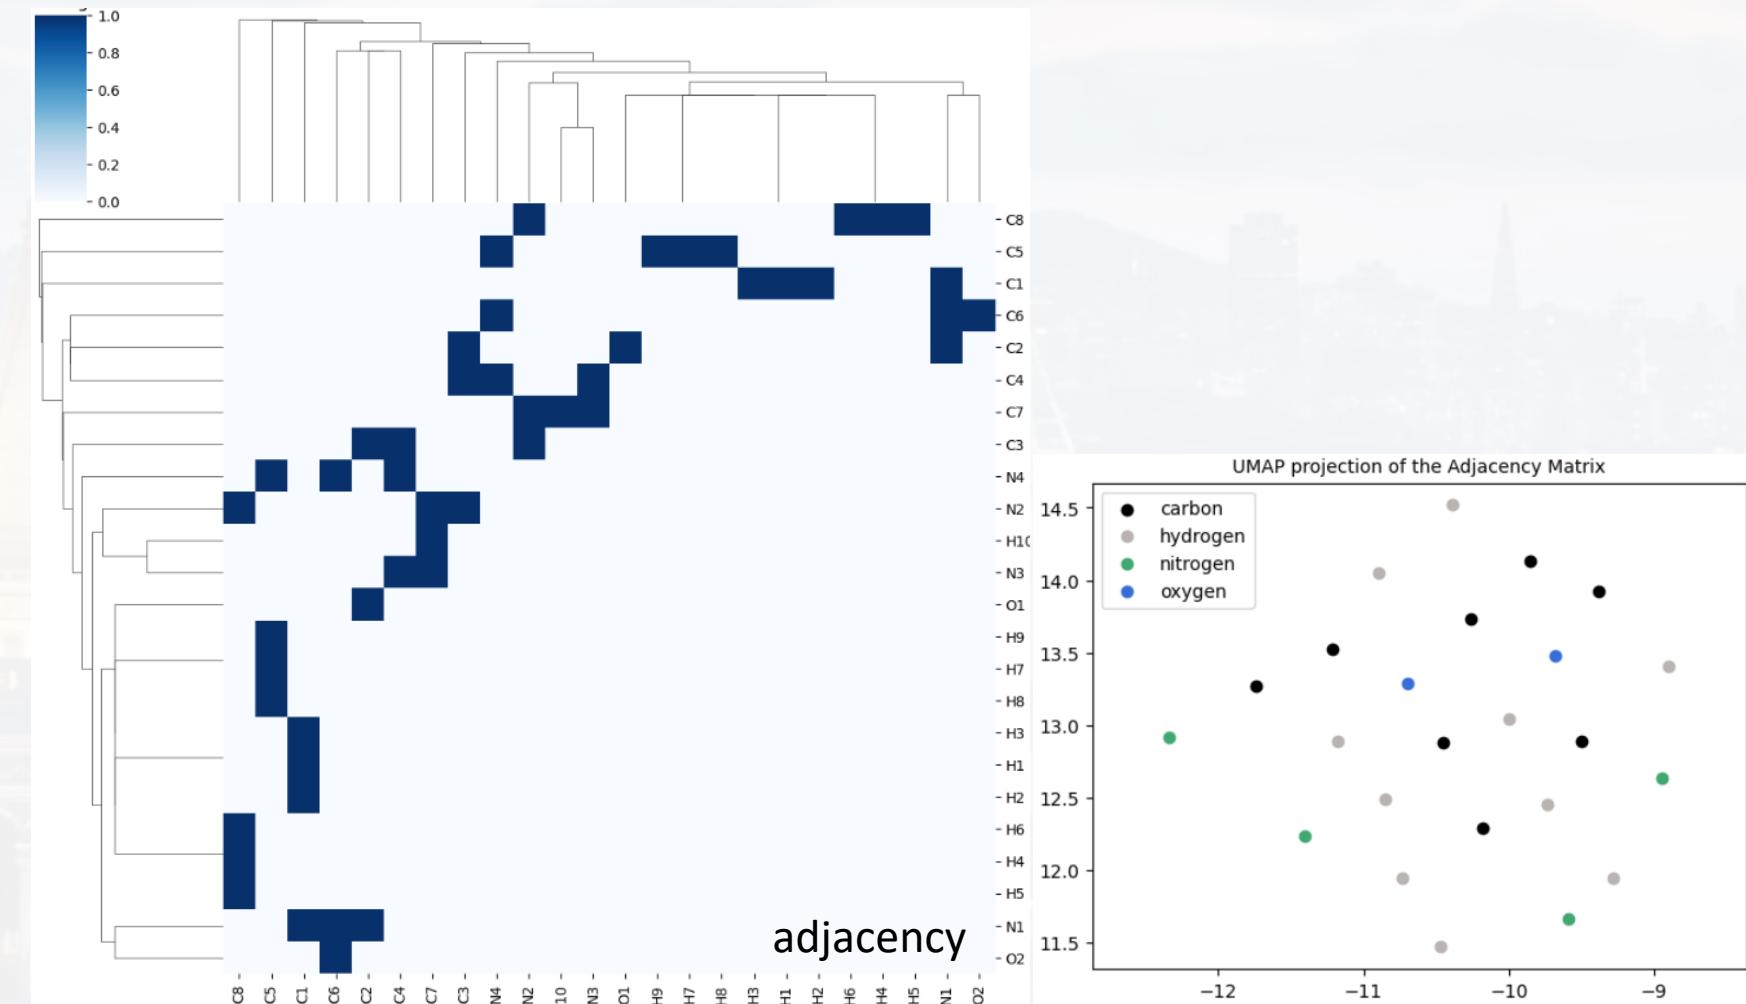


building and visualizing a **weighted** graph:

```
import networkx as nx #pip install networkx
```

see: [Graph_II.ipynb](#)

Caffein molecule





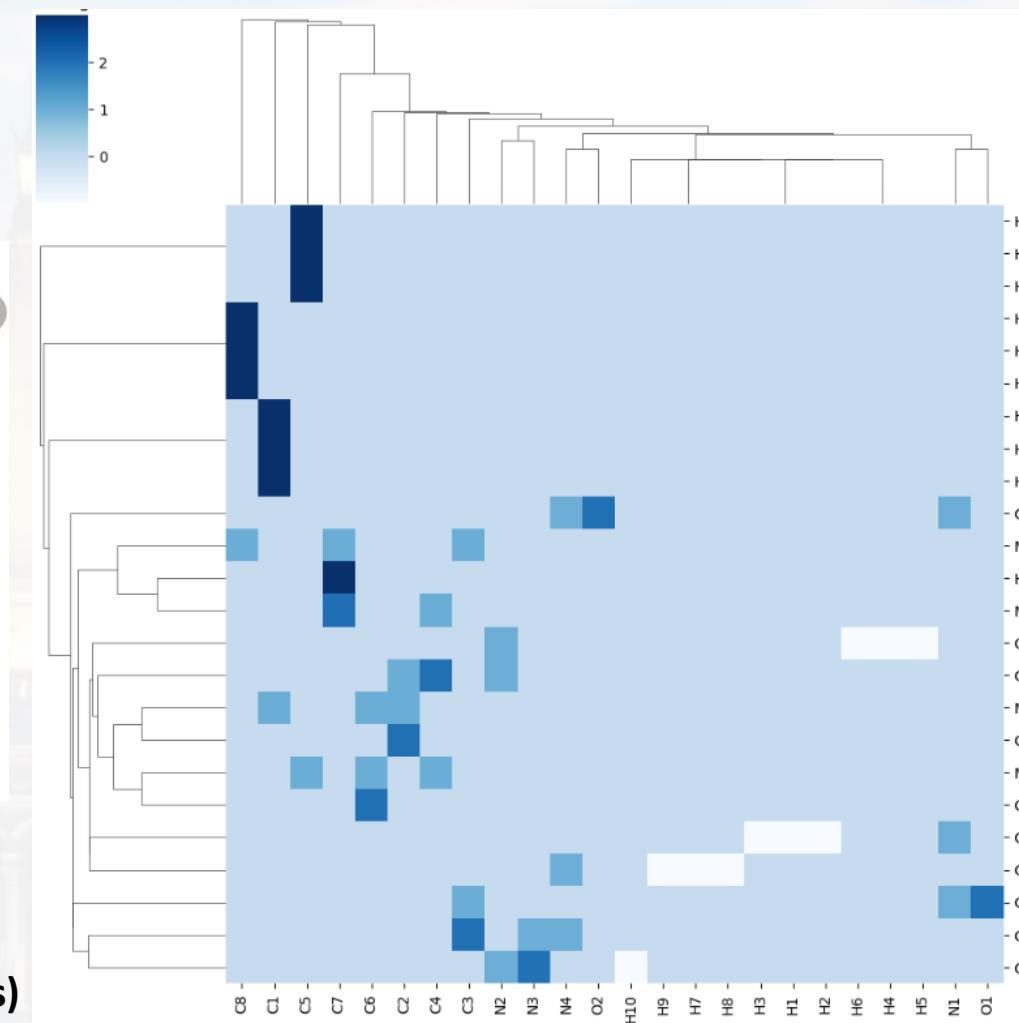
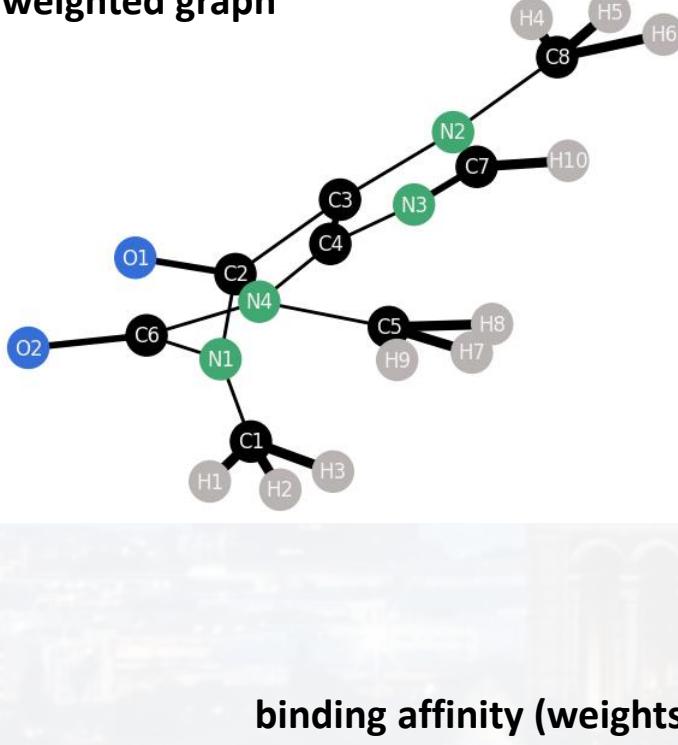
building and visualizing a **weighted** graph:

```
import networkx as nx #pip install networkx
```

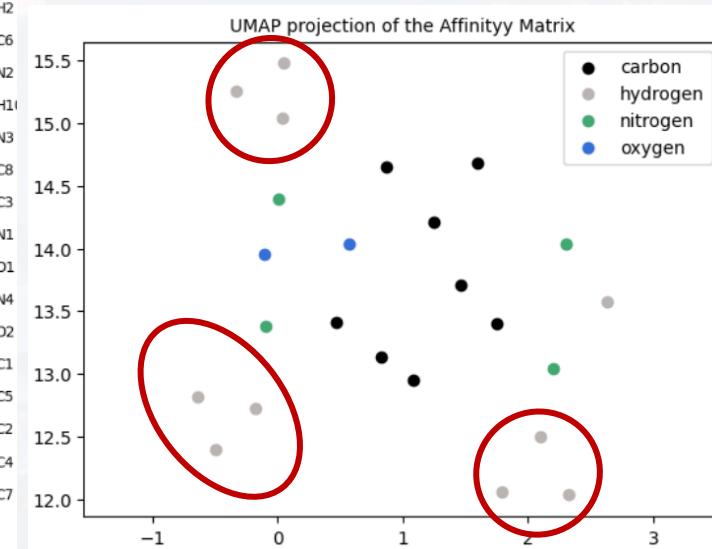
see: [Graph_II.ipynb](#)

Caffeine molecule

weighted graph



hydrogen atoms are at the edges of the molecule!





more about graphs:

$$d(n_i) = \sum_j A_{ij}$$

degree of node n_i

$$\mathcal{N}(n_i) = \{n_j \in N : (n_i, n_j) \in E\}$$

neighborhood $\mathcal{N}(n_i)$ of node n_i
for first degree neighborhood $|\mathcal{N}(n_i)| = d(n_i)$

$$S_{com} = |\mathcal{N}(n_i) \cap \mathcal{N}(n_j)|$$

number for neighbors nodes n_i and n_j have in common.

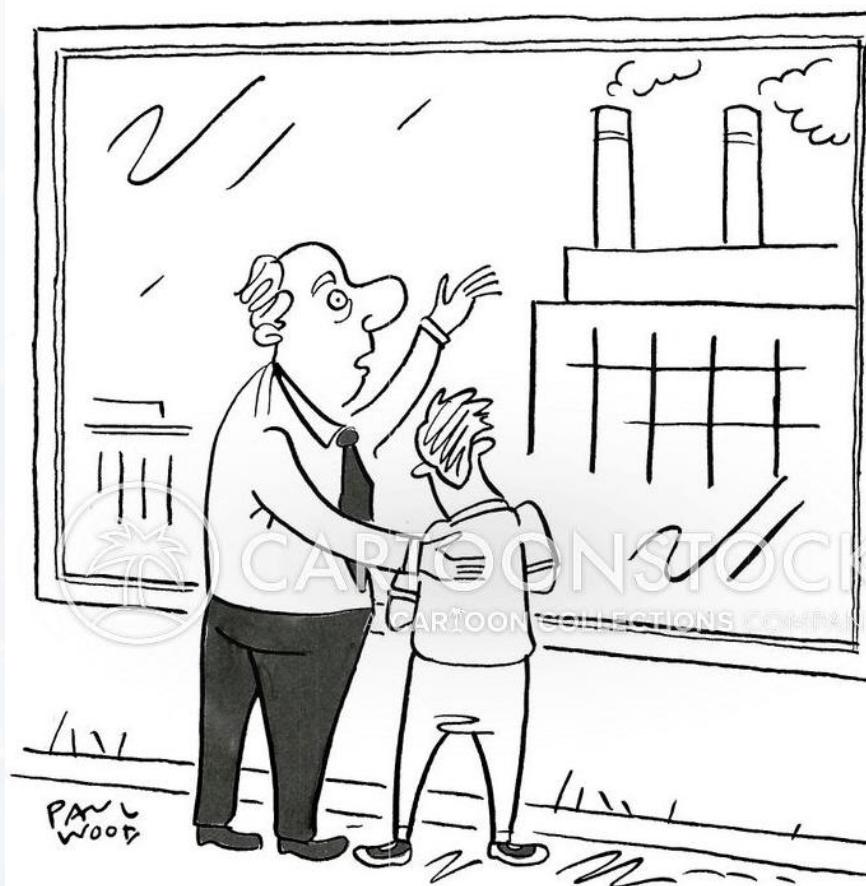
idea: nodes with many common neighbors are more likely to be similar or have a potential connection.

$$S_{rat} = \frac{|\mathcal{N}(n_i) \cap \mathcal{N}(n_j)|}{|\mathcal{N}(n_i) \cup \mathcal{N}(n_j)|}$$

ratio for neighbors nodes n_i and n_j have in common.

note:

There are more quantities (“importance”, “centrality” etc.), but they are all a function of A_{ij} .



ONE DAY SON, ALL THIS
WILL BE RUN BY ROBOTS

Outline

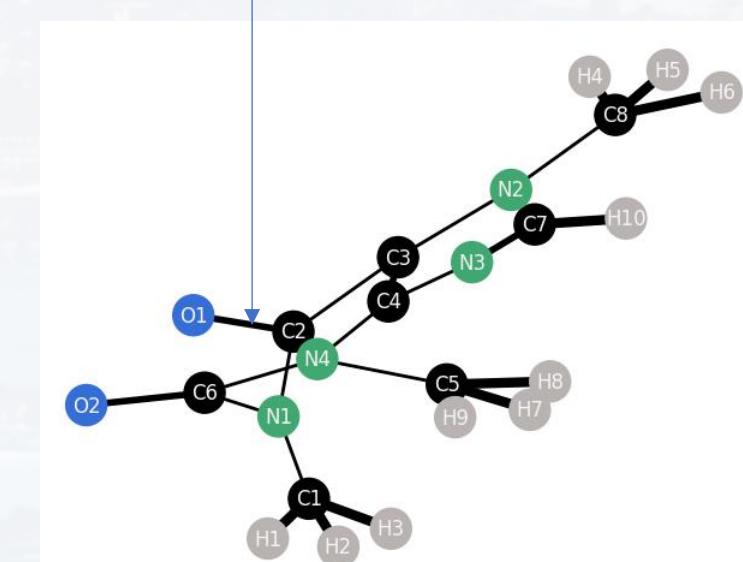
- What is a Graph
- The ANN Part
- PyTorch Example



What we can learn:

- node classification
- join nodes with similar properties to hyper nodes
- edge attributes, weights (weighted graph)
- edge prediction
- embedding (eg. 3D structure molecules)
- graph classification (is the molecule toxic y/n)
- graph regression (toxicity score)
- graph generation

weight: bond strength

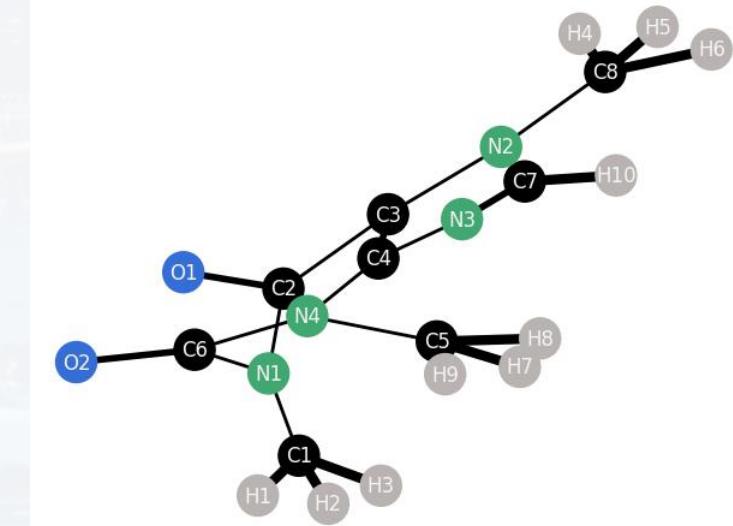




What we can learn:

- node classification
- join nodes with similar properties to hyper nodes
- edge attributes, weights (weighted graph)
- edge prediction
- embedding (eg. 3D structure molecules)
- graph classification (is the molecule toxic y/n)
- graph regression (toxicity score)
- graph generation

node (edge) level tasks





What we can learn:

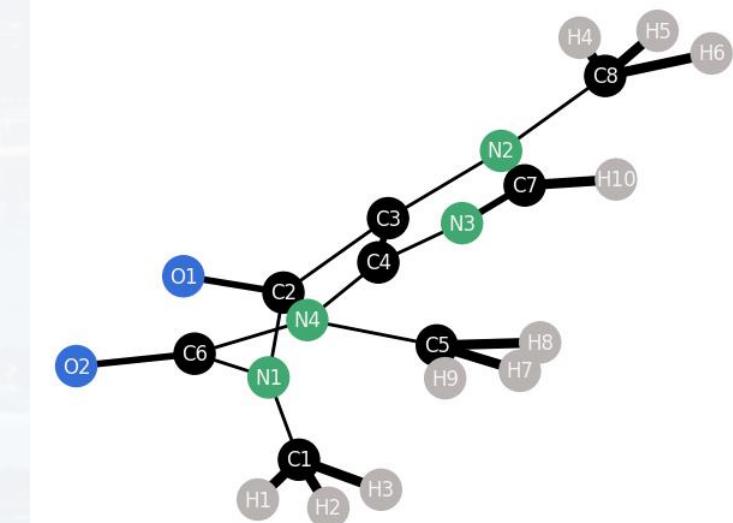
- node classification
- join nodes with similar properties to hyper nodes
- edge attributes, weights (weighted graph)
- edge prediction
- embedding (eg. 3D structure molecules)
- graph classification (is the molecule toxic y/n)
- graph regression (toxicity score)
- graph generation

information flow from one node to another:

message passing

different ways how:

- local averaging
- graph convolution (aka neighborhood aggregation)
- graph attention





What we can learn:

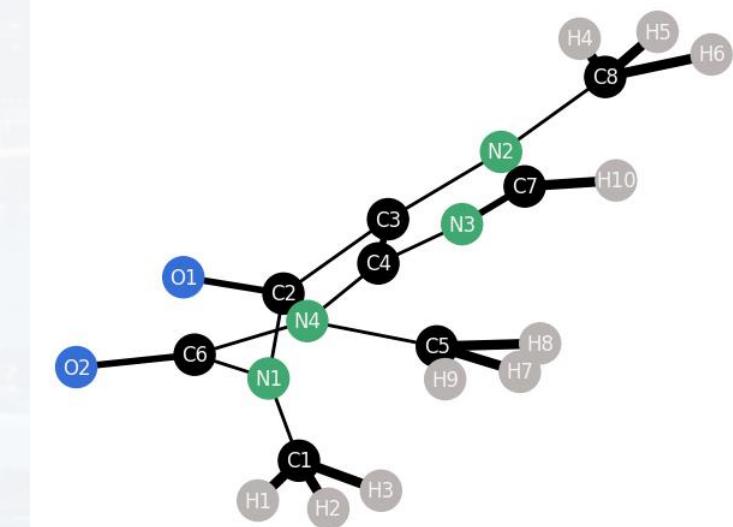
- node classification
- join nodes with similar properties to hyper nodes
- edge attributes, weights (weighted graph)
- edge prediction
- embedding (eg. 3D structure molecules)
- graph classification (is the molecule toxic y/n)
- graph regression (toxicity score)
- graph generation

information flow from one node to another:

message passing

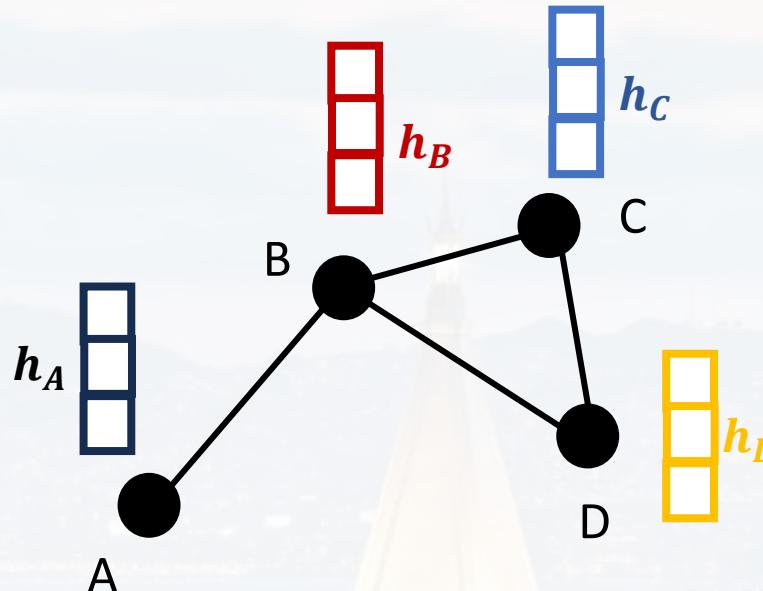
different ways how:

- local averaging
- **graph convolution** (aka neighborhood aggregation)
- **graph attention**



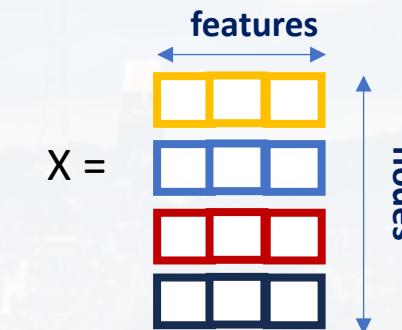


Graph Convolution

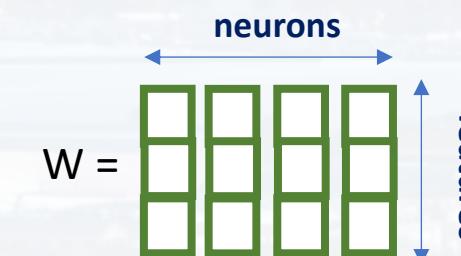


each node i has a **feature vector** h_i

matrix X of shape (number of nodes, number of node features)



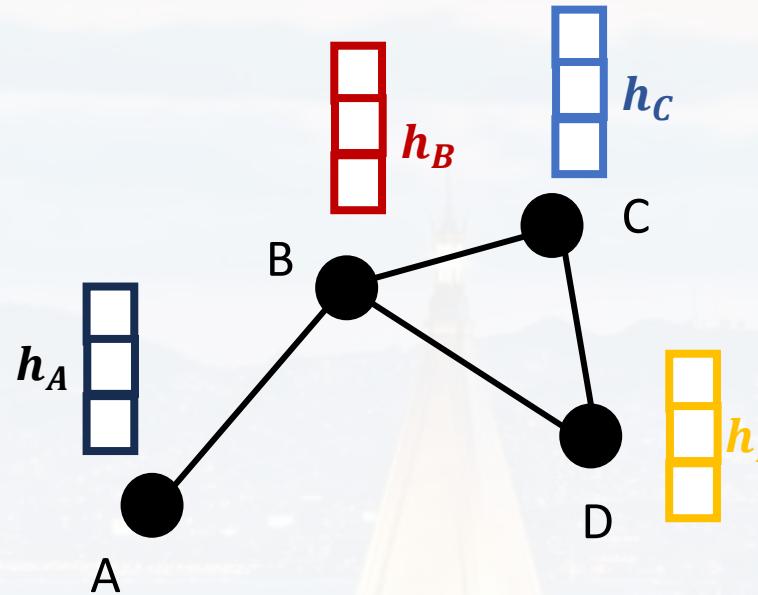
weight matrix W of shape
(number of node features, number of neurons)



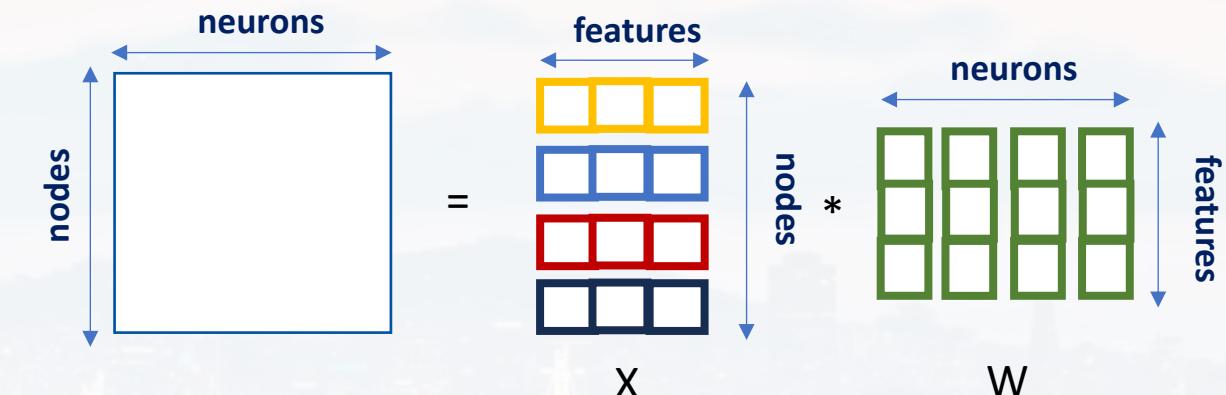
note: only **one** W for the entire graph
W is a **learnable**



Graph Convolution

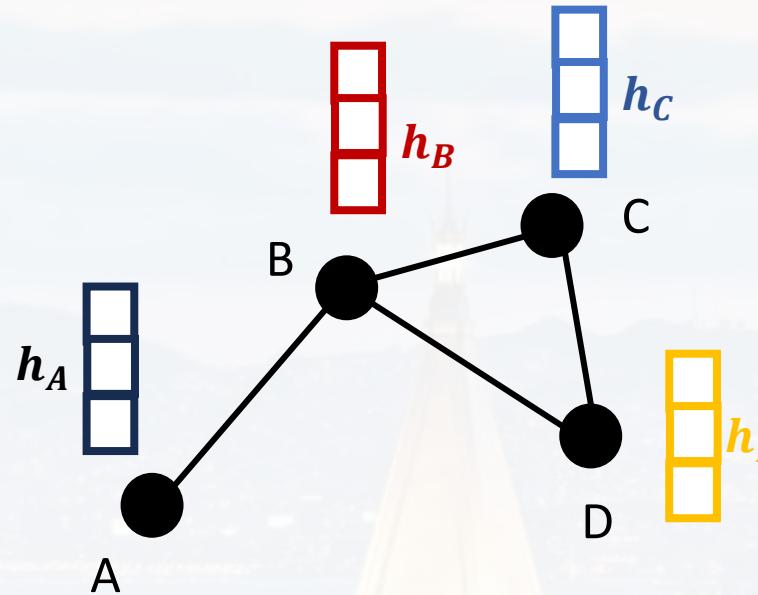


each node i has a **feature vector** h_i

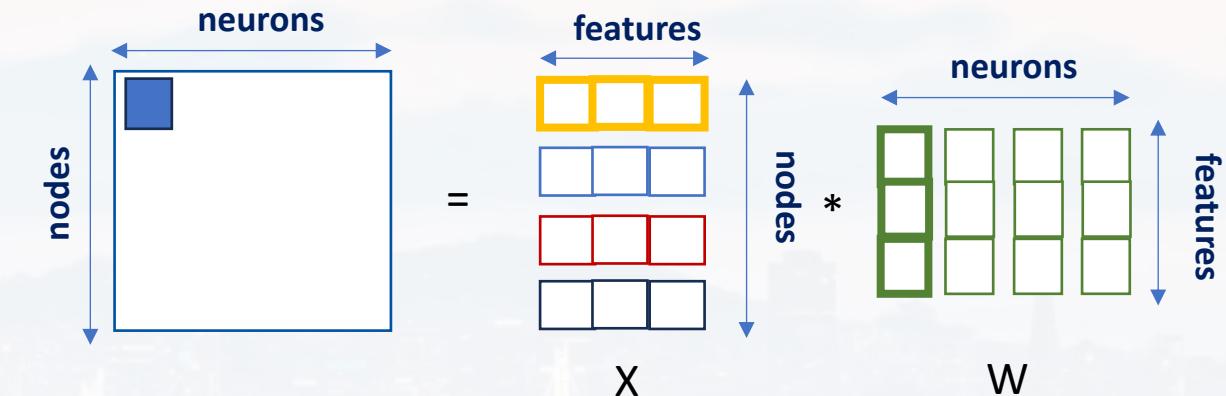




Graph Convolution

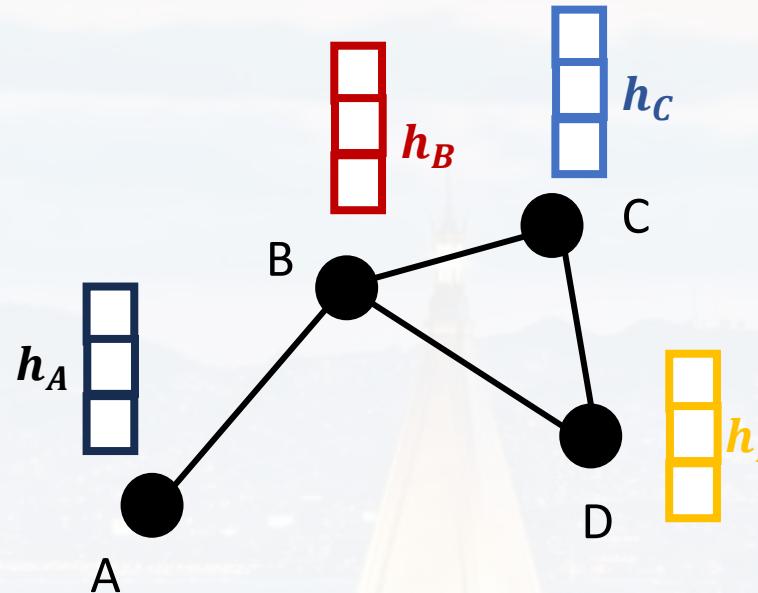


each node i has a **feature vector** h_i

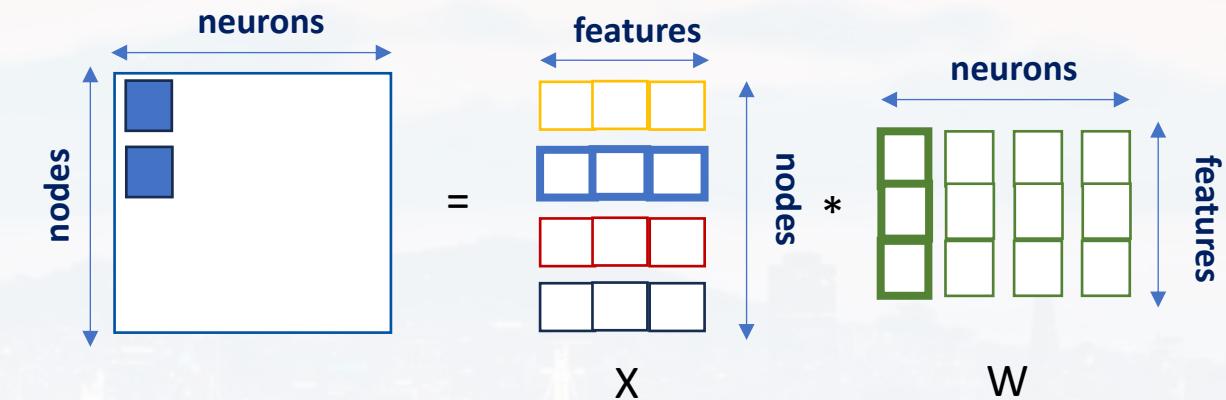




Graph Convolution

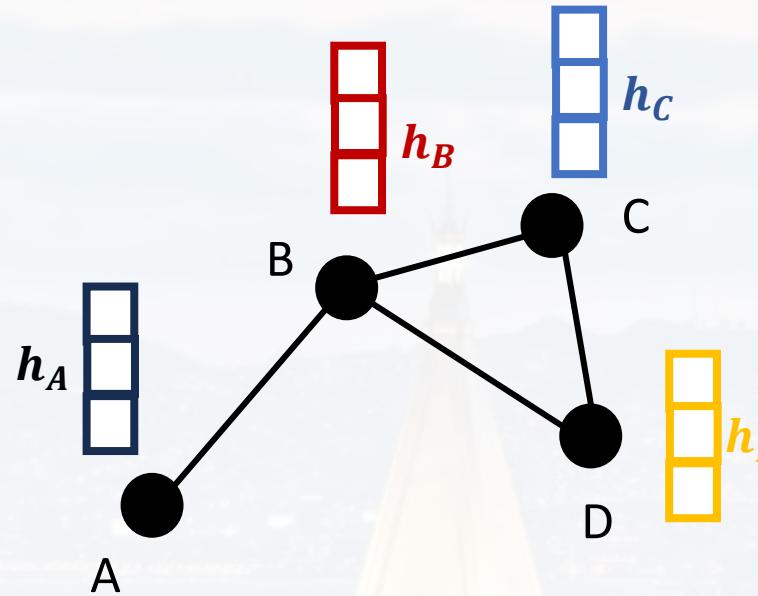


each node i has a **feature vector** h_i

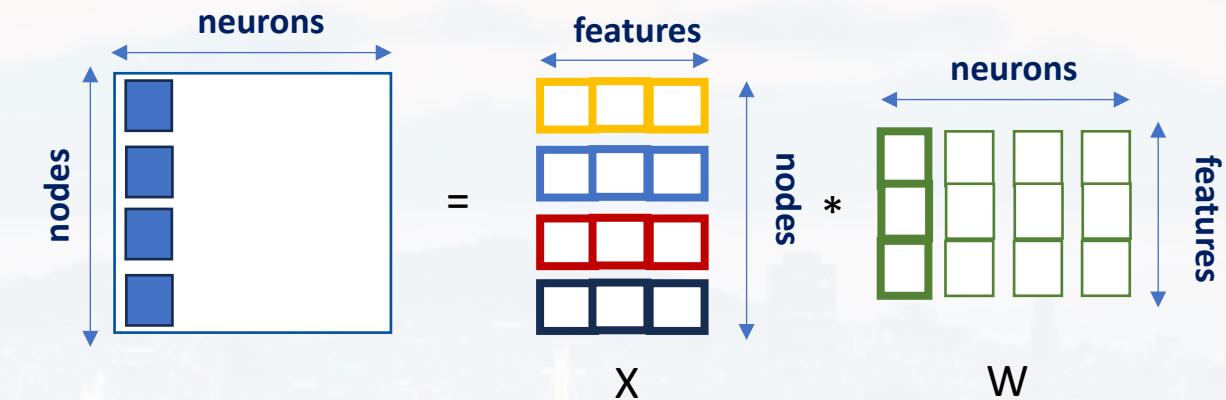




Graph Convolution

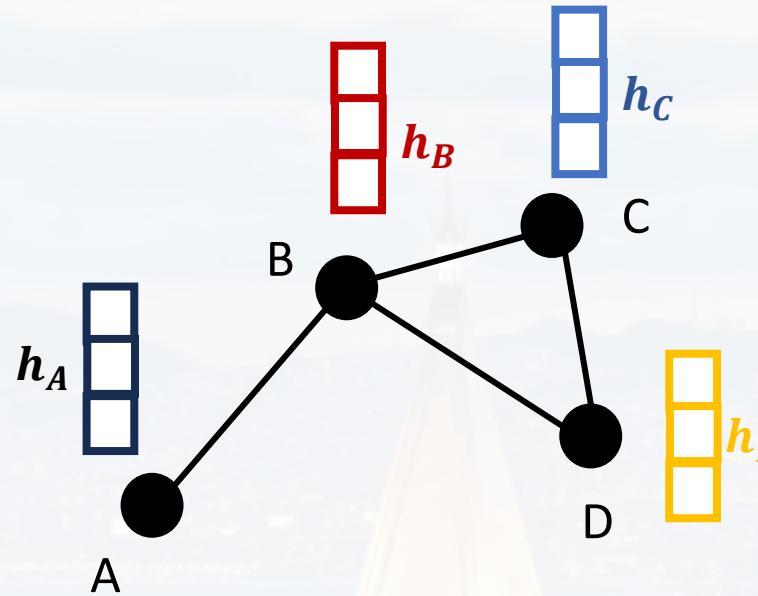


each node i has a **feature vector** h_i

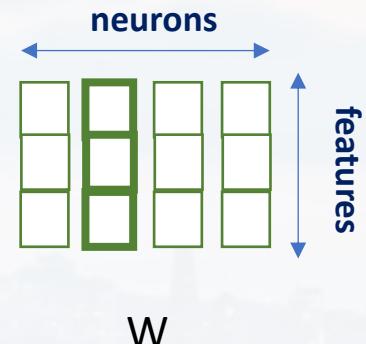
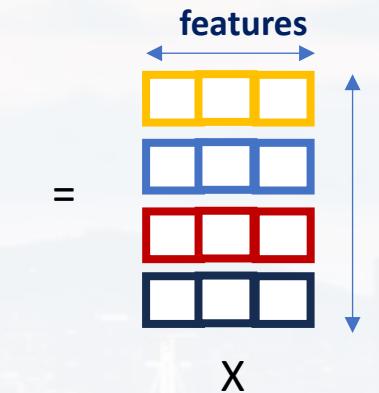
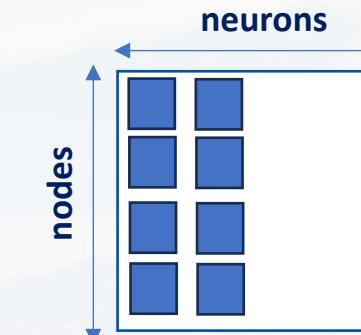




Graph Convolution

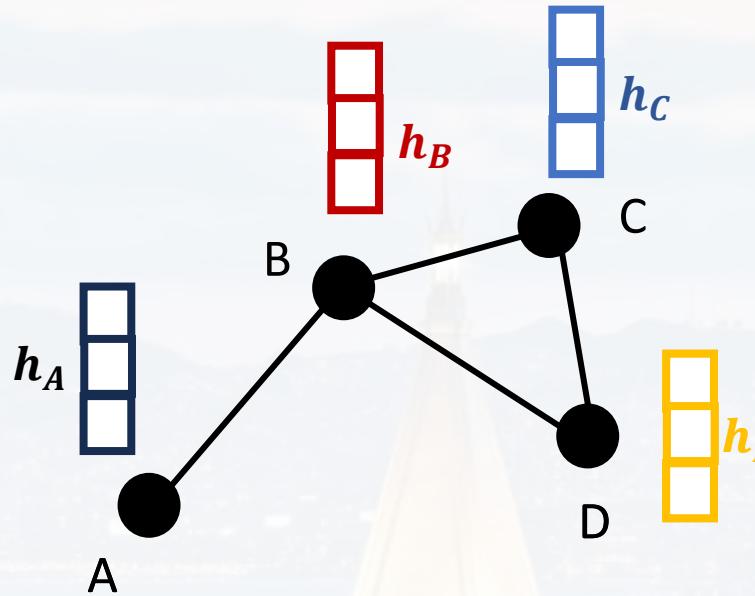


each node i has a **feature vector** h_i

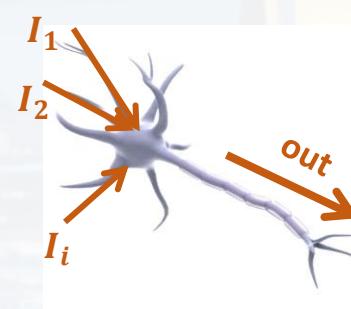
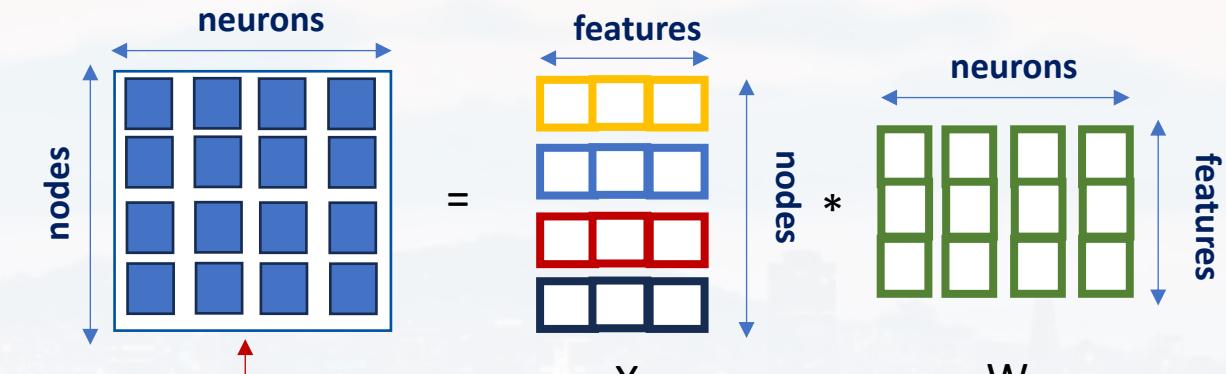




Graph Convolution



each node i has a **feature vector** h_i

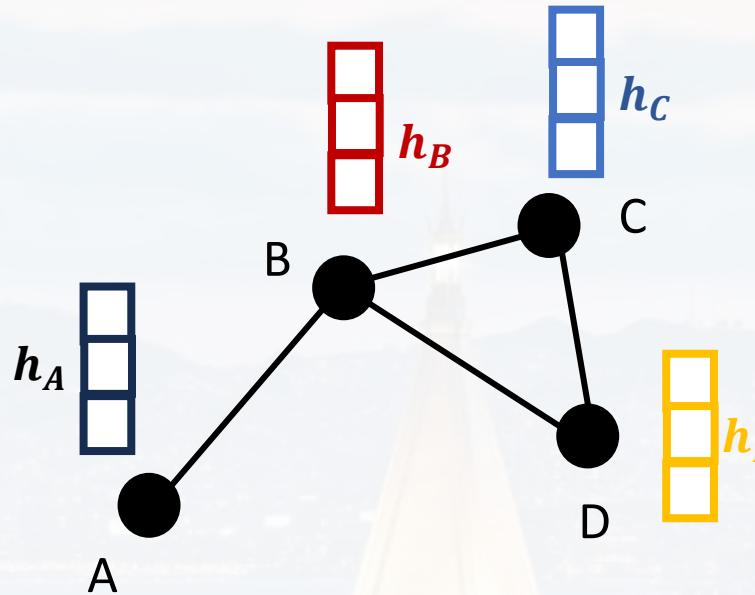


$$net = \sum_i I_i \cdot w_i + b$$

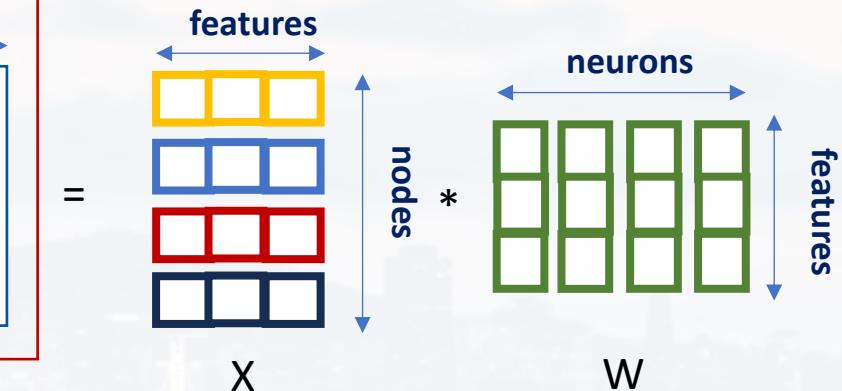
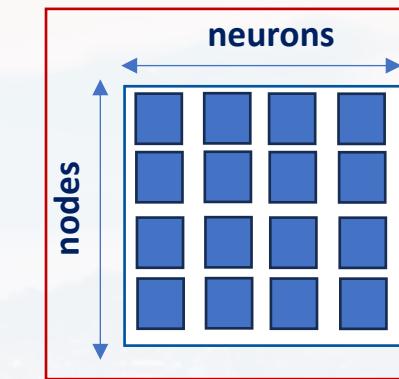
$$m_{jk} = \sum_i w_{ji} x_{ik}$$



Graph Convolution

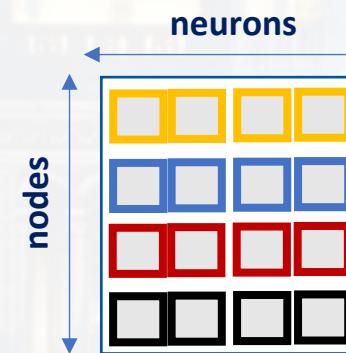


each node i has a **feature vector** h_i



$$m_{jk} = \sum_i w_{ji} x_{ik}$$

depending on W
the output features
may have different
lengths than the
input features

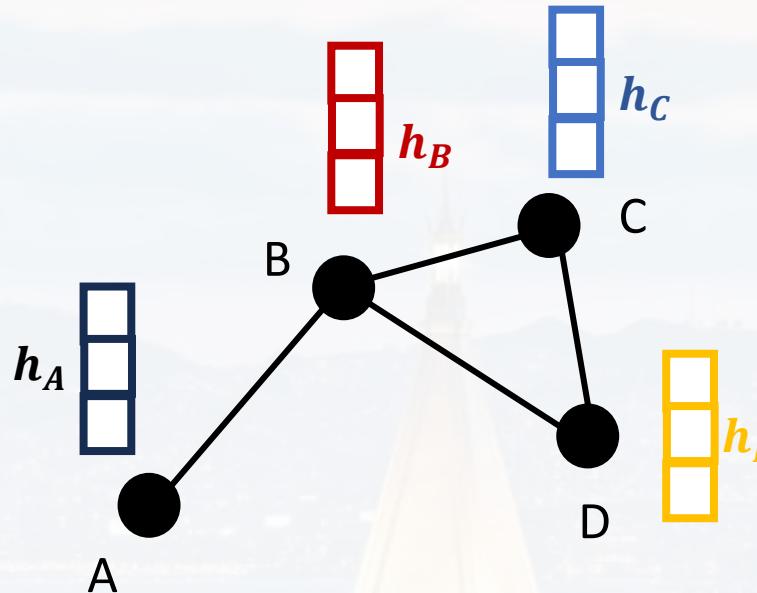


adjacency A

$$= \left[\begin{pmatrix} 0 & 1 & 0 & 0 \\ 1 & 0 & 1 & 1 \\ 0 & 1 & 0 & 1 \\ 0 & 1 & 1 & 0 \end{pmatrix} + \begin{pmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{pmatrix} \right] * M$$

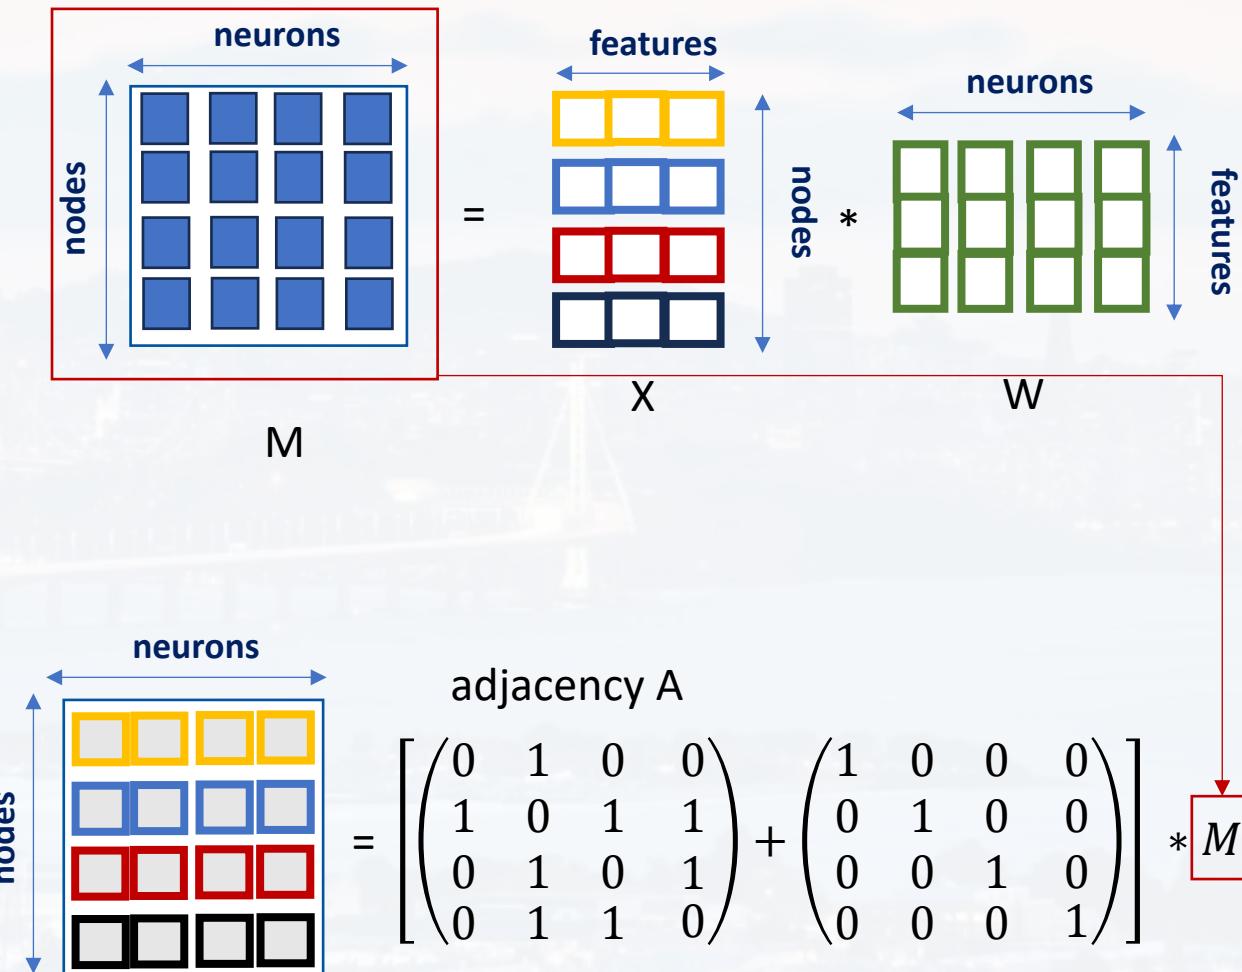


Graph Convolution



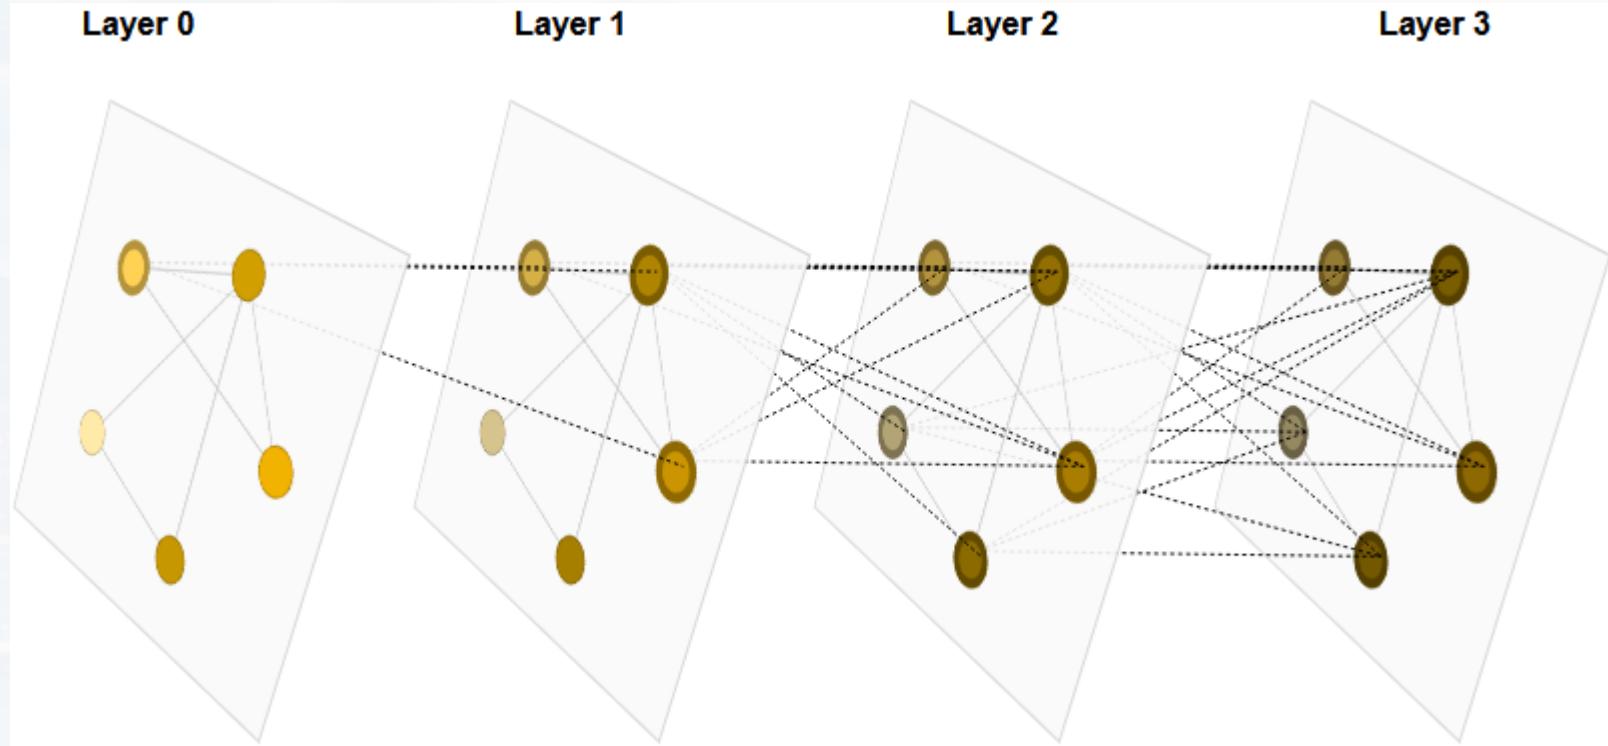
pass through a ReLU and/or
repeat the procedure with another W
(aka second convolution layer)

each node i has a **feature vector** h_i





Graph Convolution

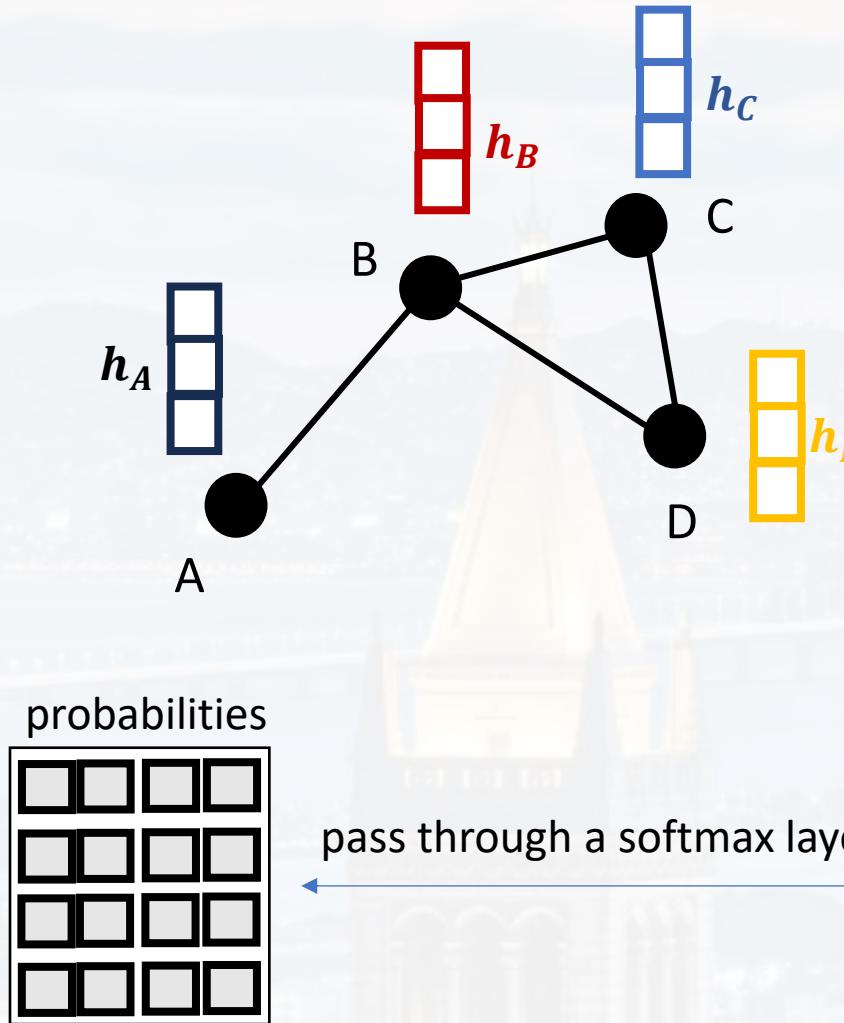


1st layer: one-hop neighborhood
2nd layer: two-hop neighborhood
etc

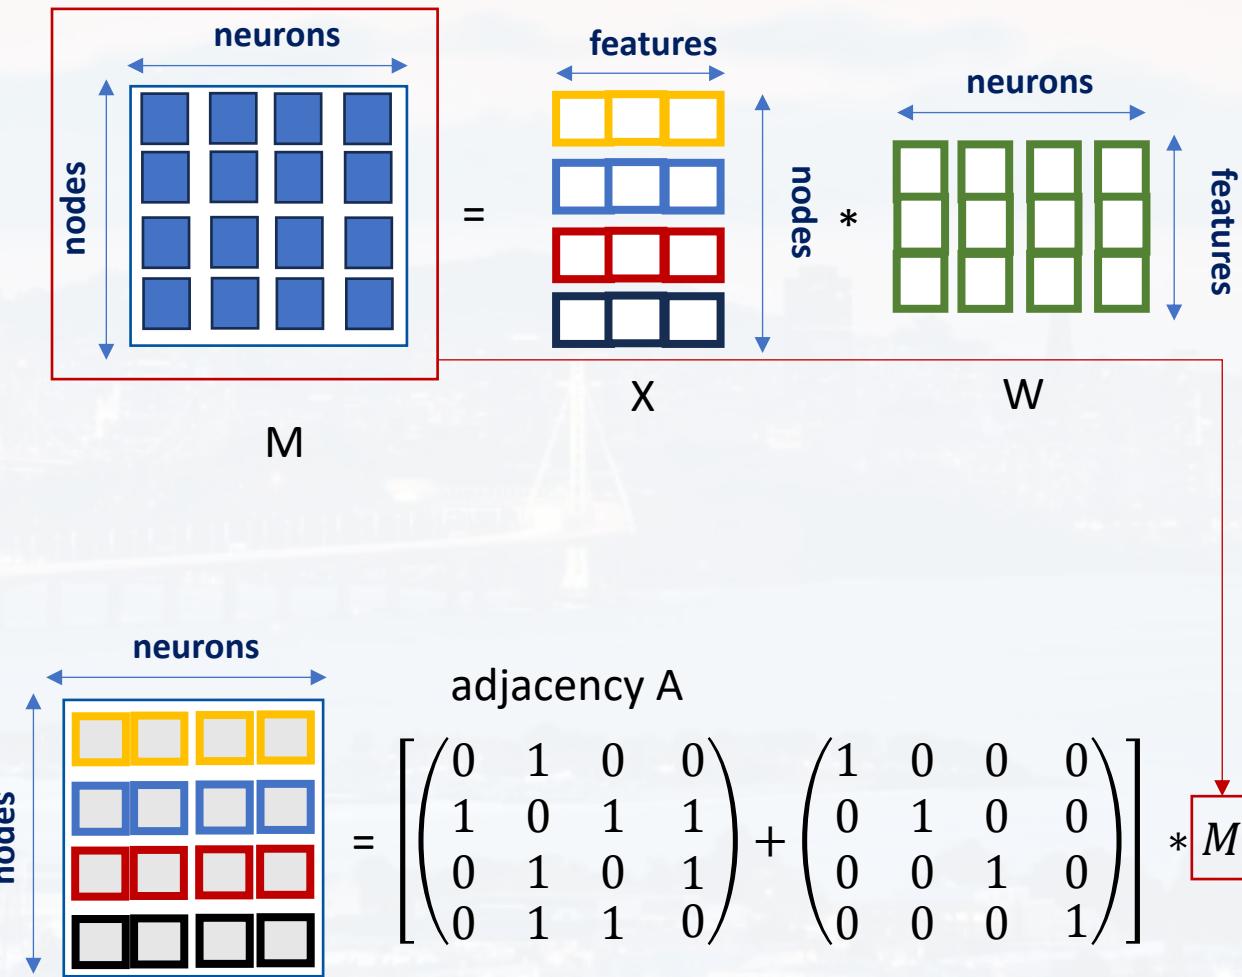
[animation here](#)



Graph Convolution

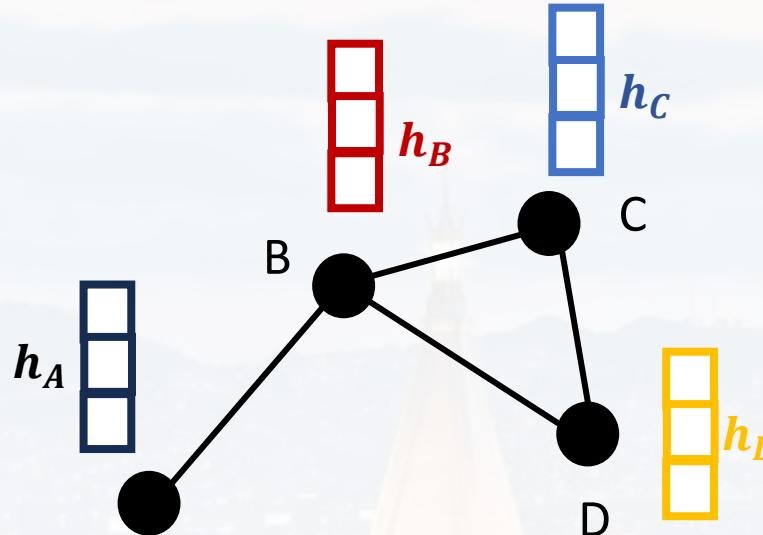


each node i has a **feature vector** h_i





Graph Convolution



summary

- A: adjacency matrix (number of nodes x number of nodes)
- I: identity matrix (number of nodes x number of nodes)
- X: node feature matrix (number of nodes x number of features)
- W: weight matrix (number of features x number of neurons)
- σ : any activation function
- $D^{-1/2}$: diagonal matrix for normalization

$$H(\text{embedding}) = \sigma [D^{-1/2} (A + I) D^{-1/2} X W]$$

However, this would give nodes with higher degree a larger weight

→ normalizing by $\frac{1}{\sqrt{d(n_i)}}$ and $\frac{1}{\sqrt{d(n_j)}}$

more information [here](#)



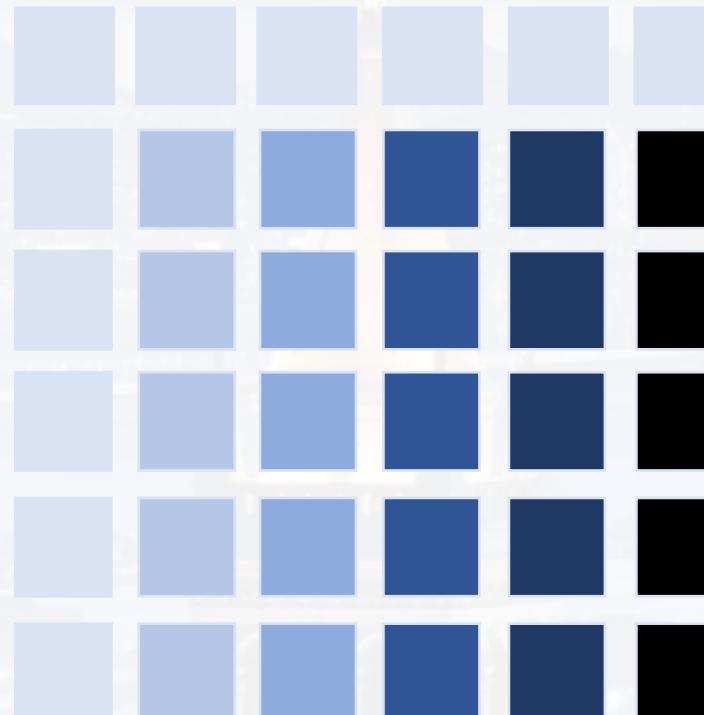
Matthew N. Bernstein



Graph Attention

see language models (lect 15)

“The cat jumped on the roof.”



how the first token influences all other token

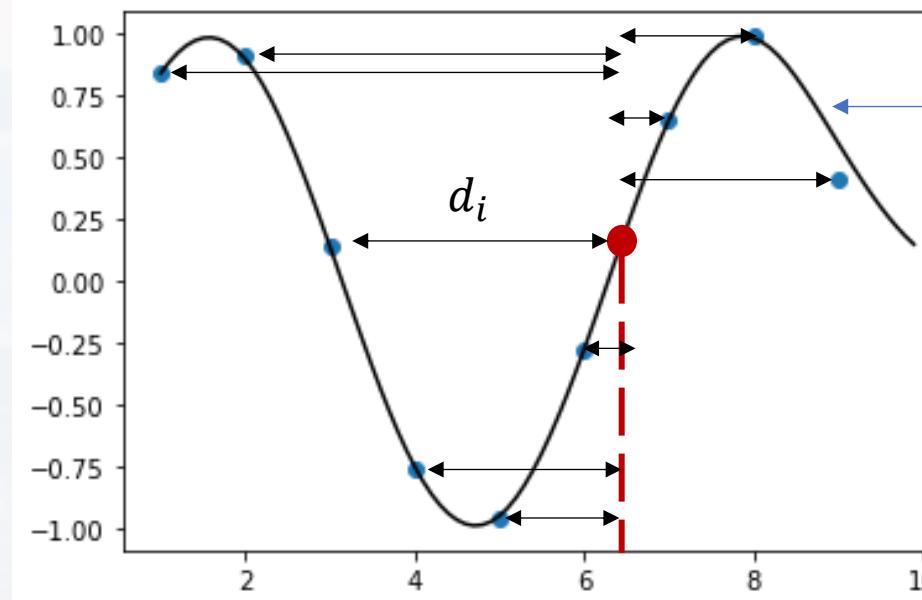


how the second token influences all other token

.... and so on



Attention



We want to interpolate between the blue dots
→ generating the black line
→ **no curve fitting!**

- idea:
- select a point for which we want the interpolation for
 - calculate distance d_i to every other point
 - each data point should influence the value of the interpolated point
 - the closer, the stronger the influence → weighted mean

$$y_{int} \sim \sum_{i=1}^I w_i y_i$$
$$w_i \sim \frac{1}{d_i}$$



Attention



```
D = np.tile(v, (1, len(v))) - np.tile(L.transpose(), (len(v), 1))
```

- each data point should influence the value of the interpolated point
- the closer, the stronger the influence → weighted mean

$$y_{int} \sim \sum_{i=1}^I w_i y_i$$

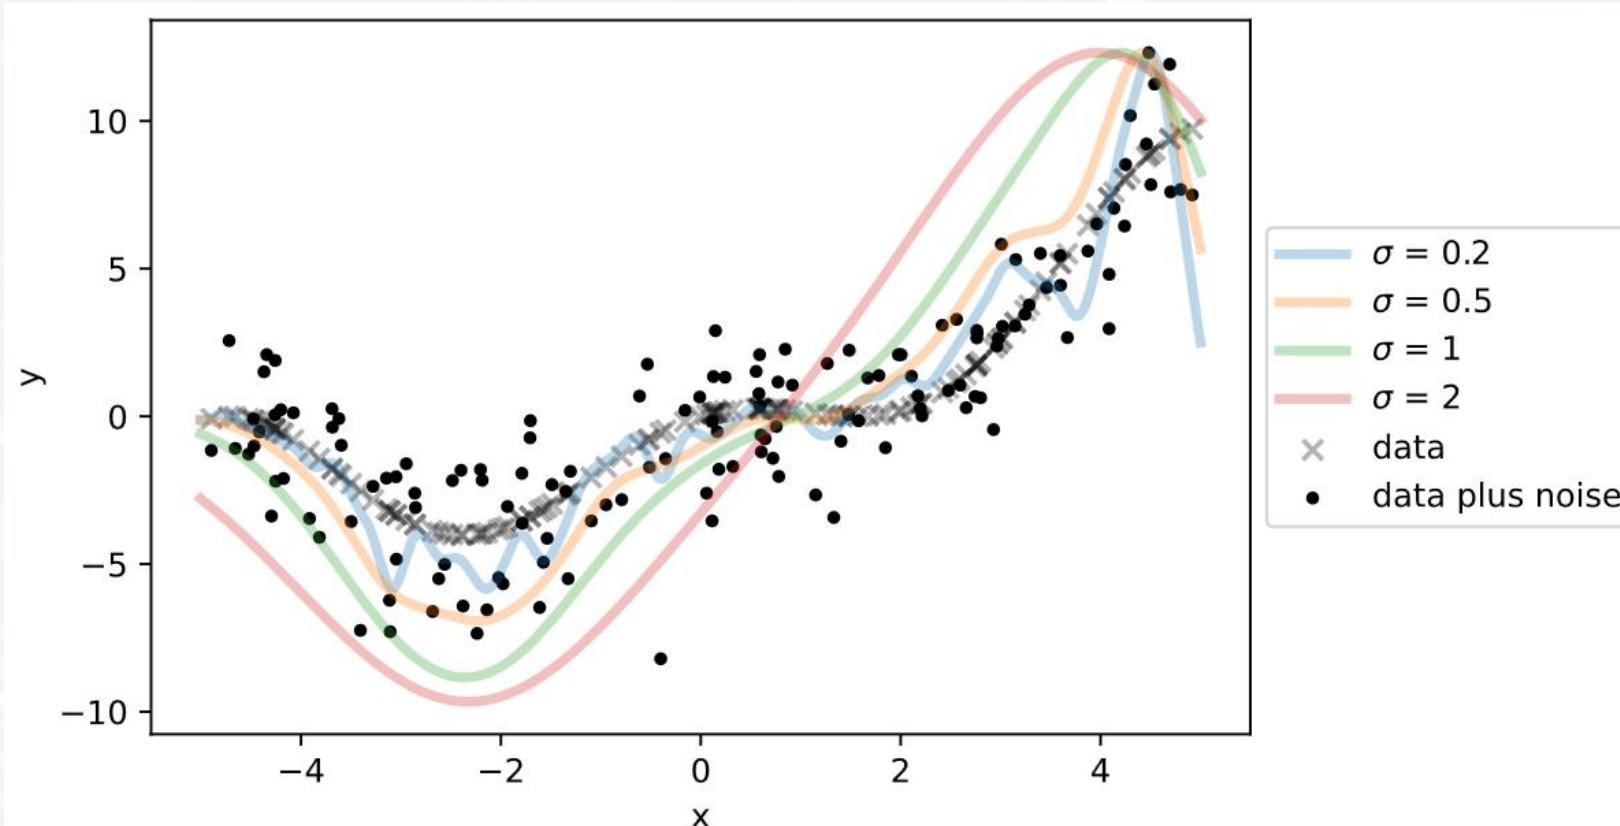
Gaussian kernel

```
w      = np.exp(-(D**2)/(sigma))
w      = w/np.sum(w + 1e-16, axis = 0)
yint  = np.dot(w.transpose(), y)
```

```
D = np.tile(V, (1, len(V))) - np.tile(L.transpose(), (len(V), 1))
```

Gaussian kernel

```
W      = np.exp(-(D**2)/(sigma))
W      = W/np.sum(W + 1e-16, axis = 0)
yint  = np.dot(W.transpose(), y)
```



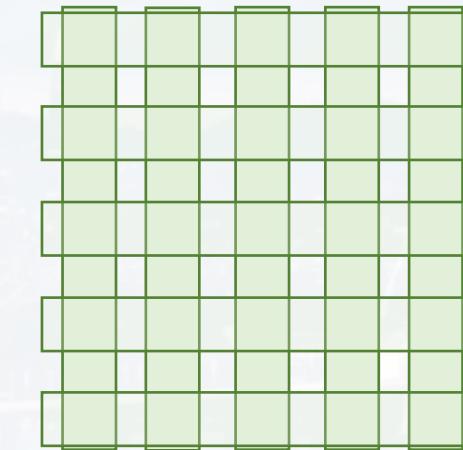
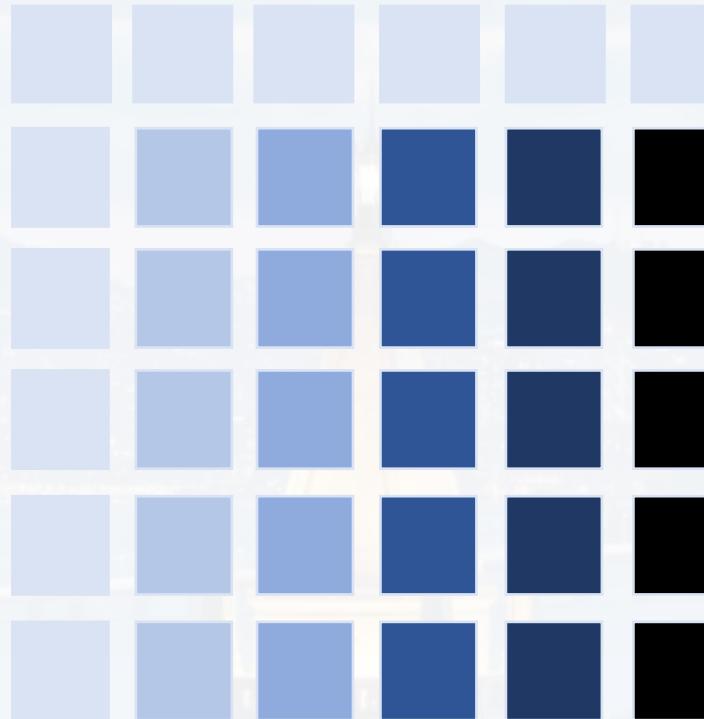
check out:

[SmoothGaussKernel.py](#)
[SmoothExamples.py](#)



Attention

"The cat jumped on the roof."



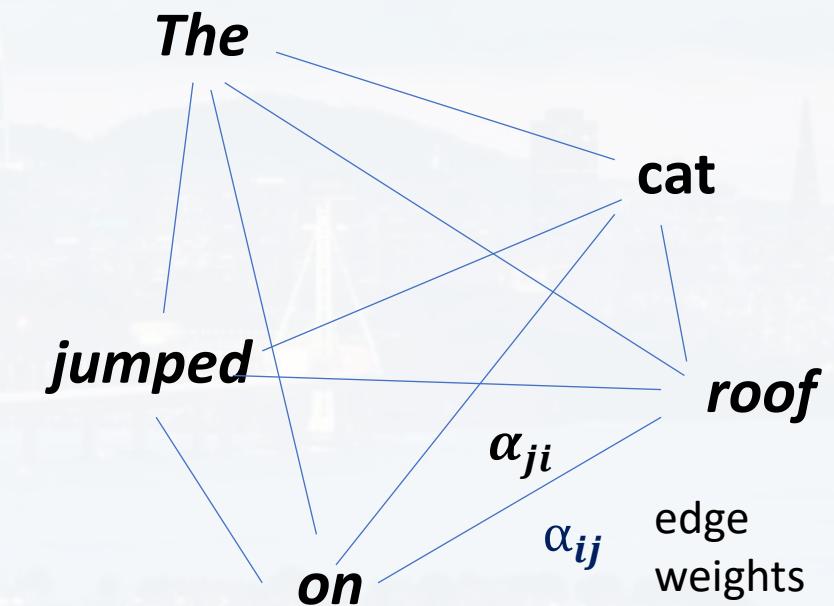
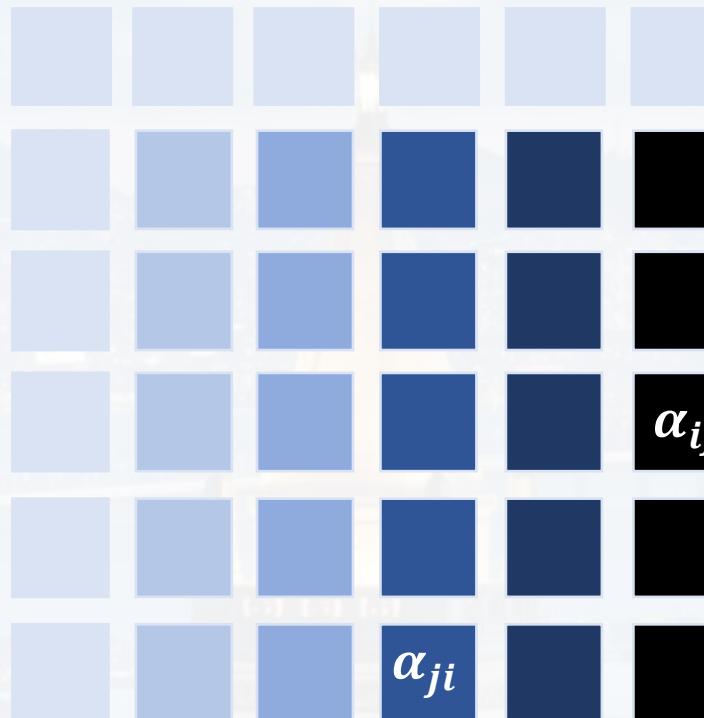
```
Gaussian kernel      W      = np.exp(-(D**2)/(sigma))  
                      W      = W/np.sum(W + 1e-16, axis = 0)  
                      yint   = np.dot(W.transpose(), y)
```

actual attention:
these weights are learnable,
no kernel assumed!



Graph Attention

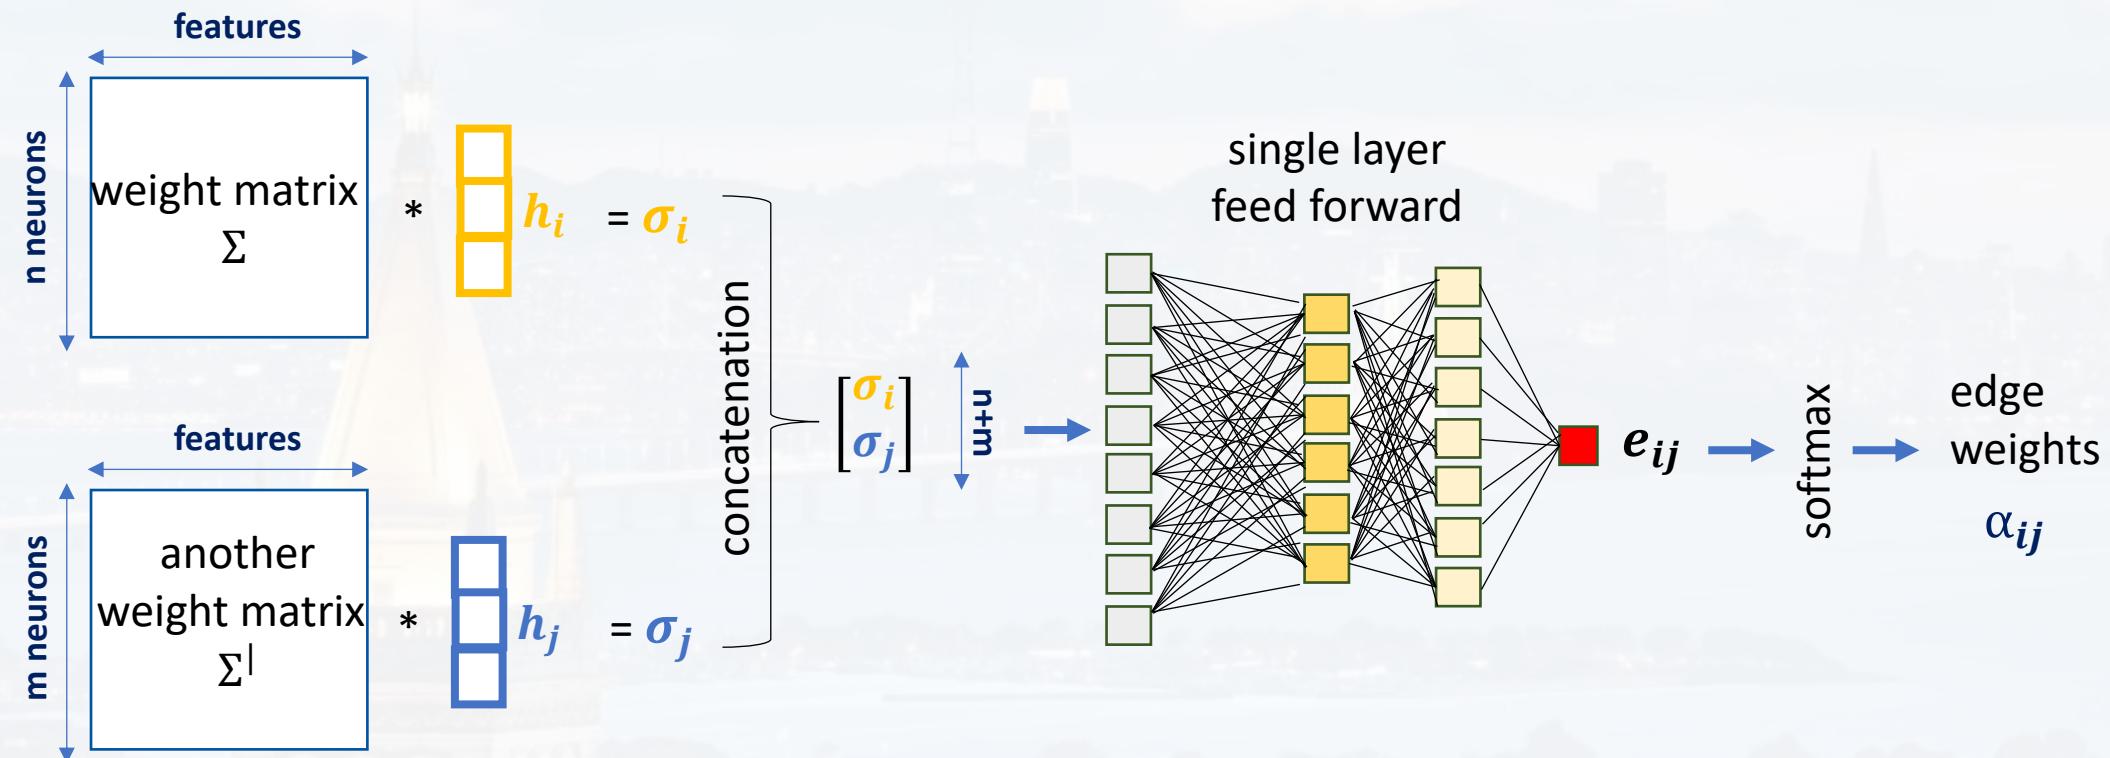
“The cat jumped on the roof.”

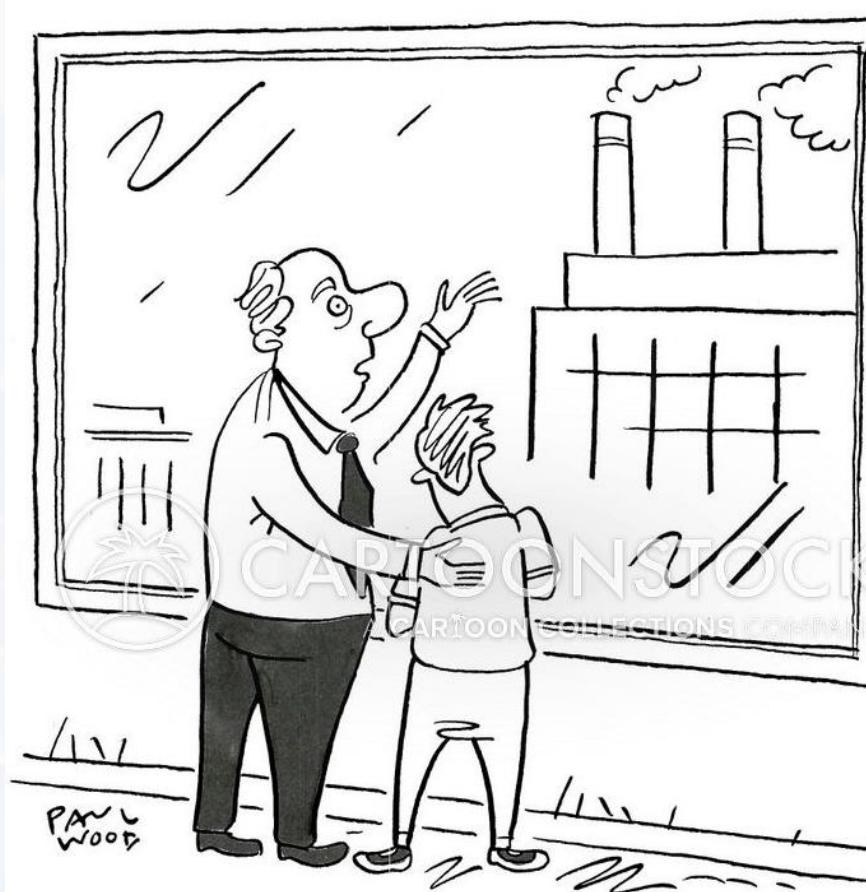




Graph Attention

Learning the weights!
(edge attributes)





ONE DAY SON, ALL THIS
WILL BE RUN BY ROBOTS

Outline

- What is a Graph
- The ANN Part
- PyTorch Example



node classification: convolution GNN

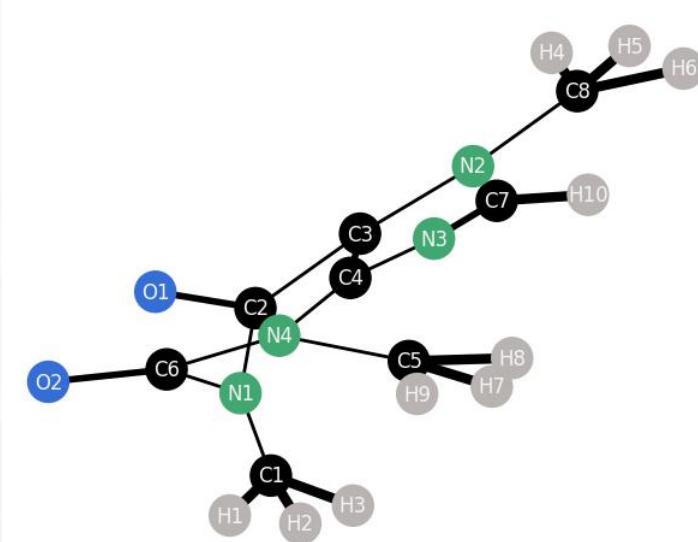
```
self.conv1 = GCNConv(n_node_features, n_neuron)  
self.conv2 = GCNConv(n_neuron, n_classes)
```

```
log_softmax(x3, dim = 1)
```

- edge weights: binding affinity

see Graph_III.ipynb

| epoch | loss | accuracy |
|-------|------|----------|
| 0 | 1.49 | 66.67% |
| 10 | 1.94 | 66.67% |
| 20 | 0.17 | 79.17% |
| 30 | 0.13 | 79.17% |
| 40 | 0.14 | 79.17% |
| 50 | 0.11 | 79.17% |
| 60 | 0.11 | 79.17% |
| 70 | 0.11 | 79.17% |
| 80 | 0.11 | 79.17% |
| 90 | 0.11 | 79.17% |
| 100 | 0.11 | 79.17% |
| 110 | 0.11 | 79.17% |
| 120 | 0.10 | 79.17% |
| 130 | 0.10 | 79.17% |
| 140 | 0.10 | 79.17% |
| 150 | 0.10 | 79.17% |
| 160 | 0.10 | 79.17% |



```
print(Y)  
print(Y_pred)
```

```
[0. 0. 0. 0. 0. 0. 0. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 2. 2. 2. 2. 3. 3.]  
tensor([0, 0, 0, 0, 0, 0, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 2, 2, 0, 0, 0])
```



Thank you very much for your attention!



ONE DAY SON, ALL THIS
WILL BE RUN BY ROBOTS