

Lecture 02:

Bayesian Methods



Markus Hohle
University California, Berkeley

Machine Learning Algorithms
MSSE 277B, 3 Units



Lecture 1: Course Overview and Introduction to Machine Learning

Lecture 2: Bayesian Methods in Machine Learning

classic ML tools & algorithms

Lecture 3: Dimensionality Reduction: Principal Component Analysis

Lecture 4: Linear and Non-linear Regression and Classification

Lecture 5: Unsupervised Learning: Clustering and Gaussian Mixture Models

Lecture 6: Adaptive Learning and Gradient Descent Optimization Algorithms

Lecture 7: Introduction to Artificial Neural Networks - The Perceptron

ANNs/AI/Deep Learning

Lecture 8: Introduction to Artificial Neural Networks - Building Multiple Dense Layers

Lecture 9: Convolutional Neural Networks (CNNs) - Part I

Lecture 10: CNNs - Part II

Lecture 11: Recurrent Neural Networks (RNNs) and Long Short-Term Memory (LSTMs)

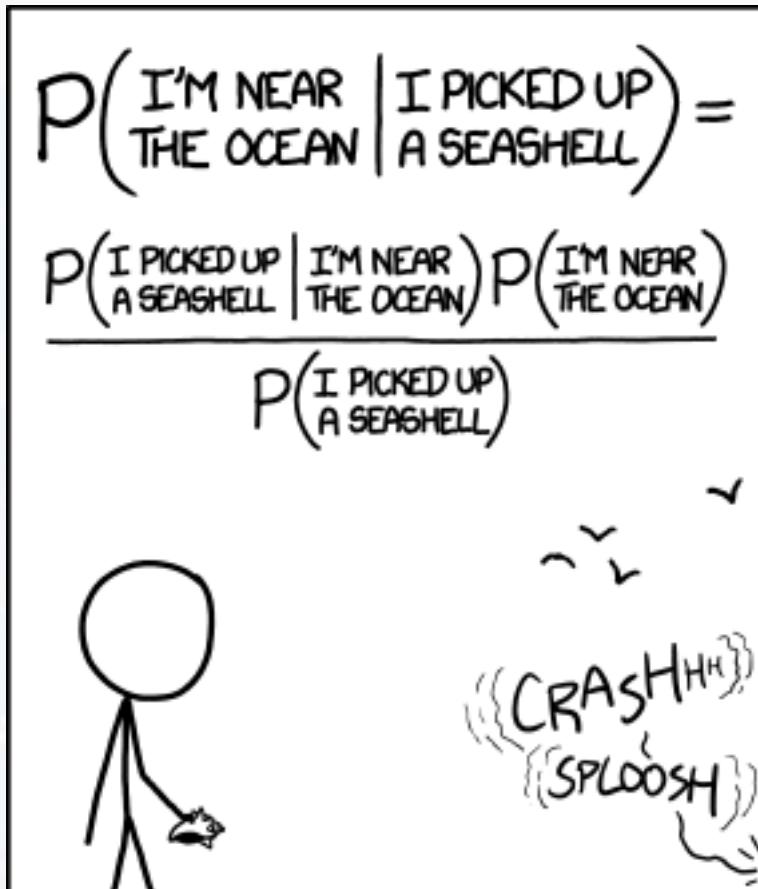
Lecture 12: Combining LSTMs and CNNs

Lecture 13: Running Models on GPUs and Parallel Processing

Lecture 14: Project Presentations

Lecture 15: Transformer

Lecture 16: GNN

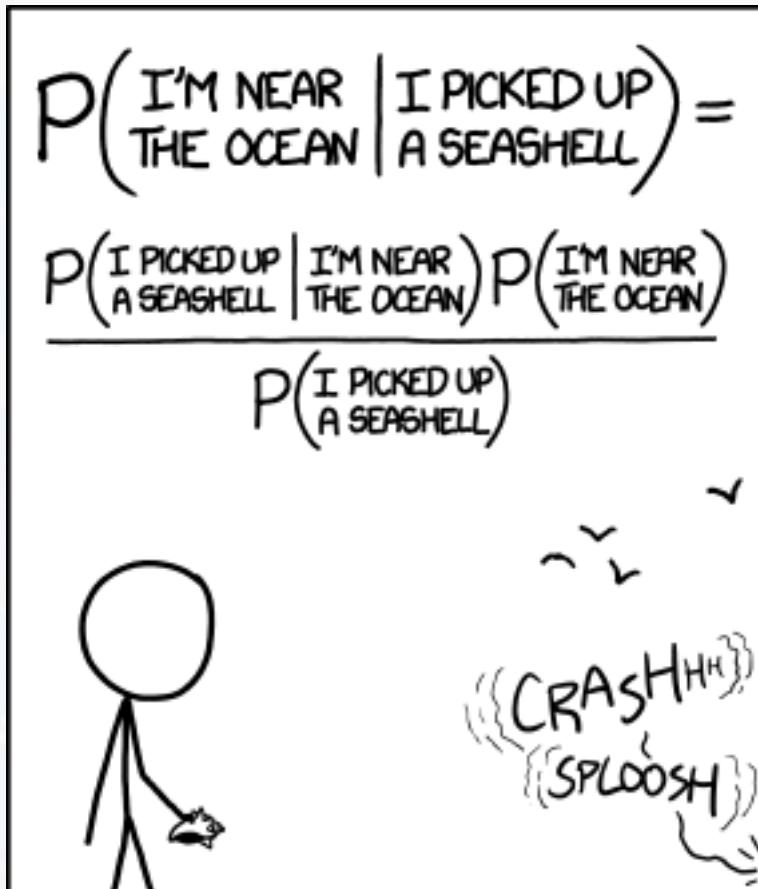


STATISTICALLY SPEAKING, IF YOU PICK UP A SEASHELL AND DON'T HOLD IT TO YOUR EAR, YOU CAN PROBABLY HEAR THE OCEAN.

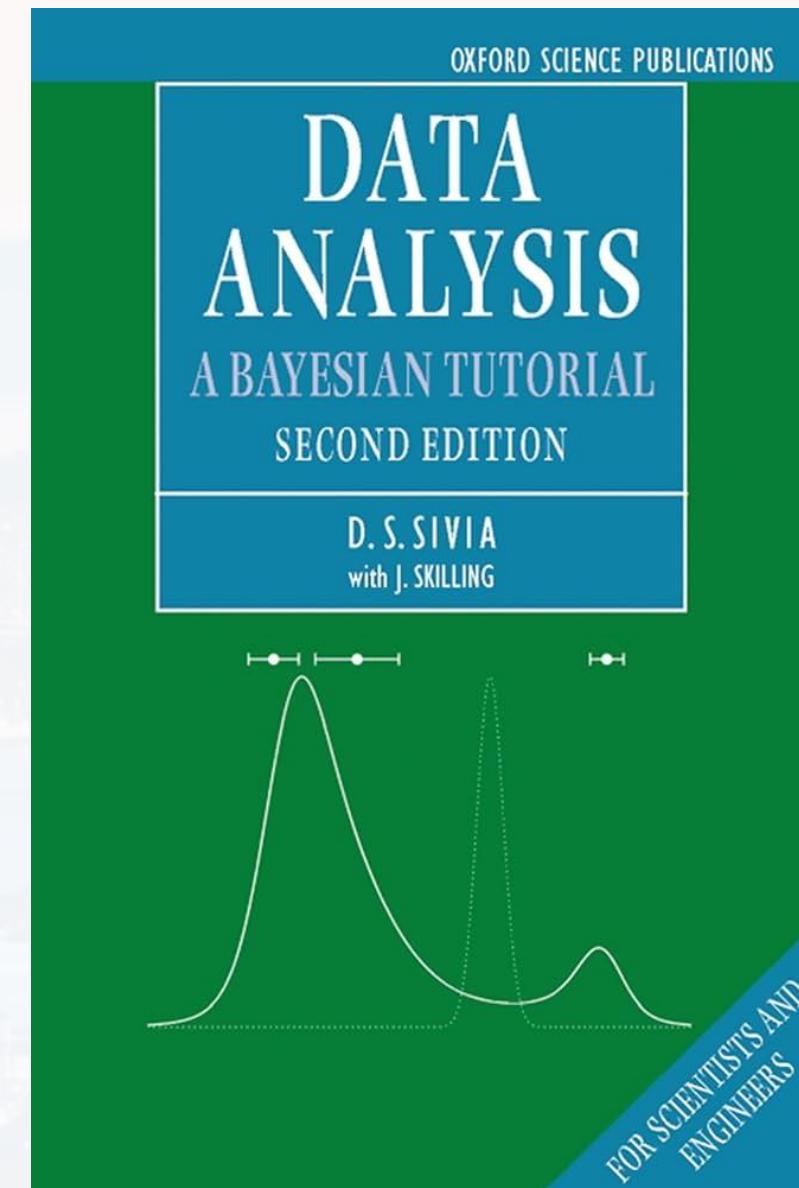
Outline

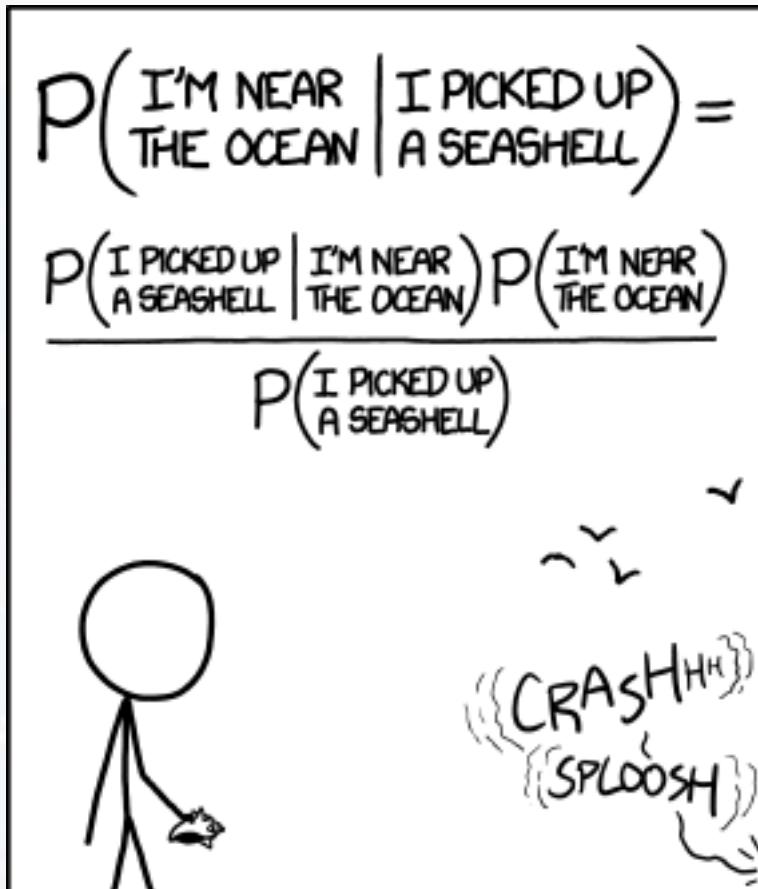
- The Idea and Bayes Theorem
(Discussion on Thursday,
BayesianExamples.ipynb)
- Naïve Bayes (today)
- Parameter Estimation (office hours, Friday)
- Model Selection (advanced, optional)
- Bayesian Networks (Graphs)
- Variational Bayes

FYI



STATISTICALLY SPEAKING, IF YOU PICK UP A SEASHELL AND DON'T HOLD IT TO YOUR EAR, YOU CAN PROBABLY HEAR THE OCEAN.





STATISTICALLY SPEAKING, IF YOU PICK UP A SEASHELL AND DON'T HOLD IT TO YOUR EAR, YOU CAN PROBABLY HEAR THE OCEAN.

Outline

- The Idea and Bayes Theorem

- Naïve Bayes

- Parameter Estimation

- Model Selection

FYI

- Bayesian Networks (Graphs)

- Variational Bayes



Why Bayesian Statistics?

frequentist: assuming sample is infinite (even tough there are corrections for small n)

vs:

Bayesian:

- taking the **exact amount** of information into account that's available
- model "**learns**" by adding more data (BPE)
- is based on **information theory & links to quantum mechanics**

→ **maximum entropy, given constrains** (prior knowledge)

→ variational calculus

- o EM algorithm (GMM, HMM etc)

- o **Variational Auto Encoder**

→ non-parametric (e. g. in contrast to MLE)

....and more



$P(A \cap B)$ probability P that the events A and B occur

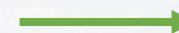
so far: A and B were independent $P(A \cap B) = P(A)P(B) = P(B)P(A)$



Thomas Bayes
(1701 - 1761)

now: **conditional probabilities** | “given” or “under the condition”

$$\begin{aligned} P(A \cap B) &= P(A|B)P(B) \\ &= P(B|A)P(A) \end{aligned}$$



$$P(A|B)P(B) = P(B|A)P(A)$$

Bayes Theorem

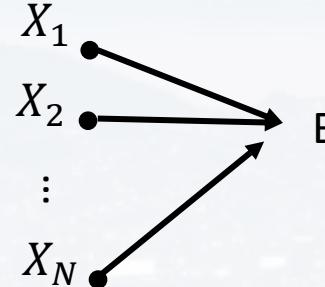
posterior $P(A|B) = \frac{P(B|A)P(A)}{P(B)}$ prior



$$P(A|B)P(B) = P(B|A)P(A)$$

Bayes Theorem

posterior $P(A|B) = \frac{P(B|A)P(A)}{P(B)}$ prior



$$P(B) = \sum_{n=1}^N P(B|X_n)P(X_n)$$
$$P(B) = \int P(B|X)P(X) dX$$

}

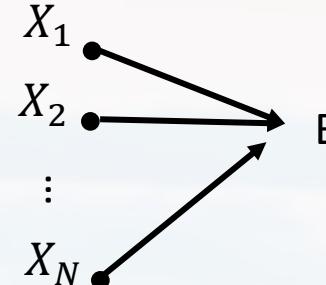
marginalization



Thomas Bayes
(1701 - 1761)

Probability $P(B)$ that I am going to be too late for a meeting:

$$P(B) = P(B|I forgot that I have a meeting) P(I forgot that I have a meeting) +$$
$$P(B|I got sick) P(I got sick) +$$
$$P(B|BART was too late) P(BART was too late) + \dots$$



$$P(B) = \sum_{n=1}^N P(B|X_n)P(X_n)$$
$$P(B) = \int P(B|X)P(X) dX$$

marginalization



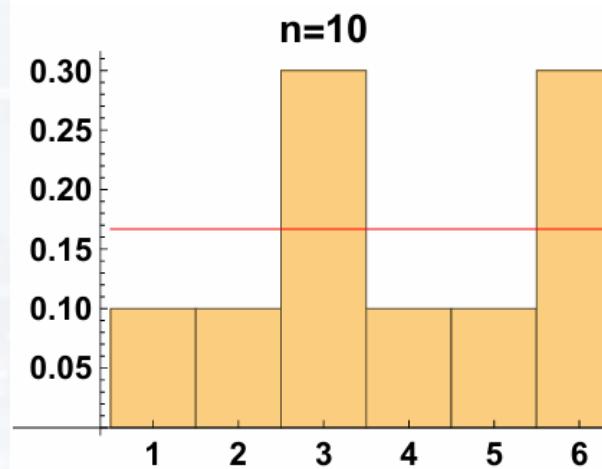
Thomas Bayes
(1701 - 1761)

model: M

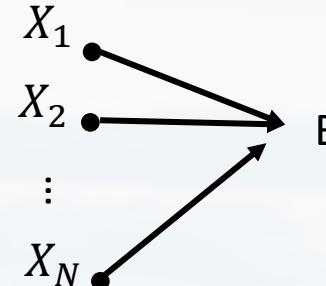
data: D

for a normal distribution M: $\mathcal{N}(\mu, \sigma)$

$$P(D|M) = \int P(D|\mu, \sigma, M) P(\mu, \sigma|M) d\Omega_{\mu, \sigma}$$



- σ = 2, μ = 3.5
- σ = 2, μ = 5.0
- σ = 1.5, μ = 3.5
- σ = 7.0, μ = 1.0 ...and so on



$$P(B) = \sum_{n=1}^N P(B|X_n)P(X_n)$$
$$P(B) = \int P(B|X)P(X) dX$$

]

marginalization



Thomas Bayes
(1701 - 1761)

example:

model: M
data: D

$$P(D|M) = \int P(D|all\ model\ param, M) P(all\ model\ param|M) d\ all\ model\ param$$

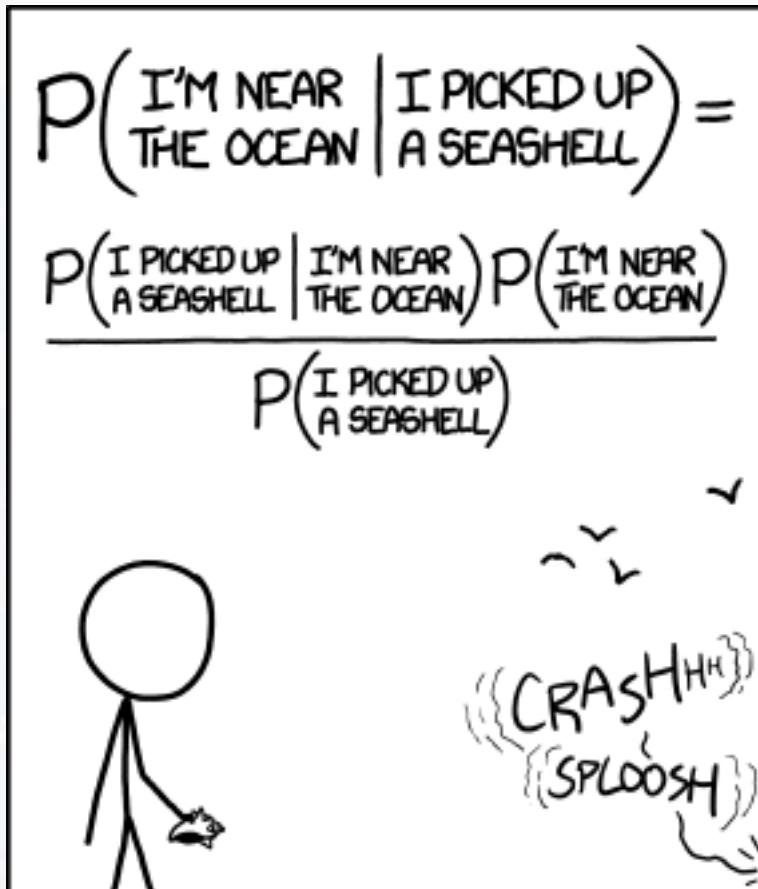
for a normal distribution $\mathcal{N}(\mu, \sigma)$

$$P(D|\mathcal{N}) = \int P(D|\mathcal{N}(\mu, \sigma)) P(\mu, \sigma|\mathcal{N}(\mu, \sigma)) d \Omega_{\mu, \sigma}$$

for a Poisson distribution $p(\lambda)$

$$P(D|p) = \int P(D|p(\lambda)) P(\lambda|p(\lambda)) d \lambda$$

and so on...



STATISTICALLY SPEAKING, IF YOU PICK UP A SEASHELL AND DON'T HOLD IT TO YOUR EAR, YOU CAN PROBABLY HEAR THE OCEAN.

Outline

- The Idea and Bayes Theorem

- Naïve Bayes

- Parameter Estimation

- Model Selection

FYI

- Bayesian Networks (Graphs)

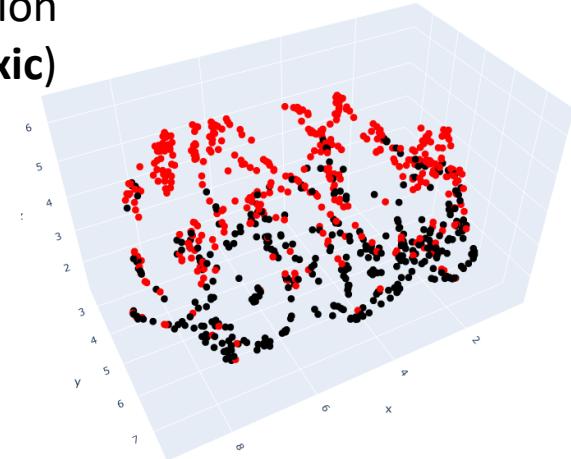
- Variational Bayes



$$P(A|B) = \frac{P(B|A)P(A)}{P(B)}$$

Bayes Theorem

UMAP projection
(toxic, non-toxic)



\vec{x} : vector with all model parameters (or features)

Index	molecular_weight	electronegativity	bond_lengths	num_hydrogen_bonds	logP	label
0	413.228	2.94416	3.41991	1	10.4335	Toxic
1	447.945	3.55371	3.66831	7	10.3475	Toxic
2	309.199	3.19761	2.84841	0	7.88825	Non-Toxic
3	382.554	3.8653	3.46237	8	9.59041	Toxic
4	310.904	3.18141	2.87774	6	7.85477	Non-Toxic
5	353.857	3.12105	3.32724	6	8.58887	Non-Toxic

K different classes
(here $K = 2$)



\vec{x} : vector with all model parameters (or features)

Index	molecular_weight	electronegativity	bond_lengths	num_hydrogen_bonds	logP	label
0	413.228	2.94416	3.41991	1	10.4335	Toxic
1	447.945	3.55371	3.66831	7	10.3475	Toxic
2	309.199	3.19761	2.84841	0	7.88825	Non-Toxic
3	382.554	3.8653	3.46237	8	9.59041	Toxic
4	310.904	3.18141	2.87774	6	7.85477	Non-Toxic
5	353.857	3.12105	3.32724	6	8.58887	Non-Toxic

$P(C_k|\vec{x})$: probability that datapoint belongs to class C_k , given \vec{x}

$P(\vec{x}|C_k)$: probability that datapoint has features \vec{x} , given class C_k

K different classes
(here $K = 2$)

$$P(A|B) = \frac{P(B|A)P(A)}{P(B)}$$

Bayes Theorem

$$P(C_k|\vec{x}) = \frac{P(\vec{x}|C_k)P(C_k)}{P(\vec{x})} = \frac{P(C_k)}{P(\vec{x})} \boxed{\prod_{i=1}^I P(x_i|C_k)} \sim P(C_k) \prod_{i=1}^I P(x_i|C_k) \quad \sum_{k=1}^K P(C_k|\vec{x}) = 1$$

Naïve Bayes:

- all features are **mutually independent**
- i. e.: no **correlation** between features
- features can be factorized



$$P(C_k|\vec{x}) = \frac{P(\vec{x}|C_k)P(C_k)}{P(\vec{x})} = \frac{P(C_k)}{P(\vec{x})} \prod_{i=1}^I P(x_i|C_k)$$

$$\sim P(C_k) \prod_{i=1}^I P(x_i|C_k)$$

$$\sum_{k=1}^K P(C_k|\vec{x}) = 1$$

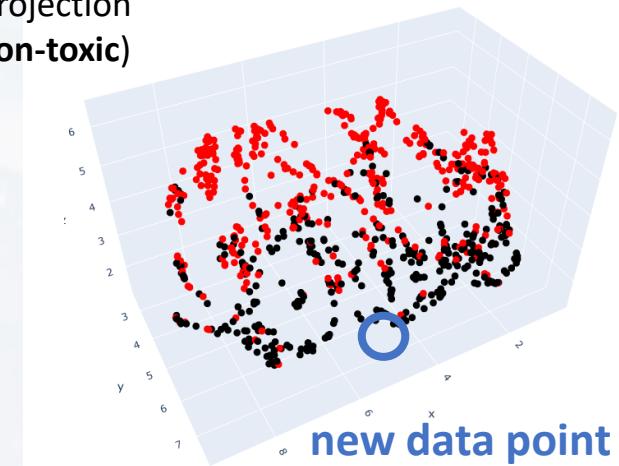
new data point: finding the class C_k , that maximizes $P(C_k|\vec{x})$

$$k_{\text{new}} = \underset{k}{\operatorname{argmax}} \left\{ P(C_k) \prod_{i=1}^I P(x_i|C_k) \right\}$$

from the training data → supervised learning

$P(C_k|\vec{x})$: probability that datapoint belongs to class C_k , given \vec{x}
 $P(\vec{x}|C_k)$: probability that datapoint has features \vec{x} , given class C_k

UMAP projection
(toxic, non-toxic)



different models for $P(x_i|C_k)$

- multinomial
- Gaussian
- ...



from the training data → supervised learning

1) creating the model:

```
my_model = library.method(argument1 = 'arg1', ... )
```

2) training the model

```
out = my_model.fit(xtrain, ytrain)
```

3) evaluation

```
ypred = out.predict(xeval)
accur = (ypred == yeval).sum()/len(yeval)
```

4) prediction (actual application)

```
ypred = out.predict(xnew)
```



Python:

```
from sklearn.naive_bayes import *
from sklearn.preprocessing import MinMaxScaler
```

importing methods for
naïve bayes

scaling/normalizing data

```
Train = pd.read_csv('molecular_train_gbc_cat.csv')
Test = pd.read_csv('molecular_test_gbc_cat.csv')
```

```
XTrain = Train.drop('Label', axis = 1).values
YTrain = Train['Label']
```

```
XTest = Test.drop('Label', axis = 1).values
YTest = Test['Label']
```

```
print(YTrain[:10])
```

0	Toxic
1	Toxic
2	Non-Toxic
3	Non-Toxic
4	Non-Toxic
5	Toxic
6	Non-Toxic
7	Non-Toxic
8	Toxic
9	Non-Toxic

Name: label, dtype: object



```
scaler = MinMaxScaler(feature_range = (0, 1))  
XTrainS = scaler.fit_transform(XTrain)
```

scaling the data to
mean = 0 and std = 1

```
gnb = GaussianNB()  
Fit = gnb.fit(XTrainS, YTrain)
```

the actual fit

applying the model to the test data set

```
XTestS = scaler.transform(XTest)
```

scaling the test set
without fitting

```
Ypred = Fit.predict(XTestS)  
Probs = Fit.predict_proba(XTestS)
```

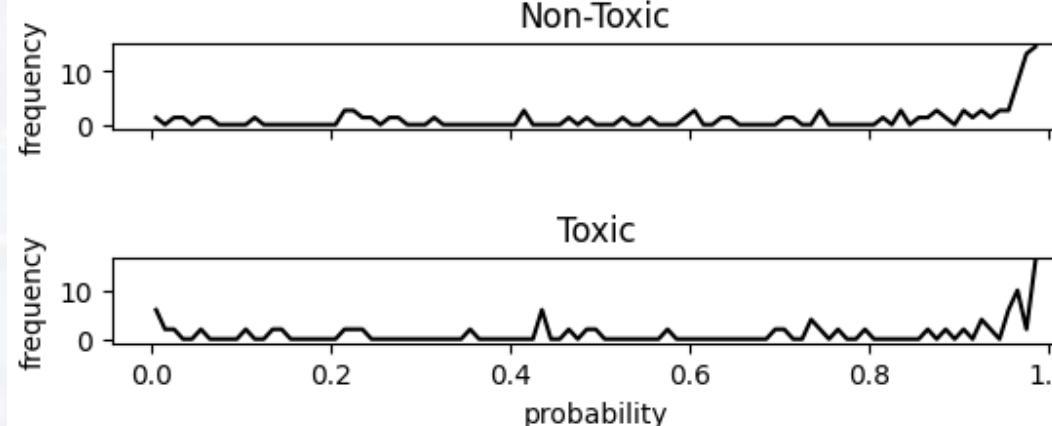
predicting the class
 C_k and calculating
the probabilities



```
XTestS = scaler.transform(XTest)
```

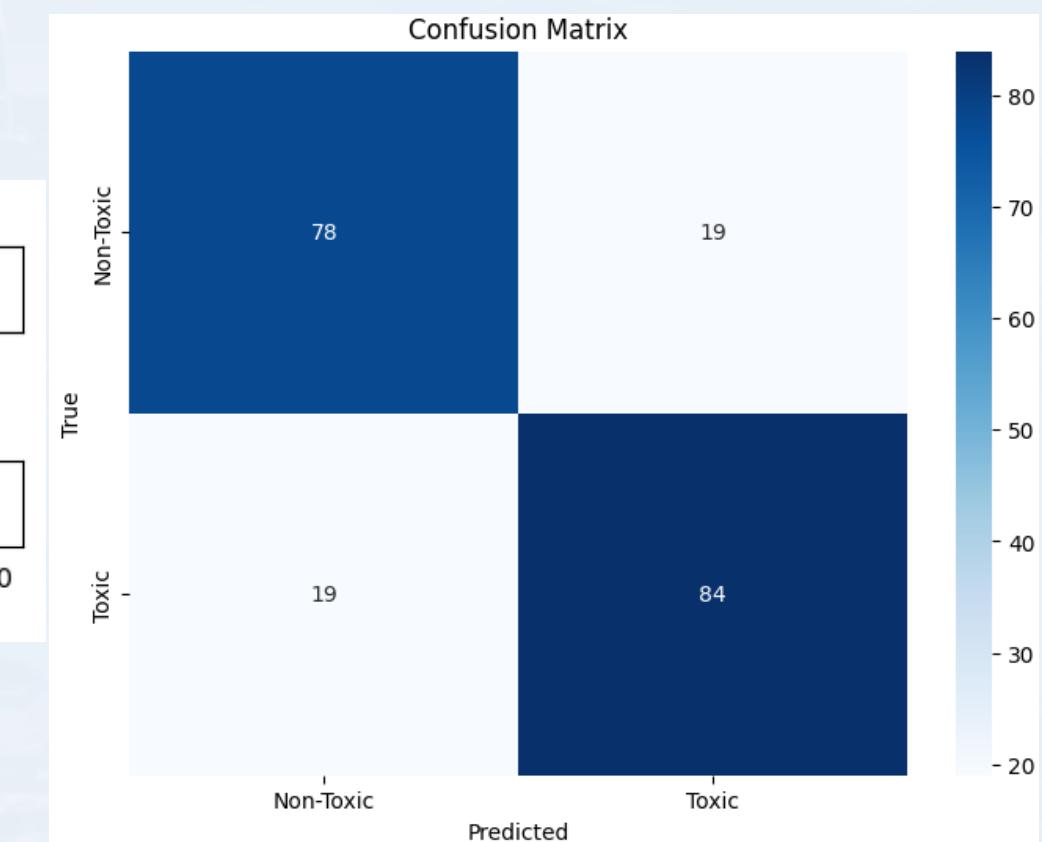
```
Ypred = Fit.predict(XTestS)
Probs = Fit.predict_proba(XTestS)
```

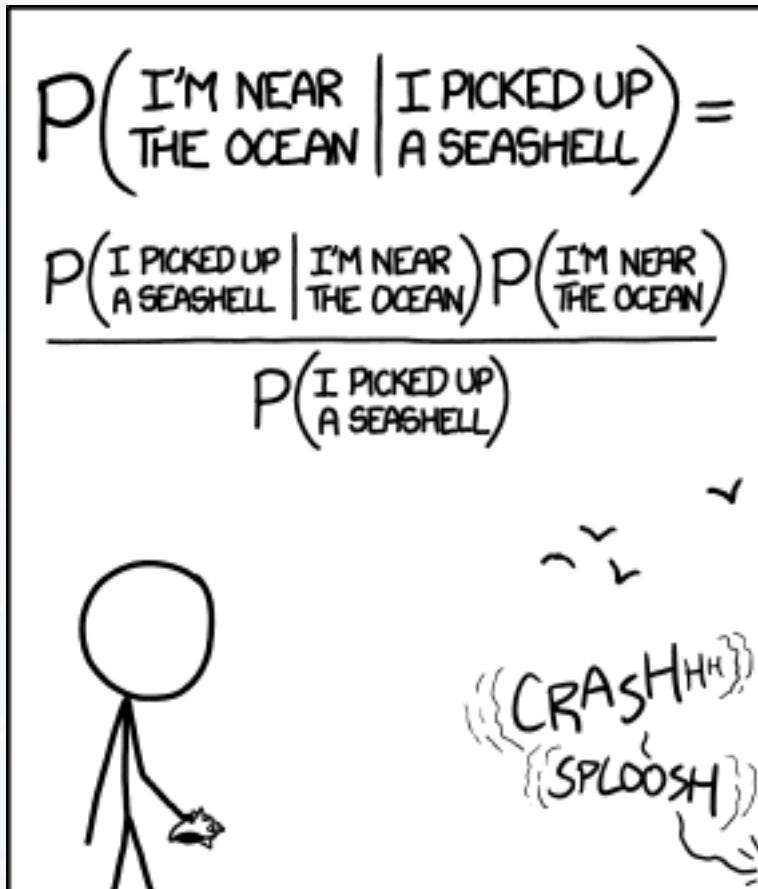
evaluation



see

[Walk_Through_NaiveBayes.ipynb](#)





STATISTICALLY SPEAKING, IF YOU PICK UP A
SEASHELL AND DON'T HOLD IT TO YOUR EAR,
YOU CAN PROBABLY HEAR THE OCEAN.

Outline

- The Idea and Bayes Theorem

- Naïve Bayes

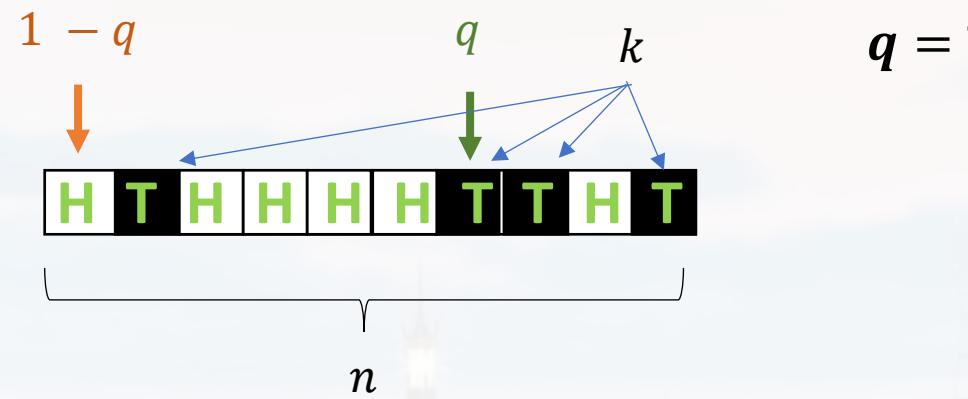
- Parameter Estimation

- Model Selection

FYI

- Bayesian Networks (Graphs)

- Variational Bayes



fair coin? $q = 0.5$???

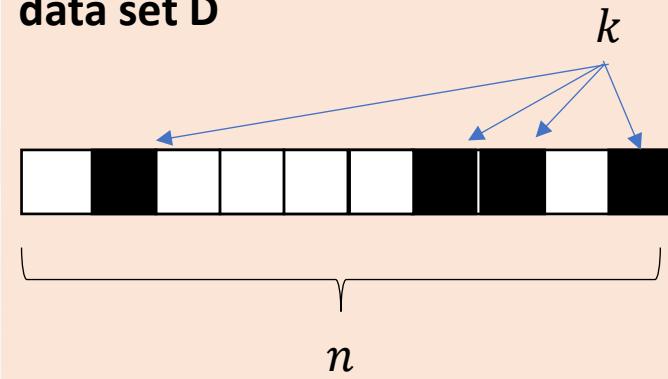


mutation $q = ??$





data set D

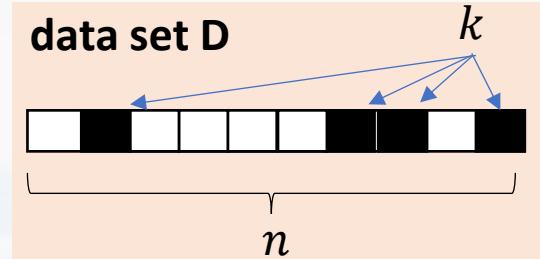


$q = ?$

goal:

- $P(q|D)$

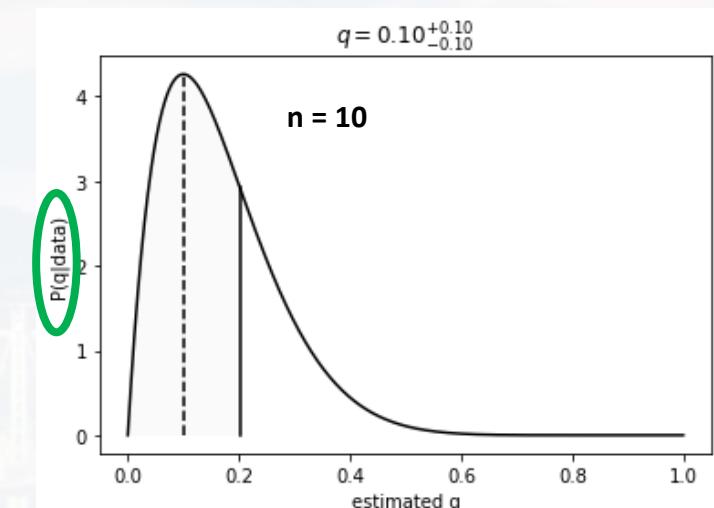
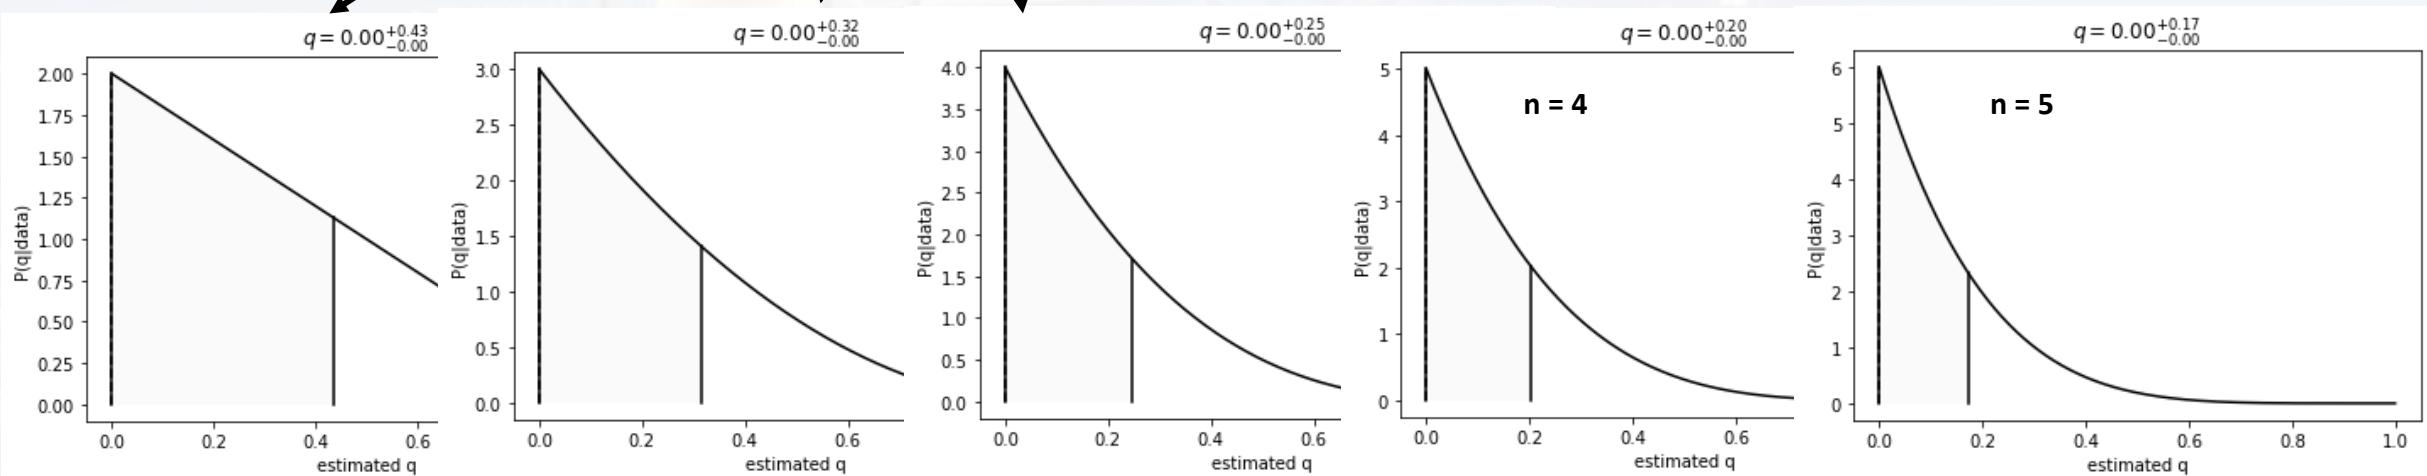
- the larger D , the more certain q
→ learning

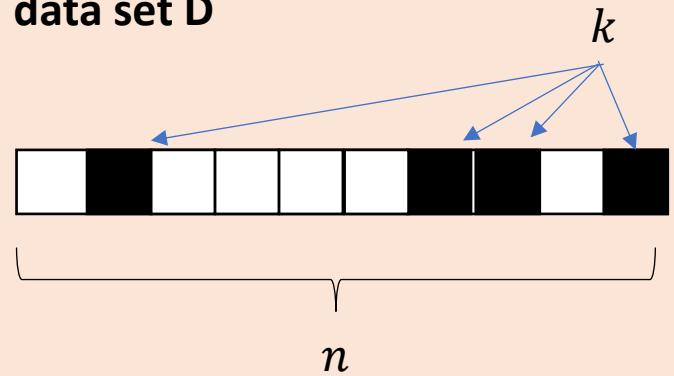
 $q = ?$

goal:

- $P(q|D)$
- the larger D , the more certain q
→ learning

$n = 1 \quad n = 2 \quad n = 3 \quad n = 4 \quad n = 5 \quad \dots \quad n = 9 \quad n = 10$



**data set D** $q = ?$ **goal:**

- $P(q|D)$
- the larger D , the more certain q
→ learning

$$P(k|n, q) = \binom{n}{k} q^k (1 - q)^{n-k}$$

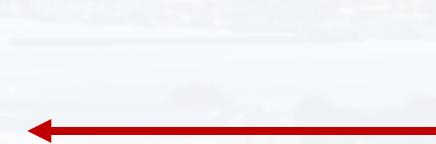
Bayes' theorem:

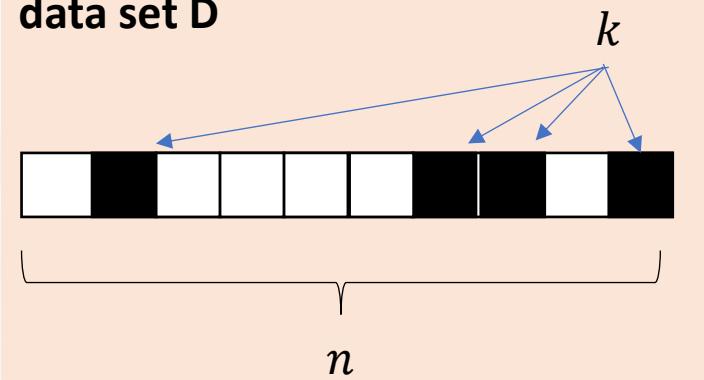
likelihood function (here: binomial)

$$P(q|data\ set) = \frac{P(data\ set|q)P(q)}{P(data\ set)} \text{ prior evidence (const wrt q)}$$

$$= \frac{\binom{n}{k} q^k (1-q)^{n-k}}{P(D)} P(q)$$

$$\sim q^k (1 - q)^{n-k} P(q)$$

 $P(D)$ and $\binom{n}{k}$ are no functions of q 

**data set D** $q = ?$ **goal:**

- $P(q|D)$
- the larger D , the more certain q
→ learning

$$P(k|n, q) = \binom{n}{k} q^k (1 - q)^{n-k}$$

$$P(q|data\ set) = \frac{P(data\ set|q)P(q)}{P(data\ set)}$$

$$= \frac{\binom{n}{k} q^k (1 - q)^{n-k}}{P(D)} P(q)$$

$$\sim q^k (1 - q)^{n-k} P(q)$$

$$\sim q^k (1 - q)^{n-k}$$

max. entropy: $P(q) = const$
if no prior information about q

$$P(q|data\ set) = \frac{q^k (1 - q)^{n-k}}{\int_0^1 q^k (1 - q)^{n-k} dq}$$



check out bayesian_bino.py

n1 = 4

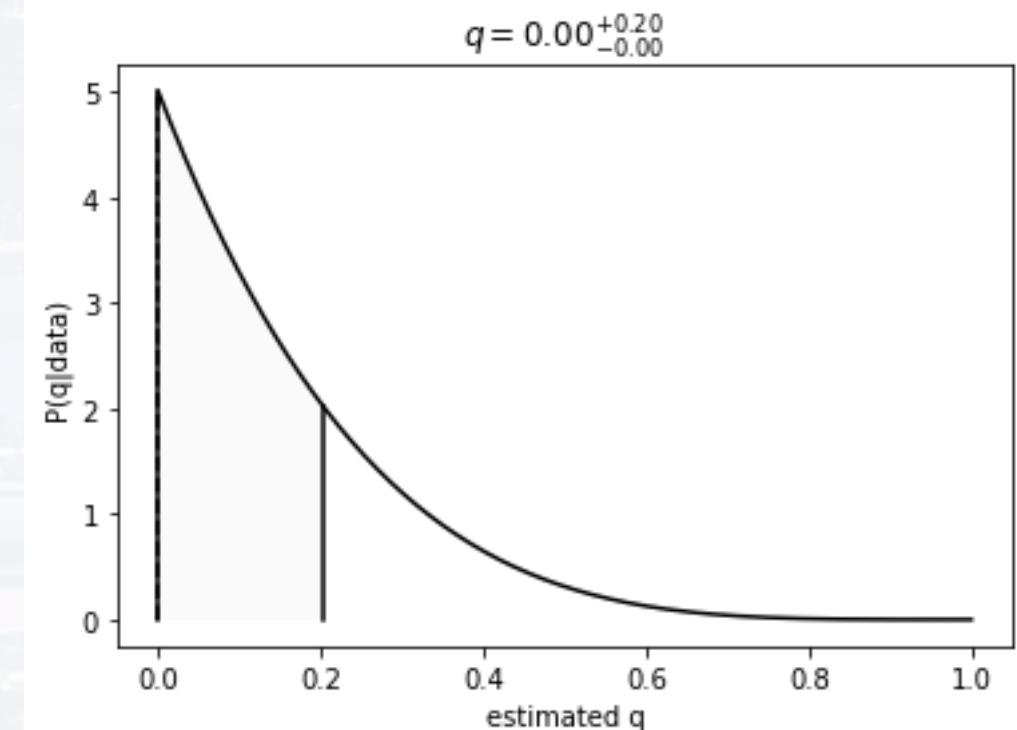
k1 = np.random.binomial(n1, 0.25)

creating artificial data set

note: in reality **q** is unknown!

[q1, b, _] = bayesian_bino(n1, k1)

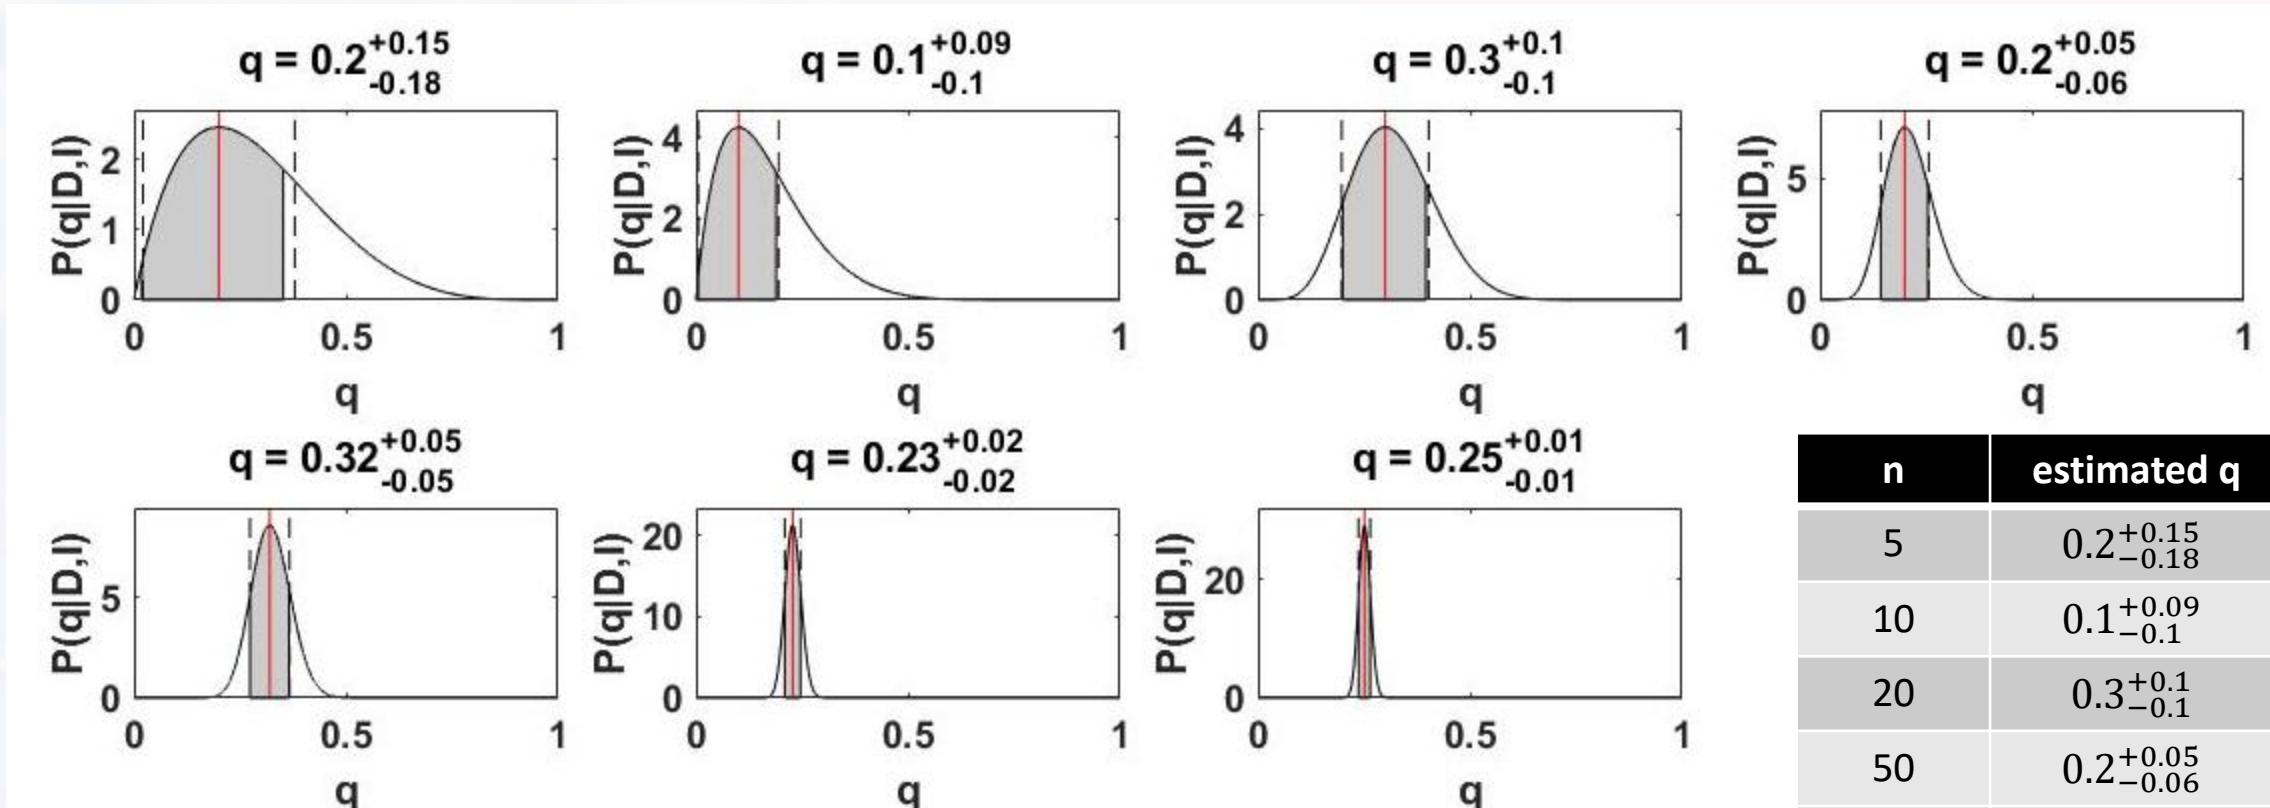
$$P(q|data\ set) = \frac{q^k(1-q)^{n-k}}{\int_0^1 q^k(1-q)^{n-k} dq}$$





check out bayesian_bino.py

$$P(q|data\ set) = \frac{q^k(1-q)^{n-k}}{\int_0^1 q^k(1-q)^{n-k} dq}$$

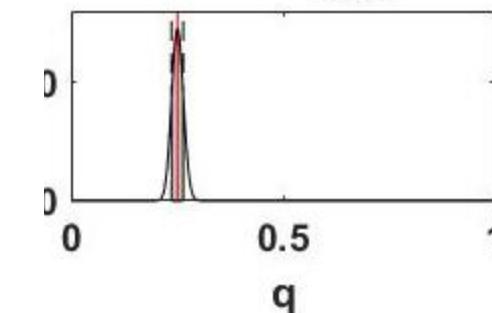
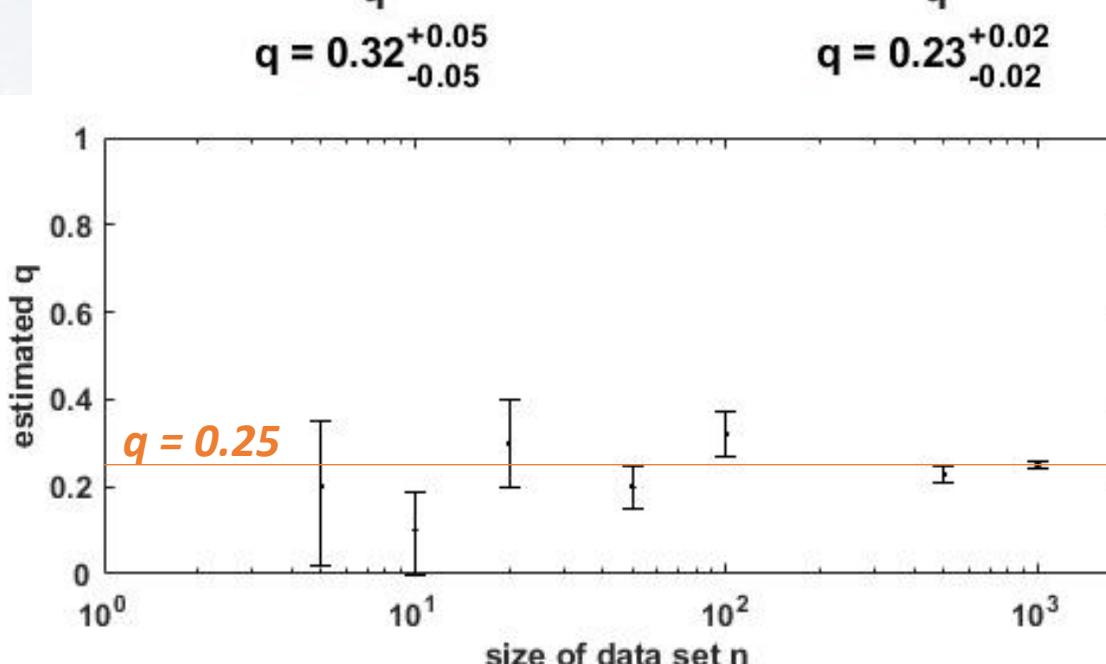
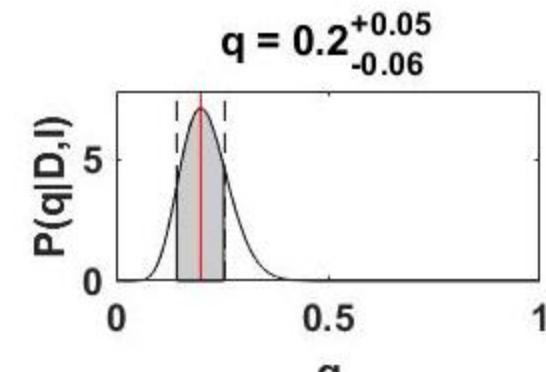
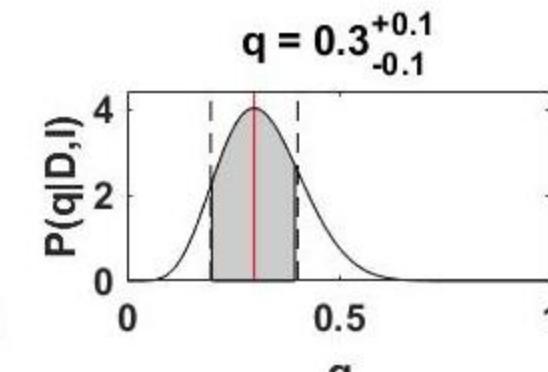
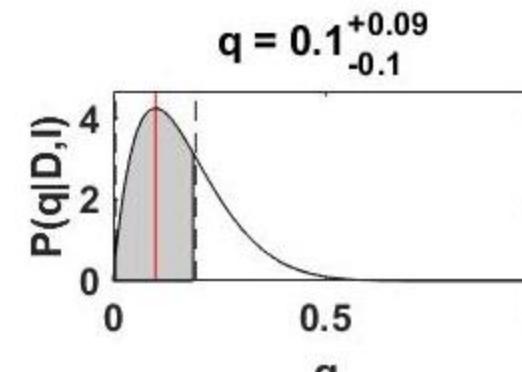
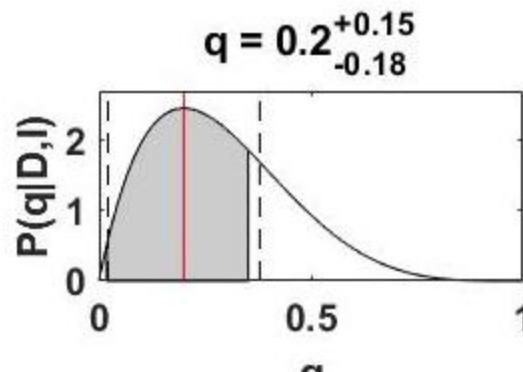


n	estimated q
5	$0.2^{+0.15}_{-0.18}$
10	$0.1^{+0.09}_{-0.1}$
20	$0.3^{+0.1}_{-0.1}$
50	$0.2^{+0.05}_{-0.06}$
100	$0.32^{+0.05}_{-0.05}$
500	$0.23^{+0.02}_{-0.02}$
1,000	$0.25^{+0.01}_{-0.01}$
infinity	0.25



check out bayesian_bino.py

$$P(q|data\ set) = \frac{q^k(1-q)^{n-k}}{\int_0^1 q^k(1-q)^{n-k} dq}$$



n	estimated q
5	$0.2^{+0.15}_{-0.18}$
10	$0.1^{+0.09}_{-0.1}$
20	$0.3^{+0.1}_{-0.1}$
50	$0.2^{+0.05}_{-0.06}$
100	$0.32^{+0.05}_{-0.05}$
500	$0.23^{+0.02}_{-0.02}$
1,000	$0.25^{+0.01}_{-0.01}$
infinity	0.25

Of course, Bayesian Parameter Estimation
works with **any other pdf**

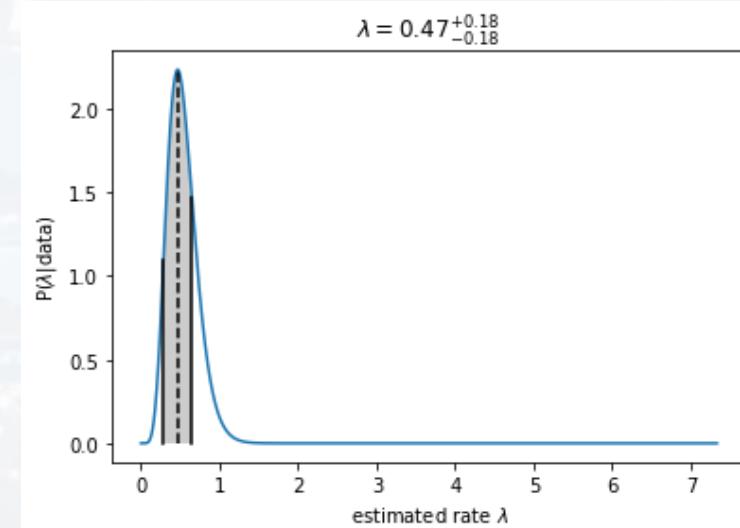
likelihood function

$$P(q|data\ set) = \frac{P(data\ set|q)P(q)}{P(data\ set)} \text{ prior} \\ P(data\ set) \text{ evidence (const wrt q)}$$

What is the average number of WhatsUp messages I get every day?

Mon:	5	event	- has no duration
Tue:	7		- is rare
Wed:	1		
Thu:	3		→ Poissonian
Fri:	9		
Sat:	2		
Sun:	5		$P(k \lambda) = \frac{\lambda^k}{k!} e^{-\lambda}$

```
data = np.random.poisson(lam = 0.4, 15)  
poissfit(data)
```





Of course, Bayesian Parameter Estimation works with **any other pdf**

goal:

- $P(q|D)$
- the larger D , the more certain q
→ learning

What is the average number of WhatsUp messages I get every day?

Mon: 5 event - has no duration

Tue: 7 - is rare

Wed: 1

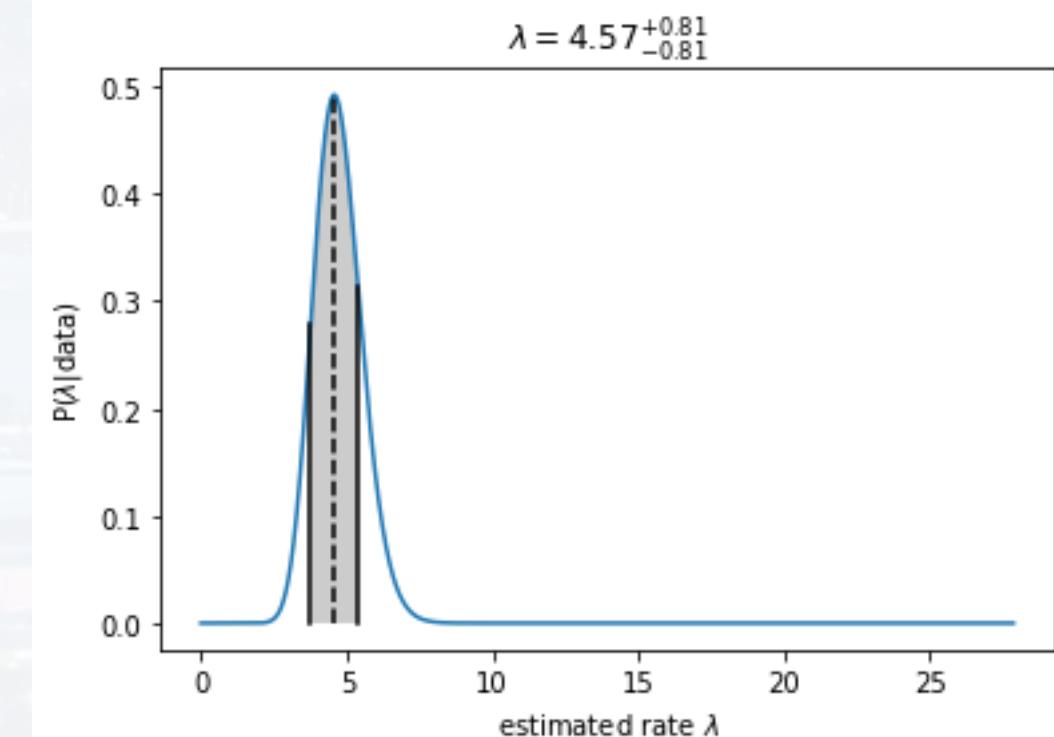
Thu: 3 → Poissonian

Fri: 9

Sat: 2 $P(k|\lambda) = \frac{\lambda^k}{k!} e^{-\lambda}$

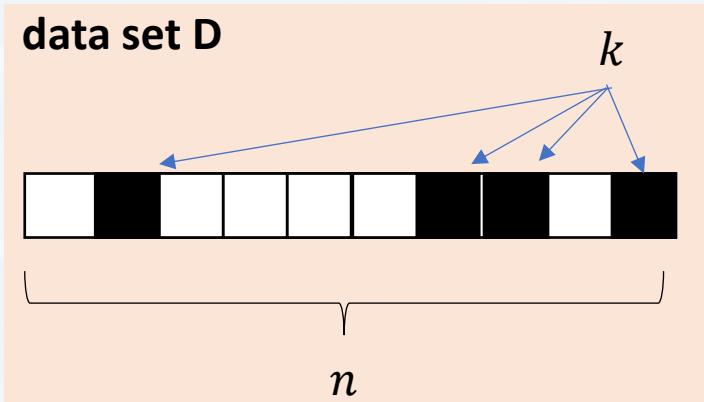
Sun: 5

`poissfit([5, 7, 1, 3, 9, 2, 5])`



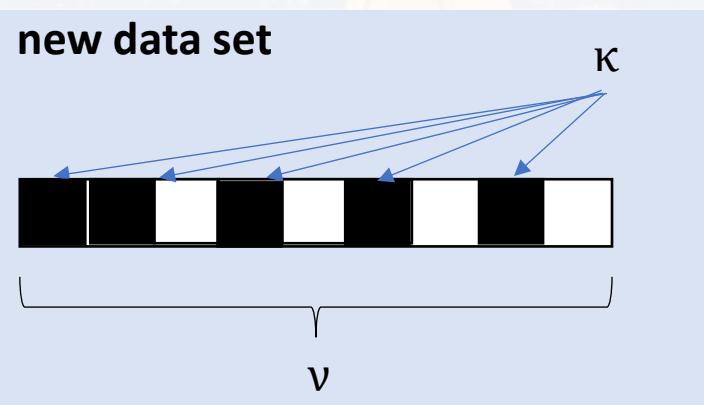
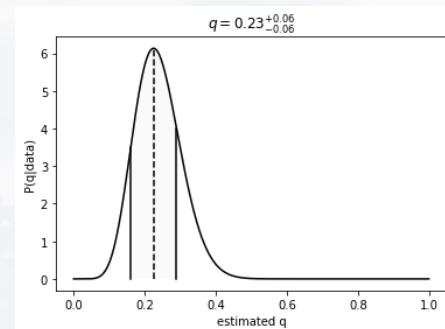


What if there is new data?



~~$q = ?$~~

$$P(q|data\ set) = \frac{q^k(1-q)^{n-k}}{\int_0^1 q^k(1-q)^{n-k} dq}$$



if there **is** prior information I about q :

$$P(q|new\ data\ set, I) = \frac{P(new\ data\ set|q, I) P(q, I)}{P(new\ data\ set)}$$



What if there is new data?

$$P(q|new\ data\ set, I) = \frac{P(new\ data\ set|q, I) P(q, I)}{P(new\ data\ set)}$$

$$P(q|data\ set) = \frac{q^k(1-q)^{n-k}}{\int_0^1 q^k(1-q)^{n-k} dq}$$

$$= \frac{q^\kappa(1-q)^{\nu-\kappa}}{\int_0^1 q^\kappa(1-q)^{\nu-\kappa} dq} \frac{q^k(1-q)^{n-k}}{q^k(1-q)^{n-k} dq}$$

$$= \frac{q^{k+\kappa}(1-q)^{\nu-\kappa+n-k}}{\int_0^1 q^{k+\kappa}(1-q)^{\nu-\kappa+n-k} dq}$$

often: $\kappa = \alpha - 1$
 $\beta = \nu - \kappa - 1$

$$= \frac{q^{k+\alpha-1}(1-q)^{n-k+\beta-1}}{\int_0^1 q^{k+\alpha-1}(1-q)^{n-k+\beta-1} dq}$$

Beta function

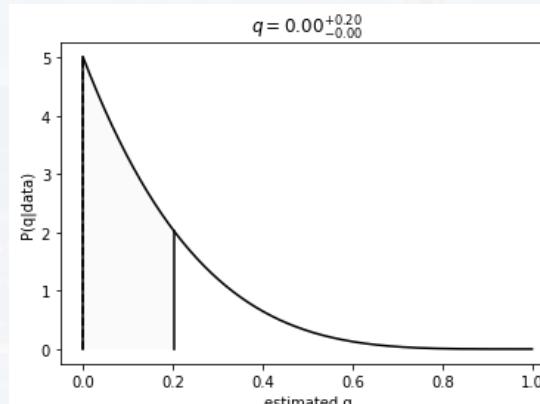


What if there is new data?

$$P(q|new\ data\ set, I) = \frac{q^k(1-q)^{n-k}}{\int_0^1 q^k(1-q)^{n-k} dq}$$

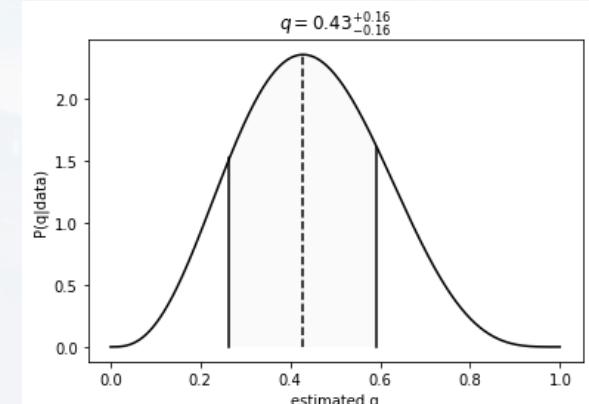
n1 = 4

```
k1 = np.random.binomial(n1, q = 0.2)
[_, _, Prior] = bayesian_bino(n1, k1)
```



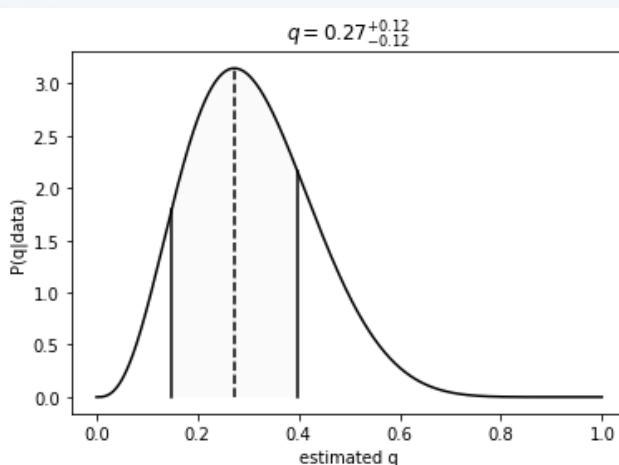
n2 = 7

```
k2 = np.random.binomial(n2, q = 0.2)
[_, _, _] = bayesian_bino(n2, k2)
```



$$P(q, I) = \frac{q^k(1-q)^{n-k}}{\int_0^1 q^k(1-q)^{n-k} dq}$$

```
[_, _, _] = bayesian_bino(n2, k2, Prior = Prior)
```

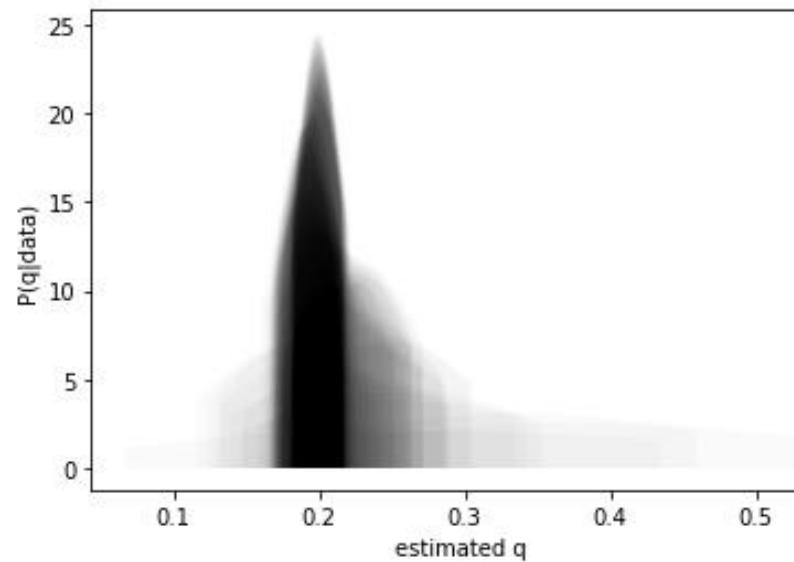




What if there is new data?

$$P(q|new\ data\ set, I) = \frac{q^{\kappa}(1-q)^{v-\kappa}}{\int_0^1 q^{\kappa}(1-q)^{v-\kappa} dq} \frac{q^k(1-q)^{n-k}}{\int_0^1 q^k(1-q)^{n-k} dq}$$

The **posterior** from the **previous** experiment is the **prior** of the **next** experiment



- we become more certain about the model parameters
- learning!
- see e.g. **Variational Auto Encoders**

2D images → 3D objects



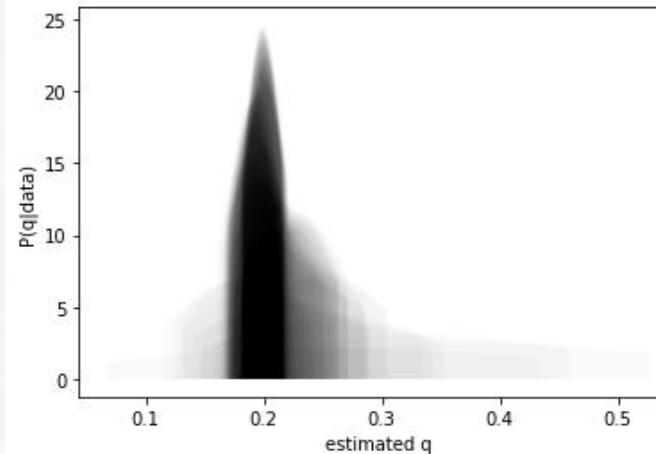
credit: StableAI



What if there is new data?

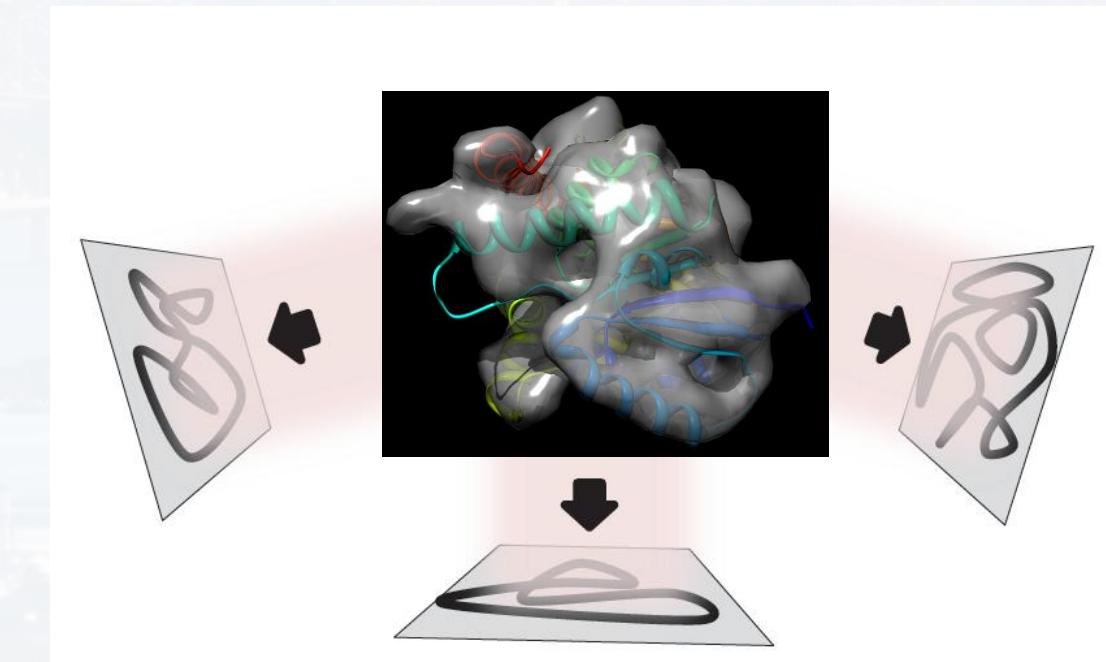
$$P(q|new\ data\ set, I) = \frac{q^{\kappa}(1-q)^{v-\kappa}}{\int_0^1 q^{\kappa}(1-q)^{v-\kappa} dq} \frac{q^k(1-q)^{n-k}}{\int_0^1 q^k(1-q)^{n-k} dq}$$

The **posterior** from the **previous** experiment is the **prior** of the **next** experiment

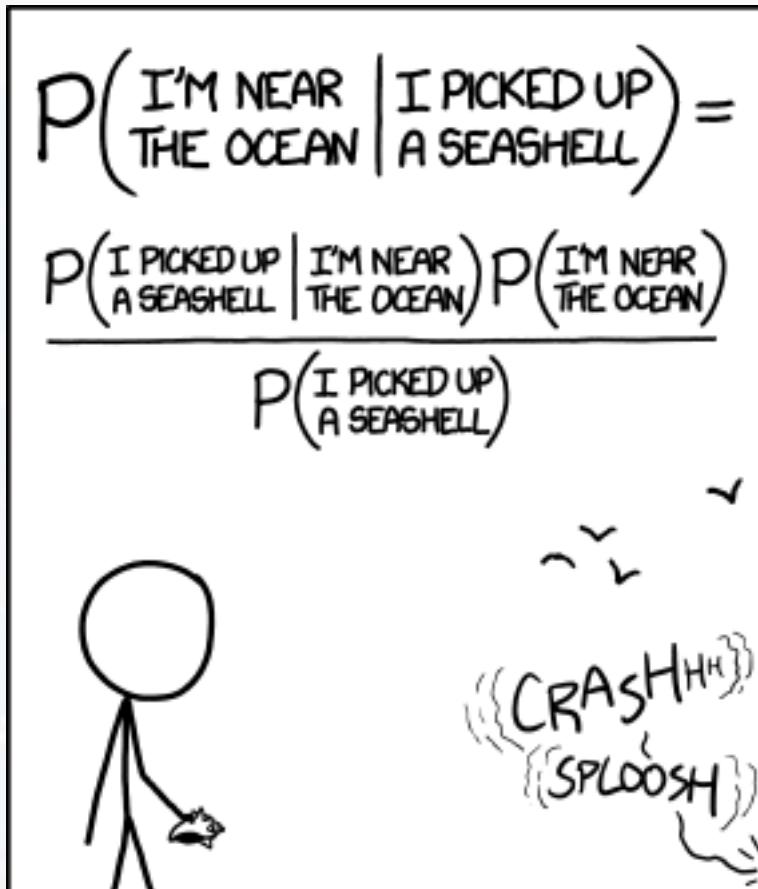


- we become more certain about the model parameters
- learning!
- see e.g. **Variational Auto Encoders** 2D images → 3D objects

Cryo – EM: 3D structure from 2D images/projections



(image courtesy: Thomas Becker, GC LMU Munich)



STATISTICALLY SPEAKING, IF YOU PICK UP A SEASHELL AND DON'T HOLD IT TO YOUR EAR, YOU CAN PROBABLY HEAR THE OCEAN.

Outline

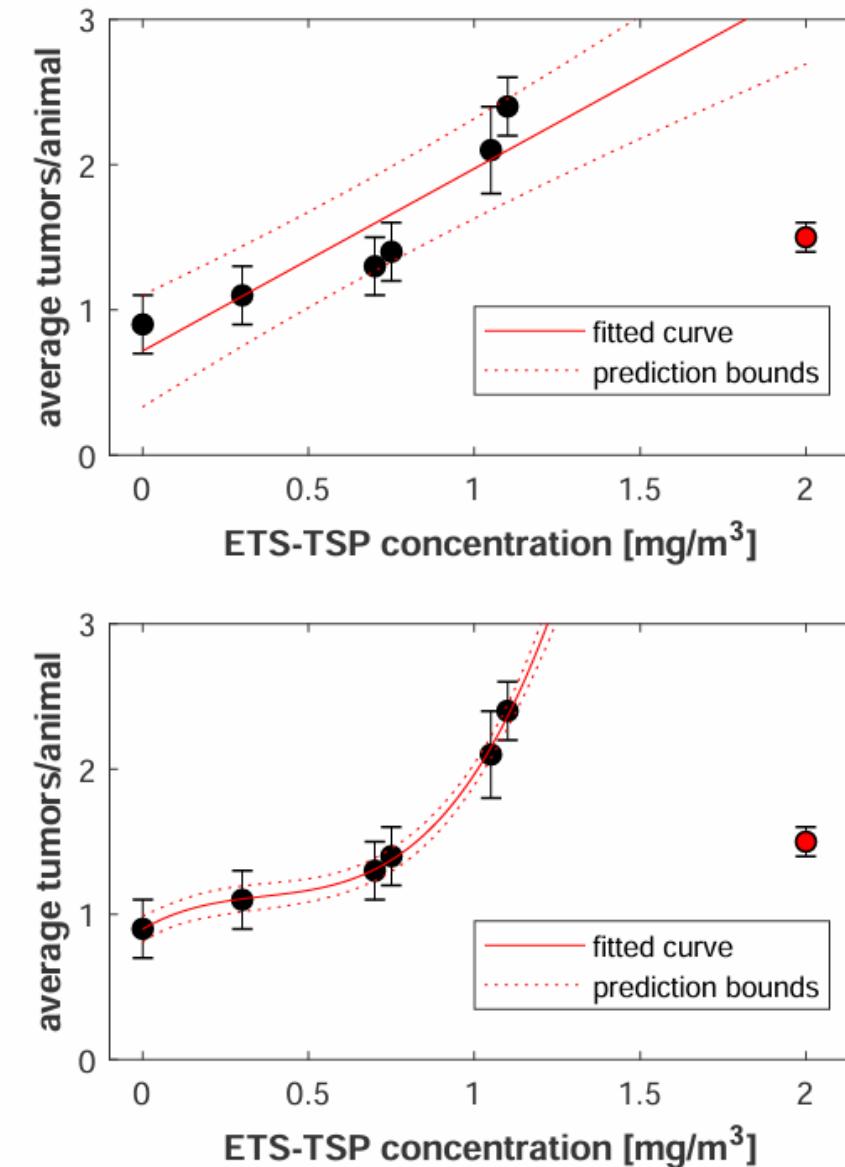
- The Idea and Bayes Theorem
- Naïve Bayes
- Parameter Estimation
- Model Selection

FYI

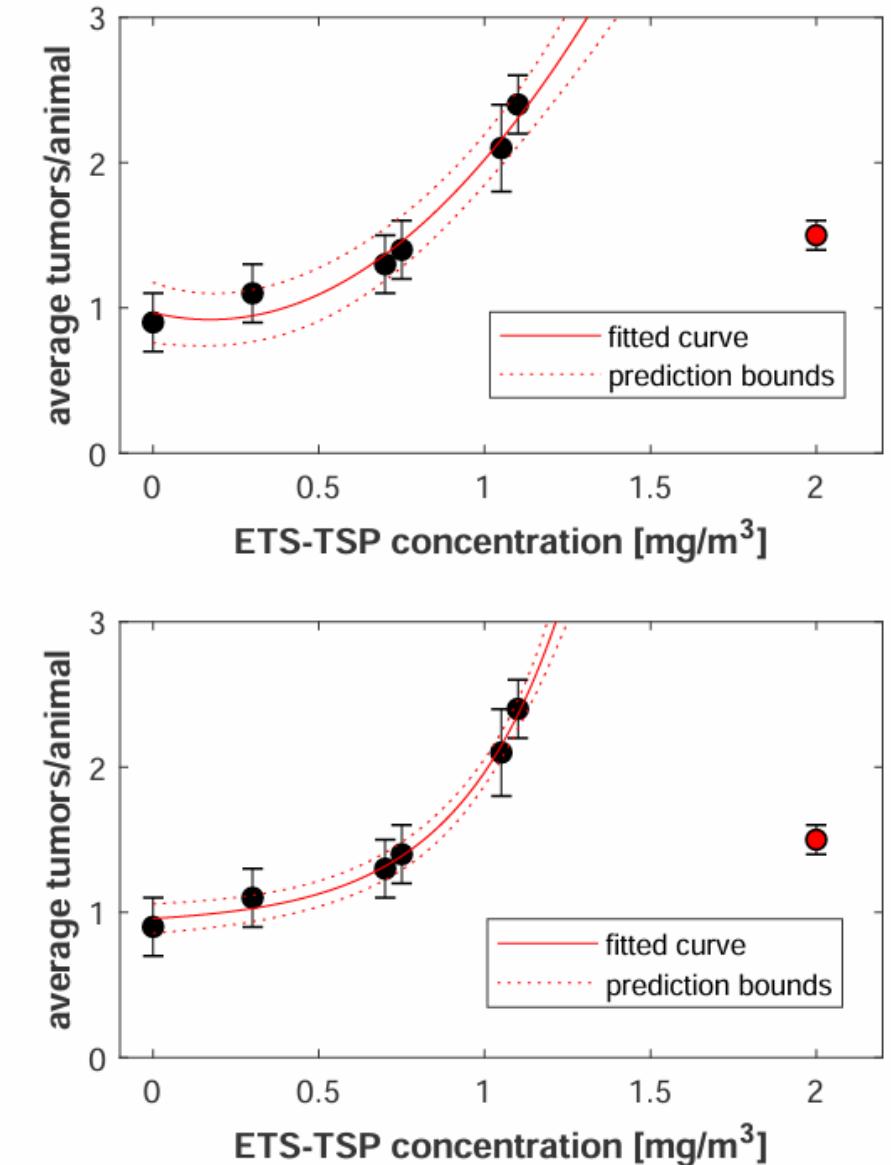
- Bayesian Networks (Graphs)
- Variational Bayes



often, we have many competing models



→ assigning probabilities if a model is correct





often, we have many competing models

→ assigning probabilities if a model is correct

goal: $\rho = \frac{P(M_A|D)}{P(M_B|D)}$

Bayes' theorem

$$= \frac{P(D|M_A) P(M_A)}{P(D)} \cdot \frac{P(D)}{P(D|M_B) P(M_B)}$$

D : data
M_A : model A
M_B : model B

marginalization:

$\{\alpha\}_i$: all parameter of model M_i

$$P(D|M_i) = \int P(D|\{\alpha\}_i M_i) P(\{\alpha\}_i | M_i) d\Omega_{\{\alpha\}_i}$$

assuming all α_{ij} are
mutually independent
(Naïve Bayes)

$$= \int P(D|\{\alpha\}_i, M_i) \prod_j P(\alpha_{ij} | M_i) d\alpha_{ij}$$



goal:

$$\rho = \frac{P(M_A|D)}{P(M_B|D)} = \frac{P(D|M_A) P(M_A)}{P(D)} \cdot \frac{P(D)}{P(D|M_B) P(M_B)}$$

 D
 M_A
 M_B
 $\{\alpha\}_i$: data
: model A
: model B
: all parameter of model M_i

marginalization:

$$P(D|M_i) = \int P(D|\{\alpha\}_i M_i) P(\{\alpha\}_i | M_i) d\Omega_{\{\alpha\}_i}$$

assuming all α_{ij} are
mutually independent
(Naïve Bayes)

$$= \int \underbrace{P(D|\{\alpha\}_i, M_i)}_{\text{likelihood function}} \prod_j \underbrace{P(\alpha_{ij} | M_i)}_{\text{prior of } \alpha_{ij} \text{ BEVORE(!) measurement}} d\alpha_{ij}$$

↓
likelihood function
→ the actual model

→ prior of α_{ij} BEVORE(!) measurement
Maximum Entropy without prior knowledge:
$$\frac{1}{\alpha_{ij}(\max) - \alpha_{ij}(\min)}$$



goal:

$$\rho = \frac{P(M_A|D)}{P(M_B|D)} = \frac{P(D|M_A) P(M_A)}{P(D)} \cdot \frac{P(D)}{P(D|M_B) P(M_B)}$$

D
 M_A
 M_B
 $\{\alpha\}_i$

: data
: model A
: model B
: all parameter of model M_i

marginalization:

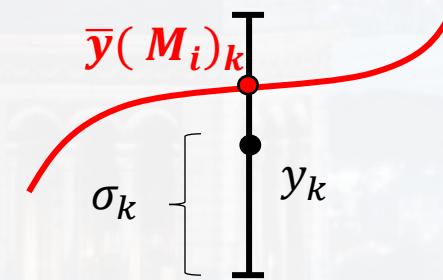
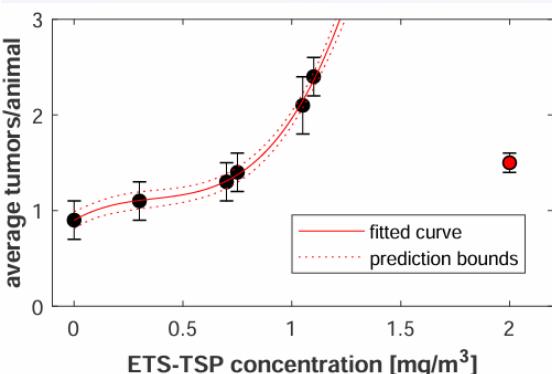
$$P(D|M_i) = \int P(D|\{\alpha\}_i M_i) P(\{\alpha\}_i | M_i) d\Omega_{\{\alpha\}_i}$$

assuming all α_{ij} are
mutually independent
(Naïve Bayes)

$$= \int P(D|\{\alpha\}_i, M_i) \prod_j P(\alpha_{ij} | M_i) d\alpha_{ij}$$

$$= \prod_j \frac{1}{\alpha_{ij}(max) - \alpha_{ij}(min)} \cdot \int P(D|\{\alpha\}_i, M_i) d\alpha_{ij}$$

likelihood function
→ the actual model



y_k : measured value
 σ_k : error
 $\bar{y}(M_i)_k$: model value (after fit)



goal:

$$\rho = \frac{P(M_A|D)}{P(M_B|D)} = \frac{P(D|M_A) P(M_A)}{P(D)} \cdot \frac{P(D)}{P(D|M_B) P(M_B)}$$

D
 M_A
 M_B
 $\{\alpha\}_i$
 y_k
 σ_k
 $\bar{y}(M_i)_k$

: data
: model A
: model B
: all parameter of model M_i
: measured value
: error
: model value (after fit)

marginalization:

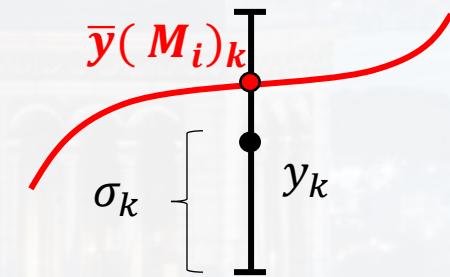
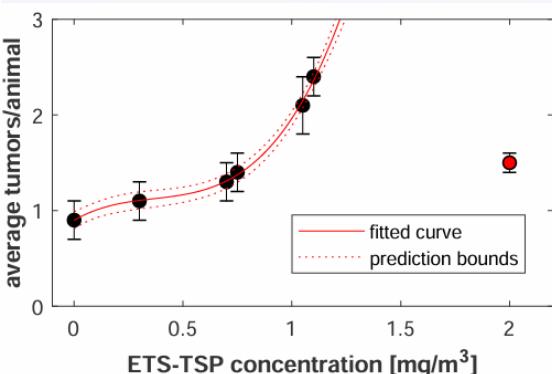
$$P(D|M_i) = \int P(D|\{\alpha\}_i M_i) P(\{\alpha\}_i | M_i) d\Omega_{\{\alpha\}_i}$$

assuming all α_{ij} are mutually independent (Naïve Bayes)

$$= \int P(D|\{\alpha\}_i, M_i) \prod_j P(\alpha_{ij} | M_i) d\alpha_{ij}$$

$$= \prod_j \frac{1}{\alpha_{ij}^{(max)} - \alpha_{ij}^{(min)}} \cdot \int P(D|\{\alpha\}_i, M_i) d\alpha_{ij}$$

likelihood function
→ the actual model



$$P(y_k | \alpha_{ij}, M_i) \approx \frac{1}{\sqrt{2\pi\sigma_k^2}} e^{-\frac{1}{2} \frac{(\bar{y}(M_i)_k - y_k)^2}{\sigma_k^2}}$$

for $\sigma_k \ll |y_k|$



marginalization:

$$P(D|M_i) = \int P(D|\{\alpha\}_i M_i) P(\{\alpha\}_i | M_i) d\Omega_{\{\alpha\}_i}$$

$$= \int P(D|\{\alpha\}_i, M_i) \prod_j P(\alpha_{ij} | M_i) d\alpha_{ij}$$

$$= \prod_j \frac{1}{\alpha_{ij}(max) - \alpha_{ij}(min)} \cdot \int P(D|\{\alpha\}_i, M_i) d\alpha_{ij}$$

D	: data
M_A	: model A
M_B	: model B
$\{\alpha\}_i$: all parameter of model M_i
y_k	: measured value
σ_k	: error
$\bar{y}(M_i)_k$: model value (after fit)

likelihood function
→ the actual model

$$P(D|\{\alpha\}_i M_i) = \prod_k P(y_k | \alpha_{ij}, M_i) = \prod_k \frac{1}{\sqrt{2\pi\sigma_k^2}} e^{-\frac{1}{2} \frac{(\bar{y}(M_i)_k - y_k)^2}{\sigma_k^2}}$$

$$= \left(\prod_k \frac{1}{\sqrt{2\pi\sigma_k^2}} \right) \cdot e^{-\frac{1}{2} \sum_k \frac{(\bar{y}(M_i)_k - y_k)^2}{\sigma_k^2}} = \left(\prod_k \frac{1}{\sqrt{2\pi\sigma_k^2}} \right) \cdot e^{-\frac{1}{2} \cancel{x_i^2}}$$



$$\rho = \frac{P(D|M_A) P(M_A)}{P(D|M_B) P(M_B)}$$

D	: data
M_A	: model A
M_B	: model B
$\{\alpha\}_i$: all parameter of model M_i
y_k	: measured value
σ_k	: error
$\bar{y}(M_i)_k$: model value (after fit)

$$= \frac{P(M_A)}{P(M_B)} \cdot \frac{\int e^{-\frac{1}{2}\chi_A^2} d\Omega_{\{\alpha\}_A}}{\int e^{-\frac{1}{2}\chi_B^2} d\Omega_{\{\alpha\}_B}} \cdot \frac{\prod_j \alpha_{jB}(\max) - \alpha_{jB}(\min)}{\prod_j \alpha_{jA}(\max) - \alpha_{jA}(\min)}$$

prior probability of each model: maximum entropy $\rightarrow 1:1$

fit quality: integral over χ^2

Occam's Razor: simple models are preferred

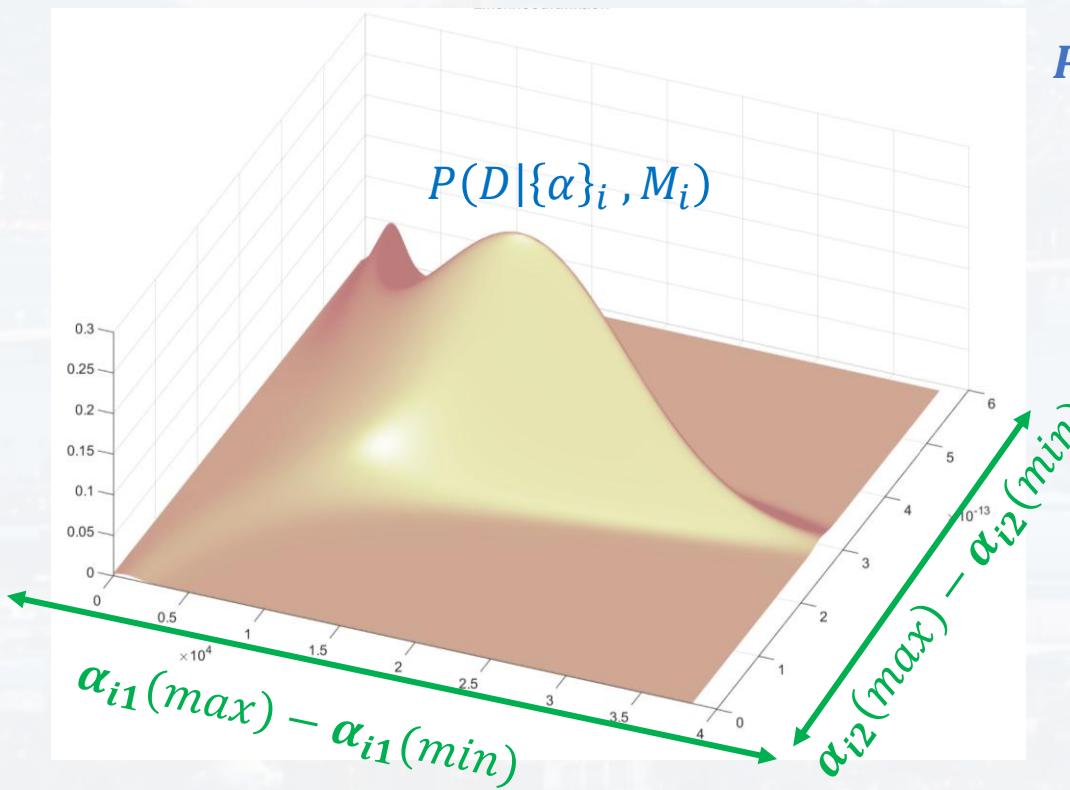


$$\rho = \frac{P(D|M_A) P(M_A)}{P(D|M_B) P(M_B)}$$

$$= \frac{P(M_A)}{P(M_B)} \cdot$$

$$\frac{\int e^{-\frac{1}{2}\chi_A^2} d\Omega_{\{\alpha\}_A}}{\int e^{-\frac{1}{2}\chi_B^2} d\Omega_{\{\alpha\}_B}} \cdot \frac{\prod_j \alpha_{jB}(max) - \alpha_{jB}(min)}{\prod_j \alpha_{jA}(max) - \alpha_{jA}(min)}$$

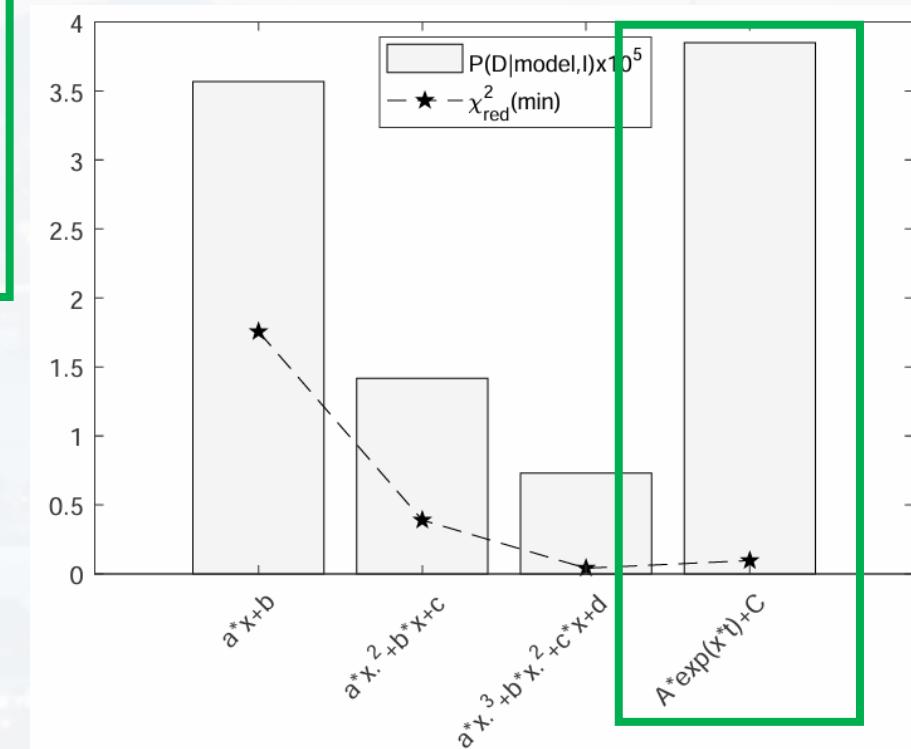
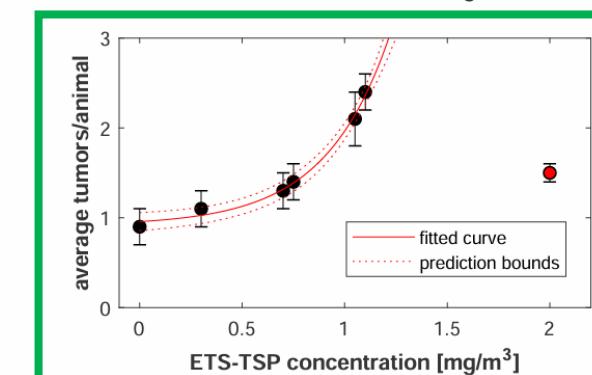
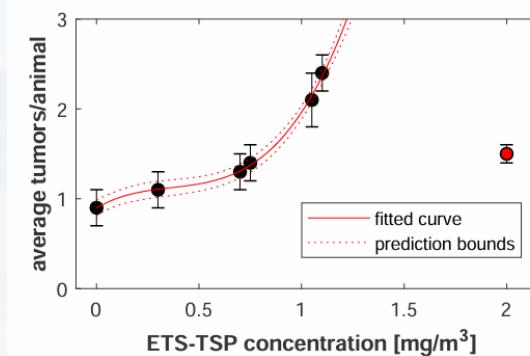
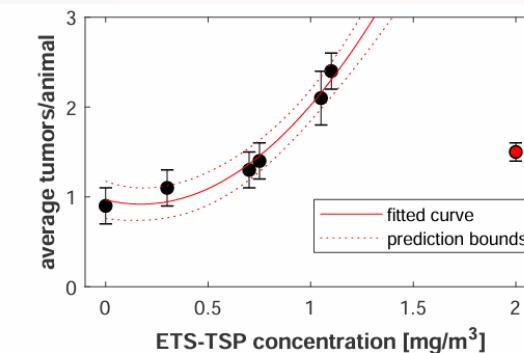
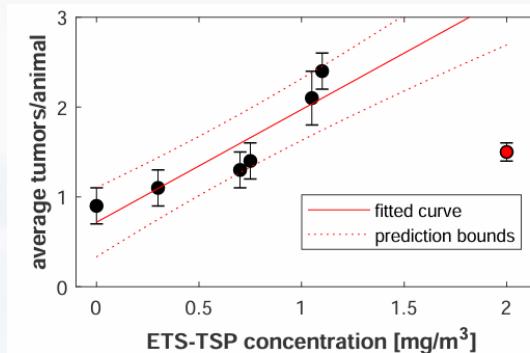
D	: data
M_A	: model A
M_B	: model B
$\{\alpha\}_i$: all parameter of model M_i
y_k	: measured value
σ_k	: error
$\bar{y}(M_i)_k$: model value (after fit)



$$P(D|M_i) = \int P(D|\{\alpha\}_i M_i) P(\{\alpha\}_i | M_i) d\Omega_{\{\alpha\}_i}$$

$$= \int P(D|\{\alpha\}_i, M_i) \prod_j P(\alpha_{ij} | M_i) d\alpha_{ij}$$

$$= \prod_j \frac{1}{\alpha_{ij}(max) - \alpha_{ij}(min)} \cdot \int P(D|\{\alpha\}_i, M_i) d\alpha_{ij}$$



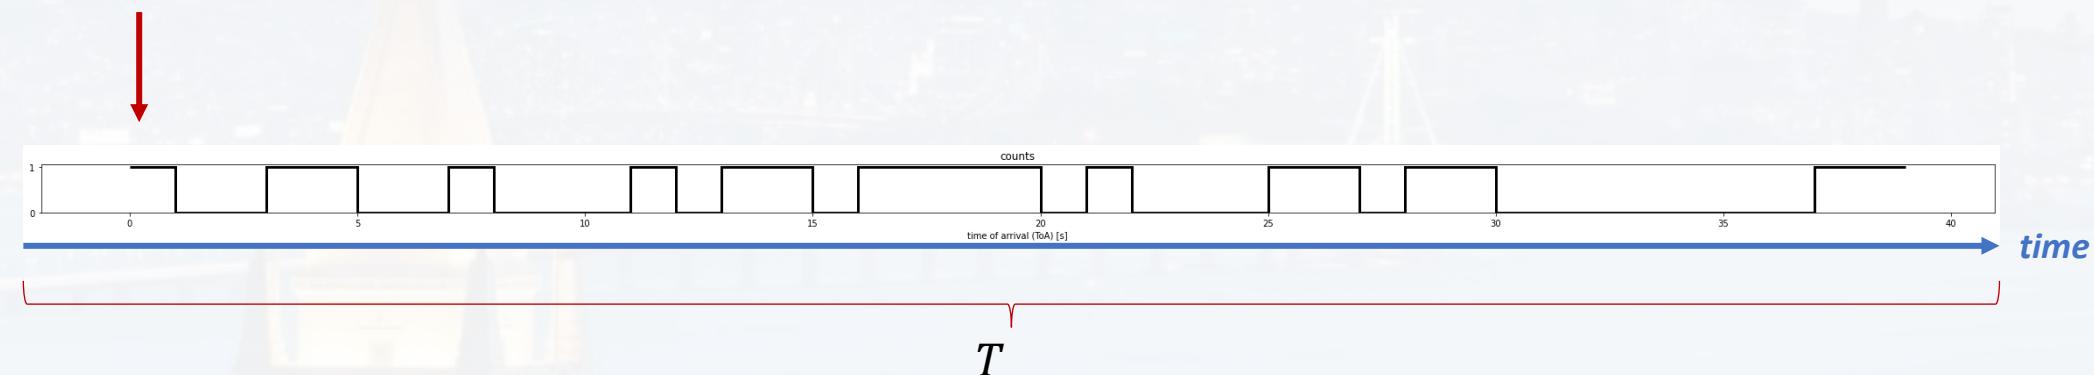
The key part is the likelihood function!

$$\rho = \frac{P(D|M_A) P(M_A)}{P(D|M_B) P(M_B)}$$

$$P(D|M_i) = \int P(D|\{\alpha\}_i M_i) P(\{\alpha\}_i | M_i) d\Omega_{\{\alpha\}_i}$$

$$P(n|r(t)) = \frac{(r(t) \cdot \Delta t)^n}{n!} e^{-r(t) \cdot \Delta t}$$

now: Poisson distribution, see also max. ent. distributions



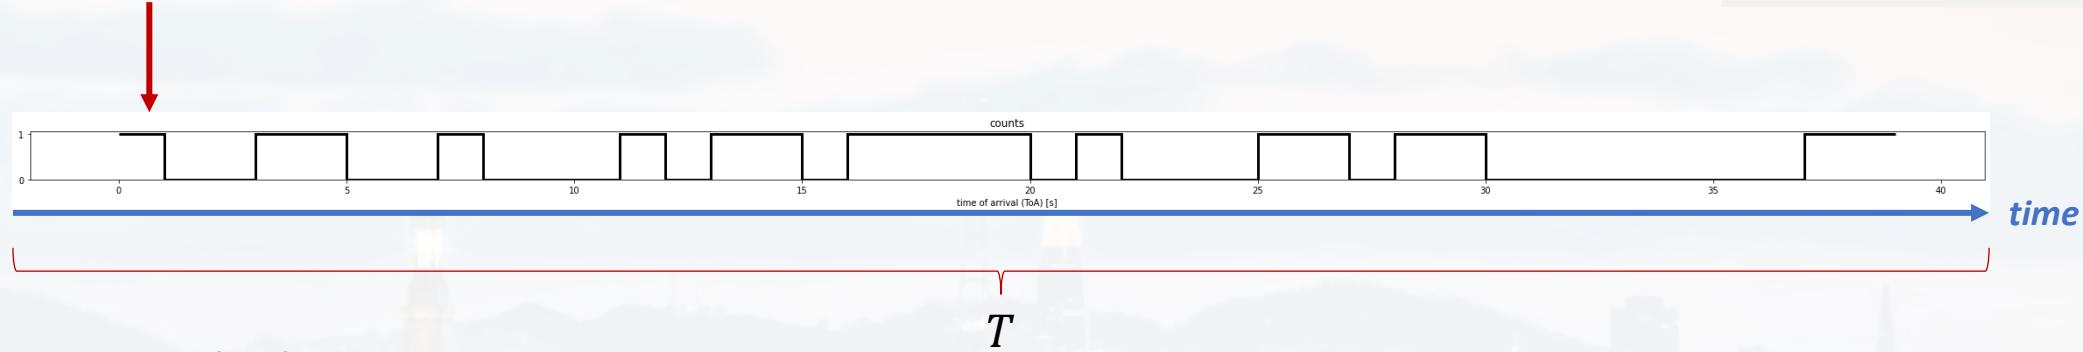
- M_A : constant model (no signal, just noise) $\rightarrow r(t) = \text{const}$
 M_B : signal of unknown phase, amplitude & frequency $\rightarrow r(t) = f(t)$



$$P(n|r(t)) = \frac{(r(t) \cdot \Delta t)^n}{n!} e^{-r(t) \cdot \Delta t}$$

Poisson distribution

$r(t)$:	rate
Δt :	time resolution
n :	number of events
T :	obs. time span
D :	data set



N

intervals with $n = 1$

Q

intervals with $n = 0$

$$(N + Q)\Delta t = T$$

$$P(D| r(t), t) = \prod_{i=1}^N r(t_i) \cdot \Delta t e^{-r(t_i) \cdot \Delta t} \cdot \prod_{i=1}^Q e^{-r(t_i) \cdot \Delta t}$$

$$P(D|M_i) = \int P(D|\{\alpha\}_i M_i) P(\{\alpha\}_i | M_i) d\Omega_{\{\alpha\}_i}$$

$$P(D| r(t), t) = (\Delta t)^N \cdot \prod_{i=1}^N r(t_i) \cdot \exp \left[- \sum_{i=1}^{Q+N} r(t_i) \cdot \Delta t \right]$$



$$P(n|r(t)) = \frac{(r(t) \cdot \Delta t)^n}{n!} e^{-r(t) \cdot \Delta t}$$

Poisson distribution

$r(t)$:	rate
Δt :	time resolution
n :	number of events
T :	obs. time span
D :	data set

N intervals with $n = 1$
 Q intervals with $n = 0$

$$(N + Q)\Delta t = T$$

$$\begin{aligned} P(D | r(t), t) &= (\Delta t)^N \cdot \prod_{i=1}^N r(t_i) \cdot \exp \left[-\sum_{i=1}^{Q+N} r(t_i) \cdot \Delta t \right] \\ &\quad \underbrace{\qquad\qquad\qquad}_{= \int_0^T r(t) dt} \end{aligned}$$

$$P(D | r(t), t) = (\Delta t)^N \cdot \prod_{i=1}^N r(t_i) \cdot \exp \left(- \int_0^T r(t) dt \right)$$



$$P(D | r(t), t) = (\Delta t)^N \cdot \prod_{i=1}^N r(t_i) \cdot \exp\left(-\int_0^T r(t) dt\right)$$

$r(t)$:	rate
Δt :	time resolution
n :	number of events
T :	obs. time span
D :	data set

m phase bins:

$$r_j$$

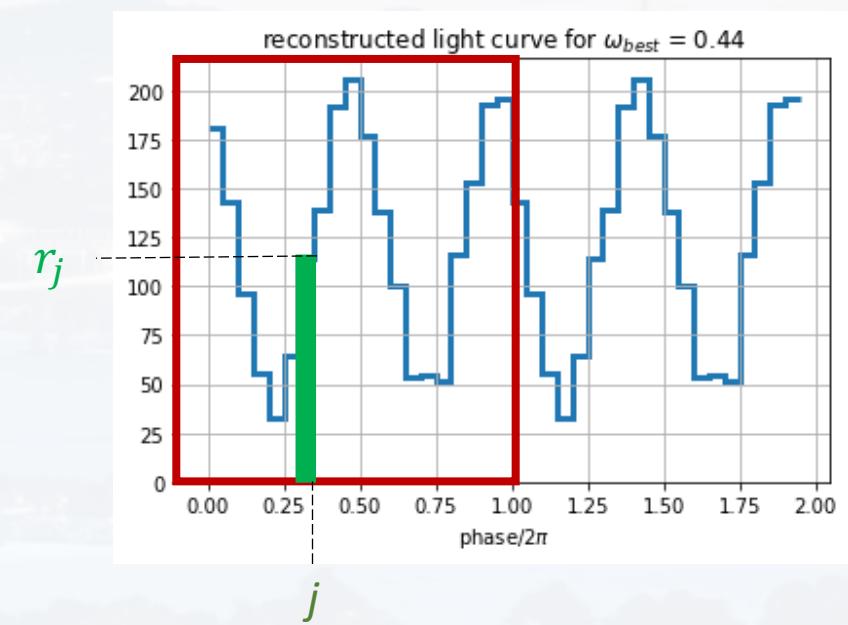
rate in each phase bin j

$$A = \frac{1}{m} \sum_{j=1}^m r_j$$

average rate

$$f_j = \frac{r_j}{\sum_{j=1}^m r_j} = \frac{r_j}{A m}$$

fraction of total rate in
each phase bin j



Each light curve of any shape is being fully described by f_j



$$P(D | r(t), t) = (\Delta t)^N \cdot \prod_{i=1}^N r(t_i) \cdot \exp\left(-\int_0^T r(t) dt\right)$$

constant model: $r_j = \text{const } \forall j$

$$\rightarrow r_j = A$$

$$P(D | r(t), t) = (\Delta t)^N \cdot A^N \cdot e^{-AT}$$

actual signal

- amplitude
- phase
- frequency
- offset

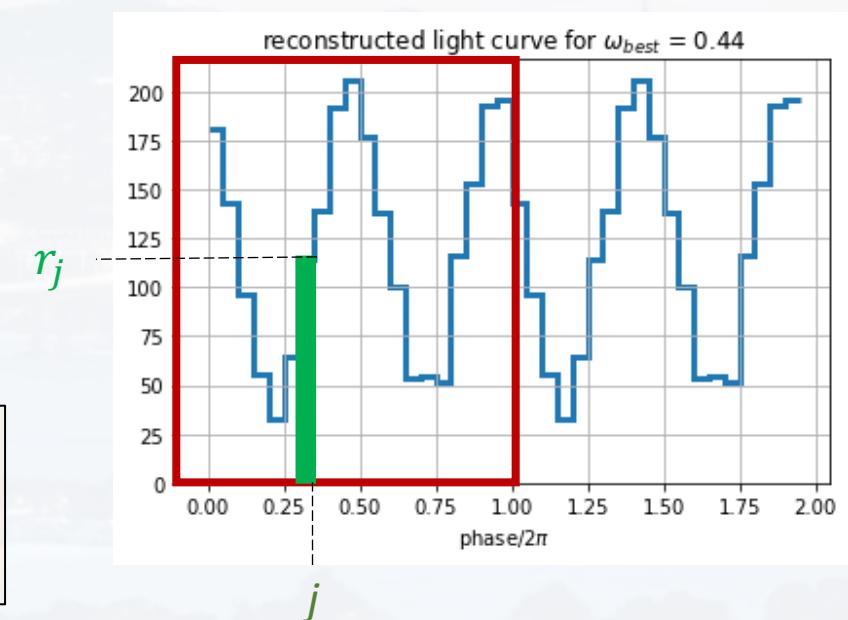
detection

$$- t_i$$

analysis

- phase
- frequency
- $f_j(A, m)$

$r(t)$:	rate
Δt :	time resolution
n :	number of events
T :	obs. time span
D :	data set
N :	number of intervals with $n=1$
r_j	rate in each phase bin j
A :	average rate
m :	number of phase bins
f_j :	fraction of total rate in j





$$M_A \text{ (constant): } P(D | r(t), t) = (\Delta t)^N \cdot A^N \cdot e^{AT}$$

$$M_B \text{ (signal): } P(D | r(t), t) = (\Delta t)^N \cdot \prod_{i=1}^N r(t_i) \cdot \exp\left(-\int_0^T r(t) dt\right)$$

$$P(D|M_i) = \int P(D|\{\alpha\}_i M_i) P(\{\alpha\}_i | M_i) d\Omega_{\{\alpha\}_i}$$

$$\rightarrow P(\omega, \varphi, A, f | M_i) = P(\omega | M_i) P(\varphi | M_i) P(A | M_i) P(f | M_i)$$

$r(t)$:	rate
Δt :	time resolution
n :	number of events
T :	obs. time span
D :	data set
N :	number of intervals with $n=1$
r_j	rate in each phase bin j
A :	average rate
m :	number of phase bins
f_j :	fraction of total rate in j
ω :	frequency
φ :	phase

max entropy:

$$P(\omega | M_i) = \frac{1}{\omega \ln(\omega_{max}/\omega_{min})}$$

$$\omega_{max} = \frac{2\pi N}{T} \quad \omega_{min} = \frac{2\pi}{T}$$

$$\text{in practice: } \omega_{min} = 10 \frac{2\pi}{T}$$

$$P(\varphi | M_i) = \frac{1}{2\pi}$$

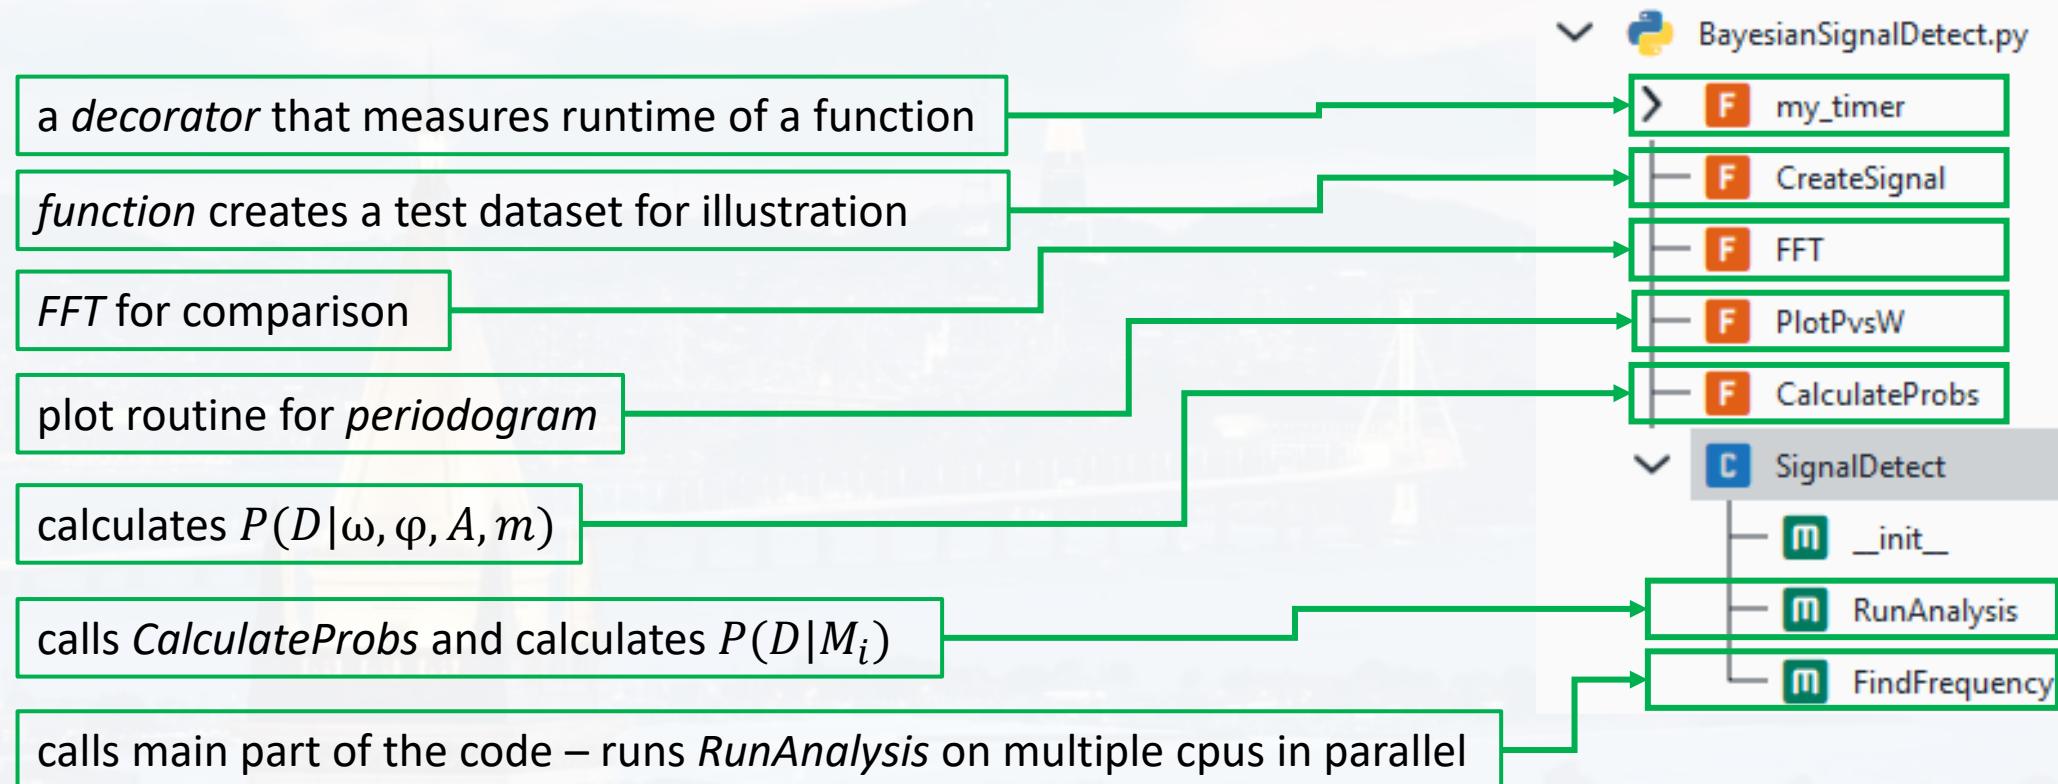
$$P(A | M_i) = \frac{1}{A_{max}}$$

$$P(f | M_i) = (m - 1)! \delta\left(1 - \sum_{j=1}^m f_j\right)$$



you find the python package BayesianSignalDetect.py [here](#)

```
from BayesianSignalDetect import *
```

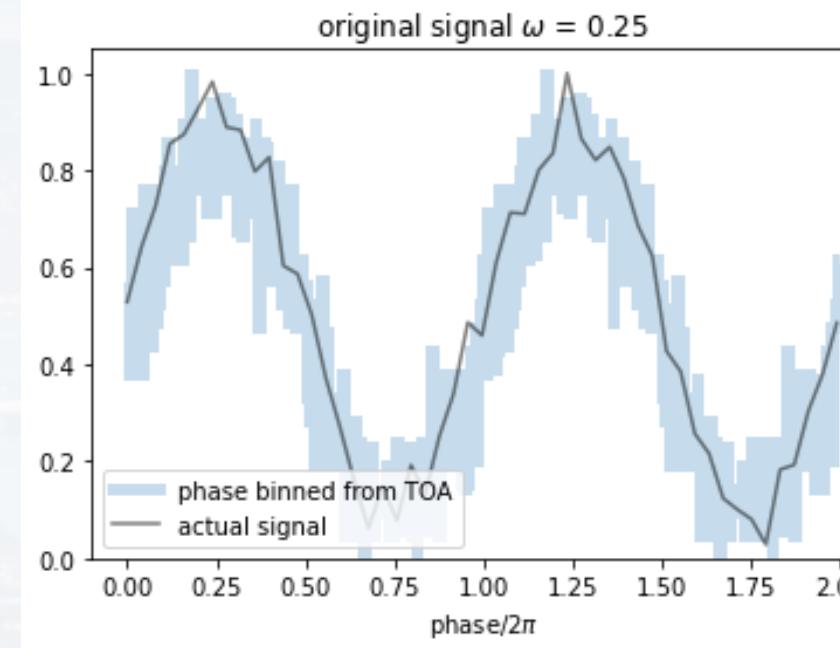




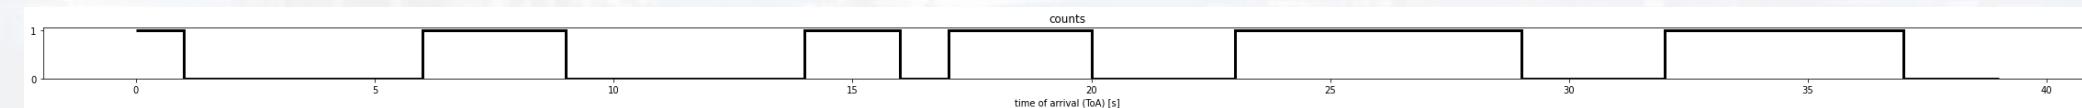
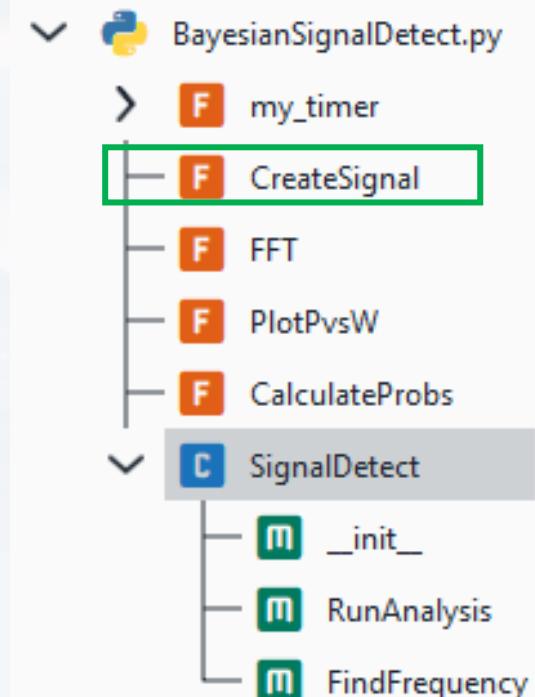
you find the python package BayesianSignalDetect.py [here](#)

```
from BayesianSignalDetect import *
```

```
T = CreateSignal(5000, 0.25, 0.1)
```



N + Q



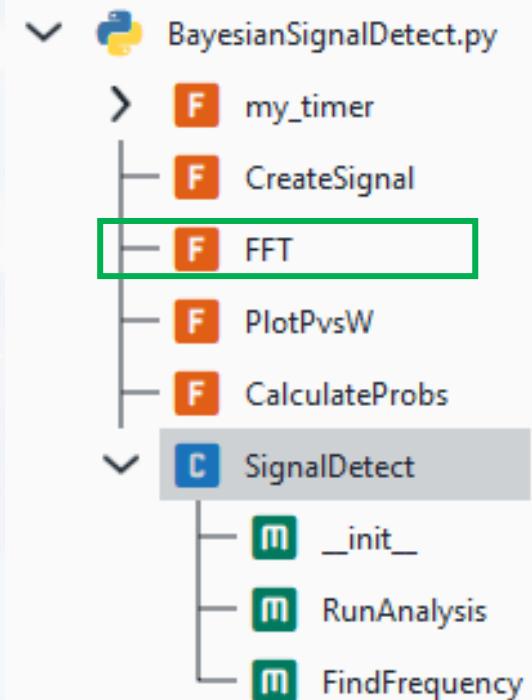
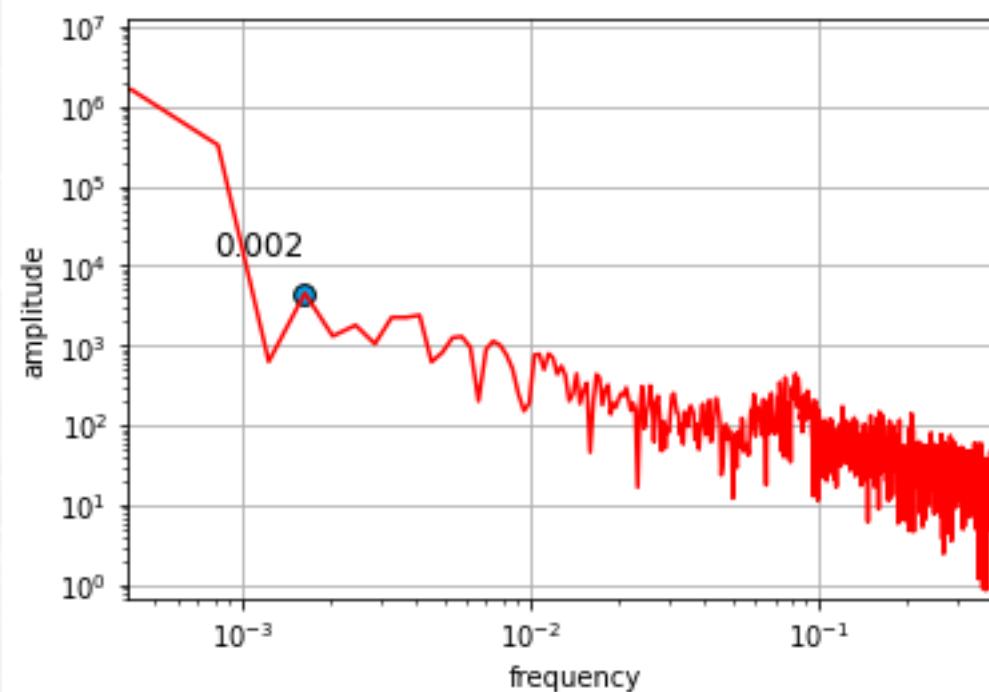


you find the python package BayesianSignalDetect.py [here](#)

```
from BayesianSignalDetect import *
```

```
T = CreateSignal(5000, 0.25, 0.1)
```

```
FFT(T)
```



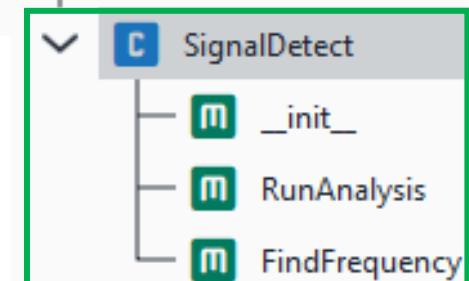
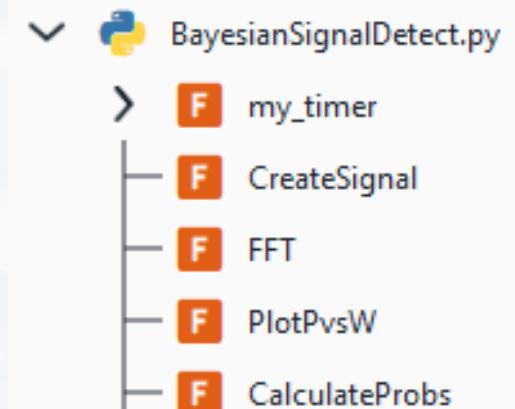
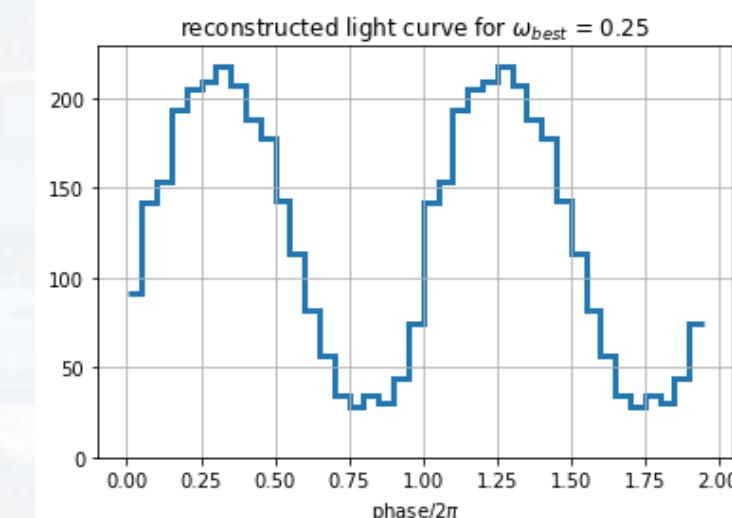
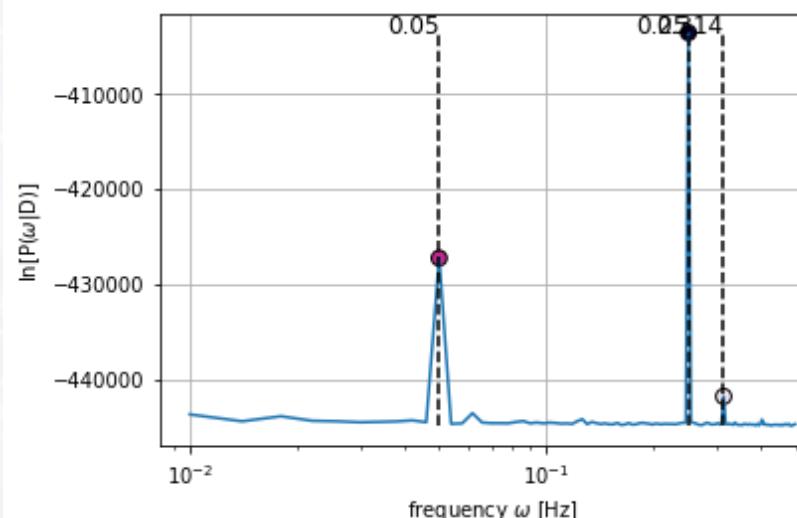


you find the python package BayesianSignalDetect.py [here](#)

```
from BayesianSignalDetect import *
```

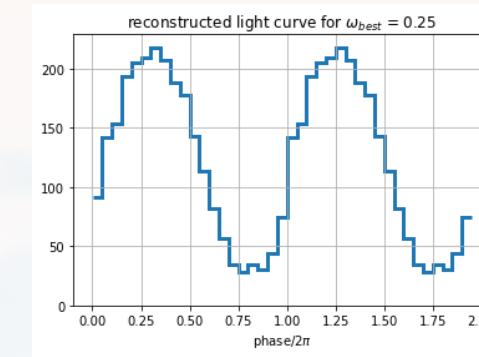
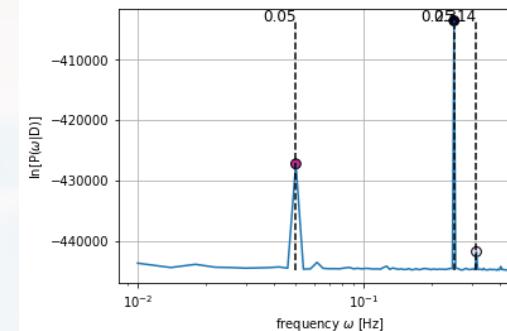
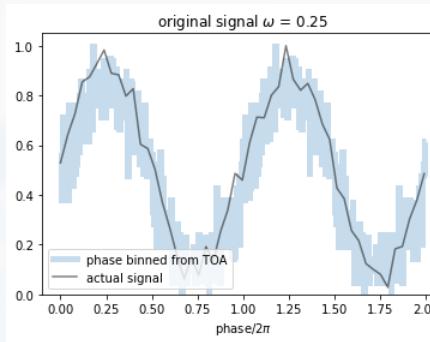
```
T = CreateSignal(5000, 0.25, 0.1)
```

```
S = SignalDetect(T, w_end = 0.5, w_start = 0.01)
[Omega, P] = S.FindFrequency()
```

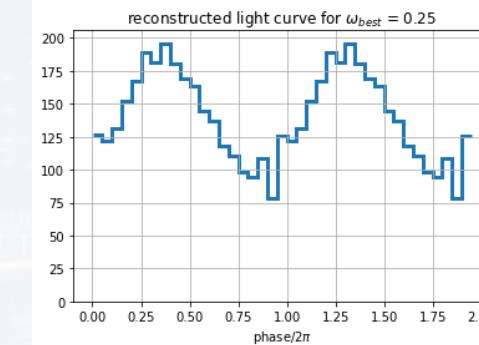
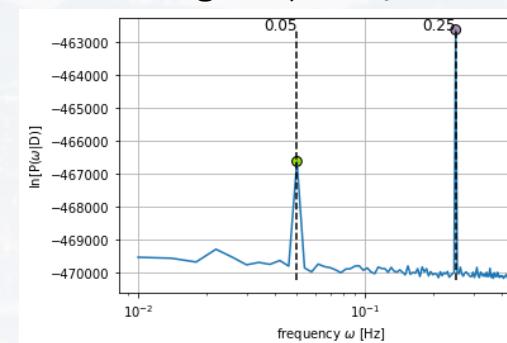
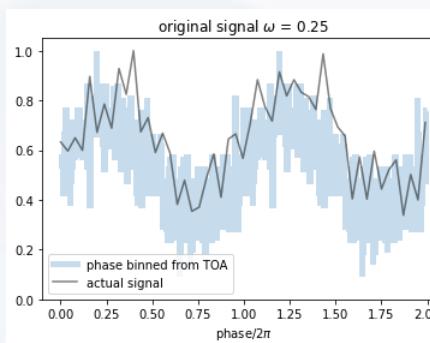




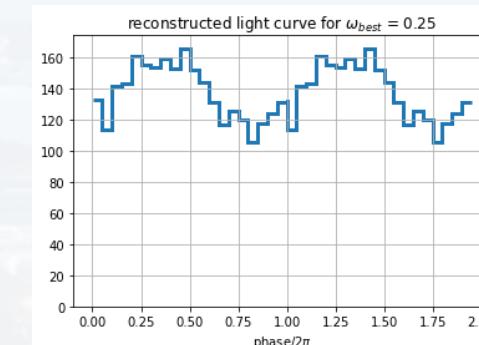
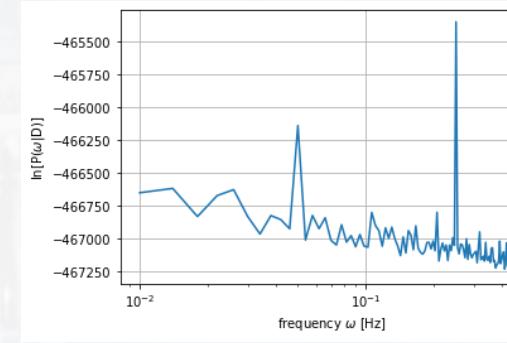
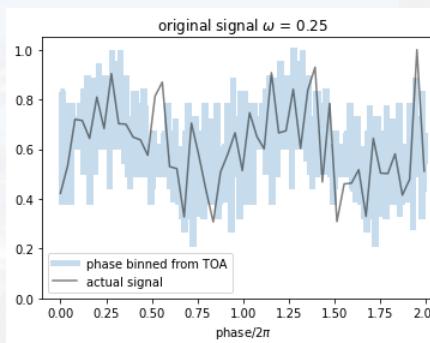
T = CreateSignal(5000, 0.25, 0.1)



T = CreateSignal(5000, 0.25, 0.5)



T = CreateSignal(5000, 0.25, 1)





Thank you for your attention