

## Lecture 10:

### ANN: Perceptron, Backpropagation, SGD



Markus Hohle

University California, Berkeley

Bayesian Data Analysis and  
Machine Learning for Physical  
Sciences



## Course Map

Module 1	Maximum Entropy and Information, Bayes Theorem
Module 2	Naive Bayes, Bayesian Parameter Estimation, MAP
Module 3	MLE, Lin Regression
Module 4	Model selection I: Comparing Distributions
Module 5	Model Selection II: Bayesian Signal Detection
Module 6	Variational Bayes, Expectation Maximization
Module 7	Hidden Markov Models, Stochastic Processes
Module 8	Monte Carlo Methods
Module 9	Machine Learning Overview, Supervised Methods & Unsupervised Methods
<b>Module 10</b>	<b>ANN: Perceptron, Backpropagation, SGD</b>
Module 11	Convolution and Image Classification and Segmentation
Module 12	RNNs and LSTMs
Module 13	RNNs and LSTMs + CNNs
Module 14	Transformer and LLMs
Module 15	Graphs & GNNs

## Outline



### - Gradient Descent

- Vanilla
- Learning Rate Schedule
- L1 and L2
- Momentum
- AdaGrad
- RMSProp

### - Perceptron

### - Backpropagation



## Outline

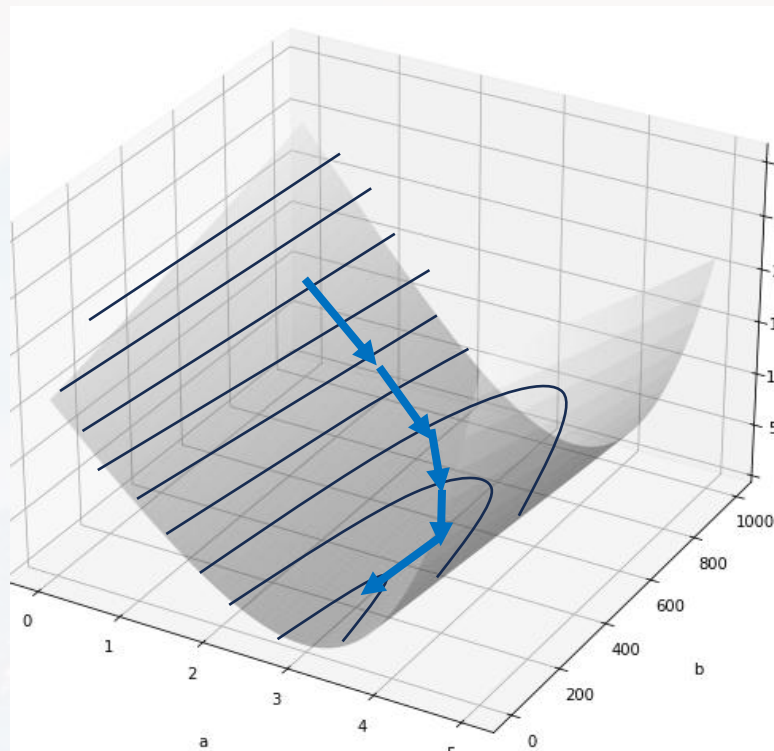
### - Gradient Descent

- Vanilla
- Learning Rate Schedule
- L1 and L2
- Momentum
- AdaGrad
- RMSProp

### - Perceptron

### - Backpropagation





**goal:** find the (global) minimum of a loss function

$$\mathcal{L}, \text{ e. g. } \mathcal{L} = \frac{1}{2} [y - \hat{y}]^2$$

**problem:** no closed form of  $\mathcal{L}$  as a function of the model parameters  $\vec{x}$

**idea:** (stochastic) gradient descent

vanilla: initialize  $\vec{x}_{t=0}$

usually  $\mathcal{N}(0,1)$

$$\vec{x}_{t+1} = \vec{x}_t - \epsilon \nabla_{\vec{x}} \mathcal{L}(\vec{x}_t)$$

problem:  $\epsilon$  too large  $\rightarrow$  leads to oscillations  
 $\epsilon$  too small  $\rightarrow$  takes too long to converge

one idea: decreasing learning rate

$$\epsilon(t) = \frac{\epsilon_{t=0}}{1 + \kappa t}$$

can also be a stepwise function (**learning rate schedule**)

$\epsilon > 0$ :	learning rate
$t$ :	iteration
$\kappa$ :	decay rate
$y$ :	true value
$\hat{y}$ :	model prediction



$$\vec{x}_{t+1} = \vec{x}_t - \epsilon \nabla_{\vec{x}} \mathcal{L}(\vec{x}_t)$$

$\epsilon > 0$ :	learning rate
$t$ :	iteration
$\kappa$ :	decay rate
$y$ :	true value
$\hat{y}$ :	model prediction

### AdaptiveGradient:

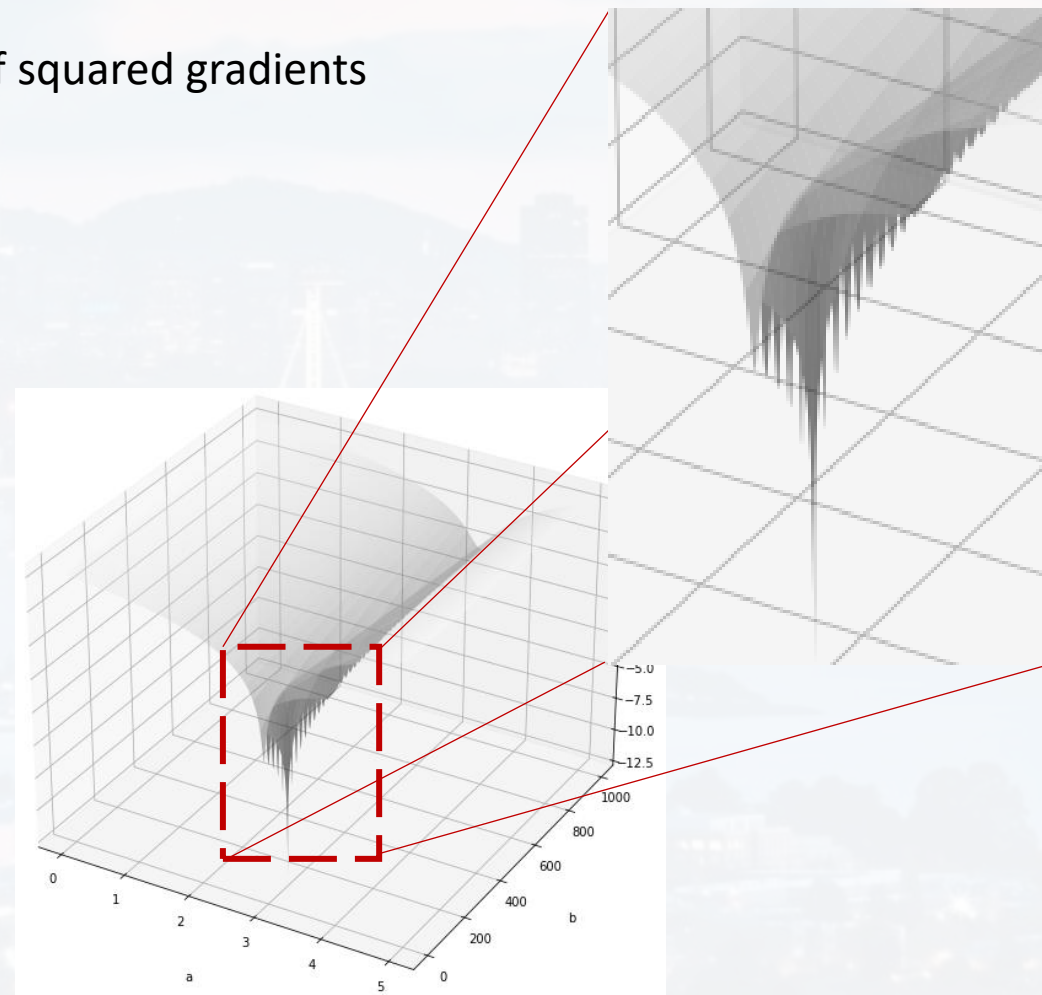
**idea:** reducing  $\epsilon$  over time by **accumulated sum** of squared gradients

$$r_{t+1} = r_t + [\nabla_{\vec{x}} \mathcal{L}(\vec{x}_t)]^2$$

$$\epsilon = \frac{\epsilon_{t=0}}{\sqrt{r_{t+1}} + \delta}$$

$\delta > 0$  for stability

**problem:** can get stuck in a local minimum







$$\vec{x}_{t+1} = \vec{x}_t - \epsilon \nabla_{\vec{x}} \mathcal{L}(\vec{x}_t)$$

$\epsilon > 0$ :	learning rate
$t$ :	iteration
$\kappa$ :	decay rate
$y$ :	true value
$\hat{y}$ :	model prediction
$\mu_{t=0}$ :	momentum

momentum:

taking the **average** of  **$N$**  previous gradients

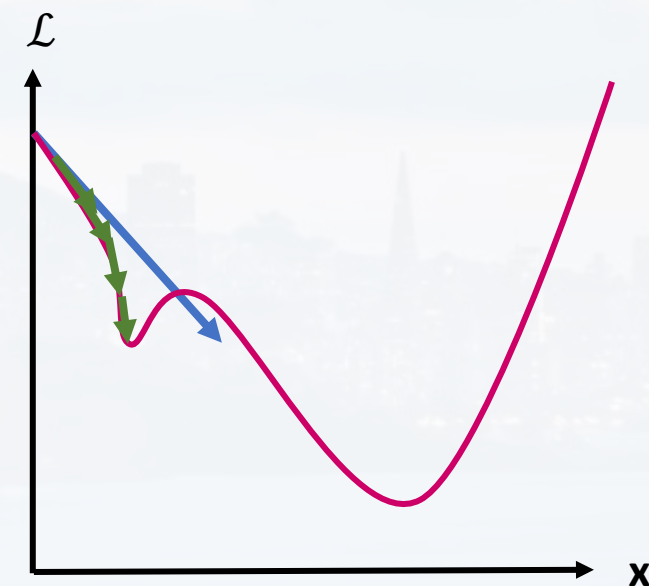
$$\langle \nabla_{\vec{x}} \mathcal{L}(\vec{x}_t) \rangle = \frac{1}{N} [\nabla_{\vec{x}} \mathcal{L}(\vec{x}_{t-1}) + \nabla_{\vec{x}} \mathcal{L}(\vec{x}_{t-2}) + \dots + \nabla_{\vec{x}} \mathcal{L}(\vec{x}_{t-N})]$$

but we want more recent gradients to contribute more than older gradients

→ **weighted average** with weighting factor  $\mu_k$

$$\langle \nabla_{\vec{x}} \mathcal{L}(\vec{x}_t) \rangle = \sum_{k=t-N}^{t-1} \mu_k \cdot \nabla_{\vec{x}} \mathcal{L}(\vec{x}_k)$$

find a clever way to adjust  $\mu_k$  during every iteration  $t$





$$\vec{x}_{t+1} = \vec{x}_t - \epsilon \nabla_{\vec{x}} \mathcal{L}(\vec{x}_t)$$

$\epsilon > 0$ :	learning rate
$t$ :	iteration
$\kappa$ :	decay rate
$y$ :	true value
$\hat{y}$ :	model prediction
$\mu_{t=0}$ :	momentum

momentum:

**weighted average** with weighting factor  $\mu_k$

find a clever way to adjust  $\mu_k$  during every iteration  $t$

$$\langle \nabla_{\vec{x}} \mathcal{L}(\vec{x}_0) \rangle = \nabla_{\vec{x}} \mathcal{L}(\vec{x}_0)$$

$$\mu_0 = (0,1)$$

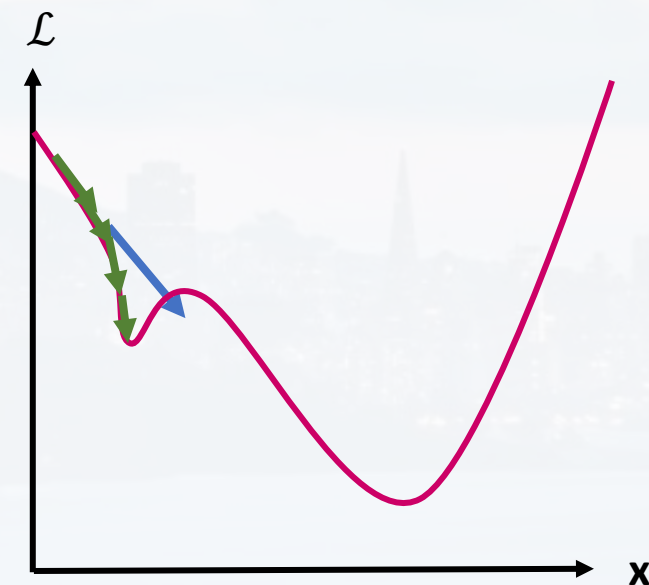
$$\langle \nabla_{\vec{x}} \mathcal{L}(\vec{x}_1) \rangle = \nabla_{\vec{x}} \mathcal{L}(\vec{x}_1) + \mu_0 \cdot \nabla_{\vec{x}} \mathcal{L}(\vec{x}_0)$$

$$\begin{aligned} \langle \nabla_{\vec{x}} \mathcal{L}(\vec{x}_2) \rangle &= \nabla_{\vec{x}} \mathcal{L}(\vec{x}_2) + \boxed{\mu_0} [\nabla_{\vec{x}} \mathcal{L}(\vec{x}_1) \\ &\quad + \boxed{\mu_0} \nabla_{\vec{x}} \mathcal{L}(\vec{x}_0)] \end{aligned}$$

$$\mu_{k=2} = \mu_0 \mu_0 = \mu_0^2$$

$$\langle \nabla_{\vec{x}} \mathcal{L}(\vec{x}_3) \rangle = \nabla_{\vec{x}} \mathcal{L}(\vec{x}_3) + \boxed{\mu_0} \cdot [\nabla_{\vec{x}} \mathcal{L}(\vec{x}_2) + \boxed{\mu_0} \cdot [\nabla_{\vec{x}} \mathcal{L}(\vec{x}_1) + \boxed{\mu_0} \cdot \nabla_{\vec{x}} \mathcal{L}(\vec{x}_0)]]$$

... and so on...







$$\vec{x}_{t+1} = \vec{x}_t - \epsilon \nabla_{\vec{x}} \mathcal{L}(\vec{x}_t)$$

momentum:

**weighted average** with weighting factor  $\mu_k$

find a clever way to adjust  $\mu_k$  during every iteration  $t$        $\mu_0 = (0,1)$

$$\langle \nabla_{\vec{x}} \mathcal{L}(\vec{x}_3) \rangle = \nabla_{\vec{x}} \mathcal{L}(\vec{x}_3) + \mu_0 \cdot [\nabla_{\vec{x}} \mathcal{L}(\vec{x}_2) + \mu_0 \cdot [\nabla_{\vec{x}} \mathcal{L}(\vec{x}_1) + \mu_0 \cdot \nabla_{\vec{x}} \mathcal{L}(\vec{x}_0)]]$$

... and so on...

**class Optimizer:**

```
def __init__(self, learning_rate = 0.1, decay = 0, momentum = 0):  
    self.learning_rate      = learning_rate  
    self.decay              = decay  
    self.current_learning_rate = learning_rate  
    self.iterations         = 0  
    self.momentum           = momentum
```

$\epsilon > 0$ :	learning rate
$t$ :	iteration
$\kappa$ :	decay rate
$y$ :	true value
$\hat{y}$ :	model prediction
$\mu_{t=0}$ :	momentum



$$\vec{x}_{t+1} = \vec{x}_t - \epsilon \nabla_{\vec{x}} \mathcal{L}(\vec{x}_t)$$

regularization: recall: linear regression

$\epsilon > 0$ :	learning rate
$t$ :	iteration
$\lambda$ :	Lagrangian Multiplier
$y$ :	true value
$\hat{y}$ :	model prediction
$\mu_{t=0}$ :	momentum

$$\hat{\beta} = \underset{\beta}{\operatorname{argmin}} \left\{ \frac{1}{N} \|Y - X\beta\|^2 + \lambda \|\beta\|^1 \right\}$$

the Loss Function  
 $L(X, Y, \lambda)$

L1 or **Least absolute shrinkage and selection operator**  
- encourages **sparsity** of  $\beta$   
- reduces **overfitting**

$$\hat{\beta} = \underset{\beta}{\operatorname{argmin}} \left\{ \frac{1}{N} \|Y - X\beta\|^2 + \lambda \|\beta\|^2 \right\}$$

L2 or **Ridge**  
- **penalizes large  $\beta$**

$$\hat{\beta} = \underset{\beta}{\operatorname{argmin}} \left\{ \frac{1}{N} \|Y - X\beta\|^2 + \lambda \max(0, -\beta) \right\} \quad - \text{penalizes negative } \beta$$

...and so on



$$\vec{x}_{t+1} = \vec{x}_t - \epsilon \nabla_{\vec{x}} \mathcal{L}(\vec{x}_t)$$

regularization:

**L1 and L2 regularization**

$$\vec{x}_{t+1} = \vec{x}_t - \epsilon \nabla_{\vec{x}} [\mathcal{L}(\vec{x}_t) + \lambda_1 \|\vec{x}_t\|^1 + \lambda_2 \|\vec{x}_t\|^2]$$

$$\vec{x}_{t+1} = \vec{x}_t - \epsilon \nabla_{\vec{x}} \mathcal{L}(\vec{x}_t) - \epsilon \lambda_1 \nabla_{\vec{x}} \|\vec{x}_t\|^1 - \epsilon \lambda_2 \nabla_{\vec{x}} \|\vec{x}_t\|^2$$

$\epsilon > 0$ :	learning rate
$t$ :	iteration
$\lambda$ :	Lagrangian Multiplier
$y$ :	true value
$\hat{y}$ :	model prediction
$\mu_{t=0}$ :	momentum





$$\vec{x}_{t+1} = \vec{x}_t - \epsilon \nabla_{\vec{x}} \mathcal{L}(\vec{x}_t)$$

regularization:

**L1 and L2 regularization**

$$\vec{x}_{t+1} = \vec{x}_t - \epsilon \nabla_{\vec{x}} [\mathcal{L}(\vec{x}_t) + \lambda_1 \|\vec{x}_t\|^1 + \lambda_2 \|\vec{x}_t\|^2]$$

$$\vec{x}_{t+1} = \vec{x}_t - \epsilon \nabla_{\vec{x}} \mathcal{L}(\vec{x}_t) - \epsilon \lambda_1 \nabla_{\vec{x}} \|\vec{x}_t\|^1 - \epsilon \lambda_2 \nabla_{\vec{x}} \|\vec{x}_t\|^2$$

$\epsilon > 0$ :	learning rate
$t$ :	iteration
$\lambda$ :	Lagrangian Multiplier
$y$ :	true value
$\hat{y}$ :	model prediction
$\mu_{t=0}$ :	momentum

- notes:**
- vanilla gradient descent does not stop if values for  $\vec{x}$  are too large
  - the derivative of  $\|x\|^1$  returns the sign (i. e. direction) and therefore encourages sparsity (**reduces overfitting**)
  - usually  $\lambda \ll \|\vec{x}_t\|^n$
  - will be important for ANNs later



$$\vec{x}_{t+1} = \vec{x}_t - \epsilon \nabla_{\vec{x}} \mathcal{L}(\vec{x}_t)$$

$\epsilon > 0$ :	learning rate
$t$ :	iteration
$\lambda$ :	Lagrangian Multiplier
$y$ :	true value
$\hat{y}$ :	model prediction
$\mu_{t=0}$ :	momentum

### Root Mean Square Propagation:

**problem:** Adagrad accumulates all previous gradients  $\rightarrow$  slows down updates too fast

$$r_{t+1} = r_t + [\nabla_{\vec{x}} \mathcal{L}(\vec{x}_t)]^2 \quad \epsilon = \frac{\epsilon_{t=0}}{\sqrt{r_{t+1}} + \delta} \quad \delta > 0 \text{ for stability}$$

**idea:** weighted average like for **momentum**

$$r_{t+1} = \beta r_t + (1 - \beta)[\nabla_{\vec{x}} \mathcal{L}(\vec{x}_t)]^2 \quad \beta = (0,1)$$

$$\langle \nabla_{\vec{x}} \mathcal{L}(\vec{x}_1) \rangle = \nabla_{\vec{x}} \mathcal{L}(\vec{x}_1) + \mu_0 \cdot \nabla_{\vec{x}} \mathcal{L}(\vec{x}_0) \quad \text{momentum}$$

$$\epsilon \rightarrow \frac{\epsilon}{\sqrt{r_{t+1}} + \delta} \quad \text{adaptive gradient, aka AdaGrad}$$

Root Mean Square Propagation **RMSProp**

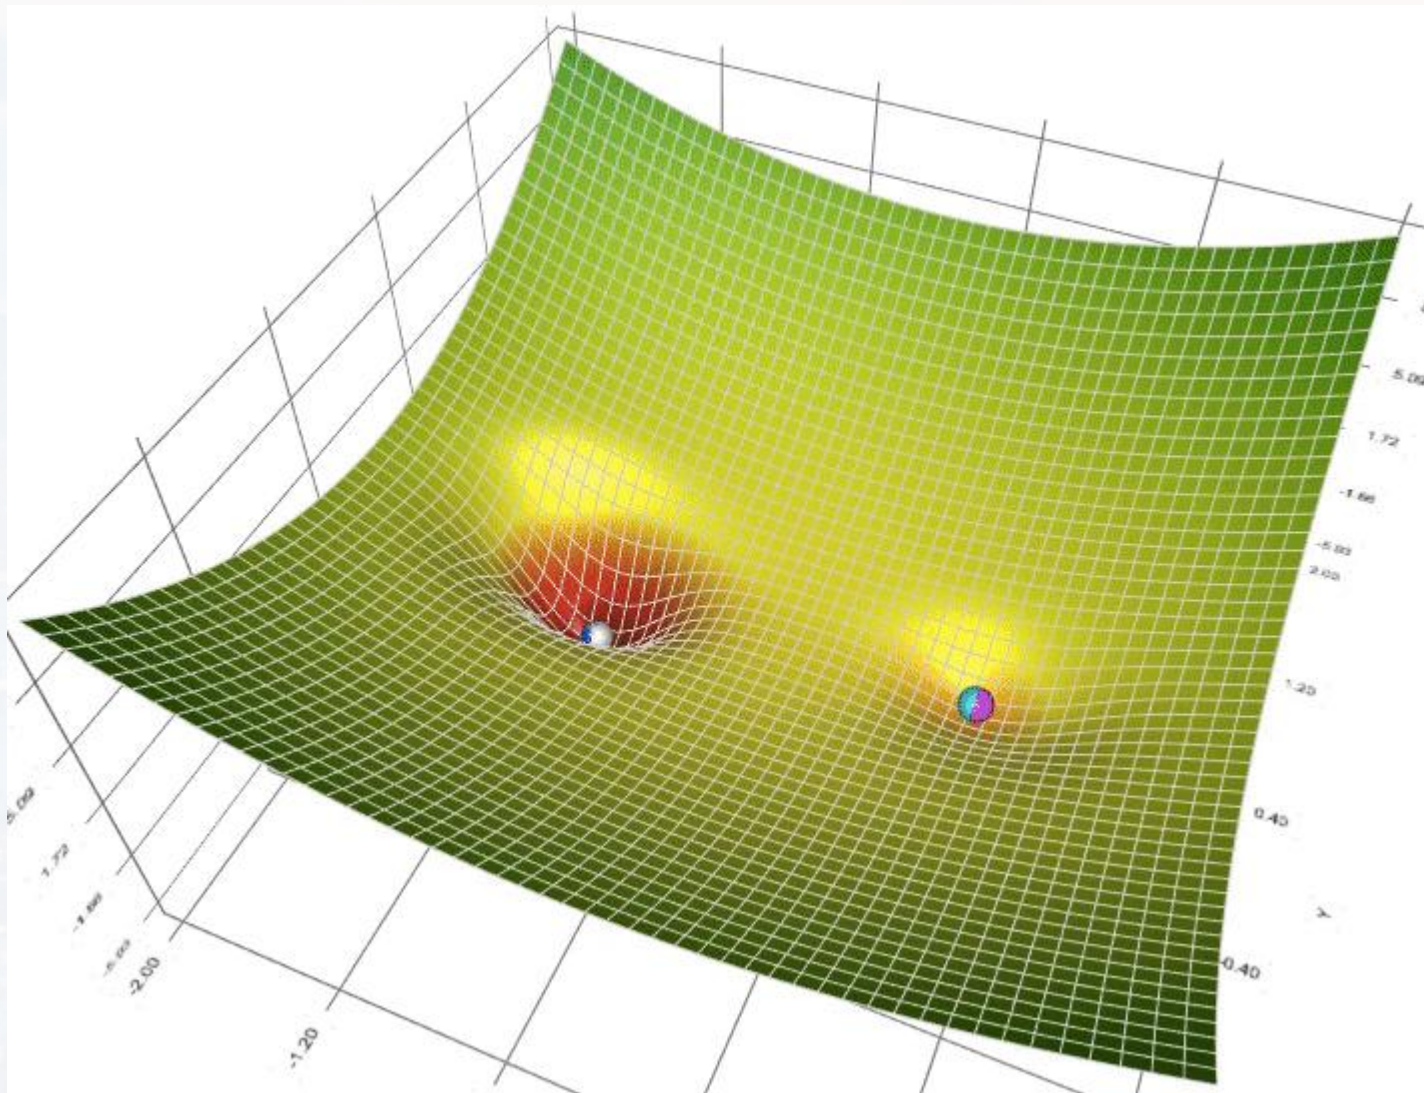
all combined:  
**Adaptive Moment Estimation**  
aka **Adam**





Lili Jiang

[TowardsDataScience](#)



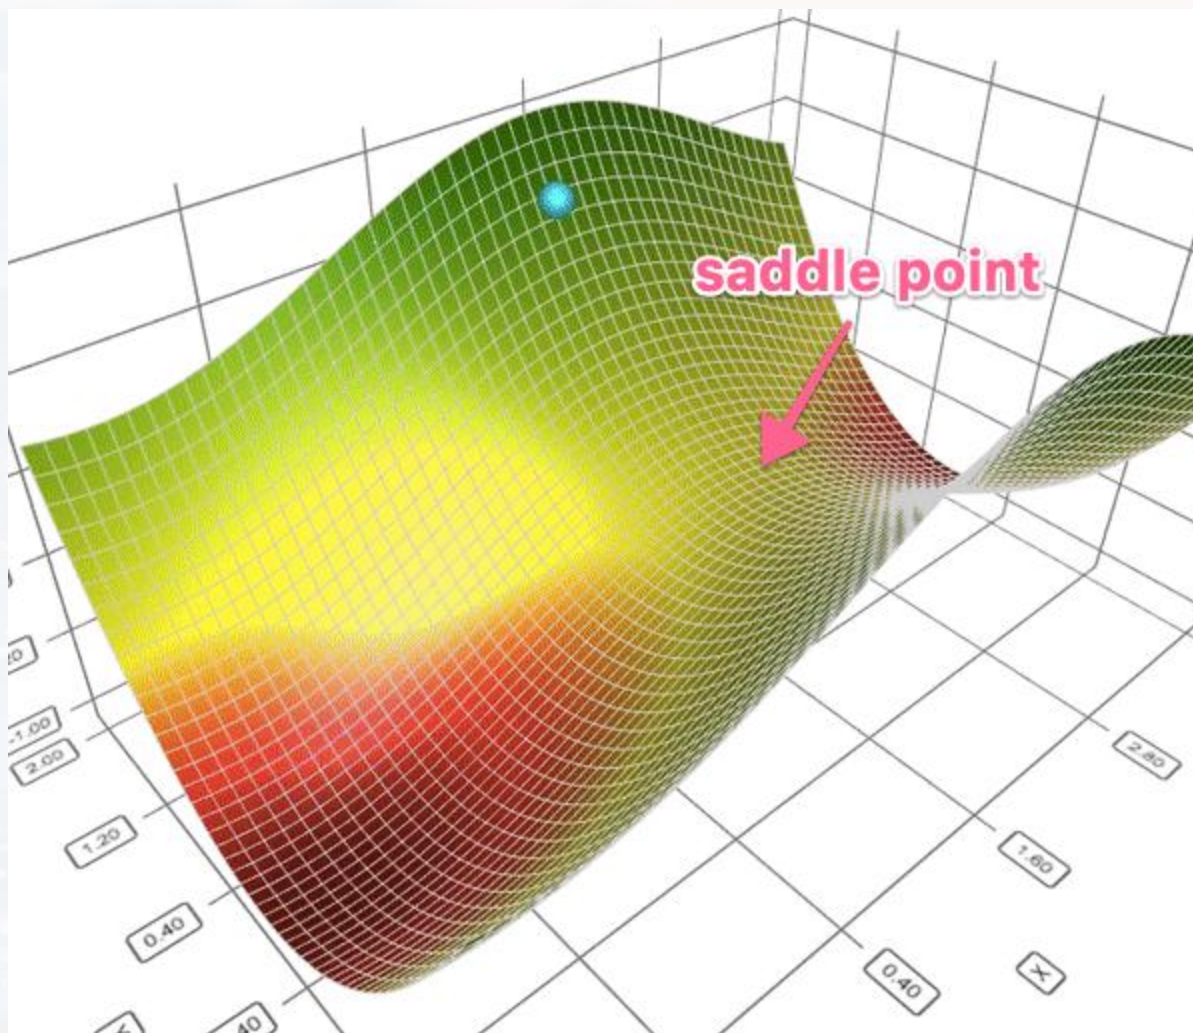
gradient descent (cyan),  
momentum (magenta),  
AdaGrad (white),  
RMSProp (green),  
Adam (blue)





Lili Jiang

[TowardsDataScience](#)



gradient descent (cyan),  
momentum (magenta),  
AdaGrad (white),  
RMSProp (green),  
Adam (blue)

see also  
`WalkThroughGradDescent.ipynb`

## Outline



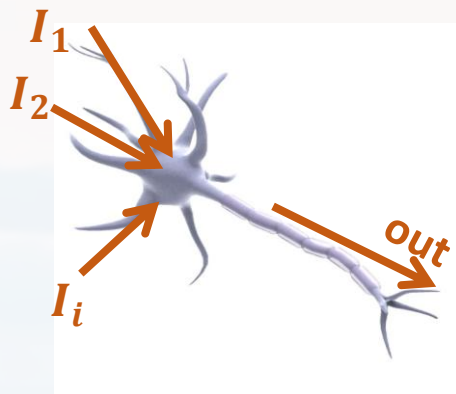
### - Gradient Descent

- Vanilla
- Learning Rate Schedule
- L1 and L2
- Momentum
- AdaGrad
- RMSProp

### - Perceptron

### - Backpropagation





$$net = \sum_i I_i \cdot w_i + b$$

weights and bias are  
“learnable”

$b$ :	bias (base potential
$I_i$ :	input $i$
$w_i$ :	corresponding weight

recall: linear models

$$y_k = \beta_0 + \sum_{n=1}^N \beta_n x_n + \epsilon$$

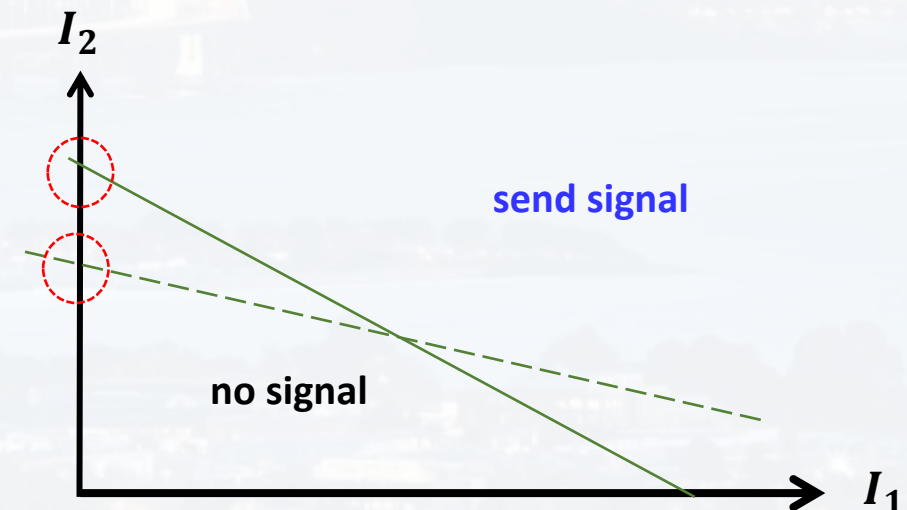
$$\begin{pmatrix} y_1 \\ \vdots \\ y_k \\ \vdots \\ y_K \end{pmatrix} = \begin{pmatrix} 1 & x_{11} & x_{12} & \dots & x_{1n} & \dots & x_{1N} \\ \vdots & \vdots & \vdots & & \vdots & & \vdots \\ 1 & x_{k1} & & & x_{kn} & & \\ \vdots & \vdots & & & \vdots & & \vdots \\ 1 & \dots & & & \dots & & \\ 1 & x_{K1} & x_{K2} & \dots & x_{Kn} & \dots & x_{KN} \end{pmatrix} \begin{pmatrix} \beta_0 \\ \beta_1 \\ \vdots \\ \beta_n \\ \vdots \\ \beta_N \end{pmatrix} + \begin{pmatrix} \epsilon_1 \\ \epsilon_2 \\ \vdots \\ \epsilon_k \\ \vdots \\ \epsilon_K \end{pmatrix}$$

simple example:

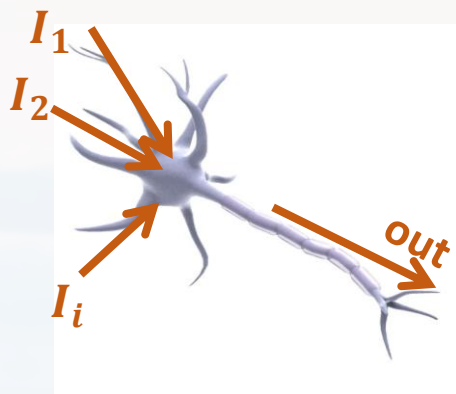
**one** neuron with a **switch**, threshold **T** and **two** input channels

fire if:  $b + I_1 w_1 + I_2 w_2 > T$

$$I_2 = -\frac{w_1}{w_2} I_1 + \frac{T - b}{w_2}$$



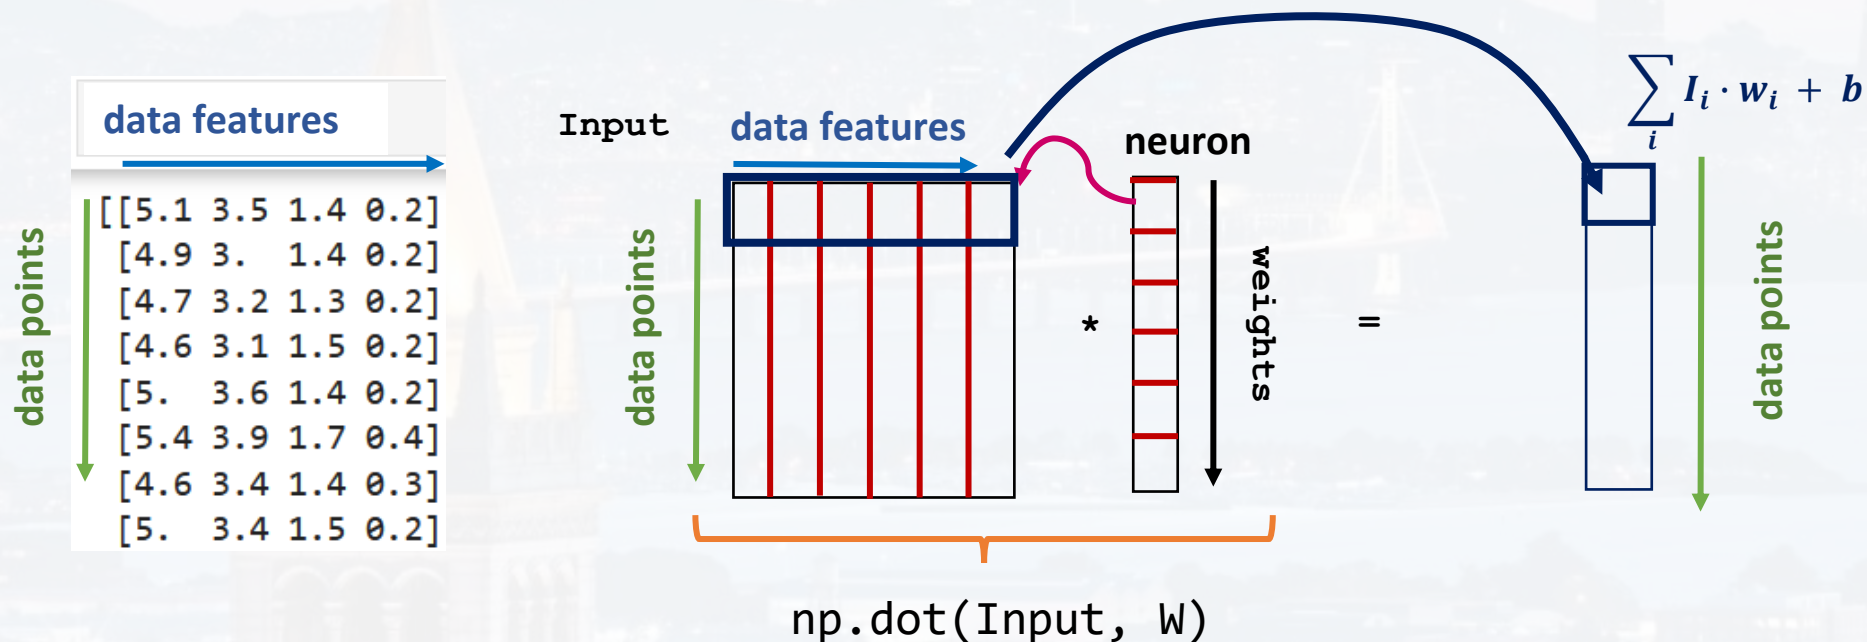


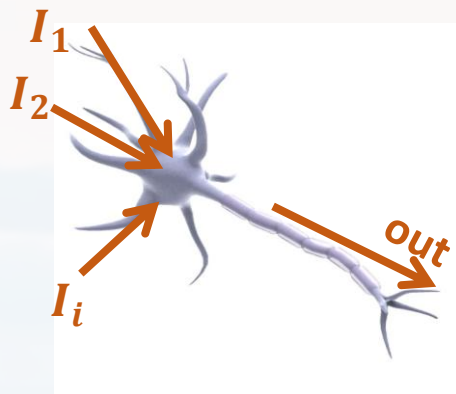


$$net = \sum_i I_i \cdot w_i + b$$

weights and bias are  
“learnable”

$b$ : bias (base potential)  
 $I_i$ : input  $i$   
 $w_i$ : corresponding weight





$$net = \sum_i I_i \cdot w_i + b$$

weights and bias are  
“**learnable**”

$b$ :	bias (base potential)
$I_i$ :	input $i$
$w_i$ :	corresponding weight

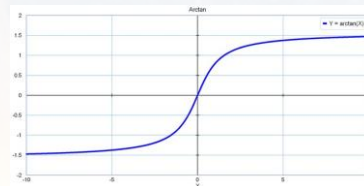
$net$  →

some activation function  $f$

→

actual output

-  $y = \arctan(net)$



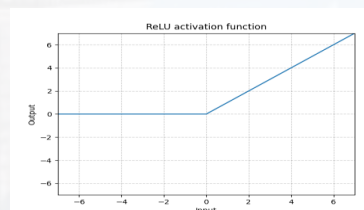
$(-\infty; +\infty) \rightarrow (-\pi/2; +\pi/2)$

-  $y = \text{sigm}(net)$



$(-\infty; +\infty) \rightarrow (0; 1)$

-  $y = \text{ReLU}(net)$   
 $= \max(0, net)$



$(-\infty; +\infty) \rightarrow (0; +\infty)$

...and many other

## Outline



### - Gradient Descent

- Vanilla
- Learning Rate Schedule
- L1 and L2
- Momentum
- AdaGrad
- RMSProp

### - Perceptron

### - **Backpropagation**





learning:

$$net = \sum_i I_i \cdot w_i + b \xrightarrow{\text{activation function } f} y = f(net) \xrightarrow{\text{loss function } \mathcal{L}} \mathcal{L} = \frac{1}{2} (t - y)^2$$

target output  $t$

$b$ : bias (base potential)  
 $I_i$ : input  $i$   
 $w_i$ : corresponding weight

finding best  $w_i$  by **minimizing**  $\mathcal{L}$

$$\Delta w_i = w_i(new) - w_i(old) = -\epsilon \frac{\partial \mathcal{L}}{\partial w_i}$$

gradient descent with learning rate  $\epsilon$

$$\frac{\partial \mathcal{L}}{\partial w_i} = \frac{\partial \mathcal{L}}{\partial y} \frac{\partial y}{\partial net} \frac{\partial net}{\partial w_i} = \boxed{-(t - y)} \boxed{f'(net)} \boxed{I_i}$$

outer derivatives      inner derivative

The required change of  $\mathcal{L}$  propagates back to the changes of  $w_i$  and  $b \rightarrow$  **backpropagation**

$$\Delta w_i = w_i(new) - w_i(old) = \epsilon (t - y) f'(net) I_i$$

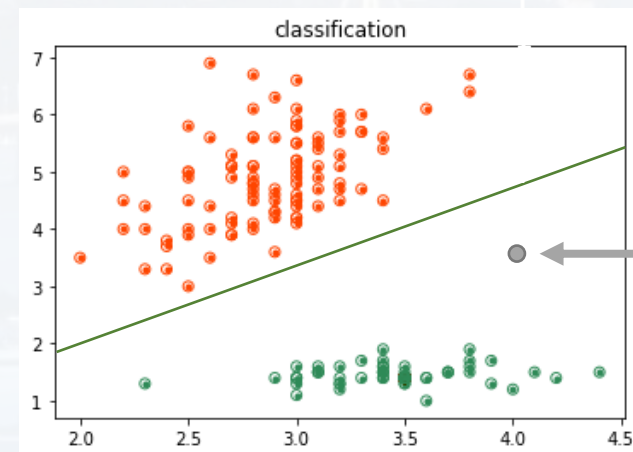
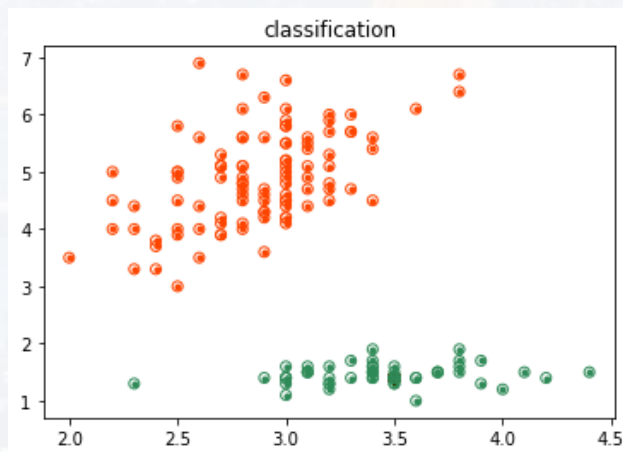
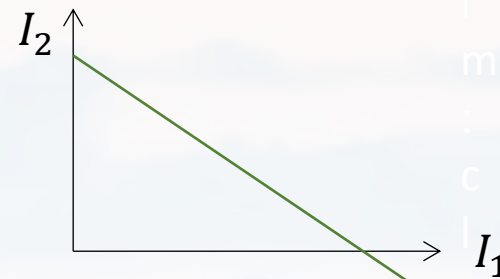
$$\Delta b = b(new) - b(old) = \epsilon (t - y) f'(net) \cdot 1$$

see Perceptron.ipynb



see `Perceptron.ipynb`

**one** neuron  
**two** features  
**two** classes  
**N** data points

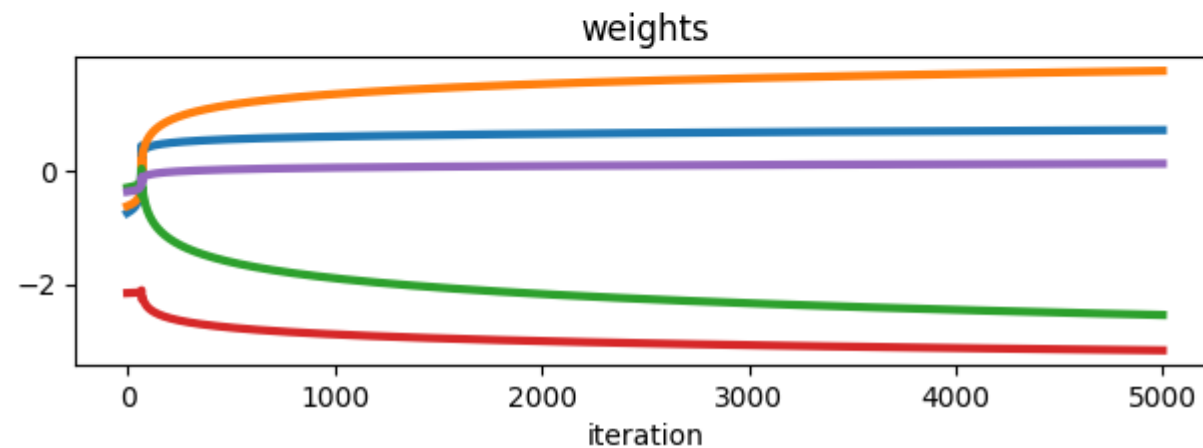
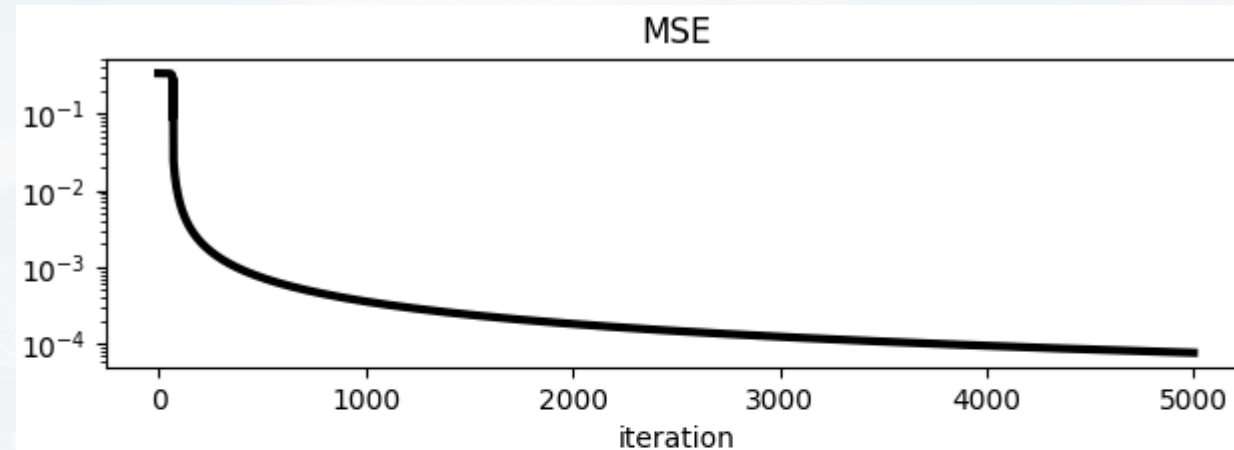
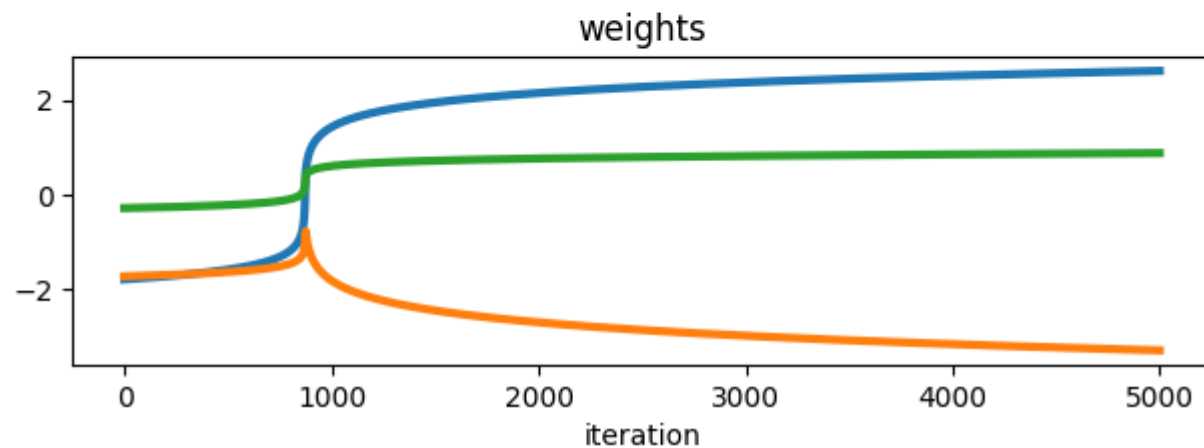
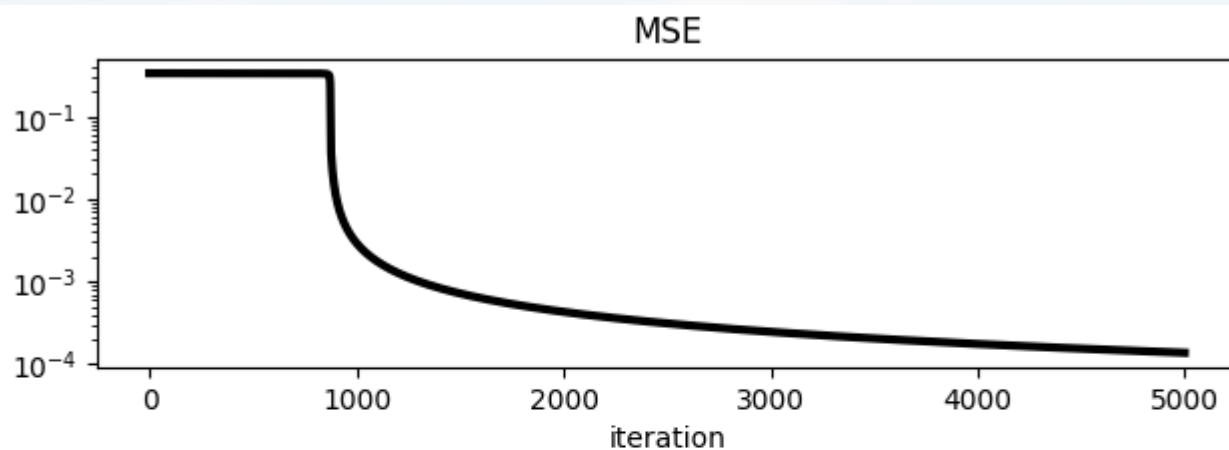


classifying a new  
data point



see `Perceptron.ipynb`

the training process

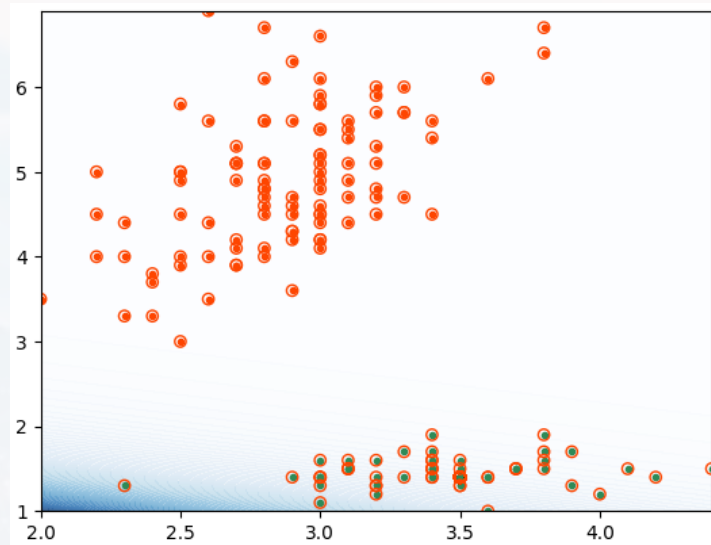




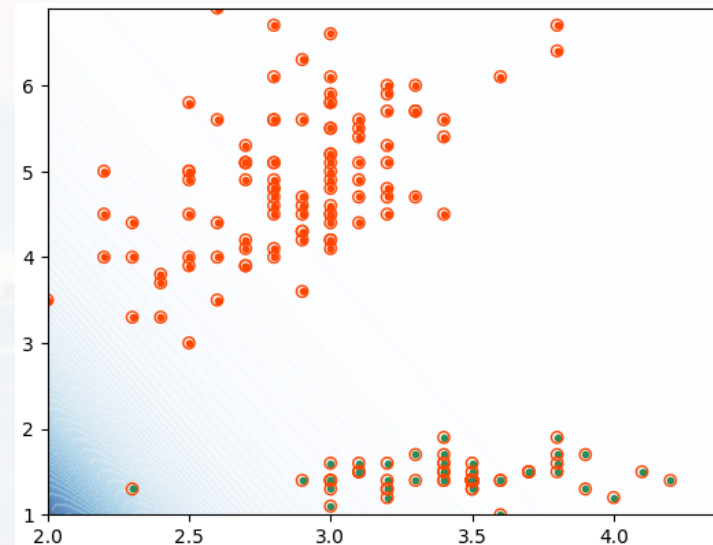


see `Perceptron.ipynb`

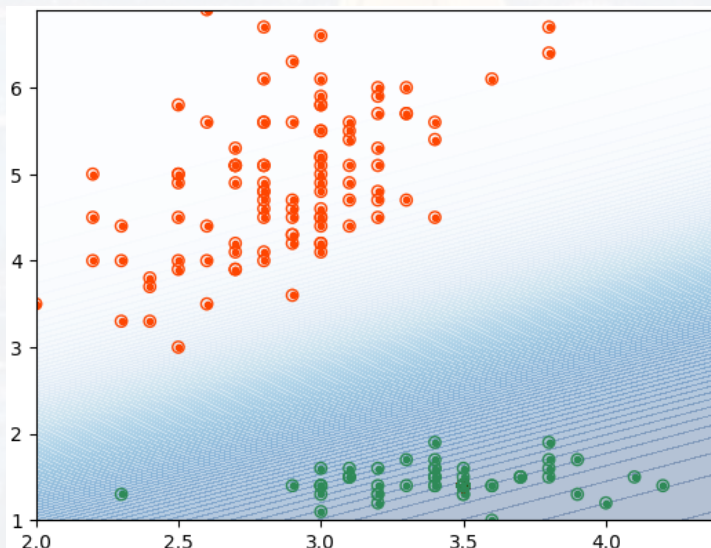
the training process



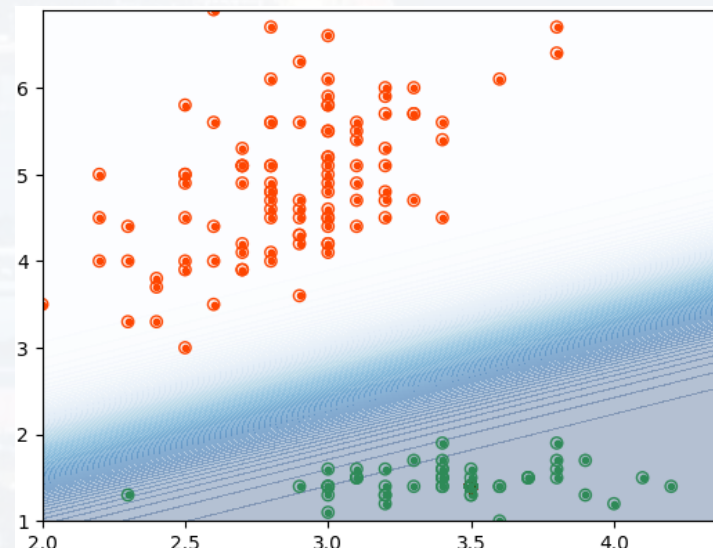
$t = 3$



$t = 5$



$t = 100$

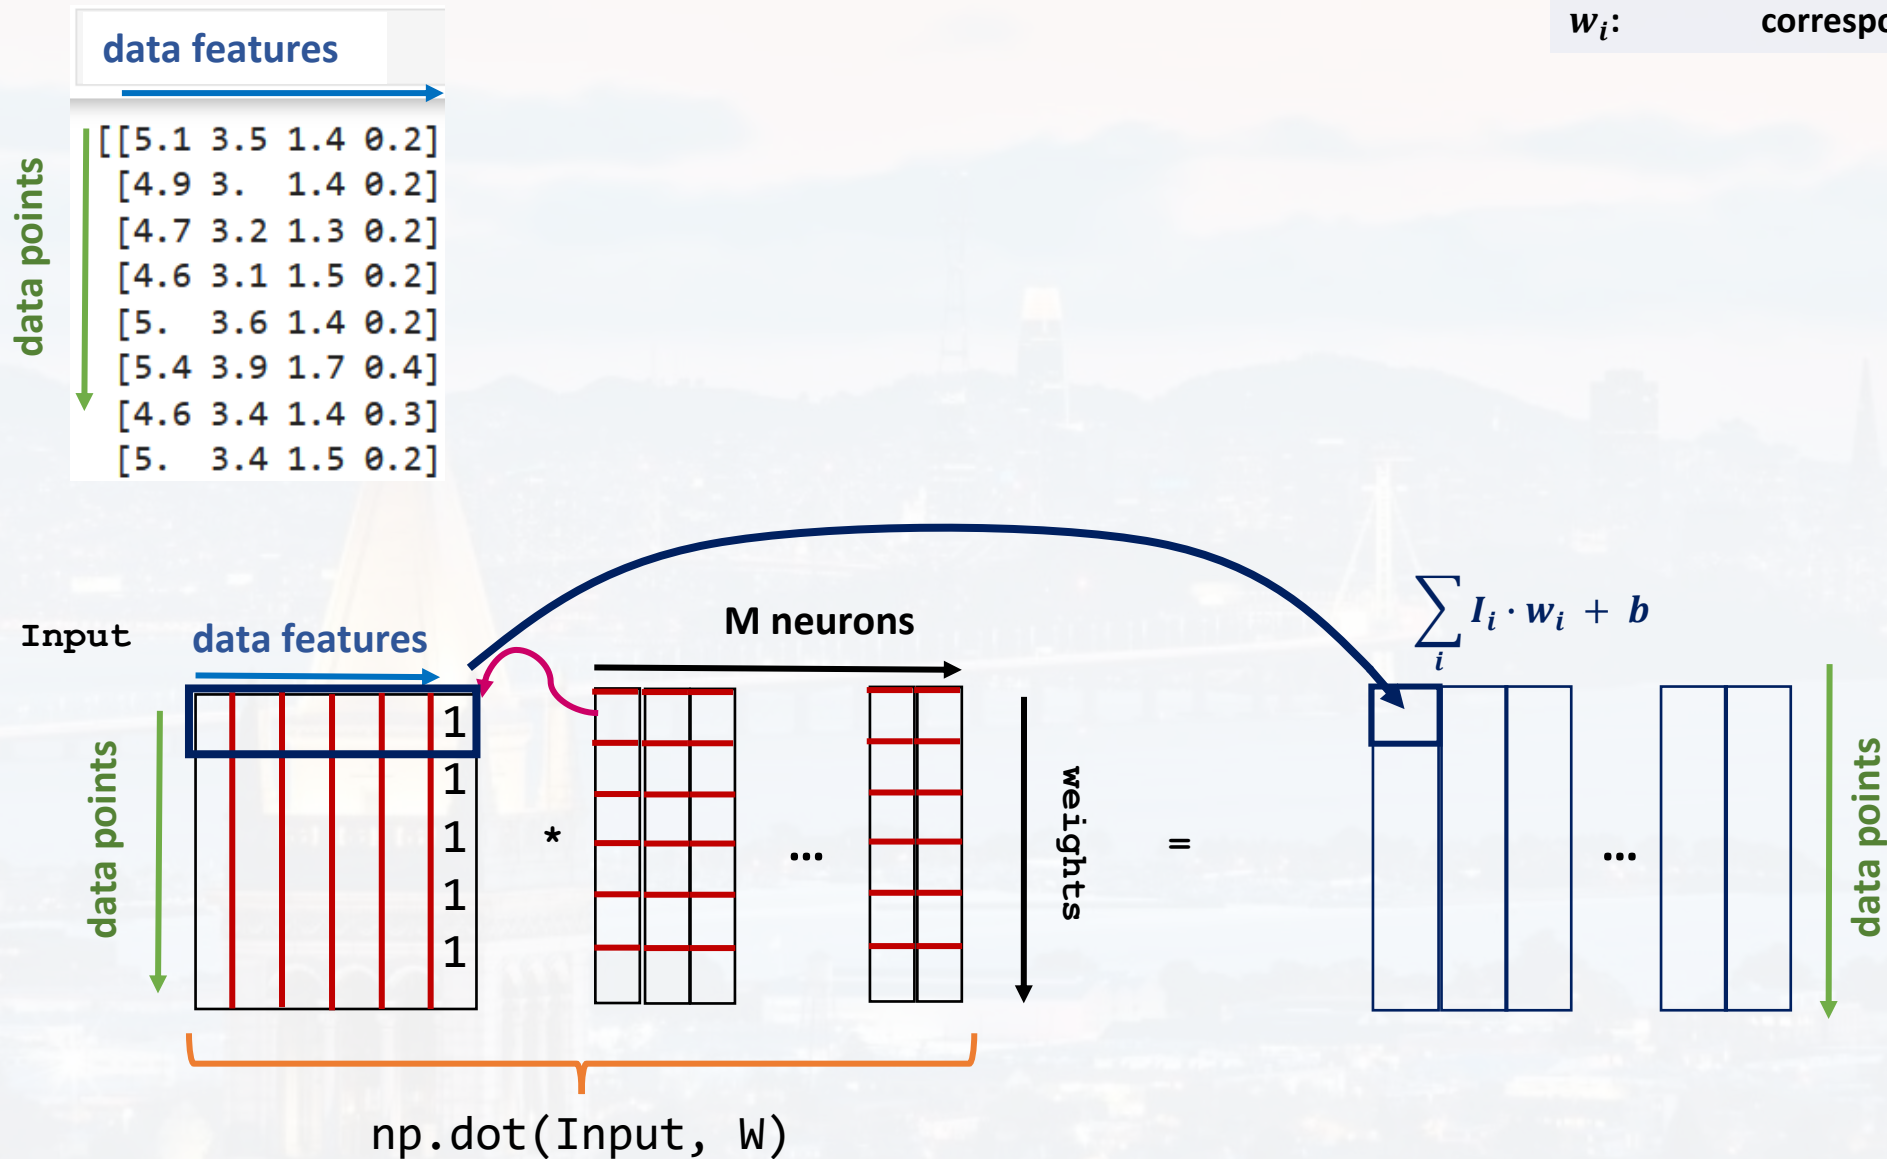


$t = 10,000$



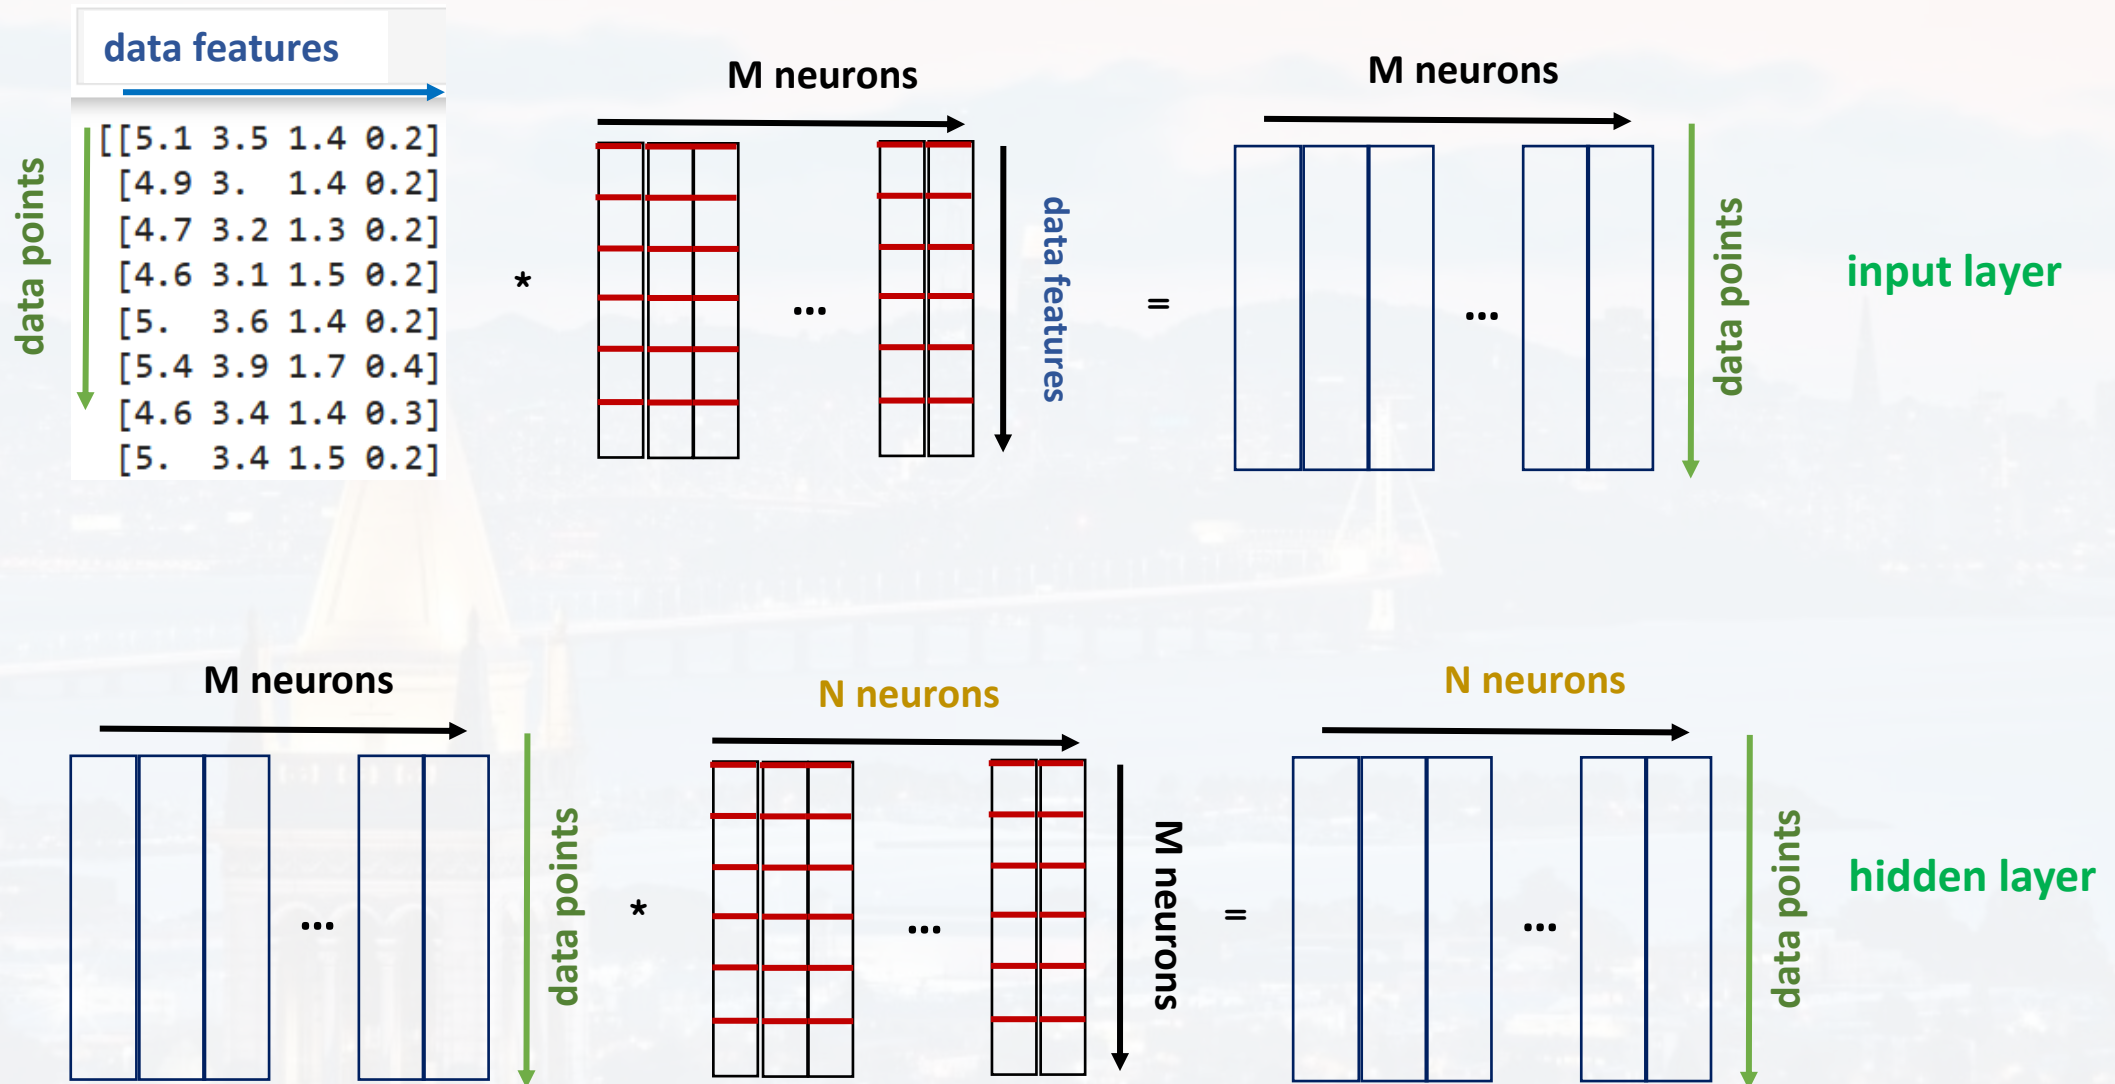
dense layer: multiple neurons

$b$ : bias (base potential)  
 $I_i$ : input  $i$   
 $w_i$ : corresponding weight



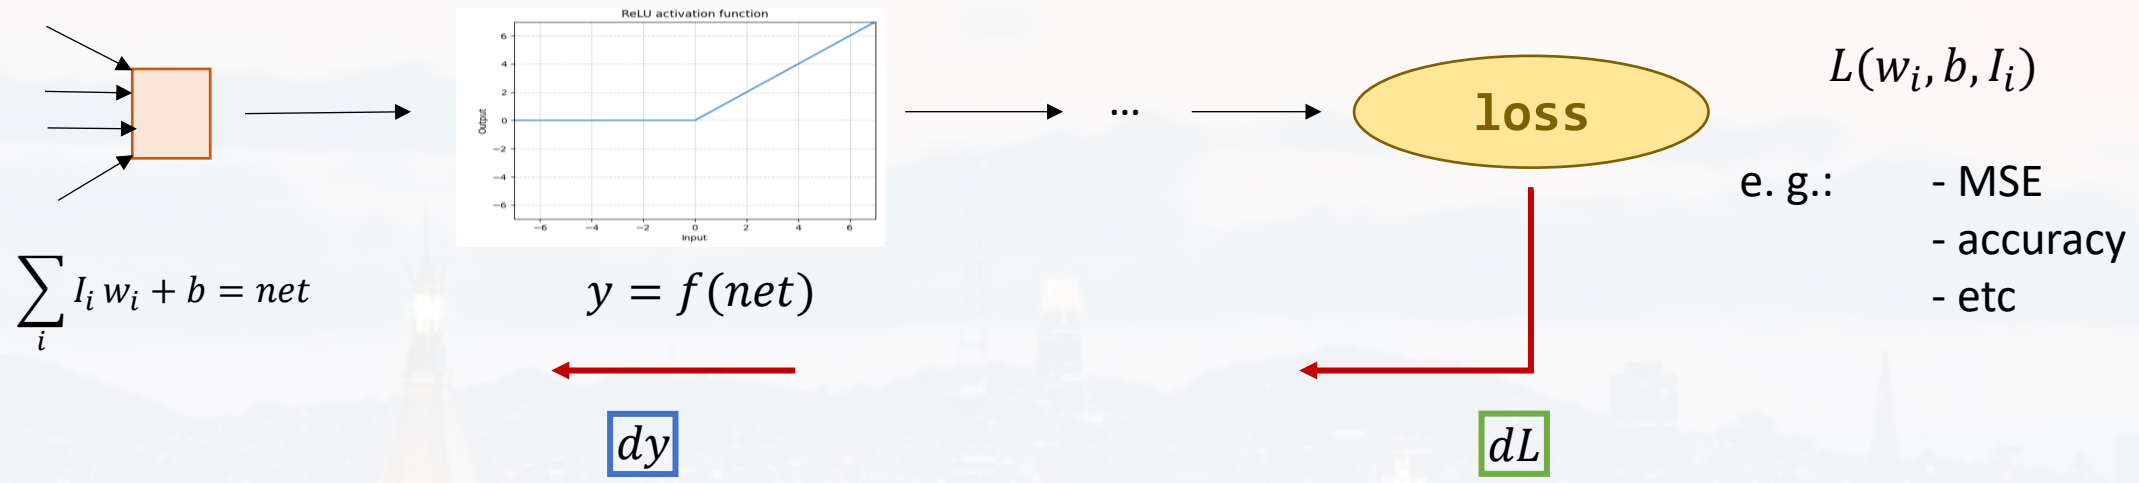
adding subsequent dense layer

$b$ : bias (base potential)  
 $I_i$ : input  $i$   
 $w_i$ : corresponding weight





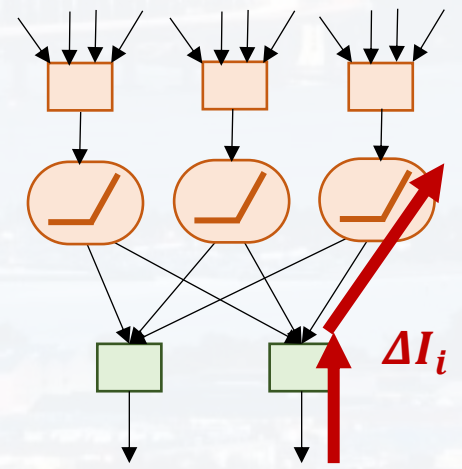
for training: building the backpropagation part



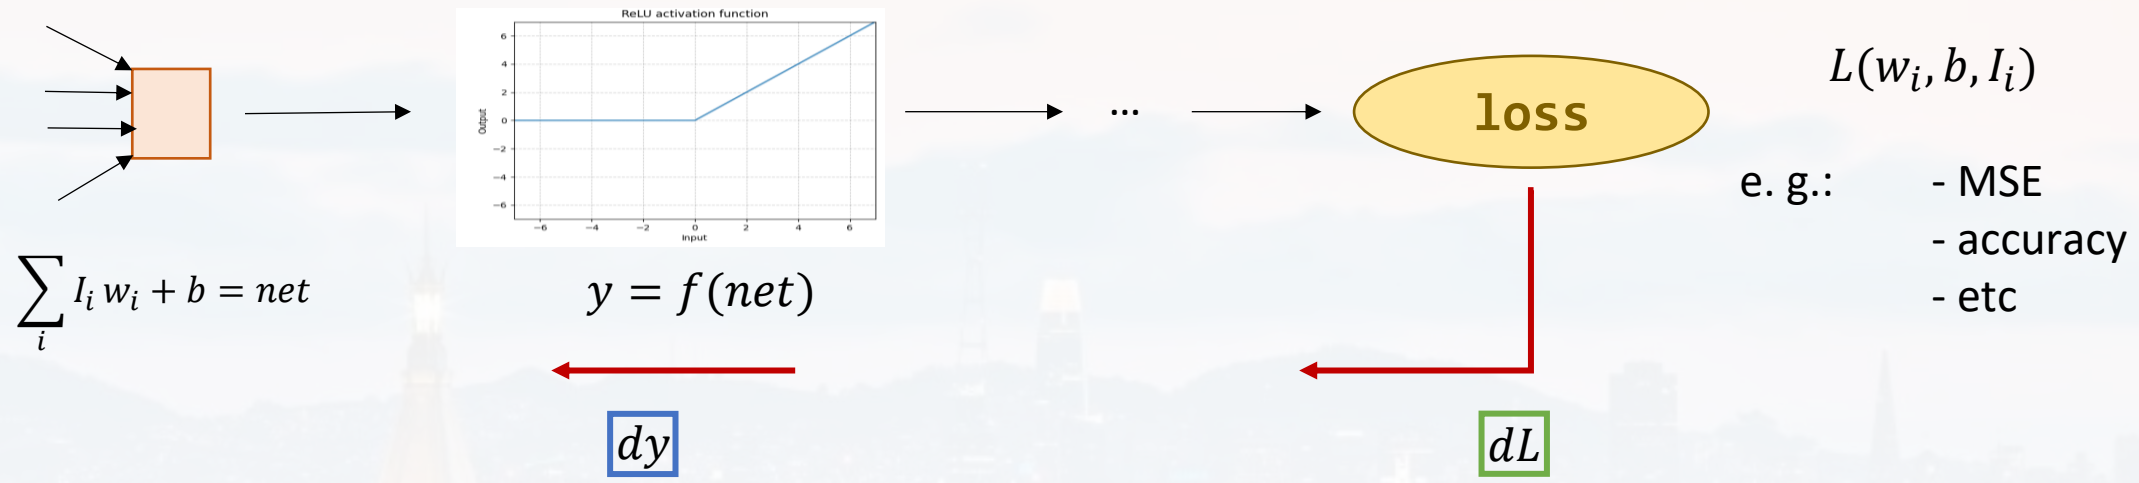
$$\Delta w_i = -\epsilon \frac{dL}{dy} \frac{dy}{dnet} I_i$$

$$\Delta I_i = -\epsilon \frac{dL}{dy} \frac{dy}{dnet} w_i \quad \text{from} \quad \frac{\partial L}{\partial I_i}$$

$$\Delta b = -\epsilon \frac{dL}{dy} \frac{dy}{dnet} 1 \quad \text{from} \quad \frac{\partial L}{\partial b}$$



for training: building the backpropagation part



$$\Delta w_i = -\epsilon \frac{dL}{dy} \frac{dy}{dnet} I_i$$

$$\Delta I_i = -\epsilon \frac{dL}{dy} \frac{dy}{dnet} w_i \quad \text{from} \quad \frac{\partial L}{\partial I_i}$$

$$\Delta b = -\epsilon \frac{dL}{dy} \frac{dy}{dnet} 1 \quad \text{from} \quad \frac{\partial L}{\partial b}$$

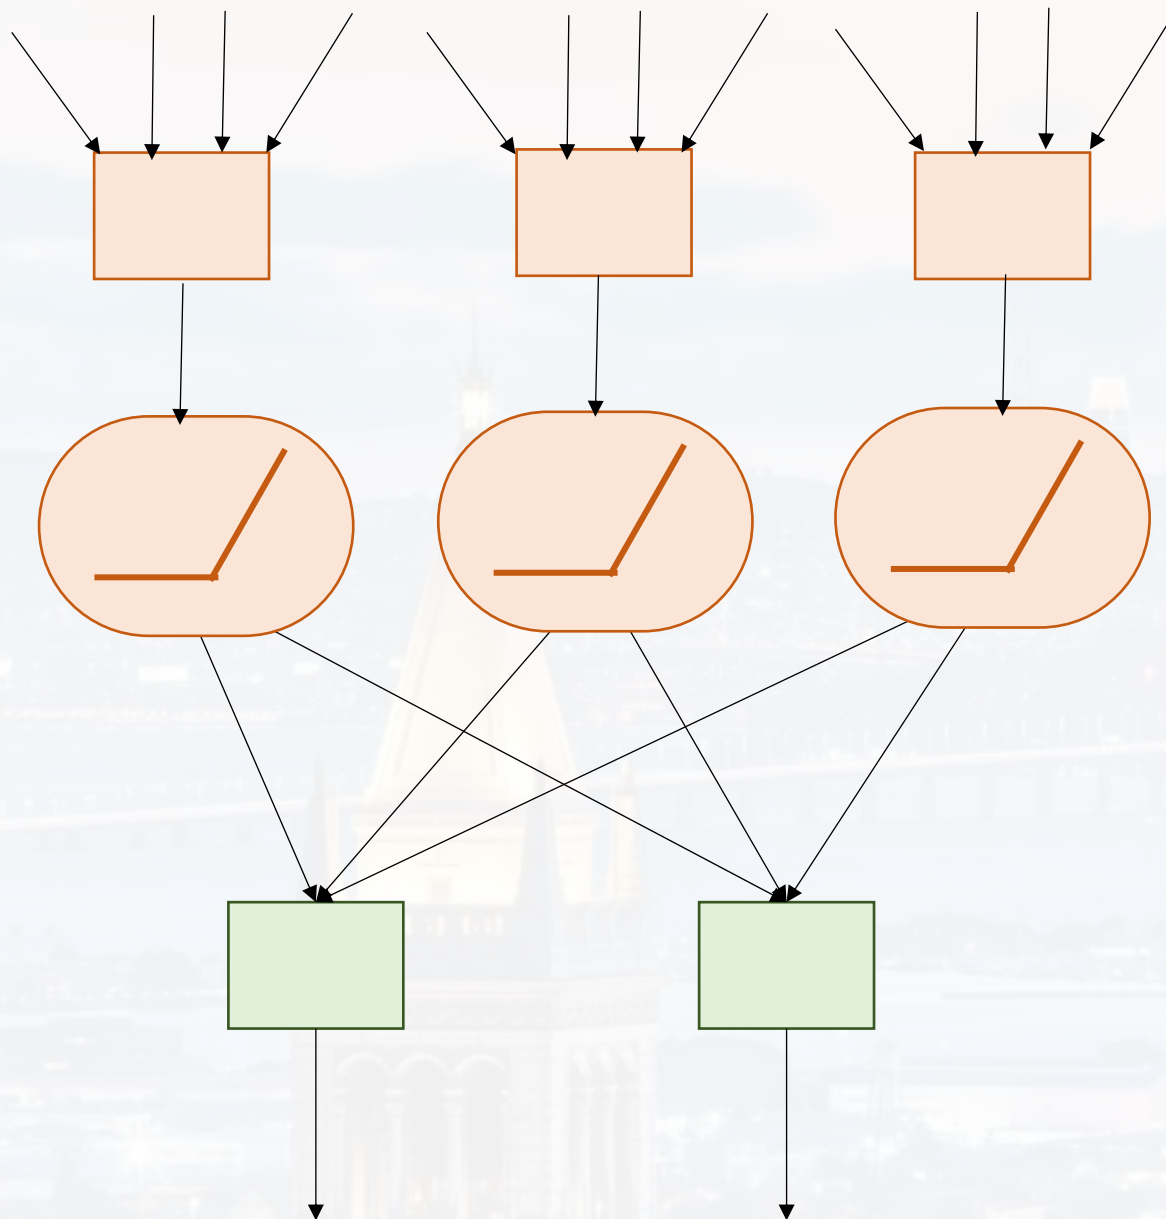
hence, we need to include the following structure for backpropagation:

```

self.dweights = inputs * dvalues
self.dinputs = weights * dvalues
self.dbiases = 1 * dvalues
  
```

inner  
derivative

product of  
outer  
derivatives



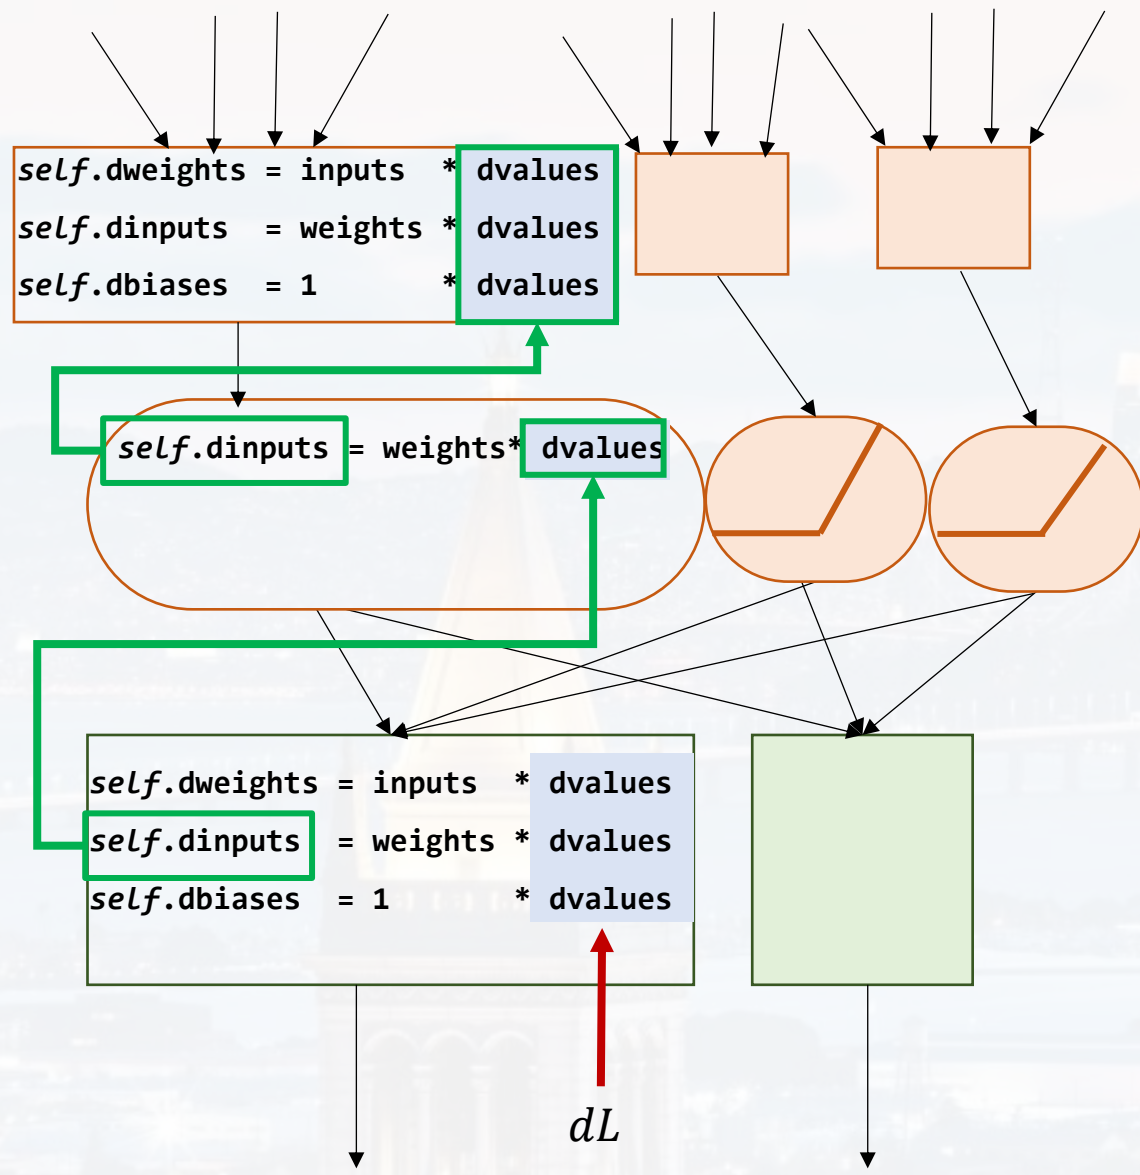
hence, we need to include the following structure for backpropagation:

```
self.dweights = inputs * dvalues
self.dinputs = weights * dvalues
self.dbiases = 1 * dvalues
```

inner  
derivative

product of  
outer  
derivatives





hence, we need to include the following structure for backpropagation:

<code>self.dweights</code>	<code>=</code>	<code>inputs</code>	<code>*</code>	<code>dvalues</code>
<code>self.dinputs</code>	<code>=</code>	<code>weights</code>	<code>*</code>	<code>dvalues</code>
<code>self.dbiases</code>	<code>=</code>	<code>1</code>	<code>*</code>	<code>dvalues</code>
		inner derivative		product of outer derivatives

$$\Delta I_i = -\epsilon \left[ \frac{dL}{dy} \right] \left[ \frac{dy}{dnet} \right] w_i$$

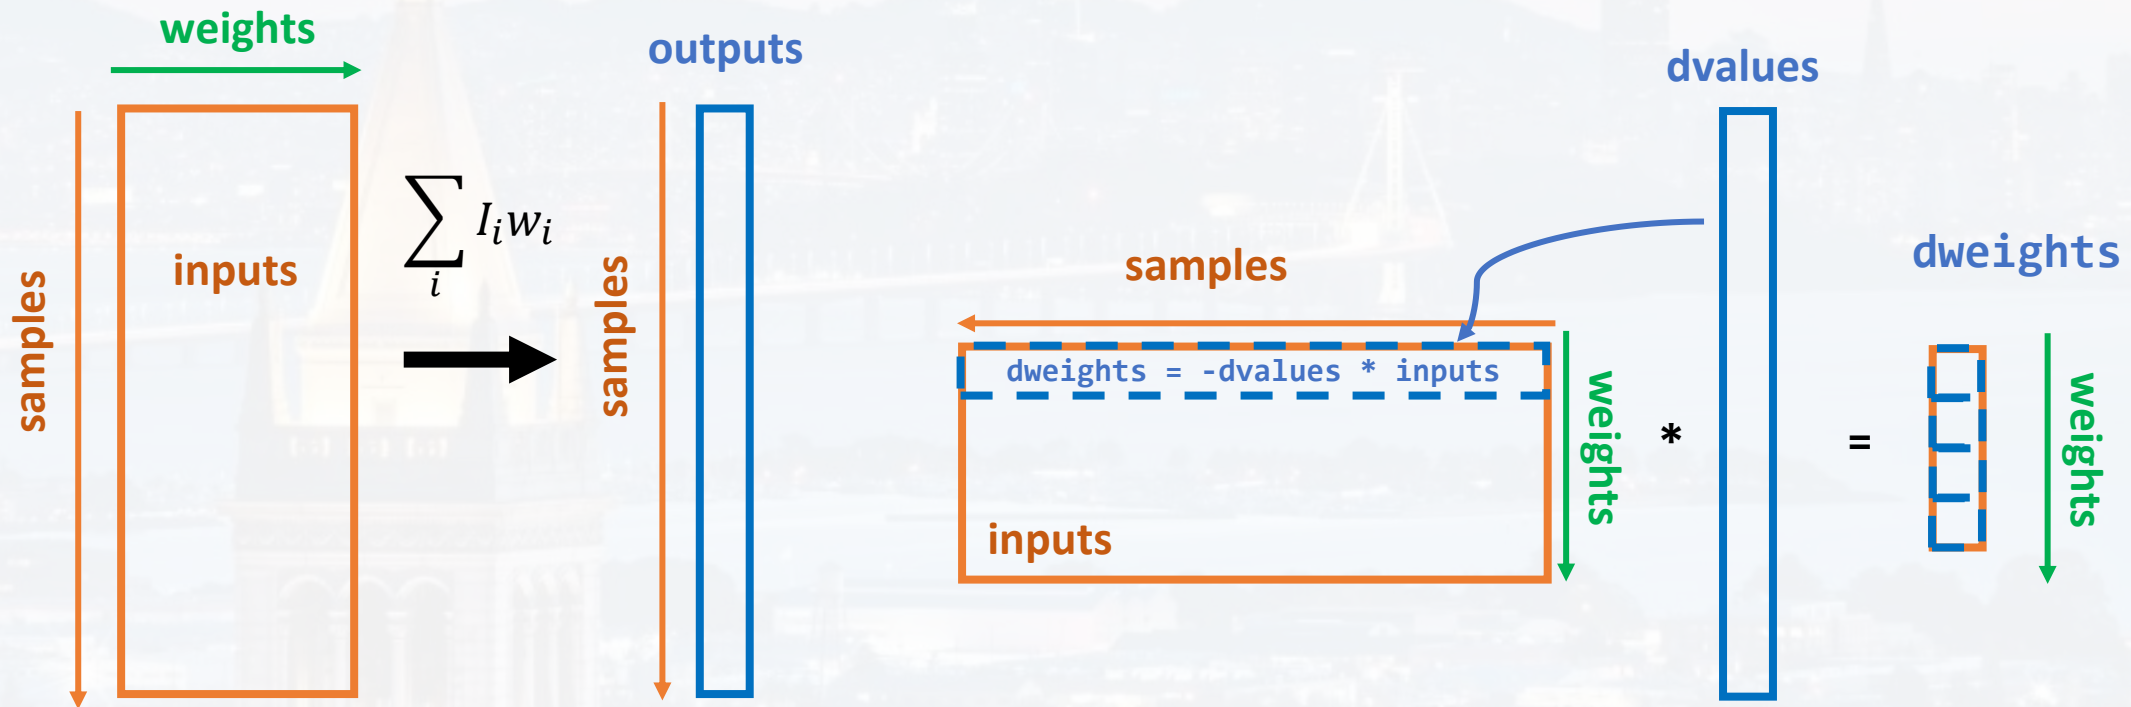
We need a **forward** part and a **backward** part!

single neuron:

forward

backward

$$dweights = -dvalues * inputs$$





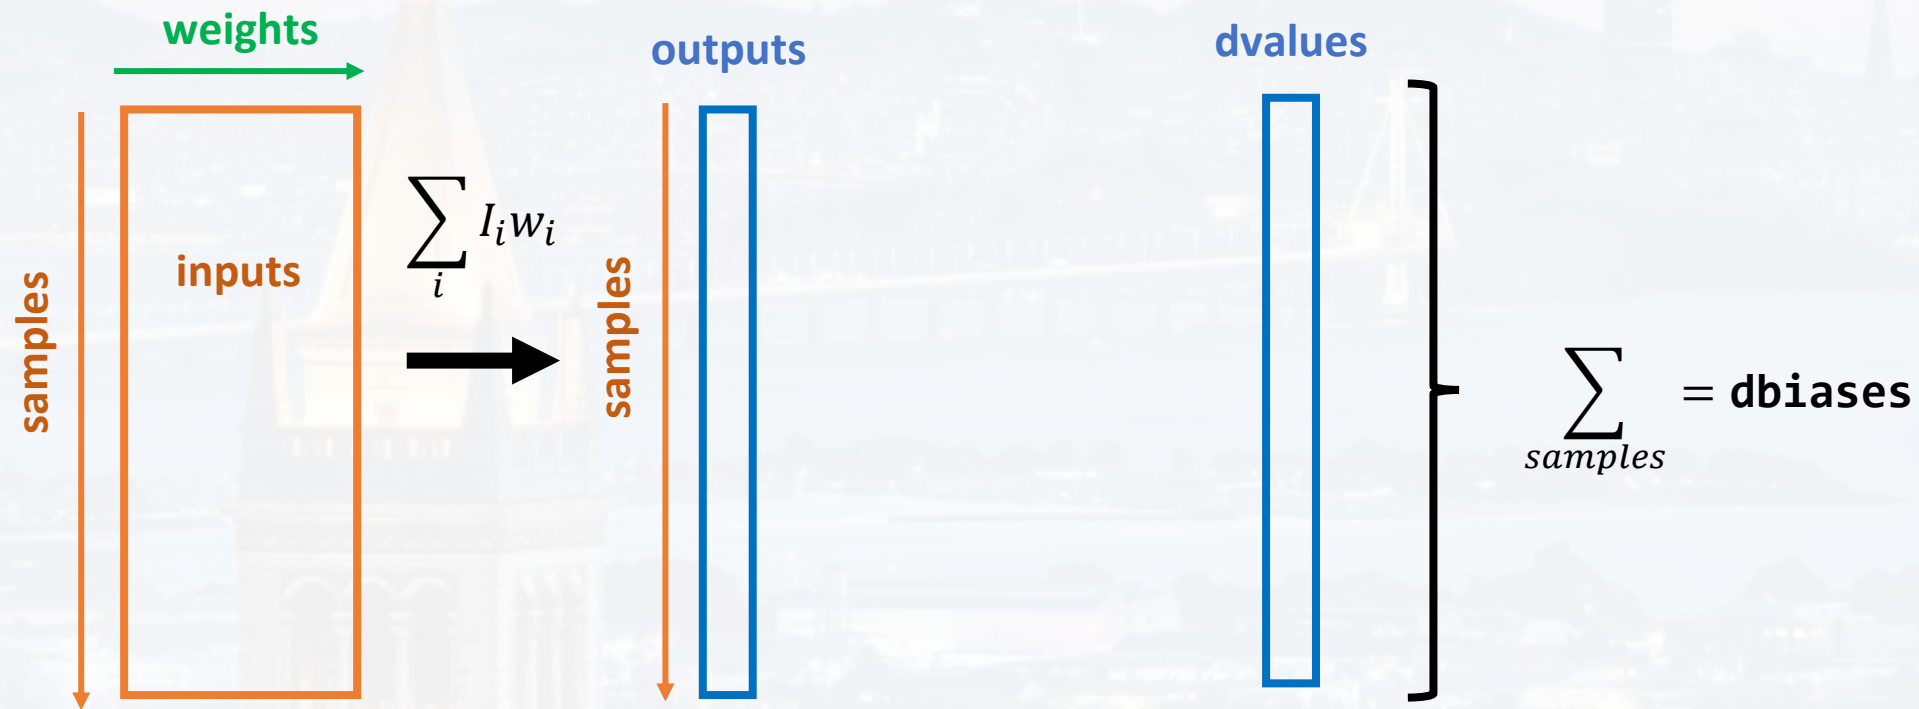
We need a **forward** part and a **backward** part!

single neuron:

forward

backward

**d**biases = -**d**values



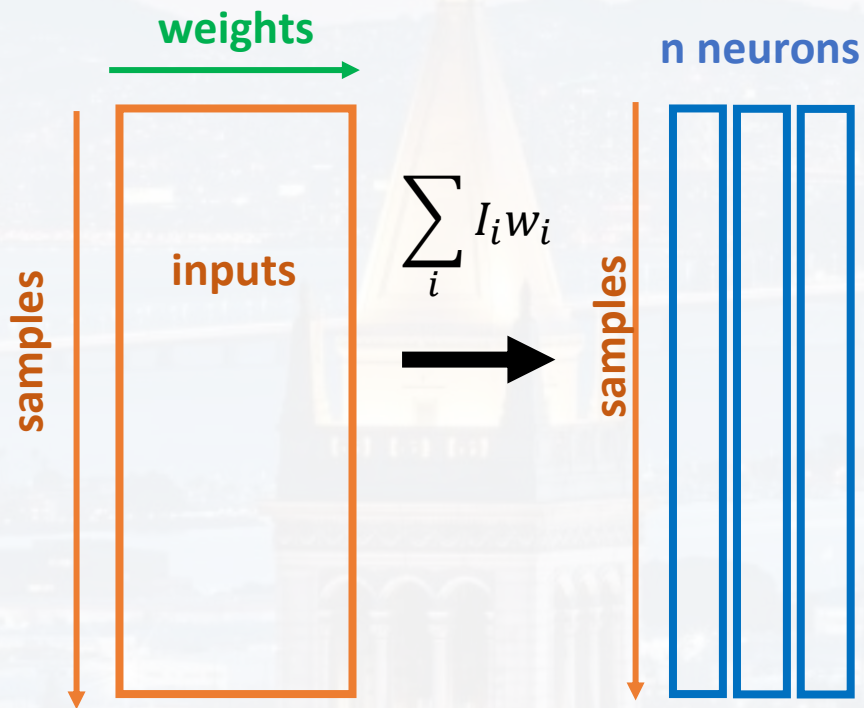




We need a **forward** part and a **backward** part!

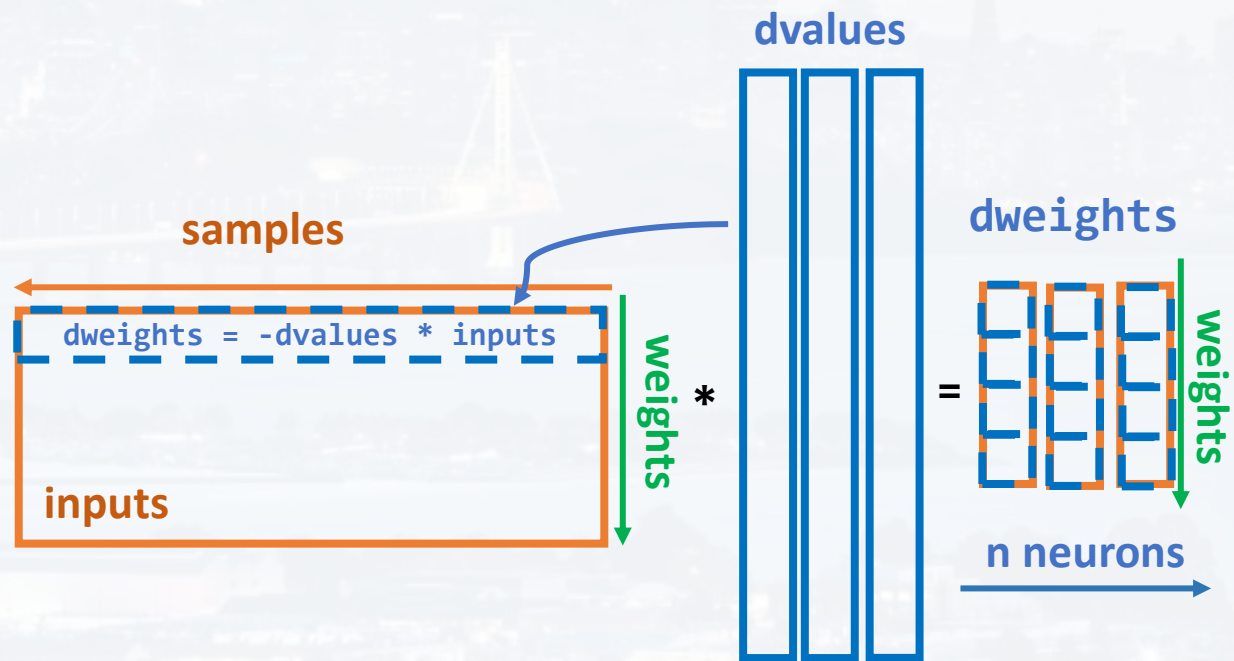
dense layer (say **three** neurons):

forward



backward

$$\text{dweights} = -\text{dvalues} * \text{inputs}$$

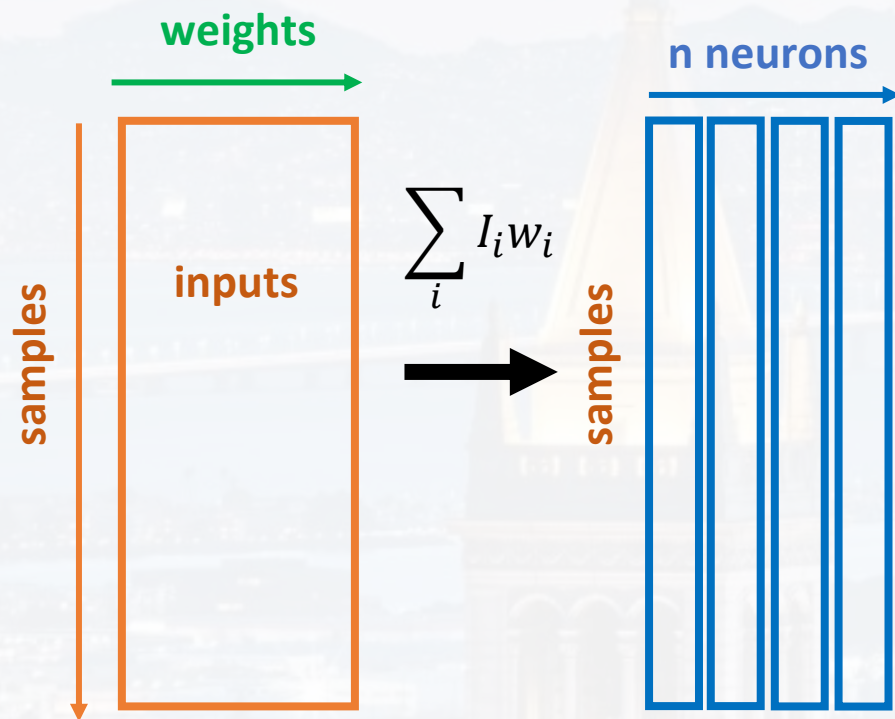




We need a **forward** part and a **backward** part!

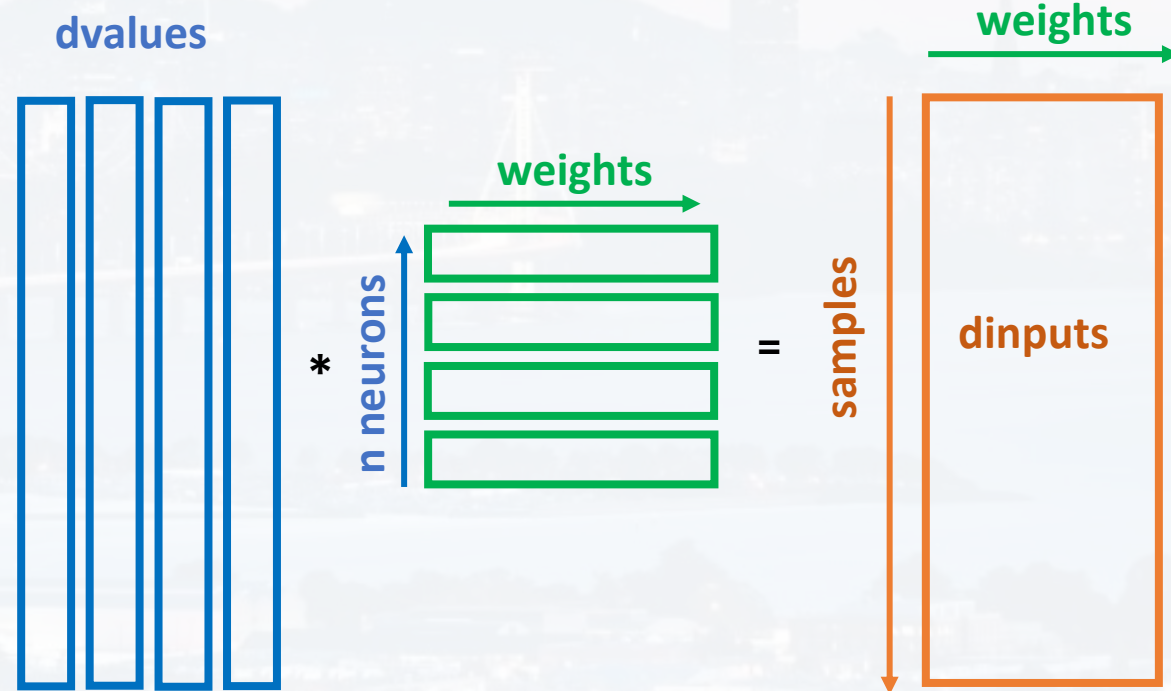
dense layer (say **four** neurons):

forward



backward

$\text{dinputs} = -\text{dvalues} * \text{weights}$  (not needed for first layer!)





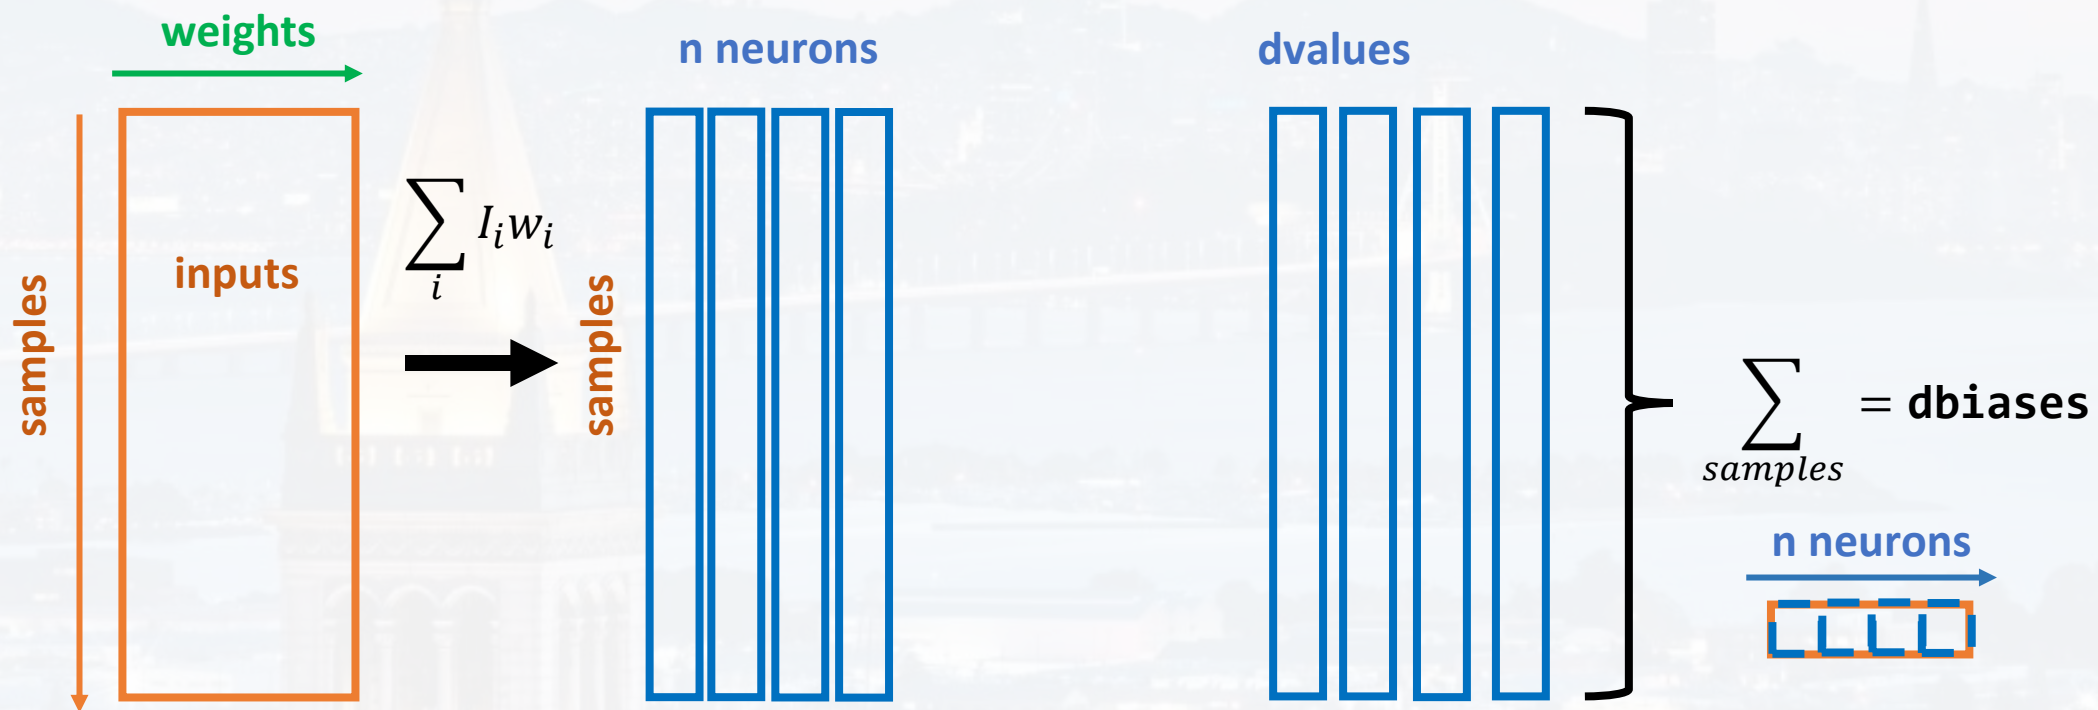
We need a **forward** part and a **backward** part!

dense layer (say **four** neurons):

forward

backward

**dbiases** = -**dvalues**







```
class Layer_Dense():
```

```
    def __init__(self, n_inputs, n_neurons):  
        self.weights = np.random.randn(n_inputs, n_neurons)  
        self.biases = np.zeros((n_neurons,))
```

number of features

```
    def forward(self, inputs):  
        self.output = np.dot(inputs, self.weights) + self.biases  
        self.inputs = inputs
```

outer derivative

```
    def backward(self, dvalues):  
  
        self.dweights = np.dot(self.inputs.T, dvalues)  
        self.dinputs = np.dot(dvalues, self.weights.T)  
        self.dbiases = np.sum(dvalues, axis = 0, keepdims = True)
```

<code>self.dweights</code>	<code>= inputs</code>	<code>*</code>	<code>dvalues</code>
<code>self.dinputs</code>	<code>= weights</code>	<code>*</code>	<code>dvalues</code>
<code>self.dbiases</code>	<code>= 1</code>	<code>*</code>	<code>dvalues</code>

inner derivative

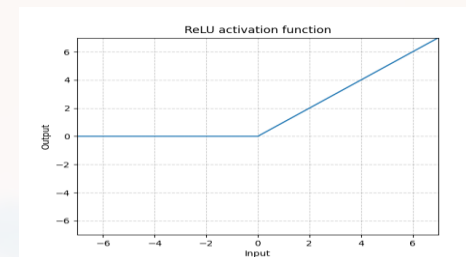
product of outer derivatives



```
class Activation_ReLU():
```

```
    def forward(self, inputs):
        self.output = np.maximum(0, inputs)
        self.inputs = inputs

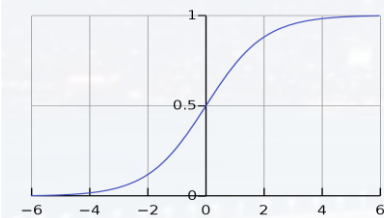
    def backward(self, dvalues):
        self.dinputs = dvalues.copy()
        self.dinputs[self.inputs <= 0] = 0 #ReLU derivative
```



```
class Activation_Sigmoid():
```

```
    def forward(self, inputs):
        self.output = np.clip(1/(1 + np.exp(-inputs)), 1e-7, 1-1e-7)
        self.inputs = inputs

    def backward(self, dvalues):
        sigm = self.output
        deriv = np.multiply(sigm, (1 - sigm))
        self.dinputs = np.multiply(deriv, dvalues)
```



basic structure:    `alpha = 0.001` *#learning rate*

```
dense1.forward(X)
ReLU.forward(dense1.output)
dense_reg.forward(ReLU.output)
```

**forward**

```
Ypred = dense_reg.output
dE     = Ypred - Target
MSE    = np.sum(abs(dE))/(Nsample*Nclasses)
print('MSE = ' + str(MSE))
```

**evaluation**

```
dense_reg.backward(dE)
ReLU.backward(dense_reg.dinputs)
dense1.backward(ReLU.dinputs)
```

**backpropagation**

```
dense_reg.weights -= alpha * dense_reg.dweights
dense_reg.biases  -= alpha * dense_reg.dbiases
dense1.weights    -= alpha * dense1.dweights
dense1.biases     -= alpha * dense1.dbiases
```

**optimization**

see ANNI.ipynb and ANNII.ipynb



next time:                    -fully functional ANN (minimal setup)  
                                     -regression vs classification

Thank you very much for your attention!

