

Lecture 11:

Convolution and Image Classification & Segmentation



Markus Hohle

University California, Berkeley

Bayesian Data Analysis and
Machine Learning for Physical
Sciences

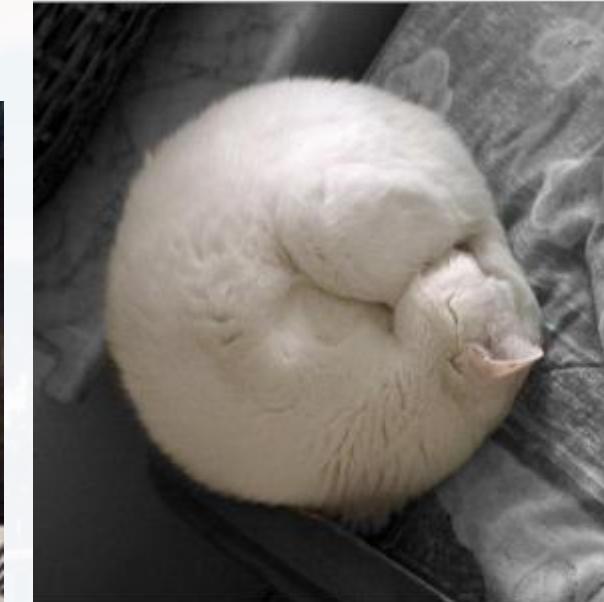
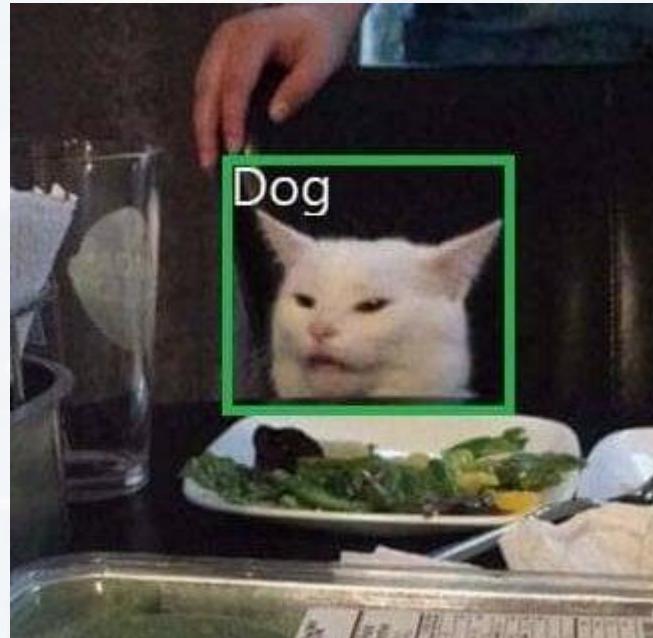


Course Map

Module 1	Maximum Entropy and Information, Bayes Theorem
Module 2	Naive Bayes, Bayesian Parameter Estimation, MAP
Module 3	MLE, Lin Regression
Module 4	Model selection I: Comparing Distributions
Module 5	Model Selection II: Bayesian Signal Detection
Module 6	Variational Bayes, Expectation Maximization
Module 7	Hidden Markov Models, Stochastic Processes
Module 8	Monte Carlo Methods
Module 9	Machine Learning Overview, Supervised Methods & Unsupervised Methods
Module 10	ANN: Perceptron, Backpropagation, SGD
Module 11	Convolution and Image Classification and Segmentation
Module 12	RNNs and LSTMs
Module 13	RNNs and LSTMs + CNNs
Module 14	Transformer and LLMs
Module 15	Graphs & GNNs



Part III



0.40 cat
0.32 frog
0.16 bird
0.06 ship
0.03 dog



0.40 cat
0.32 frog
0.16 bird
0.06 ship
0.03 dog

Outline

PyTorch & Cuda

Generative Adversarial Network (GAN)

Variational AutoEncoder (VAE)



0.40 cat
0.32 frog
0.16 bird
0.06 ship
0.03 dog

Outline

PyTorch & Cuda

Generative Adversarial Network (GAN)

Variational AutoEncoder (VAE)



problem: even for moderate setups, computational time becomes the limiting factor

self made ANN

keras/tensor flow

pytorch

pytorch on GPU (Cuda)

speed →

```
--, --, --, --, --
512/512 [=====] - 4072s 8s/step - loss: 3.4534 - accuracy: 0.3338
Epoch 4/5
512/512 [=====] - ETA: 0s - loss: 3.1926 - accuracy: 0.3980      saved  ../data/segmentation
pics/checkpoints//.3
512/512 [=====] - 3363s 7s/step - loss: 3.1926 - accuracy: 0.3980
Epoch 5/5
512/512 [=====] - ETA: 0s - loss: 3.0071 - accuracy: 0.4318      saved  ../data/segmentation
pics/checkpoints//.4
512/512 [=====] - 3616s 7s/step - loss: 3.0071 - accuracy: 0.4318
```

Lenovo T14, NVIDIA GeForce MX450: (simple LSTM)

Keras (CPU):


300 sec

PyTorch (CPU):

11 sec

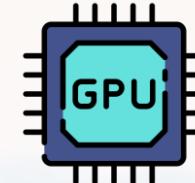
PyTorch (GPU):

3 sec



training AI

- mainly matrix operations
- GPUs are a lot better at it!



CUDA is the link of your GPU to Python (PyTorch)
check, if graphic card is on [list](#)

check your graphics device:

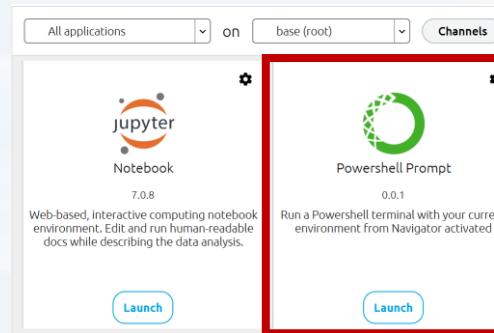
- Windows command shell prompt
- type nvidia-smi
- press *Enter*

```
(base) C:\Users\MMH_user>nvidia-smi
Tue Jan 16 20:09:26 2024
+-----+
| NVIDIA-SMI 537.79           Driver Version: 537.79      CUDA Version: 12.2 |
+-----+
| GPU  Name        TCC/WDDM     Bus-Id      Disp.A  Volatile Uncorr. ECC
| Fan  Temp  Perf  Pwr:Usage/Cap | Memory-Usage  GPU-Util  Compute M.
|          %   %       % / %   / % / % / % / % / % / % / % / % / % / % / %
|-----+
| 0  NVIDIA GeForce MX450    WDDM      00000000:01:00.0 Off  0%       N/A
| N/A  49C  P0    N/A /  9W  0MiB / 2048MiB  0%       Default
|          |          |          |          |          |          |          |
+-----+
| Processes:
| GPU  GI  CI      PID  Type  Process name
| ID   ID
|-----+
| No running processes found
|-----+
```



Installing CUDA

conda environment



```
C:\WINDOWS\System32\WindowsPowerShell\v1.0> (base) PS C:\Users\MMH_user>
```

```
C:\WINDOWS\System32\WindowsPowerShell\v1.0> (base) PS C:\Users\MMH_user> conda create --name CUDAenv |
```

```
C:\WINDOWS\System32\WindowsPowerShell\v1.0> (base) PS C:\Users\MMH_user> conda activate CUDAenv
```

```
(base) PS C:\Users\MMH_user> conda activate CUDAenv  
(CUDAenv) PS C:\Users\MMH_user> conda install -c pytorch pytorch  
Channels:  
- pytorch
```

Installing CUDA

cuda toolkit

```
(CUDAenv) PS C:\Users\MMH_user> conda install -c anaconda cudatoolkit
```

check libraries

→ type: conda list

```
(CUDAenv) PS C:\Users\MMH_user> conda list
# packages in environment at C:\Users\MMH_user\anaconda3\envs\CUDAenv:
#
#           Name          Version      Build  Channel
blas          1.0                  mkl
bzip2         1.0.8              h2bbff1b_6
ca-certificates 2024.7.2        haa95532_0
cudatoolkit    11.8.0             hd77b12b_0
expat          2.6.2              hd77b12b_0
filelock        3.13.1            py312haa95532_0
intel-openmp   2023.1.0          h59b6b97_46320
```



Installing CUDA

usually, a few libraries are missing

check again graphics card: type in anaconda **nvidia-smi**

check libraries: type in anaconda
conda list cudnn
conda list cudatoolkit
conda list torch

if not: **conda install <library>**

check Python: type in anaconda
python
import torch
torch.cuda.is_available()

open Spyder run in Spyder **pip install**
(see the commented line in CheckMyCuda.py)

Installing CUDA

usually, a few libraries are missing

CheckMyCuda.py

```
import torch

def test_cuda():
    print("PyTorch version: ", torch.__version__)
    print("CUDA version: ", torch.version.cuda)
    print("CUDA Available: ", torch.cuda.is_available())
    if torch.cuda.is_available():
        print("Number of GPUs: ", torch.cuda.device_count())
        print("GPU Name: ", torch.cuda.get_device_name(0))

if __name__ == "__main__":
    test_cuda()
```

```
PyTorch version: 2.3.1+cu118
CUDA version: 11.8
CUDA Available: True
Number of GPUs: 1
GPU Name: NVIDIA GeForce MX450
```



The key part in PyTorch is to set all matrices and the model **to the device (CPU or GPU)**

```
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print("Using device:", device)
```

```
In [13]: device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
....: print("Using device:", device)
Using device: cuda
```

Congratulation! If you see this, you are ready to go!



The key part in PyTorch is to set all matrices and the model **to the device (CPU or GPU)**

```
In [13]: device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
....: print("Using device:", device)
Using device: cuda
```

Torch objects like **model** or **torch.tensor** have the property **.to**

```
TrainX = torch.tensor(TrainX, dtype = torch.float32)
TrainY = torch.tensor(TrainY, dtype = torch.float32)
```

```
TrainX = TrainX.to(device)
TrainY = TrainY.to(device)
```

turning numpy array into torch.tensor

```
model = model.to(device)
```

allocating objects to the device



The key part in PyTorch is to set all matrices and the model **to the device (CPU or GPU)**

```
In [13]: device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
         ...: print("Using device:", device)
Using device: cuda
```

When running the training, we need to **synchronize** between GPU (for training the model) and CPU (for everything else)...

```
torch.cuda.synchronize()
```

```
#training loop
for epoch in range(n_epochs):
    outputs = model(TrainX)
```

```
    ... #do stuff...
```

```
torch.cuda.synchronize()
```



The key part in PyTorch is to set all matrices and the model **to the device (CPU or GPU)**

```
In [13]: device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
         ...: print("Using device:", device)
Using device: cuda
```

When running the training, we need to **synchronize** between GPU (for training the model) and CPU (for everything else)...

...and later detach the model from the GPU

```
PredY = model(TestX).detach().to('cpu').numpy()
```




0.40 cat
0.32 frog
0.16 bird
0.06 ship
0.03 dog


Outline

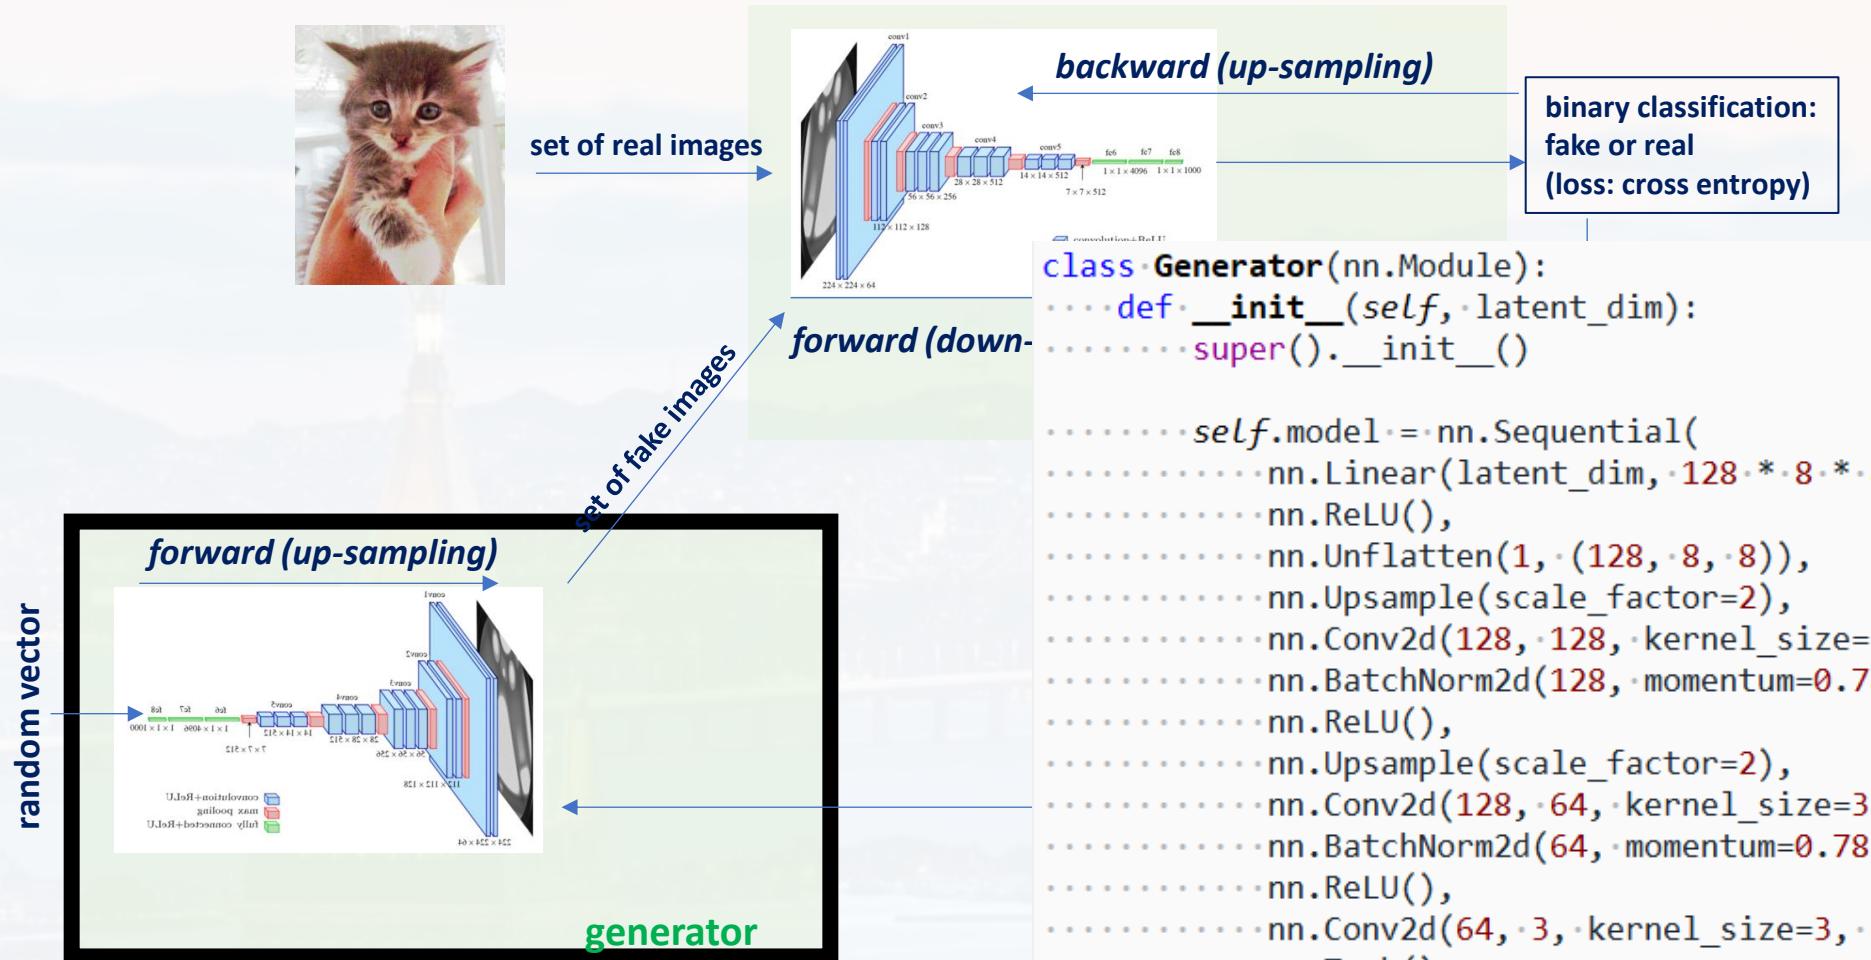
PyTorch & Cuda

Generative Adversarial Network (GAN)

Variational AutoEncoder (VAE)









x : real image

z : latent vector (random numbers)

$p_D(x)$: probability that discriminator D classifies x as real image

$G(z)$: fake image generated from generator G

\mathcal{L}_D : loss function of discriminator D

\mathcal{L}_G : loss function of generator G

(mean) binary cross entropy (BCE): $\mathcal{L}_D = -\frac{1}{N} \sum_{i=1}^N \log[p_D(x_i)] + \log[1 - p_D(G(z_i))]$

over all N images

$$\mathcal{L}_G = \frac{1}{N} \sum_{i=1}^N \log[1 - p_D(G(z_i))]$$

ideal case for D : $\mathcal{L}_D = -\log[1] - \log[1 - 0] = 0$

worst case for D : $\mathcal{L}_D = -\log[1] - \log[1 - \mathbf{0.5}] = \mathbf{0.69}$

ideal case for G : $\mathcal{L}_G = \log[1 - \mathbf{0.5}] = \mathbf{-0.69}$

worst case for G : $\mathcal{L}_G = \log[1 - 0] = 0$



Berkeley
UNIVERSITY OF CALIFORNIA

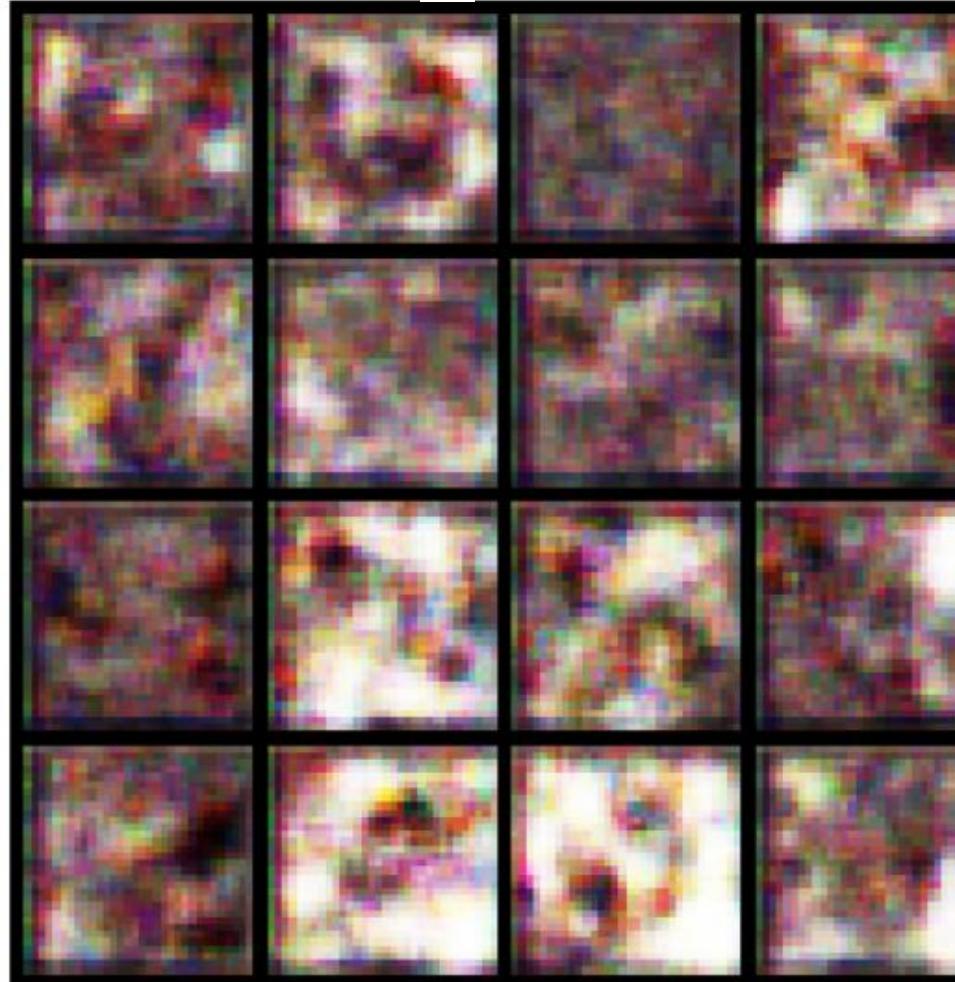
Convolution, Image Classification & Segmentation

GAN

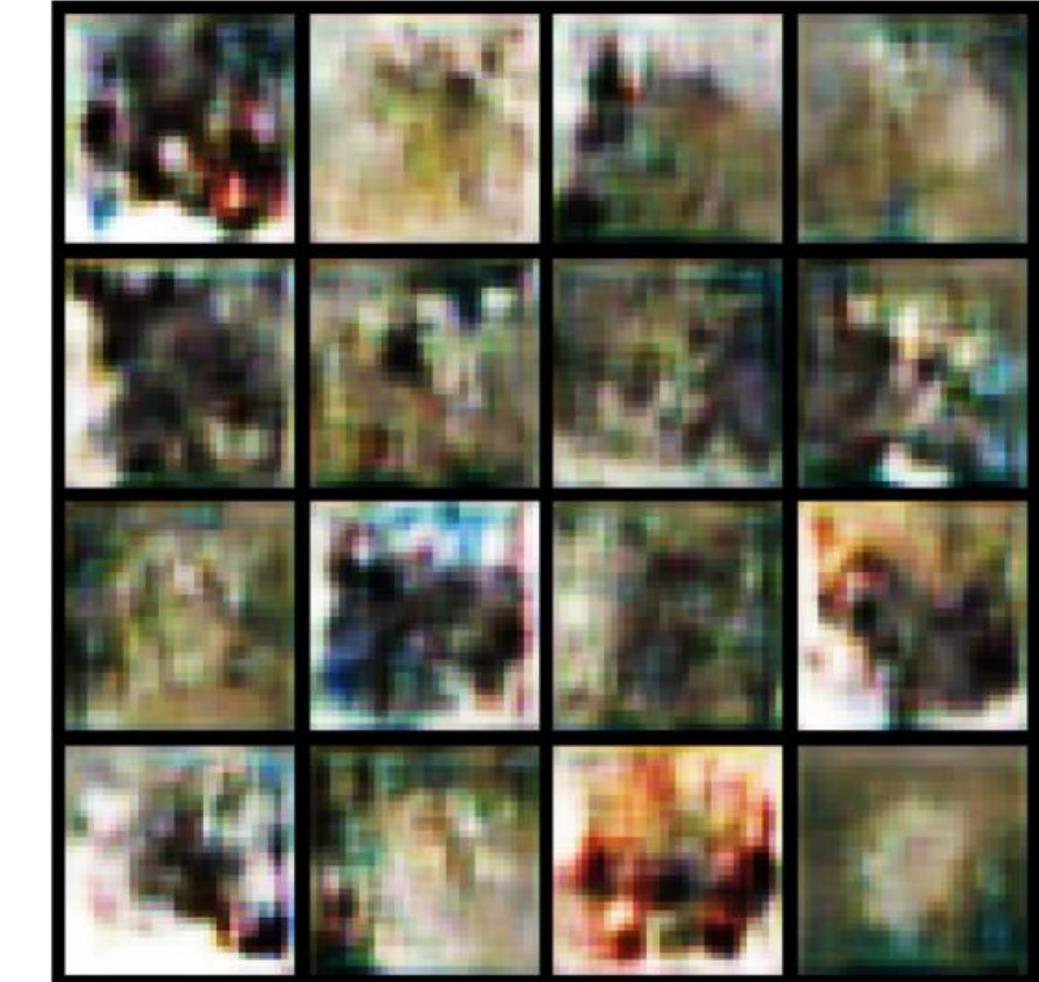
check out

GAN architecture.ipynb

after 1 epoch(s)



after 10 epoch(s)





Berkeley
UNIVERSITY OF CALIFORNIA

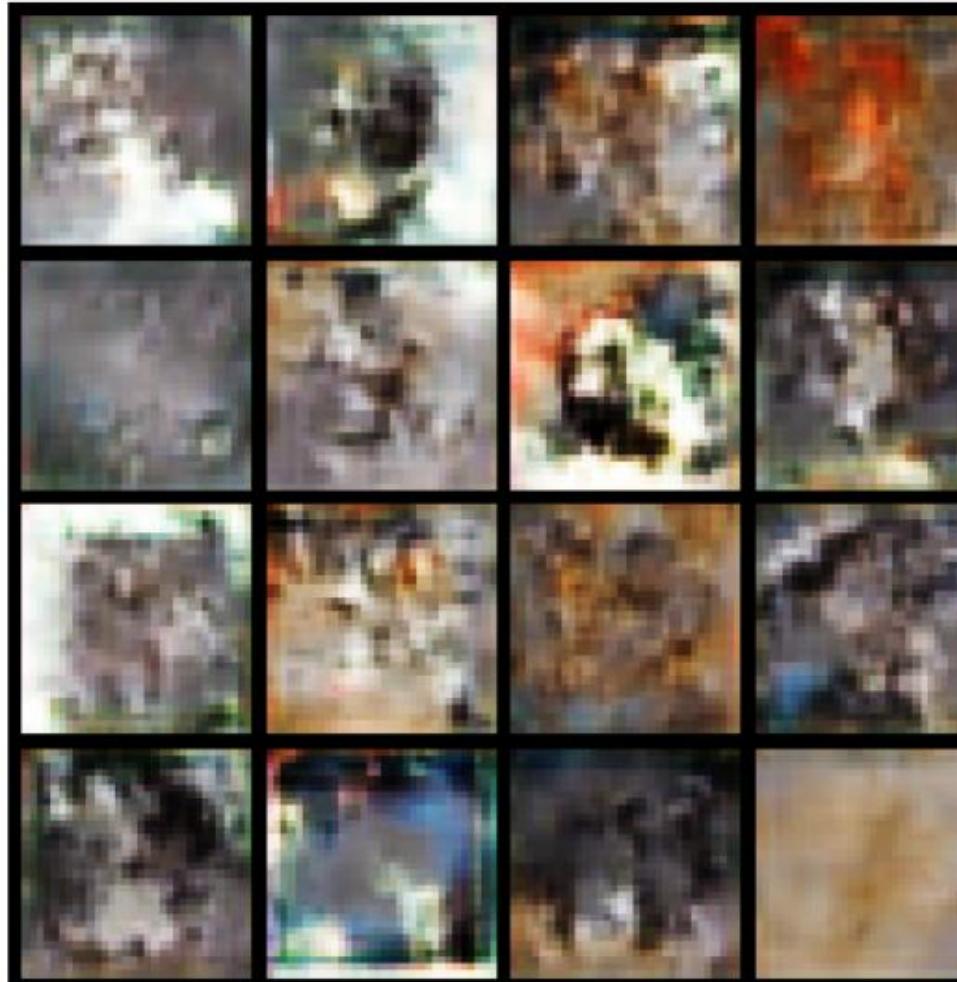
Convolution, Image Classification & Segmentation

GAN

check out


GAN architecture.ipynb

after 30 epoch(s)



after 55 epoch(s)





0.40 cat
0.32 frog
0.16 bird
0.06 ship
0.03 dog

Outline


PyTorch & Cuda

Generative Adversarial Network (GAN)

Variational AutoEncoder (VAE)

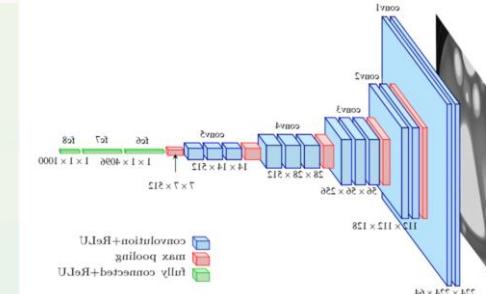
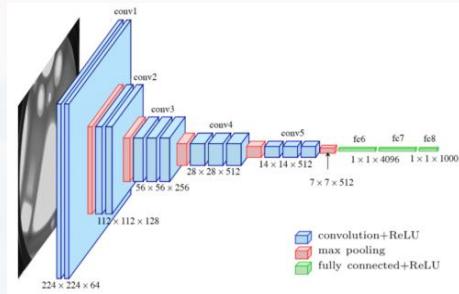


GAN:







now:





now:



GAN: each image x is represented by **one single** vector z

VAE: each image x is represented by **a distribution** we sample from in order **to get z**

i : index over dimensions of z


$$z_i \sim \mathcal{N}(\mu_i, \sigma_i^2) \quad (\text{max entropy!})$$

→ generate image from up-sampling as before

→ randomness accounts for diversity in generated data



now:



example:

output layer has shape (None, N)

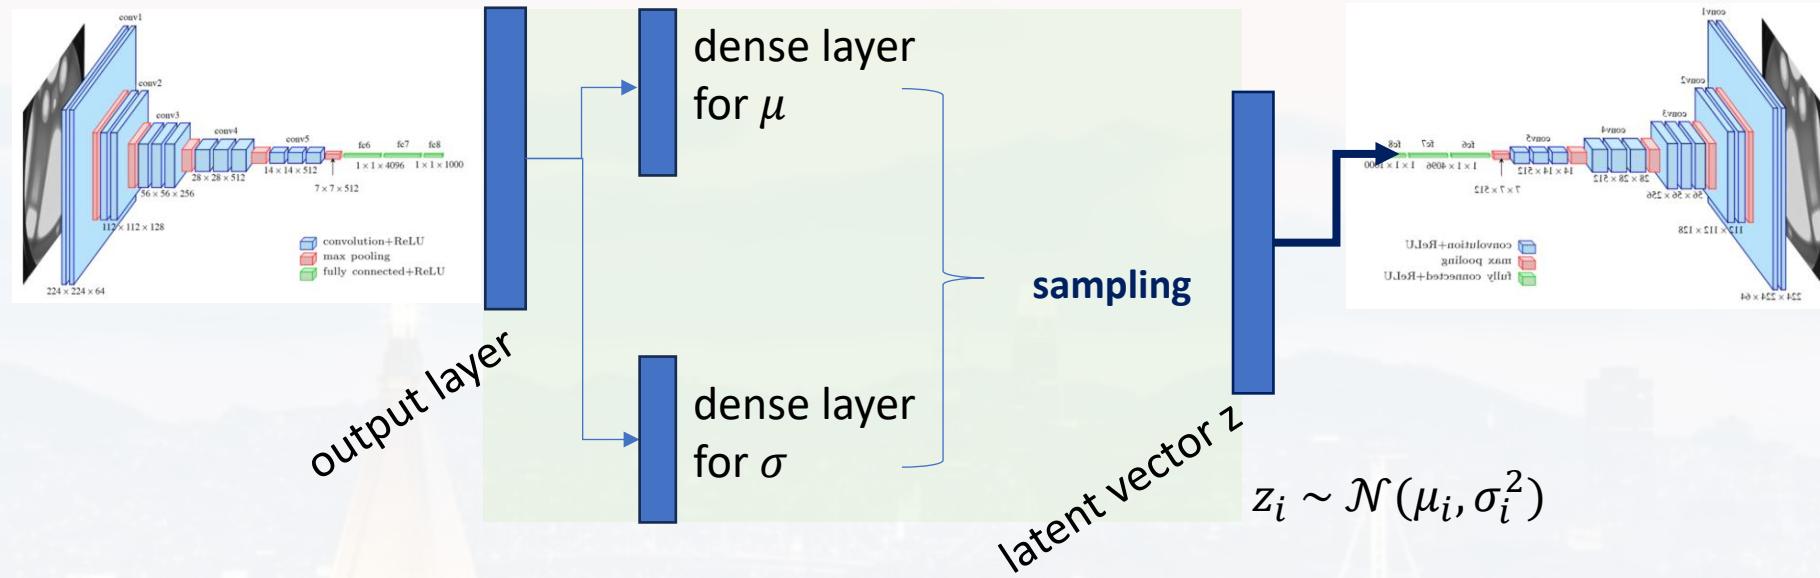
we want M dimensional encoding for z

→ weights in dense layer for σ and μ have shape (N, M)

→ latent vector has length M



now:

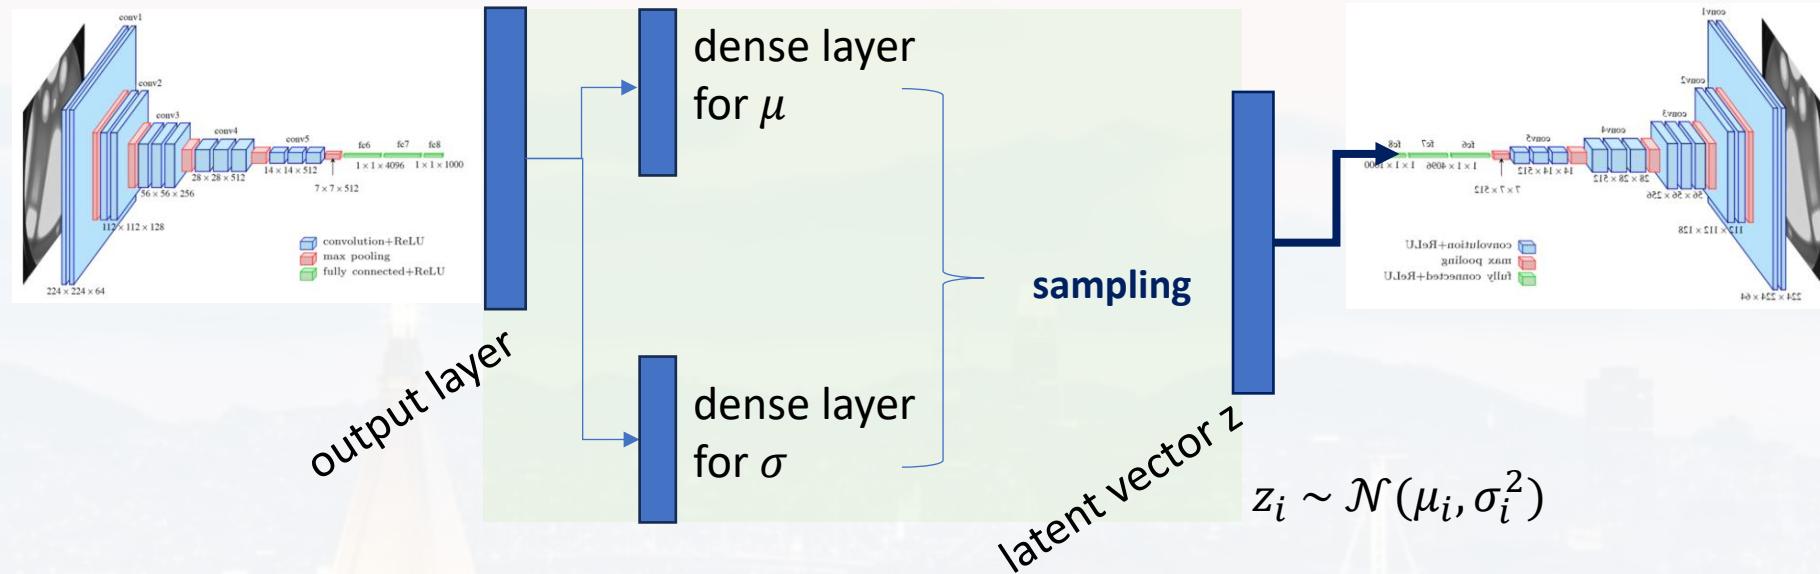


backprop:

- μ_i and σ_i^2 have been generated deterministically from a dense layer
- we need to run backpropagation through μ_i and σ_i^2 in order to adjust the weights/biases
- in order to keep randomness, we actually sample via
$$z_i = \mu_i + \sigma_i \cdot \epsilon \quad \text{where } \epsilon \sim \mathcal{N}(0,1) \text{ is not affected by backpropagation}$$



now:



loss: 1) **reconstruction loss L_R** , i. e. the difference between input image x and reconstructed image $p(x|z)$. Note: L_R alone would **lead to overfitting**, exact image gets reproduced!

2) we need a **regularization term** that counteracts L_R :

KL-divergence $L_{KL} = KL[q(z|x)||p(z)]$, which increases if learned distribution $q(z|x)$ gets too far from prior $p(z) \sim \mathcal{N}(0,1)$



loss: 1) **reconstruction loss** L_R , i. e. the difference between input image x and reconstructed image $p(x|z)$. Note: L_R alone would **lead to overfitting**, exact image gets reproduced!

2) we need a **regularization term** that counteracts L_R :

KL-divergence $L_{KL} = KL[q(z|x)||p(z)]$, which increases if learned distribution $q(z|x)$ gets too far from prior $p(z) \sim \mathcal{N}(0,1)$

see [here](#) and **module 6 (variational Bayes)**

check out VAR architecture.ipynb



Berkeley
UNIVERSITY OF CALIFORNIA

Convolution, Image Classification & Segmentation

VAE

check out

VAR architecture.ipynb

original



reconstruction



generated

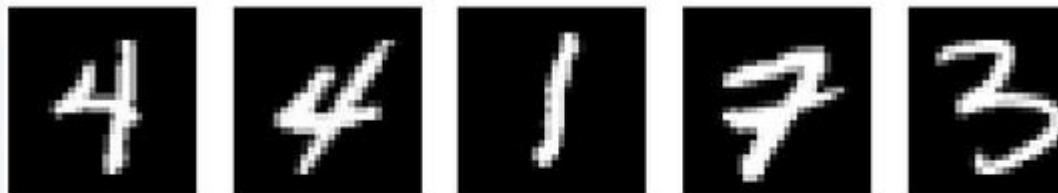




check out

VAR architecture.ipynb

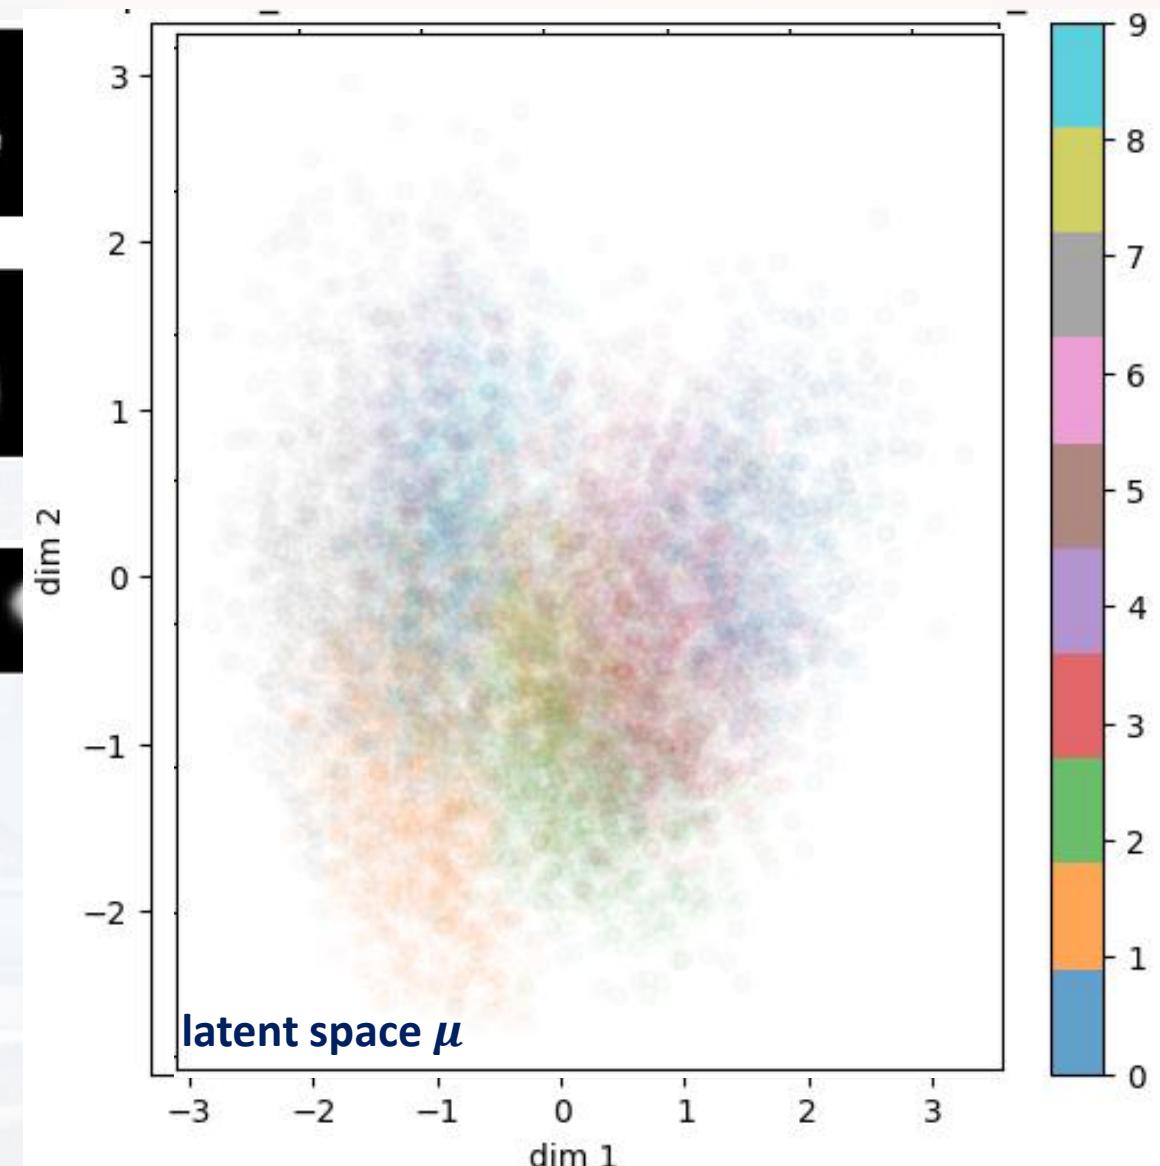
original



reconstruction

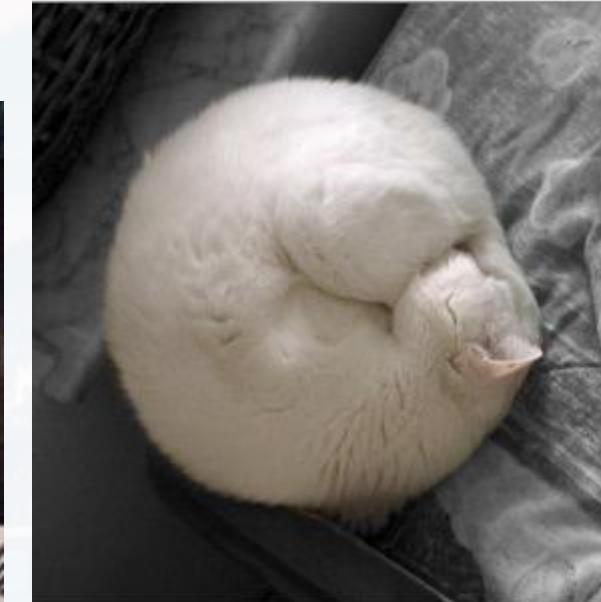


generated





Thank you very much for your attention!



0.40 cat
0.32 frog
0.16 bird
0.06 ship
0.03 dog