

## Lecture 13:

# Long Short-Term Memory Networks (LSTMs) – Part II



Markus Hohle  
University California, Berkeley

Bayesian Data Analysis and  
Machine Learning for Physical  
Sciences



## Course Map

Module 1	Maximum Entropy and Information, Bayes Theorem
Module 2	Naive Bayes, Bayesian Parameter Estimation, MAP
Module 3	MLE, Lin Regression
Module 4	Model selection I: Comparing Distributions
Module 5	Model Selection II: Bayesian Signal Detection
Module 6	Variational Bayes, Expectation Maximization
Module 7	Hidden Markov Models, Stochastic Processes
Module 8	Monte Carlo Methods
Module 9	Machine Learning Overview, Supervised Methods & Unsupervised Methods
Module 10	ANN: Perceptron, Backpropagation, SGD
Module 11	Convolution and Image Classification and Segmentation
Module 12	RNNs and LSTMs
<b>Module 13</b>	<b>RNNs and LSTMs + CNNs</b>
Module 14	Transformer and LLMs
Module 15	Graphs & GNNs



<https://www.analyticsvidhya.com>



## Outline

- LSTM for Classification
- Bidirectional LSTMs
- Stacked LSTMs
- LSTM + CNN



<https://www.analyticsvidhya.com>



## Outline

- LSTM for Classification
- Bidirectional LSTMs
- Stacked LSTMs
- LSTM + CNN



minimal model:

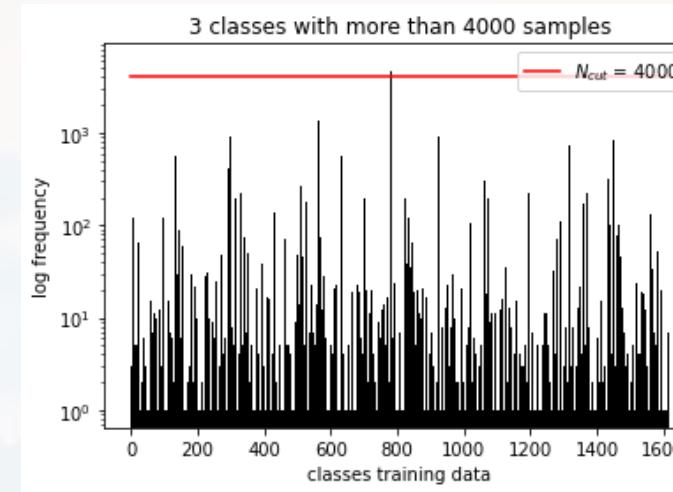
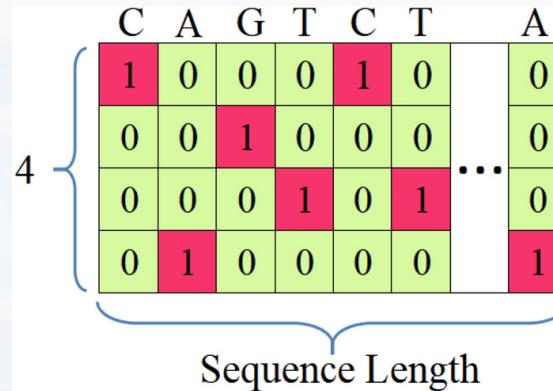
```
[N_samples, LengthSeq, N_features] = X.shape  
[N_samples, N_classes] = Y.shape
```

Y is one-hot encoded

```
model = Sequential()  
  
model.add(LSTM(n_neurons, activation = 'tanh',  
               input_shape = (LengthSeq , N_features)))  
  
model.add(Dense(N_classes, activation = 'softmax'))  
  
opt = optimizers.Adam()  
model.compile(loss = 'categorical_crossentropy', optimizer = opt,  
              metrics = [ 'accuracy' ])  
  
model.summary()
```

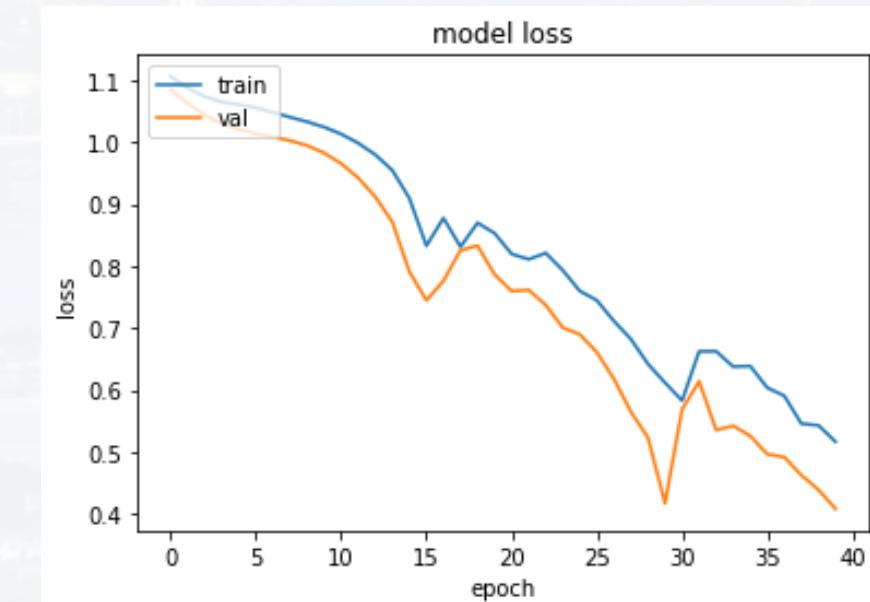
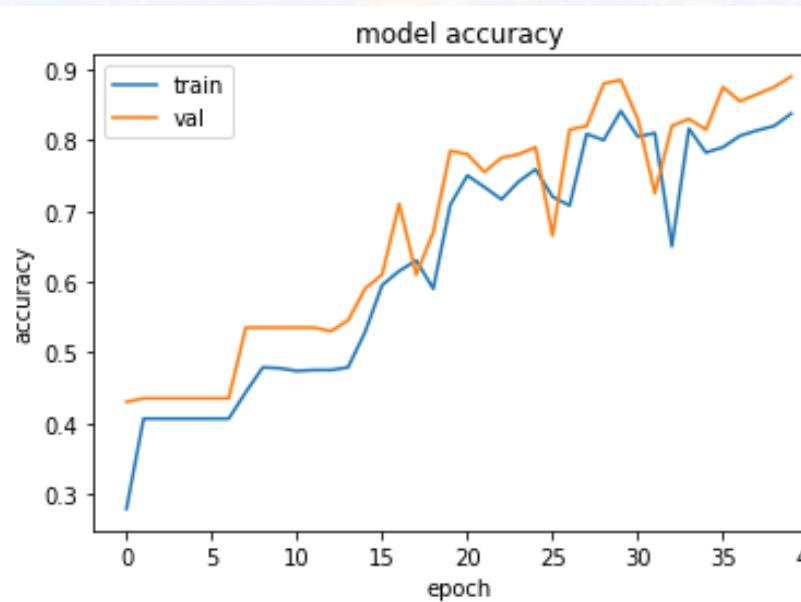


## barcode example



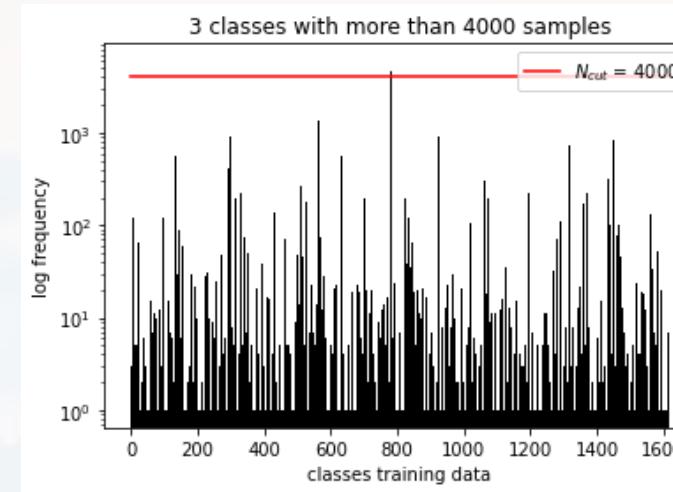
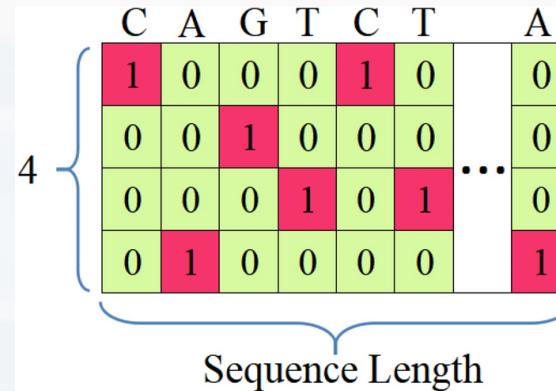
for computational reasons:

- **three** classes
- **1k** samples total
- sequences cut to length **500**



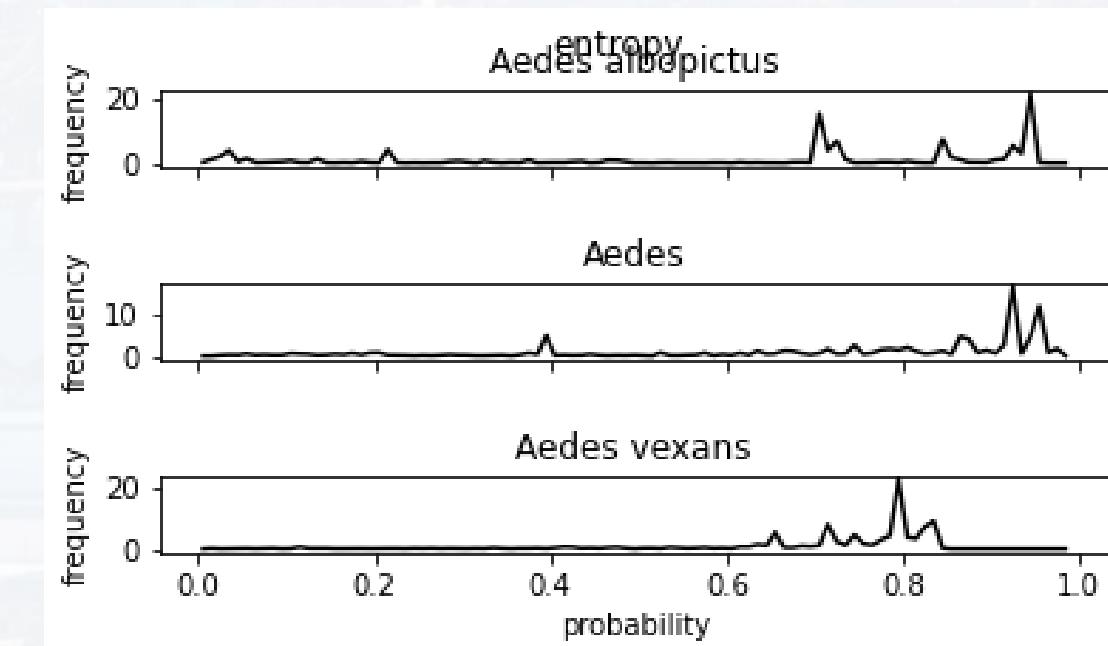
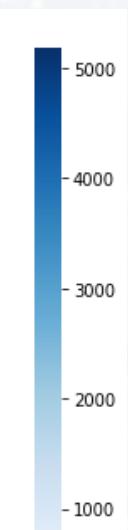
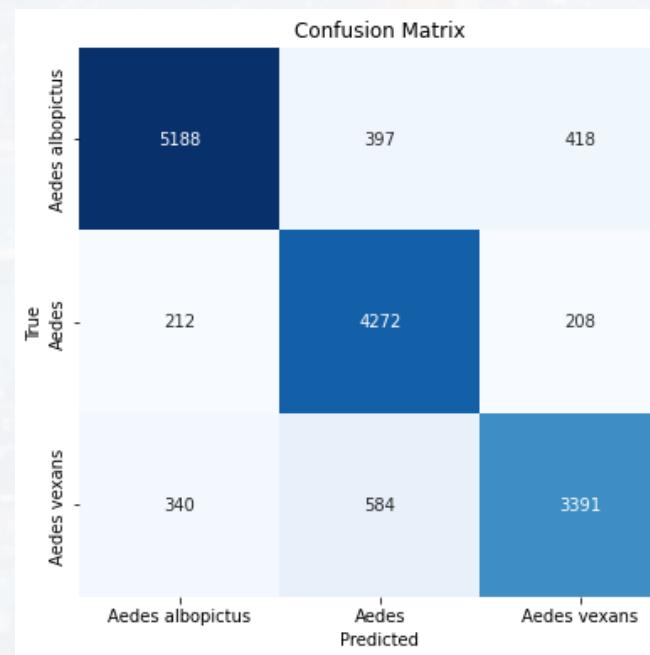


## barcode example



for computational reasons:

- **three** classes
- **1k** samples total
- sequences cut to length **500**





<https://www.analyticsvidhya.com>

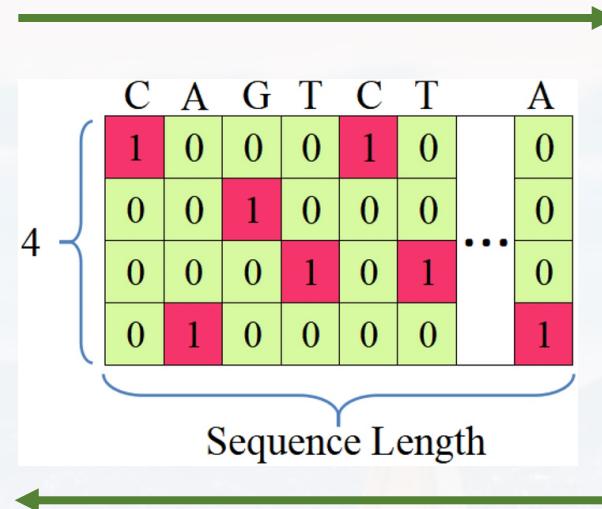


## Outline

- LSTM for Classification
- Bidirectional LSTMs
- Stacked LSTMs
- LSTM + CNN



sometimes, sequences can be read from two directions :



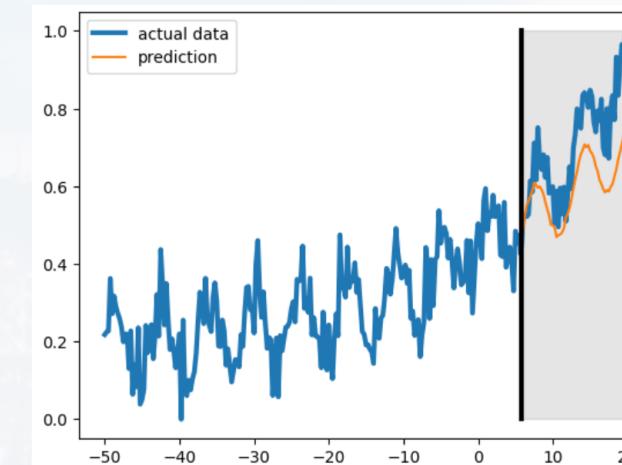
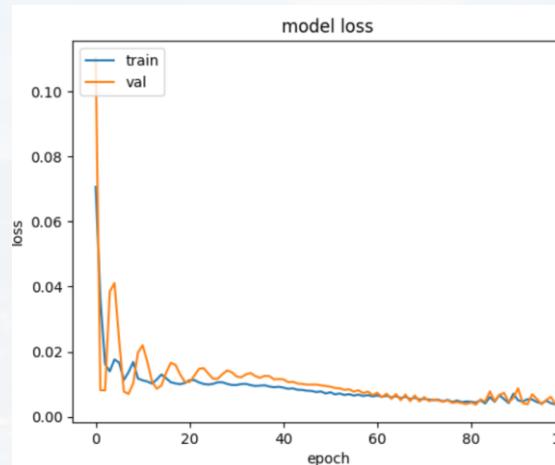
```
from keras.layers import Bidirectional
```

```
model = Sequential()  
model.add(Bidirectional(LSTM(n_neurons, activation = 'tanh'),\n                      input_shape = (dt_past, n_features)))  
model.add(Dense(dt_futu))  
  
opt = optimizers.Adam()  
model.compile(loss = 'mean_squared_error', optimizer = opt)  
  
model.summary()
```

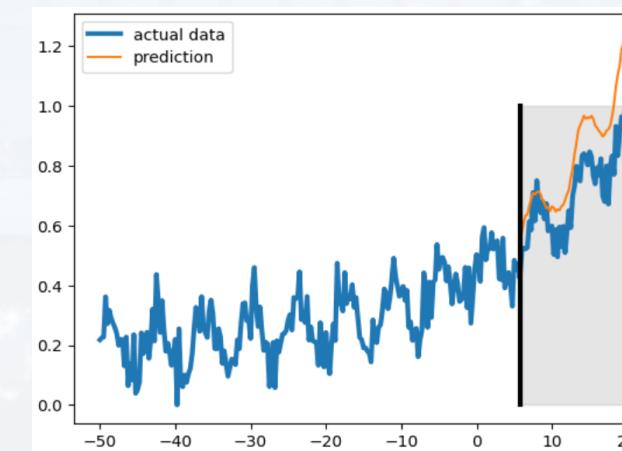
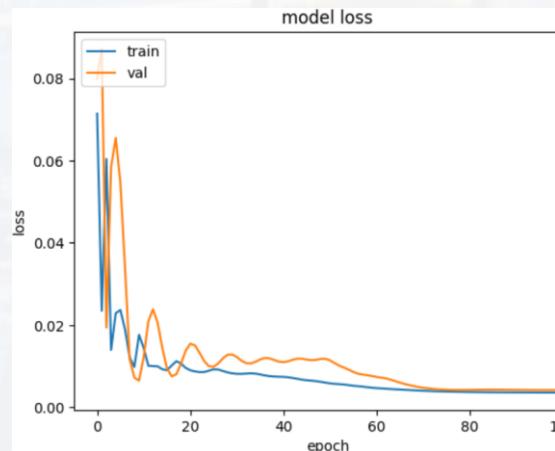


```
model = Sequential()
model.add(Bidirectional(LSTM(n_neurons, activation = 'tanh'), input_shape = (dt_past, n_features)))
model.add(Dense(dt_futu))
opt = optimizers.Adam()
model.compile(loss = 'mean_squared_error', optimizer = opt)
model.summary()
```

vanilla



bidirectional



Layer (type)	Output Shape	Param #
<hr/>		
bidirectional (Bidirectional)	(None, 800)	1286400
<hr/>		
dense_1 (Dense)	(None, 8)	6408
<hr/>		
Total params:	1292808 (4.93 MB)	
Trainable params:	1292808 (4.93 MB)	
Non-trainable params:	0 (0.00 Byte)	



<https://www.analyticsvidhya.com>

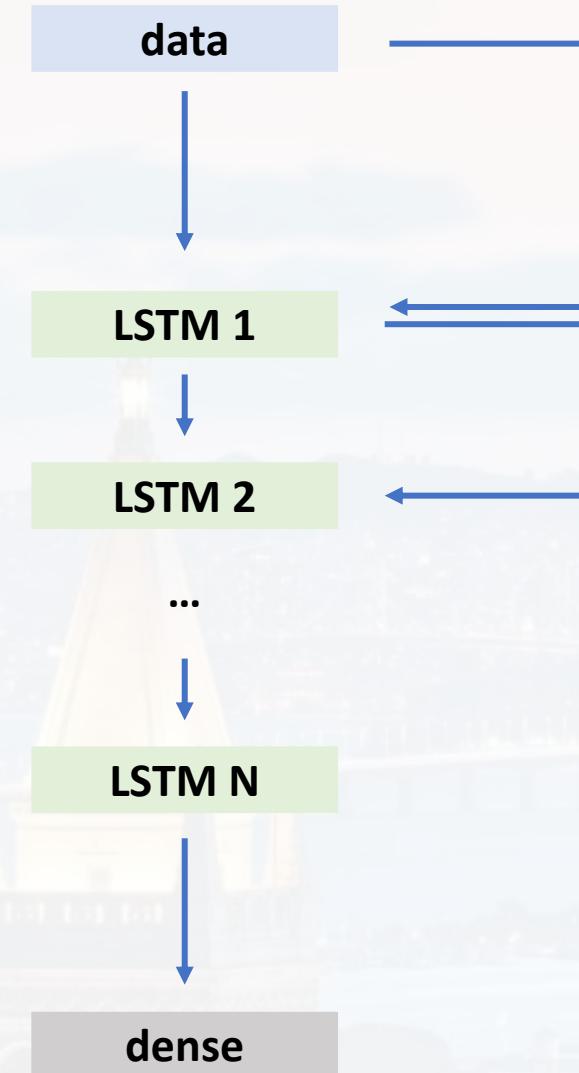


## Outline

- LSTM for Classification
- Bidirectional LSTMs
- Stacked LSTMs
- LSTM + CNN



idea:



shape: (`len(y(t)) - dt_past - dt_futu + 1`  
x `dt_past` x features)

shape: (`len(y(t)) - dt_past - dt_futu + 1`  
x `dt_past` x hidden state)

`return_sequences = True`



data



LSTM 1



LSTM 2

...

LSTM N

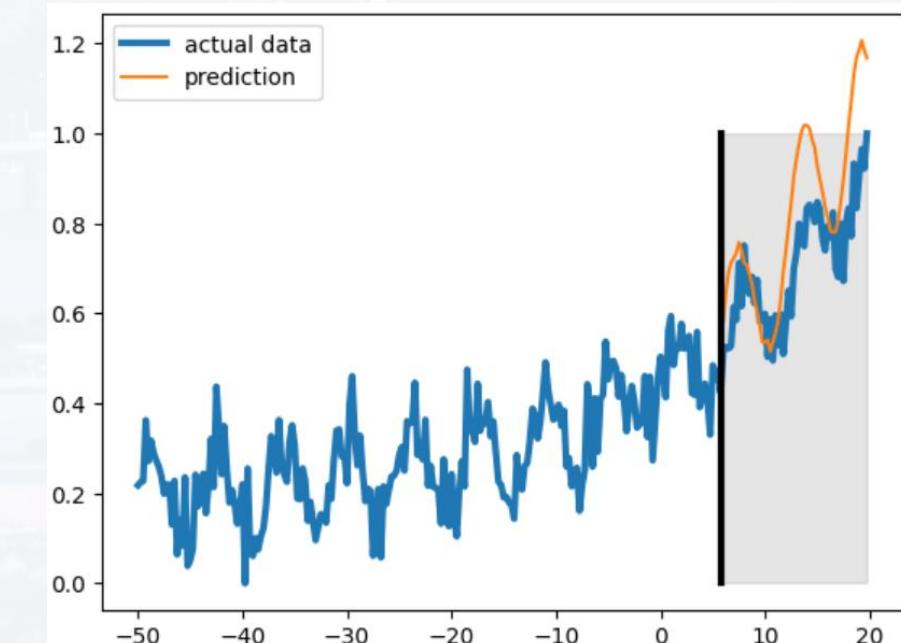
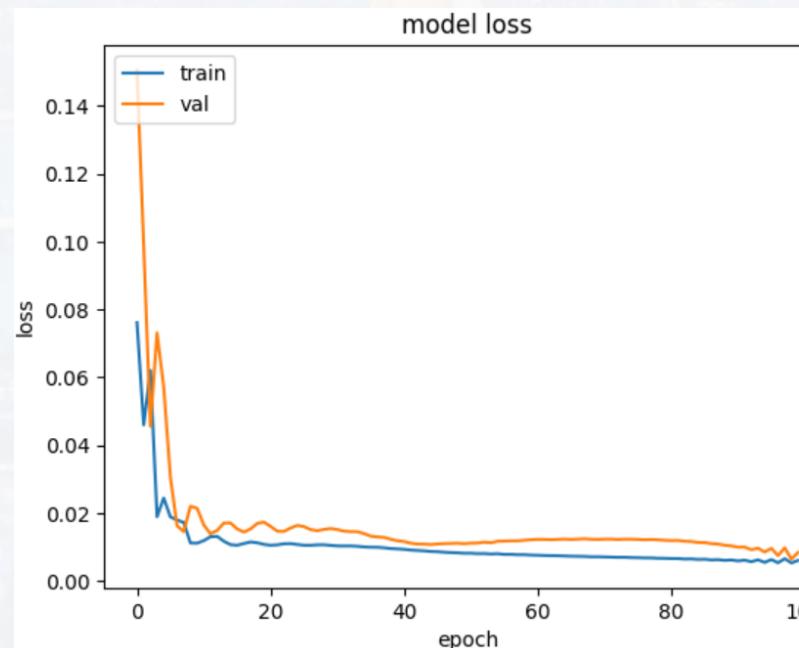


```
model = Sequential()  
  
model.add(LSTM(n_neurons, activation = 'tanh',\  
               return_sequences = True, input_shape = (dt_past, n_features)))  
  
model.add(LSTM(2*n_neurons, activation = 'relu',\  
               return_sequences = True))  
  
model.add(LSTM(n_neurons, activation = 'relu'))  
  
model.add(Dense(dt_futu))  
  
opt = optimizers.Adam()  
model.compile(loss = 'mean_squared_error', optimizer = opt)  
  
model.summary()
```

all LSTMs, **except the last** stack needs  
return\_sequences = True



Layer (type)	Output Shape	Param #
=====		
lstm_2 (LSTM)	(None, 20, 400)	643200
lstm_3 (LSTM)	(None, 20, 800)	3843200
lstm_4 (LSTM)	(None, 400)	1921600
dense_2 (Dense)	(None, 8)	3208
=====		
Total params: 6411208 (24.46 MB)		
Trainable params: 6411208 (24.46 MB)		
Non-trainable params: 0 (0.00 Byte)		





<https://www.analyticsvidhya.com>

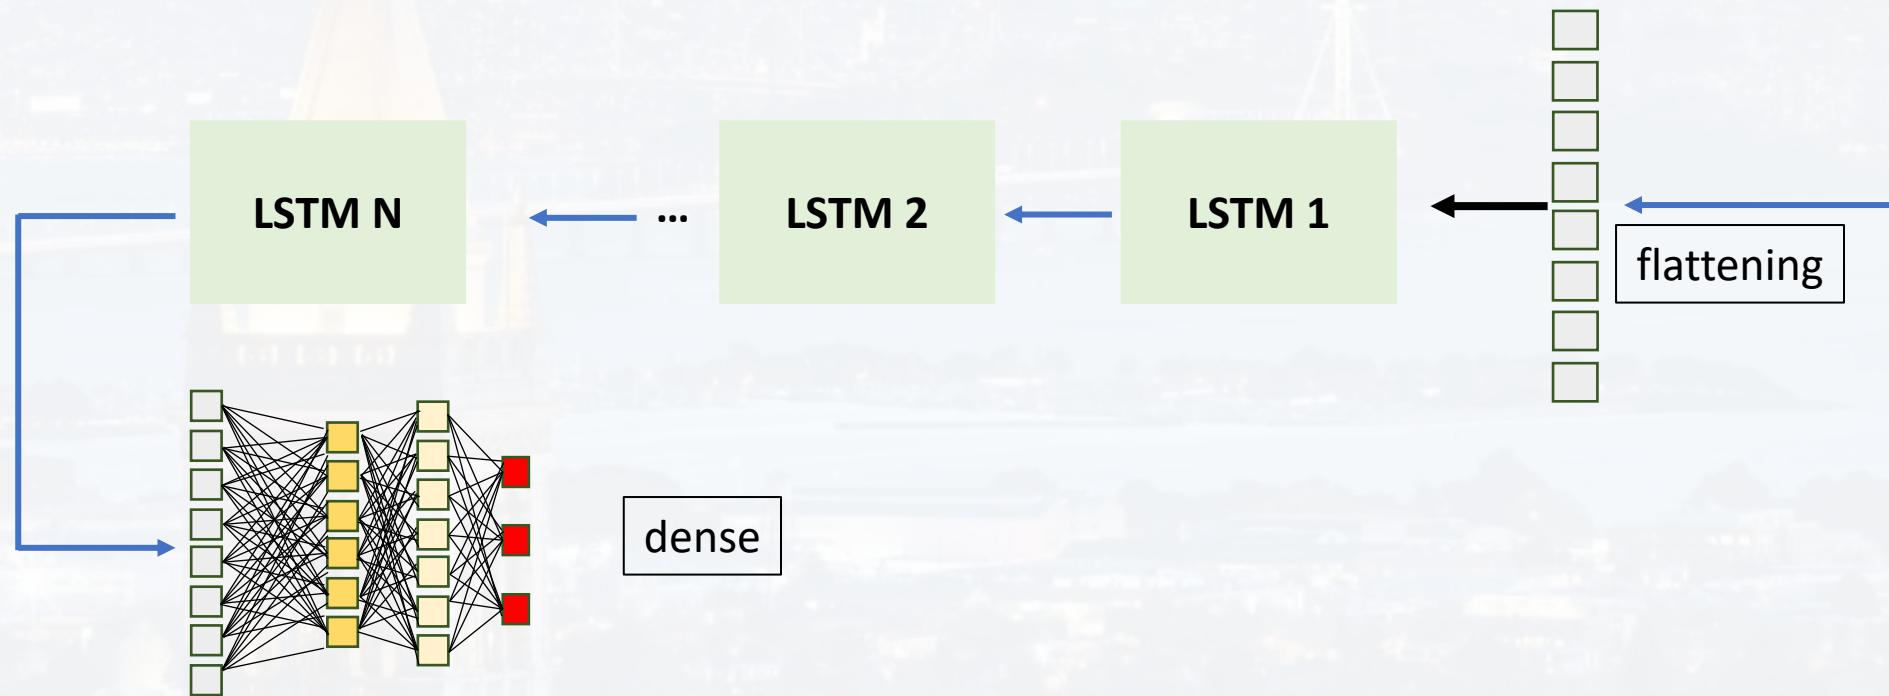
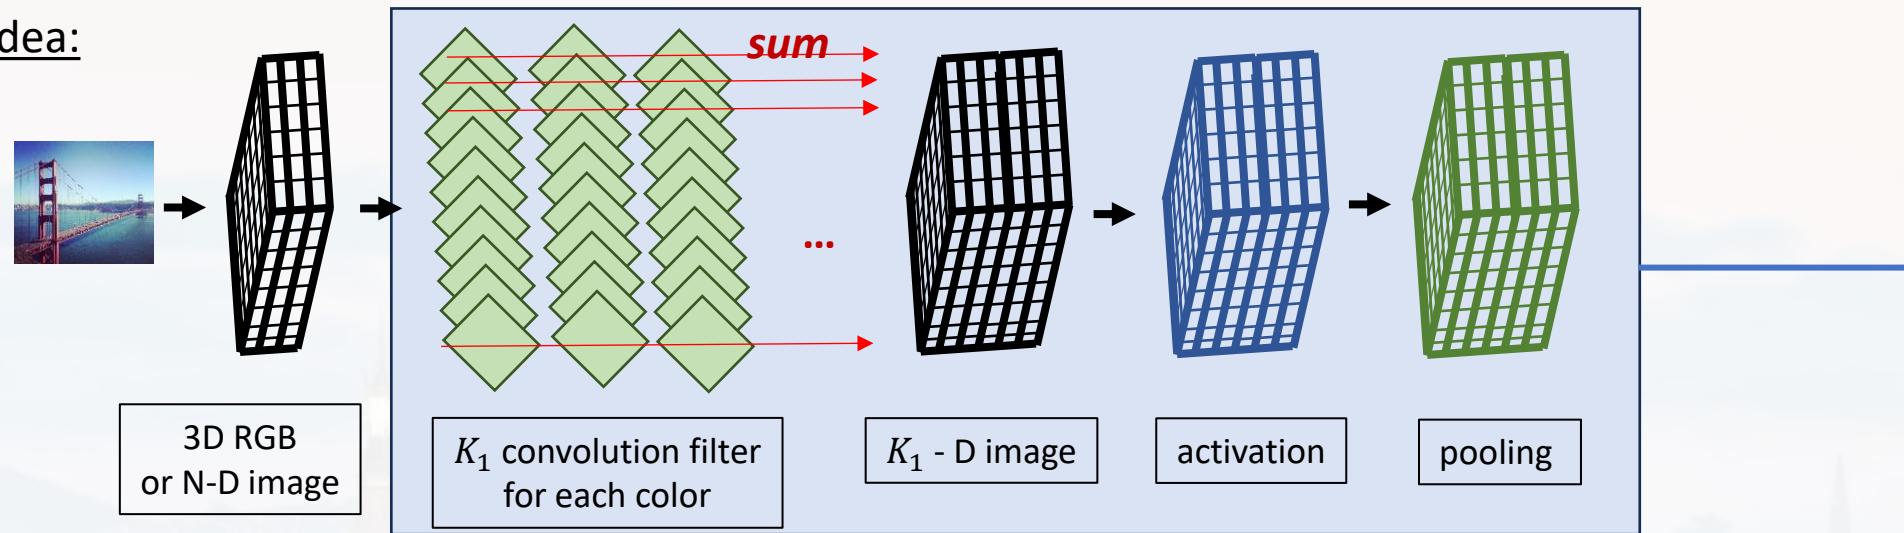


## Outline

- LSTM for Classification
- Bidirectional LSTMs
- Stacked LSTMs
- LSTM + CNN



idea:

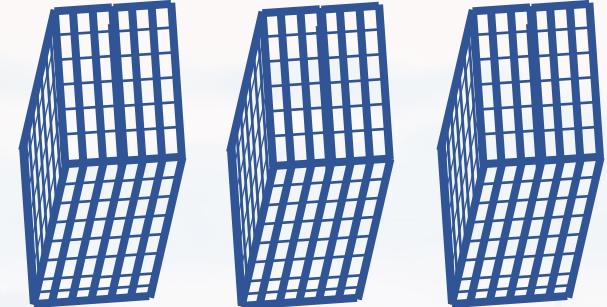




idea:

input expected by CNN (images):

$(N\_images, N\_x, N\_y, N\_color)$

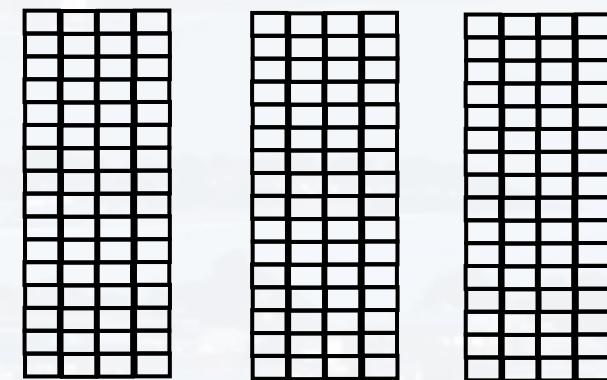


input expected by CNN (videos):

$(N\_videos, N\_t, N\_x, N\_y, N\_color)$

input expected by LSTM (sequences):

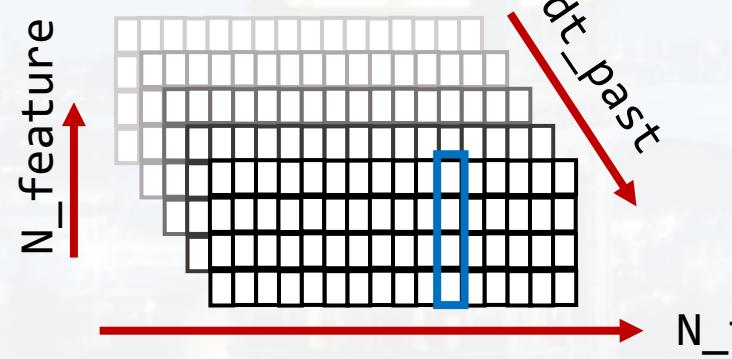
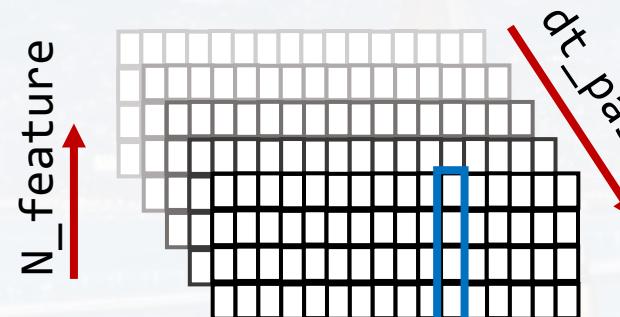
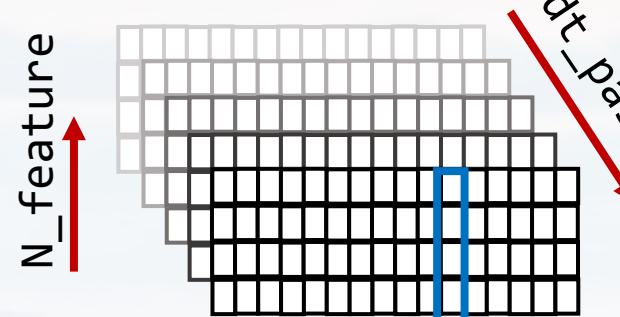
$(N\_sequences, N\_t, N\_feature)$



None



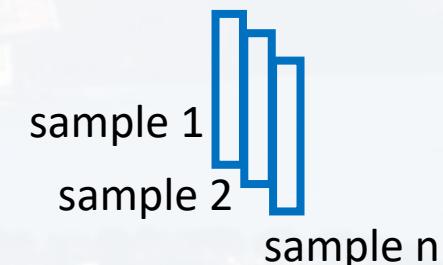
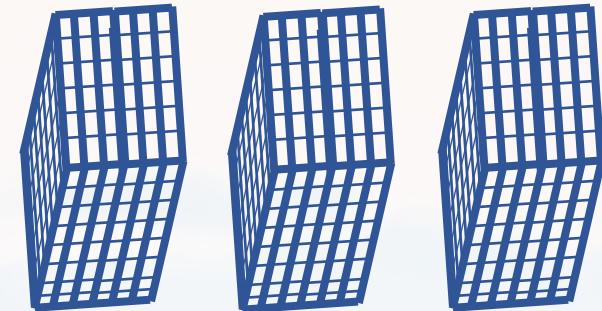
$(N_{\text{images}}, N_x, N_y, N_{\text{color}})$



sample 1

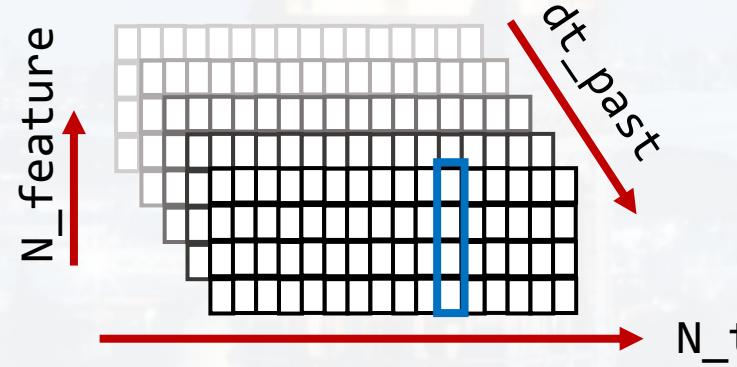
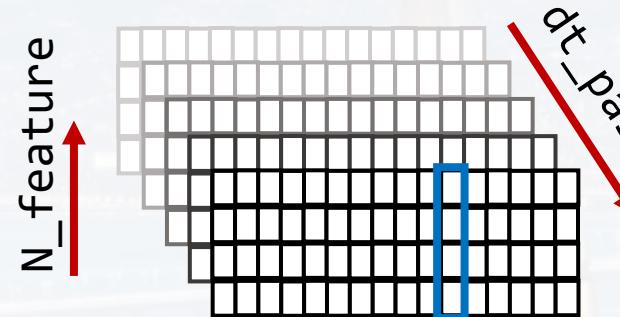
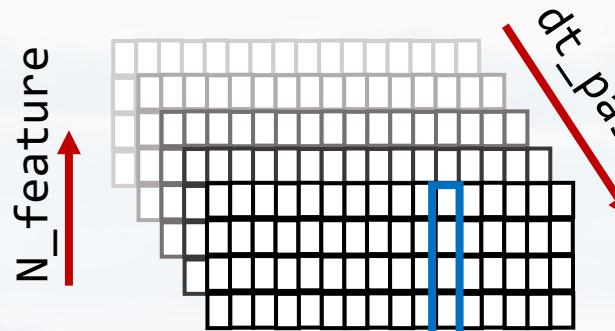
sample 2

sample n





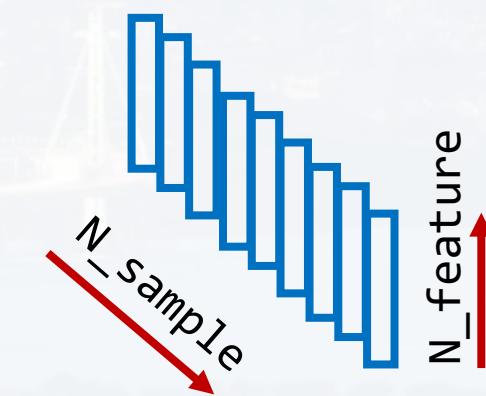
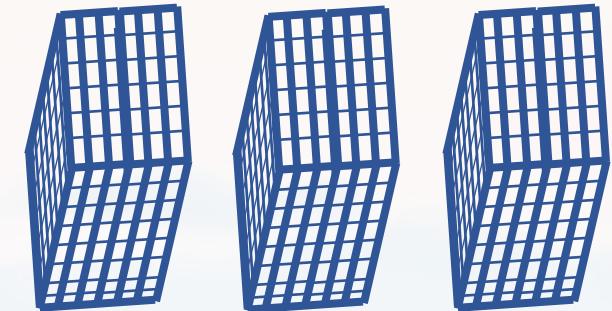
$(N_{\text{images}}, N_x, N_y, N_{\text{color}})$



sample 1

sample 2

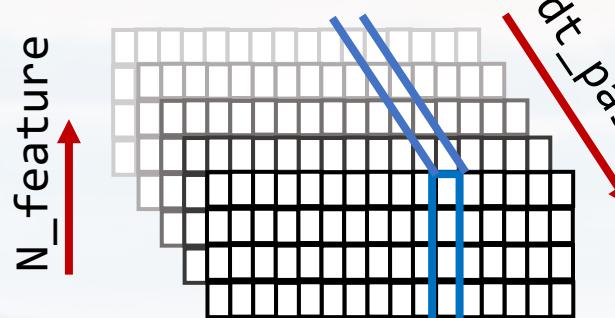
sample n



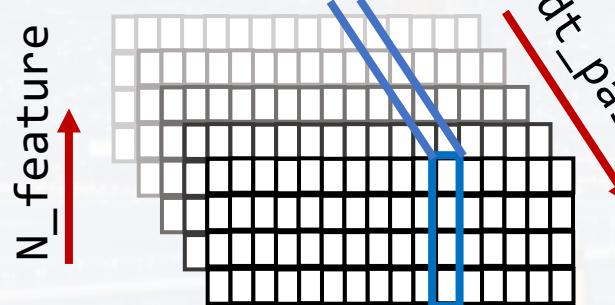
for **one** timepoint t



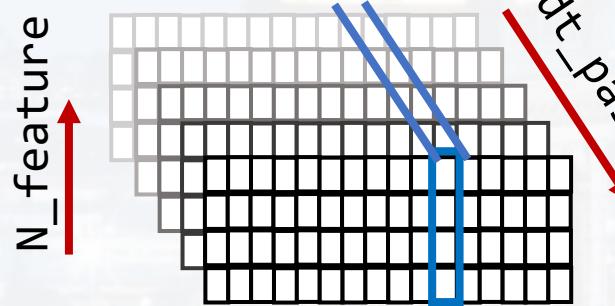
$(N_{\text{images}}, N_x, N_y, N_{\text{color}})$



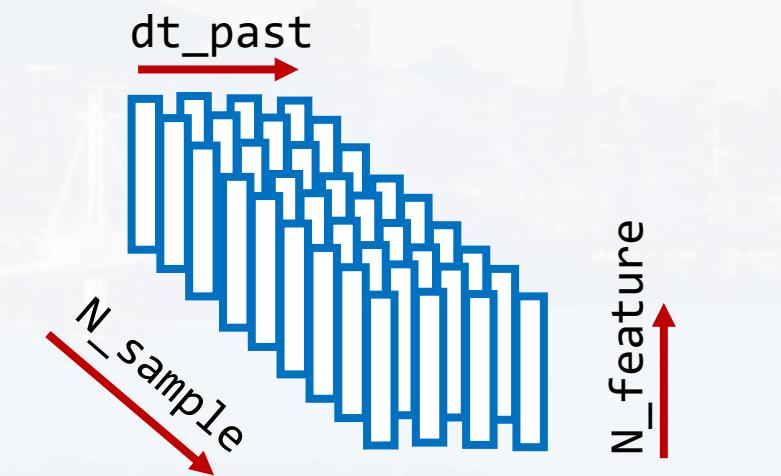
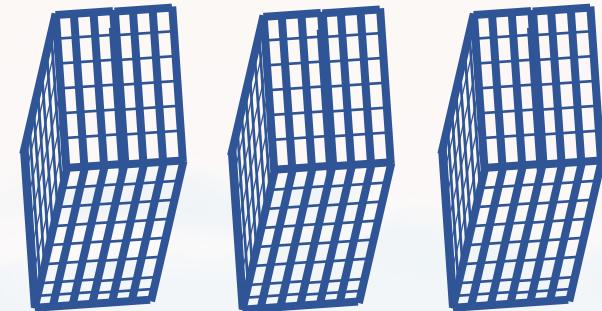
sample 1



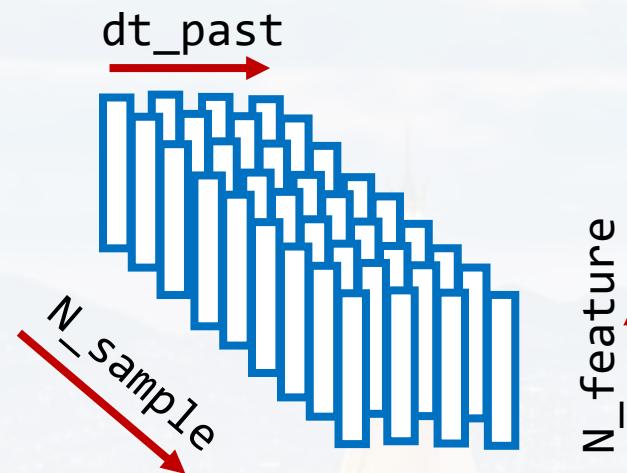
sample 2



sample n

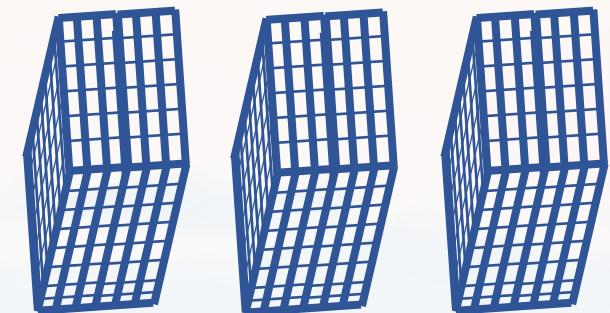


for **one** timepoint t



(*N\_images*, *N\_x*, *N\_y*, *N\_color*)

for **one** timepoint t



regression: one sample of *N\_features* and *dt\_past*

```
X = X.reshape((X.shape[0], N_samples, dt_past, N_feature))
```



```
X = X.reshape((X.shape[0], N_samples, dt_past, N_feature))
```

```
model = Sequential()
```

```
model.add(TimeDistributed(Conv1D(filters = 64, kernel_size = 3,\n                           activation = 'relu').,\n                           input_shape = (None, dt_past, N_feature))))
```

```
model.add(TimeDistributed(MaxPooling1D(pool_size = 2)))
```

```
model.add(TimeDistributed(Flatten()))
```

```
model.add(LSTM(n_neurons, input_shape = (dt_past, N_feature),\n               activation = 'tanh'))
```

```
model.add(Dense(dt_futu))
```

```
opt = optimizers.Adam()
```

```
model.compile(loss = 'mean_squared_error', optimizer = opt)
```

```
model.summary()
```

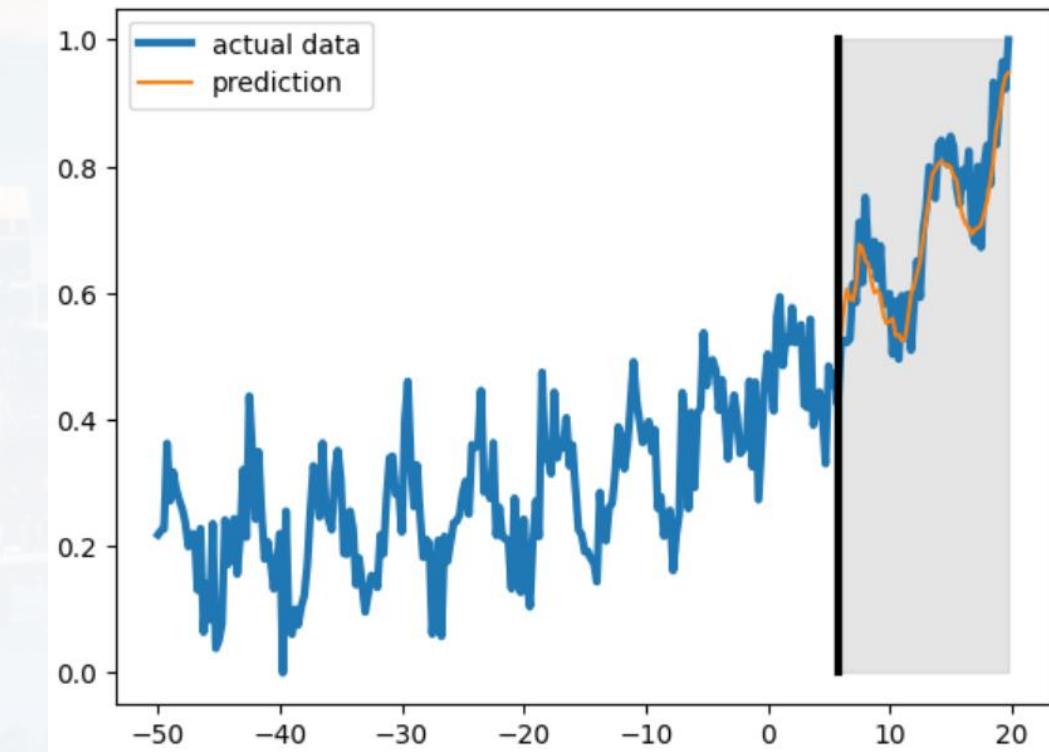
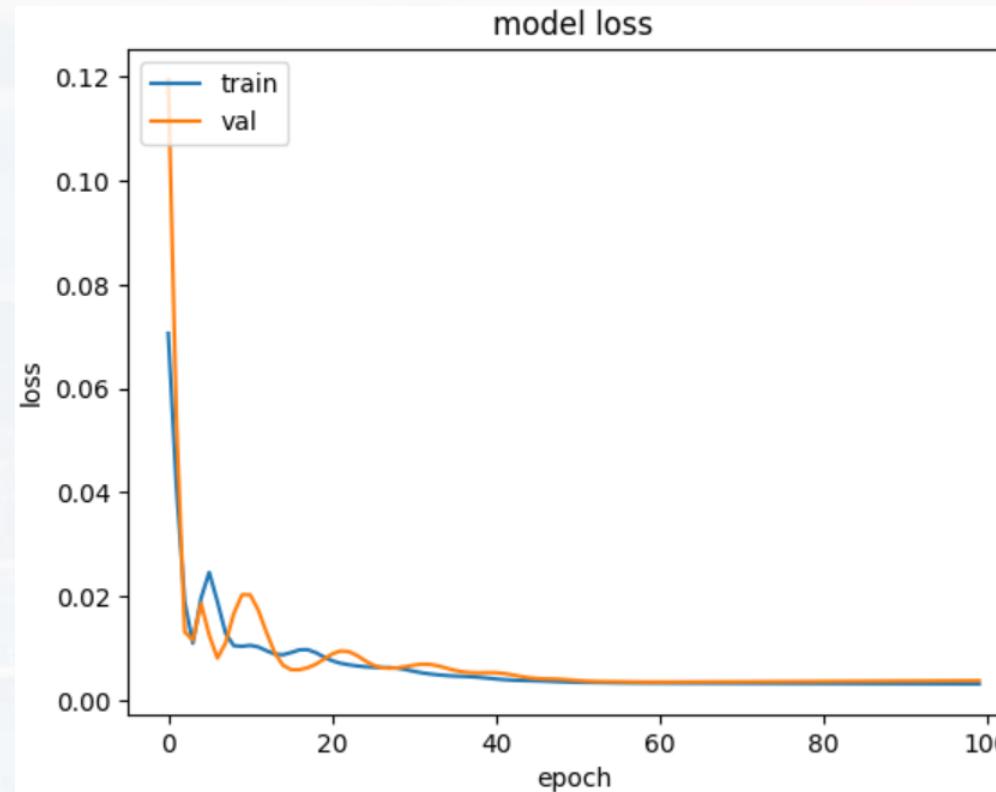
1D filter along time coordinate

takes care of  
maintaining  
matrix orientation

actual input is (None, None, dt\_past, N\_feature)

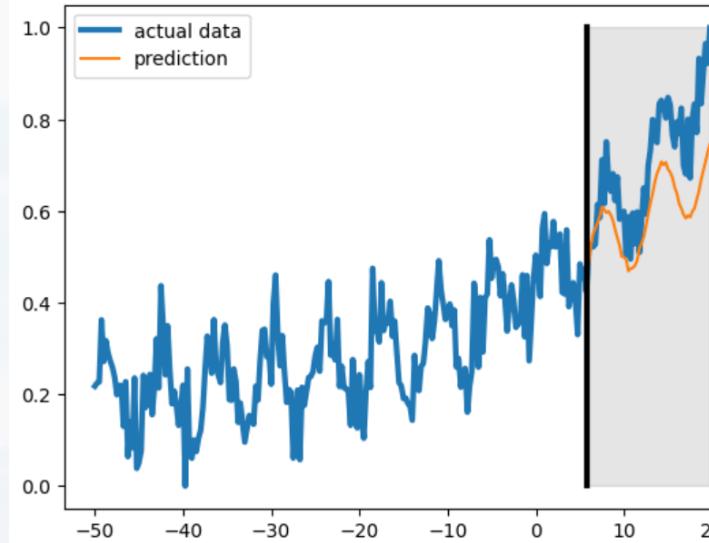


Layer (type)	Output Shape	Param #
<hr/>		
time_distributed (TimeDistributed)	(None, None, 18, 64)	256
		actual input is (None, None, dt_past, N_feature)
time_distributed_1 (TimeDistributed)	(None, None, 9, 64)	0
time_distributed_2 (TimeDistributed)	(None, None, 576)	0
lstm_5 (LSTM)	(None, 400)	1563200
dense_3 (Dense)	(None, 8)	3208
<hr/>		
Total params:	1566664 (5.98 MB)	
Trainable params:	1566664 (5.98 MB)	
Non-trainable params:	0 (0.00 Byte)	

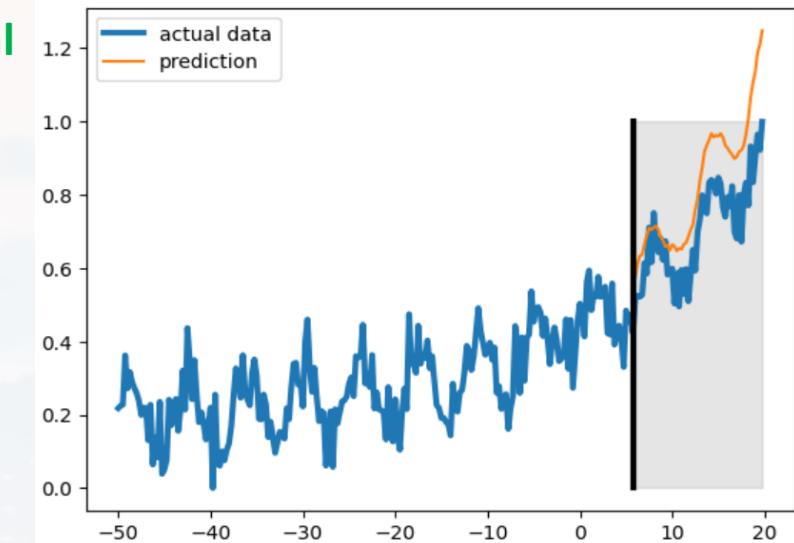




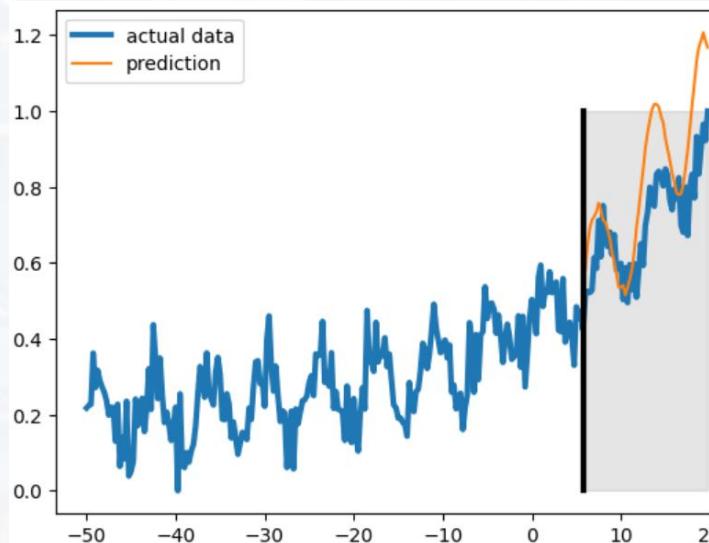
vanilla



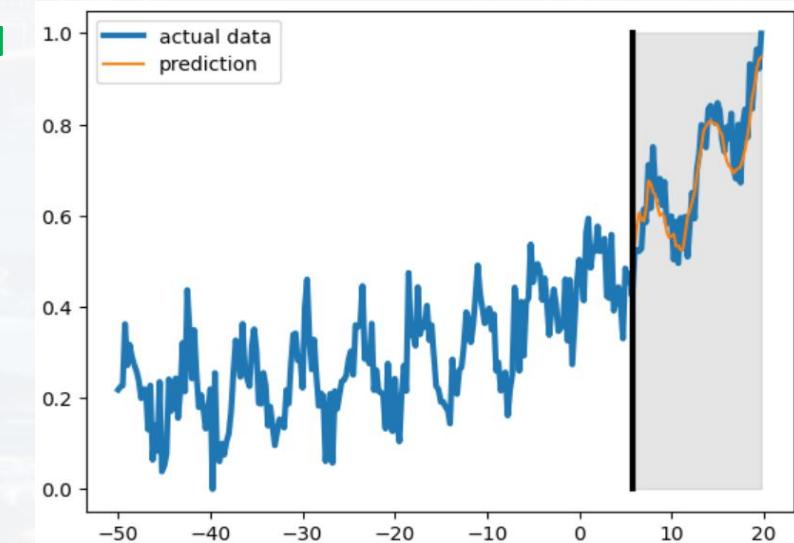
bidirectional



stacked



LSTM + CNN





classification: N samples of N\_features and dt\_past = Length\_Seq

```
[N_sample, LengthSeq, N_features] = X.shape  
model = Sequential()  
model.add(Conv1D(filters = 64, kernel_size = 3, activation = 'relu', \  
                 input_shape = (LengthSeq, N_features)))
```

1D filter along time coordinate  
= LengthSeq

```
model.add(MaxPooling1D(pool_size = 2))
```

```
model.add(LSTM(n_neurons, activation = 'tanh'))
```

```
model.add(Dense(Nclass, activation = 'softmax'))
```

```
opt = optimizers.Adam()  
model.compile(loss = 'categorical_crossentropy', optimizer = opt, \  
              metrics = ['accuracy'])
```

```
model.summary()
```



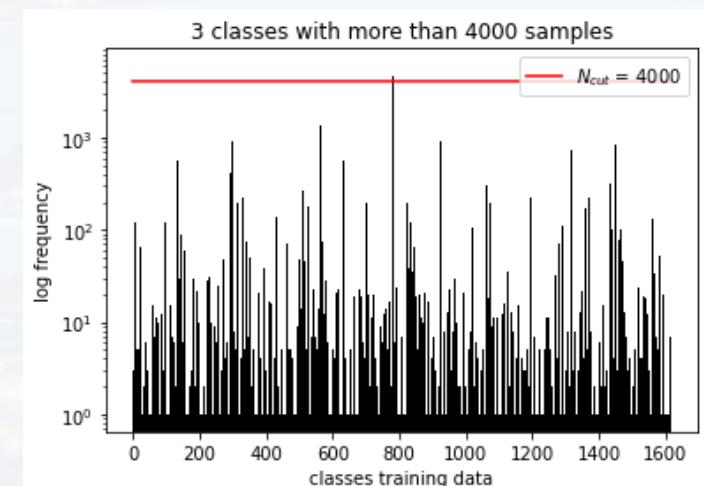
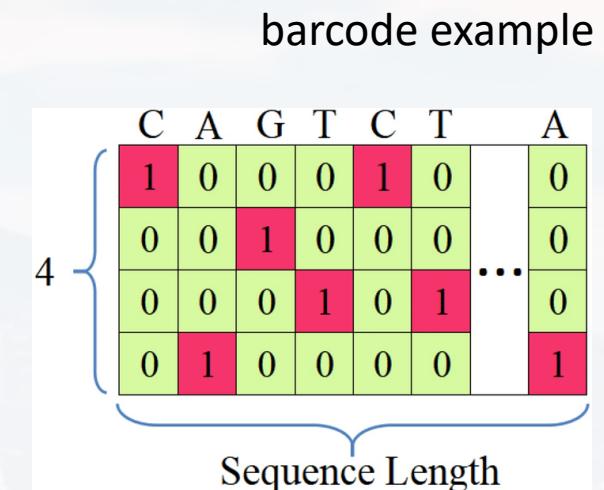
classification: N samples of N\_features and dt\_past = Length\_Seq

[N\_sample, LengthSeq, N\_features] = X.shape

Layer (type)	Output Shape	Param #
<hr/>		
conv1d_8 (Conv1D)	(None, 498, 64)	832
max_pooling1d_7 (MaxPooling1D)	(None, 249, 64)	0
lstm_7 (LSTM)	(None, 100)	66000
dense_4 (Dense)	(None, 3)	303
<hr/>		
Total params: 67135 (262.25 KB)		
Trainable params: 67135 (262.25 KB)		
Non-trainable params: 0 (0.00 Byte)		

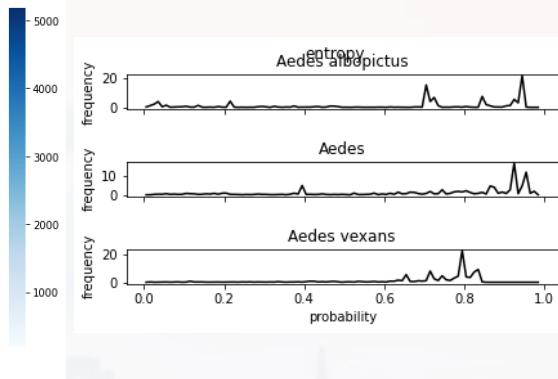
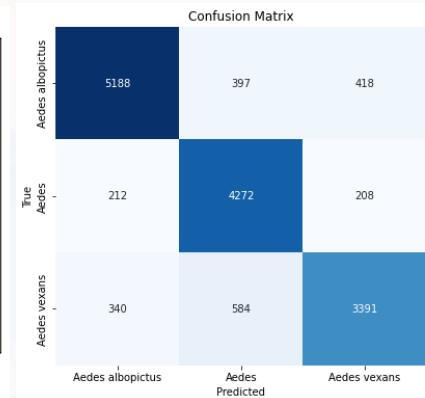
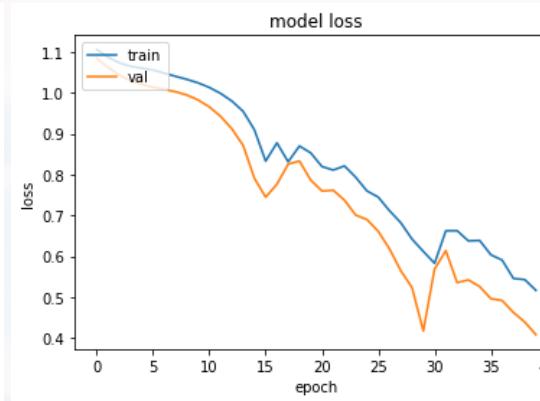
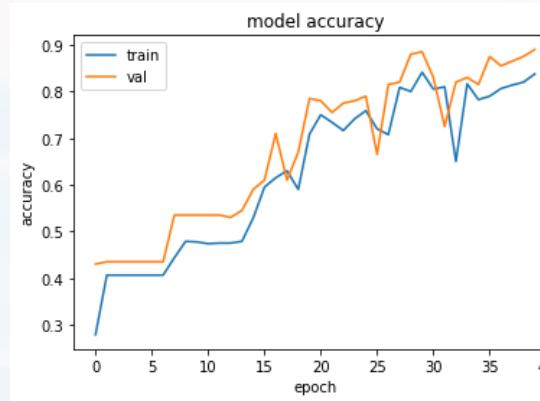
for computational reasons:

- **three** classes
- **1k** samples total
- sequences cut to length **500**

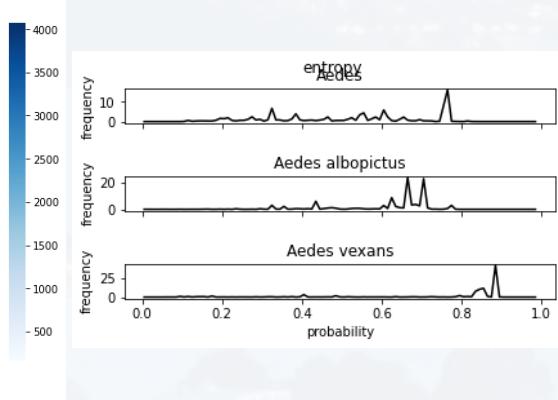
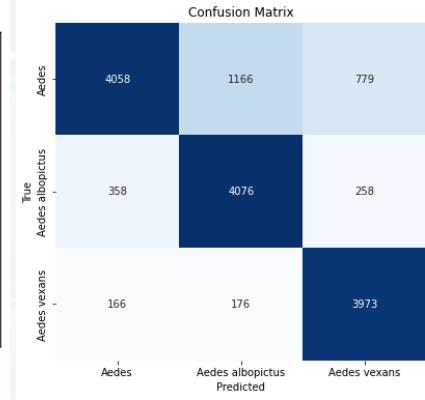
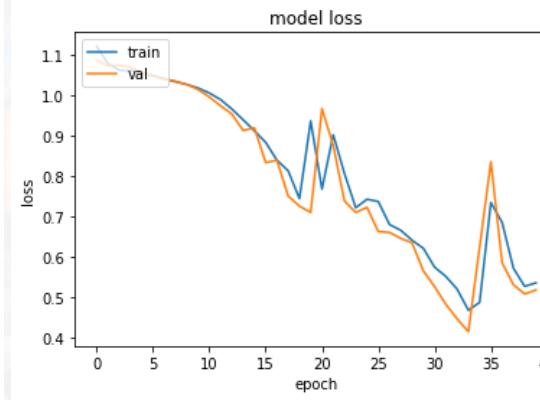
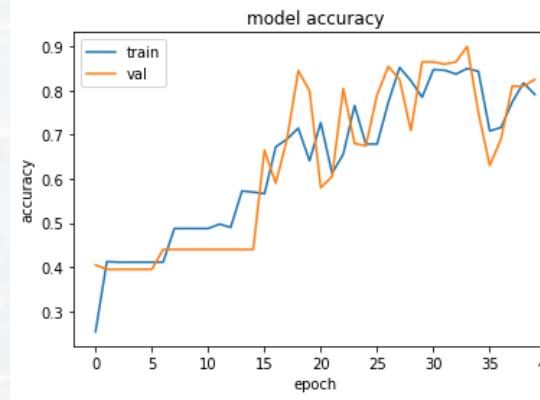




LSTM



LSTM+CNN





Thank you for your attention!

