

Lecture 15:

Graph Neural Networks (GNN)



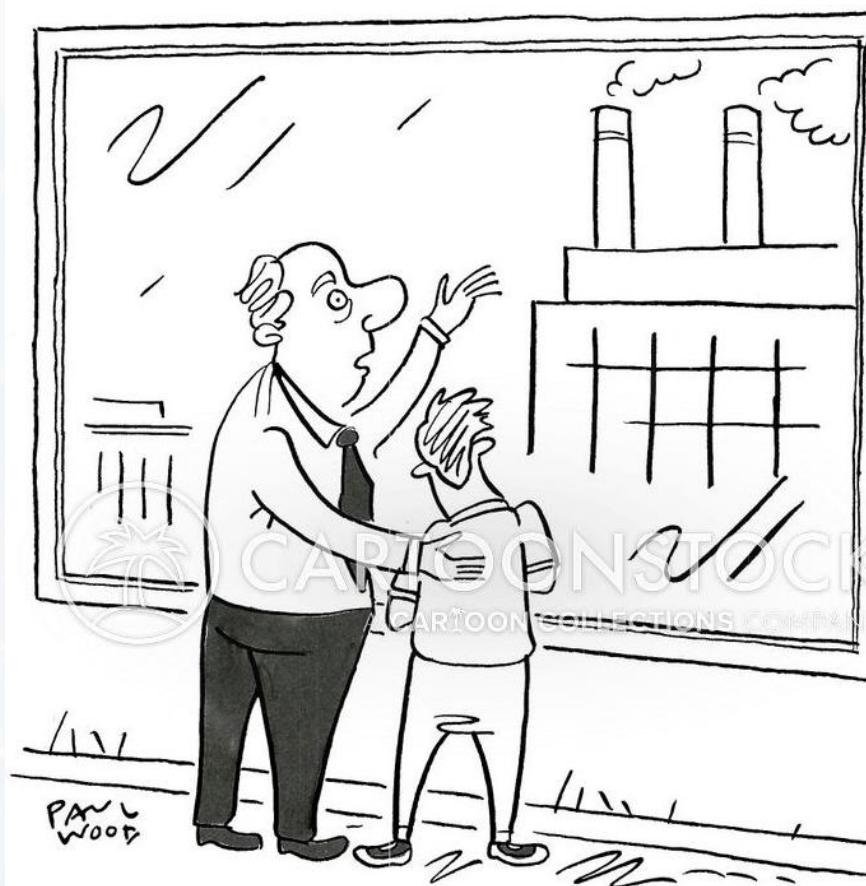
Markus Hohle
University California, Berkeley

**Bayesian Data Analysis and
Machine Learning for Physical
Sciences**



Course Map

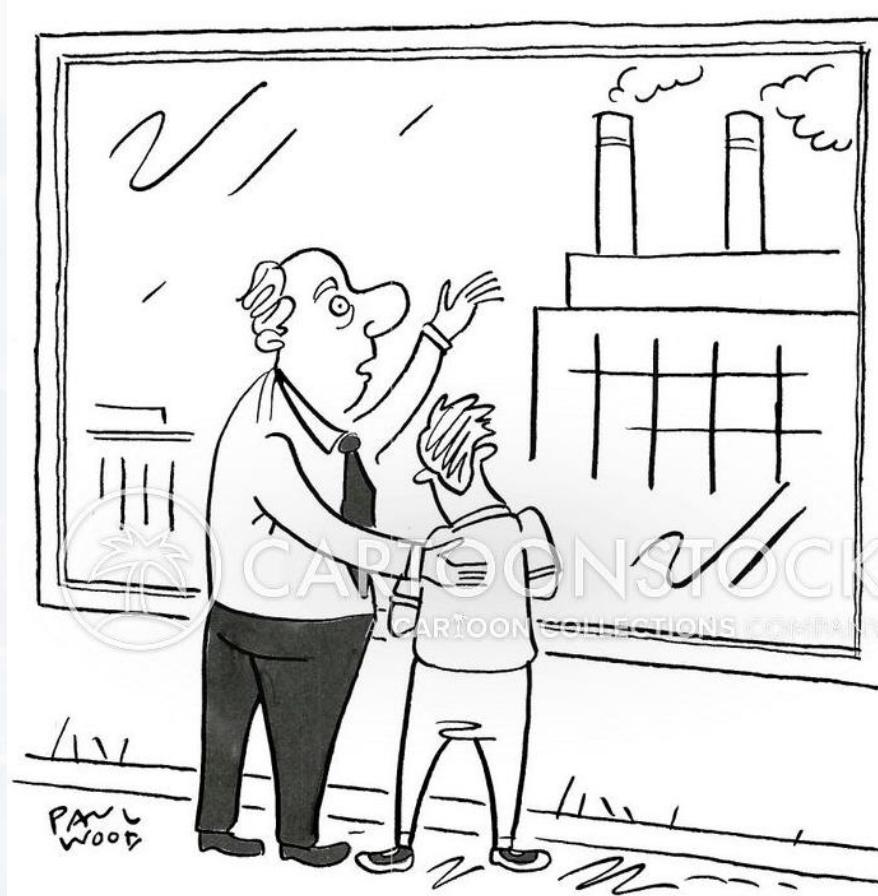
Module 1	Maximum Entropy and Information, Bayes Theorem
Module 2	Naive Bayes, Bayesian Parameter Estimation, MAP
Module 3	MLE, Lin Regression
Module 4	Model selection I: Comparing Distributions
Module 5	Model Selection II: Bayesian Signal Detection
Module 6	Variational Bayes, Expectation Maximization
Module 7	Hidden Markov Models, Stochastic Processes
Module 8	Monte Carlo Methods
Module 9	Machine Learning Overview, Supervised Methods & Unsupervised Methods
Module 10	ANN: Perceptron, Backpropagation, SGD
Module 11	Convolution and Image Classification and Segmentation
Module 12	RNNs and LSTMs
Module 13	RNNs and LSTMs + CNNs
Module 14	Transformer and LLMs
Module 15	Graphs & GNNs



ONE DAY SON, ALL THIS
WILL BE RUN BY ROBOTS

Outline

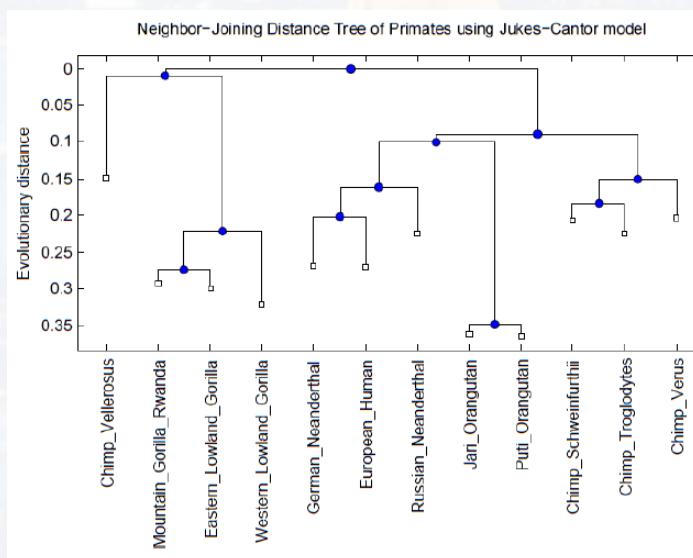
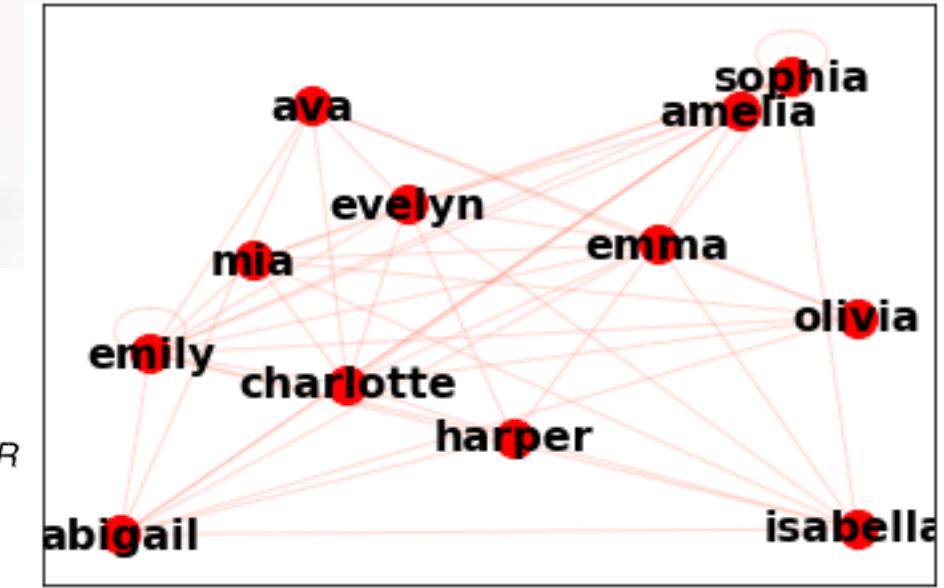
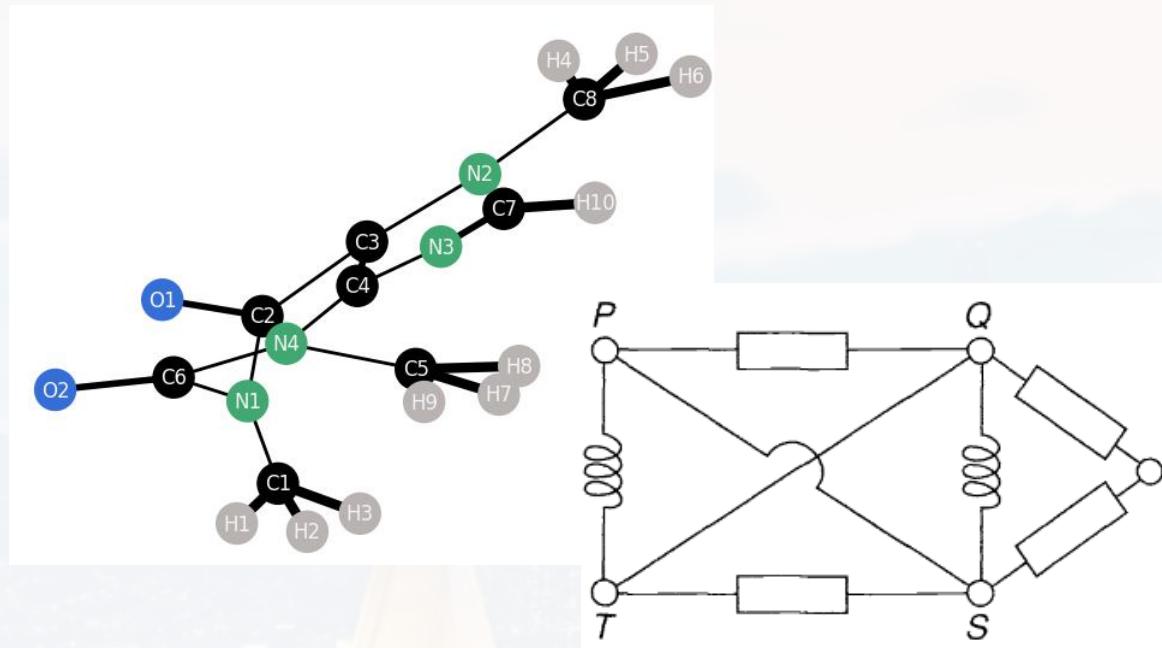
- What is a Graph
- The ANN Part
- PyTorch Example



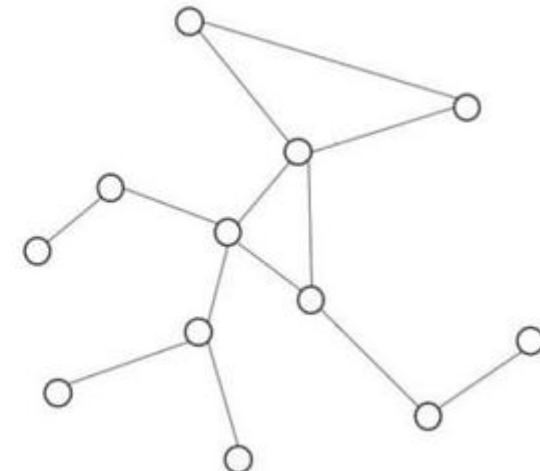
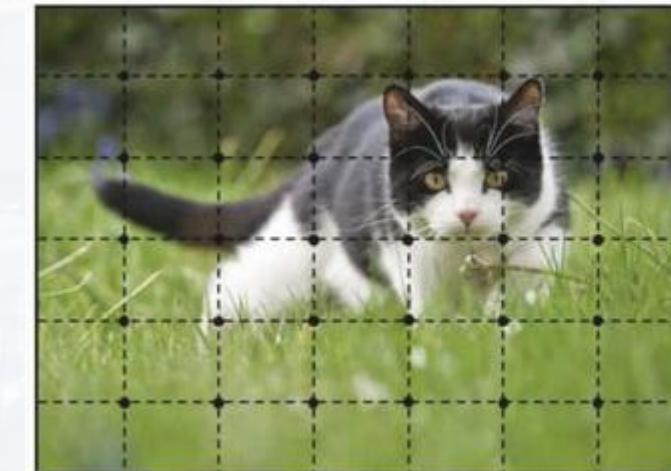
ONE DAY SON, ALL THIS
WILL BE RUN BY ROBOTS

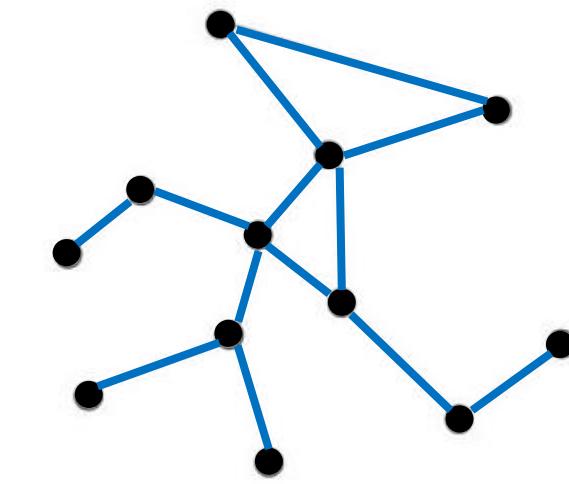
Outline

- What is a Graph
- The ANN Part
- PyTorch Example



<https://doi.org/10.1016/j.aiopen.2021.01.001>





Graph G

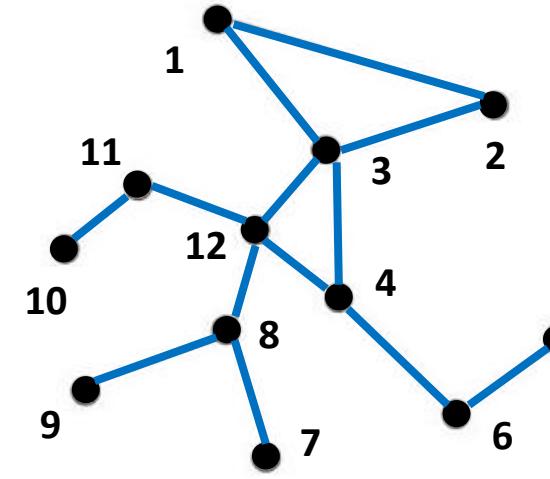
nodes N (vertices V)

edges E

$G = G(N, E)$

- social networks
- street maps
- workflows/planning
- biological signal pathways
- image processing
- diffusion processes

- nodes can have **features**
molecules: mass/ electronegativity
people: age, income, sex, ...
- edges can have **attributes**
molecules: bond length/strength
people: relations (work, friend, family)



structural information: **adjacency matrix A**

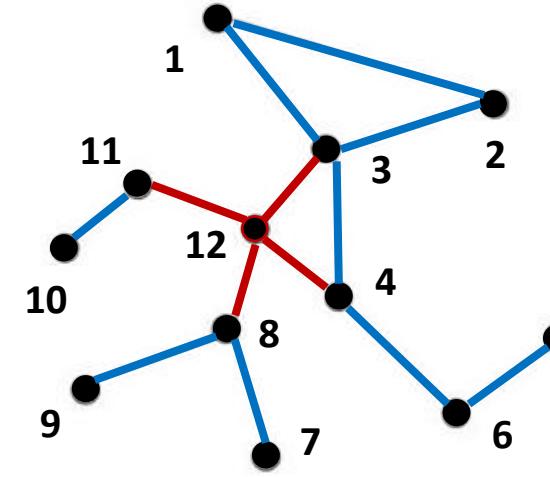
$A_{ij} = 1$ if $(n_i, n_j) \in E$
(nodes n_i and n_j have a common edge)

$A_{ij} = 0$ else

$$A = \begin{pmatrix} 0 & 1 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 1 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 1 & 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 1 & 1 & 0 & 0 & 0 & 1 & 0 & 0 & 1 & 0 \end{pmatrix}$$

Graph $G = G(N, E)$

nodes N (vertices V)
edges E



structural information: **adjacency matrix A**

$A_{ij} = 1$ if $(n_i, n_j) \in E$
(nodes n_i and n_j have a common edge)

$A_{ij} = 0$ else

node 12 has four first degree neighbors

degree d of a node

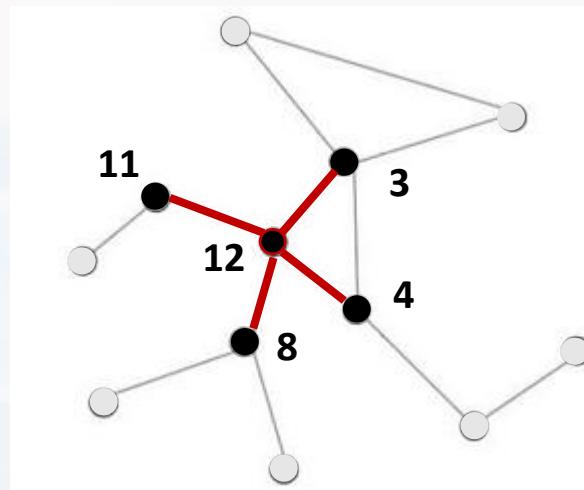
$$d(n_i) = \sum_j A_{ij}$$

$$A = \begin{pmatrix} 0 & 1 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 1 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 1 & 0 & 0 & 1 & 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 \end{pmatrix}$$

 $\boxed{\begin{matrix} 0 & 0 & \textcolor{red}{1} & \textcolor{red}{1} & 0 & 0 & 0 & \textcolor{red}{1} & 0 & 0 & \textcolor{red}{1} & 0 \end{matrix}}$

Graph $G = G(N, E)$

nodes N (vertices V)
edges E



structural information: **adjacency matrix A**

$A_{ij} = 1$ if $(n_i, n_j) \in E$
(nodes n_i and n_j have a common edge)

$A_{ij} = 0$ else

node 12 has four first degree neighbors

degree d of a node

$$d(n_i) = \sum_j A_{ij}$$

first degree neighborhood \mathcal{N}

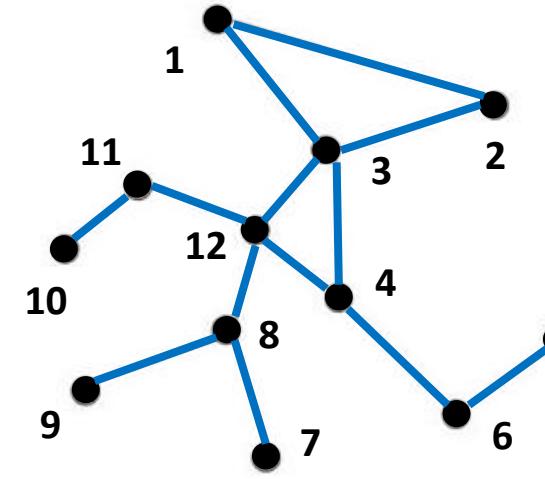
$$\mathcal{N}(n_i) = \{n_j \in N: (n_i, n_j) \in E\}$$

$$A = \begin{pmatrix} 0 & 1 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 1 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 1 & 0 & 0 & 1 & 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 1 \end{pmatrix}$$

$$\begin{array}{cccccccccc} 0 & 0 & \boxed{1} & \boxed{1} & 0 & 0 & 0 & \boxed{1} & 0 & 0 & \boxed{1} & 0 \end{array}$$

Graph $G = G(N, E)$

nodes N (vertices V)
edges E



A graph can have **loops**

structural information: **adjacency matrix A**

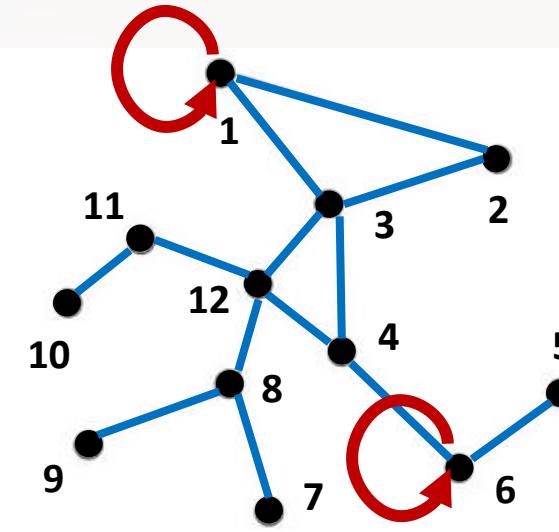
$A_{ij} = 1$ if $(n_i, n_j) \in E$
(nodes n_i and n_j have a common edge)

$A_{ij} = 0$ else

$$A = \begin{pmatrix} 0 & 1 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 1 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 \end{pmatrix}$$

Graph $G = G(N, E)$

nodes N (vertices V)
edges E



A graph can have **loops**

structural information: **adjacency matrix A**

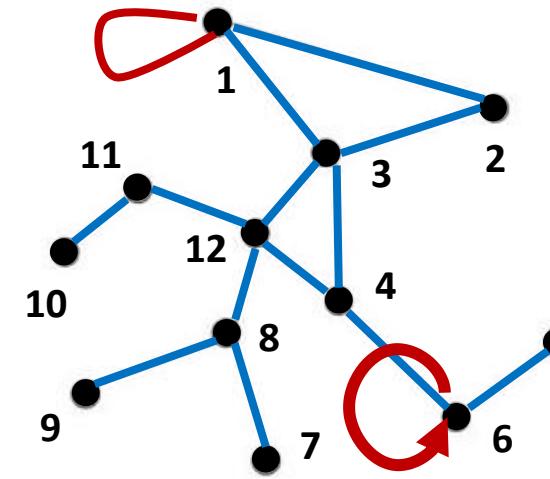
$A_{ij} = 1$ if $(n_i, n_j) \in E$
(nodes n_i and n_j have a common edge)

$A_{ij} = 0$ else

Graph $G = G(N, E)$

nodes N (vertices V)
edges E

$$A = \begin{pmatrix} 1 & 1 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 1 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 1 & 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 1 & 1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 1 & 1 & 0 & 0 & 0 & 1 & 0 & 0 & 1 & 0 \end{pmatrix}$$



A graph can have **loops**

note:

$d(n_1) = 4$, since loop is **undirected** and hits the node twice!

structural information: **adjacency matrix A**

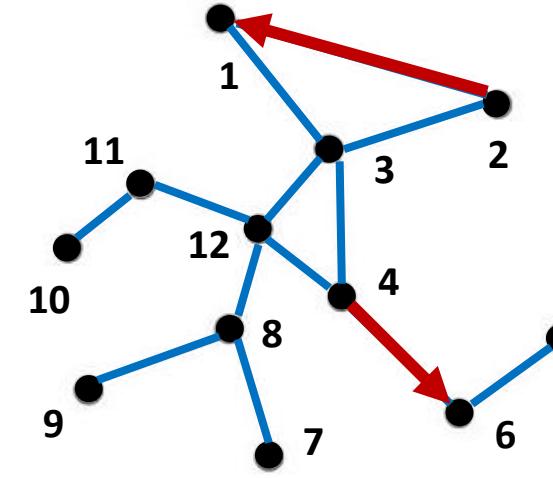
$A_{ij} = 1$ if $(n_i, n_j) \in E$
(nodes n_i and n_j have a common edge)

$$A_{ij} = 0 \text{ else}$$

Graph $G = G(N, E)$

nodes N (vertices V)
edges E

$$A = \begin{pmatrix} 2 & 1 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 1 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 1 & 0 & 0 & 1 & 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 1 & 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 1 & 0 & 0 & 0 & 1 & 0 & 1 & 0 \end{pmatrix}$$



A graph can be **directed**

structural information: **adjacency matrix A**

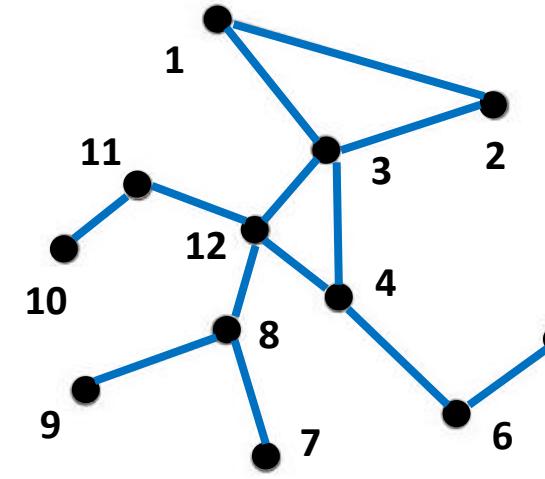
$A_{ij} = 1$ if $(n_i, n_j) \in E$
(nodes n_i and n_j have a common edge)

$A_{ij} = 0$ else

$$A = \begin{pmatrix} 0 & \textcolor{red}{0} & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ \textcolor{red}{1} & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 1 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 1 & 0 & 0 & 0 & \textcolor{red}{1} & 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & 1 & 0 & 0 & 0 & 1 & 0 & 0 & 1 & 0 \end{pmatrix}$$

Graph $G = G(N, E)$

nodes N (vertices V)
edges E



structural information: **adjacency matrix A**

$A_{ij} = 1$ if $(n_i, n_j) \in E$
(nodes n_i and n_j have a common edge)

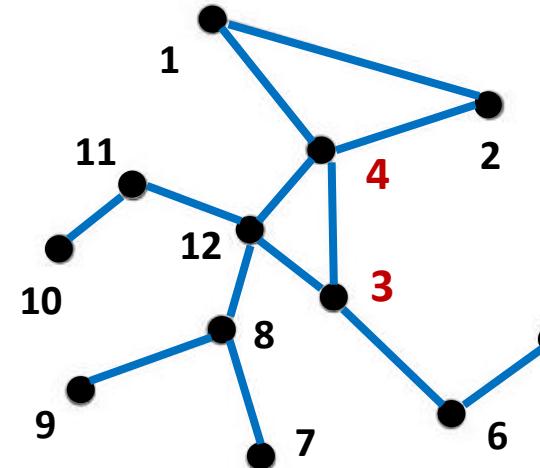
$A_{ij} = 0$ else

The order of enumerating the nodes is not relevant!
(permutation invariance)

$$A = \begin{pmatrix} 0 & 1 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 1 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 1 & 0 & 0 & 1 & 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 1 & 1 & 0 & 0 & 0 & 1 & 0 & 1 & 0 \end{pmatrix}$$

Graph $G = G(N, E)$

nodes N (vertices V)
edges E



structural information: **adjacency matrix A**

$A_{ij} = 1$ if $(n_i, n_j) \in E$
(nodes n_i and n_j have a common edge)

$A_{ij} = 0$ else

The order of enumerating the nodes is not relevant!
(permutation invariance)

Each graph can be represented by **$N!$** adjacency matrices!

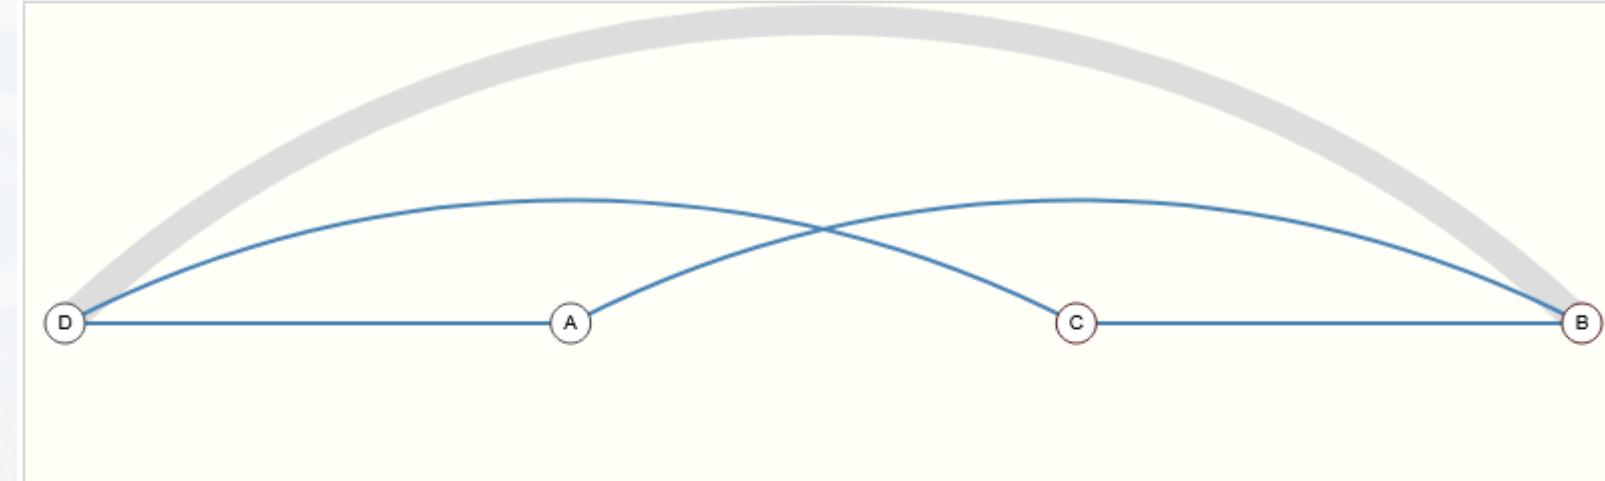
$$A = \begin{pmatrix} 0 & 1 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & \textcolor{red}{1} & 0 & 0 & \textcolor{red}{1} & 0 & 0 & 0 & 0 & \textcolor{red}{1} \\ \textcolor{red}{1} & \textcolor{red}{1} & \textcolor{red}{1} & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 1 & 1 & 0 & 0 & 0 & 1 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 & 1 & 0 & 0 & 0 & 1 & 0 & 0 & 1 \end{pmatrix}$$

Graph $G = G(N, E)$

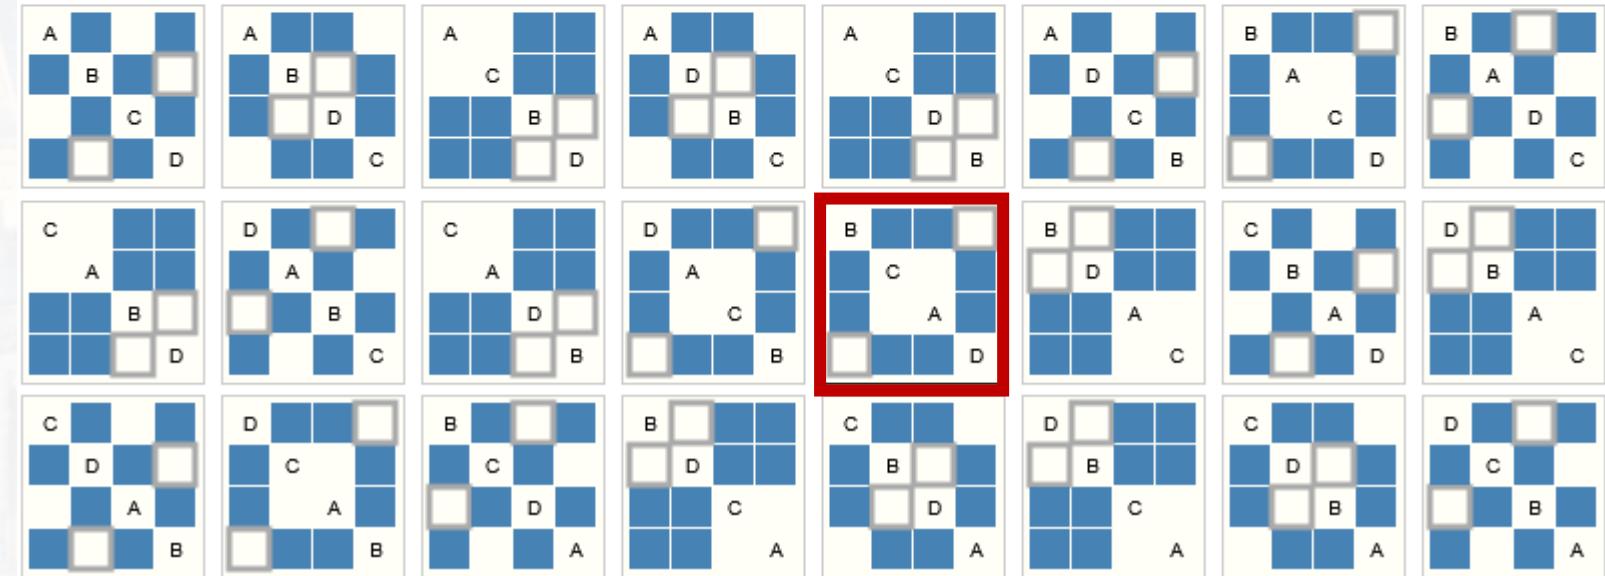
nodes N (vertices V)
edges E



Each graph can be represented by $N!$ adjacency matrices!



[animation here](#)

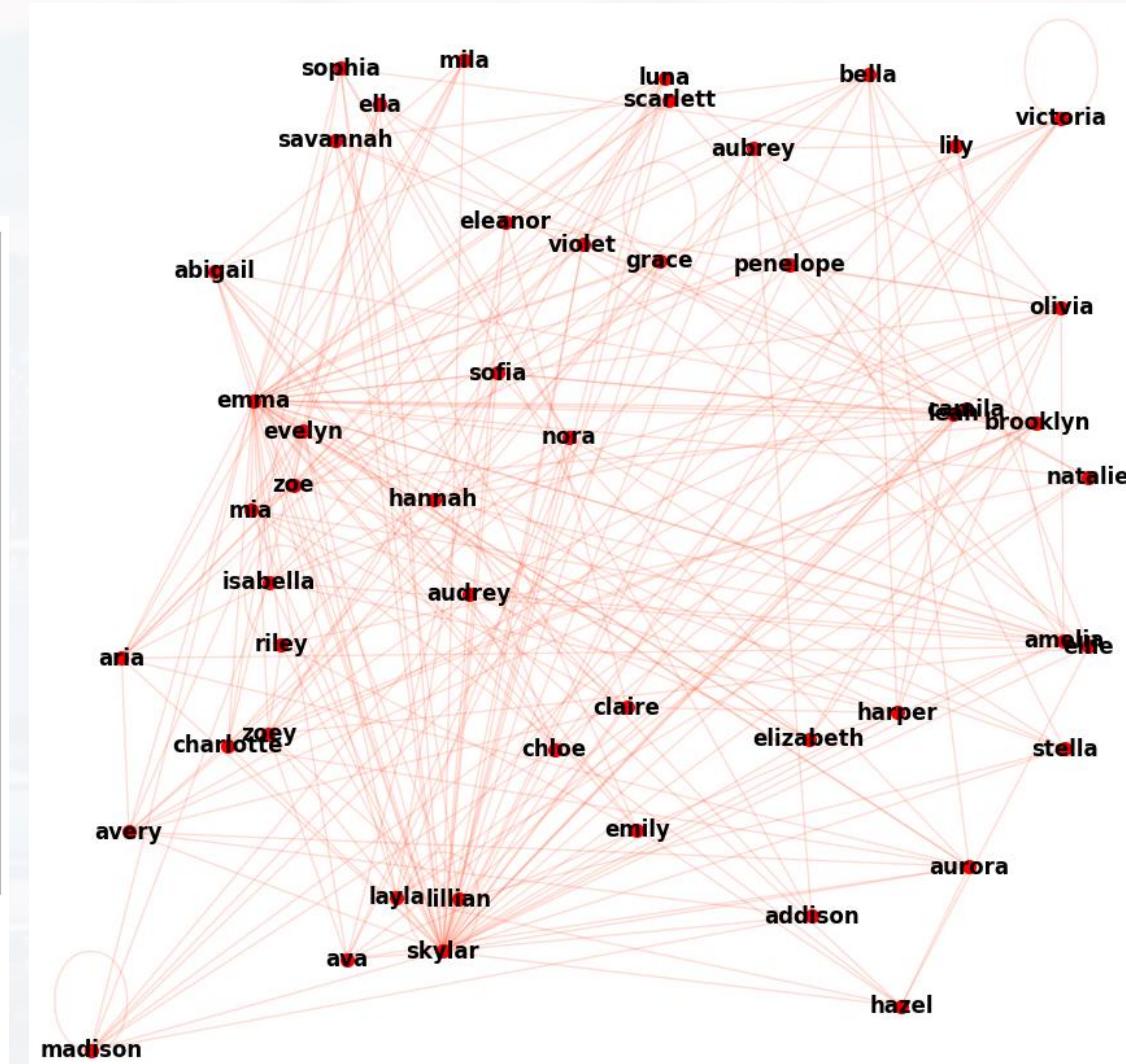
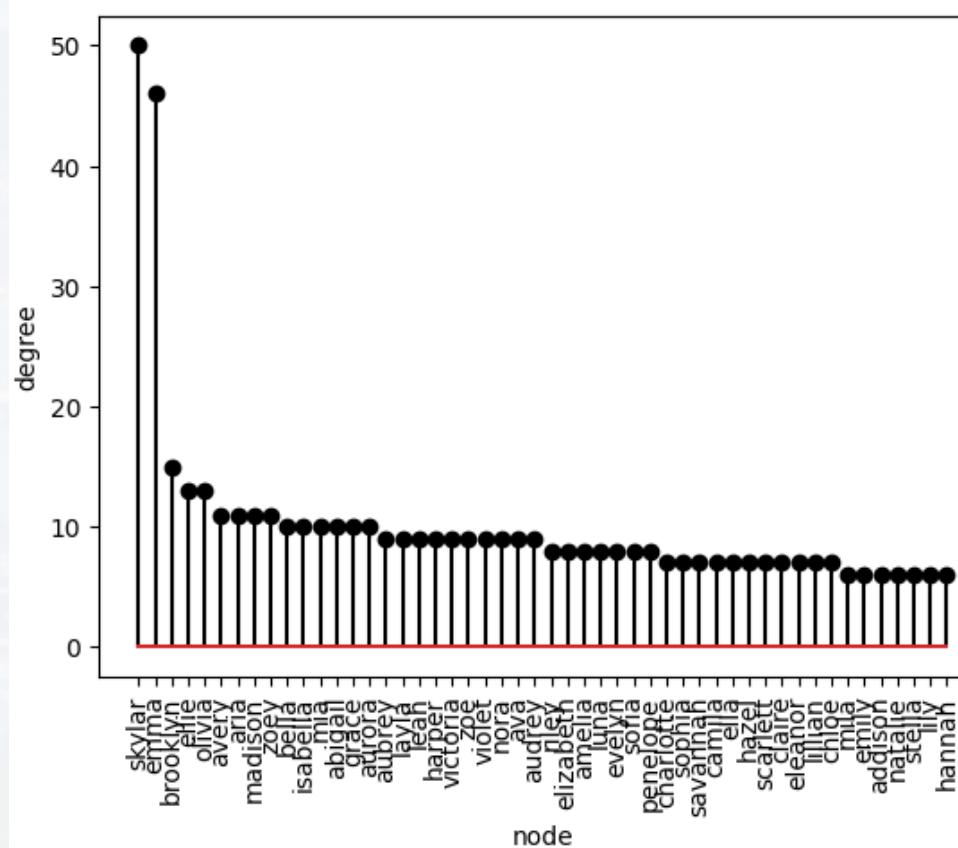




visualizing a graph:

```
import networkx as nx #pip install networkx
```

see: Graph_I.ipynb

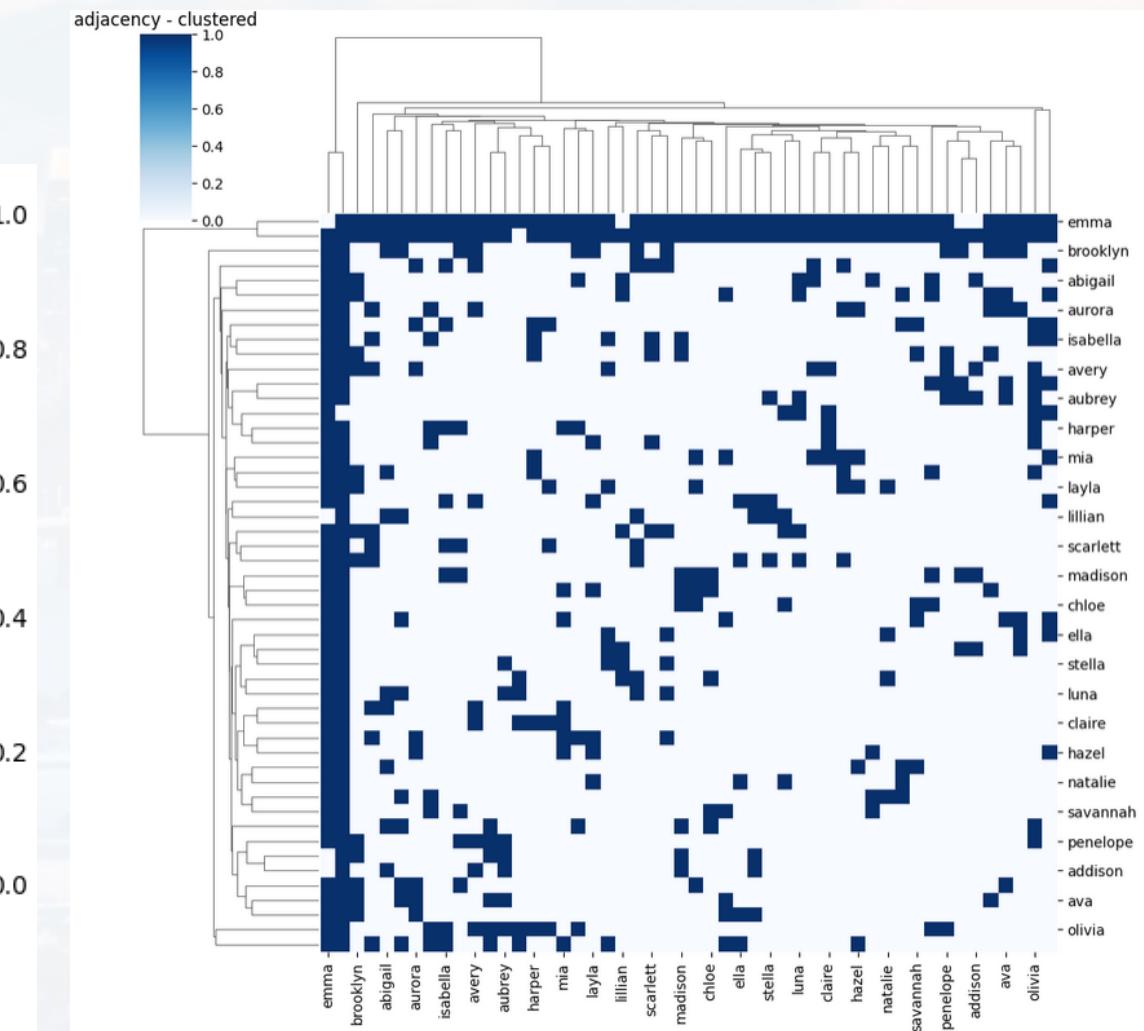
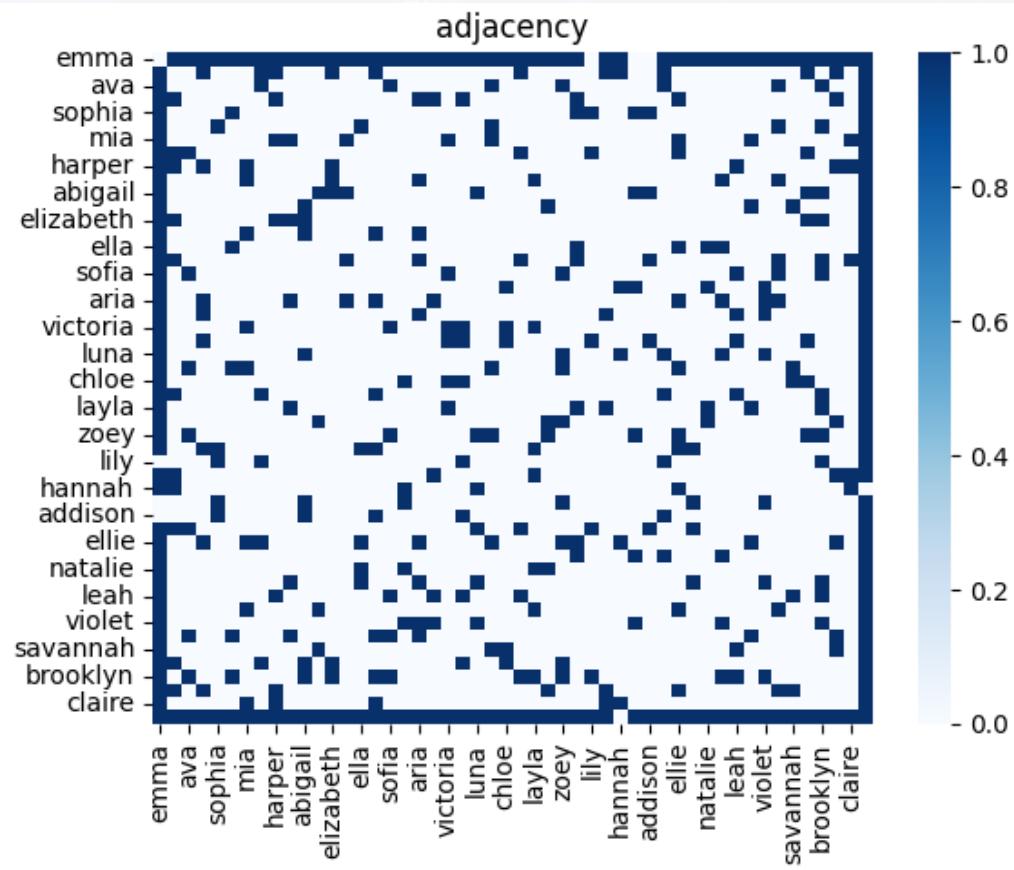




visualizing a graph:

```
import networkx as nx #pip install networkx
```

see: [Graph_I.ipynb](#)



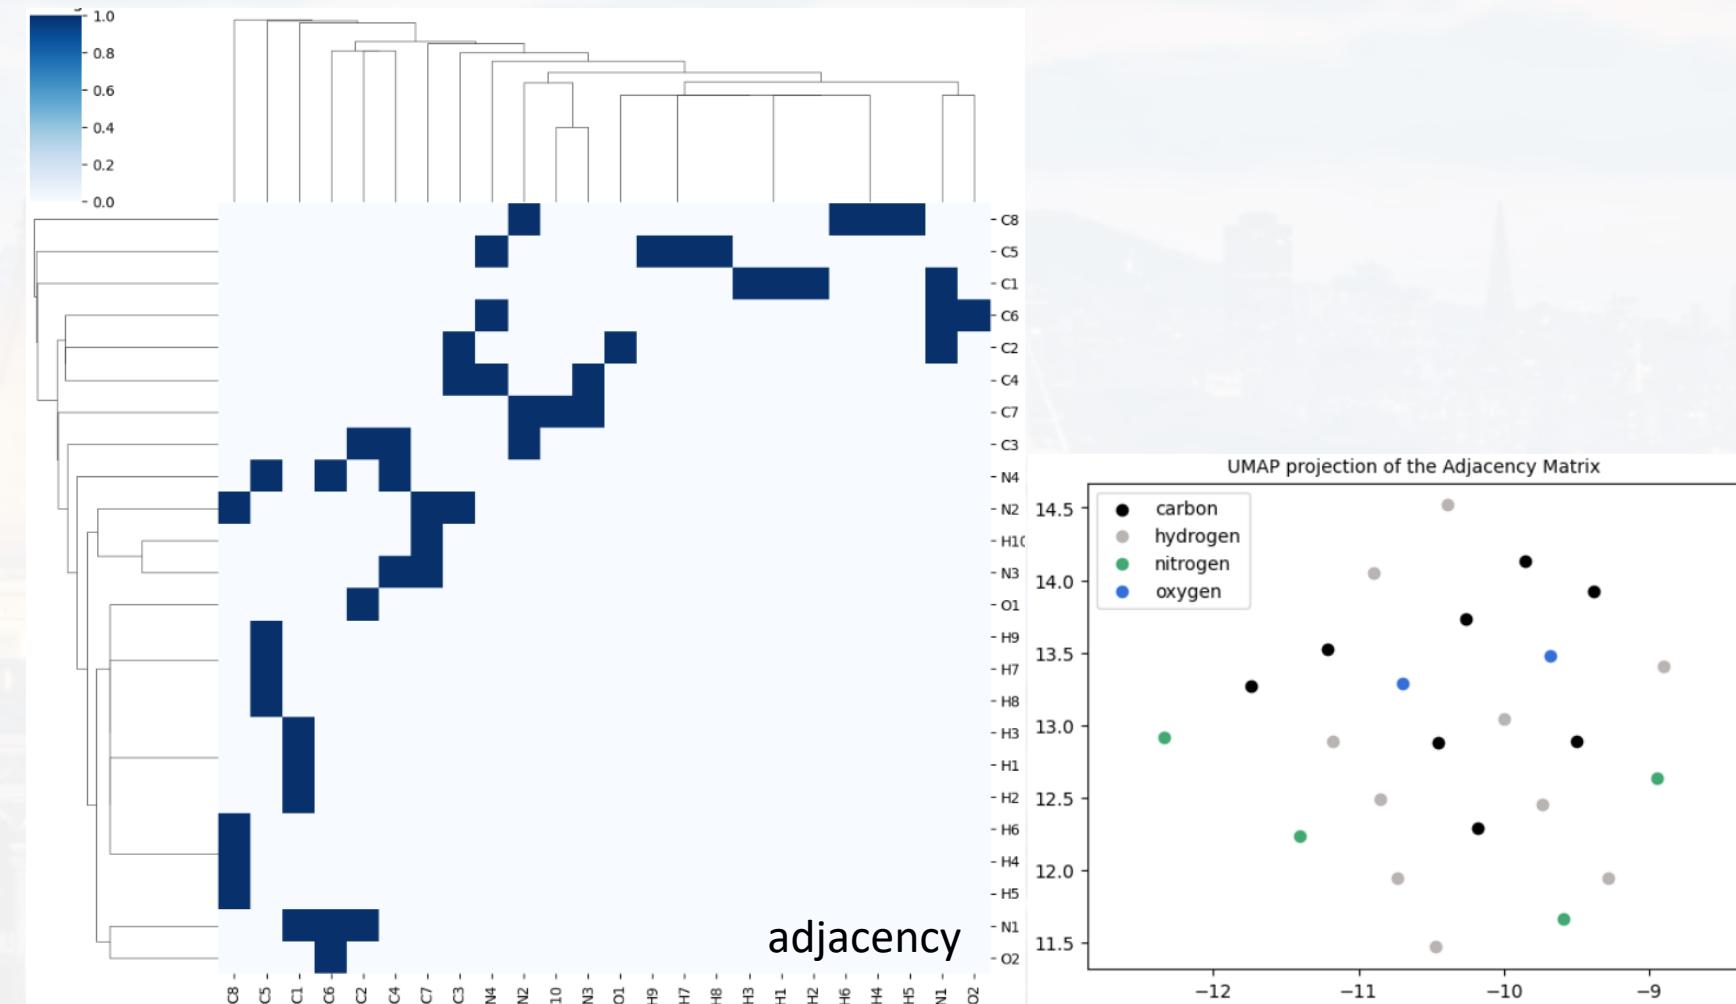


building and visualizing a **weighted** graph:

```
import networkx as nx #pip install networkx
```

see: [Graph_II.ipynb](#)

Caffein molecule





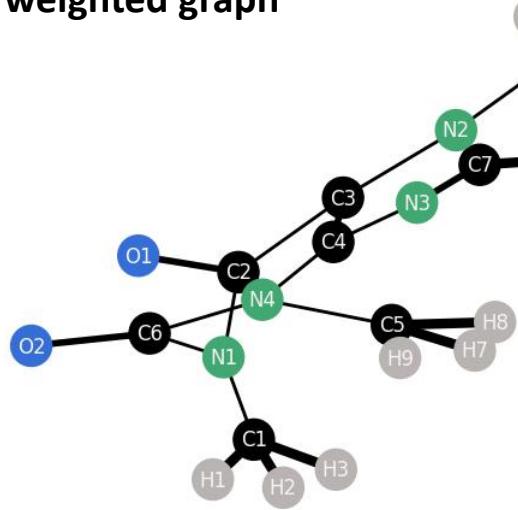
building and visualizing a **weighted** graph:

```
import networkx as nx #pip install networkx
```

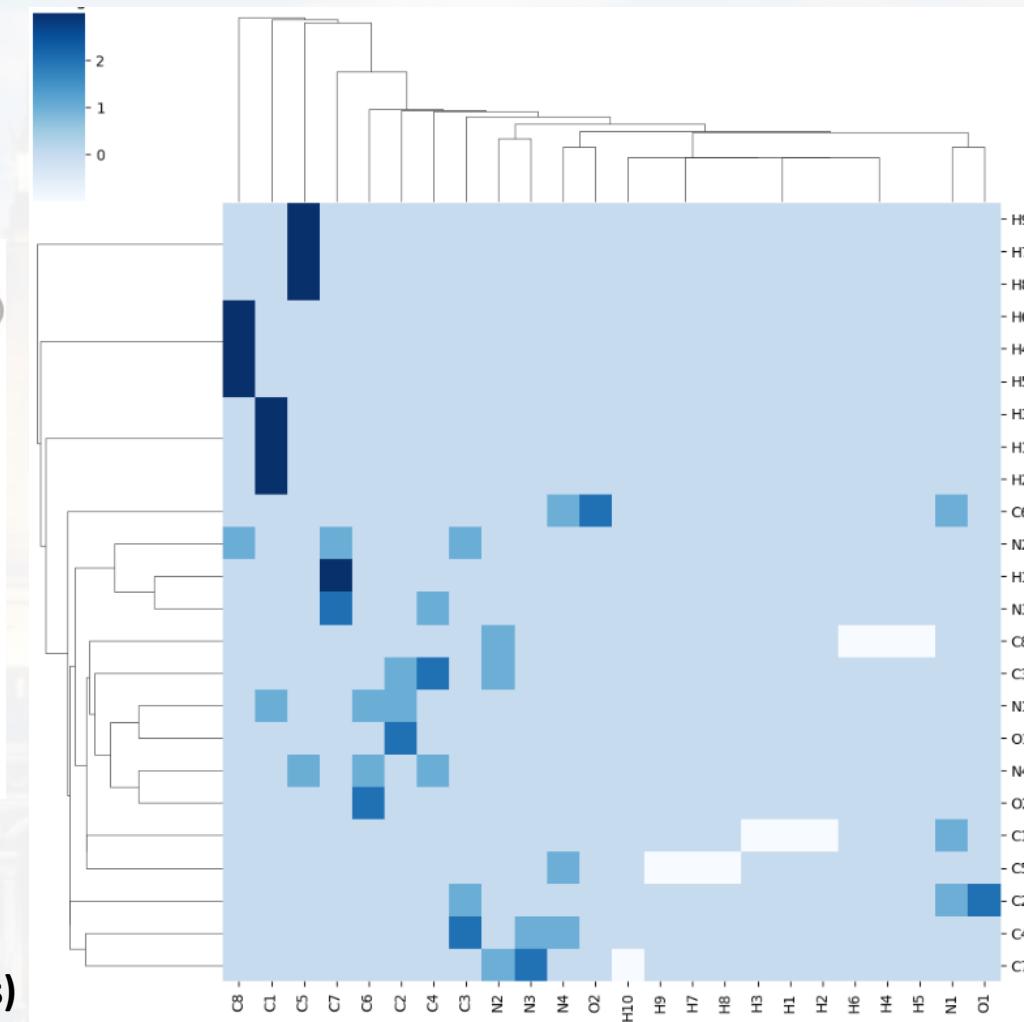
see: [Graph_II.ipynb](#)

Caffein molecule

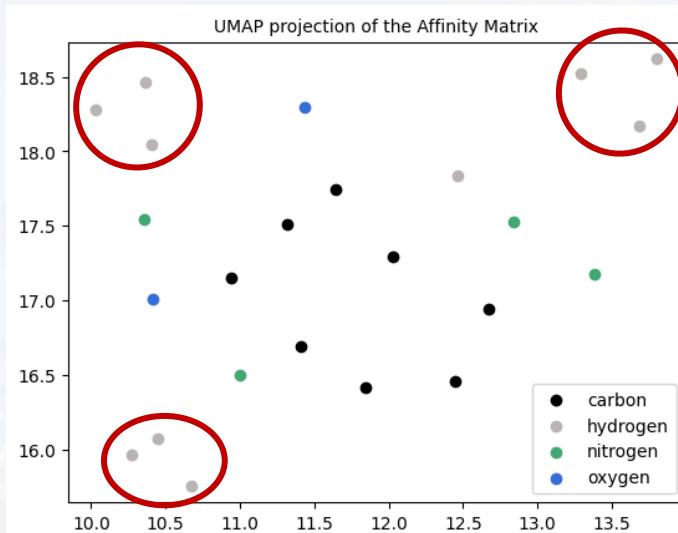
weighted graph



binding affinity (weights)



hydrogen atoms are at the edges of the molecule!





more about graphs:

$$d(n_i) = \sum_j A_{ij}$$

degree of node n_i

$$\mathcal{N}(n_i) = \{n_j \in N : (n_i, n_j) \in E\}$$

neighborhood $\mathcal{N}(n_i)$ of node n_i
for first degree neighborhood $|\mathcal{N}(n_i)| = d(n_i)$

$$S_{com} = |\mathcal{N}(n_i) \cap \mathcal{N}(n_j)|$$

number for neighbors nodes n_i and n_j have in common.

idea: nodes with many common neighbors are more likely to be similar or have a potential connection.

$$S_{rat} = \frac{|\mathcal{N}(n_i) \cap \mathcal{N}(n_j)|}{|\mathcal{N}(n_i) \cup \mathcal{N}(n_j)|}$$

ratio for neighbors nodes n_i and n_j have in common.

note:

There are more quantities (“importance”, “centrality” etc.), but they are all a function of A_{ij} .

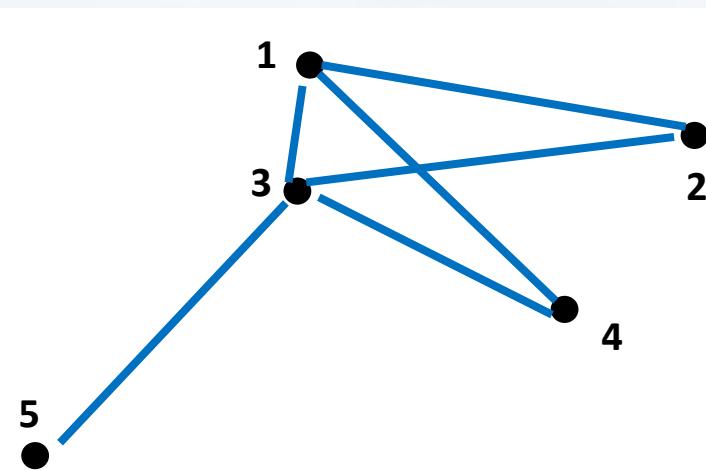


more about A_{ij} :

$$\vec{e}_0 := \begin{pmatrix} 1 \\ \vdots \\ 1 \end{pmatrix}$$

$$\boxed{\vec{e}_{t+1} = A \vec{e}_t}$$

then \vec{e}_t is the number of length t paths arriving at each node
→ dynamics of the system (diffusion, “message passing”)



$$A = \begin{pmatrix} 0 & 1 & 1 & 1 & 0 \\ 1 & 0 & 1 & 0 & 0 \\ 1 & 1 & 0 & 1 & 1 \\ 1 & 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 \end{pmatrix} \quad \vec{e}_0 := \begin{pmatrix} 1 \\ 1 \\ 1 \\ 1 \\ 1 \end{pmatrix}$$

$$\vec{e}_1 = A \vec{e}_0 = \begin{pmatrix} 3 \\ 2 \\ 4 \\ 2 \\ 1 \end{pmatrix} \quad \text{Just the degree of each node!}$$

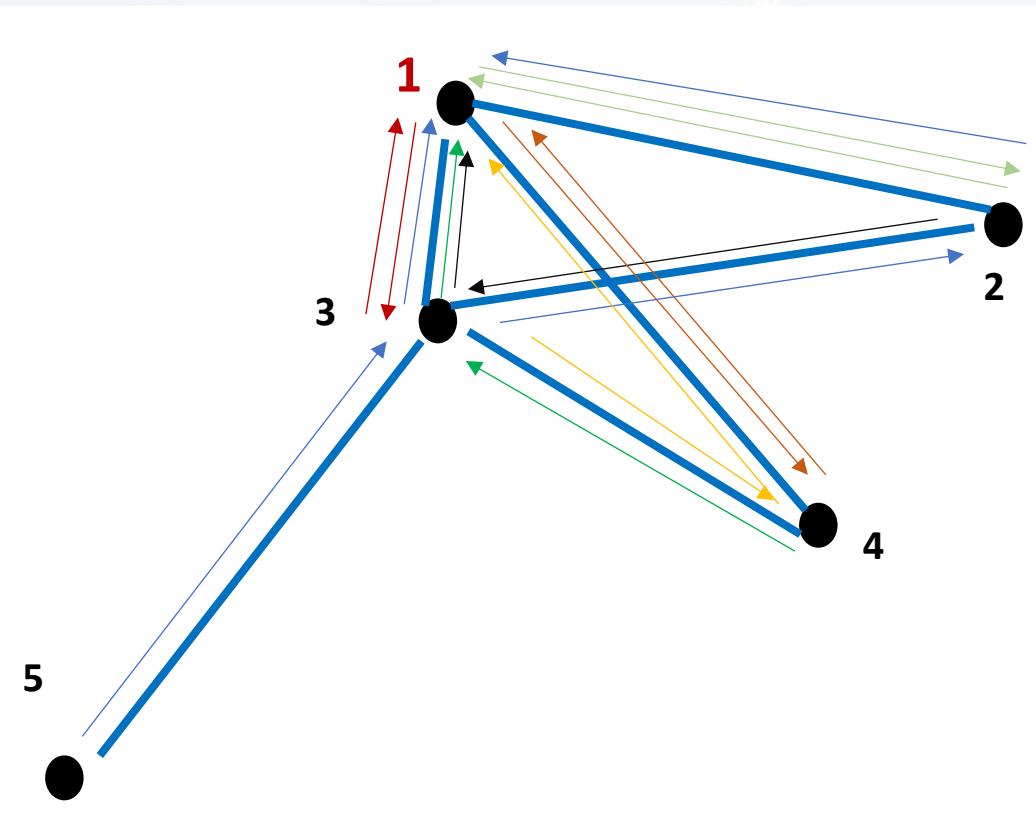
$$\vec{e}_2 = A \vec{e}_1 = \begin{pmatrix} 8 \\ 7 \\ 8 \\ 7 \\ 4 \end{pmatrix} \quad \text{number of length } t = 2 \text{ paths arriving at each node}$$



more about A_{ij} :

$$\vec{e}_0 := \begin{pmatrix} 1 \\ \vdots \\ 1 \end{pmatrix}$$

$$\boxed{\vec{e}_{t+1} = A \vec{e}_t}$$



then \vec{e}_t is the number of length t paths arriving at each node
→ dynamics of the system (diffusion, “message passing”)

$$A = \begin{pmatrix} 0 & 1 & 1 & 1 & 0 \\ 1 & 0 & 1 & 0 & 0 \\ 1 & 1 & 0 & 1 & 1 \\ 1 & 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 \end{pmatrix}$$

$$\vec{e}_2 = A \vec{e}_1 = \begin{pmatrix} 8 \\ 7 \\ 8 \\ 7 \\ 4 \end{pmatrix}$$

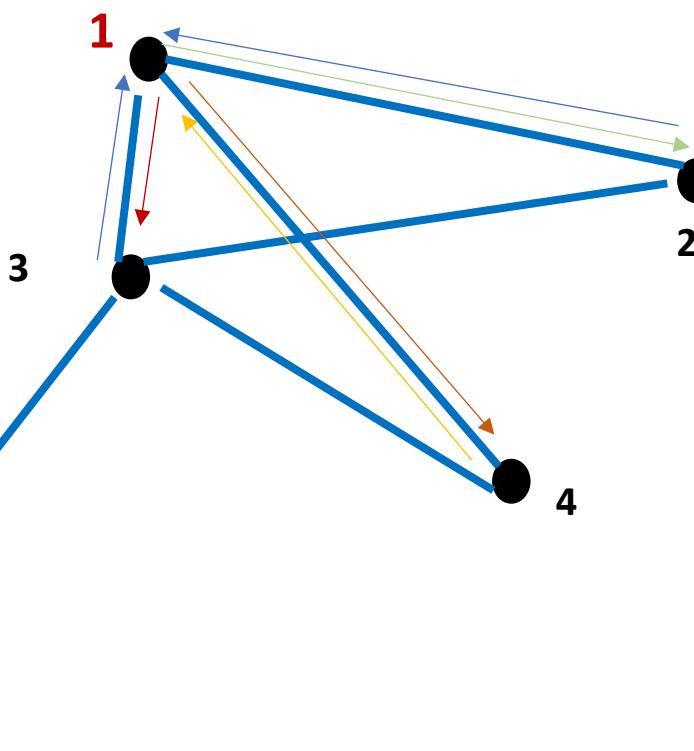
number of length $t = 2$ paths arriving at each node



more about A_{ij} :

$$\vec{e}_0 := \begin{pmatrix} 1 \\ \vdots \\ 1 \end{pmatrix}$$

$$\boxed{\vec{e}_{t+1} = A \vec{e}_t}$$



then \vec{e}_t is the number of length t paths arriving at each node
→ dynamics of the system (diffusion, “message passing”)

$$A = \begin{pmatrix} 0 & 1 & 1 & 1 & 0 \\ 1 & 0 & 1 & 0 & 0 \\ 1 & 1 & 0 & 1 & 1 \\ 1 & 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 \end{pmatrix}$$

$$\vec{e}_1 = A \vec{e}_0 = \begin{pmatrix} 3 \\ 2 \\ 4 \\ 2 \\ 1 \end{pmatrix}$$

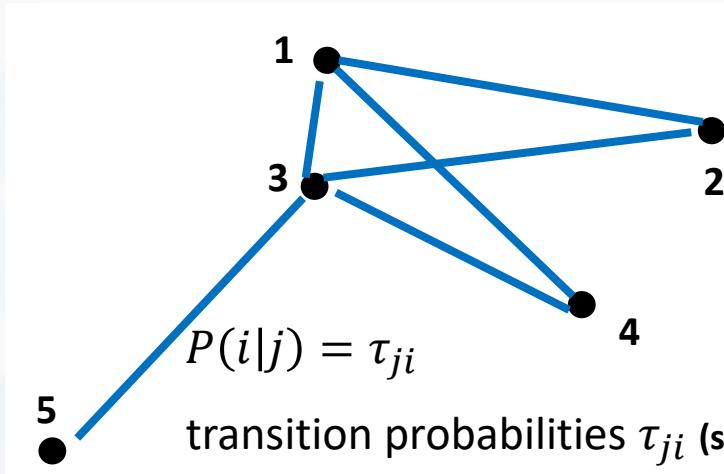
probabilistic point of view:

$$\vec{e}_1 \rightarrow \vec{P}_1 = \begin{pmatrix} P(1|2)P(2) + P(1|3)P(3) + P(1|4)P(4) \\ P(2|1)P(1) + P(2|3)P(3) \\ P(3|1)P(1) + \dots etc \\ \dots \\ \dots \end{pmatrix}$$



more about A_{ij} :

probabilistic point of view:



transition probabilities τ_{ji} (see module 7 “HMMs” and “stochastic processes”)

master equation

$$\frac{dP(n_i, t)}{dt} = \sum_{j=1}^J \tau_{ji} P(n_j, t) - \sum_{j=1}^J \tau_{ij} P(n_i, t)$$

$$= \sum_{j=1}^J \tau_{ji} P(n_j, t) - P(n_i, t) \sum_{j=1}^J \tau_{ij}$$

$$= \sum_{j=1}^J \tau_{ji} A_{ji} P(n_j, t) - P(n_i, t) \sum_{j=1}^J \tau_{ij} A_{ij}$$

$$\vec{e}_1 \rightarrow \vec{P}_1 = \begin{pmatrix} P(1|2)P(2) + P(1|3)P(3) + P(1|4)P(4) \\ P(2|1)P(1) + P(2|3)P(3) \\ P(3|1)P(1) + \dots etc \\ \dots \\ \dots \end{pmatrix}$$

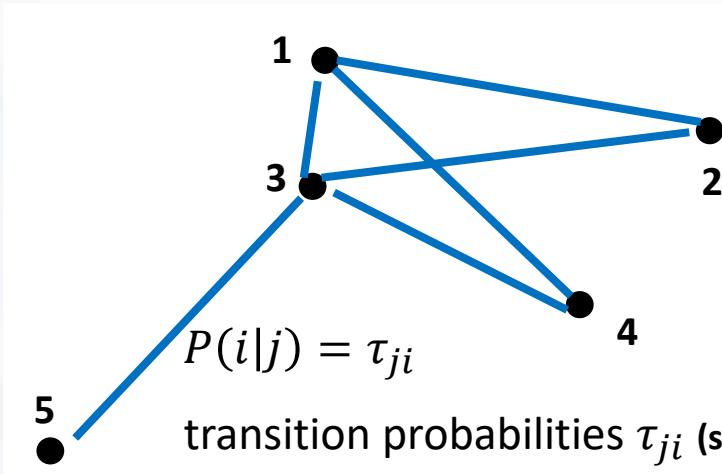
$$A_{ij} = \begin{cases} 1, & \text{if } (n_i, n_j) \in E \\ 0, & \text{else} \end{cases}$$

$$\tau_{ij} = \begin{cases} \tau_{ij}, & \text{if } (n_i, n_j) \in E \\ 0, & \text{else} \end{cases}$$



more about A_{ij} :

probabilistic point of view:



transition probabilities τ_{ji} (see module 7 “HMMs” and “stochastic processes”)

$$\vec{e}_1 \rightarrow \vec{P}_1 = \begin{pmatrix} P(1|2)P(2) + P(1|3)P(3) + P(1|4)P(4) \\ P(2|1)P(1) + P(2|3)P(3) \\ P(3|1)P(1) + \dots etc \\ \dots \\ \dots \end{pmatrix}$$

master equation

$$\frac{dP(n_i, t)}{dt} = \sum_{j=1}^J \tau_{ji} P(n_j, t) - P(n_i, t) \sum_{j=1}^J \tau_{ij}$$

$$\tau_{ij} = \begin{cases} \tau_{ij}, & \text{if } (n_i, n_j) \in E \\ 0, & \text{else} \end{cases}$$

$$\frac{d \vec{P}(n, t)}{dt} = \tau \vec{P}(n, t) - \vec{P}(n, t) * [\tau \vec{P}(n, t)]$$

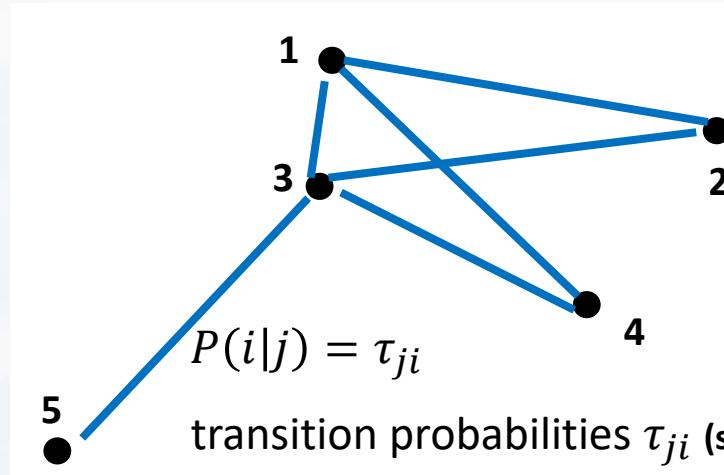
* : element wise multiplication

$$\frac{d \vec{P}(n, t)}{dt} = \vec{P}(n, t) * (\tau - [\tau \vec{P}(n, t)])$$



more about A_{ij} :

probabilistic point of view:



master equation

$$\frac{d \vec{P}(n, t)}{dt} = \tau \vec{P}(n, t) - \vec{P}(n, t) * [\tau \vec{P}(n, t)]$$

$$\frac{d \vec{P}(n, t)}{dt} = c(D - A)\vec{P}(n, t)$$

graph Laplacian \mathcal{L}

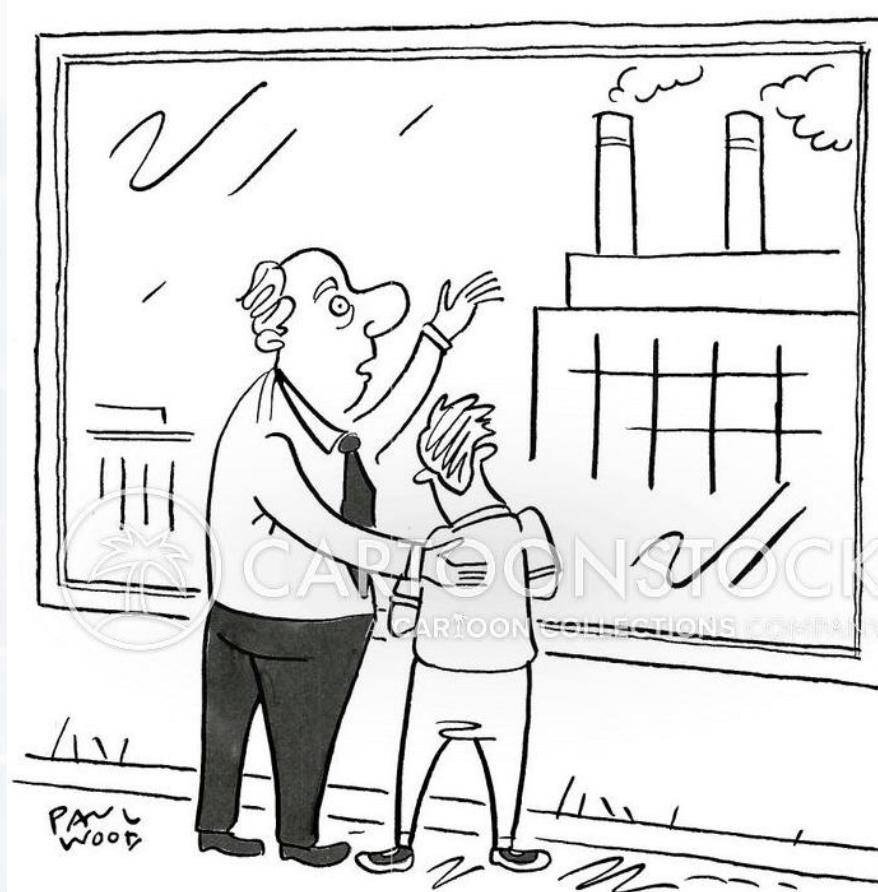
$$\vec{e}_1 \rightarrow \vec{P}_1 = \begin{pmatrix} P(1|2)P(2) + P(1|3)P(3) + P(1|4)P(4) \\ P(2|1)P(1) + P(2|3)P(3) \\ P(3|1)P(1) + \dots etc \\ \dots \\ \dots \end{pmatrix}$$

$$\tau_{ij} = \begin{cases} 1, & \text{if } (n_i, n_j) \in E \\ 0, & \text{else} \end{cases}$$

* : element wise multiplication

for $\tau_{ij} = c$

$$D = \begin{pmatrix} d(n_1) & 0 & 0 \\ 0 & d(n_2) & 0 \\ 0 & 0 & \dots \end{pmatrix}$$



ONE DAY SON, ALL THIS
WILL BE RUN BY ROBOTS

Outline

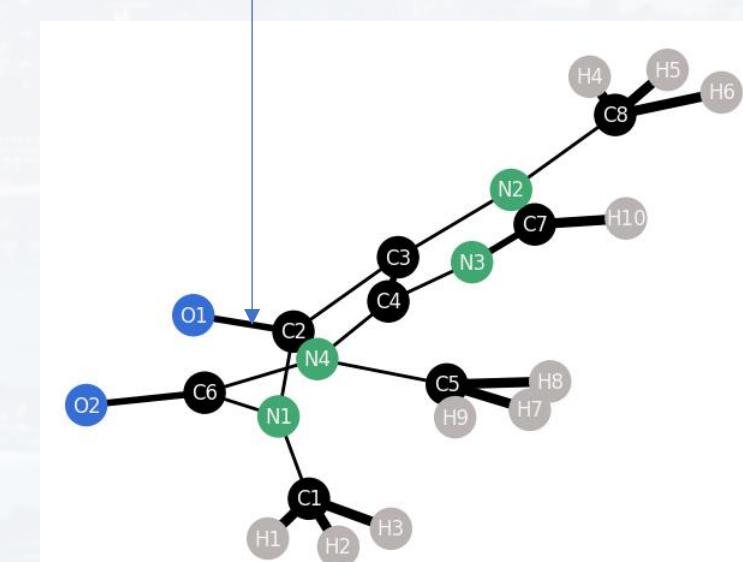
- What is a Graph
- The ANN Part
- PyTorch Example



What we can learn:

- node classification
- join nodes with similar properties to hyper nodes
- edge attributes, weights (weighted graph)
- edge prediction
- embedding (eg. 3D structure molecules)
- graph classification (is the molecule toxic y/n)
- graph regression (toxicity score)
- graph generation

weight: bond strength

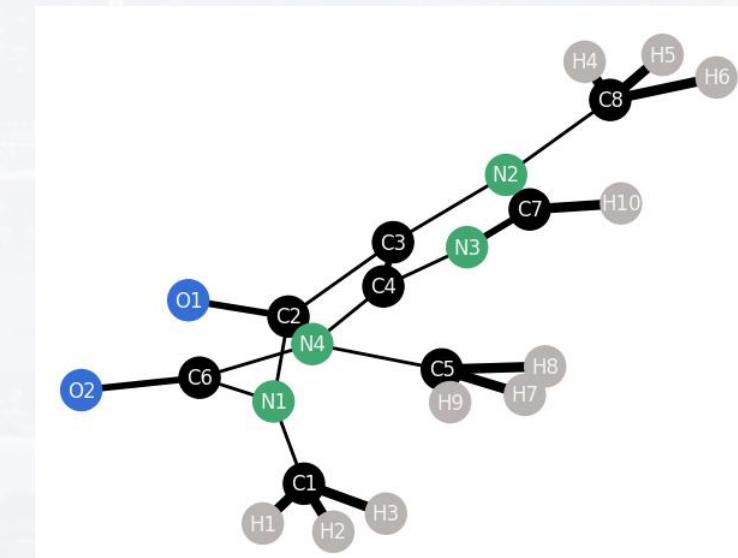




What we can learn:

- node classification
- join nodes with similar properties to hyper nodes
- edge attributes, weights (weighted graph)
- edge prediction
- embedding (eg. 3D structure molecules)
- **graph classification (is the molecule toxic y/n)**
- **graph regression (toxicity score)**
- **graph generation**

graph level tasks

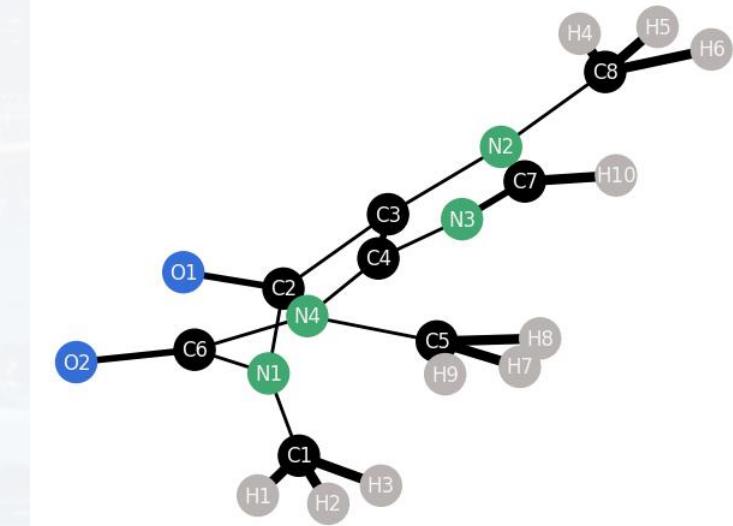




What we can learn:

- node classification
- join nodes with similar properties to hyper nodes
- edge attributes, weights (weighted graph)
- edge prediction
- embedding (eg. 3D structure molecules)
- graph classification (is the molecule toxic y/n)
- graph regression (toxicity score)
- graph generation

node (edge) level tasks





What we can learn:

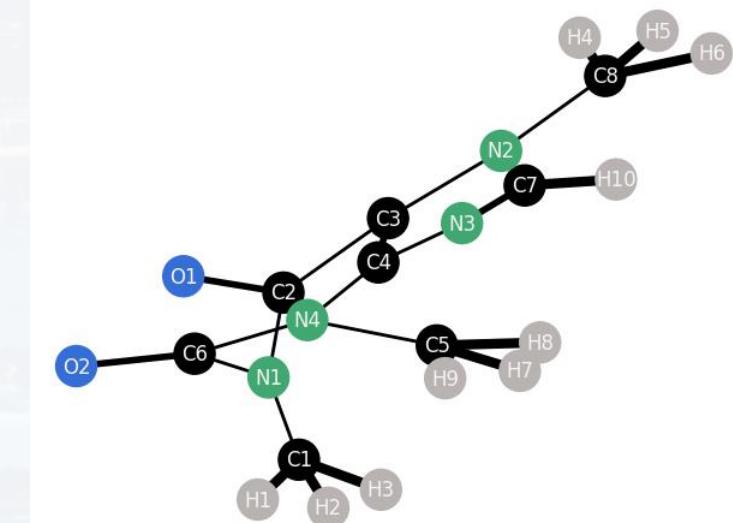
- node classification
- join nodes with similar properties to hyper nodes
- edge attributes, weights (weighted graph)
- edge prediction
- embedding (eg. 3D structure molecules)
- graph classification (is the molecule toxic y/n)
- graph regression (toxicity score)
- graph generation

information flow from one node to another:

message passing

different ways how:

- local averaging
- graph convolution (aka neighborhood aggregation)
- graph attention





What we can learn:

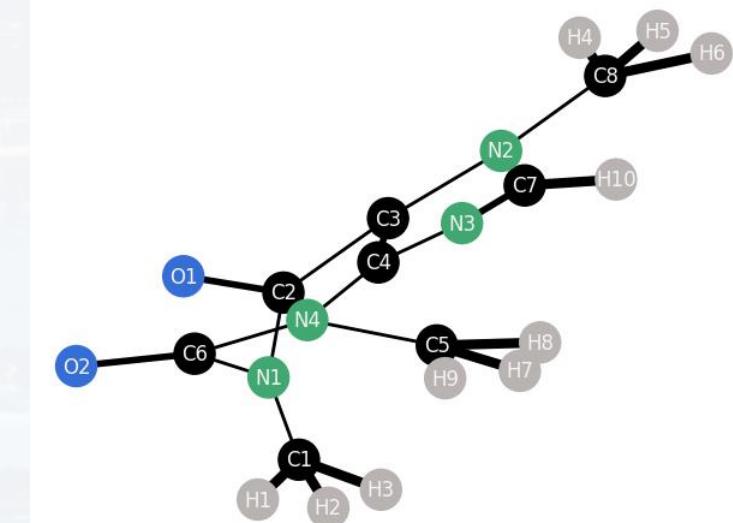
- node classification
- join nodes with similar properties to hyper nodes
- edge attributes, weights (weighted graph)
- edge prediction
- embedding (eg. 3D structure molecules)
- graph classification (is the molecule toxic y/n)
- graph regression (toxicity score)
- graph generation

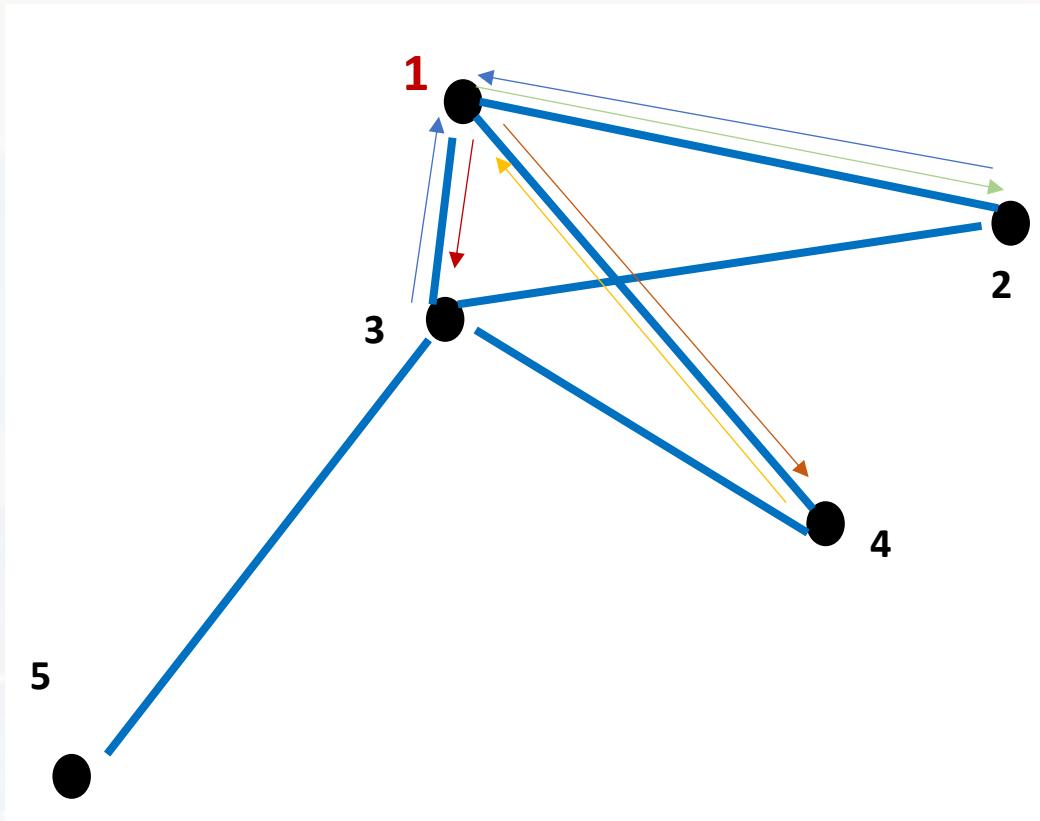
information flow from one node to another:

message passing

different ways how:

- local averaging
- **graph convolution** (aka neighborhood aggregation)
- **graph attention**





$$\vec{e}_1 \rightarrow \vec{P}_1 = \begin{pmatrix} \mathbf{P}(1|2)\mathbf{P}(2) + \mathbf{P}(1|3)\mathbf{P}(3) + \mathbf{P}(1|4)\mathbf{P}(4) \\ \mathbf{P}(2|1)\mathbf{P}(1) + \mathbf{P}(2|3)\mathbf{P}(3) \\ \mathbf{P}(3|1)\mathbf{P}(1) + \dots \text{etc} \\ \dots \\ \dots \end{pmatrix}$$

passing information from one node to the others

“message passing”

summing the information from neighbor nodes: **“aggregation”**

$h_{i,t}$: embedding vector of node i at time t

$z_{i,t}$: aggregate from i 's **neighbors**

$$z_{i,t} = \text{aggregate}(h_{m,t}: m \in \mathcal{N}(n_i))$$

updating values based on new information

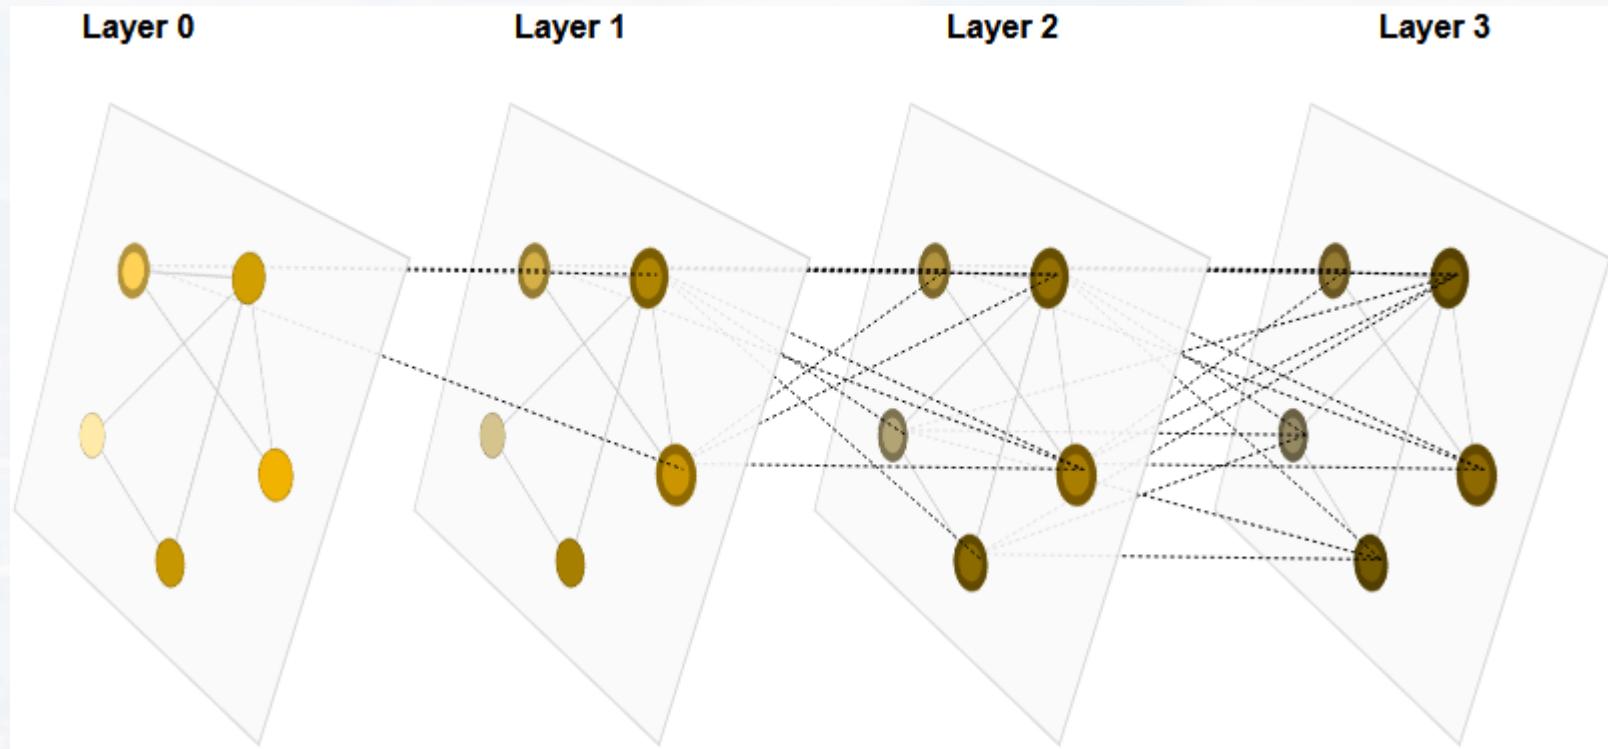
$$\text{update}(h_{i,t}, z_{i,t}) = f_{\text{nonlin}}(w_{\text{self}} h_{i,t} + w_{\text{neigh}} z_{i,t})$$

$w_{\text{self}}, w_{\text{neigh}}$: trainable weights



Graph Convolution

note: time t can be interpreted as different layers!

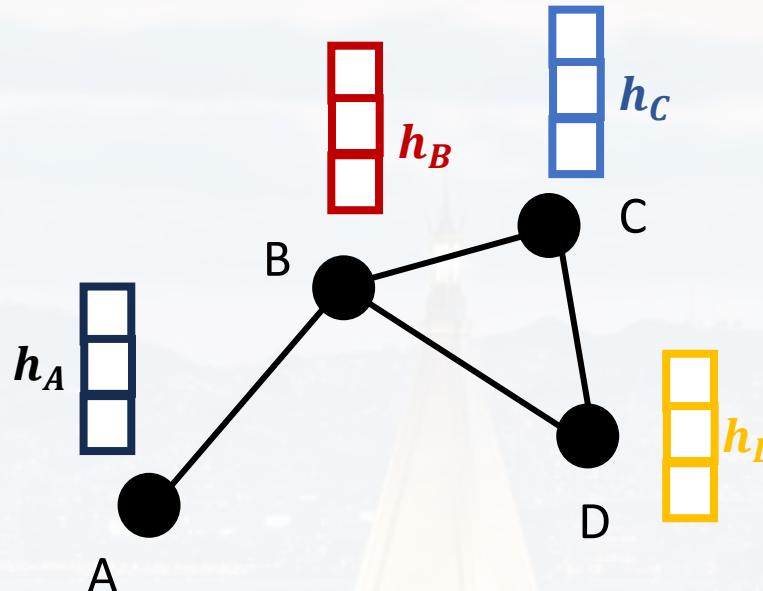


1st layer: one-hop neighborhood
2nd layer: two-hop neighborhood
etc

[animation here](#)

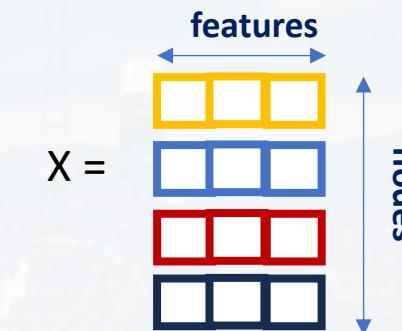


Graph Convolution

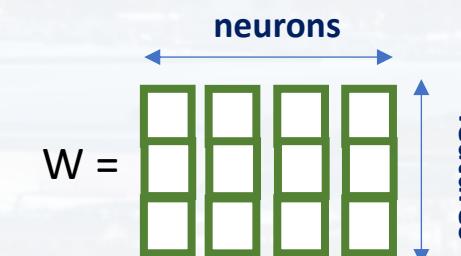


each node i has a **feature vector** h_i

matrix X of shape (number of nodes, number of node features)



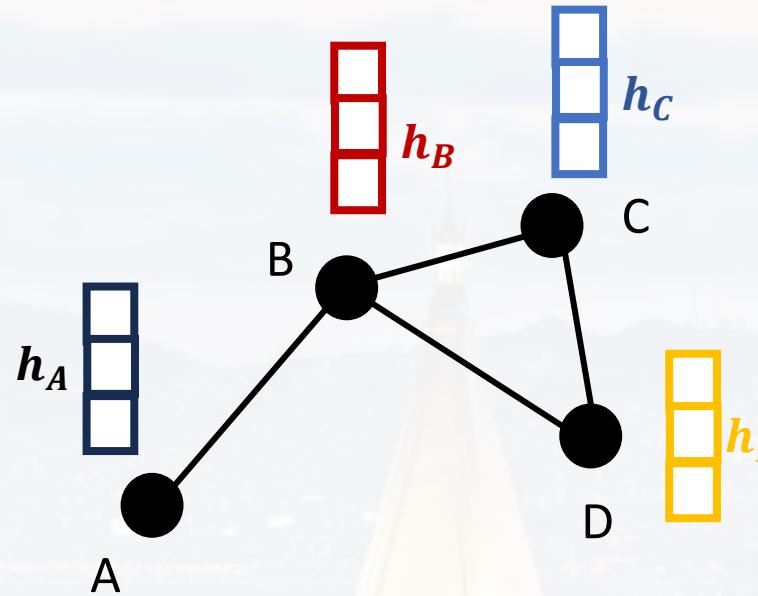
weight matrix W of shape
(number of node features, number of neurons)



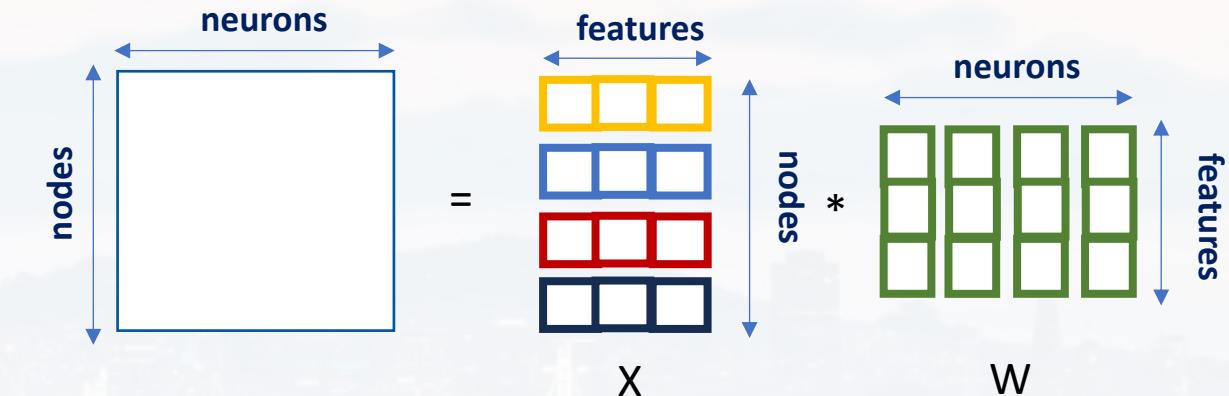
note: only **one** W for the entire graph
W is a **learnable**



Graph Convolution

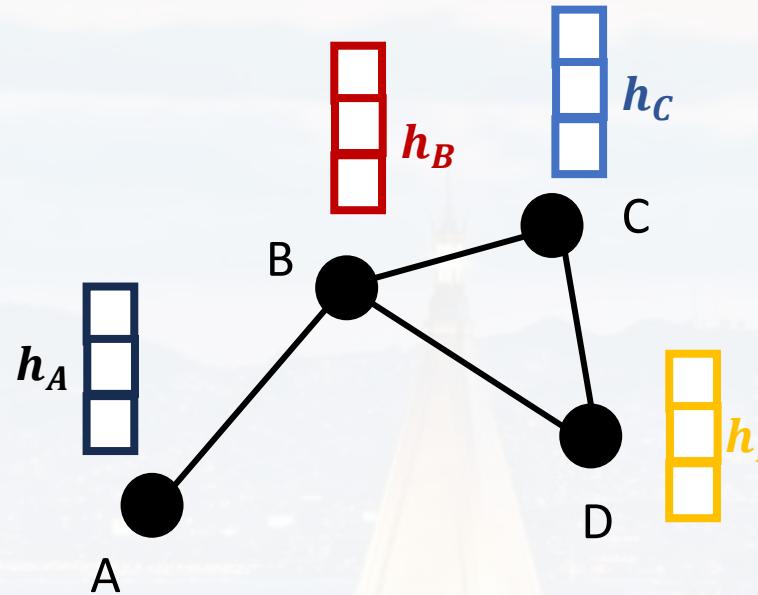


each node i has a **feature vector** h_i

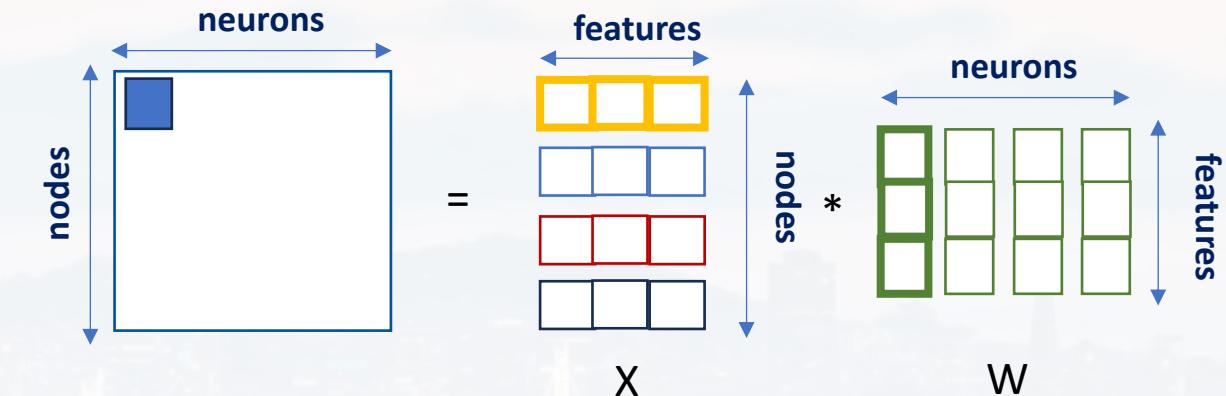




Graph Convolution

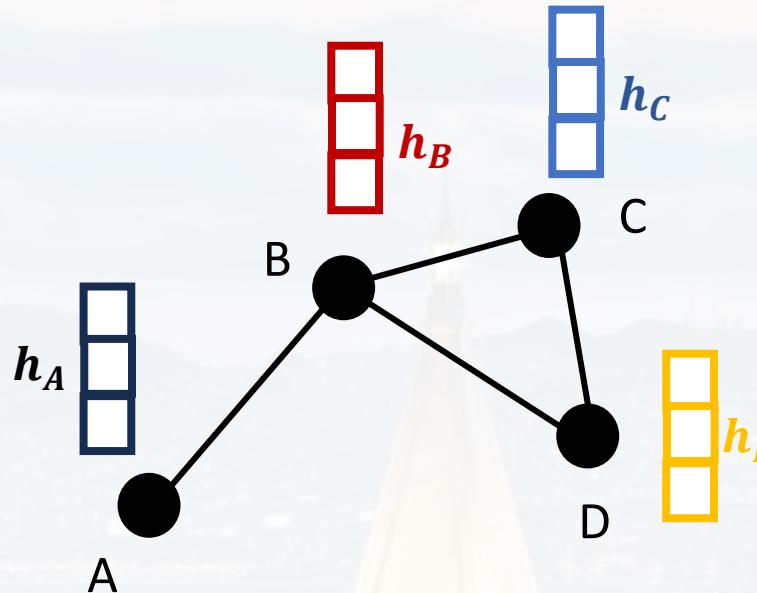


each node i has a **feature vector** h_i

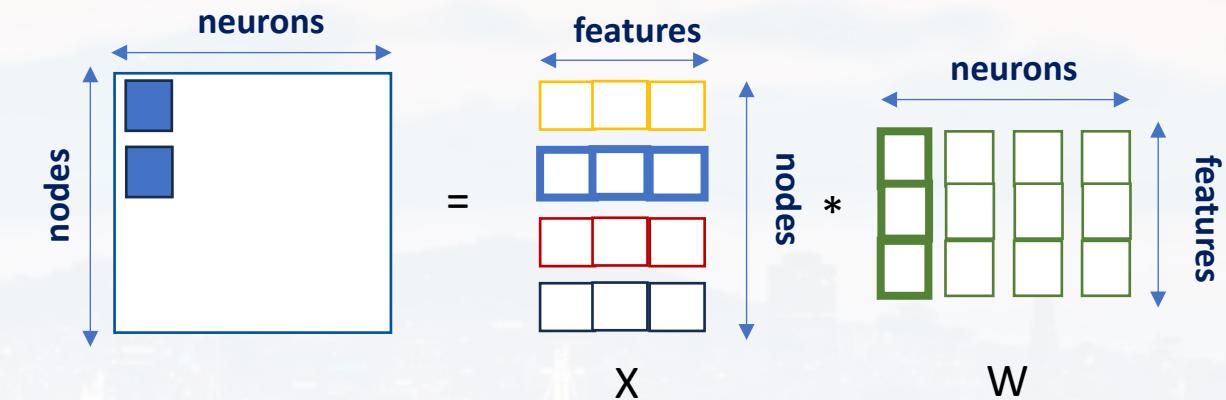




Graph Convolution

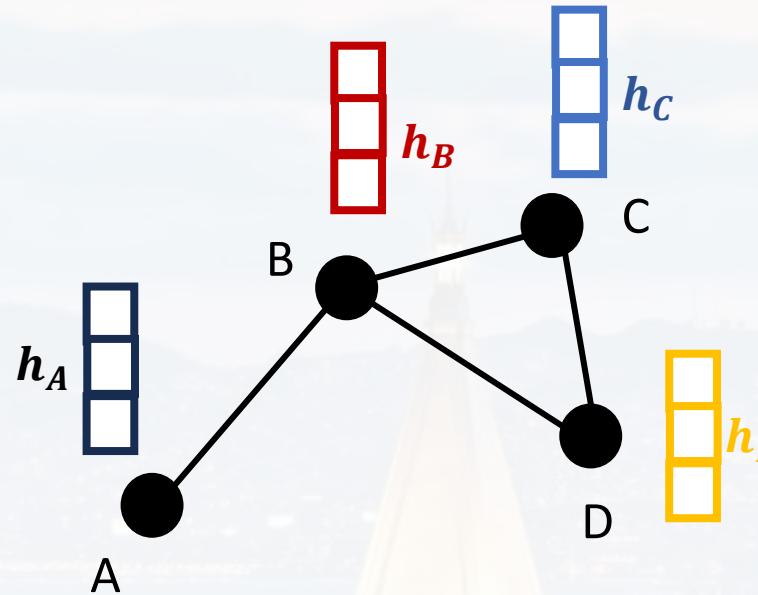


each node i has a **feature vector** h_i

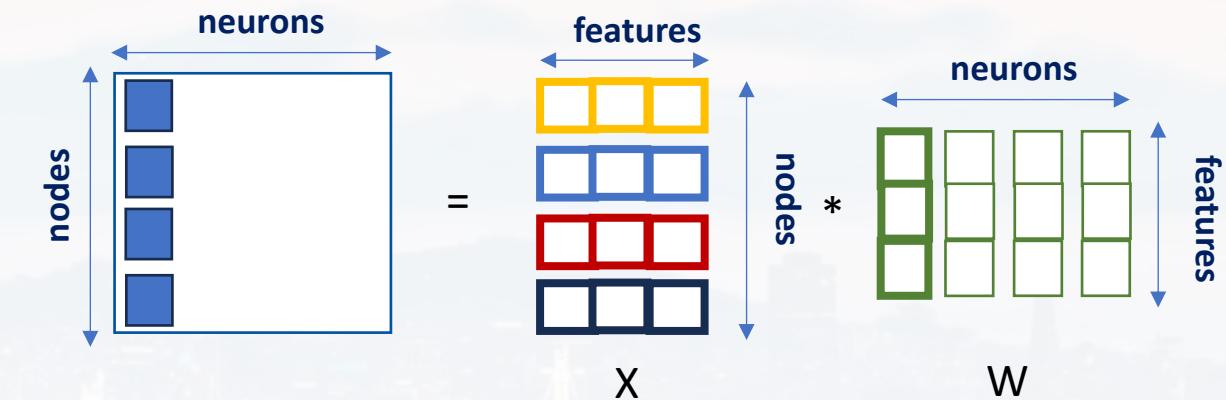




Graph Convolution

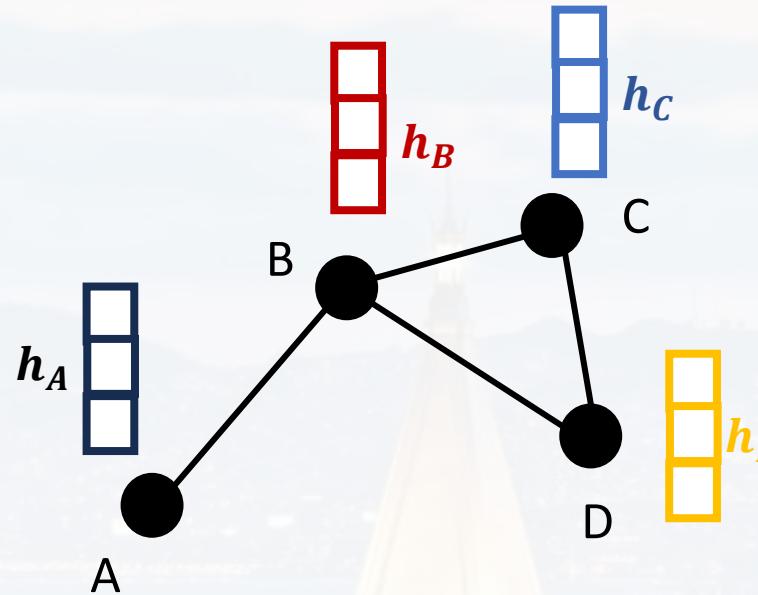


each node i has a **feature vector** h_i

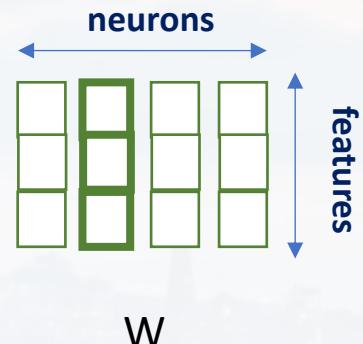
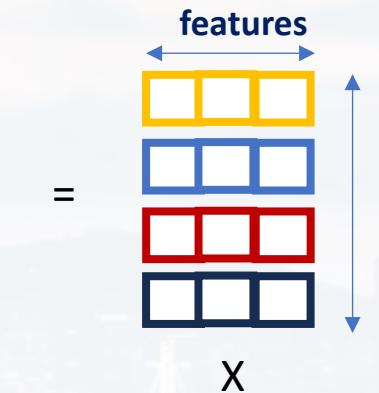
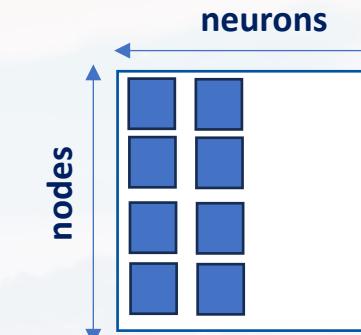




Graph Convolution

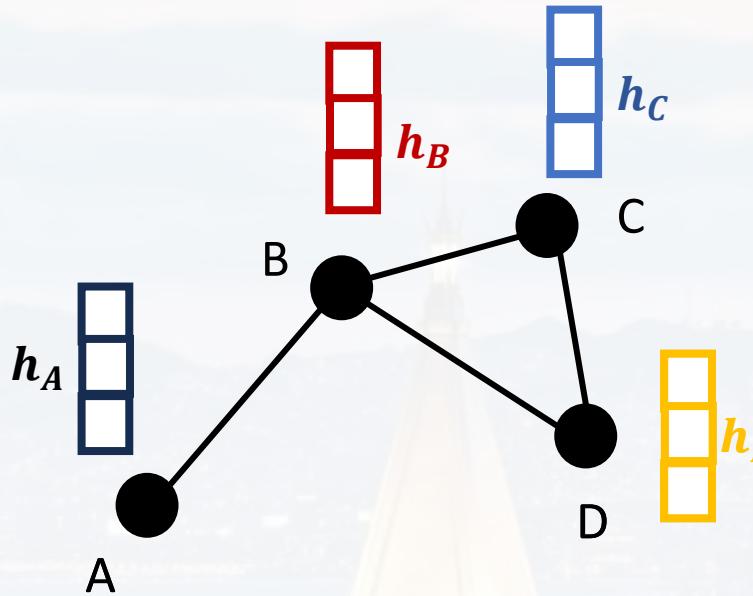


each node i has a **feature vector** h_i

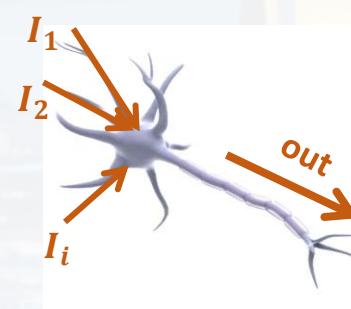
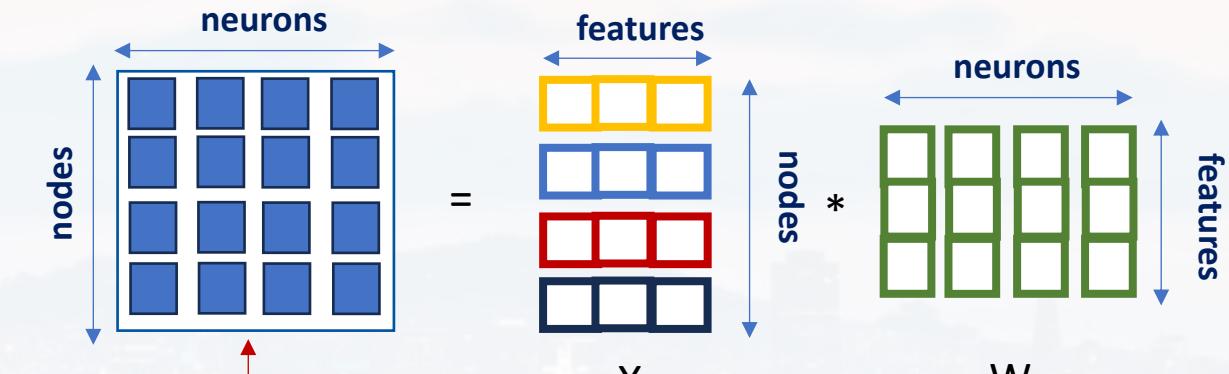




Graph Convolution



each node i has a **feature vector** h_i

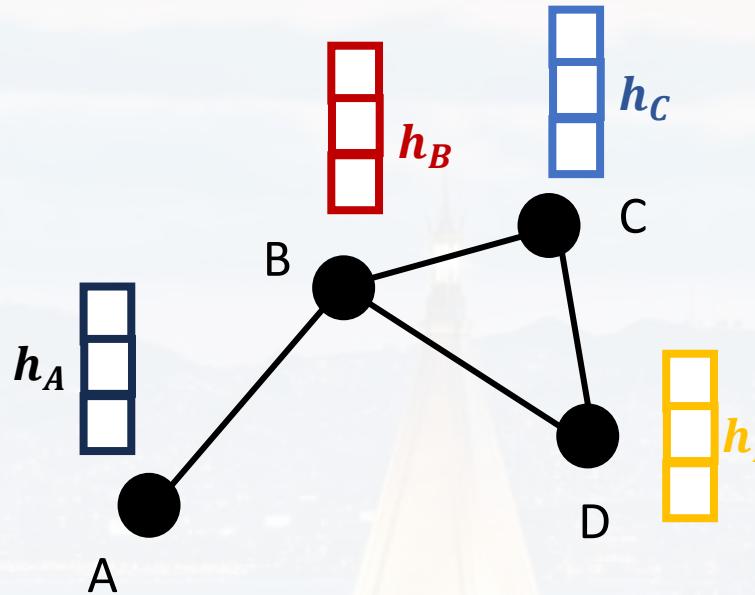


$$net = \sum_i I_i \cdot w_i + b$$

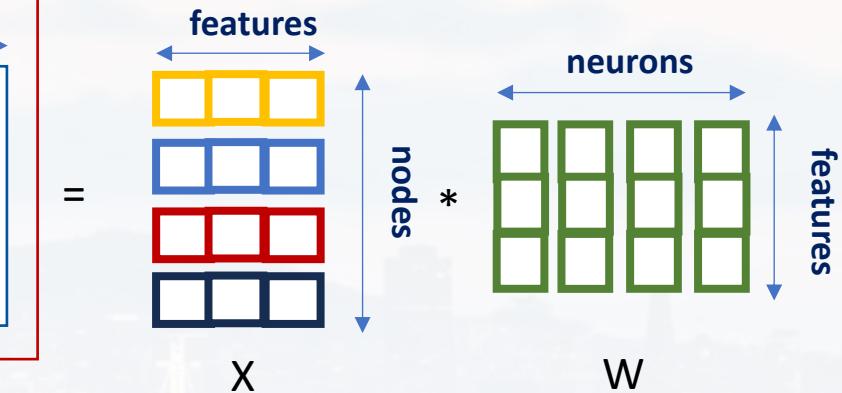
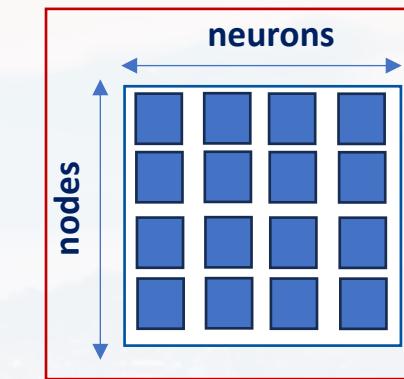
$$m_{jk} = \sum_i w_{ji} x_{ik}$$



Graph Convolution

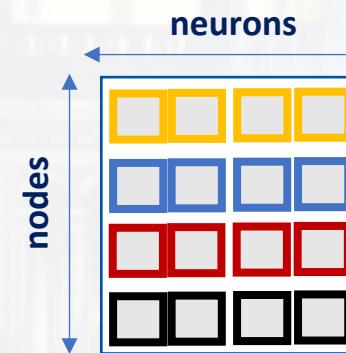


each node i has a **feature vector** h_i



$$m_{jk} = \sum_i w_{ji} x_{ik}$$

depending on W
the output features
may have different
lengths than the
input features

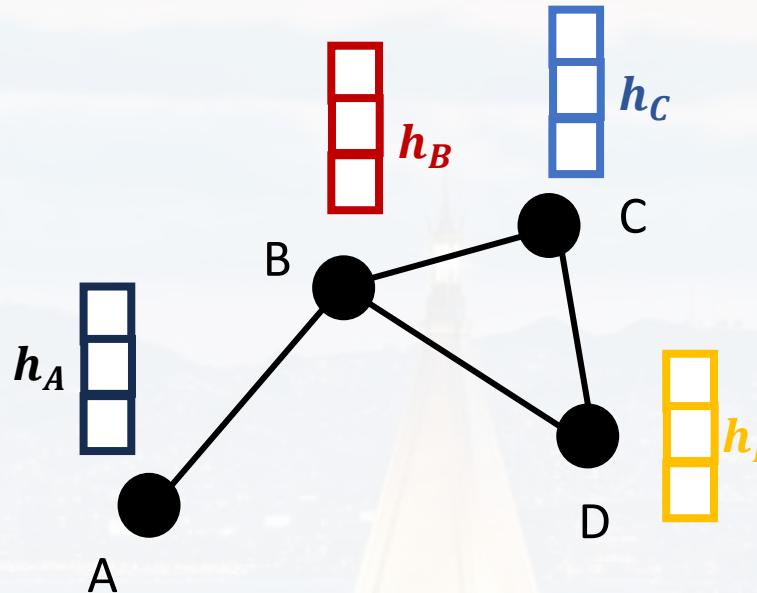


adjacency A

$$= \left[\begin{pmatrix} 0 & 1 & 0 & 0 \\ 1 & 0 & 1 & 1 \\ 0 & 1 & 0 & 1 \\ 0 & 1 & 1 & 0 \end{pmatrix} + \begin{pmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{pmatrix} \right] * M$$

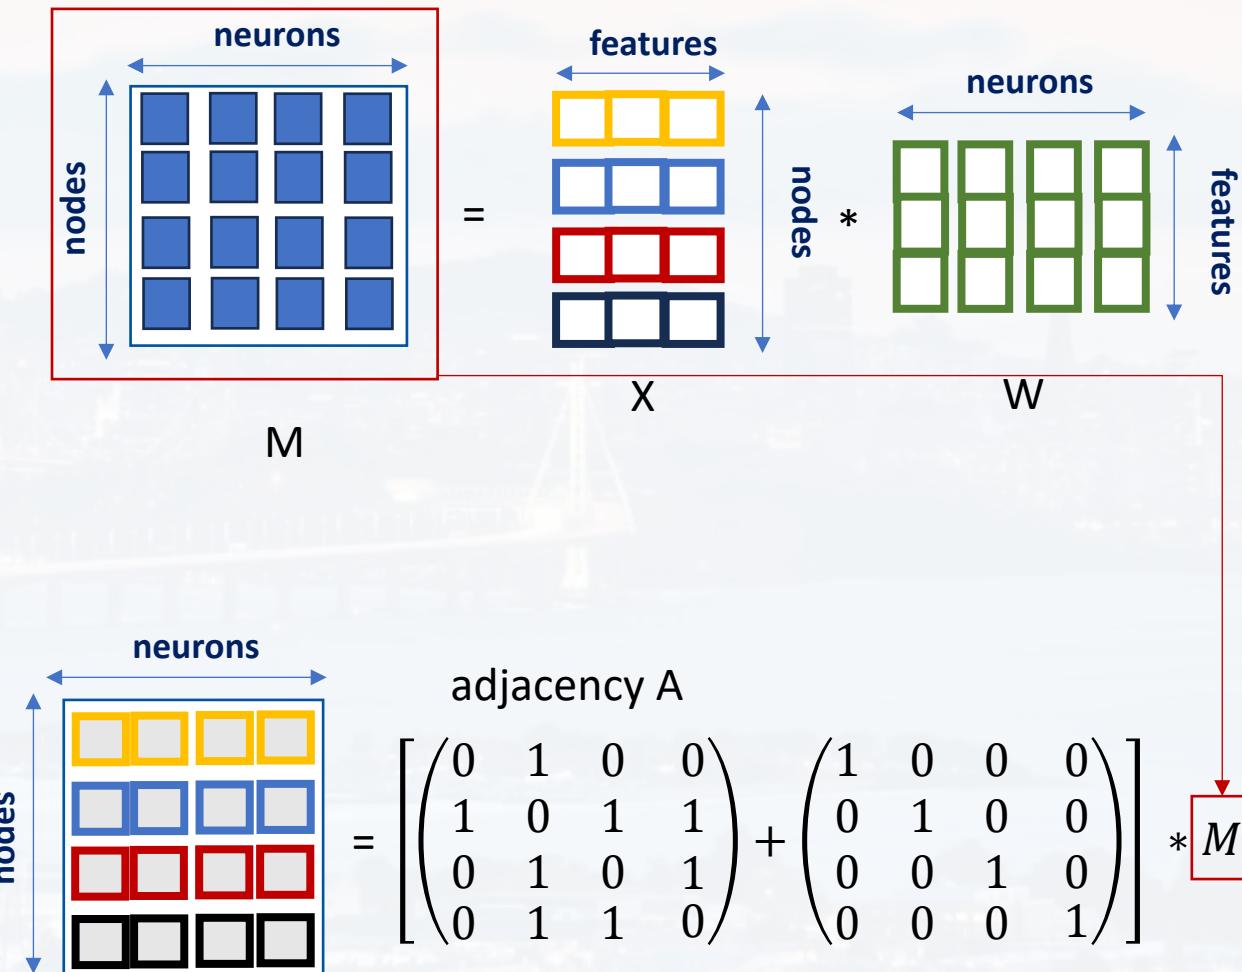


Graph Convolution



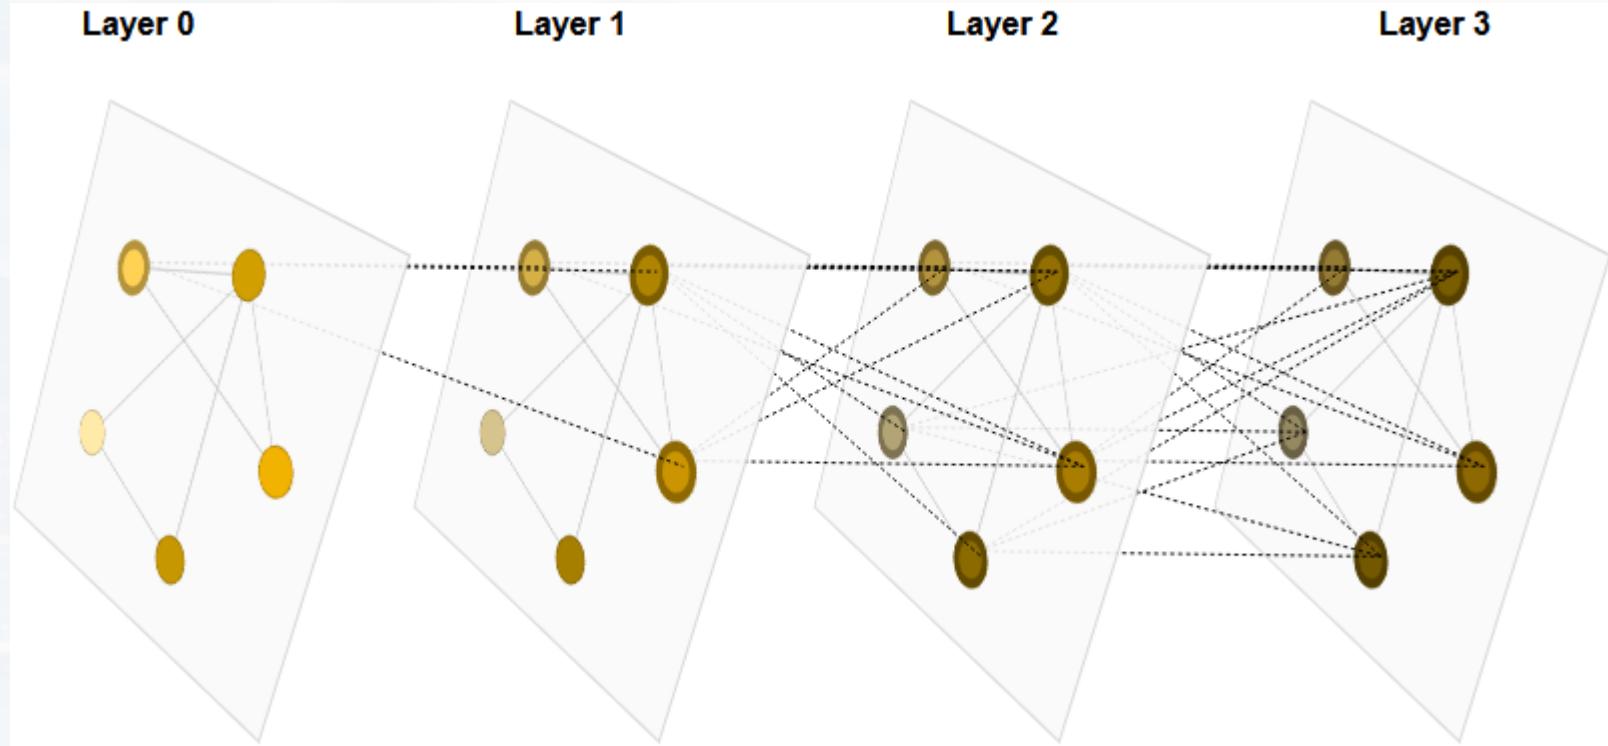
pass through a ReLU and/or
repeat the procedure with another W
(aka second convolution layer)

each node i has a **feature vector** h_i





Graph Convolution

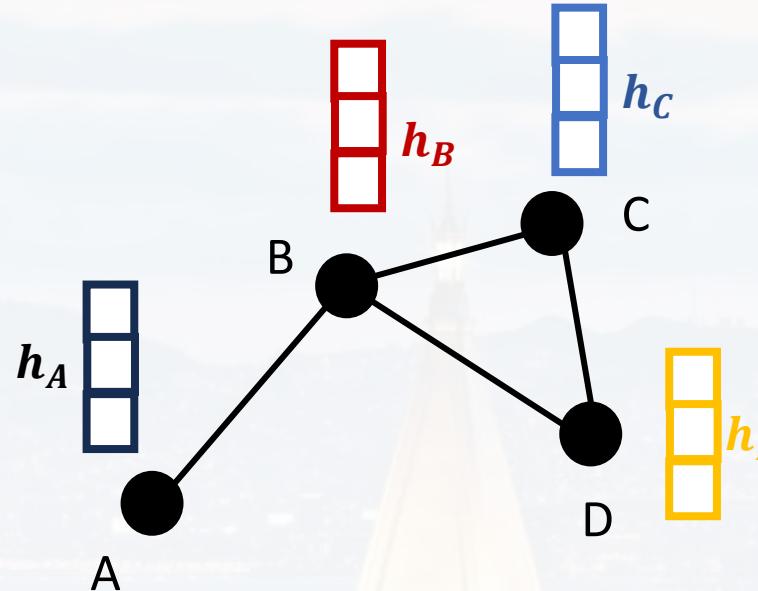


1st layer: one-hop neighborhood
2nd layer: two-hop neighborhood
etc

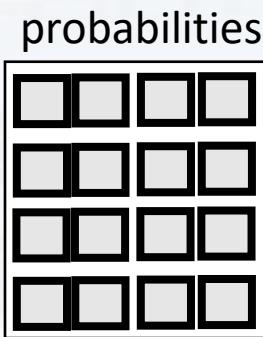
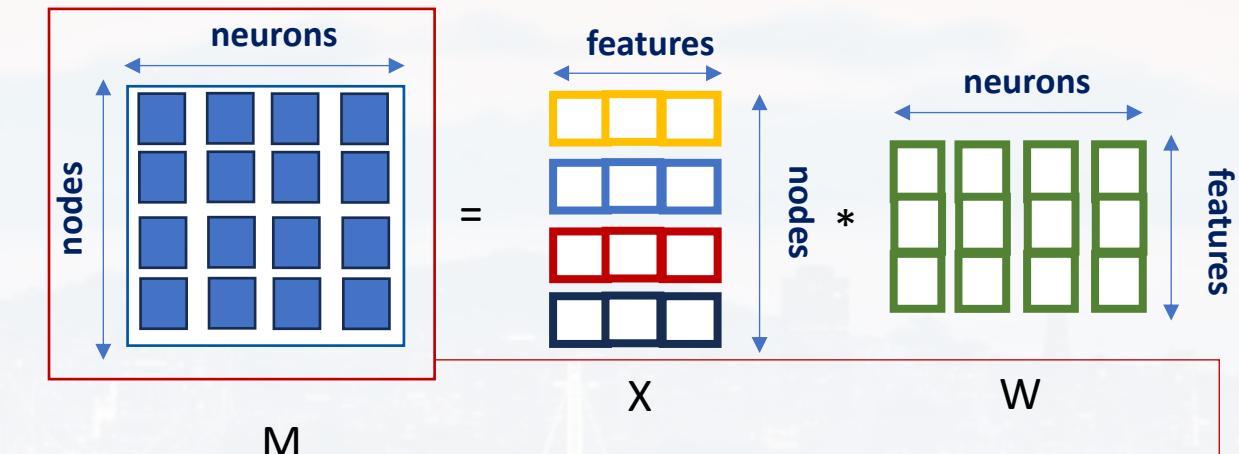
[animation here](#)



Graph Convolution



each node i has a **feature vector** h_i



node classification

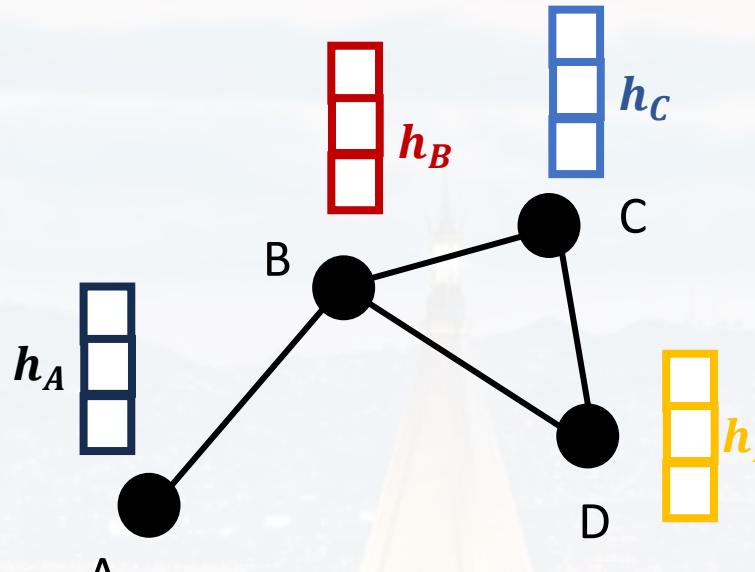
pass through a softmax layer

Diagram illustrating the computation of the adjacency matrix A . It shows the sum of two matrices: the identity matrix and a sparse adjacency matrix.

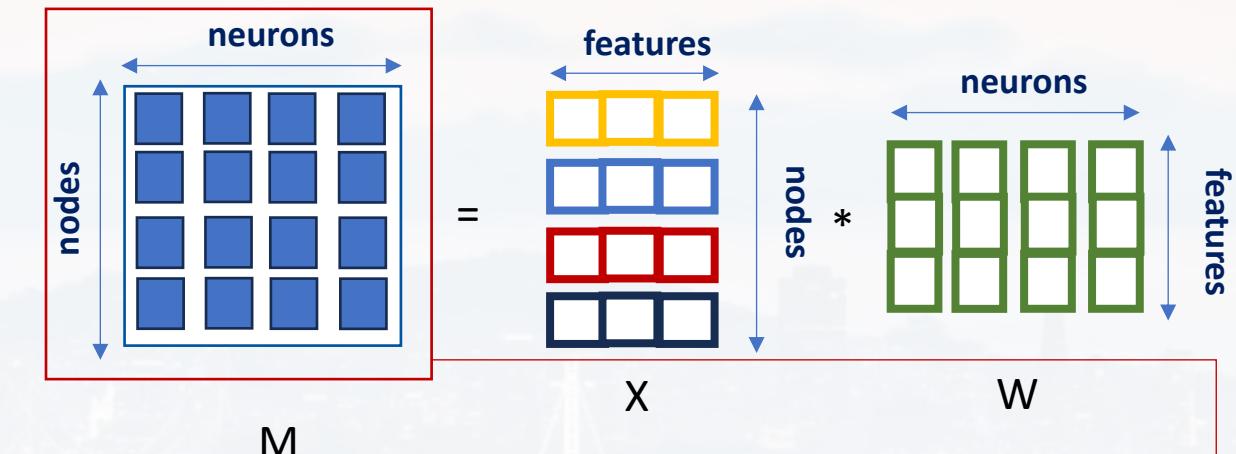
$$A = \left[\begin{pmatrix} 0 & 1 & 0 & 0 \\ 1 & 0 & 1 & 1 \\ 0 & 1 & 0 & 1 \\ 0 & 1 & 1 & 0 \end{pmatrix} + \begin{pmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{pmatrix} \right] * M$$



Graph Convolution

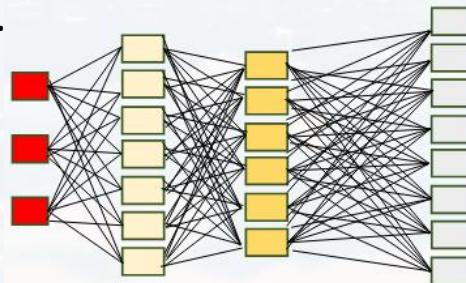


each node i has a **feature vector** h_i



node regression

e.g. 3D coordinates,
angles etc.



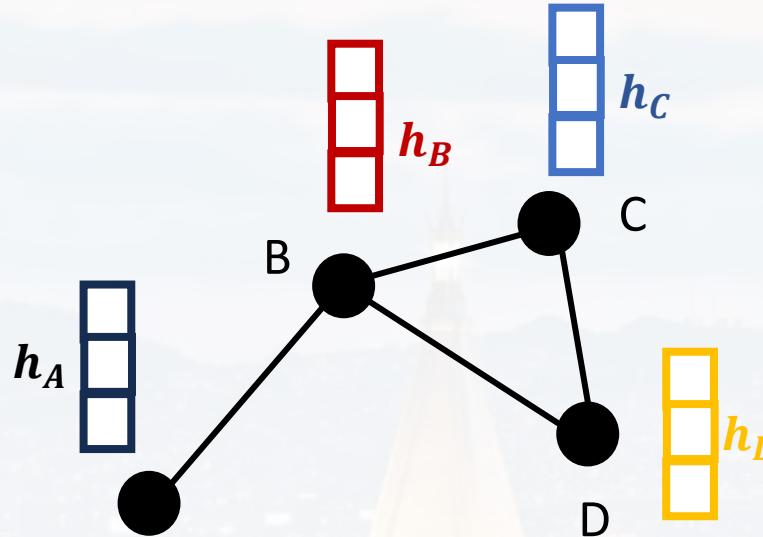
pass through a dense layer

Diagram illustrating the computation of the adjacency matrix A . It shows the sum of two matrices: $\begin{pmatrix} 0 & 1 & 0 & 0 \\ 1 & 0 & 1 & 1 \\ 0 & 1 & 0 & 1 \\ 0 & 1 & 1 & 0 \end{pmatrix}$ and $\begin{pmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{pmatrix}$, followed by multiplication with matrix M .

$$A = \left[\begin{pmatrix} 0 & 1 & 0 & 0 \\ 1 & 0 & 1 & 1 \\ 0 & 1 & 0 & 1 \\ 0 & 1 & 1 & 0 \end{pmatrix} + \begin{pmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{pmatrix} \right] * M$$



Graph Convolution



summary

A:	adjacency matrix (number of nodes x number of nodes)
I:	identity matrix (number of nodes x number of nodes)
X:	node feature matrix (number of nodes x number of features)
W:	weight matrix (number of features x number of neurons)
σ :	any activation function
$D^{-1/2}$:	diagonal matrix for normalization

$$H(\text{embedding}) = \sigma [D^{-1/2} (A + I) D^{-1/2} X W]$$

However, this would give nodes with higher degree a larger weight

→ normalizing by $\frac{1}{\sqrt{d(n_i)}}$ and $\frac{1}{\sqrt{d(n_j)}}$

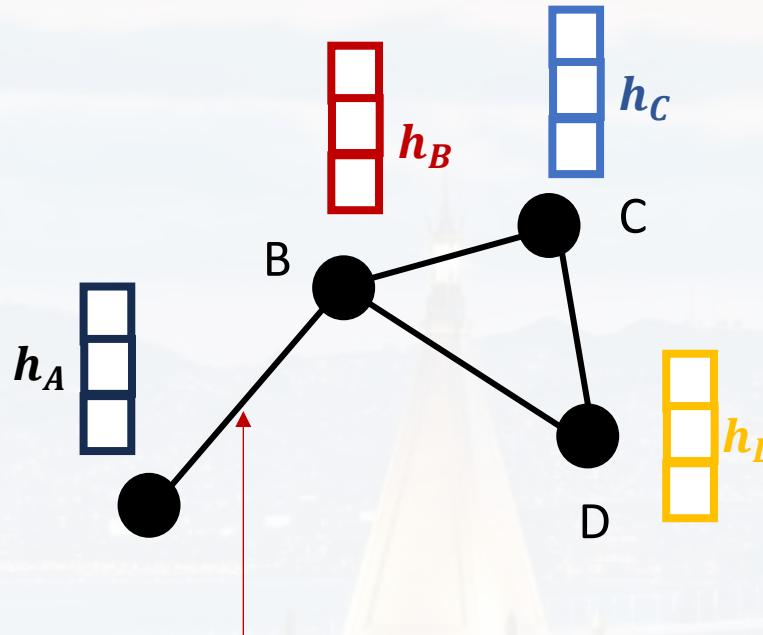
more information [here](#)



Matthew N. Bernstein



Graph Convolution



summary

- A: adjacency matrix (number of nodes x number of nodes)
I: identity matrix (number of nodes x number of nodes)
X: node feature matrix (number of nodes x number of features)
W: weight matrix (number of features x number of neurons)
 σ : any activation function
 $D^{-1/2}$: diagonal matrix for normalization

$$H(\text{embedding}) = \sigma [D^{-1/2} (A + I) D^{-1/2} X W]$$

h_i : embedding vector of node i

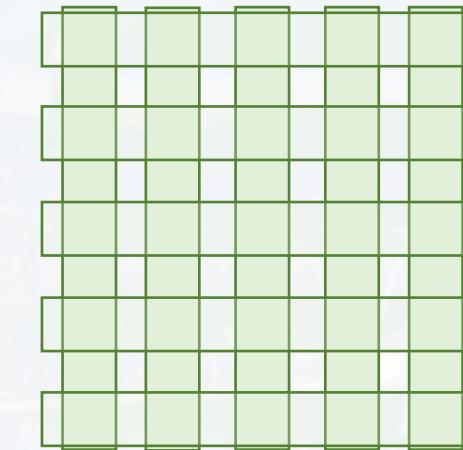
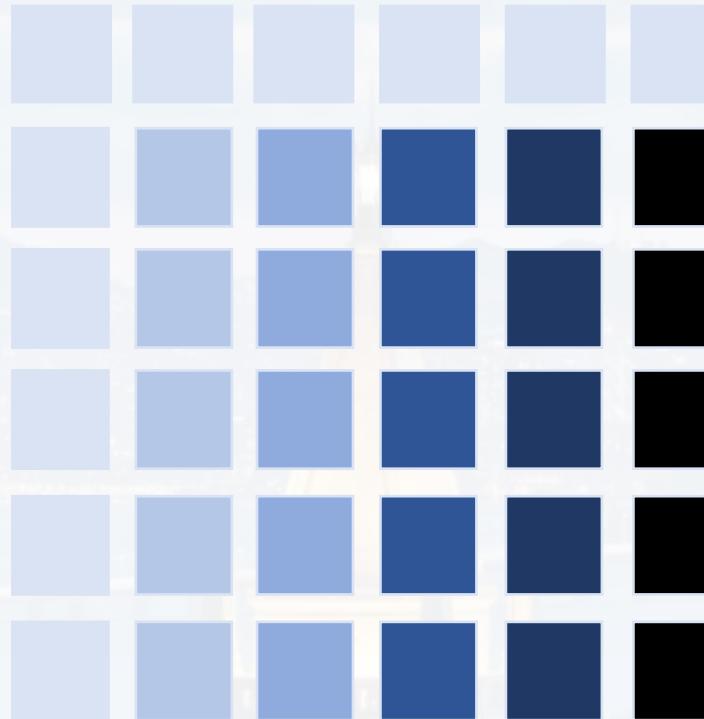
edge prediction (= classification), probability of an edge $p_{\text{edge}} = f(h_i, h_j)$

where f can be a **sigmoid** or **dense layer**



Attention

"The cat jumped on the roof."



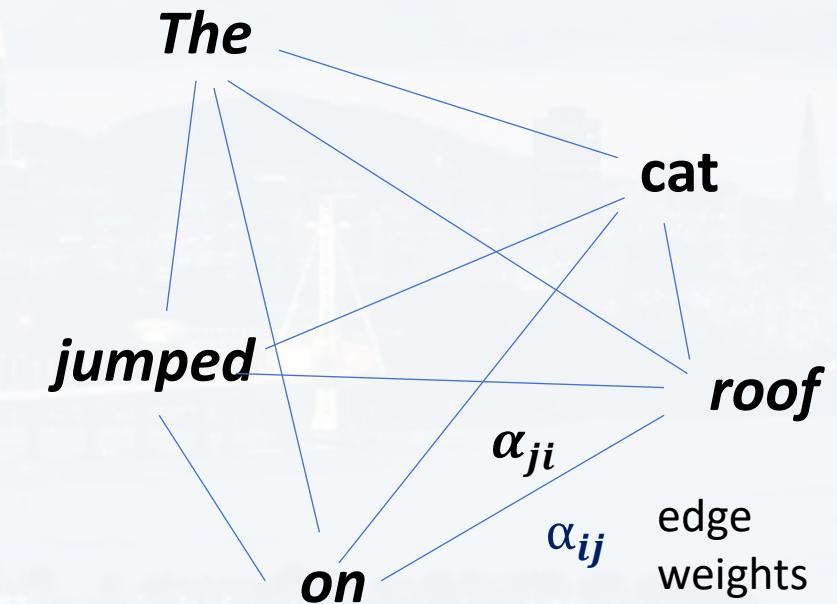
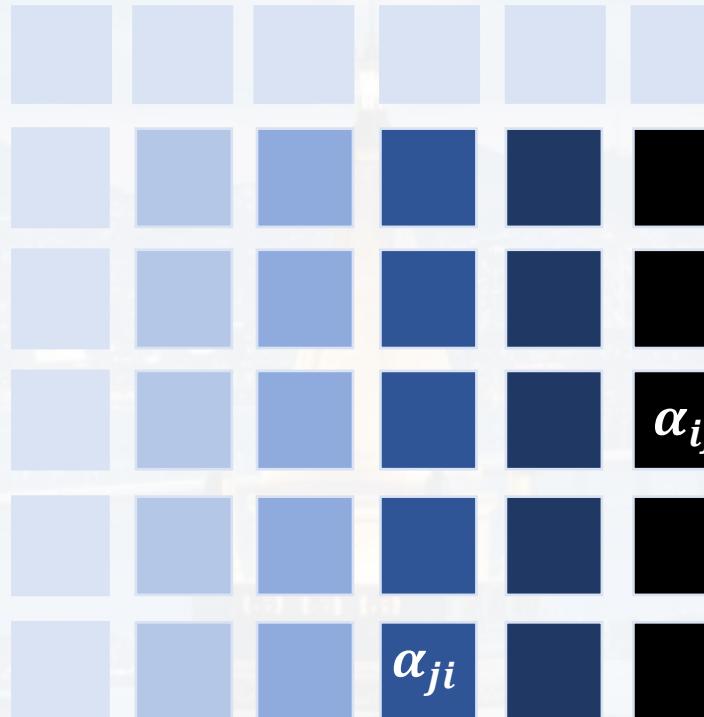
```
Gaussian kernel      W      = np.exp(-(D**2)/(sigma))  
                      W      = W/np.sum(W + 1e-16, axis = 0)  
                      yint   = np.dot(W.transpose(), y)
```

actual attention:
these weights are learnable,
no kernel assumed!



Graph Attention

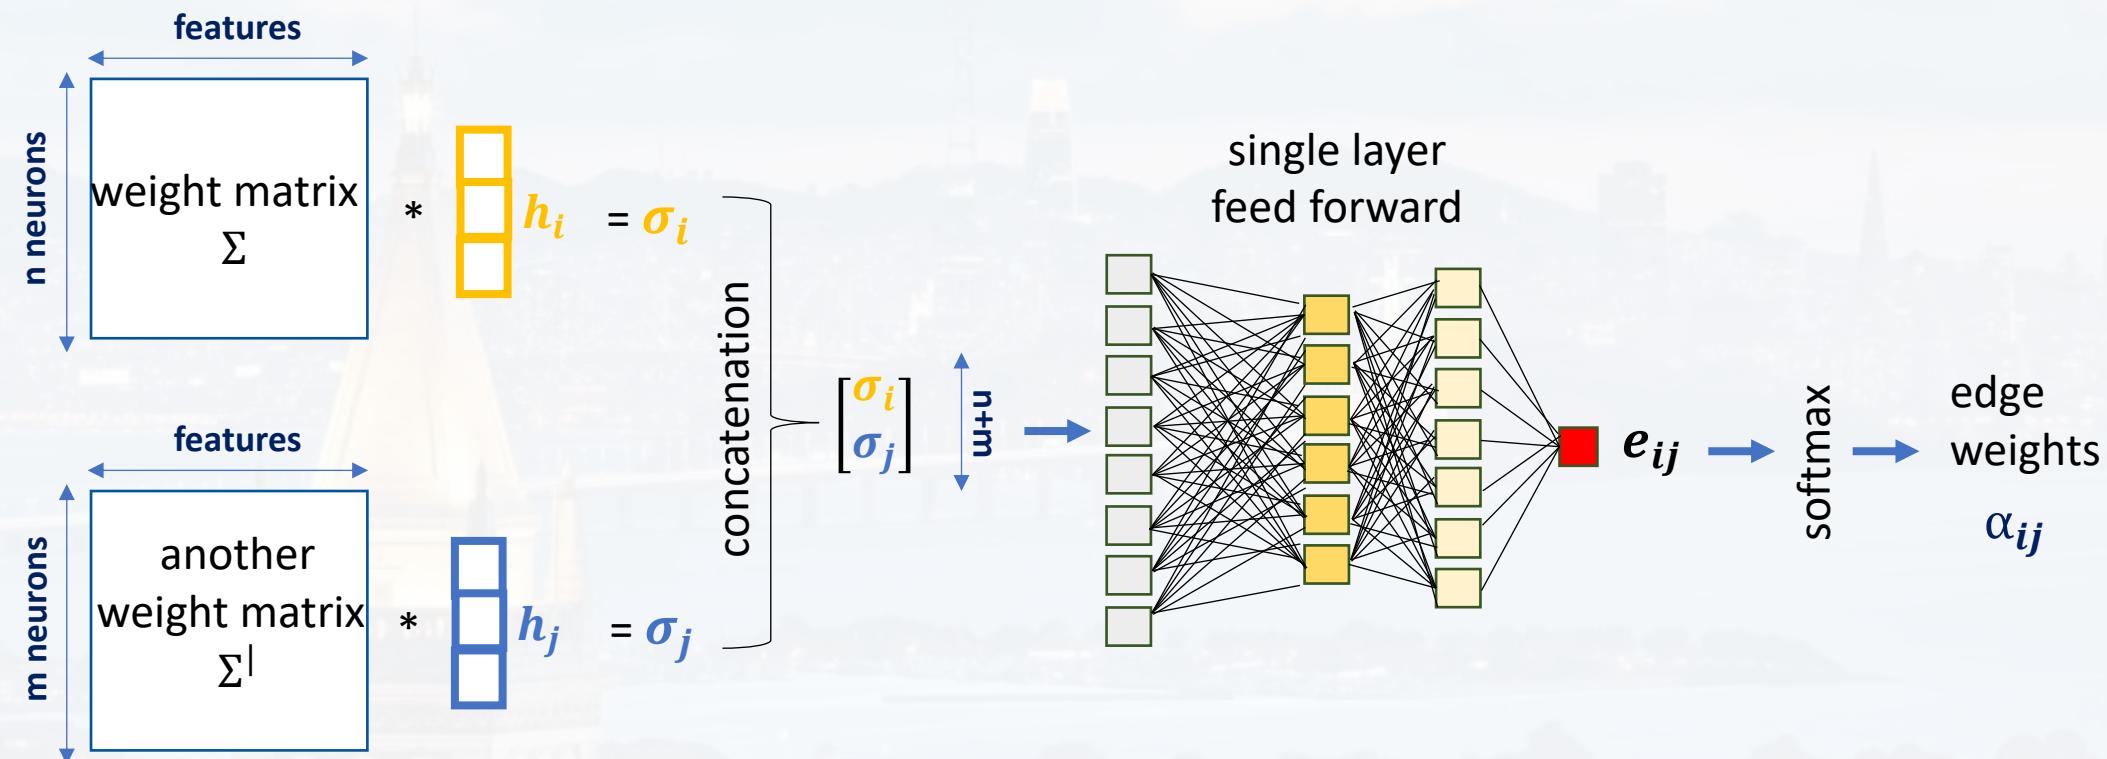
"The cat jumped on the roof."





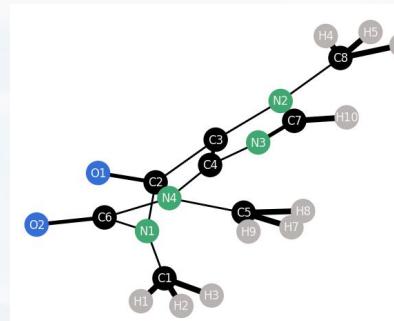
Graph Attention

Learning the weights!
(edge attributes)



Graph Attention

for graph classification/regression



convolution
GNN → node embeddings h_i

Each node contributes to the classification, but how?

→ **weighted** (= scalar) **average** of the embedding vectors h_i

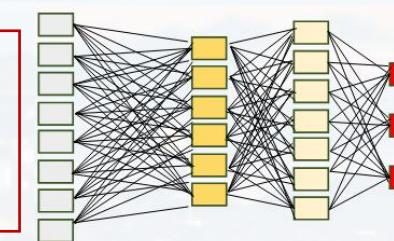
→ **number of nodes** can **differ** between different graphs!

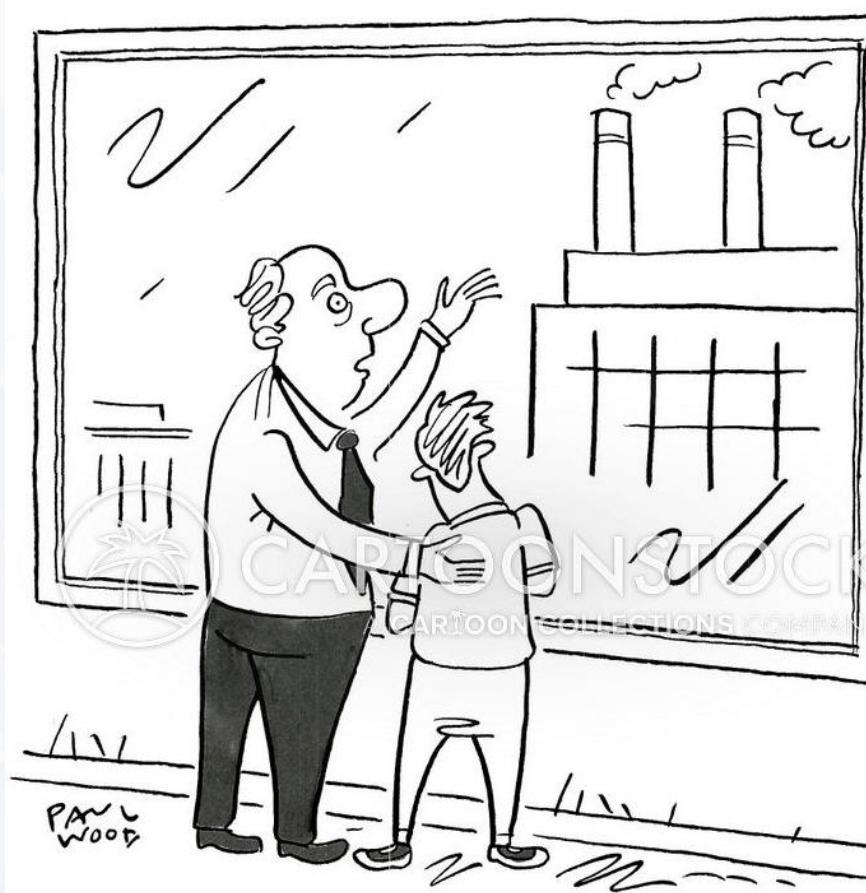
output

$$v = \sum_i \alpha_i h_i = \sum_i \text{softmax}[a^T \tanh(W h_i)] h_i$$

a^T and W are trainable

dense layer+ softmax for classification
or
dense layer for regression





ONE DAY SON, ALL THIS
WILL BE RUN BY ROBOTS

Outline

- What is a Graph
- The ANN Part
- PyTorch Example



node classification: convolution GNN

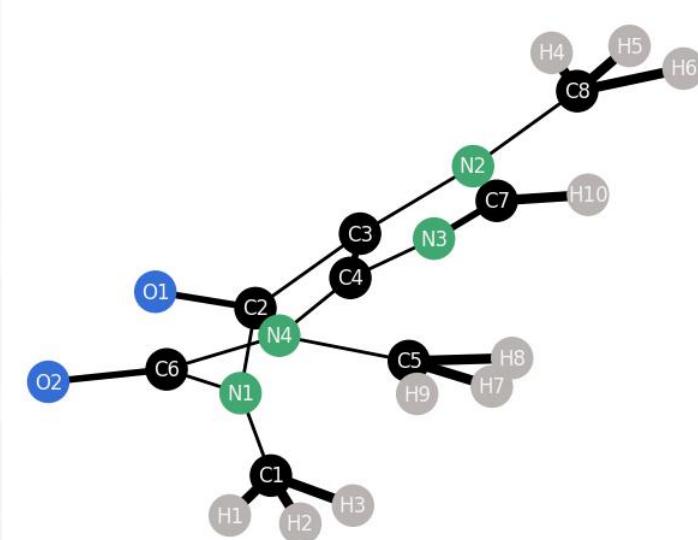
```
self.conv1 = GCNConv(n_node_features, n_neuron)  
self.conv2 = GCNConv(n_neuron, n_classes)
```

```
log_softmax(x3, dim = 1)
```

- edge weights: binding affinity

see Graph_III.ipynb

epoch	loss	accuracy
0	1.49	66.67%
10	1.94	66.67%
20	0.17	79.17%
30	0.13	79.17%
40	0.14	79.17%
50	0.11	79.17%
60	0.11	79.17%
70	0.11	79.17%
80	0.11	79.17%
90	0.11	79.17%
100	0.11	79.17%
110	0.11	79.17%
120	0.10	79.17%
130	0.10	79.17%
140	0.10	79.17%
150	0.10	79.17%
160	0.10	79.17%



```
print(Y)  
print(Y_pred)
```

```
[0. 0. 0. 0. 0. 0. 0. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 2. 2. 2. 2. 3. 3.]  
tensor([0, 0, 0, 0, 0, 0, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 2, 2, 0, 0, 0])
```



Thank you very much for your attention!

