

- **for next HW assignment:** variance, covariance, correlation
- **for next HW assignment:** Principal Component Analysis (PCA)
- regression table

- **for next HW assignment:** variance, covariance, correlation
- **for next HW assignment:** Principal Component Analysis (PCA)
- regression table



recap:

$$\mu = E(x) = \int x p(x) dx$$

$$\sigma^2 = \text{var}(x) = \int (x - \mu)^2 p(x) dx$$

$$\text{var}(x) = \int (x - \mu)^2 p(x) dx = E([x - \mu]^2)$$

variance can be interpreted as **mean of $[x - \mu]^2$**

$$= E(x^2 - 2x\mu + \mu^2)$$

$$= \int [x^2 - 2x\mu + \mu^2] p(x) dx$$

$$= \int x^2 p(x) dx - 2\mu \int x p(x) dx + \mu^2 \int p(x) dx$$

$$= E(x^2) - 2\mu E(x) + \mu^2 E(1)$$

$$\int p(x) dx = 1$$

$$= E(x^2) - 2\mu E(x) + \mu^2$$

$$\mu = E(x)$$

$$\sigma^2 = E(x^2) - E(x)^2$$



recap:

a, b = const

$$\sigma^2 = E(x^2) - E(x)^2$$

$$\mu = E(x) = \int x p(x) dx$$

$$\sigma^2 = \text{var}(x) = \int (x - \mu)^2 p(x) dx$$

$$\text{var}([a x_1 + b x_2]) = E([a x_1 + b x_2]^2) - E(a x_1 + b x_2)^2$$

$$= E(a^2 x_1^2 + 2ab x_1 x_2 + b^2 x_2^2) - E(a x_1 + b x_2)^2$$

$$= a^2 E(x_1^2) + 2ab E(x_1 x_2) + b^2 E(x_2^2) - E(a x_1 + b x_2)^2$$

$$= a^2 E(x_1^2) + 2ab E(x_1 x_2) + b^2 E(x_2^2) - [aE(x_1) + b E(x_2)]^2$$

$$= a^2 E(x_1^2) - a^2 E(x_1)^2 + b^2 E(x_2^2) - b^2 E(x_2)^2 + 2ab E(x_1 x_2) - 2ab E(x_1)E(x_2)$$

$$a^2 \text{var}(x_1)$$

$$b^2 \text{var}(x_2)$$

$$2ab \text{cov}(x_1, x_2)$$

$$= a^2 \text{var}(x_1) + b^2 \text{var}(x_2) + 2ab \text{cov}(x_1, x_2)$$

$$\text{cov}(x_1, x_2) = \text{cov}(x_2, x_1) = E(x_1 x_2) - E(x_1)E(x_2)$$

covariance



$$\sigma^2 = E(x^2) - E(x)^2$$

```
x1 = np.random.normal(0,2,(1000,))
```

```
x2 = np.random.normal(3,3,(1000,))
```

```
x3 = np.random.uniform(0,5,(1000,))
```

```
x4 = 5*np.random.uniform(3,4,(1000,))
```

```
x5 = np.sqrt(x4)
```

```
x6 = x1 + x2
```

```
x7 = 2*x3
```

```
x8 = x3*x2
```

```
All = np.vstack((x1, x2, x3, x4, x5, x6, x7, x8))
```

```
data = pd.DataFrame(All.transpose(),
```

```
                      columns = ['x1', 'x2', 'x3', 'x4', 'x5', 'x6', 'x7', 'x8'])
```

```
out = sns.pairplot(data, kind = "kde", \
```

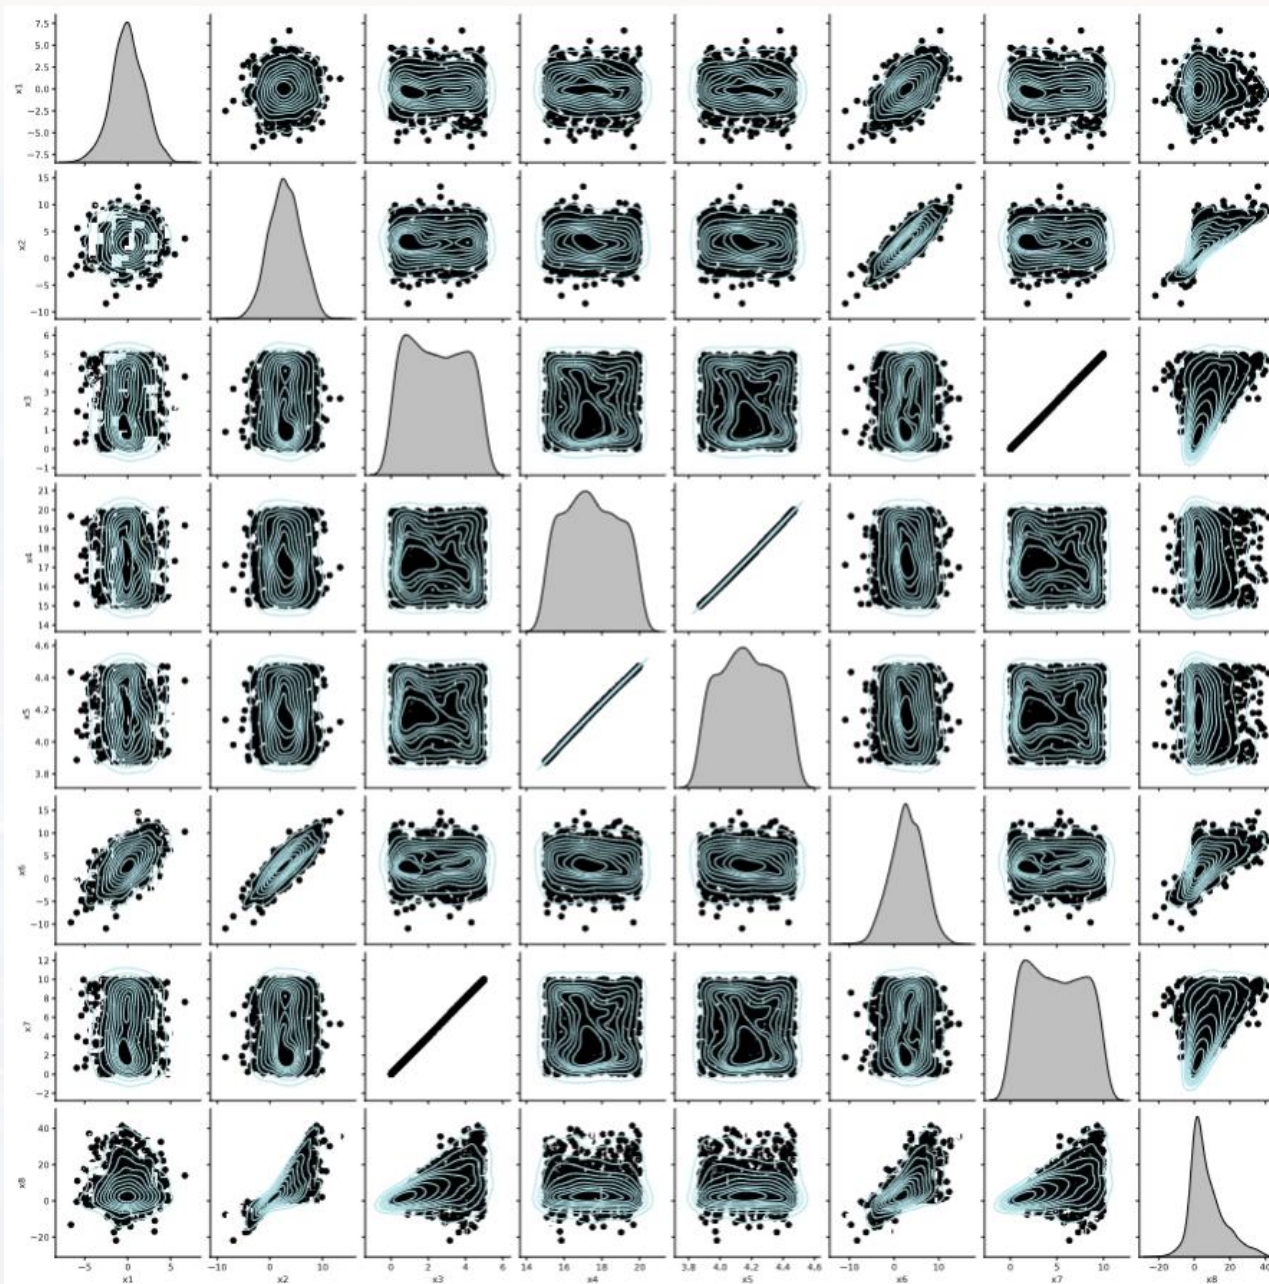
```
                      plot_kws = {'color':[176/255, 224/255, 230/255]}, \
```

```
                      diag_kws = {'color':'black'})
```

```
out.map_offdiag(plt.scatter, color = 'black')
```

$$\mu = E(x) = \int x p(x) dx$$

$$\sigma^2 = \text{var}(x) = \int (x - \mu)^2 p(x) dx$$



$$\mu = E(x) = \int x p(x) dx$$

$$\sigma^2 = \text{var}(x) = \int (x - \mu)^2 p(x) dx$$

based on the shape of the
data cloud

→ prediction how x_1 and x_2
are related, i. e.
how they **correlate**

→ how to quantify?



$$\text{cov}(x_1, x_2) = \text{cov}(x_2, x_1) = E(x_1 x_2) - E(x_1)E(x_2)$$

$$\mu = E(x) = \int x p(x) dx$$

$$\sigma^2 = \text{var}(x) = \int (x - \mu)^2 p(x) dx$$

a) x_1 and x_2 are independent

$$E(x_1 x_2) - E(x_1)E(x_2)$$

$$= \iint x_1 x_2 p(x_1) p(x_2) dx_1 dx_2 - \int x_1 p(x_1) dx_1 \int x_2 p(x_2) dx_2$$


x_1 and x_2 are independent:

x_1 is not a function of x_2 and vice versa

x_1 cannot be predicted by x_2 and vice versa

$$= \int x_1 p(x_1) dx_1 \int x_2 p(x_2) dx_2 - \int x_1 p(x_1) dx_1 \int x_2 p(x_2) dx_2 = 0$$

covariance equals **zero**
if samples are **independent!**





$$\text{cov}(x_1, x_2) = \text{cov}(x_2, x_1) = E(x_1 x_2) - E(x_1)E(x_2)$$

$$\mu = E(x) = \int x p(x) dx$$

$$\sigma^2 = \text{var}(x) = \int (x - \mu)^2 p(x) dx$$

b) x_1 and x_2 are **not** independent

$$E(x_1 x_2) - E(x_1)E(x_2)$$

$$= \iint x_1 x_2 p(x_1) p(x_2) dx_1 dx_2 - \int x_1 p(x_1) dx_1 \int x_2 p(x_2) dx_2$$


x_1 and x_2 are **not** independent:

x_1 **is** a function of x_2 and vice versa

x_1 **can** be predicted by x_2 to certain degree and vice versa

$$= \iint x_1 p(x_1) x_2(x_1) p(x_2(x_1)) dx_1 dx_2(x_1) - \int x_1 p(x_1) dx_1 \int x_2 p(x_2) dx_2$$

covariance **does not**
equal **zero**!







$$\text{cov}(x_1, x_2) = \text{cov}(x_2, x_1) = E(x_1 x_2) - E(x_1)E(x_2)$$

covariance

```
x1 = np.random.normal(0, 2, (1000,))  
x2 = 2*x1 + np.random.normal(0, 2, (1000,))
```



```
x1 = np.random.normal(0, 2, (1000,))  
x2 = 2*x1 + np.random.normal(0, 20, (1000,))
```


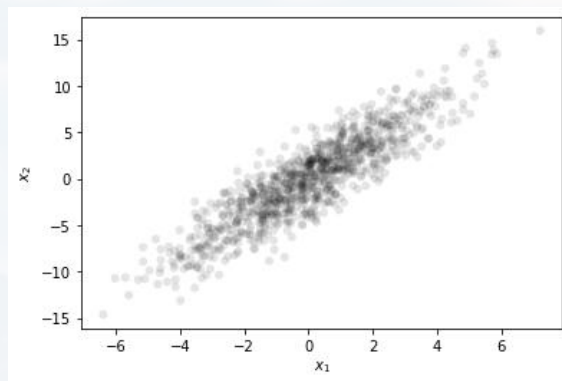


Same dependency, but different variance!



$$\text{cov}(x_1, x_2) = \text{cov}(x_2, x_1) = E(x_1 x_2) - E(x_1)E(x_2)$$

covariance



Same dependency, but different variance!

Need to scale for the variance!

Pearson's correlation
coefficient

$$\rho(x_1, x_2) = \frac{\text{cov}(x_1, x_2)}{\sqrt{\sigma_1^2 \sigma_2^2}}$$

$\rho(x_1, x_2)$:

- ranges from -1 to +1
- zero: no correlation
(completely independent)
- -1: max anti correlation
- +1: max correlation



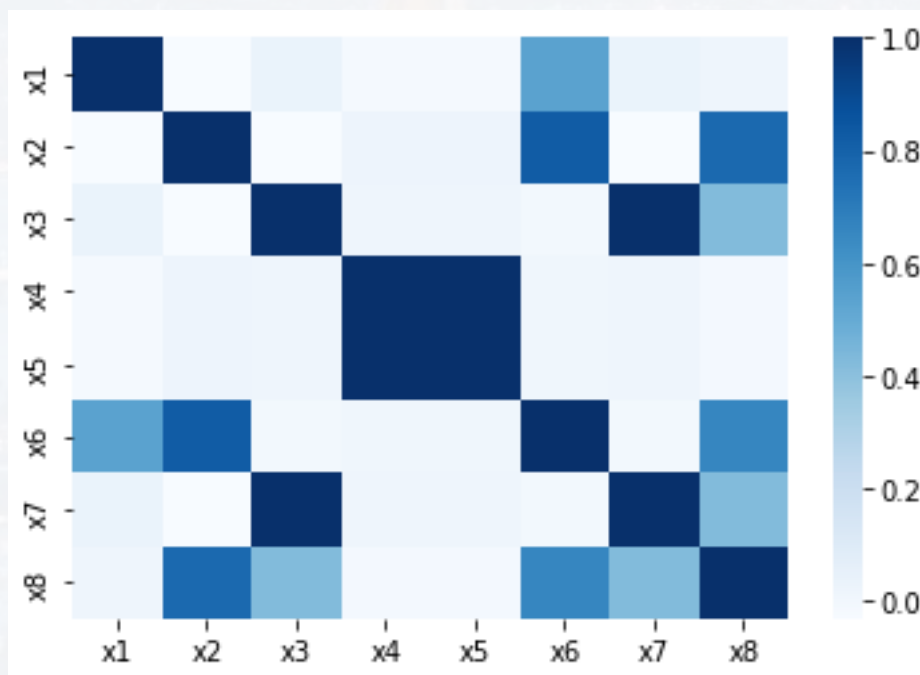
$$\text{cov}(x_1, x_2) = \text{cov}(x_2, x_1) = E(x_1 x_2) - E(x_1)E(x_2)$$

covariance

$$\rho(x_1, x_2) = \frac{\text{cov}(x_1, x_2)}{\sqrt{\sigma_1^2 \sigma_2^2}}$$

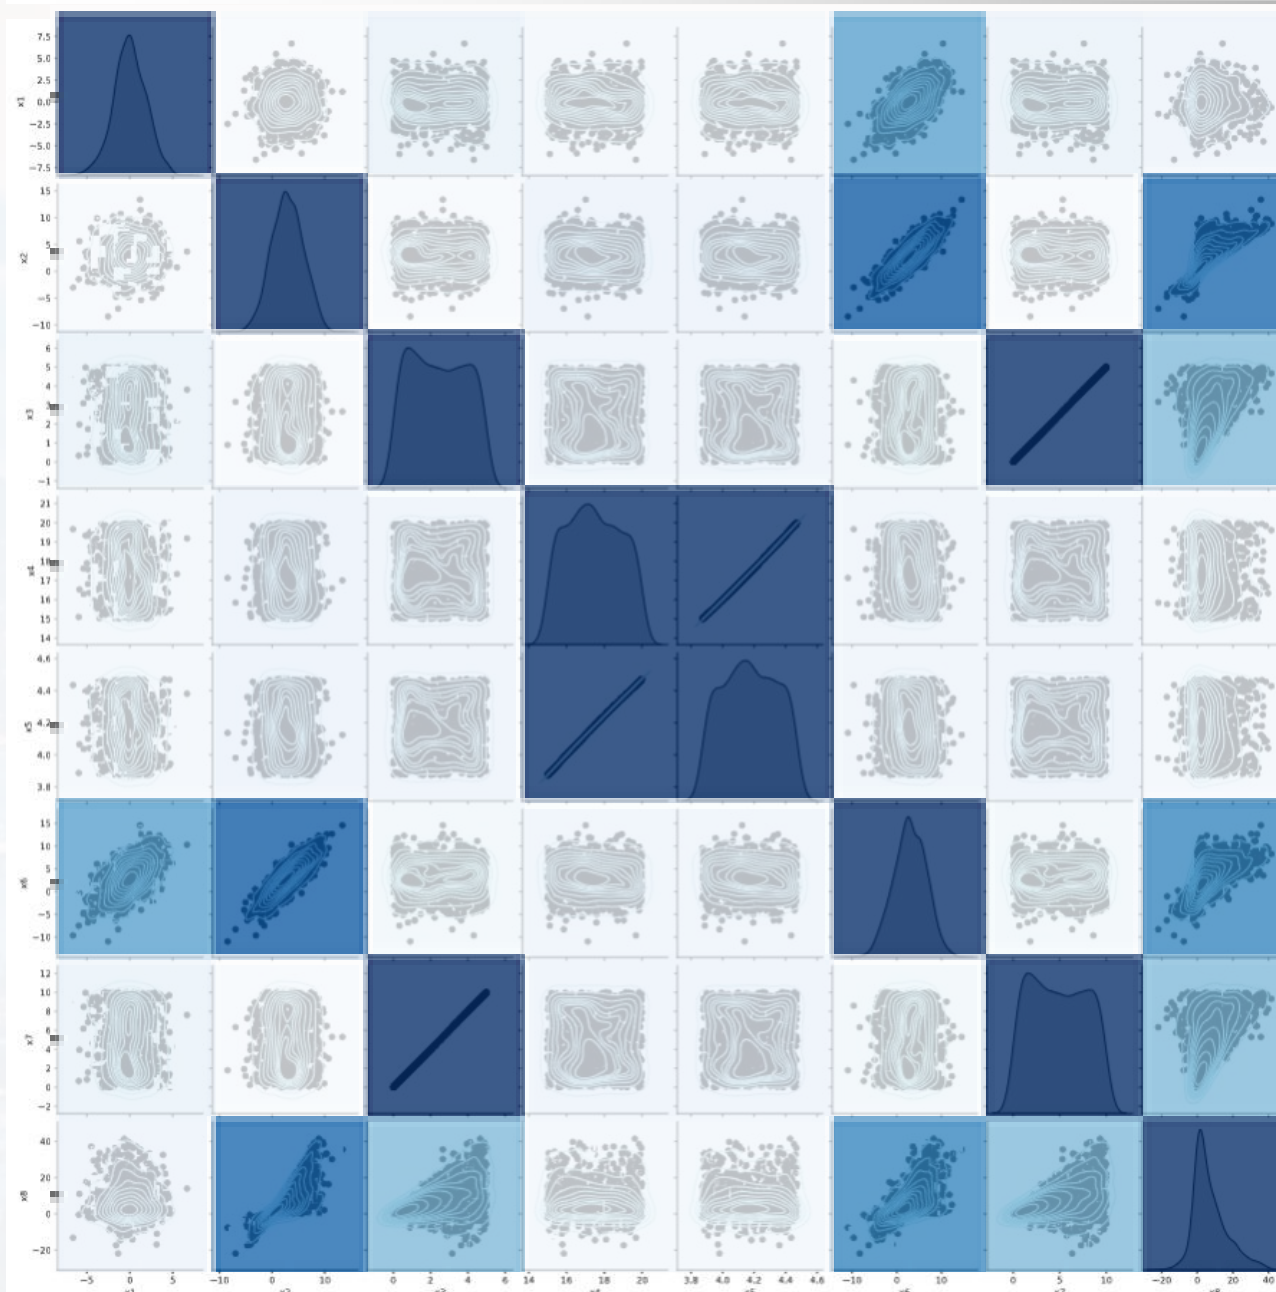
Pearson's correlation
coefficient

```
sns.heatmap(data.corr(), cmap = "Blues")
```



$\rho(x_1, x_2)$:

- ranges from -1 to +1
- zero: no correlation
(completely independent)
- -1: max anti correlation
- +1: max correlation



$\rho(x_1, x_2)$:

- ranges from -1 to +1
- zero: no correlation
(completely independent)
- -1: max anti correlation
- +1: max correlation



Important quantities you should know:

mean

$$\mu = E(x) = \int x p(x) dx$$

median m

$$\int_a^m p(x) dx = \frac{1}{2}$$

variance

$$\sigma^2 = \text{var}(x) = \int (x - \mu)^2 p(x) dx$$

$$\sigma^2 = E(x^2) - E(x)^2$$

$$\sigma_{tot}^2 = \sigma_1^2 + \sigma_2^2 + 2 \text{cov}(x_1, x_2)$$

covariance

$$\text{cov}(x_1, x_2) = E(x_1 x_2) - E(x_1)E(x_2)$$

correlation
coefficient

$$\rho(x_1, x_2) = \frac{\text{cov}(x_1, x_2)}{\sqrt{\sigma_1^2 \sigma_2^2}}$$

note:

$$\int (x - \mu)^n p(x) dx$$

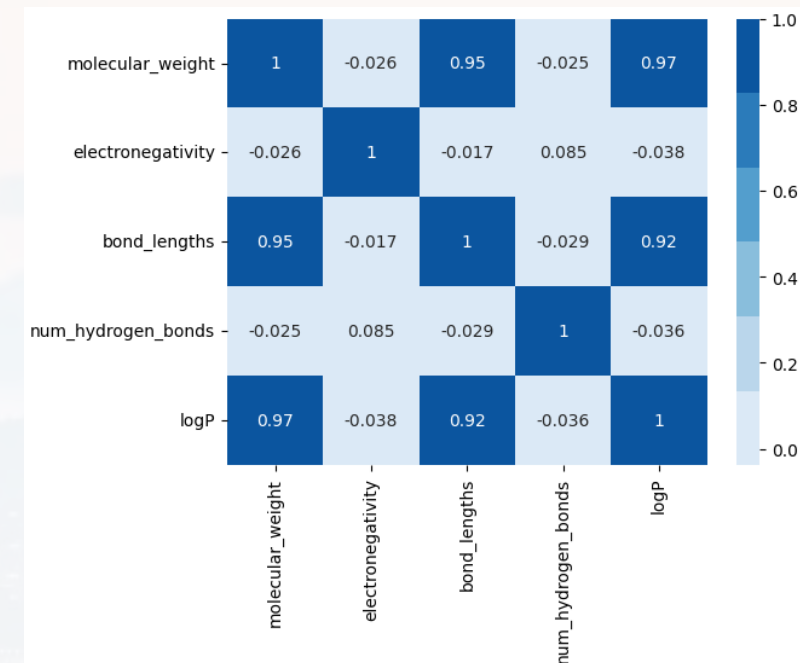
called n -th *moment*
of a pdf

- for next HW assignment: variance, covariance, correlation
- **for next HW assignment: Principal Component Analysis (PCA)**
- regression table




correlation means:

- features are **not mutually independent**
 - we can predict feature ***a*** from feature ***b*** to some extent
 - we don't need all features
- **reducing number of features** (dimensions) without losing information



$$\Sigma = \begin{pmatrix} \sigma_{11}^2 & \sigma_{12}^2 & \dots & \sigma_{1N}^2 \\ \dots & \sigma_{ii}^2 & \dots & \sigma_{iN}^2 \\ \sigma_{N1}^2 & \sigma_{Ni}^2 & \dots & \sigma_{NN}^2 \end{pmatrix} \quad \text{covariance matrix}$$

can be interpreted as a quadratic form




Sivia, “Data Analysis”
section 3.2



$$\Sigma = \begin{pmatrix} \sigma_{11}^2 & \sigma_{12}^2 & \dots & \sigma_{1N}^2 \\ \dots & \sigma_{ii}^2 & \dots & \sigma_{iN}^2 \\ \sigma_{N1}^2 & \sigma_{Ni}^2 & \dots & \sigma_{NN}^2 \end{pmatrix}$$

covariance matrix

can be interpreted as a quadratic form



Sivia, “Data Analysis”
section 3.2

idea: diagonalize Σ

$$\begin{pmatrix} \lambda_1 & 0 & \dots & 0 \\ 0 & \lambda_i & & 0 \\ 0 & 0 & & \lambda_N \end{pmatrix}$$

the diagonal are the **eigenvalues** (= variances in
new coordinate system)

principal components of the new **covariance matrix** are **parallel** to the **new coordinate axes** (= **eigenvectors**)





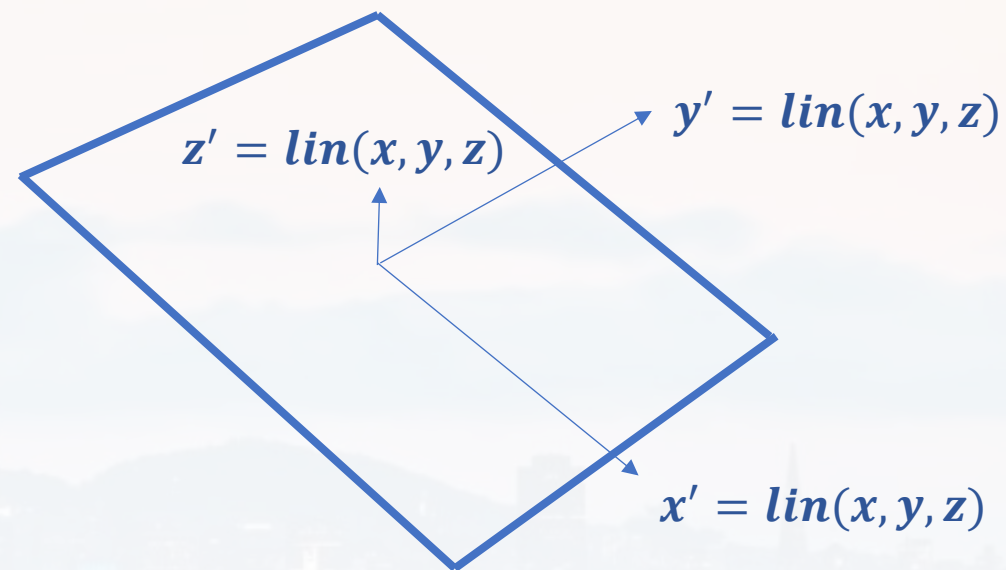
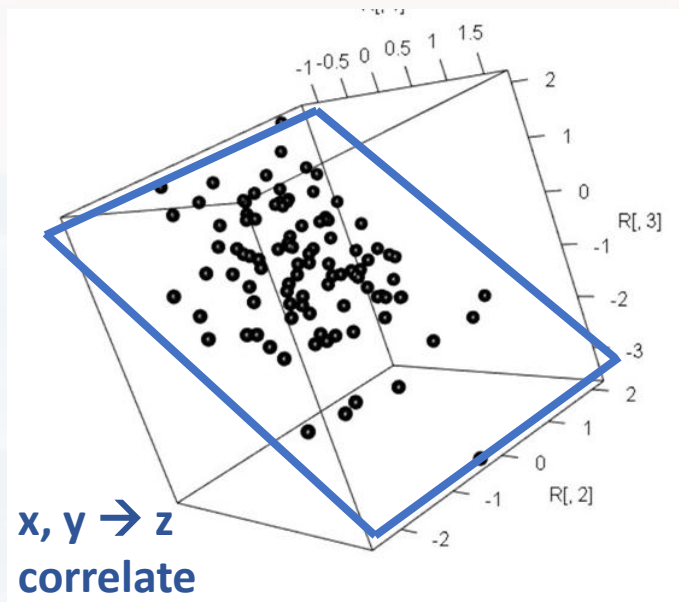
idea: diagonalize Σ

$$\begin{pmatrix} \lambda_1 & 0 & \dots & 0 \\ 0 & \lambda_i & & 0 \\ 0 & 0 & & \lambda_N \end{pmatrix}$$

the diagonal are the **eigenvalues** (= variances in new coordinate system)

principal components of the new **covariance matrix** are **parallel** to the **new coordinate axes** (= eigenvectors)

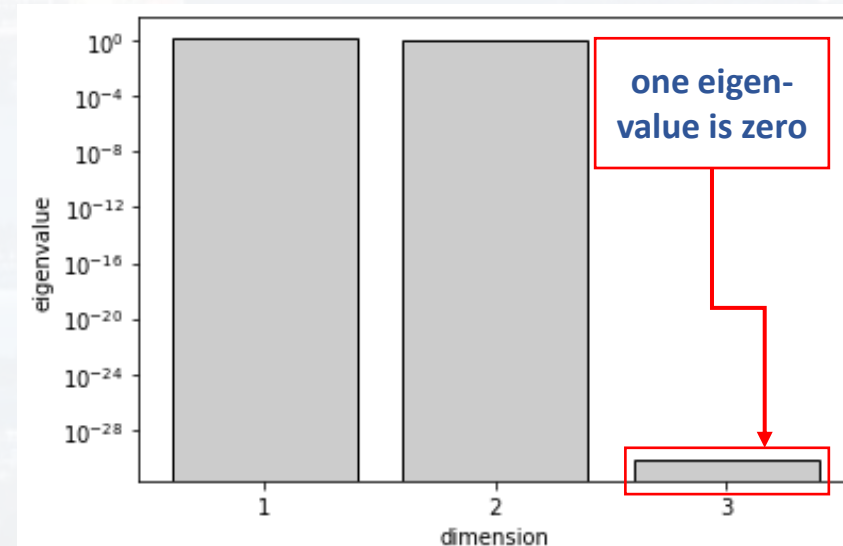




```
from sklearn.decomposition import PCA
```

```
out = PCA(n_components = 3).fit(XYZ)
```

```
eigenVec = out.components_  
eigenVal = out.explained_variance_  
eigenXYZ = out.transform(XYZ)
```



- **for next HW assignment:** variance, covariance, correlation
- **for next HW assignment:** Principal Component Analysis (PCA)
- regression table



see also discussion this week

Index	molecular_weight	electronegativity	bond_lengths	num_hydrogen_bonds	logP	toxicity_score
0	341.704	2.65585	3.09407	2	9.11147	80.9281
1	335.951	3.22262	2.89039	7	8.92848	83.4911
2	235.203	2.44115	2.48203	1	6.49731	61.8406
3	246.505	2.76656	2.71547	7	7.45089	57.0538
4	437.939	3.4801	3.59569	3	10.9156	131.326

$$y_k = \beta_0 + \sum_{n=1}^N \beta_n x_n + \epsilon$$



more accurate: determining **the p-values for the factors using ANOVA** for the corresponding residuals

```
table = sm.stats.anova_lm(my_model, typ = 1)
print(table)
```


$$y_k = \beta_0 + \sum_{n=1}^N \beta_n x_n + \epsilon$$

	df	sum_sq	mean_sq	F	PR(>F)	vs from t-test
molecular_weight	1.0	13.346285	13.346285	2847.525516	8.024085e-265	0.0000
electronegativity	1.0	0.640388	0.640388	136.631363	3.085962e-29	0.0000
bond_lengths	1.0	0.000684	0.000684	0.145954	7.025342e-01	0.6766
num_hydrogen_bonds	1.0	0.000703	0.000703	0.150055	6.985866e-01	0.6473
logP	1.0	0.013917	0.013917	2.969353	8.524510e-02	0.0852
Residual	794.0	3.721459	0.004687	NaN	NaN	



```
residuals = my_model.resid
```

```
plt.hist(residuals, color = 'w', edgecolor = 'black')  
plt.title('fit residuals')  
plt.ylabel('#')  
plt.xlabel('value')  
plt.show()
```



residuals approx.
normally distributed
around $\mu = 0$




```
residuals = my_model.resid
```

```
plt.hist(residuals, color = 'w', edgecolor = 'black')  
plt.title('fit residuals')  
plt.ylabel('#')  
plt.xlabel('value')  
plt.show()
```

residuals approx.
normally distributed
around $\mu = 0$

```
stats.probplot(residuals, dist = "norm", plot = pylab)  
pylab.show()
```

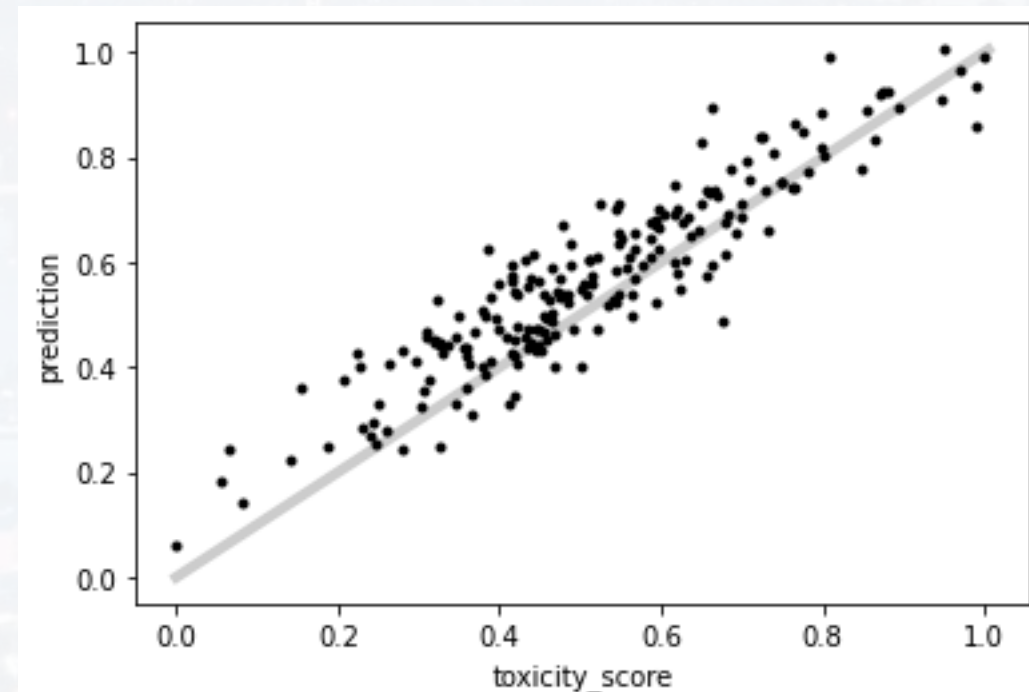




```
Ypred = my_model.predict(TestS)

higher = np.max([Ypred, TestS.toxicity_score])
lower = np.min([Ypred, TestS.toxicity_score])

plt.plot([lower, higher], [lower, higher], c = [0, 0, 0, 0.2], \
         linewidth = 4)
plt.scatter(TestS.toxicity_score, Ypred, marker = '.', c = 'k')
plt.ylabel('prediction')
plt.xlabel('toxicity score')
plt.show()
```





```
Ypred = my_model.predict(TestS)
```

```
higher = np.max([Ypred, TestS.toxicity_score])
```

```
lower = np.min([Ypred, TestS.toxicity_score])
```

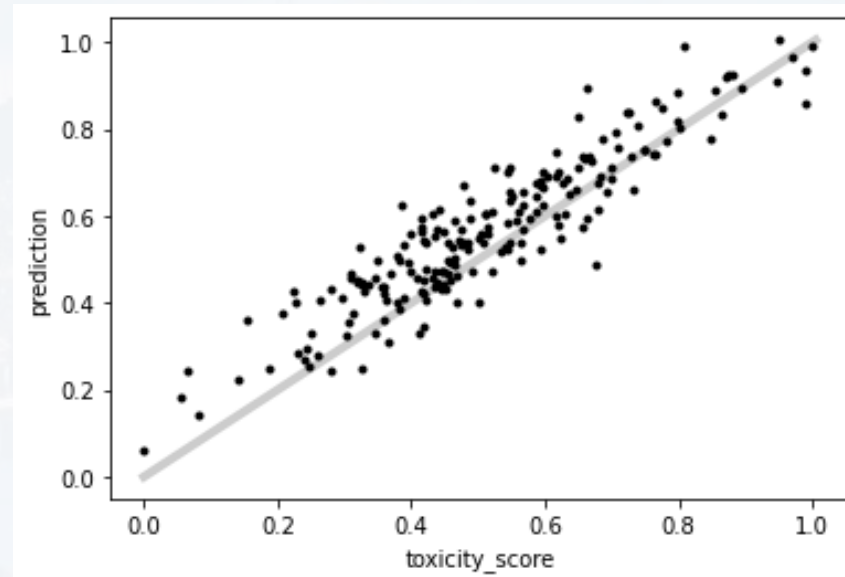
```
plt.plot([lower, higher], [lower, higher], c = [0, 0, 0, 0.2], linewidth = 4)
```

```
plt.scatter(TestS.toxicity_score, Ypred, marker = '.', c = 'k')
```

```
plt.ylabel('prediction')
```

```
plt.xlabel('toxicity score')
```

```
plt.show()
```



```
mean_dev = np.sum( abs(TestS.toxicity_score - Ypred) )/len(Ypred)
```

```
print(mean_dev)
```

5%

- **for next HW assignment:** variance, covariance, correlation
- **for next HW assignment:** Principal Component Analysis (PCA)
- regression table